


AI Adoption Starts With the Constraint, Not the Tool

A portfolio standard for turning AI interest into better recurring work



Marco Policani

Enterprise Portfolio · PMO · AI Operating Governance

policani.net · linkedin.com/in/marcpolicani

July 2026

EXECUTIVE SUMMARY

Fund the smallest change that moves a named constraint, then expand only when the evidence holds.

A working AI demo can still solve the wrong problem. Start with recurring work and name its constraint, the thing that holds it back today. Then compare AI with simpler changes before choosing a tool.

- 01** Name the recurring task, its current limit, the baseline, the owner, and the evidence that would show improvement.
- 02** Choose the least complex change that can solve the problem with enough safety, speed, reliability, and value.
- 03** Take each request through five gates: name the constraint, choose the change, design a small proof, decide whether to expand, change, stop, or try another option, and keep checking value once the change is part of normal work.

CONTENTS

Start with the work that is stuck	What a steering group should ask
Why tool-first adoption creates noise	Watch for five failure patterns
Name the constraint	Keep a small evidence pack
Choose the smallest change that can work	What to do next
Use five gates	Notes and sources
A worked example	

Start with the work that is stuck

Most AI requests start one step too late. A team sees a tool, model, or agent and asks where it could fit. Pilots begin. Usage grows. Only then do leaders ask where value should appear.

Start with the recurring work instead. Name what limits that work today. Then compare the changes that could remove the limit. AI may be the best choice. It may also be data cleanup, a clearer decision, training, a workflow change, simple automation, a vendor feature, or a decision to wait.

This order matters because AI activity is easier to count than changed work. McKinsey reported that 88 percent of survey respondents used AI in at least one business function in 2025. Only 39 percent reported an effect on earnings across the enterprise.[1] Those figures do not prove why the gap exists. They do show that widespread use and enterprise value are different things.

RAND found a related problem from the builder's side. Its review of AI project failures points to misunderstood problems, wrong measures, technology-first choices, and weak data.[2] Gartner later reported that many GenAI projects had been dropped after proof of concept because of poor data, weak controls, rising costs, or unclear value.[3] Different studies use different methods, but both point to the same operating risk: a working demo can still solve the wrong problem.

The first portfolio question should therefore be simple: Which recurring work is blocked, and what is blocking it? That question turns AI interest into a decision that leaders can compare with other uses of scarce money, people, and attention.

Why tool-first adoption creates noise

Tool-first adoption moves quickly. A vendor shows a strong result. Teams find tasks that look similar. Licenses spread. Dashboards show active users and prompts. The portfolio can report motion within a quarter.

But the business did not fund prompts. It funded faster service, fewer errors, lower rework, better decisions, broader coverage, lower cost, or less risk. Those results live in a workflow. They rarely live in the tool's own dashboard.

Tool-first reviews also mix three questions. Can the technology perform the task? Is the workflow ready for the change? Is this change worth the effort and risk? A capable model can fail in a workflow with poor data. A useful idea can fail because no one owns the result. An AI option can also lose to a cheaper change that solves the same problem.

This creates portfolio noise. One group asks for a chatbot. Another wants a summarizer. A third wants an agent. The requests use different measures and cannot be compared. The most visible idea often gets attention, even when another change would remove the real limit.

Name the constraint

A constraint is the current condition that limits value in recurring work. It may be a slow handoff, missing data, unclear approval, poor exception handling, scarce review time, or weak trust in the evidence. The constraint is not the broad goal. 'Improve productivity' is too vague. 'Reduce the two-day wait for a risk decision' is specific enough to test.

A useful intake asks five questions:

- What work repeats often enough for improvement to matter?
- Where does that work slow down, break, repeat, or consume scarce judgment?
- What is the current baseline?
- Who owns the result of the work?
- What evidence would show that the constraint moved?

These questions fit the NIST AI RMF Core. Govern, Map, Measure, and Manage form a continuing cycle for context, evidence, risk, and action.[5] Constraint-first adoption uses the same logic at the portfolio front door. It makes the work and its risk clear before a team chooses a tool.

Choose the smallest change that can work

The minimum effective intervention is the least complex change that can move the constraint with enough reliability, speed, safety, and value. 'Minimum' does not always mean cheap. A critical workflow may justify major redesign. It means the team chooses only what the problem requires.

Compare at least six response types: change the workflow, improve the data, train or reassign people, automate fixed rules, buy or configure a tool, or use AI for part of the work. Include a decision to wait when the facts are weak.

The order of work matters too. Better data may need to come before a model. Clear authority may need to come before an agent can act. A stable process may need to come before any automation. If the prerequisite is missing, fund that prerequisite first.

A McKinsey analysis argues that some AI value will require deeper changes to offerings and business models, not only small productivity gains.[7] That may be true for a high-value constraint. The portfolio still needs to state the choice, evidence, cost, and risk. A large bet should not receive a weaker decision standard because it sounds strategic.

Use five gates

A five-gate path keeps the decision visible from the first request through normal operations.

EXHIBIT 1**Five gates move from a stuck task to proven value**

- 1 Name the constraint**
Record the recurring task, current limit, baseline, affected group, and result that should improve.
- 2 Choose the change**
Compare workflow, data, people, fixed rules, vendor features, AI, and waiting.
- 3 Design the proof**
Set a small test, safeguards, use signal, workflow measure, and human review point.
- 4 Decide what happens next**
Use real workflow evidence to expand, change, stop, or try another option.
- 5 Keep checking value**
Once the change is normal, keep checking quality, cost, use, exceptions, and drift.

1. Constraint intake

The idea owner names the recurring work, the constraint, the baseline, the affected group, and the result that should improve. An idea without a named workflow remains an exploration. It does not enter the funded portfolio.

2. Intervention choice

A small group compares workflow, data, people, fixed-rule automation, vendor features, AI, and waiting. It chooses the smallest credible change and records why the other options lost.

3. Proof design

The team defines a small test, the operating safeguards, the use signal, and the workflow result that matters. McKinsey's measurement framework separates technical performance, adoption, operating measures, strategic outcomes, and financial impact.[6] An early test should not claim financial impact before the workflow result exists.

4. Reliance decision

Leaders read the evidence and choose to expand, change, stop, or try another intervention. A demo result is not enough. The change must work with real users, real data, real controls, and its full operating cost.

5. Value review

Once the change becomes part of normal work, owners keep checking quality, cost, use, exceptions, and drift. The work is successful when the workflow keeps producing the promised result, not when the tool ships.

A worked example

Consider a shared-services team that spends the first half of each month building a status view across 12 delivery streams. Leaders receive it late. The tool-first request is easy to imagine: buy a generative assistant to write the report faster.

Constraint intake changes the problem. The recurring work is the monthly portfolio review. The limit is not report writing. Leaders spend about 40 minutes rebuilding status before the first decision. The desired result is more time for decisions, with no loss of accuracy.

At the second gate, the group compares options. A summarizer would produce the same long recap sooner. A better response is to standardize the intake fields and use AI only to flag conflicts and critical risks. The team chooses workflow redesign plus AI support. It does not approve an autonomous agent.

The proof runs for two review cycles. It measures time to the first decision and checks whether the flagged issues match what leaders judge important. AI prepares the evidence. The portfolio lead owns every decision.

If the meeting reaches decisions sooner and the quality check holds, leaders can expand the test. If not, they can change the intake or stop the AI part. Either result creates useful evidence. The pilot does not need to become permanent to count as learning.

What a steering group should ask

A steering group does not need to become an AI design board. It needs a repeatable way to test whether a proposed change deserves the next unit of funding.

EXHIBIT 2**Five questions test the next funding step**

- 1 Can we observe the limit?**
Name where the task slows, breaks, repeats, or uses scarce judgment.
- 2 How should the result improve?**
Connect the change to time, quality, rework, coverage, cost, risk, or a customer result.
- 3 Did we compare simpler choices?**
Test workflow, data, people, fixed rules, vendor features, AI, and waiting.
- 4 Can the test produce evidence?**
Set the baseline, workflow measure, review point, and next decision before work starts.
- 5 Who owns the result?**
Name the process, technical, and reliance owners before the test begins.

Ask five questions:

- Is the constraint specific enough to observe?
- Is the path from the change to the result credible?
- Did the team compare AI with simpler options?
- Does the test name a baseline, workflow measure, and human review point?
- Are the process, technical, and reliance owners named?

The score matters less than the conversation. These questions expose assumptions while the choice is still easy to change.

Watch for five failure patterns

First is the jump from demo to rollout. The missing step is proof that the result survives the company's own workflow, data, people, and controls.

Second is the orphaned pilot. The experiment may work, but no process owner can change normal work if it succeeds.

Third is the usage trap. Tool activity rises while cycle time, quality, rework, or decisions stay the same. Use is supporting evidence, not the value claim.

Fourth is choosing the strongest model for every task. Many workflows need stable cost, repeatable results, and clear controls more than maximum capability.

Fifth is reliance that grows without a decision. An assistant's output quietly becomes an input to important choices. No one adds quality checks, a stop path, or a named owner. The reliance level should be explicit before convenience turns into authority.

Keep a small evidence pack

A decision should preserve the workflow and constraint, baseline, owner, options compared, chosen change, proof result, risks, and next gate. Keep the record short enough for leaders to use. Its purpose is to trace why the portfolio chose this path and what evidence would change that choice.

A McKinsey podcast on agentic organizations argues that daily workflows and working practices must change, with governance and risk controls added.[4] The paper uses that point as practitioner guidance, not proof that one design works everywhere. The evidence pack makes each local design and its limits visible.

What to do next

Take one active AI request before the next steering meeting. Replace the tool name with the recurring work it is meant to change. Write the constraint, baseline, owner, options, and one result that would justify another funding step.

EXHIBIT 3

Start with the constraint. Then choose the tool.

If the request cannot name those facts, keep it in exploration. If a simpler change can move the constraint, fund that first. If AI is the best option, run the smallest test that can prove it in real work.

Start with the constraint. Then choose the tool.

Notes and sources

[1] Alex Singla, Alexander Sukharevsky, Bryce Hall, Lareina Yee, Michael Chui, and Tara Balakrishnan, "The state of AI in 2025: Agents, innovation, and transformation," McKinsey & Company, November 5, 2025. Verified July 7, 2026. <https://www.mckinsey.com/capabilities/quantumblack/our-insights/the-state-of-ai-2025> <https://www.mckinsey.com/capabilities/quantumblack/our-insights/the-state-of-ai-2025>

- [2] James Ryseff, Brandon F. De Bruhl, and Sydne J. Newberry, "The Root Causes of Failure for Artificial Intelligence Projects and How They Can Succeed: Avoiding the Anti-Patterns of AI," RAND Corporation, 2024. Verified July 7, 2026. https://www.rand.org/pubs/research_reports/RRA2680-1.html
- [3] Arun Chandrasekaran, "Why 50% of GenAI Projects Fail - And How to Beat the Odds," Gartner, January 26, 2026. Verified July 7, 2026. <https://www.gartner.com/en/articles/genai-project-failure>
- [4] Alexis Krivkovich and Lucia Rahilly, "AI is everywhere. The agentic organization is not yet," The McKinsey Podcast, April 2, 2026. Verified July 7, 2026. <https://www.mckinsey.com/capabilities/people-and-organizational-performance/our-insights/ai-is-everywhere-the-agentic-organization-isnt-yet>
- [5] NIST, "AI Risk Management Framework Core," excerpt from AI RMF 1.0, 2023. Verified July 7, 2026. <https://airc.nist.gov/airmf-resources/airmf/5-sec-core/>
- [6] Johannes-Tobias Lorenz, Joshan Cherian Abraham, Robert Levin, and Douglas Ziman, "From promise to impact: How companies can measure and realize the full value of AI," McKinsey & Company, April 24, 2026. Verified July 7, 2026. <https://www.mckinsey.com/capabilities/quantumblack/our-insights/from-promise-to-impact-how-companies-can-measure-and-realize-the-full-value-of-ai>
- [7] Antoine Montard, Dago Diedrich, and Tanguy Catlin, "Where AI will create value - and where it will not," McKinsey Quarterly, April 29, 2026. Verified July 7, 2026. <https://www.mckinsey.com/capabilities/strategy-and-corporate-finance/our-insights/where-ai-will-create-value-and-where-it-wont>

About the author

Marco Policani is an enterprise portfolio, PMO, and AI operating-governance leader. He builds portfolio governance systems where AI carries the analytical load and named humans own every decision — an approach documented across case studies, walkthroughs, and working governance guides on his portfolio site.

Portfolio & governance library: policani.net/governance · Full portfolio: policani.net · LinkedIn: linkedin.com/in/marcpolicani

This white paper may be shared freely with attribution. © 2026 Marco Policani.