

---

WHITE PAPER · PORTFOLIO & AI GOVERNANCE

# AI Adoption Starts With the Constraint, Not the Tool

How to review one AI request before funding it

---

**Marco Policani**

Enterprise Portfolio · PMO · AI Operating Governance

[policani.net](http://policani.net) · [linkedin.com/in/marcpolicani](https://linkedin.com/in/marcpolicani)

July 2026

## EXECUTIVE SUMMARY

**Before funding an AI tool, name the recurring work it should change and the step that limits it. Then fund the smallest change that moves that step.**

When an AI request arrives as a tool, the question of where the value should appear is left for later. This paper recommends asking first which recurring work the request would change and what most limits it, which it calls the constraint. It then shows how to find the constraint, choose the smallest change, and take one request through five gates, steps that each end in a written output and a named decision-maker.

- 01** Name the recurring work, the constraint, the baseline, the owner, and the evidence that would show the constraint has eased. If nobody can name the constraint, follow about ten recent cases to find the step where cases lose the most time, are redone most often, or wait for a scarce person.
- 02** Fund the smallest change that can move the constraint. It may be a workflow change, better data, training, fixed rules, a vendor feature, or AI.
- 03** Take each request through five gates, each with a named decision-maker: name the constraint; choose the change; design a small test; decide whether to expand, revise, or stop; and keep checking value. Keep a small, capped allowance for exploration with no named workflow.

## CONTENTS

|                                              |                                              |
|----------------------------------------------|----------------------------------------------|
| A request that arrives as a tool             | Use five gates                               |
| Why a faster step may change nothing         | A worked example                             |
| Name the constraint                          | Five ways a request goes wrong anyway        |
| Find the constraint when nobody has named it | Objections, and what the sources do not show |
| Choose the smallest change that can work     | What to do next                              |
|                                              | Notes and sources                            |

---

## A request that arrives as a tool

A leader on a steering group, the group that decides where money and people go, receives an AI request. A team wants to try a tool, a model, or an agent. A pilot starts, people use it, and a dashboard reports activity. The question of where the value should appear tends to come later. That sequence is this paper's observation of how such requests arrive. No source counts how often it happens, and the sources below show why it matters.

This paper recommends reversing the order. Before anyone picks a tool, ask which recurring work the request would change and what most limits that work today. A request names a task the tool would speed up, and the limit may sit somewhere else. Speeding up a step that does not limit the work leaves the result where it was. The paper calls the limiting step the constraint: the step that most limits what the work delivers. The paper turns the question into a path for one request. Name the recurring work and its constraint, choose the smallest change that could move it, and fund each further step only on evidence.

The working belief behind a tool-first request, as this paper reads it, is that a pilot people use shows the idea is worth funding. McKinsey describes executives' focus on AI so far as largely pilots seeking incremental productivity gains,[1] which is the practice that belief supports. Use is not the same as value, though. McKinsey's 2025 survey found that 88 percent of respondents regularly use AI in at least one business function. Nearly two-thirds said their organizations had not begun to scale AI across the enterprise, and 39 percent reported an effect on earnings before interest and taxes at the enterprise level.[2] A survey cannot say why the gap exists.

Two other sources speak to how problems get chosen. RAND interviewed 65 experienced data scientists and engineers about AI projects that failed. They named five root causes. The first is that people on the business and technical sides misunderstood or miscommunicated what problem needed solving, so models were built for the wrong measures or did not fit the business workflow. The third cause is that organizations chased the latest technology instead of a real user problem.[3] Gartner reported in January 2026 that at least half of generative AI projects were abandoned after proof of concept by the end of 2025, with poor data, weak risk controls, rising cost, and unclear business value as causes. Its first failure point is lack of business value. Projects that chase flashy demos without success measures

become easy targets when budgets tighten or executives ask for proof of return. Gartner says it analyzed hundreds of implementations and shows no method.[4]

These sources use different methods. Together they show that pilots and use are common, that value does not follow automatically, and that practitioners and analysts name unclear problems and unclear value among the causes of failure.

RAND's advice is to focus on the problem, not the technology.[3] This paper adds three things that advice leaves open. The first is the unit to work on: one piece of recurring work and its constraint. The second is a fixed order of decisions, each with a named decision-maker. The third is rules for three cases: an unknown constraint, a test that gives no result, and a request that names no workflow. The title states the paper's recommendation, and no source tests it.

The first question for any request should therefore be simple: Which recurring work is blocked, and what is blocking it? Asking it turns AI interest into a decision that leaders can compare with other uses of scarce money, people, and attention.

---

## Why a faster step may change nothing

Operations management has an old name for the idea that a faster step can change nothing. Eliyahu Goldratt's Theory of Constraints starts from the observation that every system has at least one constraint, a situation that keeps it from reaching its purpose, and that the system's output is limited by that constraint. Improving any other step does not raise the output. When the constraint is fixed, another step becomes the limit, and the cycle repeats.[5] Goldratt first applied the idea to production planning and later explained it in a 1984 business novel, *The Goal*. A 2014 review describes its use in services, project management, and other settings.[5]

This paper borrows the idea and does not prove it for knowledge work or for AI adoption. A business workflow is also not a factory line, because its steps overlap and people work around them. The inference is narrower. When a team asks for AI, ask whether the tool relieves the step that limits what the work delivers. A useful test: if removing the wait at a step would only move the backlog to the next step, that step was not the constraint.

A constraint usually shows up as a slow handoff, missing data, an unclear approval, poor handling of unusual cases, scarce review time, or weak trust in the evidence. The constraint is not the broad goal. "Improve productivity" is a goal. "Reduce the two-day wait for a risk decision" names a constraint specific enough to test.

Tool-first adoption can move fast: licenses spread, and dashboards show active users and prompts. McKinsey's 2026 article on measuring AI value says many organizations deploy AI in ways that are more visible than valuable.[6] A business funds faster service, fewer errors, less rework, better decisions, broader coverage, lower cost, or less risk, not prompts. Those results show up in the workflow, not in the tool's own dashboard.

Tool-first review also mixes three questions. Can the technology do the task? Is the workflow ready for the change? Is the change worth the effort and risk? An AI option can lose to a cheaper change that solves the same problem. When one group asks for a chatbot, another for a summarizer, and a third for an agent, the requests use different measures and cannot be compared. Attention can then go to the most visible idea even when another change would remove the real constraint.

---

## Name the constraint

A useful request answers five questions. The business owner answers them, with help from whoever asked for the tool. The business owner is the leader accountable for the work's results, who would have to change how it is done.

- What work repeats often enough for improvement to matter?
- Where does that work slow down, break, repeat, or use scarce judgment?
- What is the current baseline, meaning a measurement of the result today?
- Who owns the result of the work?
- What evidence would show that the constraint has eased?

The first question asks for recurring work because RAND advises choosing enduring problems, since an AI project takes time and patience.[3] Reading "enduring" as "recurring" is this paper's own step: work that repeats keeps paying back a fix.

---

## Find the constraint when nobody has named it

A request seldom arrives with its constraint named, so the business owner has to find it. The four steps below are this paper's own method, not part of the Theory of Constraints review. The count of about ten cases is a planning estimate, not a rule.

## EXHIBIT 1

**Four steps find a constraint that nobody has named**

- ① **Ask the people who do it**  
Where does the work wait, get redone, or need a scarce person?
- ② **Follow about ten cases**  
Trace recent cases start to finish and mark each wait, redo, and scarce person.
- ③ **Pick the biggest delay**  
The step where cases lose the most time, are redone, or wait for a scarce person.
- ④ **Check it with the doers**  
Confirm it with the people who do the work before writing it down.

If the work has no baseline, measuring it for a short fixed period is the first funded step. It costs the work team's time and no tool purchase. If no step stands out, or the cases are too different to compare, the request has not named its constraint yet. It stays in exploration, a low-cost trial with no claim on funded work, which the objections below describe.

---

## Choose the smallest change that can work

The smallest change is the least complex one that can move the constraint with enough reliability, speed, safety, and value. Smallest does not always mean cheapest. A critical workflow may justify a major redesign. It means choosing only what the constraint requires.

Compare at least six kinds of response: change the workflow, improve the data, train or reassign people, automate fixed rules, buy or configure a tool, or use AI for part of the work. Waiting is also a choice. A companion paper on choosing the response once the constraint is named covers that choice in depth, including when waiting is right and how hard each response is to undo.

The order of work matters too. Better data may need to come before a model. Clear authority may need to come before an agent can act. A stable process may need to come before any automation. If the prerequisite is missing, fund that prerequisite first.

## Use five gates

A gate here is a step that ends in a written output and a named decision-maker: the business owner for gates 1, 2, 3, and 5, and the steering group for gate 4. The request moves on when the output exists and the decision-maker has signed it. The five gates are a sequence. Each takes what the one before produced and cannot start without it. Changes cannot be compared at gate 2 until the recurring work and its constraint are named at gate 1. A test cannot be designed at gate 3 until a change is chosen, and rollout cannot be decided at gate 4 without test results. Gate 5 has nothing to review until a change is in normal use.

### EXHIBIT 2

#### Five gates, each with a named decision-maker, take a request from tool to proven value

- 1 **Name the constraint**  
Business owner. Record the work, its constraint, the baseline, and the result wanted.
- 2 **Choose the change**  
The four: owner, daily worker, technical lead, analyst. Pick the smallest change.
- 3 **Design the proof**  
Business owner with the people who run it. Set the test, safeguards, and workflow measure.
- 4 **Decide what happens next**  
Steering group, who hold the budget. Expand, revise, stop, or try another.
- 5 **Keep checking value**  
Business owner. Check quality, cost, use, unusual cases, and drift.

Each gate adds to one record, the evidence pack: a one-page document the business owner starts at gate 1 and the portfolio or PMO analyst keeps current. McKinsey's measurement article recommends a single evidence pack that tracks benefits and total cost.[6] This paper's version also holds the recurring work and its constraint, the baseline, the owner, the options compared, the chosen change, the test and its result, the decision, and the date of the next review.

## 1. Constraint intake

The business owner, helped by the person who asked for the tool, answers the five questions in the evidence pack and lists who is affected. A request with no named recurring work, or with no one who will own the result, goes back to the person who asked, through the analyst. That person can find a business owner or leave it in exploration. It does not enter the funded portfolio. The cost is about an hour of the business owner's time, plus the time to measure a baseline if none exists.

## 2. Choice of change

A group of four compares the kinds of response and chooses the smallest credible change. If AI is the choice, the group also picks the smallest model or feature that can move the constraint. The business owner decides what counts as better and breaks any tie, because that person lives with the result. Someone who does the work every day knows where it actually waits. A technical lead knows the data and the tools, including AI. A portfolio or PMO analyst keeps the evidence pack and makes sure every option is compared the same way. The person who asked for the tool presents the request but is not one of the four, so the option's advocate does not also judge it. If the business owner is also the person who asked, the analyst asks the group to argue the strongest alternative first. A security or legal reviewer joins only when personal, regulated, or customer data is involved. NIST's framework asks that the people who set the context for an AI system bring broad expertise in the domain and in how it is used.[7] The choice of these four seats is this paper's own. The group records why the other options lost. The cost is one meeting of about an hour plus preparation.

## 3. Proof design

The business owner and the people who will run the test, the test team, write down a small test, its safeguards, one signal of use, and the workflow result that matters, before the test starts. Without a baseline the test cannot show change, so measuring comes first. McKinsey's framework separates technical performance, adoption, operating measures, strategic outcomes, and financial impact. It also recommends building measurement into the rollout through A/B testing or staggered deployment so the results hold up.[6] That is a consultancy's framework, not a tested standard. An early test should not claim financial impact before the workflow result exists. As a planning estimate, the test takes a few review cycles of the work. The business owner's team carries the time, and the steering group carries any spend.

#### 4. Rollout decision

The steering group, on the business owner's recommendation, chooses to expand the change, revise it, stop it, or try a different one. A demo result is not enough. The change must work with real users, real data, real controls, and its full operating cost. The decision also states how far people may depend on the tool's output and who checks it. NIST's framework likewise asks for a determination of whether a system meets its intended purposes and whether it should proceed.[7] A missing result differs from a weak one. A weak result is a reason to revise or stop. A missing result means the test did not work as a test, so nothing expands until it is rerun or redesigned. The decision itself takes one meeting. The larger cost is the rollout.

#### 5. Value review

Once the change is part of normal work, the business owner keeps checking quality, cost, use, unusual cases, and drift, meaning a result that slowly gets worse. The analyst keeps the evidence pack current, reads it against the original baseline, and reports to the steering group, so the business owner does not grade the result alone. As a planning assumption, review after the first full quarter and then whenever the work, the tool, or the data changes. If the review does not happen, the change counts as unreviewed, and the steering group funds nothing further for it until it does. The change succeeds when the workflow keeps producing the promised result, not when the tool ships. Fixing one constraint can expose the next, so this review may start a new round at gate 1.

For a request that needs only a small test, the planning estimate is a few hours of four people's time before anything is built. That is the paper's assumption, not a measurement.

---

### A worked example

Suppose a team assembles a status view across many delivery streams before each monthly portfolio review, and leaders get it late. This is an illustration, not an account of a real case, and it reports no result. The tool-first request is easy to imagine: buy a generative assistant to write the report faster.

Constraint intake changes the problem. The recurring work is the monthly portfolio review. Suppose that following recent cases showed leaders spending the first part of each review rebuilding status before the first decision. Then report writing is not the constraint. The desired result is more time for decisions, with no loss of accuracy. A faster report would leave the constraint where it is.

At the second gate, the group of four compares options. A summarizer would produce the same long recap sooner. A better response is to standardize the intake fields and use AI only to flag conflicts and critical risks. The group chooses workflow redesign plus AI support and does not approve an autonomous agent.

The test would run for two review cycles. It measures time to the first decision and checks whether the flagged issues match what leaders judge important. AI prepares the evidence, and the business owner, here the portfolio lead, owns every decision.

At the fourth gate, if decisions come sooner and the quality check holds, the steering group can expand the test. If not, it can revise the intake or stop the AI part. If the quality check was never run, nothing expands. Each result gives the next request something to build on.

---

## Five ways a request goes wrong anyway

These five are separate. No order is implied. Each maps to the gate meant to catch it. The list is this paper's judgment, built from the sources above. No source counts how often each occurs.

The first is the jump from a demo to rollout. The missing step is proof that the result survives the company's own workflow, data, people, and controls. Gate 4 asks for that proof. A companion paper on approving AI for production covers the approval record in full.

The second is the orphaned pilot. The experiment works, but no business owner can change normal work if it succeeds. Gate 1 refuses a request that has no owner.

The third is the usage trap. Tool activity rises while cycle time, quality, rework, or decisions stay the same. Use is supporting evidence, not the value claim. Gate 3 sets the workflow measure before the test starts.

The fourth is choosing the strongest model for every task. Some workflows need stable cost, repeatable results, and clear controls more than maximum capability. Gate 2 asks for the smallest change that works, which can include a smaller model.

The fifth is use that grows without a decision. An assistant's output quietly becomes an input to important choices. No one adds quality checks, a stop path, or a named owner. How much people depend on the tool should be explicit before convenience turns into authority. Gate 4 makes that decision, and gate 5 checks it afterward.

---

## Objections, and what the sources do not show

### Exploration finds uses nobody planned

A constraint-first path can miss value that nobody yet sees as a constraint. McKinsey argues that the larger prize from AI lies in reshaping offerings and business models more than in productivity gains.[1] That is a consultancy's argument, not a finding, and it cuts against a path built on today's workflow. If it is right, this path can fund efficient work that wins no lasting advantage. The path is built for requests that change existing work, so it accepts that cost. For the rest, the paper's answer is to keep exploration open and cap it: a small fixed allowance of money and time for trying tools with no named workflow, reviewed on a set date, with a rule that anything wanting more must come back through gate 1. No source says how large the allowance should be. The steering group sets it. Bets on new offerings or markets are strategy decisions, not intake decisions, and they still deserve a stated choice, evidence, cost, and risk.

### The gates take time

The gates slow adoption. For a low-risk use that is easy to reverse, such as a team trying an approved assistant on its own drafts, a short conversation at gates 1 and 2 is enough. That is this paper's suggestion. Use all five gates when a request needs money, shared data, customer-facing use, or a change to how a team works. Gartner's figure shows that many generative AI projects are dropped after proof of concept, and McKinsey's shows that most organizations have not begun to scale. Neither shows that gates would have prevented it. The paper's estimate is a few hours of four people's time. The comparison is with a pilot that runs without a decision, and no source prices that.

### A factory idea in an office

Knowledge work often has several things holding it back at once, and people route around the ones they can. The Theory of Constraints is a way to focus the question, not a law of office work. If the first constraint is fixed and the result barely moves, the group has probably named the wrong one or reached the next one. Gate 5 exists to catch that.

### What the sources do not show

No source compares constraint-first review with tool-first review, so none shows how much of the gap between use and value this path would close.

When several requests compete for the same people and money, comparing them is a different job, covered in a companion paper on reviewing several AI proposals with one brief. This paper covers the first decision about one request.

---

## What to do next

Take one active AI request before the next steering meeting. Replace the tool name with the recurring work it is meant to change. Write the constraint, the baseline, the owner, and one result that would justify another funding step. Then list the simpler changes that could move the constraint. The request becomes something the steering group can compare, fund in steps, and stop, with a record of why.

If the request cannot name those facts, keep it in exploration. If a simpler change can move the constraint, fund that first. If AI is the best option, run the smallest test that can prove it in real work.

Start with the constraint. Then choose the tool.

---

## Notes and sources

[1] Antoine Montard, Dago Diedrich, and Tanguy Catlin, “Where AI will create value—and where it won’t,” McKinsey Quarterly, April 29, 2026.

<https://www.mckinsey.com/capabilities/strategy-and-corporate-finance/our-insights/where-ai-will-create-value-and-where-it-wont>

[2] Alex Singla, Alexander Sukharevsky, Bryce Hall, Lareina Yee, Michael Chui, and Tara Balakrishnan, “The state of AI in 2025: Agents, innovation, and transformation,” McKinsey & Company, November 5, 2025.

<https://www.mckinsey.com/capabilities/quantumblack/our-insights/the-state-of-ai-2025>

[3] James Ryseff, Brandon F. De Bruhl, and Sydne J. Newberry, “The Root Causes of Failure for Artificial Intelligence Projects and How They Can Succeed: Avoiding the Anti-Patterns of AI,” RAND Corporation, 2024.

[https://www.rand.org/pubs/research\\_reports/RRA2680-1.html](https://www.rand.org/pubs/research_reports/RRA2680-1.html)

[4] Arun Chandrasekaran, “Why 50% of GenAI Projects Fail — And How to Beat the Odds,” Gartner, January 26, 2026. <https://www.gartner.com/en/articles/genai-project-failure>

- [5] Zeynep Tuğçe Şimşit, Noyan Sebla Günay, and Özalp Vayvay, "Theory of Constraints: A Literature Review," *Procedia — Social and Behavioral Sciences* 150 (2014): 930-936. <https://openaccess.marmara.edu.tr/entities/publication/ddb50b4f-d650-4d0a-b760-4c7780c10e09/full>
- [6] Johannes-Tobias Lorenz, Joshan Cherian Abraham, Robert Levin, and Douglas Ziman, "From promise to impact: How companies can measure and realize the full value of AI," McKinsey & Company, April 24, 2026. <https://www.mckinsey.com/capabilities/quantumblack/our-insights/from-promise-to-impact-how-companies-can-measure-and-realize-the-full-value-of-ai>
- [7] NIST, "AI Risk Management Framework Core," excerpt from AI RMF 1.0, 2023. <https://airc.nist.gov/airmf-resources/airmf/5-sec-core/>

### About the author

Marco Policani is an enterprise portfolio, PMO, and AI operating-governance leader. He builds portfolio governance systems where AI carries the analytical load and named humans own every decision — an approach documented across case studies, walkthroughs, and working governance guides on his portfolio site.

Portfolio & governance library: [policani.net/governance](https://policani.net/governance) · Full portfolio: [policani.net](https://policani.net) · LinkedIn: [linkedin.com/in/marcpolicani](https://linkedin.com/in/marcpolicani)

This white paper may be shared freely with attribution. © 2026 Marco Policani.