


AI Knowledge Bases Are Operating Infrastructure, Not "Second Brains"

Why enterprise AI depends on owned, current, and traceable source material



Marco Policani

Enterprise Portfolio · PMO · AI Operating Governance

policani.net · linkedin.com/in/marcpolicani

July 2026

EXECUTIVE SUMMARY

Reliable AI begins with owned, current source material.

An AI assistant can retrieve the right-looking page and still give a wrong answer because the page was stale before the model ever saw it. Part of the fix has to come before retrieval: a named owner approves each source before the assistant may rely on it.

- 01** A fluent answer can hide a stale or wrong source, so the sponsor should approve reliance only after the sources behind it pass the owner's entry check, not on the strength of the demo.
- 02** Five disciplines govern what enters the knowledge base: ownership, quality control, update discipline, access governance, and lifecycle management.
- 03** Start with the knowledge domain where a wrong answer carries the highest consequence, and require provenance for every trusted source.

The answer passed the demo

Imagine an executive sponsor reviewing an AI assistant before people rely on it for operating work. The answer appears quickly, reads well, and cites a policy page that looks authoritative. During the review, the team discovers that the page was replaced six months ago. The retrieval worked. The model followed the source. The answer was still wrong.

The sponsor must decide whether that knowledge base is reliable enough for operating use. A clean interface and a good model can make the system look ready while the source underneath it remains stale, contradictory, or unowned. The resulting errors are difficult to spot because generative AI turns weak source material into fluent, plausible prose.

The personal-notebook frame examined here assumes that the person who keeps the notes is also the person who reads them. That person remembers which page is old, what a shorthand phrase means, and which two notes disagree. More notes may be useful because one mind supplies the missing context at the moment of use.

Enterprise use removes that protection. The person or system reading a source often did not create it and cannot see the context that would expose a defect. An admitted document may shape answers, recommendations, and actions across the organization. Approval makes the collection a shared dependency used beyond its creator's context. This is the tacit-knowledge problem in operating form: a personal second brain runs on the context its owner carries but cannot fully write down — what

Michael Polanyi meant by knowing more than we can tell — and enterprise reuse strips exactly that tacit layer, leaving explicit records that cannot flag their own decay.

Why the error survives the model

A study of data quality in retrieval-augmented generation found that problems concentrate in source and ingestion stages. They then spread through retrieval and generation. The study is qualitative, based on practitioner interviews rather than a controlled estimate of how much each defect changes output. The authors still draw a clear direction: most coded challenges arise at data extraction, and they report that quality issues carry forward into later stages.[1]

Teams often respond to a weak answer at the visible layer. They adjust the prompt, tune retrieval, or compare another model. Those actions may improve how the system finds and presents content. They cannot make a superseded policy current or decide which of two conflicting pages the business accepts.

The stakes grow with use. One wrong answer may be corrected by the person who recognizes it. Repeated, plausible errors spread through decisions before anyone knows which source caused them. Users then lose trust in the AI, including answers supported by good material. Investigation also slows because an unowned source gives the reviewer no clear person, history, or validation record to inspect. The cost arrives after adoption, across many answers, while the cost of source control is visible up front.

Reliability is decided at source entry

The NIST Generative AI Profile places information integrity and data provenance within the risk-management work around generative AI.[2]

EXHIBIT 1

| A trusted retrieval set begins with an accountable source owner.

The NIST AI Risk Management Framework Core separates the work into GOVERN, MAP, MEASURE, and MANAGE.[3] Applied here, those functions require an owner and known source context. The organization checks quality and acts when the source changes.

Five disciplines at the source-entry gate

Five controls govern admission. The owner rejects a source if any control lacks an answer or if a material conflict remains.

EXHIBIT 2

Five controls govern source admission

- 1 **Ownership**
Name one owner who can approve changes and act when the source fails.
- 2 **Quality control**
Check accuracy, consistency, completeness, and timeliness for the domain.
- 3 **Update discipline**
Set the event or cadence that forces review and retires stale material.
- 4 **Access governance**
Control who may read sensitive material and who may write to the trusted set.
- 5 **Lifecycle management**
Show validation, maintenance, deprecation, removal, and source provenance.

These controls expose the problem hidden by the personal-notebook frame. When two documents disagree, the owner and lifecycle record provide a path to identify the source of record and the change that made another copy stale. When a domain cannot supply an approved answer, the absence remains visible instead of being filled by an undated draft. Rejecting a defective source at admission keeps it out of later answers that would retrieve it. These controls extend established knowledge-management practice rather than invent a new one. ISO 30401 sets requirements for establishing, implementing, maintaining, reviewing, and improving a knowledge management system,[4] and NASA governs the capture, storage, reuse, and continuity of its critical knowledge as agency practice.[5]

Provenance makes an answer inspectable

The five disciplines govern the source as an asset. Provenance connects that asset to the answer a person sees. A useful record identifies the source, owner, and last review. It also shows the basis for approval and whether a newer version replaced it. The amount of detail should match the consequence of relying on the answer.

Without provenance, an answer can sound authoritative while leaving no path back to the decision that made its source trustworthy. A reviewer cannot tell whether the content was approved, copied,

inferred, or carried forward after its context changed. The result is an answer with no defensible approval trail.

Capture provenance during admission. Reconstructing it later is slow and sometimes impossible because the owner, rationale, or prior version may be gone.

Start with the knowledge that can do the most harm

An enterprise does not need to govern its entire knowledge estate at once. The practical first move is to choose one domain where a wrong answer carries a material consequence. It may be a policy used to approve work, a procedure that controls access, or guidance that shapes an external commitment. The specific domain depends on the organization; the test is the consequence of being confidently wrong.

For that domain, define the trusted retrieval set and assign the owner who approves admission. The approval record names the quality standard, refresh trigger, write permissions, and provenance. It also states which events deprecate or remove content. The team then tests the AI with real questions whose answers can be traced to those sources.

This bounded start keeps the work proportionate. It also shows whether answers come from current material and whether conflicts are visible. The sponsor can test whether a reviewer can reconstruct the source and whether an owner responds when it changes. Expansion to another domain follows after the sponsor sees that control working.

What source governance can and cannot promise

Source control is one necessary layer of reliable AI. The system must also retrieve, test, and use the source well. The five disciplines do not make every answer correct, and one control depth will not fit every domain.

The evidence also does not provide a universal return formula. The RAG study identifies where practitioners encounter data-quality problems and reports that those problems carry across stages. It does not say how much an organization should spend on each source or how much reliability a given control will buy. Those choices require local measures and a level of governance proportionate to the decision risk.

Approve the dependency

Return to the answer that passed the demo. The model retrieved the page it had been allowed to trust. The failure began when a superseded page remained inside the trusted set without an owner, review trigger, or visible lifecycle status.

The executive decision is therefore concrete. Choose the highest-consequence knowledge domain and approve reliance only after its sources pass a source-entry gate. Use the five disciplines to keep the set current and controlled, and provenance to make each answer traceable. That is how a document collection becomes operating infrastructure the organization can trust and inspect.

Notes and sources

- [1] Leopold Müller, Joshua Holstein, Sarah Bause, Gerhard Satzger, and Niklas Kühl, "Data Quality Challenges in Retrieval-Augmented Generation," arXiv:2510.00552, submitted October 1, 2025. <https://arxiv.org/abs/2510.00552>
- [2] Chloe Autio, Reva Schwartz, Jesse Dunietz, Shomik Jain, Martin Stanley, Elham Tabassi, Patrick Hall, and Kamie Roberts, "Artificial Intelligence Risk Management Framework: Generative Artificial Intelligence Profile," NIST AI 600-1, July 26, 2024. <https://doi.org/10.6028/NIST.AI.600-1>
- [3] National Institute of Standards and Technology, "AI Risk Management Framework Core," AI RMF 1.0, 2023. <https://airc.nist.gov/airmf-resources/airmf/5-sec-core/>
- [4] International Organization for Standardization, "ISO 30401:2018 — Knowledge management systems — Requirements," November 2018. <https://www.iso.org/standard/68683.html>
- [5] National Aeronautics and Space Administration, "Knowledge Management," APPEL Knowledge Services. <https://www.nasa.gov/learning-resources/for-professionals/appel-knowledge-management/>

About the author

Marco Policani is an enterprise portfolio, PMO, and AI operating-governance leader. He builds portfolio governance systems where AI carries the analytical load and named humans own every decision — an approach documented across case studies, walkthroughs, and working governance guides on his portfolio site.

Portfolio & governance library: policani.net/governance · Full portfolio: policani.net · LinkedIn: linkedin.com/in/marcpolicani

This white paper may be shared freely with attribution. © 2026 Marco Policani.