


AI Knowledge Bases Need Owners and Current Sources

Keep the sources behind AI answers current and traceable



Marco Policani

Enterprise Portfolio · PMO · AI Operating Governance

policani.net · linkedin.com/in/marcpolicani

July 2026

EXECUTIVE SUMMARY**Approve each source and keep a path from every AI answer back to it.**

An AI assistant can give a wrong answer from a page it found and used correctly. Check who approved the page, when it was last reviewed, and whether it still applies.

- 01** Name one owner to approve each source before AI uses it.
- 02** Check quality, review triggers, access, and status when a page enters the trusted set.
- 03** Keep a path from each answer back to its approved page.

The answer passed the demo

An executive sponsor tests an AI assistant before staff use it at work. The answer comes back fast and cites a policy page. Then the team finds out that the page was replaced six months ago. The system found the page and used it correctly. The answer was still wrong.

The sponsor must decide whether staff can rely on the knowledge base. A clear screen and a capable model can make it look ready. Yet its sources may be old, at odds with each other, or owned by no one. AI can turn those weak sources into answers that sound right.

Personal notes work because their owner knows what they mean. That person remembers which page is old, what a short phrase means, and which notes disagree. That memory fills gaps in the written record.

At work, other people and systems use the notes. They may not know why a page changed or which version still applies. Once a team approves a document for AI use, that document may shape answers across the company. The owner's private memory no longer protects the reader. A shared knowledge base needs records that carry enough context for others to check the source.

Why the error survives the model

A study of AI systems that retrieve documents before answering found many data problems in the sources and the process that loads them. Those problems then carry into the answer. The researchers interviewed practitioners. They did not measure how much each defect changes an answer. Their findings still show why teams must check source quality early.[1]

When an answer is weak, teams may change the prompt, improve search, or try another model. These steps may help the system find and explain a page. They cannot make an old policy current. They also cannot decide which of two conflicting pages the company accepts.

The cost grows as more people use the tool. One person may spot a bad answer. Many may not. Wrong answers can shape decisions before anyone finds the faulty source. People then lose trust even in sound answers. Without a source owner or review history, the team cannot quickly find out what went wrong.

Reliability is decided at source entry

NIST's Generative AI Profile calls for work on information integrity and records of where content came from.[2]

EXHIBIT 1

The source owner decides which page the company accepts before AI uses it.

NIST's AI Risk Management Framework gives teams four jobs: govern, map, measure, and manage risk.[3] For a knowledge base, that means naming an owner, checking each source, and acting when it changes.

Five disciplines at the source-entry gate

Use five checks before a source enters the trusted set. The owner rejects it if a key question has no answer or two important pages still conflict.

EXHIBIT 2

Five checks before AI uses a page

- 1 **Owner**
Name who approves changes and acts when the page fails.
- 2 **Quality**
Check that it is accurate, complete, current, and consistent.
- 3 **Review**
Set a date or event that requires a new check.

4 Access

Limit who may read or change trusted pages.

5 Status

Show whether the page is under review, current, replaced, or removed.

When two pages disagree, the owner and change record show which one still applies. If no approved answer exists, the gap stays visible. The system should not fill it with an undated draft. These checks build on established practice. ISO 30401 sets requirements for knowledge management systems,[4] and NASA has a policy for preserving and using critical knowledge.[5]

Provenance makes an answer inspectable

The five checks control the source. The next step links that source to the answer a person sees. Record the page, its owner, when it was checked, why it was approved, and whether a newer version replaced it. Keep more detail when a wrong answer could do more harm.

EXHIBIT 3

Keep the source, owner, last check, approval reason, and replacement history with the answer.

Without that link, a reviewer cannot tell whether the answer came from an approved page, an old copy, or the model's own guess. A fluent answer is not enough to show why anyone should trust its source.

Keep this record when the owner approves the page. Later, the owner, the reason for approval, or the old version may be gone.

Start with the knowledge that can do the most harm

Start with one area where a wrong answer matters. It may be a policy used to approve work, a rule that controls access, or guidance used with a customer. Choose based on what could happen if the AI gives a wrong answer with confidence.

Choose which pages the AI may use and name the owner who approves them. Record the quality check, review trigger, who may change a page, and when to replace or remove it. Then ask real questions and trace each answer to an approved page.

This first test shows whether the AI uses current pages and whether conflicts are visible. The sponsor can check whether reviewers can find the source and whether the owner acts when it changes. Expand only after that works in the first area.

What source governance can and cannot promise

Source checks help, but the AI must still find and use the right page. The five checks cannot make every answer correct. A high-risk area may also need more review than a low-risk one.

The study does not say how much each check will cost or how much it will improve results. Teams need to measure that in their own work and match the review effort to the risk of a wrong answer.

Approve the dependency

Return to the answer that passed the demo. The model used a page the team had allowed it to trust. The failure began when an old page stayed in that set with no owner, review trigger, or clear status.

The next decision is clear. Pick an area where a wrong answer would matter. Let staff rely on its knowledge base only after the owner approves each source through the five checks. Keep a path from each answer back to its approved page.

Notes and sources

[1] Leopold Müller, Joshua Holstein, Sarah Bause, Gerhard Satzger, and Niklas Kühl, "Data Quality Challenges in Retrieval-Augmented Generation," arXiv:2510.00552, submitted October 1, 2025. <https://arxiv.org/abs/2510.00552>

[2] Chloe Autio, Reva Schwartz, Jesse Dunietz, Shomik Jain, Martin Stanley, Elham Tabassi, Patrick Hall, and Kamie Roberts, "Artificial Intelligence Risk Management Framework: Generative Artificial Intelligence Profile," NIST AI 600-1, July 26, 2024. <https://doi.org/10.6028/NIST.AI.600-1>

[3] National Institute of Standards and Technology, "AI Risk Management Framework Core," AI RMF 1.0, 2023. <https://airc.nist.gov/airmf-resources/airmf/5-sec-core/>

[4] International Organization for Standardization, "ISO 30401:2018 — Knowledge management systems — Requirements," November 2018. <https://www.iso.org/standard/68683.html>

[5] National Aeronautics and Space Administration, "Knowledge Management," APPEL Knowledge Services. <https://www.nasa.gov/learning-resources/for-professionals/appeal-knowledge-management/>

About the author

Marco Policani is an enterprise portfolio, PMO, and AI operating-governance leader. He builds portfolio governance systems where AI carries the analytical load and named humans own every decision — an approach documented across case studies, walkthroughs, and working governance guides on his portfolio site.

Portfolio & governance library: policani.net/governance · Full portfolio: policani.net · LinkedIn: linkedin.com/in/marcpolicani

This white paper may be shared freely with attribution. © 2026 Marco Policani.