
WHITE PAPER · PORTFOLIO & AI GOVERNANCE

The AI Replacement Boomerang Is a Governance Failure

Why analysts expect many AI-driven cuts to be abandoned or reversed, and the check to run first

Marco Policani

Enterprise Portfolio · PMO · AI Operating Governance

policani.net · linkedin.com/in/marcpolicani

July 2026

EXECUTIVE SUMMARY

Book an AI-driven saving only after a short, written check of what the role carried beyond its task, who will own the unusual cases, and how junior staff will become senior.

Gartner (for customer service) and Forrester forecast that many AI-driven workforce cuts will be abandoned or quietly reversed, but neither measures why one company's cut fails. This paper's inference is that cuts come back when no one first checks what the role carried, who will own the hard cases (the unusual cases a routine process cannot settle), and how junior staff will become senior. A redesign gate, a required check before the saving is booked, puts those three questions in front of the decision.

- 01** The sponsor answers five questions on one page before the saving is booked. Has anyone redesigned how the team works, starting from what the role carried? Will the tool replace the person's output or extend it? Who decides and answers afterward? What happens to the pipeline of future senior staff? What would show the decision was wrong? A missing answer holds the saving.
- 02** Read four warning signs together and against a baseline taken before the cut: rework, hard cases with no owner, a thinning senior bench, and backfill under a different cost code. One sign alone proves little.
- 03** Keep the decision reversible: let the tool carry part of the task while the people stay, and take the headcount decision last.

CONTENTS

Analysts expect many AI-driven cuts to be abandoned or reversed

What a demo shows and a role carries

Automate or augment: what the evidence says about the choice

The pipeline that cannot be rehired quickly

Where this argument is weaker

A redesign gate before any cut

One page, attached to the proposal

Two situations, walked through the page

The warning signs the budget hides

Keep the decision easy to undo

Treat every cut as a redesign

Notes and sources

Analysts expect many AI-driven cuts to be abandoned or reversed

Consider a cost review that receives this proposal. A pilot showed that an AI tool can do a routine part of a team's work, so the proposal cuts the team and books the saving. The saving is real in the budget on the day it is booked. What has changed, in this paper's reading, is that a pilot can now show a task done quickly and cheaply, while nothing in a pilot shows what the role carried. The open question is whether the saving is still there a year later.

Two analyst firms forecast that, for many organizations, plans like this will not hold. In June 2025, Gartner predicted that by 2027, half of the organizations that expected to significantly reduce their customer service workforce because of AI will abandon those plans. It attributes the shift to many companies struggling to reach staffing goals that have no human agents, and it cites a March 2025 poll of 163 customer service and support leaders, 95 percent of whom said they plan to keep human agents to define AI's role.[1] In November 2025, Forrester wrote in its 2026 workforce predictions that it expects half of the layoffs attributed to AI to be quietly reversed, with the jobs returning offshore or at lower wages.[2] These are forecasts, not counts of what has happened. Gartner's covers customer service only, and an organization that abandons a plan may never have made the cut.

Read the other way, the same forecasts say that about half of these plans and layoffs will hold. Neither measures why one company's cut fails and another's does not. Forrester's explanation is that replacing people with machines is not always cheaper or smarter unless a firm has a complete approach that accounts for the people AI success depends on.[2] That is a reason, not a measurement.

This paper makes one claim: an AI-driven cut is a decision about how a team's work is designed, and it should be decided that way, not only as a headcount line. Its inference is that cuts come back when the decision starts from what the tool can do. In that case nobody first checks three things: what the role carried beyond the task, who will own the hard cases (the unusual cases a routine process cannot settle), and how junior staff will become senior ones. A boomerang is a cut whose cost returns as rework, rehiring, or backfill. A governance failure, here, is a cut approved without written answers to those three questions. In plain words, the paper's position is that a likely cause of a return is that nobody checked first, not a surprise from the technology. Some cuts blamed on AI have other causes, and some fail because the

tool performs worse than the demo suggested. The paper does not cover those, and a later section says how to tell the difference.

The sections build in order: what a role carries that a demo does not show, what research says about automating versus augmenting, and what cannot be rehired once it is lost. The paper then tests its argument before it proposes a redesign gate, a required check before the saving is booked.

What a demo shows and a role carries

A task is one step of work, such as drafting a summary. A role is the job a person holds, which usually bundles many tasks. A workflow is how a team's tasks, handoffs, and decisions fit together. A demo shows that a model can do a task. A headcount line removes a role. A role often carries more than its task: the hard cases nobody wrote down, the knowledge of why things are done a certain way, the relationships, and the judgment that catches a mistake early. That is a hypothesis to test for each role, and the gate later in this paper is how to test it.

Redesign is still rare, by one survey's count. Deloitte's article on operating models for people working with AI agents reports that its State of AI in the Enterprise 2026 survey found that 84 percent of companies had not redesigned jobs to fit AI, even though expectations for automation were high.[3] The figure measures redesign, not cuts, so it does not show how many cuts were made without it. It suggests that at most companies a cut made today would rest on a job not yet redesigned for AI. Whether anyone examined the role first, the survey does not say.

Automate or augment: what the evidence says about the choice

If a role carries more than its task, the next question is what happens to the people when a tool takes the task. The tool can replace what a person produces, or it can extend it. The largest dataset this paper found on that difference is Stanford's. Researchers at the Stanford Digital Economy Lab studied monthly payroll records from ADP, the largest payroll software provider in the United States, covering millions of workers through September 2025. They found a 16 percent relative decline in employment for workers aged 22 to 25 in the most AI-exposed occupations, compared with the least exposed, after controlling for shocks at individual firms. Employment for more experienced workers in the same occupations stayed stable.

The declines were concentrated in occupations where AI use looked more like automating the work than helping a person do it.[4]

The authors classify occupations as automating or augmenting from estimates based on conversations with one AI assistant, Claude, not from what employers deployed. They tested other explanations, including exposure to interest rates, and report that the pattern held. They also caution that their estimates may be influenced by factors other than generative AI.[4] The study is early and covers U.S. payroll data, so it cannot show that leaders chose the outcome. This paper's inference is that inside a single company, whether a deployment automates a role or augments it is a design choice leaders can govern, because the same model aimed at the same job can be set up either way.

McKinsey, a consulting firm, argues that AI-first work means redesigning end-to-end processes, with people above the loop who steer outcomes. It also says agents are replacing tasks historically handled by knowledge workers and that small human teams can supervise many agents, so its design can mean smaller teams.[5] That is an argument, not a measured result. It supports the point this paper tests: when the work is redesigned, the cut should follow the redesign instead of coming before it.

The pipeline that cannot be rehired quickly

The cost that is hardest to see on a budget line is the loss of the people who would have become experienced. Entry-level work is where junior staff pick up context, judgment, and relationships. This paper uses "the pipeline" for that path from junior to senior. Stanford's sharpest decline is among workers aged 22 to 25, and it can be read forward. If organizations stop hiring and developing junior staff because AI covers the junior tasks, fewer experienced people will be available in a few years.[4] That is an inference. The study measures employment now, not future skill, and it does not show a shortage.

The mechanism is older than AI. In "Ironies of Automation," Lisanne Bainbridge argued in 1983 that automating the routine parts of a job removes the practice that builds and keeps skill. The people who remain are then less able to handle what the machine cannot, and they need more training, not less.[6] Her examples came from industrial plants, so applying them to entry-level work is an analogy.

Forrester's forecast adds a point about what a reversal looks like: the reversed jobs return offshore or at lower wages.[2] That suggests a reversal may not put the original team back, and the cost can return in a different form and place, outside the

budget that booked the saving. The reasoning here is that a pipeline is also slower to rebuild than a team is to rehire, because the experience it produced took years to form. So leaders should judge an AI-driven workforce decision against the skills the organization will need in three to five years, not only against the tasks it removes this year.

Where this argument is weaker

Five objections limit what this paper claims. The gate that follows responds to them.

EXHIBIT 1

The sources support a forecast and a mechanism, not a measured rate

- 
Forecast
 Gartner and Forrester expect abandoned plans and quietly reversed layoffs.
- 
Measured
 Stanford: a 16 percent relative decline for ages 22 to 25, most exposed jobs.
- 
Analogy
 Bainbridge shows how automating routine work erodes skill.
- 
Argument
 McKinsey and NIST argue for redesign and ongoing management.
- 
Not established
 No source counts reversals or shows why one cut fails.

Full automation is sometimes right. Some tasks are narrow, stable, and need little judgment, and refusing to automate them protects no one. A defensible automation has four features: someone understands the task instead of assuming it, the judgment and the handling of hard cases that the role carried have moved somewhere instead of disappearing, a named owner answers for the output and can stop it, and someone has weighed the effect on the pipeline. When all four hold, removing the person from a step can be a sound operating choice.

Some reversals are not governance failures. A tool that performs worse than the demo suggested is a technology shortfall that no redesign would have fixed. A cut blamed on AI that had other causes is a different problem. The title's claim therefore needs a test. When a cut comes back, ask whether anyone wrote down, before approval, what the role carried, who would own the hard cases, and how junior staff would become senior ones. If no one did, the return is a governance failure in this paper's sense. If someone did, and the cut failed anyway, it was another kind of failure, and the written answers show which kind. The test labels a returned cut. It does not prove that the answers would have prevented the return, and cuts that skipped the answers may have held. So the title states the paper's position, supported by the forecasts, the mechanism, and Forrester's own explanation. It is not a measured finding.

The forecasts may be wrong. Gartner's covers customer service only, and Forrester's full report is open to clients, so only its blog statement could be read for this paper. Neither is a count of reversals.

The evidence for the pipeline is a mechanism, not a measurement. This paper found no study of the multi-year effect of cutting entry-level roles, and the period is still short. The pipeline is a risk to govern, not a certainty.

The gate costs time: it delays the booking of the saving and takes time from the people who do the work. The section on finding what a role carried estimates that cost.

These limits narrow the claim. The paper does not say AI-driven cuts always return or that people must always stay. It says a cut is better decided with the redesigned workflow in front of the decision. The gate takes days of effort, and the forecasts put a reversal at about half for the plans and layoffs they cover.

A redesign gate before any cut

Slowing AI adoption is not the remedy, and neither is protecting roles for their own sake. Route every AI-driven workforce proposal through a redesign gate, a required check before anyone books a headcount change. The sponsor, the leader who brings the proposal, answers five questions in writing. The head of the affected function and a finance partner review the answers, and the executive who approves the budget decides. Reviewers use judgment on one point above all: an answer such as "the team will handle hard cases" is not yet an answer. It needs a name and a route.

EXHIBIT 2

Five questions to answer before the cut**Workflow redesigned?**

Write down what the role carried, what the tool does, and where a hard case goes.

Augment or automate?

If the person leaves, name the judgment and the hard-case handling that leave too.

Who decides now?

Name the person who owns the outcome and can show how it was reached.

Pipeline protected?

If junior work disappears, fund another path to grow the senior bench.

Early warning named?

Name the signals and take a baseline now, so a mistake shows while it is cheap to reverse.

The first question comes first, because the others have nothing to attach to until the workflow is written down. The second follows from it. The last three can be answered in any order. The three questions named at the start sit inside these five: what the role carried and who owns the hard cases are part of the first, and the junior path is the fourth.

First, has anyone redesigned the workflow, starting from what the role carried beyond its task? What the person still owns, what the tool now does, and where a hard case goes should be written down instead of assumed. For how to design who owns the hard cases, see the companion paper "The Exception Path Is the Real Automated Workflow." Second, is this augmentation or automation? If it removes the person, name the judgment and the handling of hard cases that leave with them. Third, who decides now? Someone must still own the outcome and be able to show how it was reached. Fourth, what does this do to the pipeline? If it cuts the entry-level work, say how the organization will grow the senior people it needs to supervise the tool. Fifth, what would show that this decision was wrong? Name the signals and take a baseline before the cut. A headcount report may not show a quiet reversal, so the baseline has to exist before the cut for a mistake to be seen while it is still cheap to reverse.

How to find out what a role carried

The first question depends on knowing what the role carried beyond the task. Three sources answer it, and they check each other. The sponsor's team runs the work, an analyst from finance or the PMO reads the cases, and the head of the affected function confirms the result.

- Ask the people who do the work. Have each person list what they handled recently that is not in the job description: the cases they fix, who asks them for help, and what they check before passing work on.
- Sample real cases. Pull recent cases, including the odd ones, and trace what would have happened to each under the new design. Mark the ones the tool would have mishandled or passed on, and who would have received them.
- Read the escalation and rework record. The log of escalations, reworked items, and complaints shows which cases left the routine path and who answered them.

Listen for what a headcount line can hide: hard cases the role absorbed, knowledge that kept a process stable, the training ground for future senior staff, judgment that caught errors early, and accountability for the tool's output. The people being asked have a stake in the answer, so the case sample and the escalation record check what they say. People who receive the work downstream are worth asking too.

The cost is time from the people who do the work, an analyst to read cases, and a delay before the saving can be booked. For a team of about ten people, this paper's estimate is days of effort over a few weeks, which is a judgment, not a measured figure. The cost grows with the number of roles. If the role cannot be examined this way, because the work is new or no records exist, treat that as the finding: the proposal is not ready for a cut. Run the reversible version instead, described below, and read the signals against a baseline.

One page, attached to the proposal

A redesign gate changes a decision only if it reaches the room where the decision is made. That is usually a planning or cost review under time pressure, not a governance forum. So the sponsor attaches one page to any proposal that cuts headcount because of AI. The page is a decision record, not a living document. It is written once, before approval, and read again on the date set for the first look at the signals. It answers the five questions, shows how the first was answered, including what the affected role carried and how that was found, and names the signals, the baseline, and the date of the first reading.

A missing answer holds the saving: the proposal is not ready, and that is the finding, not a delay. A vague answer goes back to the sponsor for names and routes. A complete answer that shows the role carried little can proceed, and the page records that finding. The page adds no committee and no new stage. Where intake and funding discipline already exist, a project names its benefit, its owner, and its risks before it is funded, and a workforce cut should name what it removes beyond a salary line.

Two situations, walked through the page

In the portfolio and delivery work I have reviewed, AI earned its place where it widened what a team could cover while the team kept the judgment. Neither situation below is a reported case, and neither carries an outcome.

Consider an analysis team under pressure to cut because a model can draft its routine output. Its page would say that the model drafts, and the analysts own the hard cases, the review, and the client relationship. It would name rework and the work handled per analyst as signals, with the level before the cut as the baseline. Under this design, a saving would show up as more work handled per analyst, not as fewer analysts.

Now suppose an organization has already cut an entry-level function and handed the work to an AI tool. A page written afterward would have to answer the pipeline question, and the gap to expect is in supervision: fewer developing staff ready to own the harder work or to check the tool. The repair would be a deliberate development path, funded as an asset the organization needs. In both situations the task may be automatable while the experience around it is not.

The warning signs the budget hides

Forrester expects reversals to be quiet, with jobs returning offshore or at lower wages.[2] A headcount report may not show that. Four signs can show that work is coming back. The first is rework that climbs after a cut. The second is a hard case that arrives with no owner. When a role is removed and nobody takes over its decisions, the hard cases it absorbed can show up somewhere else, such as a more senior and more expensive desk, and arrive late. The third is a thinning senior bench: fewer people ready for the harder work or to supervise the tool. The fourth is quiet

backfill, such as contractors, overtime, an offshore team, or a new tool subscription that reproduces the cut work under a different cost code.

EXHIBIT 3

Four warning signs, read against a baseline

Rework climbs after a cut

Output looks healthy while the share that has to be redone or corrected quietly rises.

Hard cases arrive unowned

The hard cases the role absorbed can show up elsewhere — at a more senior desk, and late.

The senior bench thins

Fewer people ready for the harder work or to supervise the AI. The pipeline effect, early.

Quiet backfill returns

Contractors, overtime, an offshore team, or a subscription reproduce the cut work under another cost code.

Each sign also appears in healthy teams for other reasons. Rework rises with volume or a new product, a bench thins through ordinary turnover, and contractors are hired for projects. No single sign proves a boomerang. Read the four together and against a baseline taken before the cut, and look for several moving past the team's normal swings. A team with no baseline should measure its current rework, hard-case handling, and bench depth before the cut, not after it.

A boomerang can go unpriced because the ledger records the saving precisely and spreads the costs across other budgets and later quarters. A decision that loses money over three years can look profitable in the quarter it is made. So measure the other side of the ledger as seriously as the saving, and judge a cut by the tasks that reappear somewhere else, not by the headcount it removed.

Keep the decision easy to undo

The long-run evidence is still forming, and the forecasts above are forecasts. So keep AI-driven workforce decisions reversible for as long as is reasonable. A cut that can be undone cheaply if the signals turn bad is a different risk from one that destroys knowledge and relationships nobody can rebuild quickly.

Reversibility comes from sequencing. Redesign the workflow and prove the augmentation case before the people leave. Hold a phase where the tool carries more of the task while the people who own the judgment are still there to catch what it gets wrong. Take the headcount decision last, after the signals have been read against the baseline for a period set in advance. The NIST AI Risk Management Framework describes this kind of ongoing management for AI risk in general, and the same idea applies here to the workforce decision.[7] The sequence has a cost: carrying both the tool and the team delays the saving. The proposal should state how long and how much.

Treat every cut as a redesign

Analysts expect many AI-driven cuts to be abandoned or reversed, and they do not say which ones. This paper's answer is that the difference is decided before the cut, when someone either takes the role apart or books the saving from a demo. Back in the cost review, a proposal that arrives with its page answered gives the reviewers something to approve, narrow, or send back before the saving is booked.

So treat every AI-driven workforce decision as a redesign. Name what the person still owns. Choose augmentation or automation on purpose. Say who decides and answers afterward. Protect the pipeline that produces the people who will supervise the tool. Keep the decision reversible. Done that way, AI can reduce work without removing capability that the organization later has to buy back. You can rehire for a quality problem. A lost pipeline can take years to rebuild.

Notes and sources

[1] Gartner, "Gartner Predicts 50% of Organizations Will Abandon Plans to Reduce Customer Service Workforce Due to AI," press release, Stamford, Conn., June 10, 2025. <https://www.gartner.com/en/newsroom/press-releases/2025-06-10-gartner-predicts-50-percent-of-organizations-will-abandon-plans-to-reduce-customer-service-workforce-due-to-ai>

[2] Betsy Summers, "Predictions 2026: The Workforce Muddles Through Ambient Disruption," Forrester blog, November 12, 2025. <https://www.forrester.com/blogs/future-of-work-predictions-2026-whats-coming-for-work-and-the-workforce>

- [3] David Mallon, Brad Kreit, and Natasha Buckley, "Rethinking operating models for humans with agents," Deloitte Insights, April 2, 2026.
<https://www.deloitte.com/us/en/insights/topics/talent/operating-models-for-humans-ai-agents.html>
- [4] Erik Brynjolfsson, Bharat Chandar, and Ruyu Chen, "Canaries in the Coal Mine? Six Facts about the Recent Employment Effects of Artificial Intelligence," Stanford Digital Economy Lab, working paper, November 13, 2025.
https://digitaleconomy.stanford.edu/app/uploads/2025/11/CanariesintheCoalMine_No v25.pdf
- [5] Alexander Sukharevsky, Alexis Krivkovich, Arne Gast, Arsen Storozhev, Dana Maor, Deepak Mahadevan, Lari Hämäläinen, and Sandra Durth, "The agentic organization: Contours of the next paradigm for the AI era," McKinsey & Company, September 26, 2025. <https://www.mckinsey.com/capabilities/people-and-organizational-performance/our-insights/the-agentic-organization-contours-of-the-next-paradigm-for-the-ai-era>
- [6] Lisanne Bainbridge, "Ironies of Automation," Automatica 19(6), 1983, 775–779.
https://ckrybus.com/static/papers/Bainbridge_1983_Automatica.pdf
- [7] National Institute of Standards and Technology, "AI Risk Management Framework Core," excerpt from Artificial Intelligence Risk Management Framework (AI RMF 1.0), January 2023. <https://airc.nist.gov/airmf-resources/airmf/5-sec-core/>

About the author

Marco Policani is an enterprise portfolio, PMO, and AI operating-governance leader. He builds portfolio governance systems where AI carries the analytical load and named humans own every decision — an approach documented across case studies, walkthroughs, and working governance guides on his portfolio site.

Portfolio & governance library: policani.net/governance · Full portfolio: policani.net · LinkedIn: linkedin.com/in/marcpolicani

This white paper may be shared freely with attribution. © 2026 Marco Policani.