


AI Usage Belongs in Workflow Governance

How to govern AI where work happens and scale controls with consequence



Marco Policani

Enterprise Portfolio · PMO · AI Operating Governance

policani.net · linkedin.com/in/marcpolicani

July 2026

EXECUTIVE SUMMARY

A license identifies the user. A usage-control record identifies the work, its limits, and the accountable owner.

Two uses of the same approved AI account can carry very different consequences. The license record cannot show reliance, authority, review, or the ability to contain a mistake.

- 01** Treat each consequential AI-enabled workflow as the governance record.
- 02** Apply the same six controls with greater intensity as consequence, reliance, or autonomy rises.
- 03** Require one short usage-control record before a workflow can influence a decision or take an action.

One AI account can carry two different risks

Your organization has approved an AI tool and wants more people to use it. The access decision looks responsible: an approved vendor, sign-on controls, an acceptable-use policy, and licenses assigned to the right employees. Lower friction can help people learn where the tool is useful. A sanctioned option is also better than forcing useful work into personal accounts that the organization cannot see.

Now follow two uses of the same account. One employee asks the tool to draft an internal meeting note. Another relies on it in a workflow that affects a customer, or allows an agent to trigger a downstream action. The license record looks the same in both cases. It cannot reveal the consequence of an error or how much authority the AI had. It also leaves the review point and stop authority unknown.

That is the problem facing the CIO or AI-governance leader asked to expand access. A blanket policy creates two bad choices. Heavy approval on every low-stakes use can drive people around the sanctioned path. Light controls on every use allow consequential work to proceed without evidence, bounds, escalation, or ownership. The access decision cannot distinguish the work that needs little friction from the work that needs a stronger control. Underneath the dilemma is the oldest principle in risk management, quietly skipped for AI: control intensity should scale with the consequence of the activity, not the tool that performs it. Consequence lives in the workflow, which is why the account — a record of access — is the wrong unit of governance.

The stakes become more concrete when AI can act. McKinsey's agentic-era analysis expands the risk from wrong statements to wrong actions, tool misuse, and operation beyond guardrails.[1] A drafting error can be corrected before use. An unbounded action may reach another system or person before anyone knows it happened.

Build the missing record around consequence

Account approval opens the capability; a separate workflow record must name the work AI may perform, the decision or action it can influence, the evidence it must leave, and the owner who can stop it.

EXHIBIT 1

Account approval opens the capability; a separate workflow record must name the work AI may perform, the decision or action it can influence, the evidence it must leave, and the owner who can stop it.

That separation resolves the apparent choice between open access and blanket restriction. The organization can keep a low-stakes drafting path light while applying stronger controls to a workflow that influences a consequential decision or performs an external action.

NIST supplies the context-of-use frame; the six controls in this paper are Marco's operating translation. The AI Risk Management Framework organizes risk work through GOVERN, MAP, MEASURE, and MANAGE. It treats context of use as material to the risk decision.[2] For each consequential workflow, leaders therefore need to define what AI may do, what evidence it must leave, and which limits contain a mistake. They must also name the point where a human takes over and the person who owns the use. The ISO/IEC 42001 management-system standard sets requirements for establishing, maintaining, and continually improving an AI management system in organizations that provide or use AI [4].

The unit also keeps governance out of individual prompt surveillance. Leaders focus visibility on consequential workflows and usage patterns. That boundary preserves room for experimentation while giving the organization a way to recognize where reliance or autonomy has raised the stakes.

Six controls make a workflow governable

Each consequential AI-enabled workflow needs six controls. Together they make the record usable.

EXHIBIT 2**Six controls for consequential workflows****Scope and limits**

Define permitted work; bound volume, rate, and autonomy.

Visibility and evidence

Monitor actual use; retain enough to reconstruct action and reliance.

Review and ownership

Name review and stop authority; revisit the use as evidence changes.

Monitoring is the practical starting point because it shows whether the controls work and whether reliance has changed. The NIST Generative AI Profile calls for continuous monitoring and structured feedback. It also treats incident response and the human-AI configuration as operating concerns across the lifecycle.[3] Without visibility, a leader cannot tell whether a limit works or whether people now rely on an output in a new way. Yesterday's drafting aid may already feed a decision.

This monitoring belongs at the workflow level. A customer-impacting use might track AI-triggered actions, human overrides, boundary breaches, and escalations. Those signals expose gaps in the sanctioned path and show when reliance or autonomy requires stronger control.

Use stronger controls when consequences rise

The six controls stay constant; their intensity changes with the work.

An internal drafting use may need an approved tool, light monitoring, and clear boundaries on the information entered. The person remains responsible for the note, and the output can be corrected before it affects anyone else. Requiring a formal approval for every draft would add friction without a proportionate benefit.

A workflow that feeds a customer-impacting decision needs more. Access should be scoped to the approved use. The organization needs evidence of what informed the decision, a human review point, a named owner, and an escalation path for uncertainty. If the AI can take an action, limits on rate, volume, and permitted actions contain a bad configuration before it travels.

For usage limits, the operating question is how much the AI can do before a person reviews the case or a ceiling stops the run.

Shadow use belongs in the same analysis. When the sanctioned path is so restrictive that it cannot meet a real workflow need, people may choose tools outside it. The approved route has to meet the workflow need and remain easy to find. Monitoring then shows where unmet demand or new reliance requires a change. A consequential use may still need to be blocked; leaders can also correct the sanctioned path that people were routing around. This is not a fringe case: in a November 2025 survey of 2,000 U.S. and U.K. employees at companies with more than 500 staff, commissioned by the security vendor BlackFog, 49% reported using AI tools their employer had not sanctioned, and 58% of those users relied on free versions [5].

Start with one usage-control record

Before expanding AI into a consequential workflow, require one short usage-control record. It should name the workflow and owner, then answer six questions:

- What work may AI perform, and what access is in scope?
- What usage and outcome signals will be monitored?
- What limits bound volume, rate, or autonomy?
- What evidence and logs must the workflow retain?
- When does a human review or escalation occur, and who can pause or stop the use?
- When will the usage pattern be reviewed again?

This record is the operating move. It avoids building a heavy approval process around every user. Low-stakes work can be documented lightly. Consequential work receives the fuller control set before reliance or autonomy increases.

The record also makes governance continuous. New workflows appear, existing uses scale, and a low-stakes drafting aid may become an input to a decision. NIST treats risk management as work that continues across the lifecycle.[2] At each review, the owner uses observed evidence to change or end the permitted use.

Keep the claim bounded

This paper proposes an operating standard. No controlled comparison establishes that workflow governance causes faster adoption or better outcomes than broad access. The McKinsey and NIST sources support context, monitoring, guardrails, and human responsibility. The six-control record remains Marco's proposal.

No universal threshold separates low-stakes from high-stakes work. Organizations must judge the data involved, the consequence of error, the degree of reliance or autonomy, and the ability to recover.

Excessive process can recreate the friction that broad access was meant to remove. A weak record leaves consequential work exposed. Proportionality requires judgment, evidence, and a named owner.

The record should also earn its place. If it never changes a control decision, simplify it. The purpose is a usable decision at the workflow boundary.

Make the next access decision differently

Return to the two employees using the same approved tool. The internal drafting use can keep a light path because a person can inspect and correct the output before it matters. The customer-impacting workflow gets the full usage-control record because the consequence is higher.

The organization still expands AI access. It now knows where that access enters real work and what the AI may do there. The control record also shows how a mistake is contained and who decides when the use changes.

For the next customer-impacting workflow, require the record before reliance or autonomy increases. Let the named owner decide whether the evidence supports a light path, stronger controls, or stopping the use.

Notes and sources

[1] Gabriel Morgan Asaftei, Roger Roberts, Abby Sticha, and Cécile Prinsen, "State of AI trust in 2026: Shifting to the agentic era," McKinsey & Company, March 25, 2026. <https://www.mckinsey.com/capabilities/tech-and-ai/our-insights/tech-forward/state-of-ai-trust-in-2026-shifting-to-the-agentic-era>

[2] National Institute of Standards and Technology, "AI Risk Management Framework Core," AI RMF 1.0, 2023. <https://airc.nist.gov/airmf-resources/airmf/5-sec-core/>

[3] Chloe Autio, Reva Schwartz, Jesse Dunietz, Shomik Jain, Martin Stanley, Elham Tabassi, Patrick Hall, and Kamie Roberts, "Artificial Intelligence Risk Management Framework: Generative Artificial Intelligence Profile," NIST AI 600-1, July 26, 2024. <https://doi.org/10.6028/NIST.AI.600-1>

[4] International Organization for Standardization and International Electrotechnical Commission, "ISO/IEC 42001:2023 — Information technology — Artificial intelligence — Management system," 2023. <https://www.iso.org/standard/42001>

[5] Taryn Plumb, "Roughly half of employees are using unsanctioned AI tools, and enterprise leaders are major culprits," CIO, January 29, 2026. <https://www.cio.com/article/4124760/roughly-half-of-employees-are-using-unsanctioned-ai-tools-and-enterprise-leaders-are-major-culprits.html>

About the author

Marco Policani is an enterprise portfolio, PMO, and AI operating-governance leader. He builds portfolio governance systems where AI carries the analytical load and named humans own every decision — an approach documented across case studies, walkthroughs, and working governance guides on his portfolio site.

Portfolio & governance library: policani.net/governance · Full portfolio: policani.net · LinkedIn: linkedin.com/in/marcpolicani

This white paper may be shared freely with attribution. © 2026 Marco Policani.