

The AI Value Test: Automate, Build, Buy, Hire, or Wait

A portfolio decision model for funding AI work only when the operating case is strong enough



Marco Policani

Enterprise Portfolio · PMO · AI Operating Governance

policani.net · linkedin.com/in/marcpolicani

July 2026

EXECUTIVE SUMMARY

Choose AI only when it moves the constraint better than automating, buying, hiring, or waiting.

Proving that AI can complete a task does not prove it is the best choice. Compare five answers before funding a pilot: automate, build, buy, hire, or wait.

- 01** Each option uses money, capacity, attention, data, and tolerance for risk in a different way. Judge all five options against the same eight questions: value, readiness, data, risk, cost, capacity, reversibility, and strategic control.
- 02** Each option also creates work after approval. Count maintenance, product ownership, vendor oversight, training, and review dates before funding it.
- 03** End each review with a clear decision: explore, pilot, scale, switch option, wait, or stop. Tie that decision to evidence.

CONTENTS

The funding question comes first	A worked example
Why automation is only one answer	Failure patterns the value test prevents
The five options	The evidence pack for the value test
The five options at a glance	Governance after the choice
Decision criteria that keep the test honest	What the review should decide
A funding-gate model	What the research does not settle
How the answers differ in practice	Make the choice before the tool does
How to compare the five choices	Notes and sources
What each answer demands after approval	

The funding question comes first

The AI value test is not whether AI can do the work. It is whether the portfolio should automate, build, buy, hire, or wait.

The question belongs in portfolio governance. Automation, building, buying, hiring, and waiting each place a different demand on money, capacity, accountability, risk tolerance, data, and time. A narrow automation frame can make the first demo look successful while hiding the operating cost of the choice.

Recent research weakens that narrow case. McKinsey found that AI use was widespread, but enterprise-level impact remained uneven.[1] Its 2026 framework says companies should measure AI as they would any capital investment. The chain runs from technical performance through adoption and operating measures to strategic and financial results.[5] In 2026, Gartner reported that at least half of GenAI projects had been abandoned after proof of concept by the end of 2025. It cited weak business value, poor data, rising total cost, gaps in responsible AI, and weak change management.[3] In 2025, Gartner also predicted that more than 40 percent of agentic AI projects would be canceled by the end of 2027. It cited rising cost, unclear value, and weak controls.[6]

The portfolio should replace one question with another. Do not ask only, "Where can we use AI?" Ask, "Which option creates the best value for this constraint, given our evidence and the capacity we can sustain?"

Waiting is not failure, buying not weakness, hiring not old-fashioned, building not prestige. Each is a legitimate answer when the constraint, value mechanism, risk, readiness, and cost point in that direction.

Why automation is only one answer

Automation is attractive because it promises a visible reduction in effort. The danger is that effort is often the easiest part of the work to see, not the most important part of the value system.

The right answer varies by work. Some work should be automated, bought, built, left to people, or deferred. Some should use AI to support people rather than replace them, because human judgment carries the value. Other work needs redesign before any AI is added.

The wrong choice can raise cost even if it appears to reduce labor. A custom agent may need integrations, monitoring, evaluation, human review, security controls, and support. A simple workflow change may avoid those costs. A vendor tool may speed delivery but add dependency, license cost, and data exposure the business did not price. Hiring or training may look slower but build lasting internal skill. Waiting can protect the portfolio from funding a workflow that has not earned trust.

McKinsey's 2026 article on where AI will create value argues that productivity gains become table stakes.[7] It says deeper value comes from changes to offerings, business models, and market structures. That does not make productivity unimportant. Leaders still need to know which economic game they are playing. Work-reduction automation, a strategic capability, and a new business model need different funding logic.

The five options

Automate when the workflow is stable and the rules are known. The data must be reliable, and the team must be able to manage exceptions. The value case should depend on speed, throughput, consistency, or lower unit cost. The solution may use AI, fixed rules, scripts, RPA, workflow tools, or a mix. Fund the simplest automation that works reliably.

Build when the capability matters to strategy, depends on proprietary context, or could become a repeatable advantage. A build creates long-term obligations. The organization must own the architecture, data pipelines, evaluation, model routing, monitoring, support, change, and future funding. Build only when control and learning are worth those obligations.

Buy when the need is common, a mature product exists, and speed matters more than custom technology. Buying still needs governance. Vendor claims do not replace proof. Procurement should review data rights, integration work, model or agent behavior, security, cost at scale, portability, and the work needed to make the tool useful.

Hire or train when the constraint is judgment, domain knowledge, relationships, adoption, change leadership, or operating capacity. AI can extend the reach of experts. It does not replace people who know what good looks like, can challenge outputs, and own decisions. If the organization underinvests in people, AI may produce work faster than anyone can assess it.

Wait when the business case is weak or the workflow is unstable. Waiting may also fit when data is not ready, the risk boundary is unclear, the vendor market is immature, or the organization cannot absorb the change. A wait should be an explicit choice with review triggers, not a vague delay. It should name what must become true before the idea returns. Real-options analysis gives this choice a specific value: the option to act later. That option matters when uncertainty is high and the commitment is hard to reverse. Acting too soon spends that option.

The five options at a glance

The five options are funding choices, and they compare most clearly side by side.

EXHIBIT 1

Five funding options, five different conditions

- 1 Automate**
The workflow is stable, decision rules are known, data is reliable, and value comes from speed, throughput, consistency, or unit cost. Fund the simplest reliable automation.
- 2 Build**
The capability matters to strategy, depends on proprietary context, or could become a repeatable advantage. Build only when control and learning are worth the long-term work.
- 3 Buy**
The need is common, a mature product exists, and speed matters. Vendor claims do not replace proof.
- 4 Hire or train**
The constraint is judgment, domain knowledge, adoption, change leadership, or capacity. AI can extend an expert's reach. It does not replace knowing what good looks like.
- 5 Wait**
The case is weak, the workflow is unstable, or the data is not ready. Name what must change before the idea returns.

Option	When it fits	What it demands after approval	Dominant failure mode
Automate	Workflow is stable, rules are known, data is reliable, and value comes from speed, throughput, or consistency	Monitoring for drift, exceptions, cost, and changing rules	Automation theater: effort falls, no outcome moves
Build	Capability is differentiating or tightly bound to proprietary context	A roadmap, evaluation method, support model, and future funding you now own	Prestige building: a custom asset where a feature would do
Buy	Need is common, mature capability exists, and time-to-value matters	Configuration, integration, adoption, data-rights, and exit management	Procurement substitution: the tool arrives, the constraint stays
Hire or train	Constraint is judgment, relationship, adoption, or capacity	Role clarity, learning paths, and evidence the human system is changing the work	Headcount avoidance: AI covers a gap only people can close
Wait	Case is weak, workflow unstable, or data and controls are not ready	A named trigger and review date, not an open-ended delay	Passive waiting: deferral with no condition to bring it back

Decision criteria that keep the test honest

The five options become useful only when the portfolio evaluates them against common criteria.

Start with value. Is the constraint tied to cost, revenue, customer experience, risk, regulatory exposure, employee capacity, decision quality, or a future advantage? Do not use a complex answer for a low-value constraint.

Then test whether the workflow is ready. Is the current process stable enough to improve? Are its handoffs, owners, exceptions, and standards for a good result clear? If the workflow is messy, AI may spread the ambiguity rather than remove it.

Then test the data and context. Does the system have the information it needs to produce reliable output? Is the data current, controlled, available, and appropriate to use? RAND places problem understanding and data readiness near the center of project risk. Gartner lists poor data and unclear business value among its GenAI failure points.[2][3]

Risk and trust determine how much reliance is acceptable. NIST's AI RMF organizes risk management around four functions: govern, map, measure, and manage.[4] That matters because reliance is not a yes-or-no switch. A tool might assist with a draft, propose a classification, prepare evidence, or execute a step. Each level needs different controls. A 2026 McKinsey survey of about 500 organizations found that security and risk concerns were the top barrier to scaling agentic AI.[8]

Count the full cost. Include licensing, token or inference use, integration, security review, evaluation, monitoring, human review, training, support, vendor management, and maintenance. Gartner's warning on agentic AI links cancellation risk to rising cost, unclear value, and weak controls.[6]

Finally, test reversibility. If the choice is wrong, how hard will it be to stop, unwind, switch vendors, rebuild trust, or recover the workflow? Hard-to-reverse choices need stronger proof before scale.

Data needs extra scrutiny because it can fail quietly. A demo may work on a clean sample and then fail on the uneven data used in production. That data may have limited permissions or unclear ownership. Ask whether the data for this workflow is fit for this purpose, not whether clean data exists somewhere else. When it is not, fund the data work before the model.

A funding-gate model

The AI value test should run before a project receives serious funding, then repeat at each gate. A good first answer can become a bad scale answer when evidence changes.

At idea intake, the portfolio asks for the constraint, option set, value hypothesis, owner, and evidence plan. It does not need a polished deck. It needs enough clarity to decide whether the idea deserves discovery.

At discovery, the team compares automate, build, buy, hire, and wait. The output is a recommended path and the rejected alternatives. Rejected alternatives matter because they show that the team considered cheaper, safer, or more durable options.

At pilot, test technical performance, adoption, and early workflow results. McKinsey's five-layer framework helps because it stops teams from jumping from model quality to financial claims.[5] A pilot should show three things: the system works, people can use it in real work, and a useful operating result is starting to appear.

At scale, review the full operating case. Include financial impact, total cost, controls, support, change readiness, and ownership. A project is not ready for daily use if no one can explain who will maintain evidence, handle exceptions, and review drift.

At run, the decision remains alive. The organization should continue asking whether the original answer is still right. A bought capability may need to be replaced by a build. A custom build may become a

vendor commodity. An automation may expose a need for training. A wait decision may become fundable after the data foundation improves.

How the answers differ in practice

Consider a recurring executive reporting workflow. If the constraint is that leaders receive too much unstructured status language, automation alone may generate prettier noise. The better move may be to redesign the intake fields, define decision categories, and use AI only to prepare exceptions for human review.

Consider a service operation with high first-contact resolution pressure. If the knowledge base is current and the resolution paths are stable, buying or configuring an AI-assisted tool may create value quickly. If the knowledge base is stale, the first investment may be content governance and ownership. If the real constraint is expert escalation capacity, hiring or training may carry more value than another model layer.

Consider a product team exploring agentic workflows. If the use case is strategically differentiating and depends on proprietary context, a build path may make sense. If the use case is a common horizontal task, buying may be faster and safer. If the agent would act without a clear review boundary, waiting is governance.

How to compare the five choices

The test works when each option faces the same eight questions: value, readiness, data, risk, cost, capacity, reversibility, and strategic control.

The portfolio already used six of these at the criteria stage: value, readiness, data, risk, cost, and reversibility. The comparison now applies them to every option, not only the preferred one.

Two questions are new. Capacity asks whether the organization has enough people and attention to carry out the choice without starving higher-value work. Strategic control asks whether the organization needs to own the capability, data, workflow logic, or learning loop.

The questions do not carry equal weight in every decision. For a low-stakes choice that is easy to reverse, readiness and cost may dominate. A hard-to-reverse or high-exposure choice needs more proof. Examples include a multi-year build, a deep vendor dependency, or a workflow that affects regulated decisions. In those cases, reversibility and strategic control move to the front. Use the eight questions as a checklist, not an average score. One failed answer can outweigh five comfortable answers when the choice is hard to undo.

What each answer demands after approval

Automation demands maintenance. Someone must monitor whether the rules, model behavior, integrations, exceptions, and cost still match the value case. A workflow that changes every month may turn a promising automation into an expensive upkeep obligation.

A build decision demands product discipline. The organization owns the roadmap, evaluation method, support model, data dependency, security posture, and future funding. Build decisions work when the capability is important enough to justify ownership.

A buy decision demands adoption discipline. Vendor capability does not equal business value. The organization still has to configure the tool, fit it into the work, train users, review outputs, manage data exposure, and track the operating result.

Hiring or training needs a clear view of the judgment required. If the constraint is expertise, trust, relationships, or decision ownership, people are the capability. They are not a cost left over after automation. AI may extend their reach, but the portfolio still needs to fund the human system.

A wait decision demands a trigger. Waiting is governable only when the organization names what must change: data readiness, vendor maturity, risk tolerance, workflow stability, budget capacity, or strategic urgency. Without a trigger, waiting becomes drift.

A worked example

The following is an illustration, not a client account. It runs one constraint through the five options.

A proposal team is missing deadlines because first drafts take too long to assemble. The obvious request is to buy an AI drafting tool. The value test asks a wider question before funding anything.

Automate: drafting is not stable, rule-based work; it depends on judgment about what each opportunity needs. Full automation would produce fast drafts the team rewrites anyway.

Build: nothing in generic proposal drafting is differentiating enough to justify owning a custom system, its evaluations, and its upkeep.

Buy: a mature AI-assisted drafting tool could help, but only if the underlying content — past proposals, approved language, current pricing — is organized and trustworthy. On a disorganized content base, the tool accelerates the wrong starting point.

Hire or train: if the team lacks people who can shape a winning argument, a tool will not fix the bottleneck. The proposal team needs more capacity or skill.

Wait: if the content base is stale, fund content governance first. Delay the purchase until the library reaches a defined level of coverage and freshness.

The strongest first move is often not the tool in the request. It may be the content and staffing work that makes a later tool dependable. The value test does not reject AI. It puts AI behind the conditions that let it work. The decision record might say, "govern the proposal library now, protect senior proposal capacity, and reassess the purchase at the coverage threshold." That is more useful than "buy a drafting assistant."

Failure patterns the value test prevents

The first pattern is automation theater. The team automates a visible task because it can, then struggles to show that a customer, process, or financial outcome improved.

Prestige building is the second pattern. A team funds a custom solution because it looks strategic, even though a mature vendor feature or workflow change could solve the constraint faster.

The third pattern is procurement substitution. A tool is bought to avoid a harder operating decision about ownership, process design, data quality, or human review. The software arrives; the constraint remains.

The fourth pattern is headcount avoidance. AI is asked to cover a judgment, relationship, or change-management gap that actually requires skilled people. The organization saves a role on paper and pays for it later in quality, trust, or adoption.

The fifth pattern is passive waiting. Leaders defer a decision without naming what evidence would bring it back. A governed wait is active. It protects scarce capacity while preserving the option to move when readiness improves.

The evidence pack for the value test

The evidence pack should show how the portfolio chose among automate, build, buy, hire, and wait. It should defend the chosen answer and show the rejected answers clearly enough for leaders to see the tradeoff.

The first page names the constraint, the value mechanism, and the current baseline. The second page compares the five options across value, readiness, data, risk, total cost, capacity, reversibility, and strategic control. The third page defines the proof needed for the next gate. The fourth page records the decision and the trigger for review.

This structure prevents a common AI funding mistake: treating one plausible option as the best one. A build can make sense and still lose to a vendor feature. A vendor feature can be fast and still lose to data cleanup. An automation can cut effort and still be wrong if the workflow needed redesign first.

The evidence pack also helps finance and operating leaders participate without becoming model experts. They can challenge the value mechanism, the total cost, the capacity assumption, the proof plan, and the reversibility of the choice. That is where many AI decisions become clearer.

Governance after the choice

The value test does not end when the portfolio chooses an option. Each choice creates different work after approval.

Automated work needs monitoring for drift, exceptions, cost, and quality. A build needs backlog ownership, support funding, evaluation, and lifecycle management. A purchase needs vendor oversight, adoption review, data-rights monitoring, and an exit path. Hiring or training needs clear roles, learning paths, and proof that the new skill is changing the work. Waiting needs named triggers and a review date.

This work after approval is easy to miss. Leaders can fund a pilot and still underfund what makes it last. The result is familiar: an impressive proof of concept, a fragile path to daily use, unclear ownership, and a quiet return to the old work.

A portfolio decision is only complete when the organization knows what it will have to own after approval. That is what turns the choice from a technology preference into an operating commitment.

What the review should decide

A value-test review should end with one of six decisions: explore, pilot, scale, switch option, wait, or stop. Each decision should be tied to evidence, not to enthusiasm or fatigue.

EXHIBIT 2**Six review decisions replace yes or no****Explore**

The constraint may be valuable, but the team has not earned a pilot; the next work is discovery.

Pilot

Enough clarity to test a defined option with limited exposure and explicit proof.

Scale

The proof is strong enough to expand reliance and operating commitment.

Switch option

The original path was wrong, but the constraint still matters.

Wait

The idea returns when a named condition changes.

Stop

The value case no longer deserves attention.

Explore means the constraint may be valuable, but the team has not earned a pilot. The next work is discovery: clearer baseline, user research, data review, vendor scan, risk boundary, or intervention comparison. Pilot means the team has enough clarity to test a defined option with limited exposure and explicit proof.

Scale means the proof is strong enough to expand reliance and operating commitment. Switch option means the team learned that the original path was wrong but the constraint still matters. Wait means the idea should return when a named condition changes. Stop means the value case no longer deserves attention.

Most AI ideas should not be killed or funded as if the answer were already known. These six decisions keep an option open without letting the portfolio commit too soon.

What the research does not settle

The research does not offer one formula for when to automate, build, buy, hire, or wait. It shows where AI programs stumble: unclear value, weak problem framing, poor data, rising cost, adoption gaps, and weak controls.

This decision model is a governance standard, not a statistical prediction. Adapt it to the organization's risk, rules, architecture, workforce, and business strategy.

Timing is another limit. Vendors, models, prices, and enterprise controls change quickly. That change strengthens the case for the value test. Revisit the decision as the market changes instead of locking a 2026 assumption into a multi-year commitment.

Make the choice before the tool does

AI gives organizations more possible moves. It does not make the choice for them.

Portfolio leaders should not ask every team to automate. They should ask which constraint matters, which option fits, what evidence earns reliance, and whether the full operating case is stronger than the alternatives.

Automate, build, buy, hire, or wait. The value is in making the choice before the tool makes it for you.

Notes and sources

1. Alex Singla, Alexander Sukharevsky, Bryce Hall, Lareina Yee, Michael Chui, and Tara Balakrishnan, "The state of AI in 2025: Agents, innovation, and transformation," McKinsey & Company, November 5, 2025. Verified July 7, 2026. <https://www.mckinsey.com/capabilities/quantumblack/our-insights/the-state-of-ai-2025>
2. James Ryseff, Brandon F. De Bruhl, and Sydne J. Newberry, "The Root Causes of Failure for Artificial Intelligence Projects and How They Can Succeed: Avoiding the Anti-Patterns of AI," RAND Corporation, 2024. Verified July 7, 2026. https://www.rand.org/pubs/research_reports/RRA2680-1.html
3. Arun Chandrasekaran, "Why 50% of GenAI Projects Fail - And How to Beat the Odds," Gartner, January 26, 2026. Verified July 7, 2026. <https://www.gartner.com/en/articles/genai-project-failure>
4. NIST, "AI Risk Management Framework Core," excerpt from AI RMF 1.0, 2023. Verified July 7, 2026. <https://airc.nist.gov/airmf-resources/airmf/5-sec-core/>
5. Johannes-Tobias Lorenz, Joshan Cherian Abraham, Robert Levin, and Douglas Ziman, "From promise to impact: How companies can measure and realize the full value of AI," McKinsey & Company, April 24, 2026. Verified July 7, 2026. <https://www.mckinsey.com/capabilities/quantumblack/our-insights/from-promise-to-impact-how-companies-can-measure-and-realize-the-full-value-of-ai>
6. Gartner, "Gartner Predicts Over 40% of Agentic AI Projects Will Be Canceled by End of 2027," June 25, 2025. Verified July 7, 2026. <https://www.gartner.com/en/newsroom/press-releases/2025-06-25-gartner-predicts-over-40-percent-of-agentic-ai-projects-will-be-canceled-by-end-of-2027>

7. Antoine Montard, Dago Diedrich, and Tanguy Catlin, "Where AI will create value - and where it will not," McKinsey Quarterly, April 29, 2026. Verified July 7, 2026. <https://www.mckinsey.com/capabilities/strategy-and-corporate-finance/our-insights/where-ai-will-create-value-and-where-it-wont>
8. Gabriel Morgan Asaftei, Roger Roberts, Abby Sticha, and Cécile Prinsen, "State of AI trust in 2026: Shifting to the agentic era," McKinsey & Company, 2026. Verified July 7, 2026. <https://www.mckinsey.com/capabilities/tech-and-ai/our-insights/tech-forward/state-of-ai-trust-in-2026-shifting-to-the-agentic-era>

About the author

Marco Policani is an enterprise portfolio, PMO, and AI operating-governance leader. He builds portfolio governance systems where AI carries the analytical load and named humans own every decision — an approach documented across case studies, walkthroughs, and working governance guides on his portfolio site.

Portfolio & governance library: policani.net/governance · Full portfolio: policani.net · LinkedIn: linkedin.com/in/marcpolicani

This white paper may be shared freely with attribution. © 2026 Marco Policani.