
WHITE PAPER · PORTFOLIO & AI GOVERNANCE

Benefits Realization Is Where Transformations Quietly Fail

Why programs that ship on time and on budget still fail to deliver value — and how to govern the end of the lifecycle, not just the middle

Marco Policani

Enterprise Portfolio · PMO · AI Operating Governance

policani.net · [linkedin.com/in/marcpolicani](https://www.linkedin.com/in/marcpolicani)

July 2026

EXECUTIVE SUMMARY

A transformation has not created value until the benefit is measured, owned, and acted on after delivery.

Most transformation failures are not delivery failures. The program ships, the milestones are checked, and the value that justified the effort is never measured. This paper defines benefits realization as a governed discipline and shows how to run the end of the lifecycle as deliberately as the middle.

- 01** Delivery answers “did we build it?” Benefits realization answers “did it produce the value we funded it for?” — a separate job most organizations skip.
- 02** Six ordinary steps close the gap: a named benefits owner, a baseline taken before the start, targets set on the operating or financial result rather than on output or adoption, a review booked for after delivery, a real decision on what the review finds, and a loop back into the next business case.
- 03** The structural fix is one added review after go-live, called a realization gate: booked at approval, owned by a named leader, and asking whether the benefit appeared and what happens next.

CONTENTS

The program shipped and the value never showed up
 Why no one notices
 What benefits realization means
 Benefits get over-promised, then never checked
 Measure the outcome, not the activity
 The steps that close the gap
 Questions a delivery report cannot answer
 Two examples: one kept the value, one lost it

AI programs are the most exposed
 Two things this paper does not claim
 Add one gate after go-live
 How this closes the funding loop
 Who owns the value once the team is gone
 Start with your biggest program
 The real obstacle is that nobody wants to know
 Govern the value, not just the delivery
 Notes and sources

The program shipped and the value never showed up

Most transformation failures are not dramatic. The program does not collapse. It does not miss its deadline or blow its budget. It ships. The new solution goes live, the milestones are checked, the team is thanked, and everyone moves on. Months later the value that justified the whole effort has not appeared. No one is quite sure why. No one is sure anyone is still measuring. This is a common way transformations fail. It happens at the end of the lifecycle, where the value is supposed to be captured and often is not. The shortfall can be large. McKinsey and the University of Oxford studied more than 5,400 large IT projects. On average those projects delivered 56 percent less value than predicted[4]. BCG studied 895 digital transformations. The ratings came largely from the executives who led them. Only about 30 percent met their target value and produced change that lasted[5].

Benefits realization is a separate job. It means capturing and checking the value a transformation was funded to produce. Most organizations skip it. That is why so many well-delivered programs show so little benefit anyone can measure. The question that decides success is whether the program produced the value it was funded for. An organization that governs delivery hard while ignoring realization is measuring the wrong end of its own programs.

This is well established in the delivery profession, even where it is poorly practiced. The Project Management Institute publishes benefits realization management as a practice guide within its standards. It describes the discipline as the thread running from strategy through what a project delivers. PMI's own survey found only one organization in three reporting high maturity at it.[1] The gap between knowing and doing is where the value leaks out. This paper is about closing it. What is benefits realization, why does it keep getting skipped, and how do you govern it as deliberately as delivery?

Why no one notices

This failure stays hidden because of what organizations choose to measure. They measure delivery closely: schedule, budget, scope, milestones, defects, go-live.

They measure benefits rarely, if at all. Measuring benefits takes work after the exciting part is over. It usually falls to a team that has already been broken up and reassigned. So a program can deliver beautifully, produce no benefit, and still look like a success on every dashboard the organization keeps.

The timing makes it worse. Teams treat go-live as the finish line. For value it is the starting line. Benefits build after launch, as the changed way of working beds in. The intended outcomes then either appear or do not. Go-live is also when attention moves on and budgets close. The people who understood the value case scatter to the next project. The organization declares victory at the exact moment it stops watching whether it won.

There is also comfort in not looking. Check the benefits and you may find that a large, celebrated program did not deliver what it promised. That is uncomfortable for everyone who sponsored it, approved it, or led it. Skipping the measurement protects the story of success. It also leaves the substance unchecked. This is the same avoidance that keeps weak projects alive in a portfolio, moved to the back end of the lifecycle. Measuring would force a reckoning, so the organization declines to measure.

What benefits realization means

You can define benefits realization precisely enough to govern it:

Benefits realization means four things done on purpose. Name the specific value the transformation is expected to produce. Record the baseline and the measure before the work starts. Put one named person on the hook for the value, not merely for the delivery. Then, after go-live, check whether the value actually appeared, and make a real decision about what you find. A program with no benefits owner, no baseline, and no check after delivery has a delivery plan and nothing more. That holds whatever its business case claimed about value at approval.

That definition splits two roles organizations routinely merge. The delivery owner answers for building the thing. The benefits owner answers for the value the thing was supposed to produce. The two jobs run on different clocks. The delivery owner's job ends at go-live. The benefits owner's job starts there. Merging them into one role that ends at delivery is a primary reason benefits go unrealized. Someone has to stay accountable after the delivery team moves on. Otherwise the value has no owner at exactly the time it is supposed to appear.

Benefits get over-promised, then never checked

Two failures sit at opposite ends of the lifecycle, and they feed each other. At the front, benefits are systematically over-promised, because optimism is what gets a program funded. Flyvbjerg and Gardner studied large projects. They find that estimates built from the inside view are prone to optimism bias. The inside view is the view of the people close to the work and invested in it. A more realistic number comes from the outside view. Compare the project to what actually happened on a reference class of similar past efforts, rather than to its own best-case story.[2] Benefits promised from the inside view are inflated. That inflation is baked into the business case that justifies the spend.

At the back there is silence. No one goes back to check whether those benefits appeared. The optimism is never corrected and never learned from. The organization keeps funding programs on optimistic value cases it never checks, and the next business case runs as high as the last. The outside view fixes the front end. It grounds estimates in what similar projects actually delivered. Verifying benefits fixes the back end. It feeds real results into the next estimate. Do neither, and the cycle repeats.

Measure the outcome, not the activity

When organizations do try to measure benefits, they often measure the wrong level. They stop at activity or output and never reach the outcome that was the point. McKinsey's model for realizing value helps here, because it keeps the levels apart: technical performance, user adoption, operational KPIs, strategic outcomes, and financial impact. A program can pass the early levels and fail the ones that matter.[3] A solution that works technically and is widely adopted has cleared the first two levels. It may still have moved no operational or financial outcome at all.

So decide, before the program starts, which layer the benefit lives in. Then measure at that layer. Reporting that the solution is live and usage is high is easy. Both are visible at go-live and both feel like success. Measuring whether the operational outcome improved and the money appeared is harder. It is also the part that counts. That is what the program was funded for, and it shows up later, if at all. McKinsey also advises building the value case up front, with evidence and staged decisions. That is the same discipline seen from the funding side. Pick the benefit and the layer, then hold the program to it through realization rather than only to go-live.[3]

The steps that close the gap

Governing benefits realization comes down to a short list. Every item is ordinary. Every item gets routinely skipped once the delivery excitement takes over.

EXHIBIT 1

The benefit ownership level

Named benefits owner

Owens the value after go-live. A different person from the delivery owner.

Baseline before you start

Record where the measure starts, before any of the work begins.

Target at the right level

Aim at the operating or financial outcome, not at output or adoption.

Book the review up front

Agree at approval to review once the benefit has had time to appear.

Decide, then learn

Adjust, invest, or stop. Feed the verified result into the next business case.

None of this is difficult. All of it needs someone to still care about the program after the ribbon is cut. The hard part is attention rather than mechanics. You have to hold a named owner to account. You have to keep a review on the calendar for a date when the organization's focus has moved on.

Questions a delivery report cannot answer

You can tell a program governed for delivery from one governed for benefits by the questions it can answer. The table pairs each delivery signal with the benefits question it leaves open.

The delivery signal	What it confirms	The benefits question it leaves open
Shipped on time and on budget	The program was well delivered	Did the value it was funded for appear
The system is live	The output exists	Did the outcome the output was meant to change move
Adoption is high	People are using it	Is the operational or financial benefit materializing
The project is closed	Delivery is complete	Who is still accountable for the value after go-live
The business case was approved	The benefit was promised	Was the promise checked against a reference class, and verified later

The right column is where transformations quietly fail. Every question in it can be answered. Most organizations answer none of them, because they governed the middle of the lifecycle and walked away from the end. Benefits realization is the work of answering that column on purpose.

Two examples: one kept the value, one lost it

The safest illustrations are typical cases rather than named programs. In portfolio and program work, the lesson is the same. The organizations that actually capture value governed the back end of the lifecycle as seriously as the front, and kept an owner and a measure alive past go-live.

In the first case, a program was set up so the value case did not close at delivery. A benefits owner, separate from the delivery lead, carried a baseline and a target through go-live. The review was booked at approval for months after launch, and it happened. It had the authority to fund further work or to stop, depending on whether the outcome moved. The value was captured because someone was still accountable for it when it was due.

In the second case, an organization found that a run of celebrated deliveries had produced little benefit anyone could measure. No one had gone back to check. More

delivery rigor was not the answer; delivery was already strong. What was missing was the back end. Baselines had to be captured before programs began. Benefits owners had to be named and kept in place. Realization reviews had to feed verified results into the next round of business cases. The delivery machine was fine. The value machine had never been built, and building it was what turned motion into outcomes.

AI programs are the most exposed

AI transformations need this discipline most. They are unusually prone to the exact failure this paper describes. AI programs throw off plenty of activity — usage, pilots, deployments — and activity is easy to mistake for value. The enthusiasm around AI makes the front-end optimism especially strong and the back-end verification especially rare. The result is a wave of AI initiatives that ship impressively and deliver benefits no one has actually measured.

The measurement layers matter even more here. An AI program can clear the early levels, technical performance and adoption. It can still move none of the operational or financial outcomes that justify it.[3] An AI deployment that works and is widely used looks like a success. It may have changed no business outcome at all. Name a benefits owner. Take the baseline. Measure at the right level. Verify after go-live. That is what separates the AI transformations that produce value from the far larger number that produce activity. None of this is new. AI has made the cost of skipping it larger and more common.

Two things this paper does not claim

Two limits are worth stating. First, this paper argues for the discipline of benefits realization, not for a single quantified failure rate. Published statistics on how much transformation value goes unrealized vary widely by source, sector, and definition. So measure your own realization gap rather than import a headline number. The profession treats benefits realization as a core discipline, and PMI devotes a practice guide to it. How much is lost by skipping it differs from one organization to the next. [1] Second, benefits realization is necessary without being sufficient on its own. Measuring and owning benefits does not guarantee a program was the right bet in the first place. It does let the organization learn the truth about its bets and act on it. Across a portfolio of many transformations, that is worth more than confidence about

any single one. Benefits realization governs how value is verified. It does not settle whether the underlying strategy was sound.

Add one gate after go-live

The single most effective structural fix for this failure is one gate most delivery lifecycles are missing: a realization gate after go-live. Lifecycles are rich with gates up to launch — approval, design, build, readiness. Then they simply stop, as if the lifecycle ended when the solution went live. The realization gate reopens the program at a scheduled point after launch, once the value has had time to appear. It asks the question that ultimately matters: did the benefit appear, and if not, what will we do about it?

The gate works because it is booked at approval, when the program is eager to commit. Leave it to be arranged after go-live and no one wants to. Commit at the outset to a specific date for checking the benefit. Commit to a named owner who has to show up with the evidence. That is what keeps the back end from being abandoned. The gate carries a real decision: accept the realized benefit as sufficient, invest further to close a gap, or stop pouring effort into an outcome that is not coming. Without that decision the review is theater. With it, the gate is where value is actually protected.

The realization gate also balances the back of the lifecycle against the front. Gate funding on a promised benefit and never gate on the realized one, and you have built a process that is rigorous about commitments and indifferent to outcomes. The realization gate closes the loop. The benefit that justified the funding gets checked against reality by the same governance that approved it. That is the minimum a serious value discipline requires.

How this closes the funding loop

Benefits realization is the missing back half of portfolio funding discipline. A portfolio can move money crisply on evidence at the front — funding, redirecting, and stopping initiatives on their merits — and still leak value if it never verifies whether the funded benefits appeared. Its evidence about what works is never updated by what actually happened. Funding discipline decides where the money goes. Benefits realization tells the organization whether those decisions were right. That is the information funding discipline needs to stay disciplined.

Together the two form a closed loop. Fund on evidence, deliver, verify the benefit, and feed the verified result into the next funding decision. Most organizations build the first half and skip the second, so their funding decisions never learn. A program is funded on an optimistic benefit, delivered, and closed. The next program is funded on an equally optimistic benefit, because no one measured the last result or fed it back. Benefits realization is what turns a portfolio of separate bets into a portfolio that learns. A portfolio that does not learn will keep making the same optimistic mistake with confidence.

That is why benefits realization belongs to portfolio governance and not merely to project management. It is how the portfolio discovers the truth about its own decisions. Govern funding without governing realization, and you are tuning the front of a process whose back end you refuse to look at.

Who owns the value once the team is gone

The organizational question underneath all of this is harder than it looks. Who owns the value after the delivery team disbands? The default answer is no one, and that is the whole problem. Delivery organizations are built to deliver and then move on to the next delivery. The benefits owner is frequently a business leader in the operational area the transformation was meant to improve, rather than the program manager who built it. That leader is still present, still accountable for the operational outcome, and still there when the benefit is due.

EXHIBIT 2

Two owners, two jobs, two time horizons

Delivery owner

Accountable for building the thing.
Ends at go-live, when the team disbands and moves to the next delivery.

Benefits owner

Accountable for the value. A business leader in the operational area the transformation was meant to improve — still present when the benefit is due.

Assign that ownership at approval. Make it real by writing it into the operational leader's own objectives. That is what keeps the benefit from becoming an orphan at go-live. It also puts the incentive in the right place. The person accountable for the operational outcome has every reason to make sure the transformation improves it,

and to raise the alarm if it does not. Give the job to the delivery team instead, whose success is measured at launch, and no one has much reason to verify the harder, later outcome. The benefit goes unowned at exactly the moment ownership matters. Put benefits ownership where the outcome lives and the accountability persists, not where the delivery happened and the accountability ends.

Start with your biggest program

You do not need to govern benefits this way across every program at once. Trying to will probably stall. Start with the organization's biggest transformation, the one whose value case is largest and most closely watched. Govern its benefits realization completely. Baseline the outcome. Name a benefits owner who outlasts delivery. Put the realization gate on the calendar at approval. Hold the review after go-live, with a real decision attached. That is where the value at stake justifies the discipline, and where a visible success sets the example.

From there the discipline spreads by example and by expectation. Leaders who watched a realization gate protect real value ask for the same on the next major program. The benefits owner role becomes a normal part of how transformations are set up, rather than an unusual imposition. The effort stays proportionate. Apply it first where the stakes are highest, and widen it as it proves it captures value that was previously leaking away unmeasured.

The real obstacle is that nobody wants to know

Benefits realization is rarely skipped out of ignorance. Most experienced leaders know that shipping an output does not by itself capture the value it was meant to produce. The reason is cultural. Verifying benefits creates accountability that many organizations are quietly structured to avoid. A realization review can conclude that a large, sponsored, celebrated program did not deliver its promised value. That conclusion comes attached to an owner, a sponsor, and a story of success. The path of least resistance is to skip the review, and the culture obliges by treating go-live as the end of the story.

Changing that means changing what a benefits shortfall means. In a punitive culture, discovering that a program underdelivered is a failure to be hidden, which guarantees no one goes looking. In a learning culture, the same discovery is information the organization already paid for. The organization can use it to correct the program,

improve the next estimate, and get better at turning investment into outcomes. The organizations that realize benefits well are the ones that made it safe to measure honestly. They treat a verified shortfall as a lesson. Reframing the review from blame to learning is what makes leaders willing to hold it at all.

This mirrors the discipline of stopping a project in portfolio governance. A portfolio that treats stopping as a failure will stop nothing. An organization that treats a benefits shortfall as a scandal will measure nothing. The fix in both cases is to make the uncomfortable act normal — stopping weak work, measuring real outcomes — and to treat it as the process working correctly. Benefits realization is a cultural commitment to look honestly at outcomes as much as it is a set of mechanics. The mechanics do not survive without the commitment. An organization that cannot bear to know whether its transformations worked will ensure, structurally, that it never finds out.

Govern the value, not just the delivery

Transformations get governed as if the hard part is delivery and the value takes care of itself once the solution is live. The reverse is closer to the truth. Delivering the output is the part the organization already knows how to govern. Realizing the value is the part it consistently abandons, at exactly the moment the value is supposed to appear. That abandonment, not any failure of delivery, is where most transformations quietly fail.

Govern the end of the lifecycle as deliberately as the middle. Name a benefits owner who outlasts the delivery team. Baseline the outcome before you start. Measure at the level where the value actually lives. Verify after go-live, and act on what you find. Feed the result back into the next optimistic business case. Do that, and transformations produce the value they promised. Skip it, and the organization will keep celebrating deliveries while the benefits it funded quietly fail to arrive.

Notes and sources

1. Project Management Institute, "Benefits Realization Management: A Practice Guide." <https://www.pmi.org/standards>
2. Bent Flyvbjerg and Dan Gardner, "How Big Things Get Done," Crown Currency, February 7, 2023. <https://www.penguinrandomhouse.com/books/672118/how-big-things-get-done-by-bent-flyvbjerg-and-dan-gardner/>

3. Johannes-Tobias Lorenz, Joshan Cherian Abraham, Robert Levin, and Douglas Ziman, "From promise to impact: How companies can measure—and realize—the full value of AI," McKinsey & Company, April 24, 2026.

<https://www.mckinsey.com/capabilities/quantumblack/our-insights/from-promise-to-impact-how-companies-can-measure-and-realize-the-full-value-of-ai>

4. Michael Bloch, Sven Blumberg, and Jürgen Laartz, "Delivering large-scale IT projects on time, on budget, and on value," McKinsey & Company, 2012.

<https://www.mckinsey.com/capabilities/tech-and-ai/our-insights/delivering-large-scale-it-projects-on-time-on-budget-and-on-value>

5. Patrick Forth, Tom Reichert, Romain de Laubier, and Saibal Chakraborty, "Flipping the Odds of Digital Transformation Success," Boston Consulting Group, October 29, 2020. <https://web-assets.bcg.com/c7/20/907821344bbb8ade98cbe10fc2b8/bcg-flipping-the-odds-of-digital-transformation-success-oct-2020.pdf>

About the author

Marco Policani is an enterprise portfolio, PMO, and AI operating-governance leader. He builds portfolio governance systems where AI carries the analytical load and named humans own every decision — an approach documented across case studies, walkthroughs, and working governance guides on his portfolio site.

Portfolio & governance library: policani.net/governance · Full portfolio: policani.net · LinkedIn: linkedin.com/in/marcpolicani

This white paper may be shared freely with attribution. © 2026 Marco Policani.