


Benefits Realization Is Where Transformations Quietly Fail

Why programs that ship on time and on budget still fail to deliver value — and how to govern the end of the lifecycle, not just the middle



Marco Policani

Enterprise Portfolio · PMO · AI Operating Governance

policani.net · linkedin.com/in/marcpolicani

July 2026

EXECUTIVE SUMMARY

A transformation has not created value until the benefit is measured, owned, and acted on after delivery.

Most transformation failures are not delivery failures. The program ships, the milestones are checked, and the value that justified the effort is never measured. This paper defines benefits realization as a governed discipline and shows how to run the end of the lifecycle as deliberately as the middle.

- 01** Delivery answers “did we build it.” Benefits realization answers “did it produce the value we funded it for” — a distinct discipline most organizations skip.
- 02** Six ordinary disciplines close the gap: a named benefits owner, a pre-start baseline, outcome-layer targets, a scheduled post-delivery review, a real decision on the result, and a feedback loop into the next business case.
- 03** The structural fix is a realization gate after go-live — scheduled at approval, owned by a named leader, asking whether the benefit materialized and what happens next.

CONTENTS

The discipline organizations skip	Two limits worth stating
Why the failure is invisible	The realization gate
What benefits realization is	Benefits realization completes funding discipline
Optimism at the front, silence at the back	Who owns benefits after go-live
Measure benefits, not activity	Starting with the flagship
The realization disciplines	The cultural obstacle
The questions delivery leaves unanswered	Govern the value, not just the delivery
Two patterns: value captured, value lost	Notes and sources
This applies squarely to AI transformations	

The discipline organizations skip

Most transformation failures are not dramatic. The program does not collapse, miss its deadline, or blow its budget in a way that makes the news inside the company. It ships. The system goes live, the milestones are checked, the team is thanked, and everyone moves on. Months later, quietly, the value that justified the

whole effort has not appeared, and no one is quite sure why, or even whether anyone is still measuring. This is a common way transformations fail: at benefits realization, the neglected end of the lifecycle where value is supposed to be captured and often is not. Large-project research shows how big the value shortfall can be: a McKinsey study with the University of Oxford of more than 5,400 large IT projects found they delivered, on average, 56 percent less value than predicted⁴, and in BCG's study of 895 digital transformations, rated largely by the executives who led them, only about 30 percent met their target value and produced sustainable change⁵.

Benefits realization, actually capturing and verifying the value a transformation was funded to produce, is a distinct discipline that most organizations skip, and skipping it is why so many well-delivered programs produce so little measurable benefit. The question that determines success is whether a program produced the value it was funded for, and an organization that governs delivery rigorously while ignoring benefits realization is measuring the wrong end of its own programs.

This is well established in the delivery profession even where it is poorly practiced. The Project Management Institute maintains benefits realization management as a practice guide within its standards, describing benefits realization as the thread that runs from strategy through project deliverables, and its own survey found only one in three organizations reporting high benefits-realization maturity.¹ The gap between the two is where the value leaks out, and this paper is about closing it: defining what benefits realization is, why it is so consistently neglected, and how to govern it as deliberately as delivery.

Why the failure is invisible

Benefits-realization failure is invisible because of an asymmetry in what organizations measure. Delivery is measured intensely: schedule, budget, scope, milestones, defects, go-live. Benefits are measured rarely, if at all, because measuring them requires effort after the exciting part is over, from a team that has usually been disbanded and reassigned. The result is that a program can be a delivery triumph and a benefits failure and appear, on every dashboard the organization actually keeps, to be a success.

The timing makes it worse. The moment a system goes live is treated as the finish line, when it is actually the starting line for value: benefits accrue after go-live, as the changed way of working beds in and the intended outcomes either materialize or do not. But go-live is exactly when attention moves on, budgets close, and the people who understood the value case scatter to the next initiative. The organization declares victory at precisely the moment it stops watching whether it won.

There is also a comfort in not looking. Verifying benefits risks discovering that a large, celebrated program did not deliver the value it promised, which is uncomfortable for everyone who sponsored, approved, or led it. Skipping the measurement protects the story of success and leaves its substance unverified. This is the same avoidance that lets weak projects live in a portfolio, applied to the back end of the lifecycle: measuring would force a reckoning, so the organization declines to measure.

What benefits realization is

Benefits realization can be defined precisely enough to govern:

Benefits realization is the discipline of defining the specific value a transformation is expected to produce, establishing the baseline and the measure before it starts, assigning a named owner accountable for the value (not merely for the delivery), and verifying after go-live whether the value actually materialized, with a real decision to act on the answer. If a program has no benefits owner, no baseline, and no post-delivery verification, it has only a delivery plan, whatever the business case claimed about value at approval.

The definition deliberately separates two roles that organizations routinely conflate: the delivery owner, accountable for building the thing, and the benefits owner, accountable for the value the thing was supposed to produce. These are different jobs with different time horizons, the delivery owner's ends at go-live, the benefits owner's begins there, and collapsing them into one role that ends at delivery is a primary reason benefits go unrealized. Someone has to still be accountable after the delivery team has moved on, or the value has no owner at exactly the time it is supposed to appear.

Optimism at the front, silence at the back

Benefits realization is undermined at both ends of the lifecycle, and the two failures compound. At the front, benefits are systematically over-promised, because optimism is what gets a program funded. Flyvbjerg and Gardner's research on large projects finds that estimates made from the inside view, by the people close to and invested in the initiative, are prone to optimism bias, and that a more realistic estimate comes from taking an outside view, comparing the initiative to the actual outcomes of a reference class of similar past efforts rather than to its own best-case story.² Benefits promised from the inside view are inflated, and the inflation is baked into the business case that justifies the spend.

At the back, there is silence: no one returns to check whether the inflated benefits materialized, so the optimism is never corrected and never learned from. Because benefits are over-promised at approval and never verified at realization, the organization systematically funds programs on optimistic value cases it never checks, so its optimism is never disciplined by evidence and its next business case is just as inflated as the last. The outside view fixes the front end by grounding estimates in real reference-class outcomes; benefits verification fixes the back end by feeding actual results back into

the next estimate. An organization that does neither repeats optimistic, unverified value cases indefinitely.

Measure benefits, not activity

When organizations do attempt to measure benefits, they often measure the wrong layer, stopping at activity or output rather than reaching the outcome that was the point. McKinsey's framework for realizing value is useful here precisely because it separates the layers: technical performance, user adoption, operational KPIs, strategic outcomes, and financial impact are distinct, and a program can succeed at the early layers while failing at the ones that matter.³ A system that works technically and is widely adopted has cleared the first two layers and may still have moved no operational or financial outcome at all.

The discipline is to define, before the program starts, which layer the benefit lives in and to measure to that layer rather than stopping at the convenient early ones. It is easy to report that the system is live and usage is high, because those are visible at go-live and feel like success. It is harder, and essential, to measure whether the operational outcome improved and the financial value appeared, because those are what the program was funded for and they show up later, if at all. McKinsey's guidance to build the value case up front, with evidence and staged decisions, is the same discipline seen from the funding side: decide what benefit at what layer, and hold the program to it through realization, not just to go-live.³

The realization disciplines

Governing benefits realization comes down to a short set of disciplines, each ordinary and each routinely skipped once the delivery excitement takes over.

EXHIBIT 1

The benefit ownership layer

Named benefits owner Owns value after go-live, separate from delivery ownership.	Pre-start baseline Capture the starting measure before work begins.
Target at the right layer Set the operating or financial outcome, not an output or adoption proxy.	Schedule the review Commit at approval to review after the benefit has had time to appear.
Decide and learn Adjust, invest, or stop; feed verified results into the next case.	

None of these is difficult; all of them require someone to still care about the program after the ribbon is cut. The real challenge of benefits realization is attention rather than mechanics: keeping a named owner accountable and a review scheduled for a time when the organization's focus has moved on.

The questions delivery leaves unanswered

The difference between a program governed for delivery and one governed for benefits shows up in what questions it can answer. The table pairs the delivery signal with the benefits-realization question it leaves unanswered.

The delivery signal	What it confirms	The benefits question it leaves open
Shipped on time and on budget	The program was well delivered	Did the value it was funded for appear
The system is live	The output exists	Did the outcome the output was meant to change move
Adoption is high	People are using it	Is the operational or financial benefit materializing
The project is closed	Delivery is complete	Who is still accountable for the value after go-live
The business case was approved	The benefit was promised	Was the promise checked against a reference class, and verified later

The right column is where transformations quietly fail. Every question in it is answerable, and most organizations do not answer any of them, because they governed the middle of the lifecycle and abandoned the end. Benefits realization is the discipline of answering the right column on purpose.

Two patterns: value captured, value lost

The safest illustrations are operating patterns, not named programs. In portfolio and program work, the durable pattern is that the organizations which actually capture value are the ones that governed the back end of the lifecycle as seriously as the front, keeping an owner and a measure alive past go-live.

In one pattern, a program was set up so that the value case was not closed at delivery. A benefits owner distinct from the delivery lead carried a baseline and a target through go-live, a realization review was scheduled at approval for months after the launch, and the review actually happened, with the authority to trigger further investment or a stop based on whether the outcome moved. The value was captured because someone was still accountable for it when it was supposed to appear.

In a second pattern, an organization discovered that a series of celebrated deliveries had produced little measurable benefit, because no one had ever gone back to check. The corrective was not more delivery rigor, which was already strong, but the addition of the missing back-end discipline: baselines captured before programs began, benefits owners named and retained, and realization reviews that fed verified results into the next round of business cases. The delivery machine was fine; the value machine had never been built, and building it was what turned motion into outcomes.

This applies squarely to AI transformations

Nowhere is benefits-realization discipline more urgently needed than in AI transformations, which are unusually prone to the exact failure this paper describes. AI programs generate abundant activity, usage, pilots, deployments, that is easy to mistake for value, and the enthusiasm around AI makes the front-end optimism especially strong and the back-end verification especially rare. The result is a wave of AI initiatives that ship impressively and deliver benefits no one has actually measured.

The measurement layers matter even more here, because AI programs can succeed at the early layers, technical performance and adoption, while failing to move the operational and financial outcomes that justify them.³ An AI deployment that works and is widely used looks like a success and may have changed no business outcome at all. Applying benefits-realization discipline to AI, naming a benefits owner, baselining the outcome, measuring to the right layer, and verifying after go-live, is what separates AI transformations that produce value from the far larger number that produce activity. The discipline is not new; AI has simply made the cost of skipping it larger and more common.

Two limits worth stating

Two limits are worth stating. First, this paper argues the discipline of benefits realization rather than a single quantified failure rate; published statistics on how much transformation value goes unrealized vary widely by source, sector, and definition, and readers should measure their own realization gap rather than import a headline number. The profession treats benefits realization as a core discipline, with a PMI practice guide devoted to it; the precise magnitude of the loss when it is skipped is organization-specific.¹ Second, benefits realization is necessary without being sufficient on its own: measuring and owning benefits does not guarantee a program was the right bet in the first place. It ensures the organization learns the truth about its bets and acts on it, which over a portfolio of many transformations is worth more than being confident about any single one. Benefits realization governs how value is verified; it does not settle whether the underlying strategy was sound.

The realization gate

The single most effective structural fix for benefits-realization failure is to add a gate that most delivery lifecycles are missing: a realization gate after go-live. Delivery lifecycles are rich with gates up to launch, approval, design, build, readiness, and then they simply stop, as if the lifecycle ended when the system went live. The realization gate reopens the program at a scheduled point after launch, when the value has had time to appear, and asks the only question that ultimately matters: did the benefit materialize, and if not, what will we do about it?

EXHIBIT 3

The realization gate reopens the program at a scheduled point after launch, when the value has had time to appear, and asks the only question that ultimately matters: did the benefit materialize, and if not, what will we do about it?

A realization gate works because it is scheduled at approval, when the program is eager to commit, rather than left to be organized after go-live, when no one wants to. Committing at the outset to a specific date on which the benefit will be checked, with a named owner responsible for showing up with the evidence, is what keeps the back end from being abandoned. The gate is not a formality; it carries a real decision, to accept the realized benefit as sufficient, to invest further to close a gap, or to stop pouring effort into an outcome that is not coming. Without that decision, the review is theater; with it, the gate is where value is actually protected.

The realization gate also completes the governance symmetry with the front of the lifecycle. An organization that gates funding on a promised benefit but never gates on the realized one has built a system that is rigorous about commitments and indifferent to outcomes. Adding the realization gate closes the loop: the benefit that justified the funding is checked against reality by the same governance that approved it, which is the minimum a serious value discipline requires.

Benefits realization completes funding discipline

Benefits realization is the missing back half of portfolio funding discipline. A portfolio that moves money crisply on evidence at the front, funding, redirecting, and stopping initiatives on their merits, still leaks value if it never verifies whether the funded benefits materialized, because its evidence about what works is never updated by what actually happened. Funding discipline decides where the money goes; benefits realization tells the organization whether those decisions were right, which is the information funding discipline needs to keep being disciplined.

Seen together, the two form a closed loop: fund on evidence, deliver, verify the benefit, and feed the verified result back into the next funding decision. Most organizations build the first half and skip the second, which means their funding decisions never learn. A program is funded on an optimistic benefit, delivered, and closed, and the next program is funded on an equally optimistic benefit because the last one's actual result was never measured and never fed back. Benefits realization is what turns a portfolio from a sequence of independent bets into a system that learns, and a portfolio that does not learn will keep making the same optimistic mistake with confidence.

This is why benefits realization belongs to portfolio governance and not merely to project management. It is the mechanism by which the portfolio discovers the truth about its own decisions, and a portfolio

that governs funding without governing realization is optimizing the front of a process whose back end it refuses to look at.

Who owns benefits after go-live

The organizational question underneath benefits realization is deceptively hard: who, exactly, owns the value after the delivery team disbands? The default answer is no one, which is the whole problem, because delivery organizations are structured around delivering and then moving on to the next delivery. The benefits owner is frequently a business leader in the operational area the transformation was meant to improve, rather than the program manager who built it, precisely because that person is still present, still accountable for the operational outcome, and still there when the benefit is supposed to appear.

EXHIBIT 2

Two owners, two jobs, two time horizons

Delivery owner

Accountable for building the thing. Ends at go-live, when the team disbands and moves to the next delivery.

Benefits owner

Accountable for the value. A business leader in the operational area the transformation was meant to improve — still present when the benefit is due.

Assigning that ownership at approval, and making it real by tying it to the operational leader's own objectives, is what keeps the benefit from becoming an orphan at go-live. It also aligns incentives correctly: the person accountable for the operational outcome has every reason to ensure the transformation actually improves it, and to raise the alarm if it does not. When the benefits owner is instead the delivery team, whose success is measured at launch, the incentive to verify the harder, later outcome simply is not there, and the benefit goes unowned at exactly the moment ownership matters. The design principle is to place benefits ownership where the outcome lives and the accountability persists, not where the delivery happened and the accountability ends.

Starting with the flagship

Building a full benefits-realization capability across every program at once is unnecessary and likely to stall. The practical entry point is to pick the organization's most consequential transformation, the flagship whose value case is largest and most scrutinized, and govern its benefits realization completely: baseline the outcome, name a benefits owner who outlasts delivery, schedule the

realization gate at approval, and hold the post-go-live review with a real decision. That is where the value at stake justifies the discipline and where a visible success establishes the pattern.

From that flagship, the discipline spreads by example and by expectation. Leaders who saw a realization gate protect real value on the flagship ask for the same on the next major program, and the benefits owner role becomes a normal part of how transformations are set up rather than an unusual imposition. This keeps the effort proportionate: applied first where the stakes are highest, and expanding as the discipline proves it captures value that was previously leaking away unmeasured.

The cultural obstacle

The reason benefits realization is skipped is rarely ignorance of its importance; most experienced leaders know that shipping an output does not by itself capture the value it was meant to produce. The reason is cultural: verifying benefits creates accountability that many organizations are quietly structured to avoid. A realization review can conclude that a large, sponsored, celebrated program did not deliver its promised value, and that conclusion has an owner, a sponsor, and a story of success attached to it. The path of least resistance is not to hold the review, and the culture obliges by treating go-live as the end of the story.

Changing this requires reframing what a benefits shortfall means. In a punitive culture, discovering that a program underdelivered is a failure to be hidden, which guarantees the discovery is avoided. In a learning culture, the same discovery is information the organization paid for and should use, to correct the program, to improve the next estimate, and to get better at turning investment into outcomes. The organizations that realize benefits well are the ones that made it safe to measure honestly, because they treat a verified shortfall as a lesson to learn from. Reframing the review from blame to learning is what makes leaders willing to hold it at all.

This mirrors the discipline of the kill decision in portfolio governance. Just as a portfolio that treats stopping a project as a failure will avoid stopping anything, an organization that treats a benefits shortfall as a scandal will avoid measuring benefits. In both cases the fix is to normalize the uncomfortable act — stopping weak work, measuring real outcomes — as the system functioning correctly. Benefits realization is as much a cultural commitment to look honestly at outcomes as it is a set of mechanics, and the mechanics do not survive without the commitment. An organization that cannot bear to know whether its transformations worked will ensure, structurally, that it never finds out.

Govern the value, not just the delivery

Transformations are governed as if the hard part is delivery and the value takes care of itself once the system is live. The reverse is closer to the truth. Delivering the output is the part the organization

already knows how to govern; realizing the value is the part it consistently abandons at exactly the moment the value is supposed to appear. That abandonment, not any failure of delivery, is where most transformations quietly fail.

Govern the end of the lifecycle as deliberately as the middle. Name a benefits owner who outlasts the delivery team, baseline the outcome before you start, measure to the layer where the value actually lives, verify after go-live, and act on what you find, including feeding the result back into the next optimistic business case. Do that, and transformations produce the value they promised. Skip it, and the organization will keep celebrating deliveries while the benefits it funded them for quietly fail to arrive.

Notes and sources

1. Project Management Institute, "Benefits Realization Management: A Practice Guide." <https://www.pmi.org/standards>
2. Bent Flyvbjerg and Dan Gardner, "How Big Things Get Done," Crown Currency, February 7, 2023. <https://www.penguinrandomhouse.com/books/672118/how-big-things-get-done-by-bent-flyvbjerg-and-dan-gardner/>
3. Johannes-Tobias Lorenz, Joshan Cherian Abraham, Robert Levin, and Douglas Ziman, "From promise to impact: How companies can measure—and realize—the full value of AI," McKinsey & Company, April 24, 2026. <https://www.mckinsey.com/capabilities/quantumblack/our-insights/from-promise-to-impact-how-companies-can-measure-and-realize-the-full-value-of-ai>
4. Michael Bloch, Sven Blumberg, and Jurgen Laartz, "Delivering large-scale IT projects on time, on budget, and on value," McKinsey & Company, 2012. <https://www.mckinsey.com/capabilities/tech-and-ai/our-insights/delivering-large-scale-it-projects-on-time-on-budget-and-on-value>
5. Patrick Forth, Tom Reichert, Romain de Laubier, and Saibal Chakraborty, "Flipping the Odds of Digital Transformation Success," Boston Consulting Group, October 29, 2020. <https://web-assets.bcg.com/c7/20/907821344bbb8ade98cbe10fc2b8/bcg-flipping-the-odds-of-digital-transformation-success-oct-2020.pdf>

About the author

Marco Policani is an enterprise portfolio, PMO, and AI operating-governance leader. He builds portfolio governance systems where AI carries the analytical load and named humans own every decision — an approach documented across case studies, walkthroughs, and working governance guides on his portfolio site.

Portfolio & governance library: policani.net/governance · Full portfolio: policani.net · LinkedIn: linkedin.com/in/marcpolicani

This white paper may be shared freely with attribution. © 2026 Marco Policani.