
WHITE PAPER · PORTFOLIO & AI GOVERNANCE

Benefits Realization Is Where Transformations Quietly Fail

Delivery governance can end before the value is checked

Marco Policani

Enterprise Portfolio · PMO · AI Operating Governance

policani.net · [linkedin.com/in/marcpolicani](https://www.linkedin.com/in/marcpolicani)

July 2026



EXECUTIVE SUMMARY

Schedule and budget controls can end at go-live while the value arrives later, so settle at approval who answers for the value, how it will be measured, and when it will be reviewed.

A program can go live with every delivery control green and still never be checked for the value it was funded to produce. The schedule, budget, and scope controls can end at go-live while the benefit arrives afterward, so a missed benefit can go unseen. Settling at approval who answers for the value, how it is measured, and when it is reviewed puts someone in charge of checking.

- 01** Name a benefit owner outside the delivery team at approval: the operating leader whose area is meant to improve, with the measure written into that leader's own objectives.
- 02** Fix a measure of the result the program was funded to change, such as fewer repeat contacts rather than logins. Take its baseline before the program starts, and say how the program's effect will be separated from other change.
- 03** Book the benefit review for a date after launch, with a decision attached: accept what the measure shows, invest more, change the approach, or stop. Keep each result so the next business case can be set beside it.

A program can deliver and still never be checked for value

Consider a customer-service platform program, funded because its business case says it will cut repeat contacts. It reaches go-live. The milestones are signed off, the delivery team is thanked and reassigned, and the dashboard shows green on schedule, budget, and scope. The head of customer service sees repeat contacts every week, but nothing on anyone's calendar compares them with the business case, asks whether the platform caused any change, or decides what follows. This is an illustration, not a real program. The rest of this paper uses it as its running example.

The argument is narrow. Delivery governance, meaning the schedule, budget, and scope controls that run a program, can end at go-live, and the value the program was funded for arrives after that. So a program can deliver and still never be checked for value. Delivery also fails. The title names the place where a failure can go unseen, because nothing is set up to look. No source reviewed here shows at which stage most transformations fail.

If the check never happens, the next program is approved on the same kind of unchecked business case, and a shortfall in one program is never found. No source measures that exact case, a program delivered to its controls that still missed its value, so the studies below do not separate delivery failure from value shortfall.

Here, a transformation is a program of changes to how an organization works, funded for a business result. Benefits realization is the work of confirming that the program produced that result. It is not a new idea. The Project Management Institute (PMI) tracks it as a maturity measure, and the U.S. Government Accountability Office (GAO) guide to managing federal technology investments includes reviews after launch that compare actual results with the estimates.[1][2] PMI's 2018 survey of project practitioners, executives, and directors of project management offices reported that one organization in three has high maturity in benefits realization. That is a self-report. It suggests that two in three organizations do not rate themselves highly at it, not that they fail to do it. PMI also names, as a challenge many companies face, managing projects on time, scope, and budget without consistently tracking whether they help reach the organization's strategic goals.[1]

Two studies show how large the shortfalls can be, though neither separates programs that delivered well from those that did not. In a 2012 article, McKinsey and the University of Oxford compared the benefits predicted at the start with actual results for more than 5,400 IT projects. Large ones, with initial price tags above \$15 million, averaged 45 percent over budget and 7 percent over time, and they delivered 56 percent less value than predicted.[3] The study shows two things. Large IT projects often run over, and value shortfalls can be measured after the fact. It does not show how many of those projects delivered well and still fell short on value. The article also says IT initiatives too often manage purely to budget and schedule targets.[3]

In 2020, Boston Consulting Group (BCG) asked 825 senior executives to rate their own companies' digital transformations and also analyzed 70 of its client programs. The score combined targets met, value created, timing, and how lasting the change was. Only about 30 percent met or exceeded their target value and produced lasting change.[4] The figure is executives rating themselves, it covers digital transformations, and timing is part of the score, so it is not a measure of value alone.

Why delivery governance can end before the value is due

In many programs, three things end at go-live. The delivery team disbands, and its people move to the next program, including the people who understood the value case. The program's budget closes, so no one has money to measure with. Attention moves to the next approval, which is already on the calendar. This is the paper's reading of how delivery organizations are usually set up, not a measured finding.

The benefit falls due after all three have gone. Benefits build as people change how they work, and how long that takes varies by program. GAO's guidance on reviewing an investment after launch says a review held too soon can miss the full benefits, and one held too late can lose what people knew about the investment.[2] So the question falls due when no one is positioned to ask it. That is the gap this paper addresses: the check has no owner at the moment it can be answered.

There may also be a reason to look away. A review could show that a celebrated program did not deliver, which falls on the sponsor, the executive who championed the program and answers for its success, and on the delivery lead who ran it. This is the paper's reading of the incentive, not a measured one. It is why the check has to be arranged while the program is being approved, before the incentive to skip it sets in.

The gap is easier to miss now that AI programs are common, because use is cheap to count: logins, prompts, pilots, deployments. A count of use says people tried the tool, not that cycle time or error rates changed. That is this paper's inference, not a measured pattern. "Felt Productivity Is Not Measured Productivity" covers how to measure what an AI tool changes.

Five things to settle at approval

PMI measures how mature organizations are at benefits realization, and GAO describes how to review benefits after launch. What this paper sets out is the order and the timing: all five are settled at approval, before anyone is reassigned. Placing the benefit owner outside the delivery team is this paper's own design. The five, in plain words, are taken one at a time below.

EXHIBIT 1

Five things to settle at approval

- ① **Benefit owner**
An operating leader outside the delivery team answers for the value after go-live.
- ② **Outcome measure**
Track the operating or financial outcome, not logins or the go-live date.
- ③ **Baseline**
Take the starting reading of the measure before the program starts.
- ④ **Benefit review**
Book a date after launch, with a decision attached.
- ⑤ **Record of results**
Keep each result so the next business case can be set beside it.

They are not a checklist in any order. The benefit owner agrees what will be measured. The baseline is the starting reading of that measure, so the measure comes first. How fast the benefit is expected to build sets the review dates. At the benefit review, the owner brings the measure against the baseline, and the funding review, the recurring meeting of the leaders who control the budget, takes a decision. The fifth follows, and it works only if the first four were done, because the record can hold only what was measured.

The benefit owner sits outside the delivery team

Name one person to answer for the value after go-live: the benefit owner. The benefit owner is the operating leader whose area the change is meant to improve, such as the head of customer service in the running example. The delivery lead, who answers for building the platform, keeps that job. The two jobs run on different clocks. The delivery lead's ends at go-live, and the benefit owner's starts there.

The operating leader is the better owner for a practical reason. That leader is still in the role when the benefit is due, and the operating result is part of the leader's own job. The fit is not automatic. One executive quoted in BCG's report said it was essential, and not easy, to connect the people accountable for delivering the technology with the line managers who had signed up for the business impact.[4] Writing the measure into the owner's own objectives is this paper's way of making the connection real. Without it, the role is a title.

Other project and program frameworks may already give benefits to business-side roles. This paper did not review them, so the owner-outside-delivery rule is presented as its design, not as a new idea.

The owner uses judgment at three points. At approval, the owner checks that the measure is the number that would move if the program worked. At the benefit review, the owner reports the measure against the baseline, along with other changes in the same period. Afterward, the owner recommends what to do. Who decides after a missed result is the subject of “Why the Kill Gate Doesn’t Fire,” which argues for someone who had no part in the approval. This paper asks only that the benefit review has a decision to make.

Some benefits belong to several areas. In that case, name one lead owner for the whole benefit and let each area’s leader carry that area’s share of the measure. If no operating leader will take the benefit, the program is not ready for approval, because no one yet intends to produce the value. That is the paper’s position, not a finding.

Owning a benefit costs time and money, and the business case should say how much. The owner spends time tracking the measure and attending reviews, someone may have to build the measure into the data the owner already receives, and the delivery team may have to stay a short while after go-live to hand it over. This paper has no source for typical amounts, and they differ by program. Write the estimate into the approval so the decision to fund includes it.

The measure and its baseline are set before the work starts

Choose a measure of the result the program was funded to change, not of the work done. Counting activity, such as training sessions held, tells you people did things. Counting output, such as a live system, tells you something was built. Counting adoption tells you people are using it. An operating result is the change in how the work performs, such as fewer repeat contacts or shorter cycle time, and a financial result is the change in cost or revenue. A program can pass the first three levels and move neither of the last two. In the running example, a live platform with high daily use has cleared the first three. It has not shown fewer repeat contacts. Track output and adoption anyway, as early signs that a program is on course, but do not accept them as proof of value.

BCG’s monitoring factor puts this in terms of operational or financial measures tied to the program’s purpose, tracked regularly at both the program and initiative levels,

with a clear mandate and one trusted data source. Only two in five organizations in its study set up all of that adequately, against 90 percent of those with winning transformations.[4] That is an association in executives' self-ratings, not a test of cause.

The baseline is the starting reading of the measure, taken before the work starts. Take it over a stretch long enough to show the measure's normal ups and downs, and record who took it and from what data. A benefit that was never measured at the start is hard to prove at the end. BCG's report lists, among its checks for tracking, a trusted source of operational and financial data that settles baselines, net versus gross benefits (after subtracting offsetting costs), and double counting (the same saving claimed twice).[4] Where the organization has no such source, the owner and the delivery lead agree at approval on which data source counts.

Separating the program's effect from everything else that changed takes a method. GAO guidance on evaluating programs explains why. Where outside factors affect an outcome, the observed result mixes the program's effect with theirs. A simple before-and-after comparison cannot show cause, because other factors may have changed in the same period. The stronger design compares the outcome with an estimate of what would have happened without the program, for example with a comparison group, and repeated readings over time strengthen a before-and-after comparison. [5] The guidance was written for public programs, so applying it to business programs is this paper's inference. Three practices follow, in the paper's order of strength:

- When the change rolls out in stages, keep teams or sites similar to the first ones switched as the comparison, compare the two groups on the same measure, and note how the rollout order was chosen.
- When it cannot be staged, take repeated readings before and after launch, not one pair, and record what else changed in the period (a price change, a new policy, a seasonal swing). Claim a contribution, not the whole difference.
- When neither works, say at approval that the benefit cannot be separated from other change. Record it as expected and unverified, and do not book savings from it.

If a program has no baseline, because it is already live, rebuild one from records for the period before launch, such as system logs, finance data, or work queues. If the records do not exist, the benefit stays unverified. A rough estimate with its method stated is still better than none. Whichever case applies, say which one at the benefit review.

The benefit review is booked at approval and carries a decision

The benefit review is a meeting held after go-live, at a date set at approval. The benefit owner brings the measure against the baseline, the other changes in the period, and the value the business case promised, and the funding review then chooses one of four decisions:

EXHIBIT 2

Four decisions the benefit review can take

Accept

The benefit is enough, and it is closed.

Invest more

A gap has a known cause that more money or time can close.

Change the approach

The measure moved less than expected, for reasons the program can fix.

Stop

The result is not coming, so spending toward it ends.

Set the date and the owner's name in the approval record, because at approval everyone is eager to commit, and after go-live nobody is. That is this paper's reading of the incentives, not a measured finding. Someone outside the program, such as the portfolio management office (PMO), the team that supports the funding review, checks the owner's figures against the agreed data source. GAO describes a central group conducting such reviews, with the owners and users of the investment taking part.[2]

Set two dates. The early check is a short look soon after launch at the direction of the measure and at output and adoption. It exists to catch a program that is clearly off course while change is still cheap. The benefit review is held when the business case says the benefit should be at its full level, and it judges the operating or financial result. GAO's guidance suggests 6 to 18 months after the investment begins operating, adjusted for when the plan expects the benefits.[2] That range is for federal technology investments. For a program whose benefit builds slowly, set the benefit review later, and say why. When delivery is continuous and there is no single

go-live, tie the early check to the first release that claims the benefit, and the benefit review to the date the business case expects the full effect.

If the owner cannot produce the measure at the date, that is a different case from a weak result, and it has its own decision. The funding review sets a date to fix the measurement, and any follow-on funding tied to the same benefit waits until it is fixed. Without a decision, the benefit review is a report.

The benefit review costs people's time. GAO's guidance says a review needs a team, with one person leading it, and that in most cases the project team should help. It says reviewing every investment may not be necessary, but an organization should review a mix of successful and less successful projects.[2] This paper's version: review every program above a size the funding review sets, and sample the rest, including programs that look successful.

The result goes into the next business case

The fifth thing is the record. Keep one line per reviewed program: what was promised, what was measured, the date, and what the benefit review decided. GAO's guidance has the same purpose, to compare actual results with estimates and carry the lessons into the investment process.[2] Over time the lines show how past programs compared with their business cases.

Business cases are written before results are known. Flyvbjerg's research on large projects, mostly transport, found that front-end estimates of benefits often differ significantly from what the projects later produced, and his due-diligence steps include checking a forecast against the track record of similar forecasts.[6] Whether transformation business cases behave the same way is untested here. The record is what any such check would draw on, because without recorded results there is no track record to compare against.

The PMO keeps the record. How a funding review uses comparable evidence to decide is the subject of "Portfolio Governance Is a Funding-Discipline Problem, Not a Reporting Problem." This paper asks only that the record exist.

Two illustrations

These are illustrations, not observed programs, and neither carries a result.

In the first, the running example is approved with all five. The head of customer service is named benefit owner, with repeat contacts in the role's objectives. The baseline for repeat contacts comes from the contact system over a period that shows normal swings. The platform rolls out region by region, so similar regions not yet switched serve as the comparison. The early check and the benefit review are booked, and the approval states the four decisions the review can take.

The second is a different program that went live without any of this. Its delivery lead has moved on and nobody owns the benefit. Its steps are the ones described above for a program already live: ask the leader whose area it affects to take the benefit, rebuild a baseline from records for the period before launch, and book a review. The program may or may not have produced its value. The aim is to find out, and to stop approving follow-on work as if the answer were known.

Limits and objections

Many changes affect the same result, so no one can credit the program. Often true, and the attribution practices above are the response. They produce a contribution, not a precise share, and some benefits cannot be separated. The benefit review says so in those cases and the benefit stays unverified.

Measurement costs money, and some benefits are small. The paper's rule is to drop any measure that would cost more than the decision it informs, and to say so at approval. A program with a small expected benefit can have a short review.

Measuring too early penalizes benefits that build slowly. That is the reason for two dates, with the benefit review tied to when the business case expects the benefit at full level.

The owner will be blamed, or will game the measure. The owner answers for reporting the result honestly and for recommending what to do, not for factors outside the area's control. Where possible, the data should come from a source the owner's team does not run, and the PMO checks the figures. A shortfall found early is information the organization has already paid for.

Some programs have no value to measure. Work required by law or contract does not claim a benefit. The approval records the requirement, the cost, and the date, and the program skips the value review.

The evidence is limited. PMI's maturity figure is self-reported. BCG's figure is executives' own ratings of digital transformations, and timing is part of it. McKinsey

and Oxford studied large IT projects in 2012, not transformations in general. Flyvbjerg's evidence is mostly transport. None of them tested whether booking a review at approval raises the value a program realizes, and none shows which stage of a program fails most often. GAO describes parts of the practice, and PMI treats benefits realization as a maturity measure. The way the five fit together, and the requirement that the benefit owner sit outside the delivery team, are this paper's design. Benefits realization governs how value is verified. It does not settle whether the strategy was sound.

What to do next

At the next approval of a program that claims business value, add one page to the approval record. The sponsor writes it, and the PMO checks that every field is filled in. It lists the benefit owner by name, the measure and the level it sits at, the baseline and how it was taken, the early check and benefit review dates, the four decisions the benefit review can take, and the estimated cost of measuring. It is a one-time record made at approval. The owner updates it at each benefit review, and the result is added to the program's line in the PMO's record.

Then test the programs already live. Pick the largest program that went live in the past year and ask four questions. Who is the benefit owner now? Is there a baseline? Is a benefit review on the calendar? Was a decision attached? A no is likely, because the five-thing design is this paper's own and PMI's figure suggests most organizations do not rate themselves highly at benefits realization. A small or required program may also never have claimed value. A no does not show that the benefit failed to appear. It shows that no one built in the check, which is the case this paper is about. A partial yes is useful too, because the page then lists which of the four is missing. Start with the single largest program, finish its page and its first benefit review, and widen to others once that review has produced a decision.

The customer-service program in the opening had a green dashboard and no one scheduled to compare repeat contacts with the business case. If it had gone to approval with this page, someone would have been named, a baseline taken, and a date booked. Whether repeat contacts fell because of the platform is something it could then answer.

Notes and sources

[1] Project Management Institute, "Pulse of the Profession 2018: Success in Disruptive Times," 2018.

<https://www.pmi.org/-/media/pmi/documents/public/pdf/learning/thought-leadership/pulse/pulse-of-the-profession-2018.pdf>

[2] U.S. General Accounting Office (now Government Accountability Office), Information Technology Investment Management: A Framework for Assessing and Improving Process Maturity, GAO-04-394G, March 2004.

<https://www.gao.gov/products/gao-04-394g>

[3] Michael Bloch, Sven Blumberg, and Jürgen Laartz, "Delivering large-scale IT projects on time, on budget, and on value," McKinsey & Company (with the University of Oxford), October 1, 2012. <https://www.mckinsey.com/capabilities/tech-and-ai/our-insights/delivering-large-scale-it-projects-on-time-on-budget-and-on-value>

[4] Patrick Forth, Tom Reichert, Romain de Laubier, and Saibal Chakraborty, "Flipping the Odds of Digital Transformation Success," Boston Consulting Group, October 29, 2020. <https://web-assets.bcg.com/c7/20/907821344bbb8ade98cbe10fc2b8/bcg-flipping-the-odds-of-digital-transformation-success-oct-2020.pdf>

[5] U.S. Government Accountability Office, Designing Evaluations: 2012 Revision, GAO-12-208G, January 2012. <https://www.gao.gov/products/gao-12-208g>

[6] Bent Flyvbjerg, "Quality Control and Due Diligence in Project Management: Getting Decisions Right by Taking the Outside View," International Journal of Project Management, 2013. <https://arxiv.org/abs/1302.2544>

About the author

Marco Policani is an enterprise portfolio, PMO, and AI operating-governance leader. He builds portfolio governance systems where AI carries the analytical load and named humans own every decision — an approach documented across case studies, walkthroughs, and working governance guides on his portfolio site.

Portfolio & governance library: policani.net/governance · Full portfolio: policani.net · LinkedIn: linkedin.com/in/marcpolicani

This white paper may be shared freely with attribution. © 2026 Marco Policani.