

Business Imagination Is the Scarce Asset Now

A portfolio-governance standard for turning cheap AI execution into measurable business value

Marco Policani

Enterprise Portfolio · PMO · AI Operating Governance

policani.net · linkedin.com/in/marcpolicani

July 2026

EXECUTIVE SUMMARY

Before funding the tool, require a testable account of the work worth changing, the value at stake, the evidence needed, the accountable owner, and the route that fits the job.

A portfolio review can see more AI outputs and more proposals while still lacking a basis for choosing which work deserves scarce capacity. This paper makes business imagination usable through one opportunity brief placed before the tool decision.

- 01** Treat a demo, usage count, or local time saving as evidence about a tool or task, not proof that an investment deserves priority.
- 02** Require one brief that names the work, baseline, value mechanism, evidence test, owner, reliance boundary, and competing routes.
- 03** Keep the brief attached after approval so the review can advance, redirect, wait, or stop on connected evidence.

The review has more ideas than decisions

You chair the portfolio review that decides which AI ideas receive scarce capacity. One team brings a working assistant. Another has a vendor demo. A third shows time saved on drafting and analysis. All three can produce something. None of that evidence tells you which work deserves to change, what result would matter, or whether AI is the right investment.

Cheaper execution can increase speed, and many leaders reasonably expect wider use to create more value. Models now handle drafting, summarizing, classification, comparison, search, translation, synthesis, and code at a cost and speed that were recently out of reach. A bounded tool can also improve real work. In a field study of customer-support agents, Brynjolfsson, Li, and Raymond found a 14 percent average productivity gain. The largest gains occurred among novice and lower-skilled workers, with little effect on the most experienced workers.[6]

The enterprise picture is less direct. McKinsey's 2025 global survey found that 88 percent of respondents reported regular AI use in at least one business function, while 39 percent reported enterprise-level EBIT impact from AI use.[1] A 2026 McKinsey discussion reported that more than 80 percent of companies were not yet seeing bottom-line impact from their AI investments.[2]

Together, these findings expose the decision in front of the portfolio leader. The field result shows what happened in one defined workflow. Enterprise value also depends on choosing and absorbing the

right changes across workflows. A demo, usage count, or local time saving can move a visible idea ahead before the business problem and value test are clear. Capacity then follows presentation quality, and better operating problems remain hidden when they arrive without a polished prototype.

The question for the review is therefore unresolved: what has become scarce when producing another plausible output is easy? It answers itself through the economics of complements: when one input to value becomes cheap and abundant, the scarce complementary input is where value migrates. Cheap execution does not make value cheap — it relocates the scarcity to the judgment that decides which customer problem is worth an assistant and which output actually changes an outcome.

The scarce judgment comes before the tool

Portfolio investment requires evidence about the work and its consequence. Higher counts of users, prompts, accepted code suggestions, documents, and pilots may leave cycle time, quality, customer experience, risk, and financial performance unchanged.

McKinsey's 2026 measurement framework separates technical performance, user adoption, operational indicators, strategic outcomes, and financial impact.[3] Each layer narrows a different uncertainty: whether the system works, whether people use it in the workflow, whether the work improves, whether that improvement matters to the business, and whether the full economics hold.

The failure can begin even earlier. Gartner predicted in 2024 that at least 30 percent of generative-AI projects would be abandoned after proof of concept by the end of 2025 because of poor data quality, inadequate risk controls, escalating costs, or unclear business value.[4] RAND's 2024 study concluded that misunderstandings and miscommunications about project intent and purpose cause more AI projects to fail than any other factor, and it advised organizations to focus on the problem rather than the technology.[5]

The evidence points to selection and operating design. The portfolio leader still decides which customer delay matters most, which handoff should disappear, which control can be simplified, which decision lacks stable evidence, and which change deserves scarce reviewers and funding.

The cited studies support three bounded facts: a defined task can improve, reported enterprise impact remains uneven, and unclear intent, weak data or controls, cost, and unclear value contribute to project abandonment. The paper makes a portfolio inference from those facts.

Because cheap execution creates more candidates than the review can frame and compare, the first limiting asset is the judgment to turn one operating problem into a testable choice; this paper calls that judgment business imagination.

This definition is practical. It asks people with domain context to notice where a workflow, decision, control, product experience, or customer interaction could work differently, state why the change matters, and make the possibility testable against other demands for capacity. The review begins with

three questions: What work should change? What would improve if it did? What is the smallest credible way to find out?

Put one opportunity brief before the tool decision

The next operating move is small: require one opportunity brief before a material AI idea receives funding, scarce delivery capacity, or permission to become part of recurring work.

EXHIBIT 1

Frame the work before funding the tool

01 Name the work Identify the workflow, decision, control, handoff, or customer interaction that should change.	02 State value and evidence Record the baseline, value mechanism, and result that earns another commitment.	03 Set ownership and limits Name the outcome owner, reliance boundary, risk, and decision that remains human-owned.	04 Compare the routes Test workflow, model, automation, data, vendor, expert, training, wait, and stop options.
--	--	--	--

The brief should answer six questions in plain business language:

- What work or decision should change? Name the workflow, handoff, control, customer interaction, or decision. Record the tool category separately.
- What is happening now? Record the known baseline: delay, rework, error, cost, risk, missed coverage, or customer friction. If no baseline exists, say what must be observed before a value claim can be tested.
- How could value appear? Name the mechanism: work reduction, shorter cycle time, better quality, broader analytical coverage, improved decision quality, risk reduction, customer improvement, or revenue enablement.
- What evidence earns the next commitment? Define the result that would justify moving beyond exploration and the evidence that would cause the team to revise or stop.
- Who owns the result, and what can AI decide? Name the person accountable for the business result, what AI may draft, flag, classify, or recommend, and which decision stays human-owned.

- What route deserves comparison? Consider workflow redesign, a frontier or lower-cost model, deterministic automation, better data, a vendor capability, human expertise, training, a policy decision, waiting, or stopping.

Together, the six answers create one record for the allocation decision. Discovery stays light: a prompt that saves an hour can enter, alongside a customer complaint, a control failure, a recurring analyst bottleneck, or a vendor proposal. The review compares the work, the evidence, and the capacity each idea would consume.

Consider a proposal to “improve portfolio reporting with AI.” That phrase is too loose to fund. The opportunity brief might reveal a narrower problem: analysts spend time reconciling inconsistent fields before an executive review, and the late reconciliation delays the decision. The value hypothesis is shorter preparation time and a more stable evidence pack. The team should compare data validation and a deterministic rule before proposing a model. The route comparison prevents premature model selection.

Compare the routes that fit the work

The brief gives the portfolio a decision it can make. The routes include far more than model access.

Some opportunities justify a short frontier-model exploration because the possible value is high and the answer is uncertain. Known, repeatable, low-risk work may fit a lower-cost model. A fixed rule may belong in a deterministic script. Weak source truth may require data cleanup before any model test. A vendor feature may be the fastest path when its operating limits fit the work. A high-stakes judgment may need a human expert and a decision record. A weak opportunity should wait or stop.

The review assigns people, models, data, automation, workflow change, governance attention, and funding according to the work's value, uncertainty, risk, repeatability, and quality of evidence. Each request receives a route that the portfolio can explain and revisit.

Tool-led discovery and low-cost experiments remain welcome. The opportunity brief becomes consequential when capacity moves or the business begins to rely on the result. Evidence and control should grow with that commitment.

Connect the pilot to the business result

The brief remains attached to the idea after approval. The team tests the value hypothesis set before delivery.

EXHIBIT 2

Five questions keep the value claim connected

- 1 **System performance**
Did it meet the task-specific acceptance criteria?
- 2 **Workflow adoption**
Did intended users put it into the work, and where did they override it?
- 3 **Operational change**
Did the named baseline move?
- 4 **Business consequence**
Did that change matter to the strategic outcome in the brief?
- 5 **Full cost and risk**
Did the result justify model, data, integration, review, change, and governance burden?

The evidence should connect five questions:

The customer-support study shows why this chain needs a boundary. Its result was credible because the researchers observed a specific workflow, a defined population, and work output.[6] The 14 percent finding applies to that setting. A portfolio should demand the same precision from its own claims.

NIST's AI Risk Management Framework organizes risk work through govern, map, measure, and manage across context and lifecycle.[7] The portfolio review should therefore record who may rely on the output, what remains human-owned, how failures will be detected, and who can pause or retire the use alongside the value evidence.

Indirect value can stay visible without being inflated. Faster decisions, broader evidence coverage, less context loss, stronger onboarding, or lower control exposure may be valuable leading indicators. They become financial claims only when the organization can show the connection. Uncertainty belongs in the record.

The next review needs one decision

At the next portfolio review, begin with the opportunity brief.

If the team can name the work, baseline, value mechanism, evidence test, owner, decision boundary, and competing routes, the review can explore, redirect, fund the next defined commitment, fold the idea into existing work, wait, or stop. An incomplete brief keeps the demonstration in discovery until the missing decision evidence is available.

The three opening teams now face the same test. The brief makes an unprototyped idea visible, keeps a demo in discovery until its business problem is clear, and advances a local gain when its evidence and limits are explicit.

The operating standard is one sentence: before funding the tool, require a testable account of the work worth changing, the value at stake, the evidence needed, the accountable owner, and the route that fits the job.

Notes and sources

[1] Alex Singla, Alexander Sukharevsky, Bryce Hall, Lareina Yee, Michael Chui, and Tara Balakrishnan, "The state of AI in 2025: Agents, innovation, and transformation," McKinsey & Company, November 5, 2025. Verified July 5, 2026. <https://www.mckinsey.com/capabilities/quantumblack/our-insights/the-state-of-ai>

[2] Alexis Krivkovich and Lucia Rahilly, "AI is everywhere. The agentic organization isn't—yet," The McKinsey Podcast, April 2, 2026. Verified July 5, 2026. <https://www.mckinsey.com/capabilities/people-and-organizational-performance/our-insights/ai-is-everywhere-the-agentic-organization-isnt-yet>

[3] Johannes-Tobias Lorenz, Joshan Cherian Abraham, Robert Levin, and Douglas Ziman, "From promise to impact: How companies can measure—and realize—the full value of AI," McKinsey & Company, April 24, 2026. Verified July 5, 2026. <https://www.mckinsey.com/capabilities/quantumblack/our-insights/from-promise-to-impact-how-companies-can-measure-and-realize-the-full-value-of-ai>

[4] Gartner, "Gartner Predicts 30% of Generative AI Projects Will Be Abandoned After Proof of Concept By End of 2025," July 29, 2024. Verified July 5, 2026. <https://www.gartner.com/en/newsroom/press-releases/2024-07-29-gartner-predicts-30-percent-of-generative-ai-projects-will-be-abandoned-after-proof-of-concept-by-end-of-2025>

[5] James Ryseff, Brandon F. De Bruhl, and Sydne J. Newberry, "The Root Causes of Failure for Artificial Intelligence Projects and How They Can Succeed: Avoiding the Anti-Patterns of AI," RAND Corporation, 2024. Verified July 5, 2026. https://www.rand.org/pubs/research_reports/RRA2680-1.html

[6] Erik Brynjolfsson, Danielle Li, and Lindsey R. Raymond, "Generative AI at Work," NBER Working Paper 31161, April 2023; revised November 2023. Verified August 3, 2026. <https://www.nber.org/papers/w31161>

[7] National Institute of Standards and Technology, "AI Risk Management Framework Core," excerpt from AI RMF 1.0, 2023. Verified July 5, 2026. <https://airc.nist.gov/airmf-resources/airmf/5-sec-core/>

About the author

Marco Policani is an enterprise portfolio, PMO, and AI operating-governance leader. He builds portfolio governance systems where AI carries the analytical load and named humans own every decision — an approach documented across case studies, walkthroughs, and working governance guides on his portfolio site.

Portfolio & governance library: policani.net/governance · Full portfolio: policani.net · LinkedIn: linkedin.com/in/marcpolicani

This white paper may be shared freely with attribution. © 2026 Marco Policani.