
WHITE PAPER · PORTFOLIO & AI GOVERNANCE

Compare AI Proposals on the Work They Change, Not the Demo

A one-page brief for comparing competing AI proposals

Marco Policani

Enterprise Portfolio · PMO · AI Operating Governance

policani.net · linkedin.com/in/marcpolicani

July 2026

EXECUTIVE SUMMARY

Before funding any of several competing AI proposals, ask each for the same one-page brief, written by the leader who answers for what it would change, and compare the proposals on the briefs, not the demos.

A demo shows that a tool works, a usage count that people tried it, and a local time saving that one task got faster. This paper reads these as answers to different questions, so a review needs one page per proposal to compare the proposals. It sets out the six questions on the page, who writes and checks it, and how a missing answer is handled. The chair gets a side-by-side page and a written reason for each funding decision.

- 01** The business owner, the leader accountable for the results of the workflow or decision being changed, writes the brief with help from the proposing team, so the problem is stated in business terms.
- 02** The six questions cover what should change, what is happening now, how value could appear, the evidence for the next step, the owner and what AI may decide, and the alternatives.
- 03** The brief does not score proposals. It shows which can be compared yet: a missing answer holds a proposal, with no people or budget, until the answer is supplied, and a weak answer counts against it.

Three proposals, one pool of capacity

You chair the portfolio review that decides which AI proposals receive scarce capacity: the people and budget that cannot cover every request. One team brings a working assistant. Another has a vendor demo. A third shows time saved on drafting and analysis. All three can produce something. None has shown which work should change, what result would matter, or why AI is the right investment.

This paper recommends that the review ask for the same one-page opportunity brief for every proposal and compare the proposals on the briefs, not on the demos. The business owner writes the brief. The business owner is the leader accountable for the results of the work the proposal would change, and the one who would have to

change how it is done. The brief answers six questions in plain business language: what work or decision should change, what is happening there now, how value could appear, what evidence would earn the next step, who owns the result and what AI may decide, and what else could solve the problem. The next step is a small funded piece of the proposal with a result stated in advance.

The brief does not produce a score. Its first job is to show which proposals can be compared yet, because a proposal with a missing answer cannot be ranked against one without. For the rest, the review chair orders the fundable proposals by judgment and writes down the reasons.

Comparing unlike proposals is an old problem in any review with competing requests, and the scene above is one example. Models can now draft, summarize, classify, search, and write code. In this paper's observation, that is why a team can show a working tool before it has said which work the tool should change. No source counts how often that happens. A prototype shows that a tool can do a task. It does not show that the task deserves more of the review's capacity than another team's problem. With no common basis for comparison, capacity can follow the best presentation, and an operating problem that arrives without a polished prototype can stay hidden. That too is this paper's inference, not a measured result. When capacity covers every request, nothing needs comparing.

Each section depends on the one before. The next shows what each kind of evidence a team brings can carry, which is why the review needs one page per proposal. The third says what goes on that page and who writes it, and the fourth says how the review reads several pages together. The objections come before the step to take next.

What each kind of evidence shows

A working demo, a usage count, and a local time saving each answer a different question. A demo shows that a tool can do a task on the cases it was shown. A usage count shows that people tried the tool. A local time saving shows that one task got faster for one group. None says how large the business problem is, who answers for it, or how the proposal compares with another. The sources below show where each of these stops.

EXHIBIT 1

What each kind of evidence can carry

- ☐ **Working demo**
A tool can do a task on the cases it was shown.
- ☐ **Usage count**
People tried the tool.
- ☐ **Local time saving**
One task got faster for one group at one company.
- ☐ **Not shown by any of them**
How large the problem is, who answers for it, or how one proposal compares with another.

A tool can help on a defined task, and a field study shows how large the gain was at one company. Researchers followed 5,179 customer-support agents at one company as an AI assistant was rolled out in stages. Access raised issues resolved per hour by 14 percent on average. The gain was 34 percent for novice and low-skilled workers and minimal for experienced and highly skilled workers.[1] The study measures what one tool did for one group doing one job at one company. It shows that a tool can help a defined workflow and that the gain differed by skill level. It does not measure which work a company should choose to change, and it does not say whether that workflow deserved capacity more than a different team's proposal. Output measured across thousands of agents is also a different kind of evidence from a team's own report of time saved.

Results across whole companies are reported less often than use. McKinsey's 2025 survey, in the field from June to July 2025, drew 1,993 respondents in 105 countries. Of them, 88 percent said their organizations regularly use AI in at least one business function, and respondents reported cost and revenue benefits from individual use cases. Just 39 percent attributed any level of earnings impact to AI at the enterprise level, and most of those said less than 5 percent of earnings before interest and taxes came from AI.[2] These are respondents' own reports, and the survey does not say why use-case benefits are not reported at the enterprise level more often. The third team's time saving could be real and still not appear in enterprise results.

The three teams in the opening are therefore not competing on one measure, and more demos, more usage, or more local savings will not put them on one. That is this paper's reading of what the sources show, and no source tests it. It is the reason the review needs one page on which every proposal answers the same questions. The six questions follow from the gaps above. The first two say how large the problem is, the fifth says who answers for it, and the rest give the review something to compare.

Practitioner research points the same way. Gartner reported in January 2026 that at least half of generative AI projects were abandoned after proof of concept by the end of 2025, citing poor data quality, inadequate risk controls, escalating costs, or unclear business value. It names lack of business value as the first failure point. Organizations that chase flashy demos or deploy everywhere at once, Gartner says, spread resources across low-impact initiatives. Without a clear way to rank use cases or measure success, their projects lack measurable business value and become easy targets when budgets tighten. Gartner's advice is to build a rigorous way to prioritize use cases.[3] Gartner says it analyzed hundreds of implementations and shows no method for the figure. It does not say that a comparison process would have changed any outcome.

One brief for every proposal

The opportunity brief is one page. The business owner writes it, with help from the proposing team, before the review meets. It is the record of one funding decision, and the business owner updates it at each later review of that proposal. When no one owns the work a proposal would change, the review chair names an owner or holds the proposal.

The business owner writes it, not the proposing team alone. RAND interviewed 65 experienced data scientists and engineers about failed AI projects, and they most often named misunderstanding or miscommunication about a project's intent and purpose as the reason. RAND tells business leaders to help technical staff understand what the project must achieve and how its result will be used.[4] Most interviewees were engineers rather than executives, so RAND says the results may be skewed toward identifying leadership failures. The inference here goes beyond RAND's finding: the person who owns the work is best placed to state the problem in business terms.

The companion paper on reviewing one AI request describes how a business owner names the work and its limiting step before a proposal reaches a review. A proposal

that has been through that intake already has most of these answers. The brief is their one-page form, and the review does not ask twice.

The brief answers six questions:

- What work or decision should change? Name the workflow, handoff, control, customer interaction, or decision. If the proposal comes with a tool, name it on a separate line after the work.
- What is happening now? Record the baseline: the known delay, rework, error, cost, risk, or customer problem today, in a measure the business already uses. If no baseline exists, say what the team must first observe.
- How could value appear? Name the expected path: less work, a shorter cycle, better quality, wider coverage, a better decision, lower risk, or added revenue.
- What evidence earns the next step? State the result that would justify more investment, and the result that would make the team revise or stop. Together these are the evidence test.
- Who owns the result, and what may AI decide? Name the person accountable for the business result: the business owner, or whoever the chair names. State what AI may draft, flag, or recommend, and which decisions people keep. The National Institute of Standards and Technology (NIST) AI Risk Management Framework asks organizations to define who is responsible when people and AI work together and who oversees the AI, and to document how human oversight works.[5]
- What else could solve the problem? Name at least one alternative that does not start with a new tool, such as a change to the process, better data, training, or waiting. A companion paper on choosing the response to an AI request compares such alternatives in depth.

On the fourth question, early evidence can show technical performance and use. Financial impact needs an operating result first. McKinsey's article on measuring AI value separates five levels of measurement: technical performance, user adoption, operating results, strategic outcomes, and financial impact. That is a consultancy's approach, not a tested standard.[6] The companion paper on reviewing one AI request covers the value review after a step is funded.

As planning estimates, the brief takes about an hour of the business owner's time, plus the time to measure a baseline when none exists. The analyst needs a few hours to assemble the page for a review, and the outside reviewer described in the next section needs about an hour for each proposal. The chair needs about ten minutes for each proposal in the meeting. No source prices these. In this paper's judgment, a

brief that takes far longer than an hour may mean the work is not yet understood, which the review needs to know.

How the review compares proposals

The review compares proposals in three steps. The portfolio analyst, usually from the portfolio or project management office, lays the briefs side by side on one page, with the six questions as rows and the proposals as columns, and marks each answer as given or missing. An answer is missing when it names no work, measure, owner, or result that a reviewer could check. A reviewer who is not on any proposing team checks each baseline and each evidence test. The review chair then gives each proposal one of four outcomes: fund the next step, hold it with the missing answers named, combine it with other work, or stop it.

EXHIBIT 2

Four outcomes for each proposal

Fund the next step

The answers are given and the evidence holds up. The step fits the available people and budget.

Hold, naming the gaps

An answer is missing. No people or budget until it is supplied; low-cost exploration can go on.

Combine with other work

Two briefs name the same work, so the review treats them as one proposal.

Stop

The evidence runs against the proposal, or the problem is too small for the people and budget it would draw.

The analyst does not judge value. The reviewer, such as a finance or operations leader, asks whether each baseline is a measurement or an estimate, who took it, and whether the evidence test could come out against the proposal. The business owner and the proposing team do not check their own page, which matters most when the business owner is also the person proposing the tool.

The chair then decides, and no formula applies. A score would have to weigh days of delay against errors, cost, and risk, and that weighing is the judgment the chair would hide inside the arithmetic. This is the paper's reasoning, and no source tests it. The

chair asks three questions of each proposal. How large is the problem in business terms, as the business owner states it in money, customers, risk, or time? How far does the evidence support the claim? How much of the capacity would the next step draw? The chair also asks each business owner what result would make them stop, because a proposal with no answer has not said what it is testing. The chair writes one line of reasoning for each decision and keeps it with the brief, so that a later review can see which judgments held up.

A missing answer differs from a weak one. A weak answer, such as a small expected value or thin evidence, counts against a proposal when the chair weighs it. A missing answer, such as no baseline, no owner, or no named evidence, means the proposal cannot be compared yet. It is held with the missing answers named, receives no scarce capacity until they are supplied, and can continue as low-cost exploration. Two briefs that name the same work are combined. A proposal is stopped when its evidence runs against it or its problem is too small for the capacity it would draw. A good proposal can also wait because capacity ran out.

Suppose the three teams file briefs. This is an illustration, not a case, and it reports no result. The assistant team's brief names the weekly status report and its owner and expects a shorter preparation cycle, but gives no baseline. The vendor-demo team's brief names customer complaint handling, gives the days a complaint waits for a first response, names its owner, says the tool would suggest a category while people decide, and names a change to routing rules as the alternative. The drafting-and-analysis team's brief names no particular workflow or measure and says only that analysts would save hours. Its owner cannot say what the freed time would be used for. The review can hold the first proposal with one named question: how long does the report take to prepare today? Answering it costs the people who do that work some days of measurement and no tool. The review can fund a next step for the second proposal. It holds the third with two missing answers named: the work to change and what the freed time would be used for. The reasoning rests on what each brief answers, and a demo does not decide the order.

Objections, and what the sources do not show

The brief slows exploration

The brief applies when scarce capacity moves or the business starts to depend on the result. A low-cost trial by one team, using an approved tool on its own work, needs no brief. Exploration can find uses that nobody could describe beforehand, and no brief can be written for work nobody has seen. The review chair sets the

threshold for scarce capacity, such as an amount of budget or a number of staff days, and no source sets it. The companion paper on reviewing one AI request describes a capped allowance for exploration, set by the same group that runs this review. The brief pays when the step it controls costs well more than the brief and its checking. For a trial that costs less, the paperwork costs more than the experiment.

A business case already does this

A review may already ask for a business case. The brief differs in three ways, and these are this paper's design, not a tested improvement. The business owner writes it before a tool is chosen. A reviewer outside the proposing teams checks it. A missing answer holds the proposal, where a scored business case would only lower its score. Asking every proposal the same six questions in the same order is ordinary practice, and it is what lets the proposals sit side by side. A business case that already answers the six questions needs no second document. The analyst can use it and mark it up.

A polished page or an optimistic baseline can win

A polished page can win instead of a polished demo, and an owner can write a baseline in a hopeful light. The analyst's marking shows whether an answer is given, not whether it is good. The check against this is the reviewer outside the proposing teams, who asks whether the baseline was measured and who took it. In a small organization with no such person, the chair asks the leader whose own proposals are not on the agenda. This reduces the risk and does not remove it.

The chair's judgment is not audited

With no formula, the chair is free to favor a preferred proposal. The written line for each decision, kept with the brief and read at the later review, is the limit. It does not remove the discretion. It makes a pattern of favoritism visible after the fact.

The estimates in a brief are guesses

A value path and an expected result are guesses before the work starts. The brief records the guess and the result that would prove it wrong, so the review can test it later. It does not make the estimate accurate.

What the sources do not show

No source tests the brief or compares reviews that use one with reviews that do not. The sources show that one tool helped one workflow at one company, and that fewer

survey respondents reported enterprise earnings effects than reported use. They show that Gartner names the lack of a clear way to rank use cases and measure success among the causes of abandoned projects, and that RAND's interviewees most often named miscommunication about a project's purpose. They do not show that comparing proposals on a common page would have prevented any failure.

That a demo, a usage count, and a local saving answer different questions is this paper's reading. The six questions, the rule on missing answers, who writes and who checks the brief, and the time estimates are this paper's design. NIST's framework is voluntary and does not prescribe a brief. A different review may weigh the questions differently, which is why the chair decides.

What to do next

Before the next portfolio review, ask the business owner of each AI proposal on the agenda to write the one-page brief. Give them the six questions, and ask that no brief name a tool before it names the work. The analyst lays the briefs side by side and marks each answer as given or missing before the meeting. At the meeting, the chair decides for each proposal whether to fund the next step, hold it with the missing answers named, combine it with other work, or stop it, and writes one line saying why.

Expect missing answers the first time. The list of them is the first result, and what it buys is a review that can say why one proposal received capacity and another did not. If the briefs arrive complete and the review decides what it would have decided anyway, narrow the requirement to the proposals that draw the most capacity.

Notes and sources

[1] Erik Brynjolfsson, Danielle Li, and Lindsey R. Raymond, "Generative AI at Work," NBER Working Paper 31161, April 2023; revised November 2023; published in *The Quarterly Journal of Economics* 140(2), 2025, 889-942.

<https://www.nber.org/papers/w31161>

[2] Alex Singla, Alexander Sukharevsky, Bryce Hall, Lareina Yee, Michael Chui, and Tara Balakrishnan, "The state of AI in 2025: Agents, innovation, and transformation," McKinsey & Company, November 5, 2025.

<https://www.mckinsey.com/capabilities/quantumblack/our-insights/the-state-of-ai-2025>

- [3] Arun Chandrasekaran, "Why 50% of GenAI Projects Fail — And How to Beat the Odds," Gartner, January 26, 2026. <https://www.gartner.com/en/articles/genai-project-failure>
- [4] James Ryseff, Brandon F. De Bruhl, and Sydne J. Newberry, "The Root Causes of Failure for Artificial Intelligence Projects and How They Can Succeed: Avoiding the Anti-Patterns of AI," RAND Corporation, 2024. https://www.rand.org/pubs/research_reports/RRA2680-1.html
- [5] National Institute of Standards and Technology, "AI Risk Management Framework Core," excerpt from AI RMF 1.0, 2023. <https://airc.nist.gov/airmf-resources/airmf/5-sec-core/>
- [6] Johannes-Tobias Lorenz, Joshan Cherian Abraham, Robert Levin, and Douglas Ziman, "From promise to impact: How companies can measure—and realize—the full value of AI," McKinsey & Company, April 24, 2026. <https://www.mckinsey.com/capabilities/quantumblack/our-insights/from-promise-to-impact-how-companies-can-measure-and-realize-the-full-value-of-ai>

About the author

Marco Policani is an enterprise portfolio, PMO, and AI operating-governance leader. He builds portfolio governance systems where AI carries the analytical load and named humans own every decision — an approach documented across case studies, walkthroughs, and working governance guides on his portfolio site.

Portfolio & governance library: policani.net/governance · Full portfolio: policani.net · LinkedIn: linkedin.com/in/marcpolicani

This white paper may be shared freely with attribution. © 2026 Marco Policani.