

Business Imagination Is the Scarce Asset Now

A portfolio-governance standard for turning cheap AI execution into measurable business value

Marco Policani

Enterprise Portfolio · PMO · AI Operating Governance

policani.net · linkedin.com/in/marcpolicani

July 2026

EXECUTIVE SUMMARY

Before funding an AI tool, require a short, testable case naming which work should change, what it is worth, how you will know, who owns it, and why this route beats the alternatives.

A portfolio review can see more AI proposals than it can judge. This paper calls the missing skill business imagination: turning an operating problem into a choice a portfolio can compare, test, and fund. It makes that skill usable through one opportunity brief placed before the tool decision.

- 01** Treat a demo, usage count, or local time saving as evidence about a tool or task, not proof that an investment deserves priority.
- 02** The brief also records today's baseline and what AI may do versus what people must still decide.
- 03** Keep the brief attached after approval so later reviews can advance, redirect, wait, or stop the investment against the test the brief set before delivery.

The review has more ideas than decisions

You chair the portfolio review that decides which AI ideas receive scarce capacity. One team brings a working assistant. Another has a vendor demo. A third shows time saved on drafting and analysis. All three can produce something. None has shown which work should change, what result would matter, or why AI is the right investment.

Cheaper execution is useful. Models can draft, summarize, classify, compare, search, translate, and code. A well-defined tool can also improve real work. In one field study, customer-support agents who used an AI assistant raised productivity by 14 percent on average. Novice and lower-skilled workers gained the most. The tool had little effect on the most experienced workers.[6]

The picture is less clear at enterprise level. McKinsey's 2025 survey found that 88 percent of respondents used AI in at least one business function. Only 39 percent reported an effect on earnings before interest and taxes.[1] In a 2026 McKinsey Podcast, senior partner Alexis Krivkovich said more than 80 percent of companies report no bottom-line gain yet from their AI investments.[2]

These findings expose the choice facing a portfolio leader. A tool may help in one defined task. That does not prove the company chose the right work to change. A demo, usage count, or local time saving can push a visible idea ahead before its business value is clear. Capacity then follows the best presentation. Better operating problems may stay hidden because they lack a polished prototype.

What becomes scarce when producing one more plausible output is easy? When one input gets cheap, value shifts to what is still hard. For AI, that hard part is often judgment: which customer problem is worth solving, which work should change, and which result would count as progress.

This paper calls that judgment business imagination. It is not free-form creativity. It is the practical skill of turning an operating problem into a choice that a portfolio can compare, test, and fund.

The scarce judgment comes before the tool

A portfolio needs evidence about the work and its result. Counts of users, prompts, documents, code suggestions, and pilots may all rise. Cycle time, quality, customer experience, risk, and financial results may not move.

McKinsey's 2026 measurement model separates five layers: technical performance, user adoption, operating results, strategic outcomes, and financial impact.[3] Each layer answers a different question. Does the system work? Do people use it in the workflow? Does the work improve? Does that improvement matter to the business? Do the full costs still support the case?

Failure can start before a pilot. In 2024, Gartner predicted that at least 30 percent of generative AI projects would be dropped after proof of concept by the end of 2025. It cited poor data, weak risk controls, rising costs, and unclear business value.[4]

RAND found a related problem. Its 2024 study interviewed 65 experienced data scientists and engineers. They most often pointed to confusion or poor communication about a project's intent and purpose as the reason AI projects failed. RAND advised leaders to focus on the problem before the technology.[5]

The sources support three limited facts. A defined task can improve. Reported enterprise impact remains uneven. Weak intent, data, controls, cost, and value can put projects at risk. The sources do not prove that business imagination causes value. That is the portfolio lesson this paper draws from them.

Cheap execution creates more ideas than a review can frame and compare. The first limit is often the judgment needed to turn one operating problem into a testable choice. That is where business imagination becomes useful.

Ask three questions before discussing a tool:

- What work should change?
- What would improve if it did?
- What is the smallest credible way to find out?

Put one opportunity brief before the tool decision

Require one short opportunity brief before a material AI idea receives funding or scarce delivery capacity. Use it as well before the result becomes part of recurring work.

EXHIBIT 1

Frame the opportunity before funding the tool

1. What should change

Name the workflow, handoff, control, or customer interaction.

2. What is happening now

Record the known delay, rework, error, cost, or risk.

3. How value could appear

Less work, a shorter cycle, better quality, wider coverage, or lower risk.

4. Evidence for next step

State what would justify more investment, and what would make the team stop.

5. Owner and AI limits

Name the outcome owner; say what AI may do and which decisions people keep.

6. Routes to compare

Workflow change, model, automation, data, vendor, expert, training, waiting, or stopping.

The brief should answer six questions in plain business language:

- What work or decision should change? Name the workflow, handoff, control, customer interaction, or decision. List the tool type separately.
- What is happening now? Record the known delay, rework, error, cost, risk, missed coverage, or customer problem. If no starting point exists, say what the team must first observe.
- How could value appear? Name the expected path. It may be less work, a shorter cycle, better quality, wider coverage, a better decision, lower risk, a better customer result, or possible revenue gain.
- What evidence earns the next commitment? State the result that would justify more investment. Also state what would make the team revise or stop.
- Who owns the result, and what can AI decide? Name the person accountable for the business result. State what AI may draft, flag, classify, or recommend. State which decisions people must still make.
- What routes should the team compare? Consider a workflow change, a model, fixed-rule automation, better data, a vendor tool, human expertise, training, a policy choice, waiting, or

stopping.

Together, these answers create one record for the funding choice. Discovery can stay light. A prompt that saves an hour can enter beside a customer complaint, a control failure, an analyst bottleneck, or a vendor proposal. The review compares the work, the evidence, and the capacity each idea would use.

Consider a proposal to “improve portfolio reporting with AI.” That phrase is too loose to fund. The brief may reveal a narrower problem. Analysts spend hours fixing inconsistent fields before an executive review. The late work delays the decision. The team expects a shorter preparation cycle and a more reliable report. It should compare better data checks and a fixed rule before it proposes a model.

Compare the routes that fit the work

The brief gives the portfolio a choice it can make. Model access is only one route.

A short test with a leading model may fit a valuable question with an uncertain answer. Stable, repeated, low-risk work may fit a cheaper model or fixed-rule automation. Poor source data may need repair before any model test. A vendor feature may be fastest when its limits fit the work. A high-stakes judgment may need a human expert and a written decision. A weak idea should wait or stop.

The review assigns people, models, data, automation, workflow change, control work, and funding to fit the job. It weighs value, uncertainty, risk, repeatability, and the strength of the evidence. Each request gets a route the portfolio can explain and revisit.

Low-cost experiments remain welcome. The brief begins to matter when capacity moves or the business starts to rely on the result. Evidence and controls should grow with that reliance.

Connect the pilot to the business result

Keep the brief with the idea after approval. The team should test the value claim it made before delivery.

EXHIBIT 2

Five questions keep the value claim connected

- 1 System performance**
Did it meet the task-specific acceptance criteria?
- 2 Workflow adoption**
Did intended users put it into the workflow, and where did they override it?
- 3 Operational change**
Did the named baseline move?
- 4 Business consequence**
Did that change matter to the strategic outcome in the brief?
- 5 Full cost and risk**
Did the business result justify model, data, integration, review, change, and control costs?

Use five questions:

The customer-support study shows why claims need clear limits. The researchers observed one workflow, one defined group, and its work output.[6] The 14 percent gain applies to that setting. A portfolio should be just as precise about its own results.

NIST groups AI risk work into four functions: govern, map, measure, and manage. Teams use them from design through retirement and adapt them to the system's setting.[7] A portfolio review should therefore record who may rely on an output, which decisions stay with people, how the team will find failures, and who can pause or retire the use.

Indirect value can remain visible without becoming a financial claim too soon. Faster decisions, wider evidence, less lost context, better onboarding, or lower control risk may be useful early signs. The company should claim financial value only when it can show the link. Keep the uncertainty in the record.

The next review needs one decision

At the next portfolio review, begin with the opportunity brief.

Can the team name the work, starting measure, path to value, evidence test, owner, human decisions, and other routes? If so, the review can explore, redirect, fund the next defined step, fold the idea into existing work, wait, or stop. If not, keep the demo in discovery until the missing evidence is ready.

The three opening teams now face the same test. The brief makes an idea without a prototype visible. It keeps a demo in discovery until the business problem is clear. It advances a local gain only when its evidence and limits are clear.

Before funding the tool, ask for a testable account of the work worth changing. Require the team to name the value, evidence, owner, human judgment, and route that fit the job.

Notes and sources

[1] Alex Singla, Alexander Sukharevsky, Bryce Hall, Lareina Yee, Michael Chui, and Tara Balakrishnan, "The state of AI in 2025: Agents, innovation, and transformation," McKinsey & Company, November 5, 2025. Verified July 5, 2026. <https://www.mckinsey.com/capabilities/quantumblack/our-insights/the-state-of-ai-2025>

[2] Alexis Krivkovich and Lucia Rahilly, "AI is everywhere. The agentic organization isn't—yet," The McKinsey Podcast, April 2, 2026. Verified July 5, 2026. <https://www.mckinsey.com/capabilities/people-and-organizational-performance/our-insights/ai-is-everywhere-the-agentic-organization-isnt-yet>

[3] Johannes-Tobias Lorenz, Joshan Cherian Abraham, Robert Levin, and Douglas Ziman, "From promise to impact: How companies can measure—and realize—the full value of AI," McKinsey & Company, April 24, 2026. Verified July 5, 2026. <https://www.mckinsey.com/capabilities/quantumblack/our-insights/from-promise-to-impact-how-companies-can-measure-and-realize-the-full-value-of-ai>

[4] Gartner, "Gartner Predicts 30% of Generative AI Projects Will Be Abandoned After Proof of Concept By End of 2025," July 29, 2024. Verified July 5, 2026. <https://www.gartner.com/en/newsroom/press-releases/2024-07-29-gartner-predicts-30-percent-of-generative-ai-projects-will-be-abandoned-after-proof-of-concept-by-end-of-2025>

[5] James Ryseff, Brandon F. De Bruhl, and Sydne J. Newberry, "The Root Causes of Failure for Artificial Intelligence Projects and How They Can Succeed: Avoiding the Anti-Patterns of AI," RAND Corporation, 2024. Verified July 5, 2026. https://www.rand.org/pubs/research_reports/RR2680-1.html

[6] Erik Brynjolfsson, Danielle Li, and Lindsey R. Raymond, "Generative AI at Work," NBER Working Paper 31161, April 2023; revised November 2023. Verified August 3, 2026. <https://www.nber.org/papers/w31161>

[7] National Institute of Standards and Technology, "AI Risk Management Framework Core," excerpt from AI RMF 1.0, 2023. Verified July 5, 2026. <https://airc.nist.gov/airmf-resources/airmf/5-sec-core/>

About the author

Marco Policani is an enterprise portfolio, PMO, and AI operating-governance leader. He builds portfolio governance systems where AI carries the analytical load and named humans own every decision — an approach documented across case studies, walkthroughs, and working governance guides on his portfolio site.

Portfolio & governance library: policani.net/governance · Full portfolio: policani.net · LinkedIn: linkedin.com/in/marcpolicani

This white paper may be shared freely with attribution. © 2026 Marco Policani.