WHITE PAPER · PORTFOLIO AND AI GOVERNANCE

The faster your review steps in, the better its rescues look

Projects picked for a bad quarter partly bounce back without the rescue. So the more readily your review steps in, the better its rescues look, whether or not they help.



Marco Policani

Portfolio · Program · PMO

October 2026

policani.net/governance

THE CENTRAL ARGUMENT

Many portfolio reviews step in each quarter where a project's numbers turn bad, and count the projects that come back as rescues that worked. But a project picked for a bad quarter is picked partly for swing that does not last. So it partly recovers, whatever the rescue did. The more readily a review steps in, the higher its turnaround rate. A stricter rule lowers it. Both hold whether or not the rescues help. The review's own line for stepping in holds the makings of a fairer test: compare projects just past the line with projects just short of it. That works as long as teams cannot steer their numbers around the line.

Each quarter your portfolio review looks at the latest numbers. When a project's numbers turn bad and the team names a cause, you step in with a recovery plan, more people or a new lead. Many of those projects come back. The share that come back is your turnaround rate, and it is part of why you trust the review. You may have moved from a half-year review to a quarterly one, or lowered the bar for stepping in. If the rate then went up, you may have taken that as proof the change worked.

It is not proof. A project's number moves around from quarter to quarter in ways that do not last. Call that swing. When you pick the projects with the worst numbers, you pick partly for swing. So those projects partly recover, whatever the rescue did. The more readily you step in, the more of your picks owe their bad number to swing. That pushes your turnaround rate up. Make the rule stricter and the rate falls, even if your rescues work just as well.

A study of English football shows how strong this effect is. Premier League clubs that replaced their manager after a poor run gained, on average, 0.21 points per match over their results before. In seasons when the same clubs had a similar poor run and changed nothing, they also gained 0.21. Control limits and rules such as "red for two quarters" help. They do not remove the effect.

You do not have to stop stepping in during the year. You have to stop judging your rescues by how many projects came back. The same goes for any change to your review. Your review already holds the makings of a fairer test. Projects just past your line for stepping in looked almost the same as projects just short of it. Yet most of the first group got help and most of the second did not. Over enough years, plot how each did afterward against where it stood when you decided. Then look for a jump at the line. That

jump shows what stepping in at the line does, compared with carrying on under your normal rule. That holds only as long as teams cannot steer which side of the line they land on.

ARGUMENT AT A GLANCE

What the turnaround record can and cannot show

What you pick

The worst number on the list, which is partly swing that will not last.

What the record shows

Recoveries that would partly have come anyway, more of them the faster you step in.

What to check

Whether projects just past your line did better than projects just short of it.

Applies to moves meant to fix a running project, chosen by a number that swings a lot compared with any lasting change.

Why you step in during the year

Moving resources can pay. Consultants at McKinsey, writing with the strategy professor Dan Lovallo, studied more than 1,600 US companies from 1990 to 2005 [1]. A third of business units got almost the same capital every year. Companies in the top third for moving capital between units earned, on average, 30 percent higher annual returns to shareholders than those in the bottom third. But over periods shorter than three years, the companies that moved more did worse. The authors suggest investor caution and how closely the moves follow strategy as possible reasons. They advise moving capital "consistently over the medium to long term in service of a clear corporate strategy" [1].

They did not study quarterly rescues. But rescues come from the same push to move resources rather than leave them in place. Once you move them during the year, you need to know which moves work. And the evidence you have is what happened after each one.

How much of a bad quarter lasts

Cost tracking on US defense contracts shows how much of a quarter's movement lasts. David Christensen and Scott Heise tracked cost efficiency on 155 contracts, read each quarter [2]. Cost efficiency is the budgeted value of the work done for each dollar spent.

Over the second half of each contract, the three-month figure's high and low were, on average, more than a full point apart (1.07). The running total since the start, on average, barely moved (0.07). It stayed within a narrow band in 153 of the 155 contracts [2].

Suppose a project's real efficiency had changed as much as its quarterly figure did, and stayed changed. Then the running total would have moved far more. So most of the quarterly movement did not last. That is this paper's reading of their figures, not the authors' claim.

Three things limit that reading. A few wild contracts pull the average spread up. A change that starts late barely shows in a running total. And these contracts were managed the whole time. So the figures show swing under normal management, not recovery with nobody doing anything.

Not all of a bad quarter is swing. The authors found that the running total usually drifted down. So part of a dip can be a lasting slide. That is why you read a dip against the project's own trend, not against a fixed level.

Why the projects you pick will partly recover anyway

Now apply the swing to a review. You look down the list and step in where the latest number is worst. A number that is low partly because of swing is more likely to land at the bottom than one that is low only because of lasting problems. So part of what makes your picks look bad is swing. Next quarter the swing usually does not repeat. On average, the number partly recovers whether or not your rescue did anything.

Statisticians call this regression to the mean: ordinary ups and downs that look like real change when you follow up only the cases you picked for a bad number. Adrian Barnett and colleagues describe it as an effect that "can make natural variation in repeated data look like real change." It grows, they write, "with increasing measurement error and when follow-up measurements are only examined on a sub-sample selected using a baseline value" [3]. That is exactly what a review does when it follows up only the projects it picked for a bad number, on the same kind of number.

What happened when football clubs changed managers

You may know the idea. Two things are less clear. How much does it matter when people make real decisions? And does a longer run of results escape it? Football is a useful test, because results are public and counted the same way every week. Lucas Besters, Jan van Ours and Martin van Tuijl studied 15 seasons of the English Premier League [4]. They matched 61 in-season manager changes with seasons of the same club in which nobody was replaced. Each pair was at a similar point in the season, with a similar shortfall against what bookmakers' odds had predicted.

After a change, points per match rose by 0.21. In the matched seasons with no change, they also rose by 0.21. Outright dismissals looked much the same: 0.28 after a dismissal against 0.26 in the matched seasons. The difference was not statistically significant. In the authors' words, "performances would have improved irrespective of the manager in charge" [4].

Two details matter for a portfolio. First, the matching used the shortfall built up over the season so far, not one bad match. Seasons with a similar shortfall and no change recovered just as much. So, on this evidence, even a shortfall built up over many matches held enough swing to produce an apparent recovery. A longer stretch of results helps only if it actually washes the swing out.

Second, the authors explain why clubs keep changing managers. Results after a change are often better than before, so the change looks like a success. A before-and-after comparison that ignores what would have happened anyway is misleading, they write. But that is how owners, fans and the press see it. And acting and failing looks better than failing without acting. From their case studies and the spread of results, the authors judge that some changes helped and some hurt. They found no pattern among the traits they tested [4].

The study has limits. The comparison seasons had no change at any point. So a season that kept sliding is less likely to be among them. That may make the seasons without a change look better than they were, and hide some benefit from the changes. Those seasons may still have had other action, such as new signings. And bookmakers' odds are a sharper forecast than most project plans, so more of a project's shortfall may be lasting.

Football is not a portfolio. The comparison holds under three conditions. A project's numbers swing a lot compared with any lasting change. The review picks projects by those numbers. And the move is meant to fix the project. Under those conditions, your turnaround rate comes from a before-and-after record, the kind the authors call misleading.

Why a quicker trigger makes the record look better

Three separate choices shape how readily a review steps in. One is how long a period each number covers. Another is how bad the number must be before you act. The third is how often you look. Each changes which projects you pick.

A number over a shorter period swings more. In the defense data, where both were read each quarter, the three-month figure spread about twice as far as a six-month average (1.07 against 0.53) [2]. Looking more often gives each project more chances to post one bad number. A six-month average updated each quarter spread 0.53. A six-month figure

read only every six months spread 0.42 [2]. A lower bar picks projects closer to acceptable, and they need less recovery to get back there. Where most projects are sound, a lower bar also adds sound projects having a bad patch.

Each of these choices raises the share of picked projects that come back. A stricter rule lowers it. That holds even when the rescues work exactly as well. No study here changed a review's trigger and watched its turnaround rate. This step is reasoning from the defense figures and the statistics above.

So suppose you moved to a quicker cycle or a lower bar and saw more turnarounds. That alone cannot show the change worked. It is what you would see if your rescues did nothing at all. If you tighten the rule and see fewer, that alone cannot show the change failed. A quicker cycle may catch real trouble sooner, and that is a benefit. The turnaround rate just cannot show it.

Control limits and a "red for two quarters" rule screen out some swing. But even when a project's trouble is real, the number that got it picked is worse than the trouble itself. Part of what follows is that extra drop wearing off. These rules reduce the effect. They do not make the record evidence.

These choices have costs, and the costs do not show in the record. People and money go to rescues that may not have been needed. The people you move leave other work behind. And since the record cannot tell you whether your rescues or your review changes help, you cannot get better at either.

"But a bad quarter can be a real warning"

It can. Some signals last. One is a change you can confirm without the numbers and expect to last: a customer lost, a test failed, a strategy changed. Another is a cost overrun to date on a contract past its early stages. In the defense data, that figure usually got worse from there [2]. Do not wait on those. The point is to tell them apart from one quarter's swing. Running totals on other measures will not always behave like the defense running total. And the football shortfall shows that a total can still hold enough swing to mislead.

You may say you act on a diagnosis, not a number. But after a bad quarter, a cause can almost always be named. Ask whether you could confirm that cause without looking at the dip.

Your review may use only red, amber and green ratings. Then the argument applies as far as the colors follow a number that swings. Colors set by judgment may lag rather than bounce around, and then less of a red rating comes from swing. That weakens the effect. It still does not make the turnaround rate evidence.

And your team may simply be good at rescues. Some teams may be. The football study allows that some owners may spot the changes that will work [4]. The before-and-after record cannot show which.

How to tell whether stepping in works

A fair test compares projects that looked alike when you decided but were handled differently. The reason for the difference must not predict recovery. Your own line for stepping in supplies that reason. Take every project just past the line and every project just short of it. They had nearly the same number, so they carry about the same swing. The line, not your guess about who would recover, decided which side each fell on.

The method is well established in economics. David Lee and Thomas Lemieux trace it to a 1960 study of students just above and just below the cutoff for a merit award [5]. Near the cutoff, they write, which side a case lands on is "randomized as though from a randomized experiment," as long as no one can "precisely manipulate" it [5]. Applying it to a portfolio review is this paper's own step.

Four choices keep the test fair. Set all of them before you look at any results.

First, group projects by side of the line, not by what you did. Most projects past the line will have been rescued and most short of it not. Projects where you overrode the line stay in their side's group. Check that rescue really was much more common just past the line. If it was not, the line did not decide, and the comparison shows nothing.

Second, count each dip once, at the first number in that dip that crossed or came near the line. Say in advance what ends a dip, such as a return to an acceptable number, so that a later dip in the same project counts as a new one. A project just short of the line may be rescued at a later review. It stays in its group. So the test compares stepping in now with carrying on under your normal rule, not rescue with no rescue. If many projects short of the line are rescued a review later, a small jump may only mean that stepping in one review earlier adds little. It does not mean rescues do nothing.

Third, fix the outcome in advance so both groups can fail it. Use the original measure, against the original baseline, at a follow-up point set before any results are in. A rescue that resets the plan cannot earn credit that way. Projects stopped or re-scoped during follow-up stay in their group, under a rule you state in advance.

Fourth, look for a jump rather than comparing two averages. You have to choose how far from the line a project can be and still count. With few projects, that band has to be wide. Then projects past the line have worse numbers on average than those short of it, and a plain comparison of the two groups is biased [5]. Instead, plot each project's later result against its number when you decided, and see whether results jump at the line. Check that the answer holds with a narrower and a wider band.

A jump shows everything that switches on at the line together: the red status, the board's attention and the rescue. Lee and Lemieux call it "the combined effect of all factors that switch on at the threshold" [5]. That combined effect is still what a board deciding whether to step in needs to know.

The test fails if teams can steer which side they land on. Numbers near a line are often partly produced by the team, such as percent complete or a self-set status. A team that does not want a rescue may hold its number just short. To check, look for a pile-up of numbers just short of the line. Also check that things set before the decision look alike on both sides, such as the size, age and type of project [5]. With few projects, these checks are weak reassurance, not proof.

Judge a proposed change to your line the same way, never by the turnaround rate. Changes to how often you look, or to how long a period each number covers, are not tested at the line. For those, this paper claims only that the turnaround rate cannot judge them. The test has limits of its own. It answers only for projects near the line, which is where a board deciding whether to move its line needs an answer. It needs a numeric trigger applied as a rule, so it does not work where the board steps in on colors set by judgment. It needs years of dips, not quarters. And each change to the line starts a new count, so a new trigger can be judged this way only years later.

Picking less swing in the first place

You can also make fewer picks on swing. Before stepping in, check a dip against the project's own trend and its usual swing. Then wait one more review, unless there is a change you can confirm without the numbers and expect to last.

That has a price. A real slide with no outside cause gets one more review before anyone acts. In the defense data, the running total settled early and then usually got worse. From that, this paper infers that a slide left alone early can lock in much of the final overrun. So, on that inference, waiting costs more early in a project. It also costs more on short projects and against fixed dates. And if a dip you waited on proves real, the sponsor who waited is exposed.

A check made case by case also puts judgment back at the line. That weakens the test unless the check is written as a fixed rule. A sponsor who judges that waiting costs too much can reasonably step in at once. That sponsor should then not read the turnaround rate as evidence.

Other ways to respond

Other responses each give something up. Holding to the annual plan gives up responding during the year. You can keep rescuing on the trigger and change only how you give credit. That keeps your quick response and fits the test at the line, but it still spends people on swing. You can leave some dips alone, at random or by a fixed rule set in advance. That gives the cleanest comparison. But it leaves those projects without help for a while, the same price as waiting. So it suits only a board with many similar projects and small dips.

When it is worth the trouble

The test costs a note in the review log for each project at each review: its number, whether it crossed the line and what was done. You run the comparison once enough projects near the line have built up. It is worth it where you step in often on similar projects. For a few very different projects, comparisons stay thin. This paper gives no figure for the cost. So where the test is worth it is a judgment, not a measured finding. A sponsor with few similar projects can reasonably judge the test too slow to keep. That sponsor should then simply stop reading the turnaround rate as evidence.

The argument has limits. It covers moves meant to fix a running project, such as a recovery plan, more people or a new lead. It covers them when they are chosen by a number that swings a lot compared with any lasting change, whether that number covers a quarter or the year to date. It does not cover cutting or stopping projects, moving money toward strong ones, or choosing among new proposals. It does not claim that rescues fail. It claims three things. The before-and-after record cannot show whether rescues work. The record overstates them on average. And it rises and falls with how readily you step in.

The review's question changes from "Did it come back after we stepped in?" to "Did projects just past our line do better than projects just short of it?"

Sources

1. Stephen Hall, Dan Lovallo and Reinier Musters, McKinsey Quarterly. How to put your money where your strategy is (2012).
2. David S. Christensen and Scott R. Heise, National Contract Management Journal. Cost Performance Index Stability (1993).
3. Adrian G. Barnett, Jolieke C. van der Pols and Annette J. Dobson, International Journal of Epidemiology. Regression to the mean: what it is and how to deal with it (2005).
4. Lucas M. Besters, Jan C. van Ours and Martin A. van Tuijl, De Economist. Effectiveness of In-Season Manager Changes in English Premier League Football (2016).
5. David S. Lee and Thomas Lemieux, NBER Working Paper 14723 (published in the Journal of Economic Literature, 2010). Regression Discontinuity Designs in Economics (2009).