

Capital Reallocation Is the Real Portfolio Decision

Why governing a portfolio means moving capacity in short cycles, not defending a plan set twelve months ago

Marco Policani

Enterprise Portfolio · PMO · AI Operating Governance

policani.net · linkedin.com/in/marcpolicani

August 2026

EXECUTIVE SUMMARY

A portfolio is governed when a sponsor can move funding and capacity on current evidence, not when the annual plan is well made.

Most organizations treat the annual plan as the portfolio decision, then defend last year's allocation until the next budget round. The decision that actually governs a portfolio is reallocation: moving funding and capacity to where the evidence now points, in short cycles, through explicit decision rights. This paper sets out why the annual cycle becomes a liability, what a short-cycle reallocation review requires, why the money so often refuses to move, and a one-page charter a sponsor can adopt without waiting for a planning season.

- 01** Run a short-cycle reallocation review — monthly or quarterly, never annual — whose only job is to decide what should move.
- 02** Give it four parts: a published trigger, a consistent evidence set, one accountable decision owner, and a required move or recorded hold.
- 03** Guard against churn: act on durable evidence over the medium term, not on every quarterly wobble.
- 04** Adopt a one-page reallocation charter so the mechanism is agreed before the evidence shifts, not argued when it does.

The plan is not the decision

In the portfolios I have governed under a fixed funding ceiling, the annual plan was usually treated as the decision. Leaders argued for months, agreed on a split of money and people across initiatives, and then defended that split until the next planning round. The plan felt like governance. It was mostly administration.

The decision that actually governs a portfolio is made continuously, not once a year. Every time the evidence shifts, a leader chooses whether to move funding and capacity toward the work that is now paying off or to leave them where last year's argument put them. An organization that can move resources on evidence governs its portfolio; one that can only wait for the next budget cycle administers a plan it has already outgrown.

The scale of that gap is measurable. McKinsey's State of Organizations 2026 surveyed more than 10,000 senior executives across 15 countries and 16 industries, fielded between June and September 2025. It found that 47 percent of leaders review budget and talent only annually or less often, while only 30 percent of organizations reallocate resources enterprise-wide.^[1] Those figures do not establish the right cadence for every portfolio. They do show that enterprise-level movement is rare

where the annual plan is still the dominant control. Moving resources is how a strategy stops being a statement and starts being an allocation.

This paper argues that reallocation, not the annual plan, is the portfolio decision worth governing: why the annual cycle becomes a liability, what the reallocation discipline looks like as a routine, why the money so often refuses to move, and a one-page charter a sponsor can adopt without waiting for a planning season.

What an annual budget actually freezes

An annual plan is a snapshot of conviction. It records what leaders believed at one moment about which initiatives deserved the most capacity. That belief was formed on the evidence available then: forecasts, business cases, and the confidence of the people making the pitch. The trouble is not that the snapshot is wrong but that the organization keeps acting on it long after the evidence behind it has changed.

EXHIBIT 1

An annual plan freezes three things in place

Funding

Money stays with initiatives instead of evidence, so a program missing every milestone keeps its capacity.

Capacity

Scarce specialists stay on work that has lost value, because reassigning them feels like admitting the plan was wrong.

Attention

The review checks conformance to plan, not whether the work is still the best use of the capacity it holds.

Front-end conviction is a weak foundation to freeze for a year. Bent Flyvbjerg's work on major projects finds that early estimates of costs and benefits are poor predictors of what actually happens. Disciplined quality control, using an outside view, improves the odds of a sound decision.[5] If the numbers that justified an allocation are unreliable at the moment of approval, defending that allocation for twelve months compounds the original error. The budget freezes an estimate that was never accurate enough to treat as fixed.

Three things get locked in place when the plan becomes the decision. Funding is locked to initiatives rather than to evidence, so a program that is missing every milestone keeps its capacity because it was funded in the plan. Capacity is locked to teams whose work has become less valuable, because

reassigning people feels like an admission that the original plan was wrong. And attention is locked to a defense posture. The review that follows an annual plan asks whether each initiative is on track against its plan, not whether it is still the best use of the capacity it holds.

The sharpest of the three is attention. A review built around the annual plan measures conformance: is the project delivering what it promised, on the promised schedule? That is a useful question, but not the governing one. A project can be perfectly on plan and still be the wrong place for the organization's next dollar and its next scarce specialist. Conformance can report status; it cannot redirect investment.

Reallocation is the portfolio decision

Portfolio management, done seriously, does not end with selection and a year of monitoring. It keeps the active portfolio under review as information changes. Stage-Gate describes the active work as subject to prioritization, acceleration, de-prioritization, termination, and reallocation when resources are constrained.[4] The point is to maximize value while keeping the mix aligned to strategy and within a manageable number of initiatives. A gate earns its place by accelerating the winners and failing fast on the losers, not by waving everything through.

EXHIBIT 2

Enterprise-level reallocation is rare

30%

of organizations reallocate resources enterprise-wide; 47% of leaders review budget and talent only annually or less often.

McKinsey & Company, The State of Organizations 2026 (survey of 10,000+ senior executives across 15 countries and 16 industries).

There is a longer-run performance case for treating reallocation as the real decision, and it should be read as background evidence for the discipline rather than as investment advice. McKinsey's study of resource reallocation, *How to Put Your Money Where Your Strategy Is*, tracked companies over 15 years. The top third of reallocators shifted an average of 56 percent of capital across their businesses over that period. That top third earned roughly 30 percent higher total shareholder returns each year than the bottom third, and were about 13 percent more likely to avoid acquisition or bankruptcy.[2] The mechanism the study points to is not financial engineering but discipline: companies that keep moving capital toward their better opportunities compound an advantage, while those that let allocations calcify slowly fund their own decline.

That evidence carries a caveat the discipline depends on, and dropping it would misread the finding. The same research is explicit that over short spans, under about three years, high reallocators underperformed. Moving money every quarter to chase the latest number is churn, and churn destroys value by never letting an investment mature. The advantage comes from consistent, evidence-led reallocation held over the medium term: moving when the evidence is real and durable, and holding when it is only noise. The governing skill is telling a signal that should move capacity from a wobble that should not.

So the standard this paper proposes is narrow and testable. A portfolio is governed when a sponsor can move funding and capacity between initiatives on current evidence, through an explicit decision, without waiting for the annual planning round. If the only mechanism to move a dollar or a specialist is next year's budget, the portfolio is administered, not governed, regardless of how disciplined the annual process looks.

The short-cycle reallocation review

The mechanism that makes the standard real is a short-cycle reallocation review: a standing routine, run monthly or quarterly, whose only job is to decide what should move. It is not a status meeting with reallocation bolted on. Status meetings are optimized to confirm that work is proceeding, which is the opposite of the disposition reallocation requires. The review needs its own cadence, its own evidence, and its own decision rights, and it rests on four elements.

A trigger. Reallocation should be provoked, not left to whoever feels bold enough to raise it. A trigger is a defined condition that puts an initiative's capacity back on the table: a missed evidence milestone, a better opportunity competing for the same specialists, a benefit case that no longer holds up against its baseline, or a shift in the environment that changes what the work is worth. Without triggers, reallocation happens only when someone feels bold enough to raise it, which is no reliable mechanism at all.

Evidence. The review runs on a small, consistent evidence set, not on advocacy. For each initiative under question it asks what the baseline was, what the leading and lagging signals now show, and what the honest outside-view read is on the remaining cost and benefit. This is where Flyvbjerg's discipline earns its place: the review discounts inside-view optimism and asks what comparable efforts actually delivered.[5] PMI's benefits-realization and portfolio guidance draws the same distinction at the standards level. It separates reporting that a project is on track from governing whether the investment still returns what it was meant to return.[3] The review governs the second thing.

A decision owner. Someone has to be able to say move, and be accountable for the call. Reallocation stalls when the authority to shift funding and capacity is diffuse, because a decision that belongs to everyone belongs to no one. The review names a single accountable owner for each reallocation decision, with the standing to redirect resources across initiative boundaries. That owner is the portfolio sponsor who governs the whole, not the project leader defending one part of it.

A move. The review's output is a change to the allocation, or an explicit, recorded decision to hold. "We discussed it" is not an output. Either capacity shifts, or the review states on the record why the current allocation is still the best available use of the resource, given the evidence. A review that produces neither has not governed. Recording why you held is itself a decision, and a legitimate one.

The cadence matters as much as the elements. McKinsey's survey found that 47 percent of leaders review budget and talent only annually or less often; 45 percent review quarterly and 7 percent monthly.[1] That evidence does not make monthly review a universal prescription. A portfolio should set its cadence to the speed at which its value assumptions can genuinely be tested. A quarterly product portfolio with slow customer-adoption evidence may need less frequent moves than an operational portfolio where a constrained team can see a missed service or cycle-time result within weeks. What is not defensible is allowing the annual plan to be the only mechanism for a material move.

The cadence also needs a decision horizon. Every charter should state the minimum evidence period before an initiative can be reduced, paused, or expanded. It should separately name the safety, compliance, and explicit stop conditions that can override that period. The horizon prevents a review from punishing work for not producing a result before the result could reasonably appear. It also forces the sponsor to distinguish an initiative whose premise has been disproved from one that is merely early. The practical question is whether the evidence has had enough time to make a different allocation more credible than continuity.

There is a public-sector version of the same discipline. The U.S. Government Accountability Office's investment-management framework says a portfolio review should operate from actual investment data, use predetermined thresholds when judging actual against expected performance, and adjust resources among investments when needed.[6] Its language is a useful corrective to a common misunderstanding. A threshold does not automate the decision or substitute a score for judgment. It makes the condition for reconsidering an allocation known before a sponsor is under pressure to defend it. The decision still requires a portfolio view: dependencies, capacity, risk, and the effect of the move on the rest of the work. But a review without a pre-agreed threshold is vulnerable to being a debate over who can make the stronger case in the room.

This is the same consistency guard the performance evidence demands: the cadence exists to catch durable shifts early, not to relitigate every allocation on a wobble. A well-run review moves rarely and decisively, and a recorded hold—the sponsor weighing a live alternative, applying the horizon, and choosing to protect an investment from premature churn—is itself a governing act, not a failure of the review.

What a review looks like when a trigger fires

The mechanics are easier to trust with a concrete shape, so here is the pattern I have seen hold up, described in operating terms rather than tied to any one organization. A specialist team — the kind of

scarce capacity every portfolio fights over — is committed across two initiatives. One initiative is hitting its evidence milestones; the other has missed two in a row and its benefit case no longer matches the baseline that was captured at approval. Under an annual plan, nothing happens until the next budget season, and the missing initiative keeps its share of the specialists because it held that share in the plan.

Under a reallocation charter, the second miss is a published trigger. The initiative's capacity is automatically on the table at the next monthly review. It arrives with the same small evidence set every initiative brings: the baseline, the current leading and lagging signals, and an outside-view read on what finishing the work would now cost and return. The team defending the struggling initiative does not get to argue from conviction; it has to argue from that evidence, and the evidence is thin. The portfolio sponsor — the single decision owner named in the charter — moves part of the specialist capacity to the initiative that is producing signal, and records the reasoning.

Two features of that sequence are worth naming. First, no one had to be brave. Because the trigger did the provoking, raising the question was simply the routine operating on schedule, not an act of political courage aimed at a colleague. Second, the move was bounded by the consistency guard. The sponsor did not strip the struggling initiative to zero on a single bad quarter, because the charter distinguishes a durable shift from a wobble. A second miss against a real baseline is a durable shift. The review moved decisively and proportionately, and it could explain both.

A conformance review would have logged the two misses, asked the struggling team for a recovery plan, and left the specialists where the plan put them: a status report and no decision. The reallocation review produced a smaller status report and an actual decision about where scarce capacity does the most good.

There is an equally important countercase. A strong initiative can miss an early milestone because a dependency arrived late, an approved market test is still running, or a change in the operating environment made the original measure temporarily uninformative. If its decision horizon has not elapsed and the premise remains credible, the sponsor should hold the allocation and record why. This is the same standard applied in the other direction, not leniency for a favored team: the evidence must be strong enough to justify a move, not merely uncomfortable enough to prompt one. A portfolio that treats every miss as a reason to cut will train teams to hide risk and will underfund work that needs a measured learning period. The charter protects both disciplines: it is harder to defend stale work on reputation, and harder to abandon promising work because the last dashboard was inconvenient.

Why the money does not move

If reallocation is this straightforward to describe, its rarity needs explaining. Only about 30 percent of organizations move resources across the enterprise. McKinsey respondents most often named organizational resistance or managerial protectionism (41 percent) and inefficient, lengthy decision-making (38 percent). A further 32 percent cited lack of willingness to make bold decisions.^[1] Those

figures are respondent views, not a causal model, but they point toward structural friction rather than a simple failure to understand the idea. Four forces hold capital in place.

The first is the sunk-cost reflex dressed as commitment. An initiative that has already consumed budget and executive credibility is defended because stopping it feels like conceding waste. The reflex is understandable and expensive. Money already spent is gone whether or not the work continues; the only live question is whether the next increment of capacity is better spent here or elsewhere. Reviews that ask "how do we get this back on track" instead of "is this still the best use of the next dollar" have already surrendered to the reflex.

The second is the political economy of the budget. In many organizations, holding a budget is a proxy for status, so giving capacity back reads as losing rather than as good governance. Leaders learn to defend their allocation regardless of its performance, because the system rewards defense. Where that is true, no better dashboard will fix reallocation, because the problem is the incentive, not the information. It will not yield until returning capacity is treated as a mark of judgment rather than a defeat.

The professional definition of a portfolio makes the trade-off plain: APM describes portfolio management as selection, prioritisation, and control in line with strategic objectives and delivery capacity, and notes that sponsors may need to sacrifice their own project priorities for the wider portfolio.[7] That is a boundary on local optimization, not a licence for a central office to second-guess every delivery decision. The sponsor closest to a project can explain its evidence; the portfolio owner decides whether that evidence still warrants the capacity it consumes against the alternatives. A review becomes credible when both roles are visible and neither impersonates the other.

The third is weak evidence at the point of decision. Reallocation requires a credible read on what each initiative is now worth, and many portfolios cannot produce one. Business cases were written to win approval, not to be revisited, and the baseline needed to judge realized benefit was never captured. This is the front-end quality problem again: if the original estimate was optimistic and no honest baseline exists, the review has nothing solid to move against, and inertia wins by default.[5] The fix is upstream. Portfolios that reallocate well are portfolios that captured a real baseline and defined what evidence would later justify a move.

The fourth is the absence of decision rights. Even where the evidence is clear and the will exists, reallocation dies if no one has the authority to shift resources across initiative or organizational boundaries. Capacity locked inside a function stays inside that function. Without an explicit, portfolio-level owner empowered to move funding and people, the review can conclude that a shift is warranted and still be unable to execute it. Decision rights are the difference between recommending a move and making one.

None of these forces is exotic. They are the ordinary physics of organizations. That is why reallocation has to be engineered as a routine: a trigger, evidence, an accountable owner, and authority to act. Left to good intentions, the money does not move.

The portfolio view that makes a move legitimate

The unit of analysis in a reallocation review is the contested resource and the set of alternatives that could use it, not the one project that happens to be in trouble. That distinction prevents a common failure: treating every performance review as a private hearing about whether one initiative deserves to survive. The portfolio owner's question is not whether a team worked hard enough or whether its sponsor can still tell a compelling story, but whether the next increment of constrained capacity will create more value here than somewhere else, given the enterprise's current objectives and the other work that depends on that capacity.

That changes the briefing material. A project-status deck is organized around activity: milestones completed, milestones delayed, issues open, and help requested. A reallocation brief must be organized around a choice. It identifies the resource at stake, the time window, and the initiative's original benefit and baseline. It then states the observed evidence, dependencies, credible alternatives, and the consequence of holding. The consequence of holding matters because inaction also allocates capacity. A sponsor who asks to keep a specialist must be able to say what the portfolio gives up by doing so.

The comparison does not need a false-precision score that turns judgment into arithmetic. A regulatory obligation, service-risk exposure, revenue opportunity, foundational dependency, and capability investment are not naturally interchangeable. Yet they may all need the same scarce architect or change lead. A mature review makes those differences visible rather than forcing them into a fabricated number. GAO's framework makes the underlying point in a different setting. It calls for decisions that consider cost, benefit, schedule, and risk, while comparing investments in the context of the whole portfolio rather than as isolated business cases.[6] The objective is not to prove that one initiative has a score of 73.4 and another 72.8. It is to make the grounds for a trade-off inspectable.

That is where the decision record earns its place. For every material move — or every material hold — the record should preserve five things: the trigger, the evidence considered, the alternatives genuinely available, the decision owner, and the expected next observation. That last element matters most: a move or a hold with no defined next observation quietly becomes permanent. Naming what must be true at the next review keeps the decision reversible on evidence rather than on politics.

The record also makes calibration possible. After several cycles, the portfolio office can ask whether its triggers are firing too late, whether decision horizons are protecting learning or sheltering weak work, and whether the organization is repeatedly moving the same bottleneck because the real constraint is structural. That is a better use of a portfolio-management function than producing increasingly elaborate status reports. It turns individual allocation decisions into evidence about the operating model itself.

A one-page reallocation charter

The discipline should be small enough to fit on one page and durable enough to survive a leadership change. A reallocation charter defines, in advance, how the portfolio moves resources, so that when the evidence shifts the mechanism is already agreed and the only question is the decision itself. A workable charter names seven things.

EXHIBIT 3

A reallocation charter names seven things before the evidence shifts

- 1 **Scope and cadence**
Which portfolio the charter governs, how often the review runs, and the evidence latency that makes that cadence credible.
- 2 **Triggers**
The defined conditions that place an initiative's capacity back on the table, published in advance.
- 3 **Evidence standard**
The small, consistent set each initiative brings: baseline, current signals, and an outside-view read.
- 4 **Decision owner**
The single accountable sponsor empowered to move funding and capacity across initiative boundaries.
- 5 **The move rule**
Every review produces a resource change or a recorded decision to hold, with the reasoning captured.
- 6 **The consistency guard**
A statement that the review acts on durable shifts, not noise, so it does not decay into churn.
- 7 **The decision horizon**
The evidence period each initiative earns before a move, plus the conditions that override it.

A charter of this shape complements the disciplines a portfolio already runs. It governs the money and capacity that are already inside the portfolio, where funding-discipline work governs what is allowed to enter and capacity work governs how much the organization can carry. Reallocation is the third leg: not what to start and not how much to hold, but where the resources already committed should now be.

The charter should be public to the people whose work it can affect. Hidden decision rules do not reduce politics; they only make every move look discretionary. Published rules let initiative leaders prepare the right evidence, challenge an error in the record, and understand the difference between a performance conversation and a portfolio decision.

It protects promises that are still earning their evidence.

The governing act

The honest test of portfolio governance is not the annual plan; any competent organization can produce a defensible one. The test is what happens in month four, when the evidence has moved and the plan has not — whether the organization can redirect funding and capacity on that new evidence, through an explicit decision made by an accountable owner, or can only note the gap and wait for the next planning round.

The evidence gives the discipline its weight without overstating it. Enterprise-level reallocation is uncommon, and leaders report organizational resistance, slow decisions, and reluctance to make hard calls as real barriers.[1] The long-run performance case for consistent, medium-term reallocation is strong enough to take seriously and bounded enough to cite honestly, including its own warning against short-term churn.[2] Together the two frame reallocation as an operating decision about where scarce capacity does the most good—not a financial maneuver—made often enough to matter and disciplined enough to last.

The plan tells you what you believed a year ago. Reallocation is what you do about what you know now. Only one of them is governance.

Notes and sources

[1]: Alexis Krivkovich, Damian Klingler, Dana Maor, Patrick Guggenberger, and Michael Anzenhofer, *The State of Organizations 2026: Three Tectonic Forces That Are Reshaping Organizations*, McKinsey & Company, February 19, 2026. Survey of more than 10,000 senior executives across 15 countries and 16 industries, fielded June–September 2025. The report says 47 percent of leaders review budget and talent only annually or less often; 45 percent review quarterly and 7 percent monthly; only 30 percent of organizations reallocate resources enterprise-wide. Reported barriers include organizational resistance or managerial protectionism (41 percent), inefficient and lengthy decision-making (38 percent), and lack of willingness to make bold decisions (32 percent). <https://www.mckinsey.com/~media/mckinsey/business%20functions/people%20and%20organizational%20performance/our%20insights/the%20state%20of%20organizations/2026/the-state-of-organizations-2026-vf.pdf>

[2]: Stephen Hall, Dan Lovallo, and Reinier Musters, "How to Put Your Money Where Your Strategy Is," McKinsey & Company, March 1, 2012. Over a 15-year study period, the top third of reallocators shifted an average of 56 percent of capital across their businesses, earned roughly 30 percent higher total shareholder returns annually than the bottom third, and were about 13 percent more likely to avoid acquisition or bankruptcy; the same research notes that over spans under three years, high

reallocators underperformed, so the advantage comes from consistent medium-term reallocation rather than annual churn. Cited as background evidence for the reallocation discipline, not as investment advice. <https://www.mckinsey.com/capabilities/strategy-and-corporate-finance/our-insights/how-to-put-your-money-where-your-strategy-is>

[3]: Project Management Institute, The Standard for Portfolio Management, Third Edition, 2013, pp. 91–92. The portfolio performance-management plan describes how resources are planned, balanced, and allocated, and benefits realization specifies how realized benefits are tracked and compared to plans. <https://www.pmi.org/-/media/pmi/documents/public/pdf/certifications/standard-for-portfolio-management-third-edition.pdf>

[4]: Stage-Gate International, "Portfolio Management: Fundamental for New Product Success." The guidance describes portfolio management as maximizing value, balancing the portfolio, maintaining strategic direction, and selecting the right number of projects under constrained resources. <https://www.stage-gate.com/blog/portfolio-management-fundamental-for-new-product-success/>

[5]: Bent Flyvbjerg, "Quality Control and Due Diligence in Project Management," 2013. The paper argues that front-end estimates of costs and benefits are often poor predictors of actual outcomes and that an outside-view discipline improves decisions. <https://arxiv.org/abs/1302.2544>

[6]: U.S. Government Accountability Office, Information Technology Investment Management: A Framework for Assessing and Improving Process Maturity, GAO-04-394G, March 2004. The framework describes a portfolio review based on actual investment data, predetermined thresholds for actual-versus-expected performance, and resource adjustments among investments as necessary. It treats portfolio decisions as enterprise-level comparisons rather than isolated project cases. <https://www.gao.gov/products/gao-04-394g>

[7]: Association for Project Management, "What is portfolio management?" APM Body of Knowledge, 8th edition. APM defines portfolio management as selection, prioritisation, and control in line with strategic objectives and delivery capacity, and notes that project sponsors may be required to sacrifice their priorities for the wider portfolio. <https://www.apm.org.uk/resources/what-is-project-management/what-is-portfolio-management/>

Marco Policani — Enterprise Portfolio · PMO · AI Operating Governance. policani.net · [linkedin.com/in/marcpolicani](https://www.linkedin.com/in/marcpolicani). This paper may be shared freely with attribution. © 2026 Marco Policani.

About the author

Marco Policani is an enterprise portfolio, PMO, and AI operating-governance leader. He builds portfolio governance systems where AI carries the analytical load and named humans own every decision — an approach documented across case studies, walkthroughs, and working governance guides on his portfolio site.

Portfolio & governance library: policani.net/governance · Full portfolio: policani.net · LinkedIn: linkedin.com/in/marcpolicani

This white paper may be shared freely with attribution. © 2026 Marco Policani.