
WHITE PAPER · AI OPERATING GOVERNANCE

Confident and Wrong

Challenge protocols for plausible AI output

Marco Policani

Enterprise Portfolio · PMO · AI Operating Governance

policani.net · linkedin.com/in/marcpolicani

July 2026

EXECUTIVE SUMMARY

Plausible output needs a challenge protocol before it changes a consequential decision or action.

Fluent AI output can enter consequential work with more confidence than evidence. A challenge protocol turns vague advice to verify into a repeatable control that states the decision, seeks disconfirmation, checks claims independently, records review, and learns from failure. This paper gives AI program owners, risk leaders, workflow owners, and executives relying on generated analysis a practical the five-part challenge protocol for making the next commitment explicit.

- 01** State the consequential decision: Review intensity should follow what the output may change, not how polished the text appears.
- 02** Seek disconfirming evidence: The first challenge should test the favored conclusion rather than ask the system to make it sound stronger.
- 03** Check claims independently: Verification means following important claims to sources outside the generated answer.
- 04** Use a named reviewer and escalation path: Human review works only when the reviewer has time, competence, authority, and a defined response to uncertainty.

CONTENTS

The argument	Learn from plausible failures
Why the conventional approach breaks	How the elements interact
The five-part challenge protocol	Put the model into the management cadence
State the consequential decision	Measures that reveal whether the control works
Seek disconfirming evidence	The strongest counterargument
Check claims independently	The shift
Use a named reviewer and escalation path	

The argument

Fluent AI output can enter consequential work with more confidence than evidence. A challenge protocol turns vague advice to verify into a repeatable control that states the decision, seeks disconfirmation, checks claims independently, records review, and learns from failure.

The hardest AI error to govern is not nonsense. It is an answer that is coherent, specific, aligned with the user's expectations, and wrong in a way that survives a quick read. Plausibility lowers the friction of reliance. Agreement can lower it further: an answer that confirms a preferred interpretation often receives less scrutiny than one that creates resistance.

OpenAI's public account of a sycophancy incident and Microsoft's lessons from red-teaming generative AI products make the operating point concrete: model behavior can fail in interaction, and broad testing still has to be translated into controls inside the workflow [1][2].

This paper is written for AI program owners, risk leaders, workflow owners, and executives relying on generated analysis. It offers the five-part challenge protocol as a management instrument: a way to connect operating evidence to the next consequential choice. The cited sources provide external context; the framework and its application are Marco Policani's synthesis. It is not a predictive score and does not replace professional, legal, security, financial, or domain judgment.

Why the conventional approach breaks

The work becomes governable when ownership, evidence, boundaries, and response are connected to a real choice. Conventional governance tends to separate planning, approval, delivery reporting, and operating ownership. That separation allows a premise to pass through several forums without any one forum owning the decision it implies. The project or workflow then carries an assumption that was acknowledged socially but never accepted operationally.

Under pressure, the workaround is predictable: capable people compensate, local decisions proliferate, and the formal record stops describing how the operation actually runs. The answer is not heavier control at every step. It is to place a clear, proportionate test at the point where the organization increases money, capacity, access, dependency, or action authority.

The five-part challenge protocol

The model has 5 connected elements. Each makes a different condition visible, but they should be reviewed together because strength in one area does not cancel a missing obligation in another. The model is designed for a short executive conversation supported by inspectable evidence, not a long maturity assessment.

EXHIBIT 1

The five-part challenge protocol connects evidence to action

- 1 **State the consequential decision**
Review intensity should follow what the output may change, not how polished the text appears.
- 2 **Seek disconfirming evidence**
The first challenge should test the favored conclusion rather than ask the system to make it sound stronger.
- 3 **Check claims independently**
Verification means following important claims to sources outside the generated answer.

4 Use a named reviewer and escalation path

Human review works only when the reviewer has time, competence, authority, and a defined response to uncertainty.

5 Learn from plausible failures

A rejected answer is operating evidence about prompts, data, interfaces, review load, and permitted use.

State the consequential decision

Review intensity should follow what the output may change, not how polished the text appears. The visible artifact is often complete before the operating condition behind it is credible.

A reviewer edits wording without realizing the answer will influence a hiring choice, investment recommendation, customer action, control decision, or public claim. The organization may call this execution risk, yet the failure began when authority and evidence were separated at the decision point.

Put the decision, affected party, reversibility, and consequence of error above the output before review begins. A useful test asks what ordinary operating pressure would break the premise and how that break would be detected before consequence spreads.

A use record that distinguishes brainstorming, drafting, recommendation, and action authority. Evidence should arrive before the latest useful decision date; perfect analysis delivered after options expire has little governance value.

Apply the protocol when reliance could create material consequence; keep low-risk assistance proportionate. The operating owner should confirm that the decision can be executed with the people, systems, and time actually available.

A tabletop review of state the consequential decision should use the observed break as its scenario: A reviewer edits wording without realizing the answer will influence a hiring choice, investment recommendation, customer action, control decision, or public claim. Participants then apply the proposed response. They put the decision, affected party, reversibility, and consequence of error above the output before review begins. The review identifies where authority, information, time, or recovery would run out. It is useful only if it changes the boundary or confirms who will act when the condition appears.

The minimum record for state the consequential decision is anchored in actual proof. It includes a use record that distinguishes brainstorming, drafting, recommendation, and action authority. The record also carries the choice supported by that proof: Apply the protocol when reliance could create material consequence; keep low-risk assistance proportionate. Keeping evidence and decision together lets a later owner distinguish a deliberate residual risk from a premise that nobody examined. It also gives the team a clear reason to reopen the choice when the underlying facts change.

Across the work governed by AI program owners, risk leaders, workflow owners, and executives relying on generated analysis, occurrences of this pattern should be compared rather than treated as unrelated project stories. The positive standard is: Review intensity should follow what the output may change, not how polished the text appears. If the same failure recurs, leaders can change delegation, capacity, intake, policy, or the intervention itself. That portfolio response is usually more valuable than asking each team to absorb another local workaround.

Two questions test state the consequential decision. What evidence would show that the condition is no longer adequate? Would that evidence arrive early enough to preserve a feasible choice? The intended choice is clear: Apply the protocol when reliance could create material consequence; keep low-risk assistance proportionate. A weak answer to either question calls for a smaller commitment, an earlier signal, or a safer fallback. A strong answer earns movement without claiming that uncertainty has disappeared.

Seek disconfirming evidence

The first challenge should test the favored conclusion rather than ask the system to make it sound stronger. The issue is not a shortage of activity; it is a shortage of inspectable commitments.

Users refine and expand an agreeable answer, creating multiple versions of the same unsupported premise. The missing condition is often treated as something the team will resolve later, even though the team lacks the authority or capacity to resolve it.

Ask what observation would make the conclusion false, what alternative explanation fits, and which stakeholder would disagree for substantive reasons. The forum should be able to distinguish an acceptable residual risk from a missing obligation that has merely been renamed.

Counterexamples, omitted cases, conflicting sources, sensitivity to changed assumptions, and unresolved uncertainty. The evidence may be compact, but it must be inspectable. A future reviewer should be able to see what was examined, what conclusion was drawn, and which limitation remained.

Narrow or withhold the conclusion when it survives only inside one framing. Leaders should prefer the smallest commitment that preserves learning and limits irreversible exposure.

A tabletop review of seek disconfirming evidence should use the observed break as its scenario: Users refine and expand an agreeable answer, creating multiple versions of the same unsupported premise. Participants then apply the proposed response. They ask what observation would make the conclusion false, what alternative explanation fits, and which stakeholder would disagree for substantive reasons. The review identifies where authority, information, time, or recovery would run out. It is useful only if it changes the boundary or confirms who will act when the condition appears.

The minimum record for seek disconfirming evidence is anchored in actual proof. It includes counterexamples, omitted cases, conflicting sources, sensitivity to changed assumptions, and unresolved uncertainty. The record also carries the choice supported by that proof: Narrow or withhold the conclusion when it survives only inside one framing. Keeping evidence and decision together lets a later owner distinguish a deliberate residual risk from a premise that nobody examined. It also gives the team a clear reason to reopen the choice when the underlying facts change.

Across the work governed by AI program owners, risk leaders, workflow owners, and executives relying on generated analysis, occurrences of this pattern should be compared rather than treated as unrelated project stories. The positive standard is: The first challenge should test the favored conclusion rather than ask the system to make it sound stronger. If the same failure recurs, leaders can change delegation, capacity, intake, policy, or the intervention itself. That portfolio response is usually more valuable than asking each team to absorb another local workaround.

Two questions test seek disconfirming evidence. What evidence would show that the condition is no longer adequate? Would that evidence arrive early enough to preserve a feasible choice? The intended choice is clear: Narrow or withhold the conclusion when it survives only inside one framing. A weak answer to either question calls for a smaller commitment, an earlier signal, or a safer fallback. A strong answer earns movement without claiming that uncertainty has disappeared.

Check claims independently

Verification means following important claims to sources outside the generated answer. The cost appears downstream, but the missing decision sits upstream.

The model supplies references, quotations, calculations, or policy interpretations that look credible and are accepted without opening the underlying material. The recurring signal is translation work: people spend time discovering what was meant, who may decide, and how to repair the unsupported assumption.

Select claims by consequence, trace them to authoritative sources, reproduce calculations, and separate facts from inference. This does not require a larger pack. It requires a traceable connection among the premise, the evidence, the boundary, and the action that follows.

Source title, author, date, relevant passage or data, calculation record, and the reviewer who performed the check. Evidence has a shelf life. The record needs the conditions or date that would make the conclusion stale.

Remove the claim or reduce reliance when independent evidence cannot be obtained. A hold is governable when the return conditions are explicit; otherwise it becomes an unmanaged queue.

A tabletop review of check claims independently should use the observed break as its scenario: The model supplies references, quotations, calculations, or policy interpretations that look credible and are accepted without opening the underlying material. Participants then apply the proposed response. They select claims by consequence, trace them to authoritative sources, reproduce calculations, and separate facts from inference. The review identifies where authority, information, time, or recovery would run out. It is useful only if it changes the boundary or confirms who will act when the condition appears.

The minimum record for check claims independently is anchored in actual proof. It includes source title, author, date, relevant passage or data, calculation record, and the reviewer who performed the check. The record also carries the choice supported by that proof: Remove the claim or reduce reliance when independent evidence cannot be obtained. Keeping evidence and decision together lets a later owner distinguish a deliberate residual risk from a premise that nobody examined. It also gives the team a clear reason to reopen the choice when the underlying facts change.

Across the work governed by AI program owners, risk leaders, workflow owners, and executives relying on generated analysis, occurrences of this pattern should be compared rather than treated as unrelated project stories. The positive standard is: Verification means following important claims to sources outside the generated answer. If the same failure recurs, leaders can change delegation, capacity, intake, policy, or the intervention itself. That portfolio response is usually more valuable than asking each team to absorb another local workaround.

Two questions test check claims independently. What evidence would show that the condition is no longer adequate? Would that evidence arrive early enough to preserve a feasible choice? The intended choice is clear: Remove the claim or reduce reliance when independent evidence cannot be obtained. A weak answer to either

question calls for a smaller commitment, an earlier signal, or a safer fallback. A strong answer earns movement without claiming that uncertainty has disappeared.

Use a named reviewer and escalation path

Human review works only when the reviewer has time, competence, authority, and a defined response to uncertainty. The distinction matters because leaders authorize commitments, while teams inherit whatever the authorization left unresolved.

A nominal approver clicks through large volumes of output, lacks domain context, or cannot stop the downstream action. What looks like a team-level miss is frequently a portfolio choice that nobody made explicitly. Scarce attention is then spent reconciling consequences rather than choosing among alternatives.

Match reviewers to consequence and subject matter; define sampling, separation of duties, and escalation for contested or novel cases. A named owner must be able to change the work; receiving a notification after the boundary is crossed is insufficient.

Reviewer identity, disposition, exceptions, reasons for override, and evidence that stop authority is real. The evidence owner is accountable for integrity, while the decision owner remains accountable for the commitment. Keeping those roles distinct improves challenge.

Route high-consequence ambiguity to a qualified owner instead of treating any human touch as sufficient. The residual uncertainty belongs in the decision record so later outcomes are compared with what leaders actually knew at the time.

A tabletop review of use a named reviewer and escalation path should use the observed break as its scenario: A nominal approver clicks through large volumes of output, lacks domain context, or cannot stop the downstream action. Participants then apply the proposed response. They match reviewers to consequence and subject matter; define sampling, separation of duties, and escalation for contested or novel cases. The review identifies where authority, information, time, or recovery would run out. It is useful only if it changes the boundary or confirms who will act when the condition appears.

The minimum record for use a named reviewer and escalation path is anchored in actual proof. It includes reviewer identity, disposition, exceptions, reasons for override, and evidence that stop authority is real. The record also carries the choice supported by that proof: Route high-consequence ambiguity to a qualified owner instead of treating any human touch as sufficient. Keeping evidence and decision together lets a later owner distinguish a deliberate residual risk from a premise that nobody examined. It also gives the team a clear reason to reopen the choice when the underlying facts change.

Across the work governed by AI program owners, risk leaders, workflow owners, and executives relying on generated analysis, occurrences of this pattern should be compared rather than treated as unrelated project stories. The positive standard is: Human review works only when the reviewer has time, competence, authority, and a defined response to uncertainty. If the same failure recurs, leaders can change delegation, capacity, intake, policy, or the intervention itself. That portfolio response is usually more valuable than asking each team to absorb another local workaround.

Two questions test use a named reviewer and escalation path. What evidence would show that the condition is no longer adequate? Would that evidence arrive early enough to preserve a feasible choice? The intended choice is

clear: Route high-consequence ambiguity to a qualified owner instead of treating any human touch as sufficient. A weak answer to either question calls for a smaller commitment, an earlier signal, or a safer fallback. A strong answer earns movement without claiming that uncertainty has disappeared.

Learn from plausible failures

A rejected answer is operating evidence about prompts, data, interfaces, review load, and permitted use. The control belongs at the point where reliance expands, not in a retrospective explanation after the outcome is known.

Corrections stay in private conversations, so the same failure pattern returns across users and workflows. Momentum makes the weakness self-reinforcing. Each additional commitment raises the political and economic cost of revisiting the premise.

Classify the failure, preserve representative cases, change controls, and feed the case into evaluation and training. The record should travel with the work so that later participants can see the basis for reliance instead of reconstructing it from memory.

Failure taxonomy, recurrence, near misses, detection point, control change, and later test results. The strongest record shows both what supports the choice and what could overturn it.

Reduce or suspend permitted use when failures outrun the organization's ability to detect and correct them. Decision quality improves when dissent and competing evidence are preserved instead of flattened into a unanimous-looking status.

A tabletop review of learn from plausible failures should use the observed break as its scenario: Corrections stay in private conversations, so the same failure pattern returns across users and workflows. Participants then apply the proposed response. They classify the failure, preserve representative cases, change controls, and feed the case into evaluation and training. The review identifies where authority, information, time, or recovery would run out. It is useful only if it changes the boundary or confirms who will act when the condition appears.

The minimum record for learn from plausible failures is anchored in actual proof. It includes failure taxonomy, recurrence, near misses, detection point, control change, and later test results. The record also carries the choice supported by that proof: Reduce or suspend permitted use when failures outrun the organization's ability to detect and correct them. Keeping evidence and decision together lets a later owner distinguish a deliberate residual risk from a premise that nobody examined. It also gives the team a clear reason to reopen the choice when the underlying facts change.

Across the work governed by AI program owners, risk leaders, workflow owners, and executives relying on generated analysis, occurrences of this pattern should be compared rather than treated as unrelated project stories. The positive standard is: A rejected answer is operating evidence about prompts, data, interfaces, review load, and permitted use. If the same failure recurs, leaders can change delegation, capacity, intake, policy, or the intervention itself. That portfolio response is usually more valuable than asking each team to absorb another local workaround.

Two questions test learn from plausible failures. What evidence would show that the condition is no longer adequate? Would that evidence arrive early enough to preserve a feasible choice? The intended choice is clear: Reduce or suspend permitted use when failures outrun the organization's ability to detect and correct them. A

weak answer to either question calls for a smaller commitment, an earlier signal, or a safer fallback. A strong answer earns movement without claiming that uncertainty has disappeared.

How the elements interact

State the consequential decision and Seek disconfirming evidence protect different parts of the same commitment. Review intensity should follow what the output may change, not how polished the text appears. The first challenge should test the favored conclusion rather than ask the system to make it sound stronger. Teams move faster when uncertainty is bounded honestly than when it is concealed inside an optimistic plan. If the operation encounters this failure—a reviewer edits wording without realizing the answer will influence a hiring choice, investment recommendation, customer action, control decision, or public claim.—the response for seek disconfirming evidence cannot compensate for the missing state the consequential decision obligation. Reviewers should reconcile the evidence for both elements before they accept the next action: narrow or withhold the conclusion when it survives only inside one framing.

Seek disconfirming evidence and Check claims independently protect different parts of the same commitment. The first challenge should test the favored conclusion rather than ask the system to make it sound stronger. Verification means following important claims to sources outside the generated answer. The framework in this paper is designed to make that boundary usable in ordinary management cadence. If the operation encounters this failure—users refine and expand an agreeable answer, creating multiple versions of the same unsupported premise.—the response for check claims independently cannot compensate for the missing seek disconfirming evidence obligation. Reviewers should reconcile the evidence for both elements before they accept the next action: remove the claim or reduce reliance when independent evidence cannot be obtained.

Check claims independently and Use a named reviewer and escalation path protect different parts of the same commitment. Verification means following important claims to sources outside the generated answer. Human review works only when the reviewer has time, competence, authority, and a defined response to uncertainty. The distinction matters because leaders authorize commitments, while teams inherit whatever the authorization left unresolved. If the operation encounters this failure—the model supplies references, quotations, calculations, or policy interpretations that look credible and are accepted without opening the underlying material.—the response for use a named reviewer and escalation path cannot compensate for the missing check claims independently obligation. Reviewers should reconcile the evidence for both elements before they accept the next action: route high-consequence ambiguity to a qualified owner instead of treating any human touch as sufficient.

Use a named reviewer and escalation path and Learn from plausible failures protect different parts of the same commitment. Human review works only when the reviewer has time, competence, authority, and a defined response to uncertainty. A rejected answer is operating evidence about prompts, data, interfaces, review load, and permitted use. The practical failure occurs between the formal approval and the ordinary pressure of running the work. If the operation encounters this failure—a nominal approver clicks through large volumes of output, lacks domain context, or cannot stop the downstream action.—the response for learn from plausible failures cannot compensate for the missing use a named reviewer and escalation path obligation. Reviewers should reconcile the evidence for both elements before they accept the next action: reduce or suspend permitted use when failures outrun the organization's ability to detect and correct them.

Learn from plausible failures and State the consequential decision protect different parts of the same commitment. A rejected answer is operating evidence about prompts, data, interfaces, review load, and permitted use. Review

intensity should follow what the output may change, not how polished the text appears. Viewed from the portfolio, this is a resource-allocation problem before it becomes a delivery problem. If the operation encounters this failure—corrections stay in private conversations, so the same failure pattern returns across users and workflows.—the response for state the consequential decision cannot compensate for the missing learn from plausible failures obligation. Reviewers should reconcile the evidence for both elements before they accept the next action: apply the protocol when reliance could create material consequence; keep low-risk assistance proportionate.

Put the model into the management cadence

A framework becomes useful when it changes the normal cadence of work. It should shape intake, funding, design, delivery, and operating review without creating a separate governance industry around itself. The following sequence is deliberately small:

EXHIBIT 2

Start with one decision, then calibrate from evidence

- 1 **Move 1**
Put the five questions into the interface or work record where reliance occurs.
- 2 **Move 2**
Define which decisions require independent checking and which can use sampling.
- 3 **Move 3**
Calibrate reviewer capacity before increasing generated volume.
- 4 **Move 4**
Review failure patterns monthly and change prompts, data, permissions, workflow, and user guidance together.

- Put the five questions into the interface or work record where reliance occurs.
- Define which decisions require independent checking and which can use sampling.
- Calibrate reviewer capacity before increasing generated volume.
- Review failure patterns monthly and change prompts, data, permissions, workflow, and user guidance together.

The first cycle should be treated as calibration. Reviewers should note which questions changed a decision, which evidence was unavailable, where authority was unclear, and which controls cost out of proportion to the exposure. That record improves the five-part challenge protocol without pretending the first version will fit every workflow or investment.

Ownership should remain close to the consequence. The PMO or AI-governance function can maintain the method, surface cross-portfolio patterns, and challenge weak evidence. It should not absorb the business owner's accountability for the result or the executive's accountability for the commitment.

Measures that reveal whether the control works

- Time from a material signal to a consequential decision.
- Commitments narrowed, redirected, paused, or stopped before avoidable exposure grew.
- Exceptions or assumptions discovered after the latest useful decision date.
- Scarce specialist and executive attention consumed by recurring unresolved conditions.
- Decisions reopened because the original evidence, boundary, or rationale was unclear.
- Operating outcomes and recovery performance after reliance expanded.

These measures are diagnostic, not a universal scorecard. Their purpose is to identify where the operating model fails repeatedly and whether governance is preserving options earlier. A lower count can indicate improvement, but it can also indicate weak detection. Leaders should read the measures with case evidence and frontline observation.

The strongest counterargument

The protocol does not assume every model answer is false or that a human is automatically right. It creates structured friction where consequence warrants it. Independent evidence, qualified judgment, and recoverable action are the safeguards; generic skepticism is not.

A reasonable critic could also argue that experienced leaders already do this through judgment. Many do. The weakness is not the absence of judgment; it is the absence of a durable operating record when people, pressure, and conditions change. The five-part challenge protocol makes the basis for judgment visible enough to challenge, execute, and revisit.

The shift

Fluent AI output can enter consequential work with more confidence than evidence. A challenge protocol turns vague advice to verify into a repeatable control that states the decision, seeks disconfirmation, checks claims independently, records review, and learns from failure. The practical shift is from reporting activity to governing the conditions under which the organization relies on that activity.

That discipline does not make uncertainty disappear. It gives leaders a better place to stand: a named decision, evidence proportionate to consequence, an explicit boundary, a responsible owner, and a response when reality departs from the assumption. Those elements are enough to turn hidden operating risk into a choice the portfolio can make deliberately.

Notes and sources

[1] OpenAI, "Sycophancy in GPT-4o: What Happened and What We're Doing About It," April 2025.
<https://openai.com/index/sycophancy-in-gpt-4o/>

[2] Blake Bullwinkel et al., "Lessons From Red Teaming 100 Generative AI Products," Microsoft Research, January 2025. <https://www.microsoft.com/en-us/research/publication/lessons-from-red-teaming-100-generative-ai-products/>

About the author

Marco Policani is an enterprise portfolio, PMO, and AI operating-governance leader. He builds portfolio governance systems where AI carries the analytical load and named humans own every decision — an approach documented across case studies, walkthroughs, and working governance guides on his portfolio site.

Portfolio & governance library: policani.net/governance · Full portfolio: policani.net · LinkedIn: linkedin.com/in/marcpolicani

This white paper may be shared freely with attribution. © 2026 Marco Policani.