13

WHITE PAPER · WORK, ADOPTION & JUDGMENT

The Coordination Tax

Design the coordination topology before declaring work automated

Marco Policani

Enterprise Portfolio · PMO · AI Operating Governance

policani.net · linkedin.com/in/marcpolicani

July 2026

EXECUTIVE SUMMARY

Automation changes where coordination happens; leaders must design the handoffs, exceptions, and evidence that remain.

Automation can remove task effort while leaving—or increasing—the work required to route, reconcile, approve, explain, and recover. Leaders should map the coordination topology and measure the total workflow before declaring productivity or autonomy. This paper gives operating executives, transformation leaders, workflow owners, and AI program leaders a practical the coordination topology map for making the next commitment explicit.

01 Map work between tasks: The unit of analysis is the path across people, systems, decisions, and queues—not the isolated activity the tool performs.

02 Distinguish productive coordination: Some interaction creates judgment, learning, trust, safety, or necessary integration; eliminating it can degrade the outcome.

03 Find boundary work: Automation often creates new translation between generated output and the systems or people that must use it.

04 Design handoffs deliberately: Every transfer needs a clear object, receiving condition, owner, and response expectation.

CONTENTS

The argument

Automation can remove task effort while leaving—or increasing—the work required to route, reconcile, approve, explain, and recover. Leaders should map the coordination topology and measure the total workflow before declaring productivity or autonomy.

A task can become dramatically faster while the process around it becomes no faster at all. Someone still finds the input, decides which version is authoritative, waits for an approval, explains an exception, reconciles systems, and repairs the cases that do not fit. That surrounding work is easy to miss because it is distributed across roles and rarely appears in the automation business case.

Recent model-based research on task exposure offers a useful hypothesis about how AI may change work, while NIST provides a more durable governance frame [1][2]. Neither relieves leaders of measuring what happens in their own workflow.

This paper is written for operating executives, transformation leaders, workflow owners, and AI program leaders. It offers the coordination topology map as a management instrument: a way to connect operating evidence to the next consequential choice. The cited sources provide external context; the framework and its application are Marco Policani's synthesis. It is not a predictive score and does not replace professional, legal, security, financial, or domain judgment.

Why the conventional approach breaks

The practical failure occurs between the formal approval and the ordinary pressure of running the work. Conventional governance tends to separate planning, approval, delivery reporting, and operating ownership. That separation allows a premise to pass through several forums without any one forum owning the decision it implies. The project or workflow then carries an assumption that was acknowledged socially but never accepted operationally.

Momentum makes the weakness self-reinforcing. Each additional commitment raises the political and economic cost of revisiting the premise. The answer is not heavier control at every step. It is to place a clear, proportionate test at the point where the organization increases money, capacity, access, dependency, or action authority.

The coordination topology map

The model has 6 connected elements. Each makes a different condition visible, but they should be reviewed together because strength in one area does not cancel a missing obligation in another. The model is designed for a short executive conversation supported by inspectable evidence, not a long maturity assessment.

EXHIBIT 1

The coordination topology map connects evidence to action

1. **Map work between tasks**
   The unit of analysis is the path across people, systems, decisions, and queues—not the isolated activity the tool performs.
2. **Distinguish productive coordination**
   Some interaction creates judgment, learning, trust, safety, or necessary integration; eliminating it can degrade the outcome.
3. **Find boundary work**
   Automation often creates new translation between generated output and the systems or people that must use it.

4. **Design handoffs deliberately**
   Every transfer needs a clear object, receiving condition, owner, and response expectation.
5. **Own exceptions and recovery**
   A faster happy path can increase the absolute number or complexity of cases requiring human coordination.
6. **Measure the whole outcome**
   Value should appear in end-to-end cycle time, quality, capacity, coverage, decision speed, or risk; tool usage and task duration are supporting measures.

Map work between tasks

The unit of analysis is the path across people, systems, decisions, and queues—not the isolated activity the tool performs. The control belongs at the point where reliance expands, not in a retrospective explanation after the outcome is known.

An automated drafting or classification step looks successful while handoff volume, waiting, clarification, and reconciliation remain untouched. Momentum makes the weakness self-reinforcing. Each additional commitment raises the political and economic cost of revisiting the premise.

Trace the case from demand to completed outcome and mark every transfer, queue, approval, dependency, and return loop. The record should travel with the work so that later participants can see the basis for reliance instead of reconstructing it from memory.

Touch count, queue time, rework, responsible role, system changes, and reasons cases move backward. The strongest record shows both what supports the choice and what could overturn it.

Automate only after leaders can see whether the proposed change removes or relocates coordination. Decision quality improves when dissent and competing evidence are preserved instead of flattened into a unanimous-looking status.

A tabletop review of map work between tasks should use the observed break as its scenario: An automated drafting or classification step looks successful while handoff volume, waiting, clarification, and reconciliation remain untouched. Participants then apply the proposed response. They trace the case from demand to completed outcome and mark every transfer, queue, approval, dependency, and return loop. The review identifies where authority, information, time, or recovery would run out. It is useful only if it changes the boundary or confirms who will act when the condition appears.

The minimum record for map work between tasks is anchored in actual proof. It includes touch count, queue time, rework, responsible role, system changes, and reasons cases move backward. The record also carries the choice supported by that proof: Automate only after leaders can see whether the proposed change removes or relocates coordination. Keeping evidence and decision together lets a later owner distinguish a deliberate residual risk from a premise that nobody examined. It also gives the team a clear reason to reopen the choice when the underlying facts change.

Across the work governed by operating executives, transformation leaders, workflow owners, and AI program leaders, occurrences of this pattern should be compared rather than treated as unrelated project stories. The positive standard is: The unit of analysis is the path across people, systems, decisions, and queues—not the isolated activity the tool performs. If the same failure recurs, leaders can change delegation, capacity, intake, policy, or the intervention itself. That portfolio response is usually more valuable than asking each team to absorb another local workaround.

Two questions test map work between tasks. What evidence would show that the condition is no longer adequate? Would that evidence arrive early enough to preserve a feasible choice? The intended choice is clear: Automate only after leaders can see whether the proposed change removes or relocates coordination. A weak answer to either question calls for a smaller commitment, an earlier signal, or a safer fallback. A strong answer earns movement without claiming that uncertainty has disappeared.

Distinguish productive coordination

Some interaction creates judgment, learning, trust, safety, or necessary integration; eliminating it can degrade the outcome. This is where a management ritual either becomes a real control or reveals itself as administrative theater.

All meetings, reviews, and handoffs are classified as waste without asking what uncertainty or accountability they resolve. Because responsibility is distributed, no single event appears large enough to stop the work. The accumulated operating load is the real exposure.

For each coordination point, name the decision, knowledge transfer, control, or relationship it produces and test whether a lighter mechanism can preserve it. Evidence should be selected for decision value. Volume is not rigor when the decisive uncertainty remains untested.

Decision quality, defect detection, learning, escalation prevented, and the consequence of removing the interaction. A later outcome should feed back into the evidence standard so repeated surprises improve the next commitment.

Keep valuable coordination, redesign avoidable routing, and automate only the stable remainder. A decision without consequence invites the old behavior to continue. Capacity, funding, permissions, scope, or sequence must move with the answer.

A tabletop review of distinguish productive coordination should use the observed break as its scenario: All meetings, reviews, and handoffs are classified as waste without asking what uncertainty or accountability they resolve. Participants then apply the proposed response. They for each coordination point, name the decision, knowledge transfer, control, or relationship it produces and test whether a lighter mechanism can preserve it. The review identifies where authority, information, time, or recovery would run out. It is useful only if it changes the boundary or confirms who will act when the condition appears.

The minimum record for distinguish productive coordination is anchored in actual proof. It includes decision quality, defect detection, learning, escalation prevented, and the consequence of removing the interaction. The record also carries the choice supported by that proof: Keep valuable coordination, redesign avoidable routing, and automate only the stable remainder. Keeping evidence and decision together lets a later owner distinguish a deliberate residual risk from a premise that nobody examined. It also gives the team a clear reason to reopen the choice when the underlying facts change.

Across the work governed by operating executives, transformation leaders, workflow owners, and AI program leaders, occurrences of this pattern should be compared rather than treated as unrelated project stories. The positive standard is: Some interaction creates judgment, learning, trust, safety, or necessary integration; eliminating it can degrade the outcome. If the same failure recurs, leaders can change delegation, capacity, intake, policy, or the intervention itself. That portfolio response is usually more valuable than asking each team to absorb another local workaround.

Two questions test distinguish productive coordination. What evidence would show that the condition is no longer adequate? Would that evidence arrive early enough to preserve a feasible choice? The intended choice is clear: Keep valuable coordination, redesign avoidable routing, and automate only the stable remainder. A weak answer to either question calls for a smaller commitment, an earlier signal, or a safer fallback. A strong answer earns movement without claiming that uncertainty has disappeared.

Find boundary work

Automation often creates new translation between generated output and the systems or people that must use it. The work becomes governable when ownership, evidence, boundaries, and response are connected to a real choice.

Users copy, normalize, explain, and verify output manually while those activities remain outside the measured process. When the consequence finally becomes visible, leaders face a narrower choice set and a more expensive recovery.

Observe real users and record formatting, source checking, exception explanation, data preparation, and downstream adaptation. The control should be short enough to use at the decision point and specific enough that a skeptical owner can challenge it.

Minutes and touches per case, correction categories, role seniority, and variability across case types. Ranges and contradictory observations are often more useful than a precise point estimate built on weak assumptions.

Include boundary work in the value case and redesign the interface or process when it dominates. The decision should identify the sacrifice. If every priority remains protected, the tradeoff has been delegated rather than resolved.

A tabletop review of find boundary work should use the observed break as its scenario: Users copy, normalize, explain, and verify output manually while those activities remain outside the measured process. Participants then apply the proposed response. They observe real users and record formatting, source checking, exception explanation, data preparation, and downstream adaptation. The review identifies where authority, information, time, or recovery would run out. It is useful only if it changes the boundary or confirms who will act when the condition appears.

The minimum record for find boundary work is anchored in actual proof. It includes minutes and touches per case, correction categories, role seniority, and variability across case types. The record also carries the choice supported by that proof: Include boundary work in the value case and redesign the interface or process when it dominates. Keeping evidence and decision together lets a later owner distinguish a deliberate residual risk from a premise that nobody examined. It also gives the team a clear reason to reopen the choice when the underlying facts change.

Across the work governed by operating executives, transformation leaders, workflow owners, and AI program leaders, occurrences of this pattern should be compared rather than treated as unrelated project stories. The positive standard is: Automation often creates new translation between generated output and the systems or people that must use it. If the same failure recurs, leaders can change delegation, capacity, intake, policy, or the intervention itself. That portfolio response is usually more valuable than asking each team to absorb another local workaround.

Two questions test find boundary work. What evidence would show that the condition is no longer adequate? Would that evidence arrive early enough to preserve a feasible choice? The intended choice is clear: Include boundary work in the value case and redesign the interface or process when it dominates. A weak answer to either question calls for a smaller commitment, an earlier signal, or a safer fallback. A strong answer earns movement without claiming that uncertainty has disappeared.

Design handoffs deliberately

Every transfer needs a clear object, receiving condition, owner, and response expectation. The framework in this paper is designed to make that boundary usable in ordinary management cadence.

Automation accelerates production into a downstream queue that lacks context, capacity, or authority to act. The problem may remain invisible for several review cycles because effort masks the missing condition. By the time variance appears, inexpensive options have already expired.

Specify the artifact, metadata, evidence, acceptance criteria, route, service expectation, and rejection path for each transfer. Proportionality is essential. Reversible, low-consequence work can proceed with a lighter record than a commitment that is costly to undo.

Queue age, incomplete handoffs, clarification requests, acceptance rate, and downstream idle or overload. Sampling should include the awkward cases. Averages built from the easy path are least informative where governance is most needed.

Throttle or reshape upstream automation when the receiving system cannot absorb the volume. The review should end with a changed commitment or a reasoned decision to keep it—not a request for more color commentary.

A tabletop review of design handoffs deliberately should use the observed break as its scenario: Automation accelerates production into a downstream queue that lacks context, capacity, or authority to act. Participants then apply the proposed response. They specify the artifact, metadata, evidence, acceptance criteria, route, service expectation, and rejection path for each transfer. The review identifies where authority, information, time, or recovery would run out. It is useful only if it changes the boundary or confirms who will act when the condition appears.

The minimum record for design handoffs deliberately is anchored in actual proof. It includes queue age, incomplete handoffs, clarification requests, acceptance rate, and downstream idle or overload. The record also carries the choice supported by that proof: Throttle or reshape upstream automation when the receiving system cannot absorb the volume. Keeping evidence and decision together lets a later owner distinguish a deliberate residual risk from a premise that nobody examined. It also gives the team a clear reason to reopen the choice when the underlying facts change.

Across the work governed by operating executives, transformation leaders, workflow owners, and AI program leaders, occurrences of this pattern should be compared rather than treated as unrelated project stories. The positive standard is: Every transfer needs a clear object, receiving condition, owner, and response expectation. If the same failure recurs, leaders can change delegation, capacity, intake, policy, or the intervention itself. That portfolio response is usually more valuable than asking each team to absorb another local workaround.

Two questions test design handoffs deliberately. What evidence would show that the condition is no longer adequate? Would that evidence arrive early enough to preserve a feasible choice? The intended choice is clear: Throttle or reshape upstream automation when the receiving system cannot absorb the volume. A weak answer to either question calls for a smaller commitment, an earlier signal, or a safer fallback. A strong answer earns movement without claiming that uncertainty has disappeared.

Own exceptions and recovery

A faster happy path can increase the absolute number or complexity of cases requiring human coordination. Viewed from the portfolio, this is a resource-allocation problem before it becomes a delivery problem.

Exceptions arrive through informal channels, bounce between teams, and consume scarce experts without visibility. Reporting usually captures the symptom after it has been translated into schedule, cost, quality, or risk. It rarely captures the original unresolved choice.

Classify exceptions, assign response ownership, provide evidence, set escalation time, and define return to a known state. Good governance preserves options. It creates a legitimate route to narrow, stage, redirect, pause, or stop without labeling learning as failure.

Exception frequency, severity, queue, resolution time, repeat cause, and specialist capacity consumed. A lack of evidence is itself useful information if the next decision allows a smaller experiment, tighter boundary, or deliberate wait.

Constrain scale when recovery demand exceeds the capacity saved on ordinary cases. The response should occur while it can still change the outcome, not after the issue has been translated into an unrecoverable variance.

A tabletop review of own exceptions and recovery should use the observed break as its scenario: Exceptions arrive through informal channels, bounce between teams, and consume scarce experts without visibility. Participants then apply the proposed response. They classify exceptions, assign response ownership, provide evidence, set escalation time, and define return to a known state. The review identifies where authority, information, time, or recovery would run out. It is useful only if it changes the boundary or confirms who will act when the condition appears.

The minimum record for own exceptions and recovery is anchored in actual proof. It includes exception frequency, severity, queue, resolution time, repeat cause, and specialist capacity consumed. The record also carries the choice supported by that proof: Constrain scale when recovery demand exceeds the capacity saved on ordinary cases. Keeping evidence and decision together lets a later owner distinguish a deliberate residual risk from a premise that nobody examined. It also gives the team a clear reason to reopen the choice when the underlying facts change.

Across the work governed by operating executives, transformation leaders, workflow owners, and AI program leaders, occurrences of this pattern should be compared rather than treated as unrelated project stories. The positive standard is: A faster happy path can increase the absolute number or complexity of cases requiring human coordination. If the same failure recurs, leaders can change delegation, capacity, intake, policy, or the intervention itself. That portfolio response is usually more valuable than asking each team to absorb another local workaround.

Two questions test own exceptions and recovery. What evidence would show that the condition is no longer adequate? Would that evidence arrive early enough to preserve a feasible choice? The intended choice is clear: Constrain scale when recovery demand exceeds the capacity saved on ordinary cases. A weak answer to either question calls for a smaller commitment, an earlier signal, or a safer fallback. A strong answer earns movement without claiming that uncertainty has disappeared.

Measure the whole outcome

Value should appear in end-to-end cycle time, quality, capacity, coverage, decision speed, or risk; tool usage and task duration are supporting measures. A useful governance design therefore begins with the consequence leaders are prepared to accept.

The business case counts seconds saved at one step while total elapsed time and workload across functions do not improve. A local fix can keep the work moving while worsening the portfolio: another specialist is borrowed, another exception is tolerated, and another dependency is promised twice.

Baseline the complete path, measure displaced and new work, and sample quality before and after the change. The control is complete only when the response is known: who acts, how quickly, within what authority, and with which safe fallback.

End-to-end time, total touches, rework, queue distribution, outcome quality, adoption, and operating cost. The receiving operation must recognize the evidence as relevant. Technical success alone cannot prove that a business workflow can carry the change.

Scale when the workflow result improves and the remaining coordination is both visible and owned. That choice should be recorded in operational terms: what may proceed, what remains outside the boundary, who owns the next move, and what would reopen the decision.

A tabletop review of measure the whole outcome should use the observed break as its scenario: The business case counts seconds saved at one step while total elapsed time and workload across functions do not improve. Participants then apply the proposed response. They baseline the complete path, measure displaced and new work, and sample quality before and after the change. The review identifies where authority, information, time, or recovery would run out. It is useful only if it changes the boundary or confirms who will act when the condition appears.

The minimum record for measure the whole outcome is anchored in actual proof. It includes end-to-end time, total touches, rework, queue distribution, outcome quality, adoption, and operating cost. The record also carries the choice supported by that proof: Scale when the workflow result improves and the remaining coordination is both visible and owned. Keeping evidence and decision together lets a later owner distinguish a deliberate residual risk from a premise that nobody examined. It also gives the team a clear reason to reopen the choice when the underlying facts change.

Across the work governed by operating executives, transformation leaders, workflow owners, and AI program leaders, occurrences of this pattern should be compared rather than treated as unrelated project stories. The positive standard is: Value should appear in end-to-end cycle time, quality, capacity, coverage, decision speed, or risk; tool usage and task duration are supporting measures. If the same failure recurs, leaders can change delegation, capacity, intake, policy, or the intervention itself. That portfolio response is usually more valuable than asking each team to absorb another local workaround.

Two questions test measure the whole outcome. What evidence would show that the condition is no longer adequate? Would that evidence arrive early enough to preserve a feasible choice? The intended choice is clear: Scale when the workflow result improves and the remaining coordination is both visible and owned. A weak answer to either question calls for a smaller commitment, an earlier signal, or a safer fallback. A strong answer earns movement without claiming that uncertainty has disappeared.

How the elements interact

Map work between tasks and Distinguish productive coordination protect different parts of the same commitment. The unit of analysis is the path across people, systems, decisions, and queues—not the isolated activity the tool performs. Some interaction creates judgment, learning, trust, safety, or necessary integration; eliminating it can degrade the outcome. The control belongs at the point where reliance expands, not in a retrospective explanation after the outcome is known. If the operation encounters this failure—an automated drafting or classification step looks successful while handoff volume, waiting, clarification, and reconciliation remain untouched.—the response for distinguish productive coordination cannot compensate for the missing map work between tasks obligation. Reviewers should reconcile the evidence for both elements before they accept the next action: keep valuable coordination, redesign avoidable routing, and automate only the stable remainder.

Distinguish productive coordination and Find boundary work protect different parts of the same commitment. Some interaction creates judgment, learning, trust, safety, or necessary integration; eliminating it can degrade the outcome. Automation often creates new translation between generated output and the systems or people that must use it. The visible artifact is often complete before the operating condition behind it is credible. If the operation encounters this failure—all meetings, reviews, and handoffs are classified as waste without asking what uncertainty or accountability they resolve.—the response for find boundary work cannot compensate for the missing distinguish productive coordination obligation. Reviewers should reconcile the evidence for both elements before they accept the next action: include boundary work in the value case and redesign the interface or process when it dominates.

Find boundary work and Design handoffs deliberately protect different parts of the same commitment. Automation often creates new translation between generated output and the systems or people that must use it. Every transfer needs a clear object, receiving condition, owner, and response expectation. A useful governance design therefore begins with the consequence leaders are prepared to accept. If the operation encounters this failure—users copy, normalize, explain, and verify output manually while those activities remain outside the measured process.—the response for design handoffs deliberately cannot compensate for the missing find boundary work obligation. Reviewers should reconcile the evidence for both elements before they accept the next action: throttle or reshape upstream automation when the receiving system cannot absorb the volume.

Design handoffs deliberately and Own exceptions and recovery protect different parts of the same commitment. Every transfer needs a clear object, receiving condition, owner, and response expectation. A faster happy path can increase the absolute number or complexity of cases requiring human coordination. This is where a management ritual either becomes a real control or reveals itself as administrative theater. If the operation encounters this failure—automation accelerates production into a downstream queue that lacks context, capacity, or authority to act.—the response for own exceptions and recovery cannot compensate for the missing design handoffs deliberately obligation. Reviewers should reconcile the evidence for both elements before they accept the next action: constrain scale when recovery demand exceeds the capacity saved on ordinary cases.

Own exceptions and recovery and Measure the whole outcome protect different parts of the same commitment. A faster happy path can increase the absolute number or complexity of cases requiring human coordination. Value should appear in end-to-end cycle time, quality, capacity, coverage, decision speed, or risk; tool usage and task duration are supporting measures. The issue is not a shortage of activity; it is a shortage of inspectable commitments. If the operation encounters this failure—exceptions arrive through informal channels, bounce between teams, and consume scarce experts without visibility.—the response for measure the whole outcome cannot compensate for the missing own exceptions and recovery obligation. Reviewers should reconcile the evidence for both elements before they accept the next action: scale when the workflow result improves and the remaining coordination is both visible and owned.

Measure the whole outcome and Map work between tasks protect different parts of the same commitment. Value should appear in end-to-end cycle time, quality, capacity, coverage, decision speed, or risk; tool usage and task duration are supporting measures. The unit of analysis is the path across people, systems, decisions, and queues—not the isolated activity the tool performs. A strong operating model makes the hidden negotiation explicit while options are still available. If the operation encounters this failure—the business case counts seconds saved at one step while total elapsed time and workload across functions do not improve.—the response for map work between tasks cannot compensate for the missing measure the whole outcome obligation. Reviewers should reconcile the evidence for both elements before they accept the next action: automate only after leaders can see whether the proposed change removes or relocates coordination.

Put the model into the management cadence

A framework becomes useful when it changes the normal cadence of work. It should shape intake, funding, design, delivery, and operating review without creating a separate governance industry around itself. The following sequence is deliberately small:

EXHIBIT 2

Start with one decision, then calibrate from evidence

1. **Move 1**
   Map one high-volume workflow using observed cases rather than process documentation alone.
2. **Move 2**
   Label every coordination point as valuable, necessary but reducible, avoidable, or created by the intervention.
3. **Move 3**
   Pilot changes at the boundaries and exception path as well as the automated task.
4. **Move 4**
   Review end-to-end measures and specialist load before expanding volume or authority.

- Map one high-volume workflow using observed cases rather than process documentation alone.
- Label every coordination point as valuable, necessary but reducible, avoidable, or created by the intervention.

- Pilot changes at the boundaries and exception path as well as the automated task.
- Review end-to-end measures and specialist load before expanding volume or authority.

The first cycle should be treated as calibration. Reviewers should note which questions changed a decision, which evidence was unavailable, where authority was unclear, and which controls cost out of proportion to the exposure. That record improves the coordination topology map without pretending the first version will fit every workflow or investment.

Ownership should remain close to the consequence. The PMO or AI-governance function can maintain the method, surface cross-portfolio patterns, and challenge weak evidence. It should not absorb the business owner's accountability for the result or the executive's accountability for the commitment.

Measures that reveal whether the control works

- Exceptions or assumptions discovered after the latest useful decision date.
- Scarce specialist and executive attention consumed by recurring unresolved conditions.
- Decisions reopened because the original evidence, boundary, or rationale was unclear.
- Operating outcomes and recovery performance after reliance expanded.
- Time from a material signal to a consequential decision.
- Commitments narrowed, redirected, paused, or stopped before avoidable exposure grew.

These measures are diagnostic, not a universal scorecard. Their purpose is to identify where the operating model fails repeatedly and whether governance is preserving options earlier. A lower count can indicate improvement, but it can also indicate weak detection. Leaders should read the measures with case evidence and frontline observation.

The strongest counterargument

Coordination is not inherently waste. Complex work needs consultation and cross-functional judgment. The tax is coordination whose purpose is unclear, whose route is avoidable, or whose cost was shifted out of view. Redesign should preserve necessary collaboration.

A reasonable critic could also argue that experienced leaders already do this through judgment. Many do. The weakness is not the absence of judgment; it is the absence of a durable operating record when people, pressure, and conditions change. The coordination topology map makes the basis for judgment visible enough to challenge, execute, and revisit.

The shift

Automation can remove task effort while leaving—or increasing—the work required to route, reconcile, approve, explain, and recover. Leaders should map the coordination topology and measure the total workflow before declaring productivity or autonomy. The practical shift is from reporting activity to governing the conditions under which the organization relies on that activity.

That discipline does not make uncertainty disappear. It gives leaders a better place to stand: a named decision, evidence proportionate to consequence, an explicit boundary, a responsible owner, and a response when reality departs from the assumption. Those elements are enough to turn hidden operating risk into a choice the portfolio can make deliberately.

Notes and sources

[1] Harang Ju, "When Coordination Is Avoidable: A Monotonicity Analysis of Organizational Tasks," 2026 preprint. https://arxiv.org/abs/2602.18673

[2] NIST, "Artificial Intelligence Risk Management Framework: Generative Artificial Intelligence Profile," July 2024. https://nvlpubs.nist.gov/nistpubs/ai/nist.ai.600-1.pdf

About the author

Marco Policani is an enterprise portfolio, PMO, and AI operating-governance leader. He builds portfolio governance systems where AI carries the analytical load and named humans own every decision — an approach documented across case studies, walkthroughs, and working governance guides on his portfolio site.

Portfolio & governance library: policani.net/governance · Full portfolio: policani.net · LinkedIn: linkedin.com/in/marcpolicani