
WHITE PAPER · PORTFOLIO & DELIVERY GOVERNANCE

Decision Friction Is a Portfolio Risk

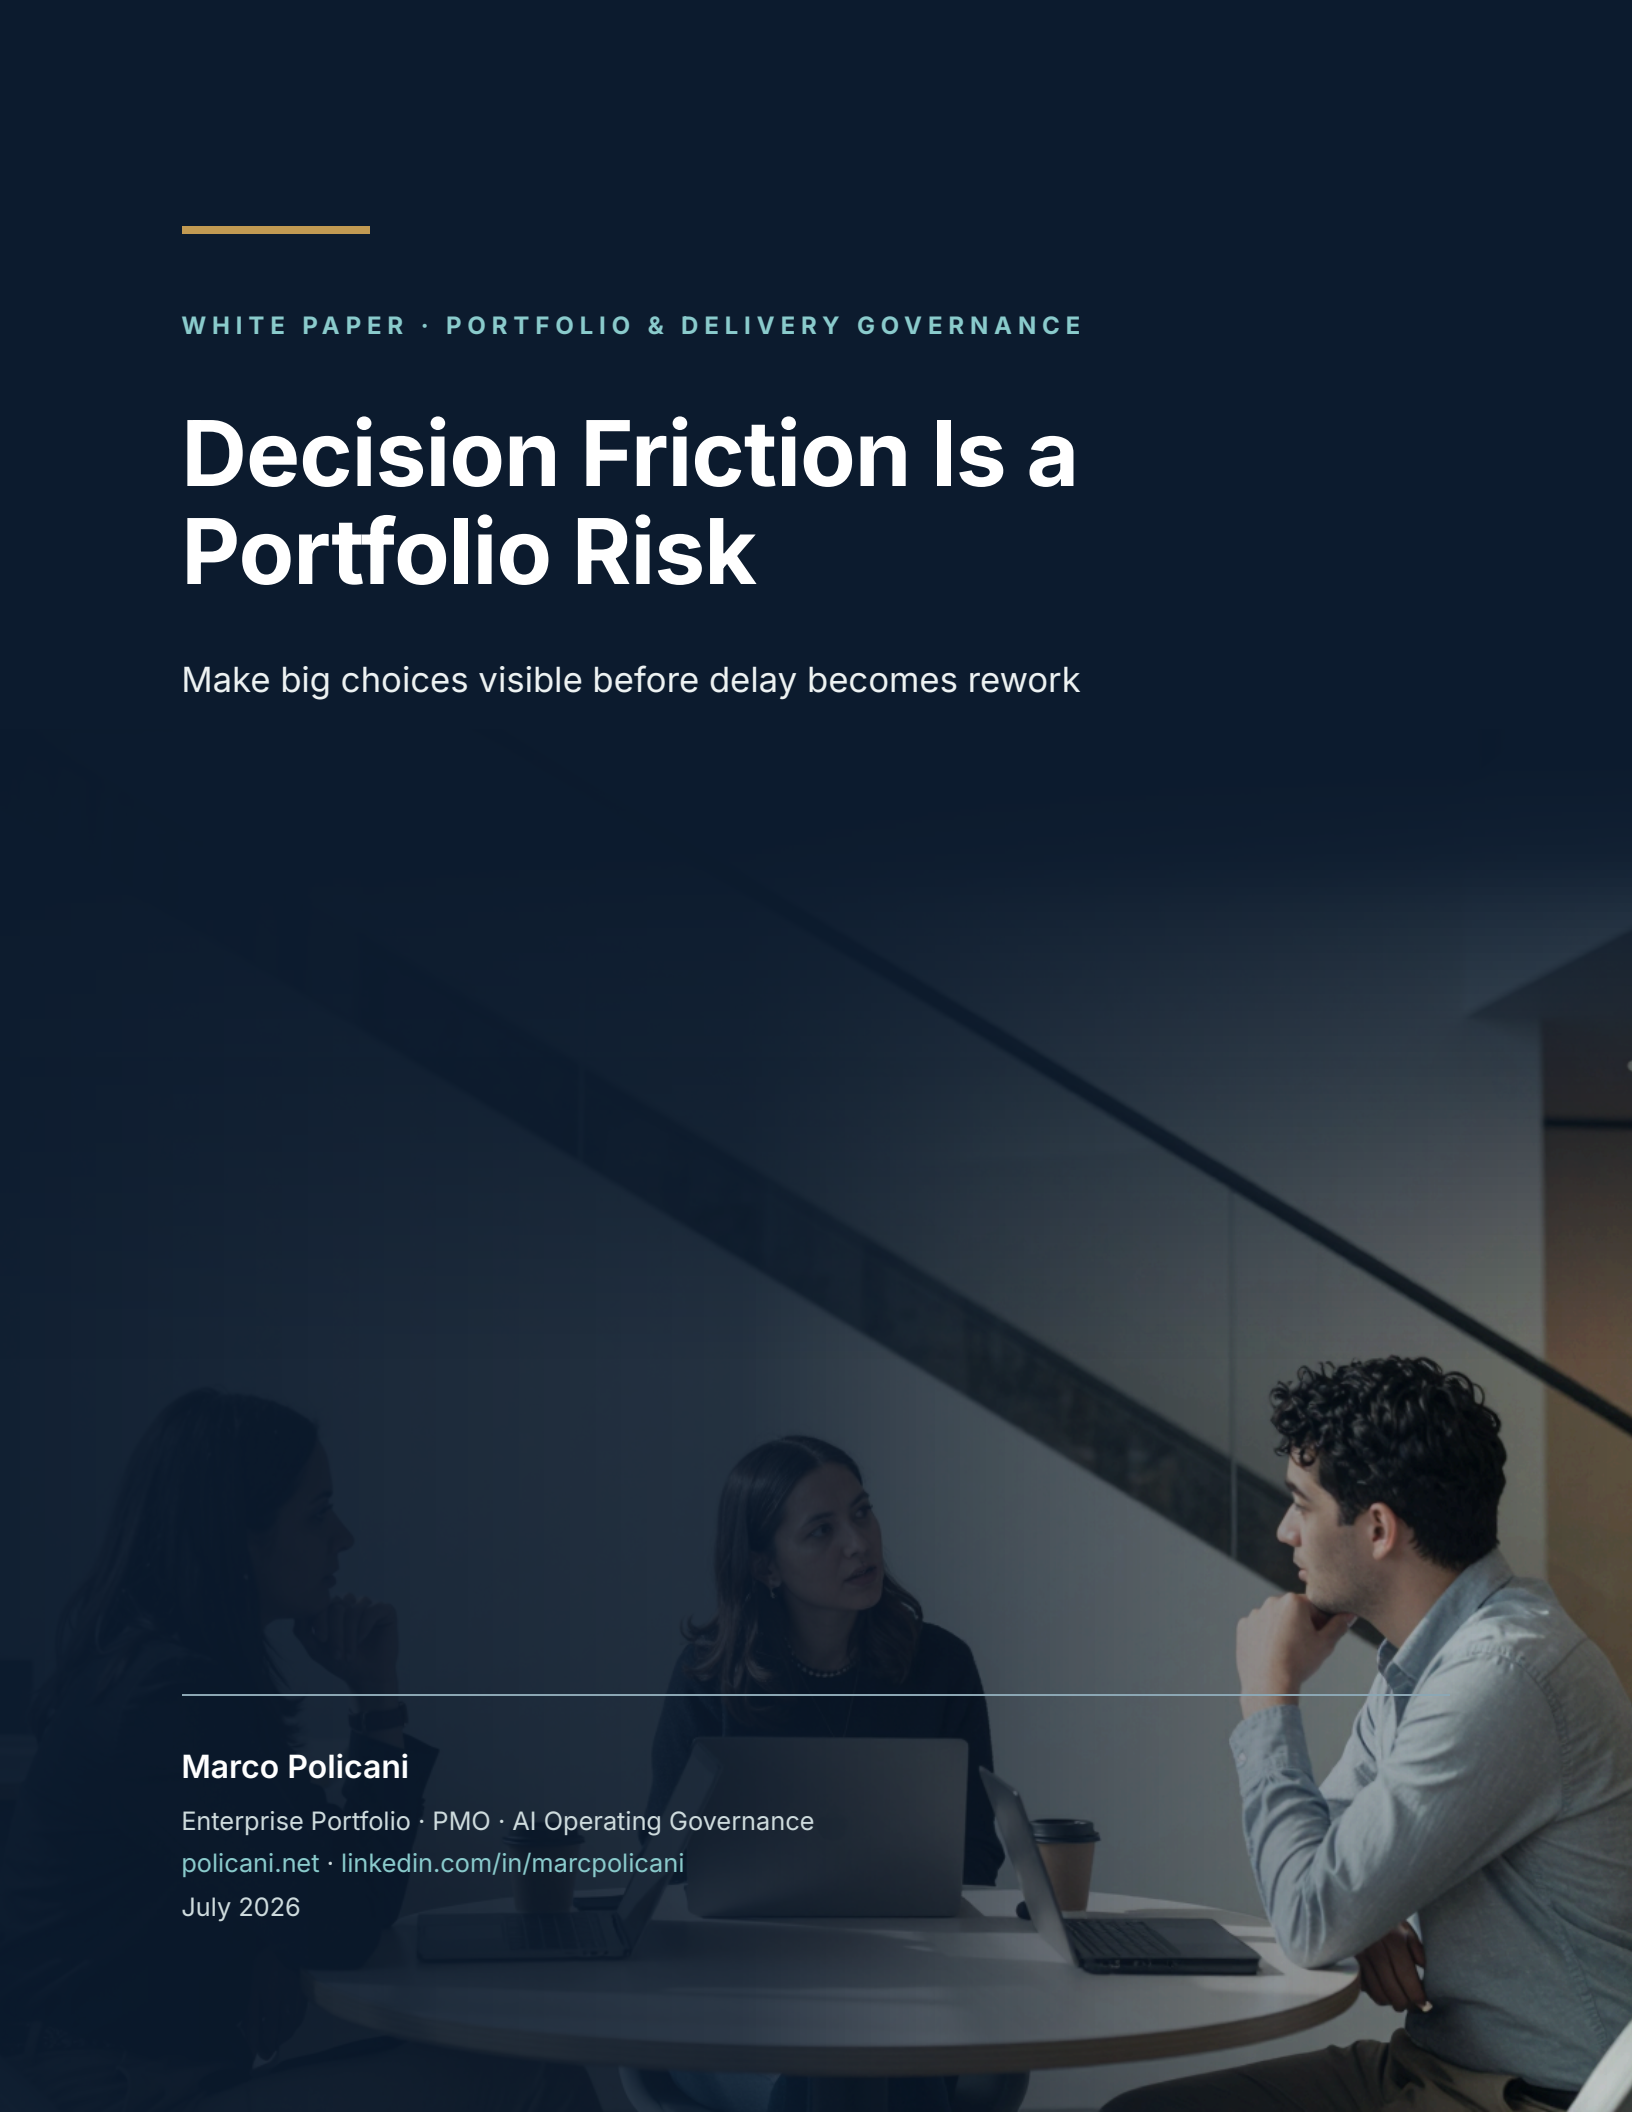
Make big choices visible before delay becomes rework

Marco Policani

Enterprise Portfolio · PMO · AI Operating Governance

policani.net · linkedin.com/in/marcpolicani

July 2026



EXECUTIVE SUMMARY

Decision friction becomes manageable when each choice that matters has an owner, evidence, a latest useful date, and a recorded outcome.

Portfolios track work, dates, spend, and risk. The decisions holding up that work often sit in meeting notes and inboxes. A decision inventory, a short list of the choices the work is waiting on, makes the queue visible before delay turns into rework or a more expensive default choice. The research cited is general, so the link from waiting to portfolio loss is a reading this paper offers.

- 01** Frame each entry as a choice someone can answer: the real options and a recommendation the deciding group can act on.
- 02** Give every decision an accountable owner and a latest useful date: the last point at which the answer can still change the work.
- 03** Add up delays across programs. The same policy question holding up three teams is one portfolio bottleneck.

The argument

Portfolios track work, dates, spend, and risk. The decisions that control the work often remain scattered across meetings, inboxes, and personal notes. When no one sees the queue, no one manages it.

Consider an invented late program, offered as an illustration and not as a reported case. Its post-mortem starts with the usual suspects: estimation, complexity, or vendor performance. Then someone walks through the calendar and finds that five weeks went to a policy question, a tradeoff, a funding answer, or an owner who was never named. The team raised the question early. It then moved from forum to forum, collecting opinions without producing a decision.

A late milestone can reflect poor execution. It can also reflect a choice the delivery team was never authorized to make: which policy applies, which tradeoff is acceptable, or who owns a process that crosses functions. A schedule report usually reveals that distinction only after the delay has become visible.

Decision friction is the time and effort required to turn a recognized question into an answer the work can use. It becomes a portfolio risk when teams wait, re-explain the question, or build against assumptions while the answer is still pending. This paper gives portfolio executives, PMO leaders, sponsors, and program leaders a practical control: a decision inventory with six disciplines. PMI's complex-project research finds that complexity erodes project value and often shows up as decision friction and rework,[1] and practitioner IT-governance writing treats good decision-making as what governance must embed in everyday work.[2] Managers report a related drag beyond project delivery. In a McKinsey survey, 61 percent said at least half of their decision-making time was used ineffectively.[3] Bain's research across more than a thousand companies found a clear correlation between decision effectiveness (quality, speed, yield, and appropriate effort) and business performance.[4][5] These sources are general: none shows that a waiting decision caused a portfolio loss, so the claim that decision friction is a portfolio risk is this paper's reading of them.

Speed is not the point. Some choices need more evidence, challenge, consultation, or legal review. The problem is unmanaged waiting: no owner is gathering evidence, no forum has been asked to decide, and the practical options narrow as the calendar moves on. "Cost of delay" is a phrase from lean product development for the economic loss that builds while something waits. This paper borrows the phrase as vocabulary only and makes no claim from product-development measurements. Many portfolios already meter the cost of delayed work; decision friction is that same meter turned on the choices the work is waiting for.

Why decisions escape governance

Work is easy to govern because it is visible. It has a name, a plan, a budget line, and someone who will be asked about it. A pending decision may appear only as a sentence in a status note or a risk that no one owns. Portfolio reporting captures the work that waits; it rarely captures the answer that would release it.

Escalation culture makes the blind spot worse. In this paper's reading, not a measured finding, teams may treat escalation as an admission that they failed or as a political attack on another leader. They try workarounds, then send a carefully worded update without a requested answer or a date. Executives receive background instead of a choice they can make.

The cost is portfolio-wide. Waiting teams spend time restating the question for each new forum. Some build against an assumed answer and later redo the work. Procurement windows close, specialists are reassigned, and the remaining options

become more expensive. The organization still makes a choice, but the calendar has made much of it already.

The portfolio decision inventory

The inventory applies the same visibility to choices that matter that portfolios already apply to work. Each entry states the choice, names the owner, identifies the evidence, sets a latest useful date, shows what is at risk, and records the outcome.

EXHIBIT 1

The decision record path

1. State the decision

Write down the choice and the options still on the table.

2. Name the owner

Assign it to the person authorized to make the decision.

3. Specify the evidence

Identify the information that could change the choice.

4. Latest useful date

Mark when delay closes options.

5. Expose downstream risk

Show the work waiting on the answer.

6. Close the loop

Record the answer, change the plan, and name what could reopen it.

A status report describes progress against a plan; a decision entry describes the answer needed to keep the plan credible. Keeping the two records beside one another makes a useful distinction visible: some work is late because delivery is struggling, while other work is late because the organization has not settled what delivery requires.

The record has to be useful to both sides of the conversation. Delivery teams need a way to raise a blocked choice without turning every issue into an executive escalation. Leaders need enough context to decide without asking the team to create a new business case each time. The six fields create that shared format: compact enough for a portfolio review, but specific enough to show who must act and what delay will affect.

State the decision

Start with one answerable choice. "Architecture concerns" is a theme. "Extend the billing platform or adopt the vendor module for the mid-market by March" gives a forum something it can decide. The difference is practical: one can be resolved in a meeting; the other creates a meeting about what the meeting is for.

Broad framing can be politically useful. It lets each participant hear a different question and delays the point at which someone must support a recommendation. The inventory requires the real alternatives, the recommended path, and the consequence of each option in operating terms.

Return themes for framing before they take up forum time. That takes PMO effort, but it gives decision-makers a usable request and rewards teams that do the work of defining it.

The recommendation should be visible, even when the requester is not the decision-maker. It tells the forum what tradeoff the team sees and gives the discussion a starting point. A request that lists only options may feel neutral, but it often asks senior leaders to reconstruct the analysis in the room. That slows the decision and makes the eventual reasoning harder to trace.

Name the accountable owner

Each entry needs the person who can bind the organization to an answer, or a defined route to that person. The owner needs authority over the area in dispute, not just proximity to the topic. A director whose remit ends at one function cannot settle a conflict with the neighboring function.

Naming the owner also improves the request. Teams have to locate the decision right before they escalate, and that often shows that they already have it. Some escalations are requests for political cover rather than authority; those should return to the team that can act.

A committee may advise. It is the owner only when its charter gives it authority to decide. A distribution list notifies people; it cannot own the decision. The useful question is simple: who can make this decision for this area? Record a real authority gap when the answer is unclear, then use the repeated gaps to fix the operating model.

An authority gap should not disappear into the next escalation. If the same kind of choice repeatedly moves between functions, the portfolio has learned something about its decision design. The practical response may be a clearer delegation, a standing forum with a defined remit, or a change to a sponsor's role. The inventory supplies the cases needed to make that change without depending on anecdotes.

Specify the evidence need

The evidence field distinguishes a decision that needs information from one that is avoiding a hard tradeoff. Record what is known, what analysis is underway, and the one question that matters: could the new information change the choice?

Before commissioning another study, ask which option the study could rule out and whether that result is worth the time. If no plausible result would alter the decision, analysis is only extending the queue.

Some uncertainty is real and should stay visible: ranges, dissent, and unknowns. In those cases, leaders must decide whether the value of more information exceeds the cost of waiting. Recording that tradeoff makes the judgment reviewable later.

This field also keeps evidence work proportional. A two-week spike may be the right response when migration effort is the one uncertainty that could reverse a recommendation. A six-month study is harder to justify when the likely result would only refine a decision that the team already needs to make. The aim is to make the purpose and time cost of analysis visible, not to discourage it.

Set the latest useful date

A due date marks when someone wants an answer. A latest useful date marks the last point at which the answer can still change the work. After that, downstream commitments have already narrowed the options.

Work backward from the first commitment that depends on the answer. When must procurement start to preserve the vendor window? When does the architect roll off? When do two design paths begin to diverge in cost? Use the earliest of those points, less the time needed to decide.

That date turns escalation from a vague request into an operational fact. "This choice loses its cheapest option on April 12" tells a forum what delay will cost. Track how

often decisions arrive before their latest useful date, then investigate the queues that do not.

Useful dates need to be tested after the fact. A date that repeatedly proves early may be a bid for attention. A date that repeatedly proves late may conceal slack until the last moment. Comparing the stated date with the actual downstream commitment helps the PMO improve both planning and the quality of future escalation.

Expose downstream risk

A delayed decision may look like one project's problem. When the same policy question blocks three programs, it is a portfolio bottleneck. The inventory makes that connection by showing each decision's investments, constrained roles, and dependencies.

The rollup should show what decision-makers need to act on: choices blocking the most value, owners with an aging queue, subjects that keep returning, and forums whose throughput is too slow for the work they govern.

One late answer can sound like a team's impatience. Forty days of delay against a fourteen-day useful window, repeated for three quarters, points to a mechanism that leadership owns.

Recurring choices deserve special attention. A data-residency question decided case by case five times should become a standing policy. A vendor-versus-build choice that recurs across programs may point to an architecture principle that has never been made explicit.

The rollup should also separate volume from consequence. A crowded queue can be inconvenient without putting much at risk. A smaller group of decisions may hold up a major investment, a scarce specialist, or a compliance commitment. Leaders need both views: where the queue is long and where delay is expensive. Treating them as the same problem produces busy reporting without better priorities.

Close the loop

A decision is complete when the people doing the work have changed their plans to reflect it. Meeting minutes alone do not achieve that. Teams can keep building against an old assumption for weeks when the answer reaches a distribution list but never reaches the plan.

Distribution is only part of the problem. Two functions can leave the same meeting with different readings: “we approved the vendor module” and “we approved it pending the security review.” Record the answer in operating language, along with its caveats, so every team is working from the same version.

Closure needs a short checklist: record the answer and reasoning, confirm that affected owners received it, update plans and budgets, and state what evidence would justify reopening the choice. That last item prevents both routine relitigation and decisions that outlive their assumptions.

The reasoning matters when circumstances change. “We chose the vendor module because volumes were under the threshold” gives a later volume surge a clear reopening path. Over time, recorded reasoning becomes useful institutional memory.

The record should name who must act after the decision, not merely who attended the meeting. A funding choice may require a finance partner to release a budget, a program lead to change a schedule, and an architecture team to stop design work on an alternative. Closure is complete only when those follow-on actions have an owner and the relevant plan reflects the decision.

One entry, end to end

A complete entry can stay short. Decision: adopt the vendor’s usage-based billing module for the mid-market, or extend the current platform. Owner: the revenue-technology VP, with escalation to the investment committee if the answer changes this year’s platform budget. Evidence: reference checks are complete and volume projections are available; the remaining uncertainty is migration effort, which a two-week spike is sizing. Latest useful date: April 12, when the vendor’s pricing lapses and the parallel design work forks. Downstream: the choice blocks the pricing pilot and keeps two integration engineers idle from May. Recommendation: adopt if the spike confirms migration under 400 hours.

A forum can make that decision or delegate it in fifteen minutes because the requesting team has already done the framing. The same question, presented as “we need alignment on billing architecture,” could take three meetings before anyone knows what choice is being requested.

The inventory does not require a long business case, a scoring model, or a new set of templates. It sets a quality bar for a decision conversation the organization already needs to have.

The example also shows why the fields belong together. A date without the downstream consequence can look arbitrary. Evidence without a named owner can sit in a queue forever. A recommendation without a reopening trigger can later look like an unexamined preference. The entry gives a forum enough context to decide while keeping the record short enough to use in a real meeting.

The disciplines as one instrument

The six disciplines work as a sequence. Weak framing creates an owner problem; an unclear owner makes dates meaningless; poor closure leaves the work running on competing assumptions. When the inventory fails, look first at the decision statement and the authority behind it.

EXHIBIT 2

Which choices belong in the inventory

Crosses teams

The answer requires authority or coordination beyond one delivery team.

Consumes scarce capacity

Delay idles a constrained role or forces work onto an assumption.

Sets a precedent

The choice will recur and may deserve a standing rule.

Is hard to reverse

Commitments will soon narrow the viable alternatives.

Use the inventory beside RACI charts and delegation matrices. Those tools show who should decide in the abstract. The inventory shows whether specific decisions are being made on time. A dated queue exposes who-decides rules that look clear on paper but do not work in practice.

Keep the scope narrow. Teams make hundreds of decisions inside their own authority every week, and an inventory that captures all of them becomes paperwork. Include choices that cross team lines, tie up constrained capacity, set a precedent, or are difficult to reverse. A portfolio should hold only the few dozen that meet that test.

The inventory does not replace local judgment. A delivery lead should still make ordinary sequencing choices without waiting for a portfolio forum. The test is whether

the choice needs authority, evidence, or coordination that the delivery team cannot supply on its own. If it does, leaving it unrecorded creates a queue that the portfolio cannot see.

Starting the inventory by Friday

The inventory starts small and earns its place:

- Begin with decisions that affect multiple investments or scarce roles.
- Put the inventory beside schedule and risk reporting, not in another disconnected log.
- Review queue age, missed latest-useful dates, repeated owners, and decisions reopened because records were weak.
- Use repeated gaps to change delegations, evidence standards, or forum design.

Build the first inventory by asking every program lead for the three questions holding up work today. Several will have been open for months. Some will already be past the point where the original alternatives were useful. That first review reveals where the organization has been absorbing delay without naming it.

Do not begin by asking teams to inventory every uncertainty. Start with decisions that have a real downstream commitment and a plausible owner. The first list should be short enough for leaders to discuss rather than admire. Once the forum resolves a few well-framed entries, teams can see the difference between reporting a problem and bringing a decision forward.

Leaders establish the escalation norm through their response to the first few dated requests. When a forum thanks the team, makes the decision, or gives a clear route to a decision-maker, other teams learn to raise issues early. When the response is silence or blame, dates disappear from future requests.

A missed useful date deserves a short review, not a blame ritual. Ask what prevented the decision: unclear ownership, a late evidence request, a forum that did not meet, or a choice that needed a different level of authority. The answer should change the queue or the operating model. Otherwise the organization is simply collecting another overdue metric.

Use the existing portfolio forum for a brief review of aging decisions and approaching dates. The PMO can improve the framing, owners can decide, and the forum can

handle what exceeds their authority. A separate decision board often adds another queue to the process it was meant to improve.

What the record starts to show

A few quarters of inventory data answer questions the organization could previously only argue about:

- Median days from decision framed to decision closed, by forum and by subject.
- Decisions resolved inside their latest useful dates, and the trend.
- Queue age, and which owners or forums hold the aging tail.
- Recurring subjects converted into standing policy, per quarter.
- Rework and idle capacity caused by late answers, priced where possible.
- Decisions reopened, and whether the reopening matched a recorded trigger.

The record also becomes institutional memory. New leaders inherit portfolios in the middle of decisions already made. A year of closed entries with their reasoning gives them a faster way to understand what was chosen and when the choice should be revisited.

Pricing delay is difficult, but it gives the inventory weight. Suppose, in an invented example, a team can show that a five-week policy wait created eleven person-weeks of rework. Leaders could then see the cost as more than an anecdote. Two or three well-supported cases each year may be enough to keep attention on the queue.

Metrics should support a decision, not decorate a dashboard. Median time from framed to closed can reveal a slow forum. The share of decisions resolved inside their useful dates can show whether escalation is working. Reopened decisions can reveal weak reasoning or a change in circumstances. Use each measure to ask what needs to change, rather than treating the measure as a score for the teams involved.

Over several quarters, the record can also guide capacity decisions. If the same small group of leaders appears on every aged item, the issue may be a span-of-control problem rather than individual responsiveness. If one forum sees a consistent surge before a funding cycle, the portfolio may need an earlier intake point. The inventory is useful because it grounds those changes in the decisions that teams were actually waiting on.

The honest limits

The inventory should not pressure people to rush choices that require deliberation, consultation, or legal review. Its purpose is to show the difference between necessary decision work and unmanaged waiting.

EXHIBIT 3

What the record starts to show

Latency

Median days from framed to closed, by forum and subject.

Reliability

Decisions resolved inside their latest useful dates, and the trend.

Bottlenecks

Which owners and forums hold the aging tail of the queue.

Policy conversion

Recurring decisions turned into standing rules each quarter.

It can also be gamed. Dates can be exaggerated to force attention, and owners can leave items “awaiting evidence” indefinitely. Review the record against outcomes. Dates that consistently prove wrong, and evidence requests that never eliminate an option, are signals the PMO can take back to the forum.

The inventory also has a practical limit: it cannot resolve a conflict that senior leaders prefer to avoid. In that situation, the record may become uncomfortable because it shows the same priority conflict aging week after week. That is still useful. It distinguishes a delivery-management problem from an executive choice that only executives can make.

That limit matters because the inventory is a control, not a substitute for leadership. It can improve framing, bring forward the real decision-maker, and document the cost of waiting. It cannot turn an unresolved conflict into agreement. When the record reaches that point, its value is clarity: the portfolio can stop disguising an executive priority call as a delivery delay.

From governing work to governing decisions

Portfolios already govern work. The next step is to govern the choices that control it with the same discipline: visible owner, evidence, useful date, downstream consequence, and recorded outcome.

Start with the two or three programs whose updates have been waiting for direction for months. Name the decision, the owner, the latest useful date, and the work being held up. Those cases give the inventory an immediate purpose.

As the queue becomes more reliable, fewer teams need to depend on informal escalation or personal memory to keep a question that matters alive, and fewer decisions disappear into a thread of status updates. Leaders have a shared place to intervene early, before commitments harden.

The result is less rework, fewer schedule slips attributed to the wrong cause, and fewer viable options lost while a question waits. A decisive organization is not one where leaders answer every question quickly. It is one where the necessary answers arrive before the work has moved past them.

Notes and sources

[1] Project Management Institute, "Pulse of the Profession® 2026: Driving Success in Complex Projects," May 2026. <https://www.pmi.org/learning/thought-leadership/driving-success-in-complex-projects>

[2] Mary K. Pratt, "IT Governance: Best Practices for Aligning IT and Business Strategy," CIO, July 2, 2026. <https://www.cio.com/article/272051/it-governance-a-formal-way-to-align-it-and-business-strategy.html>

[3] Iskandar Aminov, Aaron De Smet, Gregor Jost, and David Mendelsohn, "Decision making in the age of urgency," McKinsey & Company, April 30, 2019. <https://www.mckinsey.com/capabilities/people-and-organizational-performance/our-insights/decision-making-in-the-age-of-urgency>

[4] Marcia W. Blenko and Michael C. Mankins, "Measuring Decision Effectiveness," Bain & Company, June 5, 2012. <https://www.bain.com/insights/measuring-decision-effectiveness/>

[5] Marcia Blenko, Michael Mankins, and Paul Rogers, "The Five Steps to Better Decisions," Bain & Company, August 6, 2013. <https://www.bain.com/insights/the-five->

[steps-to-better-decisions/](#)

About the author

Marco Policani is an enterprise portfolio, PMO, and AI operating-governance leader. He builds portfolio governance systems where AI carries the analytical load and named humans own every decision — an approach documented across case studies, walkthroughs, and working governance guides on his portfolio site.

Portfolio & governance library: policani.net/governance · Full portfolio: policani.net · LinkedIn: linkedin.com/in/marcpolicani

This white paper may be shared freely with attribution. © 2026 Marco Policani.