

Decision Friction Is a Portfolio Risk

Make consequential choices visible before delay becomes rework

Marco Policani

Enterprise Portfolio · PMO · AI Operating Governance

policani.net · linkedin.com/in/marcpolicani

July 2026

EXECUTIVE SUMMARY

Decision friction is managed when the owner, evidence, and latest useful date are visible.

Portfolios track work in motion but often leave the decisions controlling that work scattered across meetings, inboxes, and individual memory. A decision inventory makes delay governable before ambiguity becomes rework, idle capacity, or an accidental choice. This paper gives portfolio executives, PMO leaders, sponsors, and program leaders a practical the portfolio decision inventory for making the next commitment explicit.

- 01** State the decision: A decision record begins with one answerable choice, not a topic for discussion.
- 02** Name the accountable owner: The owner must have authority over the choice or a defined route to the person who does.
- 03** Specify the evidence need: The inventory should show what is missing and whether it can realistically change the choice.
- 04** Set the latest useful date: A due date describes administration; a latest useful date describes when the decision can still shape the work.

CONTENTS

The argument	Expose downstream risk
Why the conventional approach breaks	Close the loop
The portfolio decision inventory	How the elements interact
State the decision	Put the model into the management cadence
Name the accountable owner	Measures that reveal whether the control works
Specify the evidence need	The strongest counterargument
Set the latest useful date	The shift

The argument

Portfolios track work in motion but often leave the decisions controlling that work scattered across meetings, inboxes, and individual memory. A decision inventory makes delay governable before ambiguity becomes rework, idle capacity, or an accidental choice.

A milestone can be late because a team executed poorly. It can also be late because the organization never settled a choice the team could not make: which policy applies, what tradeoff is acceptable, who owns the process, whether a dependency will be funded, or which outcome takes precedence. Conventional reporting usually discovers the distinction after schedule variance appears.

Decision friction is the elapsed time, repeated translation, and organizational energy consumed between recognizing that a choice is needed and making it in a form the work can use.

This paper is written for portfolio executives, PMO leaders, sponsors, and program leaders. It offers the portfolio decision inventory as a management instrument: a way to connect operating evidence to the next consequential choice. The cited sources provide external context; the framework and its application are Marco Policani's synthesis. It is not a predictive score and does not replace professional, legal, security, financial, or domain judgment.

Why the conventional approach breaks

The issue is not a shortage of activity; it is a shortage of inspectable commitments. Conventional governance tends to separate planning, approval, delivery reporting, and operating ownership. That separation allows a premise to pass through several forums without any one forum owning the decision it implies. The project or workflow then carries an assumption that was acknowledged socially but never accepted operationally.

When the consequence finally becomes visible, leaders face a narrower choice set and a more expensive recovery. The answer is not heavier control at every step. It is to place a clear, proportionate test at the point where the organization increases money, capacity, access, dependency, or action authority.

The portfolio decision inventory

The model has 6 connected elements. Each makes a different condition visible, but they should be reviewed together because strength in one area does not cancel a missing obligation in another. The model is designed for a short executive conversation supported by inspectable evidence, not a long maturity assessment.

EXHIBIT 1

The portfolio decision inventory connects evidence to action

- 1 **State the decision**
A decision record begins with one answerable choice, not a topic for discussion.
- 2 **Name the accountable owner**
The owner must have authority over the choice or a defined route to the person who does.
- 3 **Specify the evidence need**
The inventory should show what is missing and whether it can realistically change the choice.
- 4 **Set the latest useful date**
A due date describes administration; a latest useful date describes when the decision can still shape the work.

5

Expose downstream risk

Decision delay should be connected to capacity, dependencies, cost of change, and benefits—not hidden inside schedule commentary.

6

Close the loop

A decision is complete only when its rationale, conditions, and implications reach the people and records that govern execution.

State the decision

A decision record begins with one answerable choice, not a topic for discussion. The issue is not a shortage of activity; it is a shortage of inspectable commitments.

Teams escalate broad themes such as scope, architecture, or resources; leaders discuss them without knowing what commitment is requested. The missing condition is often treated as something the team will resolve later, even though the team lacks the authority or capacity to resolve it.

Write the alternatives, recommendation, and consequence in language that a named owner can accept, reject, or modify. The forum should be able to distinguish an acceptable residual risk from a missing obligation that has merely been renamed.

A concise decision statement linked to the work, the alternatives still feasible, and the downstream commitments affected. The evidence may be compact, but it must be inspectable. A future reviewer should be able to see what was examined, what conclusion was drawn, and which limitation remained.

Return vague items for framing before they consume executive review time. Leaders should prefer the smallest commitment that preserves learning and limits irreversible exposure.

A tabletop review of state the decision should use the observed break as its scenario: Teams escalate broad themes such as scope, architecture, or resources; leaders discuss them without knowing what commitment is requested. Participants then apply the proposed response. They write the alternatives, recommendation, and consequence in language that a named owner can accept, reject, or modify. The review identifies where authority, information, time, or recovery would run out. It is useful only if it changes the boundary or confirms who will act when the condition appears.

The minimum record for state the decision is anchored in actual proof. It includes a concise decision statement linked to the work, the alternatives still feasible, and the downstream commitments affected. The record also carries the choice supported by that proof: Return vague items for framing before they consume executive review time. Keeping evidence and decision together lets a later owner distinguish a deliberate residual risk from a premise that nobody examined. It also gives the team a clear reason to reopen the choice when the underlying facts change.

Across the work governed by portfolio executives, PMO leaders, sponsors, and program leaders, occurrences of this pattern should be compared rather than treated as unrelated project stories. The positive standard is: A decision record begins with one answerable choice, not a topic for discussion. If the same failure recurs, leaders can change delegation, capacity, intake, policy, or the intervention itself. That portfolio response is usually more valuable than asking each team to absorb another local workaround.

Two questions test state the decision. What evidence would show that the condition is no longer adequate? Would that evidence arrive early enough to preserve a feasible choice? The intended choice is clear: Return vague items for framing before they consume executive review time. A weak answer to either question calls for a smaller commitment, an earlier signal, or a safer fallback. A strong answer earns movement without claiming that uncertainty has disappeared.

Name the accountable owner

The owner must have authority over the choice or a defined route to the person who does. The cost appears downstream, but the missing decision sits upstream.

Responsibility is assigned to a committee, copied distribution list, or sponsor whose formal authority does not cover the contested boundary. The recurring signal is translation work: people spend time discovering what was meant, who may decide, and how to repair the unsupported assumption.

Map the decision to delegated rights and separate contributors, recommenders, and informed parties from the accountable decider. This does not require a larger pack. It requires a traceable connection among the premise, the evidence, the boundary, and the action that follows.

Owner acceptance, quorum where collective authority is required, and an escalation path for reserved matters. Evidence has a shelf life. The record needs the conditions or date that would make the conclusion stale.

Escalate promptly when the named owner cannot bind the organization. A hold is governable when the return conditions are explicit; otherwise it becomes an unmanaged queue.

A tabletop review of name the accountable owner should use the observed break as its scenario: Responsibility is assigned to a committee, copied distribution list, or sponsor whose formal authority does not cover the contested boundary. Participants then apply the proposed response. They map the decision to delegated rights and separate contributors, recommenders, and informed parties from the accountable decider. The review identifies where authority, information, time, or recovery would run out. It is useful only if it changes the boundary or confirms who will act when the condition appears.

The minimum record for name the accountable owner is anchored in actual proof. It includes owner acceptance, quorum where collective authority is required, and an escalation path for reserved matters. The record also carries the choice supported by that proof: Escalate promptly when the named owner cannot bind the organization. Keeping evidence and decision together lets a later owner distinguish a deliberate residual risk from a premise that nobody examined. It also gives the team a clear reason to reopen the choice when the underlying facts change.

Across the work governed by portfolio executives, PMO leaders, sponsors, and program leaders, occurrences of this pattern should be compared rather than treated as unrelated project stories. The positive standard is: The owner must have authority over the choice or a defined route to the person who does. If the same failure recurs, leaders can change delegation, capacity, intake, policy, or the intervention itself. That portfolio response is usually more valuable than asking each team to absorb another local workaround.

Two questions test name the accountable owner. What evidence would show that the condition is no longer adequate? Would that evidence arrive early enough to preserve a feasible choice? The intended choice is clear: Escalate promptly when the named owner cannot bind the organization. A weak answer to either question calls for

a smaller commitment, an earlier signal, or a safer fallback. A strong answer earns movement without claiming that uncertainty has disappeared.

Specify the evidence need

The inventory should show what is missing and whether it can realistically change the choice. The distinction matters because leaders authorize commitments, while teams inherit whatever the authorization left unresolved.

Requests for extra analysis become a safe way to defer, even when no new evidence would alter preferences or constraints. What looks like a team-level miss is frequently a portfolio choice that nobody made explicitly. Scarce attention is then spent reconciling consequences rather than choosing among alternatives.

List the decisive facts, assumptions, and uncertainties; distinguish evidence work from stakeholder negotiation. A named owner must be able to change the work; receiving a notification after the boundary is crossed is insufficient.

Sources, tests, estimates, ranges, dissent, and a clear statement of residual uncertainty. The evidence owner is accountable for integrity, while the decision owner remains accountable for the commitment. Keeping those roles distinct improves challenge.

Decide with bounded uncertainty when further analysis costs exceed its decision value. The residual uncertainty belongs in the decision record so later outcomes are compared with what leaders actually knew at the time.

A tabletop review of specify the evidence need should use the observed break as its scenario: Requests for extra analysis become a safe way to defer, even when no new evidence would alter preferences or constraints. Participants then apply the proposed response. They list the decisive facts, assumptions, and uncertainties; distinguish evidence work from stakeholder negotiation. The review identifies where authority, information, time, or recovery would run out. It is useful only if it changes the boundary or confirms who will act when the condition appears.

The minimum record for specify the evidence need is anchored in actual proof. It includes sources, tests, estimates, ranges, dissent, and a clear statement of residual uncertainty. The record also carries the choice supported by that proof: Decide with bounded uncertainty when further analysis costs exceed its decision value. Keeping evidence and decision together lets a later owner distinguish a deliberate residual risk from a premise that nobody examined. It also gives the team a clear reason to reopen the choice when the underlying facts change.

Across the work governed by portfolio executives, PMO leaders, sponsors, and program leaders, occurrences of this pattern should be compared rather than treated as unrelated project stories. The positive standard is: The inventory should show what is missing and whether it can realistically change the choice. If the same failure recurs, leaders can change delegation, capacity, intake, policy, or the intervention itself. That portfolio response is usually more valuable than asking each team to absorb another local workaround.

Two questions test specify the evidence need. What evidence would show that the condition is no longer adequate? Would that evidence arrive early enough to preserve a feasible choice? The intended choice is clear: Decide with bounded uncertainty when further analysis costs exceed its decision value. A weak answer to either question calls for a smaller commitment, an earlier signal, or a safer fallback. A strong answer earns movement without claiming that uncertainty has disappeared.

Set the latest useful date

A due date describes administration; a latest useful date describes when the decision can still shape the work. The control belongs at the point where reliance expands, not in a retrospective explanation after the outcome is known.

A choice is eventually made after procurement, design, staffing, or dependency commitments have already removed the viable alternatives. Momentum makes the weakness self-reinforcing. Each additional commitment raises the political and economic cost of revisiting the premise.

Work backward from the first downstream commitment and show the loss created by each period of delay. The record should travel with the work so that later participants can see the basis for reliance instead of reconstructing it from memory.

Decision date, lead times, option-expiry points, idle or rework exposure, and the escalation threshold. The strongest record shows both what supports the choice and what could overturn it.

Escalate before optionality disappears rather than after the milestone turns red. Decision quality improves when dissent and competing evidence are preserved instead of flattened into a unanimous-looking status.

A tabletop review of set the latest useful date should use the observed break as its scenario: A choice is eventually made after procurement, design, staffing, or dependency commitments have already removed the viable alternatives. Participants then apply the proposed response. They work backward from the first downstream commitment and show the loss created by each period of delay. The review identifies where authority, information, time, or recovery would run out. It is useful only if it changes the boundary or confirms who will act when the condition appears.

The minimum record for set the latest useful date is anchored in actual proof. It includes decision date, lead times, option-expiry points, idle or rework exposure, and the escalation threshold. The record also carries the choice supported by that proof: Escalate before optionality disappears rather than after the milestone turns red. Keeping evidence and decision together lets a later owner distinguish a deliberate residual risk from a premise that nobody examined. It also gives the team a clear reason to reopen the choice when the underlying facts change.

Across the work governed by portfolio executives, PMO leaders, sponsors, and program leaders, occurrences of this pattern should be compared rather than treated as unrelated project stories. The positive standard is: A due date describes administration; a latest useful date describes when the decision can still shape the work. If the same failure recurs, leaders can change delegation, capacity, intake, policy, or the intervention itself. That portfolio response is usually more valuable than asking each team to absorb another local workaround.

Two questions test set the latest useful date. What evidence would show that the condition is no longer adequate? Would that evidence arrive early enough to preserve a feasible choice? The intended choice is clear: Escalate before optionality disappears rather than after the milestone turns red. A weak answer to either question calls for a smaller commitment, an earlier signal, or a safer fallback. A strong answer earns movement without claiming that uncertainty has disappeared.

Expose downstream risk

Decision delay should be connected to capacity, dependencies, cost of change, and benefits—not hidden inside schedule commentary. This is where a management ritual either becomes a real control or reveals itself as administrative theater.

Several projects wait on the same policy or architecture choice, but each reports a local issue and the shared bottleneck stays invisible. Because responsibility is distributed, no single event appears large enough to stop the work. The accumulated operating load is the real exposure.

Link decisions across investments and aggregate recurring owners, subjects, and queues. Evidence should be selected for decision value. Volume is not rigor when the decisive uncertainty remains untested.

Affected commitments, constrained roles, delay propagation, workarounds, and repeated decision patterns. A later outcome should feed back into the evidence standard so repeated surprises improve the next commitment.

Resolve systemic bottlenecks at portfolio level instead of asking each team to chase the same answer. A decision without consequence invites the old behavior to continue. Capacity, funding, permissions, scope, or sequence must move with the answer.

A tabletop review of expose downstream risk should use the observed break as its scenario: Several projects wait on the same policy or architecture choice, but each reports a local issue and the shared bottleneck stays invisible. Participants then apply the proposed response. They link decisions across investments and aggregate recurring owners, subjects, and queues. The review identifies where authority, information, time, or recovery would run out. It is useful only if it changes the boundary or confirms who will act when the condition appears.

The minimum record for expose downstream risk is anchored in actual proof. It includes affected commitments, constrained roles, delay propagation, workarounds, and repeated decision patterns. The record also carries the choice supported by that proof: Resolve systemic bottlenecks at portfolio level instead of asking each team to chase the same answer. Keeping evidence and decision together lets a later owner distinguish a deliberate residual risk from a premise that nobody examined. It also gives the team a clear reason to reopen the choice when the underlying facts change.

Across the work governed by portfolio executives, PMO leaders, sponsors, and program leaders, occurrences of this pattern should be compared rather than treated as unrelated project stories. The positive standard is: Decision delay should be connected to capacity, dependencies, cost of change, and benefits—not hidden inside schedule commentary. If the same failure recurs, leaders can change delegation, capacity, intake, policy, or the intervention itself. That portfolio response is usually more valuable than asking each team to absorb another local workaround.

Two questions test expose downstream risk. What evidence would show that the condition is no longer adequate? Would that evidence arrive early enough to preserve a feasible choice? The intended choice is clear: Resolve systemic bottlenecks at portfolio level instead of asking each team to chase the same answer. A weak answer to either question calls for a smaller commitment, an earlier signal, or a safer fallback. A strong answer earns movement without claiming that uncertainty has disappeared.

Close the loop

A decision is complete only when its rationale, conditions, and implications reach the people and records that govern execution. The work becomes governable when ownership, evidence, boundaries, and response are connected to a real choice.

Leaders decide in a meeting, but the team continues under an older assumption or different functions interpret the outcome differently. When the consequence finally becomes visible, leaders face a narrower choice set and a more expensive recovery.

Record the answer, rationale, conditions, expiry, and changes required in plans, controls, and communications. The control should be short enough to use at the decision point and specific enough that a skeptical owner can challenge it.

Acknowledgment by affected owners, updated baselines or boundaries, and later evidence about whether the choice worked. Ranges and contradictory observations are often more useful than a precise point estimate built on weak assumptions.

Reopen the decision only when a named condition changes, not whenever disagreement resurfaces. The decision should identify the sacrifice. If every priority remains protected, the tradeoff has been delegated rather than resolved.

A tabletop review of close the loop should use the observed break as its scenario: Leaders decide in a meeting, but the team continues under an older assumption or different functions interpret the outcome differently. Participants then apply the proposed response. They record the answer, rationale, conditions, expiry, and changes required in plans, controls, and communications. The review identifies where authority, information, time, or recovery would run out. It is useful only if it changes the boundary or confirms who will act when the condition appears.

The minimum record for close the loop is anchored in actual proof. It includes acknowledgment by affected owners, updated baselines or boundaries, and later evidence about whether the choice worked. The record also carries the choice supported by that proof: Reopen the decision only when a named condition changes, not whenever disagreement resurfaces. Keeping evidence and decision together lets a later owner distinguish a deliberate residual risk from a premise that nobody examined. It also gives the team a clear reason to reopen the choice when the underlying facts change.

Across the work governed by portfolio executives, PMO leaders, sponsors, and program leaders, occurrences of this pattern should be compared rather than treated as unrelated project stories. The positive standard is: A decision is complete only when its rationale, conditions, and implications reach the people and records that govern execution. If the same failure recurs, leaders can change delegation, capacity, intake, policy, or the intervention itself. That portfolio response is usually more valuable than asking each team to absorb another local workaround.

Two questions test close the loop. What evidence would show that the condition is no longer adequate? Would that evidence arrive early enough to preserve a feasible choice? The intended choice is clear: Reopen the decision only when a named condition changes, not whenever disagreement resurfaces. A weak answer to either question calls for a smaller commitment, an earlier signal, or a safer fallback. A strong answer earns movement without claiming that uncertainty has disappeared.

How the elements interact

State the decision and Name the accountable owner protect different parts of the same commitment. A decision record begins with one answerable choice, not a topic for discussion. The owner must have authority over the choice or a defined route to the person who does. The work becomes governable when ownership, evidence, boundaries, and response are connected to a real choice. If the operation encounters this failure—teams escalate broad themes such as scope, architecture, or resources; leaders discuss them without knowing what commitment is requested.—the response for name the accountable owner cannot compensate for the missing state the decision obligation. Reviewers should reconcile the evidence for both elements before they accept the next action: escalate promptly when the named owner cannot bind the organization.

Name the accountable owner and Specify the evidence need protect different parts of the same commitment. The owner must have authority over the choice or a defined route to the person who does. The inventory should show what is missing and whether it can realistically change the choice. The cost appears downstream, but the missing decision sits upstream. If the operation encounters this failure—responsibility is assigned to a committee, copied distribution list, or sponsor whose formal authority does not cover the contested boundary.—the response for specify the evidence need cannot compensate for the missing name the accountable owner obligation. Reviewers should reconcile the evidence for both elements before they accept the next action: decide with bounded uncertainty when further analysis costs exceed its decision value.

Specify the evidence need and Set the latest useful date protect different parts of the same commitment. The inventory should show what is missing and whether it can realistically change the choice. A due date describes administration; a latest useful date describes when the decision can still shape the work. Teams move faster when uncertainty is bounded honestly than when it is concealed inside an optimistic plan. If the operation encounters this failure—requests for extra analysis become a safe way to defer, even when no new evidence would alter preferences or constraints.—the response for set the latest useful date cannot compensate for the missing specify the evidence need obligation. Reviewers should reconcile the evidence for both elements before they accept the next action: escalate before optionality disappears rather than after the milestone turns red.

Set the latest useful date and Expose downstream risk protect different parts of the same commitment. A due date describes administration; a latest useful date describes when the decision can still shape the work. Decision delay should be connected to capacity, dependencies, cost of change, and benefits—not hidden inside schedule commentary. The framework in this paper is designed to make that boundary usable in ordinary management cadence. If the operation encounters this failure—a choice is eventually made after procurement, design, staffing, or dependency commitments have already removed the viable alternatives.—the response for expose downstream risk cannot compensate for the missing set the latest useful date obligation. Reviewers should reconcile the evidence for both elements before they accept the next action: resolve systemic bottlenecks at portfolio level instead of asking each team to chase the same answer.

Expose downstream risk and Close the loop protect different parts of the same commitment. Decision delay should be connected to capacity, dependencies, cost of change, and benefits—not hidden inside schedule commentary. A decision is complete only when its rationale, conditions, and implications reach the people and records that govern execution. The distinction matters because leaders authorize commitments, while teams inherit whatever the authorization left unresolved. If the operation encounters this failure—several projects wait on the same policy or architecture choice, but each reports a local issue and the shared bottleneck stays invisible.—the response for close the loop cannot compensate for the missing expose downstream risk obligation. Reviewers should reconcile

the evidence for both elements before they accept the next action: reopen the decision only when a named condition changes, not whenever disagreement resurfaces.

Close the loop and State the decision protect different parts of the same commitment. A decision is complete only when its rationale, conditions, and implications reach the people and records that govern execution. A decision record begins with one answerable choice, not a topic for discussion. The practical failure occurs between the formal approval and the ordinary pressure of running the work. If the operation encounters this failure—leaders decide in a meeting, but the team continues under an older assumption or different functions interpret the outcome differently.—the response for state the decision cannot compensate for the missing close the loop obligation. Reviewers should reconcile the evidence for both elements before they accept the next action: return vague items for framing before they consume executive review time.

Put the model into the management cadence

A framework becomes useful when it changes the normal cadence of work. It should shape intake, funding, design, delivery, and operating review without creating a separate governance industry around itself. The following sequence is deliberately small:

EXHIBIT 2

Start with one decision, then calibrate from evidence

- 1 **Move 1**
Begin with decisions that affect multiple investments or constrained roles.
- 2 **Move 2**
Add the inventory beside schedule and risk reporting, not as another disconnected log.
- 3 **Move 3**
Review queue age, missed latest-useful dates, repeated owners, and decisions reopened because records were weak.
- 4 **Move 4**
Use recurring patterns to change delegations, evidence standards, or forum design.

- Begin with decisions that affect multiple investments or constrained roles.
- Add the inventory beside schedule and risk reporting, not as another disconnected log.
- Review queue age, missed latest-useful dates, repeated owners, and decisions reopened because records were weak.
- Use recurring patterns to change delegations, evidence standards, or forum design.

The first cycle should be treated as calibration. Reviewers should note which questions changed a decision, which evidence was unavailable, where authority was unclear, and which controls cost out of proportion to the exposure. That record improves the portfolio decision inventory without pretending the first version will fit every workflow or investment.

Ownership should remain close to the consequence. The PMO or AI-governance function can maintain the method, surface cross-portfolio patterns, and challenge weak evidence. It should not absorb the business owner's accountability for the result or the executive's accountability for the commitment.

Measures that reveal whether the control works

- Commitments narrowed, redirected, paused, or stopped before avoidable exposure grew.
- Exceptions or assumptions discovered after the latest useful decision date.
- Scarce specialist and executive attention consumed by recurring unresolved conditions.
- Decisions reopened because the original evidence, boundary, or rationale was unclear.
- Operating outcomes and recovery performance after reliance expanded.
- Time from a material signal to a consequential decision.

These measures are diagnostic, not a universal scorecard. Their purpose is to identify where the operating model fails repeatedly and whether governance is preserving options earlier. A lower count can indicate improvement, but it can also indicate weak detection. Leaders should read the measures with case evidence and frontline observation.

The strongest counterargument

Speed is not the sole objective. Some choices deserve deliberation, challenge, consultation, or legal review. The control distinguishes necessary decision work from unmanaged waiting and makes the consequence of elapsed time visible.

A reasonable critic could also argue that experienced leaders already do this through judgment. Many do. The weakness is not the absence of judgment; it is the absence of a durable operating record when people, pressure, and conditions change. The portfolio decision inventory makes the basis for judgment visible enough to challenge, execute, and revisit.

The shift

Portfolios track work in motion but often leave the decisions controlling that work scattered across meetings, inboxes, and individual memory. A decision inventory makes delay governable before ambiguity becomes rework, idle capacity, or an accidental choice. The practical shift is from reporting activity to governing the conditions under which the organization relies on that activity.

That discipline does not make uncertainty disappear. It gives leaders a better place to stand: a named decision, evidence proportionate to consequence, an explicit boundary, a responsible owner, and a response when reality departs from the assumption. Those elements are enough to turn hidden operating risk into a choice the portfolio can make deliberately.

Notes and sources

[1] Project Management Institute, "Pulse of the Profession 2026: Driving Success in Complex Projects," May 2026.
<https://www.pmi.org/learning/thought-leadership/driving-success-in-complex-projects>

[2] CIO, "IT Governance: Best Practices for Aligning IT and Business Strategy."
<https://www.cio.com/article/272051/it-governance-a-formal-way-to-align-it-and-business-strategy.html>

About the author

Marco Policani is an enterprise portfolio, PMO, and AI operating-governance leader. He builds portfolio governance systems where AI carries the analytical load and named humans own every decision — an approach documented across case studies, walkthroughs, and working governance guides on his portfolio site.

Portfolio & governance library: policani.net/governance · Full portfolio: policani.net · LinkedIn: linkedin.com/in/marcpolicani

This white paper may be shared freely with attribution. © 2026 Marco Policani.