
WHITE PAPER · PORTFOLIO & AI OPERATING GOVERNANCE

Who Gets to Decide

A clear path from evidence to action for portfolios and agent systems

Marco Policani

Enterprise Portfolio · PMO · AI Operating Governance

policani.net · linkedin.com/in/marcpolicani

August 2026

EXECUTIVE SUMMARY

Evidence changes work only when a named person can close the choice and a named owner carries it out.

Evidence does not change a portfolio or agent system by itself. For each choice that matters, name who builds the case, who can flag a legal, safety, or policy limit that blocks an option, who decides, and who makes the recorded change.

- 01** Name the choice before naming the owner.
- 02** Give each choice four jobs: one person recommends, the owner of a legal, safety, or policy limit agrees or blocks, one person or group decides, and one owner performs the change.
- 03** Build those jobs into intake, stage gates, the decision record, and the delivery system, not only into an org chart.
- 04** For AI agents, let them act inside approved limits while named people set the limits, decide exceptions, and own recovery.

The evidence is clear. No one can act.

A portfolio review can show that a project is late, its case has weakened, and another priority needs the same people. Everyone sees the conflict. The meeting still ends with a request for more analysis because no one can redirect, pause, or stop the work.

The dashboard was not the problem. The missing part was a decision path: who builds the case, who can flag a binding legal, safety, or policy constraint, who closes the choice, and who changes the work after the choice is made.

Rogers and Blenko's RAPID model separates recommending, agreeing, performing, giving input, and deciding.[1] A team needs a visible path from evidence to action.

Neilson, Martin, and Powers studied more than 125,000 employee profiles from more than 1,000 organizations. Decision rights and information flow were the two strongest execution levers in their analysis.[2] Stewart's readable account reports the same finding.[3]

Name the choice before naming the owner

An org chart says who leads a function. It does not say who can release capacity, accept a risk, approve a funding change, or stop an agent action. Start with the choice in plain words: release six months of capacity; allow an agent to issue a credit up to a limit; pause work after a missed target.

A useful statement names the choice, what it covers, what puts it on the table, and what changes after it is recorded. That turns a vague responsibility into a decision that a team can inspect.

Use four jobs

Recommend

One person prepares the case. That person shows the options, evidence, limits, and due date, then makes a recommendation. If the case is not ready, the decision is not ready.

EXHIBIT 1

Four jobs take evidence to action

- ① **Recommend**
Build the case, options, evidence, and due date.
- ② **Agree**
Flag a legal, safety, or policy limit that blocks an option, and offer an option or a route forward.
- ③ **Decide**
Close the choice using a clear person or group rule.
- ④ **Perform**
Change the funding, work, permission, or record.

Agree

Use an agree role only where someone carries a real obligation: a legal duty, safety rule, binding architecture limit, or similar constraint. The role can block an option, but it must name the constraint and offer an option or a route to escalation.

Decide

Most operating choices need one named decider. Some choices rightly belong to a board, safety body, works council, or other required group. In that case, name how the group reaches closure, handles a tie, and treats a missed deadline.

Perform

One owner changes the funding, backlog, staffing, permission, or record. This job confirms that the decision became work. PMI names unclear decision authority as one source of decision friction and rework.[4]

Put the jobs in the workflow

A role chart in a slide deck is documentation. The same four jobs in intake, a stage gate, a decision record, and the delivery system become a working control. A request should not cross an important boundary until the case, required constraints, decider, and performer are visible.

If an agree role rejects a plan, record the constraint, an acceptable option, and where the case goes next. If the decider misses the due date, pause, escalate, or allow a short extension with a visible reason.

Apply the rule to portfolio choices

Portfolio management includes choosing, ranking, speeding up, pausing, stopping, and moving resources among projects.[5] Each is a separate choice. Calling a portfolio committee the owner of all of them can create a bottleneck or push real decisions into side meetings.

Consider two projects that need the same specialist team. The portfolio office recommends a sequence using deadlines, dependencies, and capacity. A regulatory owner can block any option that breaks a binding rule. The executive who owns the combined operating result chooses. Delivery leaders change staffing, scope, and dates.

Apply the rule to agents

The same questions apply when agents work continuously. A human-in-the-loop statement does not say which person decides, at what limit, or what happens when that person does not respond. McKinsey and Deloitte advise leaders to set human accountability, limits, escalation paths, stop mechanisms, and clear treatment of exceptions.[6][7] That is practitioner guidance, not proof that this four-job model works everywhere.

EXHIBIT 2

An agent action still needs a decision path

- ① **Set the limit**
Name the action, threshold, and evidence the agent needs.
- ② **Act inside it**
Let the agent prepare a case or make an approved small change.
- ③ **Route the exception**
Send higher-risk actions to the named decider.
- ④ **Recover**
Name who can stop, restart, and repair the outcome.

An agent can gather evidence, recommend a refund, or perform a small credit inside an approved limit. Named people still set the limit, decide larger exceptions, and own recovery. A policy rule may stop an action outside the limit without gaining the right to restart it.

The limit belongs to the action, not to the agent as a whole. Several small actions can also add up to a large problem. Use a shared threshold for repeated credits, linked data changes, or repeated retries. Name the person who decides whether the fleet restarts. Deloitte also calls for override privileges and escalation paths in AI decision systems.[8]

Match the detail to the risk

Do not build a heavy contract for every small choice. Amazon describes hard-to-reverse choices as one-way doors and easier choices as two-way doors.[9] A low-

cost action inside a tested limit can stay local. A change that moves shared capacity, affects customers, or expands agent authority needs a clearer path.

Collective authority can also be the right design. The contract then names the group, quorum, decision rule, challenge path, and record. Eisenhardt's study of eight firms found that fast teams sought consensus first and then used a CEO-led choice when consensus did not come.[10]

What to do next

Pick five recent decisions. For each, ask: Who built the case? Who could block it for a real reason? Who closed it? Who changed the work? If an answer is missing, repair the role, limit, deadline, escalation path, or system record before the next decision arrives.

Clear roles do not guarantee a good result. They make the route from evidence to action visible, so the organization can see where it broke and fix the right part.

Notes and sources

[1] Paul Rogers and Marcia Blenko, "Who Has the D?" Bain & Company, June 28, 2006. <https://www.bain.com/insights/who-has-the-d/>

[2] Gary L. Neilson, Karla L. Martin, and Elizabeth Powers, "The Secrets to Successful Strategy Execution," Harvard Business Review 86, no. 6 (June 2008): 60–70. No reader-facing link is included because the full article is subscription-only.

[3] Thomas A. Stewart, "Putting Strategy into Practice," strategy+business, April 19, 2010. <https://www.strategy-business.com/article/00030>

[4] Project Management Institute, Pulse of the Profession 2026: Driving Success in Complex Projects — From Navigating Tasks to Navigating Systems, May 2026. <https://www.pmi.org/learning/thought-leadership/driving-success-in-complex-projects>

[5] Robert G. Cooper and Scott J. Edgett, "Portfolio Management: Fundamental for New Product Success," Stage-Gate International. <https://www.stage-gate.com/blog/portfolio-management-fundamental-for-new-product-success/>

- [6] Alexander Sukharevsky, Alexis Krivkovich, Arne Gast, Arsen Storozhev, Dana Maor, Deepak Mahadevan, Lari Hämäläinen, and Sandra Durth, "The Agentic Organization: Contours of the Next Paradigm for the AI Era," McKinsey & Company, September 26, 2025. <https://www.mckinsey.com/capabilities/people-and-organizational-performance/our-insights/the-agentic-organization-contours-of-the-next-paradigm-for-the-ai-era>
- [7] David Mallon, Brad Kreit, and Natasha Buckley, "Rethinking Operating Models for Humans with Agents," Deloitte Insights, April 2, 2026. <https://www.deloitte.com/us/en/insights/topics/talent/operating-models-for-humans-ai-agents.html>
- [8] David Mallon, Julie Duda, Stefano Besana, and Maya Bodan, "AI and the Future of Human Decision-Making," Deloitte Insights, March 3, 2026. <https://www.deloitte.com/us/en/insights/topics/talent/human-capital-trends/2026/decision-making-with-ai.html>
- [9] Jeffrey P. Bezos, Amazon 2015 Letter to Shareholders, 2016. https://s2.q4cdn.com/299287126/files/doc_financials/annual/2015-Letter-to-Shareholders.PDF
- [10] Kathleen M. Eisenhardt, "Making Fast Strategic Decisions in High-Velocity Environments," Academy of Management Journal 32, no. 3 (1989): 543–576. <https://doi.org/10.5465/256434>

About the author

Marco Policani is an enterprise portfolio, PMO, and AI operating-governance leader. He builds portfolio governance systems where AI carries the analytical load and named humans own every decision — an approach documented across case studies, walkthroughs, and working governance guides on his portfolio site.

Portfolio & governance library: policani.net/governance · Full portfolio: policani.net · LinkedIn: linkedin.com/in/marcpolicani

This white paper may be shared freely with attribution. © 2026 Marco Policani.