14

WHITE PAPER · AI OPERATING GOVERNANCE

The Demo-to-Production Gap

A demonstration is not operating evidence

Marco Policani

Enterprise Portfolio · PMO · AI Operating Governance

policani.net · linkedin.com/in/marcpolicani

July 2026

EXECUTIVE SUMMARY

Scale only when a governed workflow has earned evidence beyond the happy-path demonstration.

A demonstration proves that a capability can produce an appealing result under selected conditions. Production requires evidence that a real workflow can rely on it repeatedly—with accountable owners, representative cases, exceptions, controls, support, and a stop mechanism. This paper gives AI investment owners, product and technology leaders, workflow owners, and risk executives a practical the seven production questions for making the next commitment explicit.

01 Workflow fit: Production evidence begins with the complete job, including upstream inputs, downstream decisions, handoffs, and timing.

02 Operating ownership: A model owner cannot substitute for the process owner accountable for the operating consequence.

03 Representative cases: Happy-path accuracy says little about rare, ambiguous, adversarial, stale, or incomplete inputs.

04 Exception recovery: Production readiness includes detection, routing, human intervention, rollback, and return to a known state.

CONTENTS

The argument

A demonstration proves that a capability can produce an appealing result under selected conditions. Production requires evidence that a real workflow can rely on it repeatedly—with accountable owners, representative cases, exceptions, controls, support, and a stop mechanism.

Demos are built to compress uncertainty. They use selected inputs, knowledgeable operators, prepared prompts, and an audience focused on the successful path. That is appropriate for showing possibility. Trouble begins when possibility is treated as proof that ordinary users, ordinary data, and ordinary operational pressure will produce the same result.

Enterprise surveys from McKinsey and Deloitte describe broad experimentation alongside the continuing difficulty of turning AI activity into scaled operating value [1][2]. The gap is not solved by a better stage presentation. It is solved by different evidence.

This paper is written for AI investment owners, product and technology leaders, workflow owners, and risk executives. It offers the seven production questions as a management instrument: a way to connect operating evidence to the next consequential choice. The cited sources provide external context; the framework and its application are Marco Policani's synthesis. It is not a predictive score and does not replace professional, legal, security, financial, or domain judgment.

Why the conventional approach breaks

The cost appears downstream, but the missing decision sits upstream. Conventional governance tends to separate planning, approval, delivery reporting, and operating ownership. That separation allows a premise to pass through several forums without any one forum owning the decision it implies. The project or workflow then carries an assumption that was acknowledged socially but never accepted operationally.

The problem may remain invisible for several review cycles because effort masks the missing condition. By the time variance appears, inexpensive options have already expired. The answer is not heavier control at every step. It is to place a clear, proportionate test at the point where the organization increases money, capacity, access, dependency, or action authority.

The seven production questions

The model has 7 connected elements. Each makes a different condition visible, but they should be reviewed together because strength in one area does not cancel a missing obligation in another. The model is designed for a short executive conversation supported by inspectable evidence, not a long maturity assessment.

EXHIBIT 1

The seven production questions connects evidence to action

1. **Workflow fit**
   Production evidence begins with the complete job, including upstream inputs, downstream decisions, handoffs, and timing.
2. **Operating ownership**
   A model owner cannot substitute for the process owner accountable for the operating consequence.
3. **Representative cases**
   Happy-path accuracy says little about rare, ambiguous, adversarial, stale, or incomplete inputs.

4. **Exception recovery**
   Production readiness includes detection, routing, human intervention, rollback, and return to a known state.
5. **Observable controls**
   Leaders need evidence that boundaries and approvals work during actual use.
6. **Lifecycle support**
   Models, vendors, prompts, data, policies, and users change after launch.
7. **Stop authority**
   Production needs explicit pause, rollback, and retirement authority.

Workflow fit

Production evidence begins with the complete job, including upstream inputs, downstream decisions, handoffs, and timing. The work becomes governable when ownership, evidence, boundaries, and response are connected to a real choice.

The demo automates an attractive task while users still reconcile systems, rewrite output, chase approvals, or carry exceptions manually. When the consequence finally becomes visible, leaders face a narrower choice set and a more expensive recovery.

Map the end-to-end workflow and test whether the intervention removes work, moves it, or creates new coordination. The control should be short enough to use at the decision point and specific enough that a skeptical owner can challenge it.

Cycle time, touches, rework, queue time, adoption, and user observation from representative cases. Ranges and contradictory observations are often more useful than a precise point estimate built on weak assumptions.

Proceed only when the workflow outcome improves without shifting unacceptable burden elsewhere. The decision should identify the sacrifice. If every priority remains protected, the tradeoff has been delegated rather than resolved.

A tabletop review of workflow fit should use the observed break as its scenario: The demo automates an attractive task while users still reconcile systems, rewrite output, chase approvals, or carry exceptions manually. Participants then apply the proposed response. They map the end-to-end workflow and test whether the intervention removes work, moves it, or creates new coordination. The review identifies where authority, information, time, or recovery would run out. It is useful only if it changes the boundary or confirms who will act when the condition appears.

The minimum record for workflow fit is anchored in actual proof. It includes cycle time, touches, rework, queue time, adoption, and user observation from representative cases. The record also carries the choice supported by that proof: Proceed only when the workflow outcome improves without shifting unacceptable burden elsewhere. Keeping evidence and decision together lets a later owner distinguish a deliberate residual risk from a premise that nobody examined. It also gives the team a clear reason to reopen the choice when the underlying facts change.

Across the work governed by AI investment owners, product and technology leaders, workflow owners, and risk executives, occurrences of this pattern should be compared rather than treated as unrelated project stories. The positive standard is: Production evidence begins with the complete job, including upstream inputs, downstream decisions, handoffs, and timing. If the same failure recurs, leaders can change delegation, capacity, intake, policy, or the intervention itself. That portfolio response is usually more valuable than asking each team to absorb another local workaround.

Two questions test workflow fit. What evidence would show that the condition is no longer adequate? Would that evidence arrive early enough to preserve a feasible choice? The intended choice is clear: Proceed only when the workflow outcome improves without shifting unacceptable burden elsewhere. A weak answer to either question calls for a smaller commitment, an earlier signal, or a safer fallback. A strong answer earns movement without claiming that uncertainty has disappeared.

Operating ownership

A model owner cannot substitute for the process owner accountable for the operating consequence. The framework in this paper is designed to make that boundary usable in ordinary management cadence.

Technology teams own the pilot while no business leader accepts quality, policy, customer, or performance outcomes after launch. The problem may remain invisible for several review cycles because effort masks the missing condition. By the time variance appears, inexpensive options have already expired.

Name process, technical, data, risk, and support ownership and make the process owner decide the reliance boundary. Proportionality is essential. Reversible, low-consequence work can proceed with a lighter record than a commitment that is costly to undo.

Accepted responsibilities, service expectations, escalation routes, and decisions each owner can make. Sampling should include the awkward cases. Averages built from the easy path are least informative where governance is most needed.

Hold production use until the receiving operation owns both benefit and failure response. The review should end with a changed commitment or a reasoned decision to keep it—not a request for more color commentary.

A tabletop review of operating ownership should use the observed break as its scenario: Technology teams own the pilot while no business leader accepts quality, policy, customer, or performance outcomes after launch. Participants then apply the proposed response. They name process, technical, data, risk, and support ownership and make the process owner decide the reliance boundary. The review identifies where authority, information, time, or recovery would run out. It is useful only if it changes the boundary or confirms who will act when the condition appears.

The minimum record for operating ownership is anchored in actual proof. It includes accepted responsibilities, service expectations, escalation routes, and decisions each owner can make. The record also carries the choice supported by that proof: Hold production use until the receiving operation owns both benefit and failure response. Keeping evidence and decision together lets a later owner distinguish a deliberate residual risk from a premise that nobody examined. It also gives the team a clear reason to reopen the choice when the underlying facts change.

Across the work governed by AI investment owners, product and technology leaders, workflow owners, and risk executives, occurrences of this pattern should be compared rather than treated as unrelated project stories. The positive standard is: A model owner cannot substitute for the process owner accountable for the operating consequence. If the same failure recurs, leaders can change delegation, capacity, intake, policy, or the intervention itself. That portfolio response is usually more valuable than asking each team to absorb another local workaround.

Two questions test operating ownership. What evidence would show that the condition is no longer adequate? Would that evidence arrive early enough to preserve a feasible choice? The intended choice is clear: Hold production use until the receiving operation owns both benefit and failure response. A weak answer to either question calls for a smaller commitment, an earlier signal, or a safer fallback. A strong answer earns movement without claiming that uncertainty has disappeared.

Representative cases

Happy-path accuracy says little about rare, ambiguous, adversarial, stale, or incomplete inputs. Viewed from the portfolio, this is a resource-allocation problem before it becomes a delivery problem.

Evaluation uses a clean sample created by the build team and excludes the cases that consume most operating judgment. Reporting usually captures the symptom after it has been translated into schedule, cost, quality, or risk. It rarely captures the original unresolved choice.

Construct the case set from production history, user segments, exceptions, edge conditions, and foreseeable misuse. Good governance preserves options. It creates a legitimate route to narrow, stage, redirect, pause, or stop without labeling learning as failure.

Case provenance, coverage rationale, subgroup results, failure severity, and unresolved gaps. A lack of evidence is itself useful information if the next decision allows a smaller experiment, tighter boundary, or deliberate wait.

Bound permitted use to the tested case population and expand only as evidence expands. The response should occur while it can still change the outcome, not after the issue has been translated into an unrecoverable variance.

A tabletop review of representative cases should use the observed break as its scenario: Evaluation uses a clean sample created by the build team and excludes the cases that consume most operating judgment. Participants then apply the proposed response. They construct the case set from production history, user segments, exceptions, edge conditions, and foreseeable misuse. The review identifies where authority, information, time, or recovery would run out. It is useful only if it changes the boundary or confirms who will act when the condition appears.

The minimum record for representative cases is anchored in actual proof. It includes case provenance, coverage rationale, subgroup results, failure severity, and unresolved gaps. The record also carries the choice supported by that proof: Bound permitted use to the tested case population and expand only as evidence expands. Keeping evidence and decision together lets a later owner distinguish a deliberate residual risk from a premise that nobody examined. It also gives the team a clear reason to reopen the choice when the underlying facts change.

Across the work governed by AI investment owners, product and technology leaders, workflow owners, and risk executives, occurrences of this pattern should be compared rather than treated as unrelated project stories. The positive standard is: Happy-path accuracy says little about rare, ambiguous, adversarial, stale, or incomplete inputs. If the same failure recurs, leaders can change delegation, capacity, intake, policy, or the intervention itself. That portfolio response is usually more valuable than asking each team to absorb another local workaround.

Two questions test representative cases. What evidence would show that the condition is no longer adequate? Would that evidence arrive early enough to preserve a feasible choice? The intended choice is clear: Bound permitted use to the tested case population and expand only as evidence expands. A weak answer to either question calls for a smaller commitment, an earlier signal, or a safer fallback. A strong answer earns movement without claiming that uncertainty has disappeared.

Exception recovery

Production readiness includes detection, routing, human intervention, rollback, and return to a known state. A useful governance design therefore begins with the consequence leaders are prepared to accept.

The pilot celebrates average performance while failed cases disappear into a queue or require improvised technical rescue. A local fix can keep the work moving while worsening the portfolio: another specialist is borrowed, another exception is tolerated, and another dependency is promised twice.

Define exception classes, response owners, time limits, safe fallbacks, and recovery tests before volume grows. The control is complete only when the response is known: who acts, how quickly, within what authority, and with which safe fallback.

Exception rate and severity, detection latency, recovery time, backlog, repeat failures, and successful rollback exercises. The receiving operation must recognize the evidence as relevant. Technical success alone cannot prove that a business workflow can carry the change.

Pause expansion when the exception path consumes more capacity than the operating model can supply. That choice should be recorded in operational terms: what may proceed, what remains outside the boundary, who owns the next move, and what would reopen the decision.

A tabletop review of exception recovery should use the observed break as its scenario: The pilot celebrates average performance while failed cases disappear into a queue or require improvised technical rescue. Participants then apply the proposed response. They define exception classes, response owners, time limits, safe fallbacks, and recovery tests before volume grows. The review identifies where authority, information, time, or recovery would run out. It is useful only if it changes the boundary or confirms who will act when the condition appears.

The minimum record for exception recovery is anchored in actual proof. It includes exception rate and severity, detection latency, recovery time, backlog, repeat failures, and successful rollback exercises. The record also carries the choice supported by that proof: Pause expansion when the exception path consumes more capacity than the operating model can supply. Keeping evidence and decision together lets a later owner distinguish a deliberate residual risk from a premise that nobody examined. It also gives the team a clear reason to reopen the choice when the underlying facts change.

Across the work governed by AI investment owners, product and technology leaders, workflow owners, and risk executives, occurrences of this pattern should be compared rather than treated as unrelated project stories. The positive standard is: Production readiness includes detection, routing, human intervention, rollback, and return to a known state. If the same failure recurs, leaders can change delegation, capacity, intake, policy, or the intervention itself. That portfolio response is usually more valuable than asking each team to absorb another local workaround.

Two questions test exception recovery. What evidence would show that the condition is no longer adequate? Would that evidence arrive early enough to preserve a feasible choice? The intended choice is clear: Pause expansion when the exception path consumes more capacity than the operating model can supply. A weak answer to either question calls for a smaller commitment, an earlier signal, or a safer fallback. A strong answer earns movement without claiming that uncertainty has disappeared.

Observable controls

Leaders need evidence that boundaries and approvals work during actual use. A strong operating model makes the hidden negotiation explicit while options are still available.

A policy says sensitive data or consequential actions are restricted, but the system cannot show when the boundary was approached or crossed. A polished status can coexist with this failure because status describes activity; it does not prove the premise that authorizes reliance.

Instrument access, data movement, action authority, review, override, incidents, and changes. The aim is not certainty. It is a justified next commitment with an observable basis and a viable way back.

Logs tied to users and cases, review records, alerts, control tests, and accountable follow-up. Measures should expose changes in the operation rather than reward production of artifacts. Activity can support the story; it cannot substitute for the result.

Reduce authority when control performance cannot be inspected. When the premise is weak, narrowing is often a better decision than forcing a binary choice between full approval and cancellation.

A tabletop review of observable controls should use the observed break as its scenario: A policy says sensitive data or consequential actions are restricted, but the system cannot show when the boundary was approached or crossed. Participants then apply the proposed response. They instrument access, data movement, action authority, review, override, incidents, and changes. The review identifies where authority, information, time, or recovery would run out. It is useful only if it changes the boundary or confirms who will act when the condition appears.

The minimum record for observable controls is anchored in actual proof. It includes logs tied to users and cases, review records, alerts, control tests, and accountable follow-up. The record also carries the choice supported by that proof: Reduce authority when control performance cannot be inspected. Keeping evidence and decision together lets a later owner distinguish a deliberate residual risk from a premise that nobody examined. It also gives the team a clear reason to reopen the choice when the underlying facts change.

Across the work governed by AI investment owners, product and technology leaders, workflow owners, and risk executives, occurrences of this pattern should be compared rather than treated as unrelated project stories. The positive standard is: Leaders need evidence that boundaries and approvals work during actual use. If the same failure recurs, leaders can change delegation, capacity, intake, policy, or the intervention itself. That portfolio response is usually more valuable than asking each team to absorb another local workaround.

Two questions test observable controls. What evidence would show that the condition is no longer adequate? Would that evidence arrive early enough to preserve a feasible choice? The intended choice is clear: Reduce authority when control performance cannot be inspected. A weak answer to either question calls for a smaller commitment, an earlier signal, or a safer fallback. A strong answer earns movement without claiming that uncertainty has disappeared.

Lifecycle support

Models, vendors, prompts, data, policies, and users change after launch. Teams move faster when uncertainty is bounded honestly than when it is concealed inside an optimistic plan.

The launch team disbands while monitoring, evaluation maintenance, vendor changes, user support, and incident response have no funded capacity. Under pressure, the workaround is predictable: capable people compensate, local decisions proliferate, and the formal record stops describing how the operation actually runs.

Estimate standing operating work and assign lifecycle owners, budgets, service levels, and change triggers. The design should state what remains unknown. Acknowledged uncertainty can be bounded; disguised uncertainty is inherited by the operation.

Support volume, evaluation cadence, update history, vendor notices, issue age, and capacity consumed. Where consequence is high, independent checking or cross-functional challenge reduces the risk that advocacy becomes its own proof.

Do not scale a capability whose continuing obligations are unfunded. Escalation is appropriate when authority is missing, not as a routine substitute for clear delegation.

A tabletop review of lifecycle support should use the observed break as its scenario: The launch team disbands while monitoring, evaluation maintenance, vendor changes, user support, and incident response have no funded capacity. Participants then apply the proposed response. They estimate standing operating work and assign lifecycle owners, budgets, service levels, and change triggers. The review identifies where authority, information, time, or recovery would run out. It is useful only if it changes the boundary or confirms who will act when the condition appears.

The minimum record for lifecycle support is anchored in actual proof. It includes support volume, evaluation cadence, update history, vendor notices, issue age, and capacity consumed. The record also carries the choice supported by that proof: Do not scale a capability whose continuing obligations are unfunded. Keeping evidence and decision together lets a later owner distinguish a deliberate residual risk from a premise that nobody examined. It also gives the team a clear reason to reopen the choice when the underlying facts change.

Across the work governed by AI investment owners, product and technology leaders, workflow owners, and risk executives, occurrences of this pattern should be compared rather than treated as unrelated project stories. The positive standard is: Models, vendors, prompts, data, policies, and users change after launch. If the same failure recurs, leaders can change delegation, capacity, intake, policy, or the intervention itself. That portfolio response is usually more valuable than asking each team to absorb another local workaround.

Two questions test lifecycle support. What evidence would show that the condition is no longer adequate? Would that evidence arrive early enough to preserve a feasible choice? The intended choice is clear: Do not scale a capability whose continuing obligations are unfunded. A weak answer to either question calls for a smaller commitment, an earlier signal, or a safer fallback. A strong answer earns movement without claiming that uncertainty has disappeared.

Stop authority

Production needs explicit pause, rollback, and retirement authority. The practical failure occurs between the formal approval and the ordinary pressure of running the work.

No one knows which failure threshold justifies interruption, so weak performance persists while leaders debate. The happy path can survive this weakness. Ordinary variation exposes it: a disputed input, absent owner, competing commitment, or case that falls outside the assumed rule.

Predefine stop conditions, responsible owner, communication path, fallback process, and evidence required to resume. The smallest credible control is the one that changes behavior when the evidence is weak. Anything else is commentary.

Thresholds tested in exercises, successful fallback, incident decisions, and time to reduce exposure. Qualitative observation belongs beside quantitative measures when workflow behavior, adoption, or judgment explains why the number moved.

Stop or narrow use when evidence crosses the agreed boundary, regardless of adoption momentum. A stop or rollback is evidence that the control works when it protects greater value elsewhere in the portfolio.

A tabletop review of stop authority should use the observed break as its scenario: No one knows which failure threshold justifies interruption, so weak performance persists while leaders debate. Participants then apply the proposed response. They predefine stop conditions, responsible owner, communication path, fallback process, and evidence required to resume. The review identifies where authority, information, time, or recovery would run out. It is useful only if it changes the boundary or confirms who will act when the condition appears.

The minimum record for stop authority is anchored in actual proof. It includes thresholds tested in exercises, successful fallback, incident decisions, and time to reduce exposure. The record also carries the choice supported by that proof: Stop or narrow use when evidence crosses the agreed boundary, regardless of adoption momentum. Keeping evidence and decision together lets a later owner distinguish a deliberate residual risk from a premise that nobody examined. It also gives the team a clear reason to reopen the choice when the underlying facts change.

Across the work governed by AI investment owners, product and technology leaders, workflow owners, and risk executives, occurrences of this pattern should be compared rather than treated as unrelated project stories. The positive standard is: Production needs explicit pause, rollback, and retirement authority. If the same failure recurs, leaders can change delegation, capacity, intake, policy, or the intervention itself. That portfolio response is usually more valuable than asking each team to absorb another local workaround.

Two questions test stop authority. What evidence would show that the condition is no longer adequate? Would that evidence arrive early enough to preserve a feasible choice? The intended choice is clear: Stop or narrow use when evidence crosses the agreed boundary, regardless of adoption momentum. A weak answer to either question calls for a smaller commitment, an earlier signal, or a safer fallback. A strong answer earns movement without claiming that uncertainty has disappeared.

How the elements interact

Workflow fit and Operating ownership protect different parts of the same commitment. Production evidence begins with the complete job, including upstream inputs, downstream decisions, handoffs, and timing. A model owner cannot substitute for the process owner accountable for the operating consequence. The framework in this paper is designed to make that boundary usable in ordinary management cadence. If the operation encounters this failure—the demo automates an attractive task while users still reconcile systems, rewrite output, chase approvals, or carry exceptions manually.—the response for operating ownership cannot compensate for the missing workflow fit obligation. Reviewers should reconcile the evidence for both elements before they accept the next action: hold production use until the receiving operation owns both benefit and failure response.

Operating ownership and Representative cases protect different parts of the same commitment. A model owner cannot substitute for the process owner accountable for the operating consequence. Happy-path accuracy says little about rare, ambiguous, adversarial, stale, or incomplete inputs. The distinction matters because leaders authorize commitments, while teams inherit whatever the authorization left unresolved. If the operation encounters this failure—technology teams own the pilot while no business leader accepts quality, policy, customer, or performance outcomes after launch.—the response for representative cases cannot compensate for the missing operating ownership obligation. Reviewers should reconcile the evidence for both elements before they accept the next action: bound permitted use to the tested case population and expand only as evidence expands.

Representative cases and Exception recovery protect different parts of the same commitment. Happy-path accuracy says little about rare, ambiguous, adversarial, stale, or incomplete inputs. Production readiness includes detection, routing, human intervention, rollback, and return to a known state. The practical failure occurs between the formal approval and the ordinary pressure of running the work. If the operation encounters this failure—evaluation uses a clean sample created by the build team and excludes the cases that consume most operating judgment.—the response for exception recovery cannot compensate for the missing representative cases obligation. Reviewers should reconcile the evidence for both elements before they accept the next action: pause expansion when the exception path consumes more capacity than the operating model can supply.

Exception recovery and Observable controls protect different parts of the same commitment. Production readiness includes detection, routing, human intervention, rollback, and return to a known state. Leaders need evidence that boundaries and approvals work during actual use. Viewed from the portfolio, this is a resource-allocation problem before it becomes a delivery problem. If the operation encounters this failure—the pilot celebrates average performance while failed cases disappear into a queue or require improvised technical rescue.—the response for observable controls cannot compensate for the missing exception recovery obligation. Reviewers should reconcile the evidence for both elements before they accept the next action: reduce authority when control performance cannot be inspected.

Observable controls and Lifecycle support protect different parts of the same commitment. Leaders need evidence that boundaries and approvals work during actual use. Models, vendors, prompts, data, policies, and users change after launch. The control belongs at the point where reliance expands, not in a retrospective explanation after the outcome is known. If the operation encounters this failure—a policy says sensitive data or consequential actions are restricted, but the system cannot show when the boundary was approached or crossed.—the response for lifecycle support cannot compensate for the missing observable controls obligation. Reviewers should reconcile the evidence for both elements before they accept the next action: do not scale a capability whose continuing obligations are unfunded.

Lifecycle support and Stop authority protect different parts of the same commitment. Models, vendors, prompts, data, policies, and users change after launch. Production needs explicit pause, rollback, and retirement authority. The visible artifact is often complete before the operating condition behind it is credible. If the operation encounters this failure—the launch team disbands while monitoring, evaluation maintenance, vendor changes, user support, and incident response have no funded capacity.—the response for stop authority cannot compensate for the missing lifecycle support obligation. Reviewers should reconcile the evidence for both elements before they accept the next action: stop or narrow use when evidence crosses the agreed boundary, regardless of adoption momentum.

Stop authority and Workflow fit protect different parts of the same commitment. Production needs explicit pause, rollback, and retirement authority. Production evidence begins with the complete job, including upstream inputs, downstream decisions, handoffs, and timing. A useful governance design therefore begins with the consequence leaders are prepared to accept. If the operation encounters this failure—no one knows which failure threshold justifies interruption, so weak performance persists while leaders debate.—the response for workflow fit cannot compensate for the missing stop authority obligation. Reviewers should reconcile the evidence for both elements before they accept the next action: proceed only when the workflow outcome improves without shifting unacceptable burden elsewhere.

Put the model into the management cadence

A framework becomes useful when it changes the normal cadence of work. It should shape intake, funding, design, delivery, and operating review without creating a separate governance industry around itself. The following sequence is deliberately small:

EXHIBIT 2

Start with one decision, then calibrate from evidence

1. **Move 1**
   Replace the demo approval with a production-evidence plan.
2. **Move 2**
   Pilot inside the complete workflow with representative users and cases.
3. **Move 3**
   Fund exception handling, evaluation, support, and incident response as standing work.
4. **Move 4**
   Require a reliance decision that states permitted use, boundaries, controls, owner, and stop conditions.

- Replace the demo approval with a production-evidence plan.
- Pilot inside the complete workflow with representative users and cases.
- Fund exception handling, evaluation, support, and incident response as standing work.
- Require a reliance decision that states permitted use, boundaries, controls, owner, and stop conditions.

The first cycle should be treated as calibration. Reviewers should note which questions changed a decision, which evidence was unavailable, where authority was unclear, and which controls cost out of proportion to the exposure. That record improves the seven production questions without pretending the first version will fit every workflow or investment.

Ownership should remain close to the consequence. The PMO or AI-governance function can maintain the method, surface cross-portfolio patterns, and challenge weak evidence. It should not absorb the business owner's accountability for the result or the executive's accountability for the commitment.

Measures that reveal whether the control works

- Commitments narrowed, redirected, paused, or stopped before avoidable exposure grew.
- Exceptions or assumptions discovered after the latest useful decision date.
- Scarce specialist and executive attention consumed by recurring unresolved conditions.
- Decisions reopened because the original evidence, boundary, or rationale was unclear.
- Operating outcomes and recovery performance after reliance expanded.
- Time from a material signal to a consequential decision.

These measures are diagnostic, not a universal scorecard. Their purpose is to identify where the operating model fails repeatedly and whether governance is preserving options earlier. A lower count can indicate improvement, but it can also indicate weak detection. Leaders should read the measures with case evidence and frontline observation.

The strongest counterargument

Not every experiment needs production-grade controls. Exploration can be fast and bounded when output cannot create real consequence. The error is allowing experimental evidence to authorize production reliance without crossing a deliberate operating gate.

A reasonable critic could also argue that experienced leaders already do this through judgment. Many do. The weakness is not the absence of judgment; it is the absence of a durable operating record when people, pressure, and conditions change. The seven production questions makes the basis for judgment visible enough to challenge, execute, and revisit.

The shift

A demonstration proves that a capability can produce an appealing result under selected conditions. Production requires evidence that a real workflow can rely on it repeatedly—with accountable owners, representative cases, exceptions, controls, support, and a stop mechanism. The practical shift is from reporting activity to governing the conditions under which the organization relies on that activity.

That discipline does not make uncertainty disappear. It gives leaders a better place to stand: a named decision, evidence proportionate to consequence, an explicit boundary, a responsible owner, and a response when reality departs from the assumption. Those elements are enough to turn hidden operating risk into a choice the portfolio can make deliberately.

Notes and sources

[1] Alex Singla et al., "The state of AI in 2025: Agents, innovation, and transformation," McKinsey, November 2025. https://www.mckinsey.com/capabilities/quantumblack/our-insights/the-state-of-ai

[2] Deloitte AI Institute, "The State of AI in the Enterprise," 2026. https://www.deloitte.com/us/en/what-we-do/capabilities/applied-artificial-intelligence/content/state-of-ai-in-the-enterprise.html

About the author

Marco Policani is an enterprise portfolio, PMO, and AI operating-governance leader. He builds portfolio governance systems where AI carries the analytical load and named humans own every decision — an approach documented across case studies, walkthroughs, and working governance guides on his portfolio site.

Portfolio & governance library: policani.net/governance · Full portfolio: policani.net · LinkedIn: linkedin.com/in/marcpolicani