


The Write-Up You Skipped Never Gets Blamed

The next team stops or guesses, and the delay gets logged as something else



Marco Policani

Enterprise Portfolio · PMO · AI Operating Governance

policani.net · linkedin.com/in/marcpolicani

July 2026

EXECUTIVE SUMMARY

Skipping the write-up moves the cost into the next team's queue, where it gets logged as a clarification or a redesign — never as the write-up nobody did.

Everyone in the meeting already knows what happened. The next team does not, and the write-up was what should have carried it across.

- 01** The cost lands in another team's queue under another name, so nobody traces it back to the meeting and nobody fixes it.
- 02** The next team decides when the handoff is done: accept it when they can act from the record without reconstructing your meeting.
- 03** A record people rely on needs a name attached to it and a date to revisit it; an AI system inherits whatever the page left out.

The trade that looks free

The schedule is slipping, so someone offers to start the handoff now and write up the decision later.

If you run the program, that sounds fine. The meeting did its job — everyone in the room agrees. The delivery work is urgent. Writing it up feels like the part that can wait.

But the next team wasn't in the room. They need to know what changed, why it changed, what you agreed to bend, and who can change it again. Without that, the person picking up the work has two choices: stop and ask, or start and guess. Both put more pressure on the schedule you were trying to protect.

A handoff moves more than work. It has to move what the next team needs to know in order to act without guessing, and the written decision is how that travels. That is why writing it later does not save time. The hour comes back as rework.

So does starting sooner actually save time, once the decision has to cross into someone else's work?

Where the time actually goes

Writing it up is easy to skip because the cost is visible and immediate — an hour, right now, while the schedule slips. Starting the handoff is immediate too. The write-up looks like a separate job you can do any time.

EXHIBIT 1

Ask who has ever seen a delay attributed to a meeting that was never written up. The record has no line for it.

Then a product team gets the scope change and comes back asking why. Operations designs the standard path and then finds out an exception was already agreed. A new manager reopens a settled decision because nobody wrote down what would have to change for it to be revisited.

Each of those gets logged as something else — a clarification, a redesign, a reopened decision. The original choice to skip the write-up never appears in the account of the delay. That is why it never gets fixed: nothing in the record points back to it.

Which means the place to fix it is before the work leaves the room.

Let the next team decide when it is done

The best judge is the person receiving the work, because they are the one who has to act without having been in the room. A missing reason, a missing limit, a missing name — all of it shows up the moment someone tries to use the record.

Records matter most at the seams — a decision moving from a steering group to a delivery team, a design moving between functions, a program moving into operations, one owner handing off to the next. This is where knowledge has to move from the people who have it to people who do not, and survive in a form those people can use.

A record that passes does four jobs. It lets the next team act. It keeps the decision and the reasoning together. It stays useful after people leave. And it gives a person, or an AI system, something reliable to look up later. All four depend on the record working without its author standing next to it.

"Complete" is the wrong test. A long record does not help if the decision is buried, the exception is missing, or nobody can tell you whether it is still true. What matters is whether the next team can find the answer, understand where it stops applying, and know who can change it.

The same gap shows up when an AI system does the looking up. Research on data quality in retrieval systems finds that problems starting at the source carry through to the answer.[1] NIST's guidance for generative AI lists where data came from and how good it is among the things organizations should write down, and suggests checking that retrieval data is actually grounded.[2]

A colleague who hits an incomplete page knows who to call. A retrieval system sees the page, whatever is attached to it, and nothing else. A better model cannot supply context that was never written down — though bad documentation is only one of the things that can make an AI answer wrong. So if people or systems are going to rely on a record, it has to say who owns it, how current it is, and where it stops being reliable.

Give the next team five things

At the next handoff that matters, do not hand it over until the person receiving it can act from a short record with a name on it, without reconstructing your meeting.

EXHIBIT 2

Five things the next team needs

- 1 **What we decided**
What did we decide, and what is it meant to achieve?
- 2 **Why we picked it**
What evidence or constraint made the case?
- 3 **What we bent**
What did we agree to bend, and what cannot change?
- 4 **Who does what next**
What happens next, and who owns it?
- 5 **When to revisit**
What date or event should bring this back up?

Apply it wherever missing context would land in somebody else's work. A second question kills documents nobody reads. Ask what depends on this one: a decision, a handoff, a team that comes later, or a system someone approved to rely on. If the answer is nothing, it probably is not worth writing.

The person who made the decision stays responsible for the content even if someone else writes the words. They confirm the reasoning, the limits, and when to revisit. The person receiving it confirms they can actually use it.

Then keep it true

A record can pass the handoff and turn dangerous six months later if nobody maintains it. ISO 30401 treats knowledge management as a system you set up, run, maintain, review, and improve.[3] NASA's practice focuses on continuity — capturing, storing, reusing, and sharing what people know.[4] Research on organizational memory follows the same cycle: how knowledge gets recorded, used, and used again.[5]

None of them tells you how much to write, because that follows from what the record is for. Too much buries the answer. Out-of-date guidance produces confident mistakes. Neither gets caught unless someone is responsible for the record.

That owner keeps the record accurate for the use it claims to support. When a review date passes, a policy or system changes, an exception comes up, or a new decision lands, they fix it, retire it, or flag the part where someone now has to decide for themselves.

The meeting ending is not the handoff

Back to where we started: the schedule is slipping, the meeting is over, the next team is waiting. You decide whether to hand it over.

If the record passes the next team's test, they can start without guessing, and someone is on the hook for keeping it accurate. If it fails, you know exactly what to do — fill in what is missing before you hand it over.

The schedule will still be urgent. What changes is when you call it done: not when the meeting ends, but when the person receiving the work can use what you wrote.

Notes and sources

[1] Leopold Müller, Joshua Holstein, Sarah Bause, Gerhard Satzger, and Niklas Kühn, "Data Quality Challenges in Retrieval-Augmented Generation," arXiv:2510.00552, submitted October 1, 2025; accepted for ICIS 2025. <https://arxiv.org/abs/2510.00552>

[2] National Institute of Standards and Technology, "Artificial Intelligence Risk Management Framework: Generative Artificial Intelligence Profile," NIST AI 600-1, July 2024. <https://doi.org/10.6028/NIST.AI.600-1>

[3] International Organization for Standardization, "ISO 30401:2018 — Knowledge management systems — Requirements," November 2018. <https://www.iso.org/standard/68683.html>

[4] National Aeronautics and Space Administration, "Knowledge Management," APPEL Knowledge Services, updated June 12, 2026. <https://www.nasa.gov/learning-resources/for-professionals/appel-knowledge-management/>

[5] Maria Olívia Ferreira Pereira, Helena de Fátima Nunes Silva, and José Simão de Paula Pinto, "A Memória organizacional nos processos de gestão do conhecimento: um estudo na Universidade Federal do Paraná," *Informação & Informação* 21, no. 1 (2016): 348–374. <https://ojs.uel.br/revistas/uel/index.php/informacao/article/view/18253>

About the author

Marco Policani is an enterprise portfolio, PMO, and AI operating-governance leader. He builds portfolio governance systems where AI carries the analytical load and named humans own every decision — an approach documented across case studies, walkthroughs, and working governance guides on his portfolio site.

Portfolio & governance library: policani.net/governance · Full portfolio: policani.net · LinkedIn: linkedin.com/in/marcpolicani

This white paper may be shared freely with attribution. © 2026 Marco Policani.