


Documentation Is Cross-Functional Leverage, Not Clerical Output

What a slipping schedule hides when the decision record is deferred until after a cross-functional handoff



Marco Policani

Enterprise Portfolio · PMO · AI Operating Governance

policani.net · linkedin.com/in/marcpolicani

July 2026

EXECUTIVE SUMMARY

Require an owned record at consequential handoffs and test whether the receiving owner can act without reconstructing the meeting.

A decision record deferred at a cross-functional handoff moves reconstruction and risk into the receiving function, where the cost returns as delay, rework, lost continuity, or an unreliable source for AI retrieval.

- 01** The cost appears in another queue: an hour saved at the handoff can become questions, reopened decisions, and mistaken assumptions later.
- 02** The receiver supplies the practical test: accept the handoff when the owner can act from the record without reconstructing the meeting.
- 03** A relied-on source needs an owner and refresh trigger; AI retrieval inherits missing context and stale guidance from its source.

The tempting schedule trade

Your steering group has approved a change. The next team is ready to act. The schedule is slipping, so someone proposes a practical trade: start the handoff now and write up the decision later.

If you lead the portfolio or program, that choice can look sensible. The meeting aligned the people in the room. Delivery work is urgent. The record looks like work that can wait. Yet the receiving team does not share the room's context. It needs to know what changed, why it changed, which exception was accepted, and who owns the next decision.

When that context is missing, the receiving owner has to choose between pausing the handoff and proceeding on an assumption. Either choice puts a slipping schedule under more pressure. A handoff transfers more than work; it must transfer what coordination research calls common ground — the shared context that lets the receiving team act without guessing. The decision record is how that common ground travels, which is why deferring it does not save time: the handoff proceeds on an assumption or stalls, and the schedule “saving” returns as rework.

Does starting sooner save time once the decision has to cross into another function? The answer depends on what happens to the context left in the room.

Follow the handoff into delivery

Documentation is easy to defer because a program leader can see the hour needed to record a decision at the same moment the schedule is slipping. Starting the handoff is immediate. Recording the context looks like a separate task.

EXHIBIT 1

A deferred record disappears from the delay account when downstream reconstruction is logged as delivery work.

A product team receives a scope change and returns with questions about its rationale. Operations discovers an accepted exception after designing the standard path. A new owner reopens a settled choice because the condition that would change it is unavailable.

The downstream work is classified as a clarification, redesign, or reopened decision; the original choice to omit context disappears from the account of the delay.

That is where the program leader needs a control: before shared context ends.

The handoff is the point of control

Records become especially consequential at seams: a decision moving from a steering group to a delivery team, a design moving from one function to another, a program moving into operations, or responsibility moving to a new owner. These are the moments when knowledge leaves the people who created it and has to survive in a form someone else can use.

A useful seam record coordinates the receiving work and preserves the decision with its rationale. It also keeps critical knowledge available after people move on and gives people or AI a source to retrieve later. All four functions depend on use without the original author standing beside the record.

Completeness is a weak standard for the handoff. Length does not help when the decision is hidden, the exception is omitted, or no owner can confirm the current answer. Usability requires the receiving team to find that answer, understand its boundary, and know who can change it.

The recipient supplies the best test because that person has to work beyond the original room's shared context. A missing rationale, boundary, or source of authority becomes visible as soon as the receiver tries to use the record.

The retrieval function exposes the same source constraint. Research on data quality in retrieval-augmented generation reports that quality problems can originate at source and ingestion stages and

propagate through retrieval and generation.[1] NIST's Generative AI Profile treats data provenance and data quality as lifecycle concerns.[2]

A colleague may know whom to call when a page is incomplete. A retrieval system sees the page, its metadata, and the access it has been given. Model quality cannot repair missing source context, although documentation defects are only one possible cause of an AI error. Ownership, currency, and source boundaries therefore have to be explicit whenever people or AI are expected to rely on the record.

Require one usable handoff record

The operating move is to put a use test on the next consequential cross-functional handoff. Do not accept the transfer until the receiving owner can act from a short, owned record without reconstructing the meeting.

EXHIBIT 2

A usable handoff record answers five questions

- ① **Decision and outcome**
What was decided, and what outcome does it support?
- ② **Reason and constraint**
Why was the choice made, including material evidence or constraint?
- ③ **Exception and boundary**
What uncertainty or boundary must the receiver preserve?
- ④ **Action and owner**
What action is required now, and who owns it?
- ⑤ **Review trigger**
What event or date should trigger review or refresh?

The program leader applies this requirement where missing context would enter someone else's work. The same test supports stopping artifacts that no recipient uses. If no decision, handoff, continuity need, or approved source depends on an artifact, its value is doubtful.

The knowledge holder remains accountable for the content even when someone else prepares the text. The person who understands the decision has to confirm the rationale, boundary, and refresh trigger, while the receiving owner confirms that the result is usable.

Keep the record fit for use

A record can pass the handoff and later become dangerous if no one maintains it. ISO 30401 frames knowledge management as a system that is established, implemented, maintained, reviewed, and improved.[3] NASA's knowledge-management practice emphasizes continuity through capture, storage, reuse, and sharing.[4] Organizational-memory research similarly examines how knowledge is identified, recorded, used, and reused.[5]

These sources support lifecycle discipline while leaving the needed quantity to the use case. Excess volume can hinder retrieval, and stale guidance can produce confident mistakes. Ownership becomes the safeguard against preserving yesterday's answer after the operating condition has changed.

A named owner keeps the relied-on source current for its declared use. When a review date, policy or system change, material exception, or new decision makes the record inaccurate, that owner corrects it, retires it, or marks a boundary that requires direct judgment.

Acceptance completes the handoff

Return to the opening choice: the schedule is slipping, the meeting has ended, and the next team is waiting. The program leader decides whether the handoff is ready for acceptance.

Passing the receiver test establishes that the next team can begin without guessing and that a named owner can keep the source current. A failed test leaves a visible next step: supply the missing context before accepting the handoff.

The schedule may still demand urgency. With this acceptance rule, the decision crosses the boundary when the receiving owner can use its record. The meeting's end no longer stands in for a completed handoff.

Notes and sources

[1] Leopold Müller, Joshua Holstein, Sarah Bause, Gerhard Satzger, and Niklas Kühl, "Data Quality Challenges in Retrieval-Augmented Generation," arXiv:2510.00552, submitted October 1, 2025; accepted for ICIS 2025. <https://arxiv.org/abs/2510.00552>

[2] National Institute of Standards and Technology, "Artificial Intelligence Risk Management Framework: Generative Artificial Intelligence Profile," NIST AI 600-1, July 2024. <https://doi.org/10.6028/NIST.AI.600-1>

[3] International Organization for Standardization, "ISO 30401:2018 — Knowledge management systems — Requirements," November 2018. <https://www.iso.org/standard/68683.html>

[4] National Aeronautics and Space Administration, "Knowledge Management," APPEL Knowledge Services, updated June 12, 2026. <https://www.nasa.gov/learning-resources/for-professionals/appel-knowledge-management/>

[5] Maria Olívia Ferreira Pereira, Helena de Fátima Nunes Silva, and José Simão de Paula Pinto, "A Memória organizacional nos processos de gestão do conhecimento: um estudo na Universidade Federal do Paraná," *Informação & Informação* 21, no. 1 (2016): 348–374. <https://ojs.uel.br/revistas/uel/index.php/informacao/article/view/18253>

About the author

Marco Policani is an enterprise portfolio, PMO, and AI operating-governance leader. He builds portfolio governance systems where AI carries the analytical load and named humans own every decision — an approach documented across case studies, walkthroughs, and working governance guides on his portfolio site.

Portfolio & governance library: policani.net/governance · Full portfolio: policani.net · LinkedIn: linkedin.com/in/marcpolicani

This white paper may be shared freely with attribution. © 2026 Marco Policani.