13

WHITE PAPER · AI OPERATING GOVERNANCE

AI Autonomy Should Be Earned Through Control Evidence

Action authority expands only when evidence earns it

Marco Policani

Enterprise Portfolio · PMO · AI Operating Governance

policani.net · linkedin.com/in/marcpolicani

July 2026

EXECUTIVE SUMMARY

Autonomy grows when control evidence supports it; it is not a default permission tier.

Autonomy is permission to affect work, systems, money, customers, or records—not a property a model possesses. Authority should expand through explicit levels only when boundary adherence, exception handling, recovery, and accountable intervention have produced sufficient evidence. This paper gives AI program leaders, workflow owners, risk executives, and technology leaders a practical the five-level authority ladder for making the next commitment explicit.

01 Level 1 — Recommend: The system proposes an answer or action; a person evaluates it before any downstream commitment.

02 Level 2 — Prepare: The system may assemble a transaction, communication, configuration, or record but cannot release it.

03 Level 3 — Act inside a reversible boundary: The system may execute low-consequence actions that can be detected and undone within a known period.

04 Level 4 — Act with approval: The system may orchestrate consequential work after a qualified owner approves the case and conditions.

CONTENTS

The argument

Autonomy is permission to affect work, systems, money, customers, or records—not a property a model possesses. Authority should expand through explicit levels only when boundary adherence, exception handling, recovery, and accountable intervention have produced sufficient evidence.

A model can be technically capable of taking an action long before an organization is ready to permit it. Capability demonstrations answer whether the system can complete a path. Autonomy decisions answer whether it may do so under ordinary variation, with a defined consequence, an accountable owner, and a recovery mechanism.

The NIST AI Risk Management Framework and its generative-AI profile provide a useful risk-management foundation [1][2]. The authority ladder in this paper is an operating synthesis for turning that foundation into progressive reliance.

This paper is written for AI program leaders, workflow owners, risk executives, and technology leaders. It offers the five-level authority ladder as a management instrument: a way to connect operating evidence to the next consequential choice. The cited sources provide external context; the framework and its application are Marco Policani's synthesis. It is not a predictive score and does not replace professional, legal, security, financial, or domain judgment.

Why the conventional approach breaks

The framework in this paper is designed to make that boundary usable in ordinary management cadence. Conventional governance tends to separate planning, approval, delivery reporting, and operating ownership. That separation allows a premise to pass through several forums without any one forum owning the decision it implies. The project or workflow then carries an assumption that was acknowledged socially but never accepted operationally.

The happy path can survive this weakness. Ordinary variation exposes it: a disputed input, absent owner, competing commitment, or case that falls outside the assumed rule. The answer is not heavier control at every step. It is to place a clear, proportionate test at the point where the organization increases money, capacity, access, dependency, or action authority.

The five-level authority ladder

The model has 6 connected elements. Each makes a different condition visible, but they should be reviewed together because strength in one area does not cancel a missing obligation in another. The model is designed for a short executive conversation supported by inspectable evidence, not a long maturity assessment.

EXHIBIT 1

The five-level authority ladder connects evidence to action

1. **Level 1 — Recommend**
   The system proposes an answer or action; a person evaluates it before any downstream commitment.
2. **Level 2 — Prepare**
   The system may assemble a transaction, communication, configuration, or record but cannot release it.
3. **Level 3 — Act inside a reversible boundary**
   The system may execute low-consequence actions that can be detected and undone within a known period.

4 **Level 4 — Act with approval**
The system may orchestrate consequential work after a qualified owner approves the case and conditions.

5 **Level 5 — Act under preauthorized policy**
The system may execute within a narrow policy envelope without per-case approval, while monitoring and stop authority remain active.

6 **De-escalation is part of autonomy**
Authority must contract when evidence weakens, context changes, or oversight capacity falls.

Level 1 — Recommend

The system proposes an answer or action; a person evaluates it before any downstream commitment. This is where a management ritual either becomes a real control or reveals itself as administrative theater.

Recommendations look advisory but are copied directly into decisions because review is rushed or the interface makes acceptance easier than challenge. Because responsibility is distributed, no single event appears large enough to stop the work. The accumulated operating load is the real exposure.

Separate suggestion from execution, show sources and uncertainty, and record consequential acceptance or rejection. Evidence should be selected for decision value. Volume is not rigor when the decisive uncertainty remains untested.

Recommendation quality, reviewer disposition, correction patterns, decision impact, and sampled independent checks. A later outcome should feed back into the evidence standard so repeated surprises improve the next commitment.

Remain at recommendation when reviewers cannot reliably detect material errors. A decision without consequence invites the old behavior to continue. Capacity, funding, permissions, scope, or sequence must move with the answer.

A tabletop review of level 1 — recommend should use the observed break as its scenario: Recommendations look advisory but are copied directly into decisions because review is rushed or the interface makes acceptance easier than challenge. Participants then apply the proposed response. They separate suggestion from execution, show sources and uncertainty, and record consequential acceptance or rejection. The review identifies where authority, information, time, or recovery would run out. It is useful only if it changes the boundary or confirms who will act when the condition appears.

The minimum record for level 1 — recommend is anchored in actual proof. It includes recommendation quality, reviewer disposition, correction patterns, decision impact, and sampled independent checks. The record also carries the choice supported by that proof: Remain at recommendation when reviewers cannot reliably detect material errors. Keeping evidence and decision together lets a later owner distinguish a deliberate residual risk from a premise that nobody examined. It also gives the team a clear reason to reopen the choice when the underlying facts change.

Across the work governed by AI program leaders, workflow owners, risk executives, and technology leaders, occurrences of this pattern should be compared rather than treated as unrelated project stories. The positive standard is: The system proposes an answer or action; a person evaluates it before any downstream commitment.

If the same failure recurs, leaders can change delegation, capacity, intake, policy, or the intervention itself. That portfolio response is usually more valuable than asking each team to absorb another local workaround.

Two questions test level 1 — recommend. What evidence would show that the condition is no longer adequate? Would that evidence arrive early enough to preserve a feasible choice? The intended choice is clear: Remain at recommendation when reviewers cannot reliably detect material errors. A weak answer to either question calls for a smaller commitment, an earlier signal, or a safer fallback. A strong answer earns movement without claiming that uncertainty has disappeared.

Level 2 — Prepare

The system may assemble a transaction, communication, configuration, or record but cannot release it. The work becomes governable when ownership, evidence, boundaries, and response are connected to a real choice.

Prepared work silently crosses into action through auto-send, bulk approval, or ambiguous interface behavior. When the consequence finally becomes visible, leaders face a narrower choice set and a more expensive recovery.

Require an affirmative release by an authorized person and display the fields, systems, and parties affected. The control should be short enough to use at the decision point and specific enough that a skeptical owner can challenge it.

Approval identity, edits before release, rejected preparations, policy checks, and prevented boundary violations. Ranges and contradictory observations are often more useful than a precise point estimate built on weak assumptions.

Expand only when preparation is accurate and review capacity remains credible at expected volume. The decision should identify the sacrifice. If every priority remains protected, the tradeoff has been delegated rather than resolved.

A tabletop review of level 2 — prepare should use the observed break as its scenario: Prepared work silently crosses into action through auto-send, bulk approval, or ambiguous interface behavior. Participants then apply the proposed response. They require an affirmative release by an authorized person and display the fields, systems, and parties affected. The review identifies where authority, information, time, or recovery would run out. It is useful only if it changes the boundary or confirms who will act when the condition appears.

The minimum record for level 2 — prepare is anchored in actual proof. It includes approval identity, edits before release, rejected preparations, policy checks, and prevented boundary violations. The record also carries the choice supported by that proof: Expand only when preparation is accurate and review capacity remains credible at expected volume. Keeping evidence and decision together lets a later owner distinguish a deliberate residual risk from a premise that nobody examined. It also gives the team a clear reason to reopen the choice when the underlying facts change.

Across the work governed by AI program leaders, workflow owners, risk executives, and technology leaders, occurrences of this pattern should be compared rather than treated as unrelated project stories. The positive standard is: The system may assemble a transaction, communication, configuration, or record but cannot release it. If the same failure recurs, leaders can change delegation, capacity, intake, policy, or the intervention itself. That portfolio response is usually more valuable than asking each team to absorb another local workaround.

Two questions test level 2 — prepare. What evidence would show that the condition is no longer adequate? Would that evidence arrive early enough to preserve a feasible choice? The intended choice is clear: Expand only when preparation is accurate and review capacity remains credible at expected volume. A weak answer to either question calls for a smaller commitment, an earlier signal, or a safer fallback. A strong answer earns movement without claiming that uncertainty has disappeared.

Level 3 — Act inside a reversible boundary

The system may execute low-consequence actions that can be detected and undone within a known period. The framework in this paper is designed to make that boundary usable in ordinary management cadence.

Reversibility is assumed because a technical rollback exists, while customer, legal, timing, or data consequences cannot actually be reversed. The problem may remain invisible for several review cycles because effort masks the missing condition. By the time variance appears, inexpensive options have already expired.

Define transaction limits, eligible cases, observation window, rollback owner, and maximum exposure. Proportionality is essential. Reversible, low-consequence work can proceed with a lighter record than a commitment that is costly to undo.

Boundary adherence, intervention frequency, rollback tests, loss or harm avoided, and time to restore known state. Sampling should include the awkward cases. Averages built from the easy path are least informative where governance is most needed.

Reduce authority immediately when practical reversibility differs from the design assumption. The review should end with a changed commitment or a reasoned decision to keep it—not a request for more color commentary.

A tabletop review of level 3 — act inside a reversible boundary should use the observed break as its scenario: Reversibility is assumed because a technical rollback exists, while customer, legal, timing, or data consequences cannot actually be reversed. Participants then apply the proposed response. They define transaction limits, eligible cases, observation window, rollback owner, and maximum exposure. The review identifies where authority, information, time, or recovery would run out. It is useful only if it changes the boundary or confirms who will act when the condition appears.

The minimum record for level 3 — act inside a reversible boundary is anchored in actual proof. It includes boundary adherence, intervention frequency, rollback tests, loss or harm avoided, and time to restore known state. The record also carries the choice supported by that proof: Reduce authority immediately when practical reversibility differs from the design assumption. Keeping evidence and decision together lets a later owner distinguish a deliberate residual risk from a premise that nobody examined. It also gives the team a clear reason to reopen the choice when the underlying facts change.

Across the work governed by AI program leaders, workflow owners, risk executives, and technology leaders, occurrences of this pattern should be compared rather than treated as unrelated project stories. The positive standard is: The system may execute low-consequence actions that can be detected and undone within a known period. If the same failure recurs, leaders can change delegation, capacity, intake, policy, or the intervention itself. That portfolio response is usually more valuable than asking each team to absorb another local workaround.

Two questions test level 3 — act inside a reversible boundary. What evidence would show that the condition is no longer adequate? Would that evidence arrive early enough to preserve a feasible choice? The intended choice is clear: Reduce authority immediately when practical reversibility differs from the design assumption. A weak answer to either question calls for a smaller commitment, an earlier signal, or a safer fallback. A strong answer earns movement without claiming that uncertainty has disappeared.

Level 4 — Act with approval

The system may orchestrate consequential work after a qualified owner approves the case and conditions. Viewed from the portfolio, this is a resource-allocation problem before it becomes a delivery problem.

Approval becomes ceremonial under volume, or the approver cannot see the evidence needed to judge the proposed action. Reporting usually captures the symptom after it has been translated into schedule, cost, quality, or risk. It rarely captures the original unresolved choice.

Design the approval packet, sampling, separation of duties, escalation, and service capacity as part of the system. Good governance preserves options. It creates a legitimate route to narrow, stage, redirect, pause, or stop without labeling learning as failure.

Approval latency, override reasons, reviewer workload, exception severity, and actions stopped before execution. A lack of evidence is itself useful information if the next decision allows a smaller experiment, tighter boundary, or deliberate wait.

Hold volume below the point where oversight quality deteriorates. The response should occur while it can still change the outcome, not after the issue has been translated into an unrecoverable variance.

A tabletop review of level 4 — act with approval should use the observed break as its scenario: Approval becomes ceremonial under volume, or the approver cannot see the evidence needed to judge the proposed action. Participants then apply the proposed response. They design the approval packet, sampling, separation of duties, escalation, and service capacity as part of the system. The review identifies where authority, information, time, or recovery would run out. It is useful only if it changes the boundary or confirms who will act when the condition appears.

The minimum record for level 4 — act with approval is anchored in actual proof. It includes approval latency, override reasons, reviewer workload, exception severity, and actions stopped before execution. The record also carries the choice supported by that proof: Hold volume below the point where oversight quality deteriorates. Keeping evidence and decision together lets a later owner distinguish a deliberate residual risk from a premise that nobody examined. It also gives the team a clear reason to reopen the choice when the underlying facts change.

Across the work governed by AI program leaders, workflow owners, risk executives, and technology leaders, occurrences of this pattern should be compared rather than treated as unrelated project stories. The positive standard is: The system may orchestrate consequential work after a qualified owner approves the case and conditions. If the same failure recurs, leaders can change delegation, capacity, intake, policy, or the intervention itself. That portfolio response is usually more valuable than asking each team to absorb another local workaround.

Two questions test level 4 — act with approval. What evidence would show that the condition is no longer adequate? Would that evidence arrive early enough to preserve a feasible choice? The intended choice is clear: Hold volume below the point where oversight quality deteriorates. A weak answer to either question calls for a smaller commitment, an earlier signal, or a safer fallback. A strong answer earns movement without claiming that uncertainty has disappeared.

Level 5 — Act under preauthorized policy

The system may execute within a narrow policy envelope without per-case approval, while monitoring and stop authority remain active. A useful governance design therefore begins with the consequence leaders are prepared to accept.

Broad goals substitute for enforceable boundaries and rare cases create consequences outside the training or evaluation set. A local fix can keep the work moving while worsening the portfolio: another specialist is borrowed, another exception is tolerated, and another dependency is promised twice.

Encode eligible actions, limits, exclusions, monitoring, incident thresholds, and automatic safe states. The control is complete only when the response is known: who acts, how quickly, within what authority, and with which safe fallback.

Production evaluation, policy-conformance logs, exception detection, recovery exercises, drift signals, and owner response. The receiving operation must recognize the evidence as relevant. Technical success alone cannot prove that a business workflow can carry the change.

Grant this level by workflow and action class, never as a blanket status for the model. That choice should be recorded in operational terms: what may proceed, what remains outside the boundary, who owns the next move, and what would reopen the decision.

A tabletop review of level 5 — act under preauthorized policy should use the observed break as its scenario: Broad goals substitute for enforceable boundaries and rare cases create consequences outside the training or evaluation set. Participants then apply the proposed response. They encode eligible actions, limits, exclusions, monitoring, incident thresholds, and automatic safe states. The review identifies where authority, information, time, or recovery would run out. It is useful only if it changes the boundary or confirms who will act when the condition appears.

The minimum record for level 5 — act under preauthorized policy is anchored in actual proof. It includes production evaluation, policy-conformance logs, exception detection, recovery exercises, drift signals, and owner response. The record also carries the choice supported by that proof: Grant this level by workflow and action class, never as a blanket status for the model. Keeping evidence and decision together lets a later owner distinguish a deliberate residual risk from a premise that nobody examined. It also gives the team a clear reason to reopen the choice when the underlying facts change.

Across the work governed by AI program leaders, workflow owners, risk executives, and technology leaders, occurrences of this pattern should be compared rather than treated as unrelated project stories. The positive standard is: The system may execute within a narrow policy envelope without per-case approval, while monitoring and stop authority remain active. If the same failure recurs, leaders can change delegation, capacity, intake, policy, or the intervention itself. That portfolio response is usually more valuable than asking each team to absorb another local workaround.

Two questions test level 5 — act under preauthorized policy. What evidence would show that the condition is no longer adequate? Would that evidence arrive early enough to preserve a feasible choice? The intended choice is clear: Grant this level by workflow and action class, never as a blanket status for the model. A weak answer to either question calls for a smaller commitment, an earlier signal, or a safer fallback. A strong answer earns movement without claiming that uncertainty has disappeared.

De-escalation is part of autonomy

Authority must contract when evidence weakens, context changes, or oversight capacity falls. A strong operating model makes the hidden negotiation explicit while options are still available.

Autonomy rises through pilots but has no ordinary route downward short of a major incident. A polished status can coexist with this failure because status describes activity; it does not prove the premise that authorizes reliance.

Predefine downgrade triggers tied to model changes, data shifts, exception rates, control failures, staffing, and new consequences. The aim is not certainty. It is a justified next commitment with an observable basis and a viable way back.

Trigger history, time to reduce permissions, fallback performance, unresolved incidents, and evidence required to restore authority. Measures should expose changes in the operation rather than reward production of artifacts. Activity can support the story; it cannot substitute for the result.

Step down first and investigate from a safe boundary rather than preserving status while evidence is uncertain. When the premise is weak, narrowing is often a better decision than forcing a binary choice between full approval and cancellation.

A tabletop review of de-escalation is part of autonomy should use the observed break as its scenario: Autonomy rises through pilots but has no ordinary route downward short of a major incident. Participants then apply the proposed response. They predefine downgrade triggers tied to model changes, data shifts, exception rates, control failures, staffing, and new consequences. The review identifies where authority, information, time, or recovery would run out. It is useful only if it changes the boundary or confirms who will act when the condition appears.

The minimum record for de-escalation is part of autonomy is anchored in actual proof. It includes trigger history, time to reduce permissions, fallback performance, unresolved incidents, and evidence required to restore authority. The record also carries the choice supported by that proof: Step down first and investigate from a safe boundary rather than preserving status while evidence is uncertain. Keeping evidence and decision together lets a later owner distinguish a deliberate residual risk from a premise that nobody examined. It also gives the team a clear reason to reopen the choice when the underlying facts change.

Across the work governed by AI program leaders, workflow owners, risk executives, and technology leaders, occurrences of this pattern should be compared rather than treated as unrelated project stories. The positive standard is: Authority must contract when evidence weakens, context changes, or oversight capacity falls. If the same failure recurs, leaders can change delegation, capacity, intake, policy, or the intervention itself. That portfolio response is usually more valuable than asking each team to absorb another local workaround.

Two questions test de-escalation is part of autonomy. What evidence would show that the condition is no longer adequate? Would that evidence arrive early enough to preserve a feasible choice? The intended choice is clear: Step down first and investigate from a safe boundary rather than preserving status while evidence is uncertain. A weak answer to either question calls for a smaller commitment, an earlier signal, or a safer fallback. A strong answer earns movement without claiming that uncertainty has disappeared.

How the elements interact

Level 1 — Recommend and Level 2 — Prepare protect different parts of the same commitment. The system proposes an answer or action; a person evaluates it before any downstream commitment. The system may assemble a transaction, communication, configuration, or record but cannot release it. The practical failure occurs between the formal approval and the ordinary pressure of running the work. If the operation encounters this failure—recommendations look advisory but are copied directly into decisions because review is rushed or the interface makes acceptance easier than challenge.—the response for level 2 — prepare cannot compensate for the missing level 1 — recommend obligation. Reviewers should reconcile the evidence for both elements before they accept the next action: expand only when preparation is accurate and review capacity remains credible at expected volume.

Level 2 — Prepare and Level 3 — Act inside a reversible boundary protect different parts of the same commitment. The system may assemble a transaction, communication, configuration, or record but cannot release it. The system may execute low-consequence actions that can be detected and undone within a known period. Viewed from the portfolio, this is a resource-allocation problem before it becomes a delivery problem. If the operation encounters this failure—prepared work silently crosses into action through auto-send, bulk approval, or ambiguous interface behavior.—the response for level 3 — act inside a reversible boundary cannot compensate for the missing level 2 — prepare obligation. Reviewers should reconcile the evidence for both elements before they accept the next action: reduce authority immediately when practical reversibility differs from the design assumption.

Level 3 — Act inside a reversible boundary and Level 4 — Act with approval protect different parts of the same commitment. The system may execute low-consequence actions that can be detected and undone within a known period. The system may orchestrate consequential work after a qualified owner approves the case and conditions. The control belongs at the point where reliance expands, not in a retrospective explanation after the outcome is known. If the operation encounters this failure—reversibility is assumed because a technical rollback exists, while customer, legal, timing, or data consequences cannot actually be reversed.—the response for level 4 — act with approval cannot compensate for the missing level 3 — act inside a reversible boundary obligation. Reviewers should reconcile the evidence for both elements before they accept the next action: hold volume below the point where oversight quality deteriorates.

Level 4 — Act with approval and Level 5 — Act under preauthorized policy protect different parts of the same commitment. The system may orchestrate consequential work after a qualified owner approves the case and conditions. The system may execute within a narrow policy envelope without per-case approval, while monitoring and stop authority remain active. The visible artifact is often complete before the operating condition behind it is credible. If the operation encounters this failure—approval becomes ceremonial under volume, or the approver cannot see the evidence needed to judge the proposed action.—the response for level 5 — act under preauthorized policy cannot compensate for the missing level 4 — act with approval obligation. Reviewers should reconcile the evidence for both elements before they accept the next action: grant this level by workflow and action class, never as a blanket status for the model.

Level 5 — Act under preauthorized policy and De-escalation is part of autonomy protect different parts of the same commitment. The system may execute within a narrow policy envelope without per-case approval, while monitoring and stop authority remain active. Authority must contract when evidence weakens, context changes, or oversight capacity falls. A useful governance design therefore begins with the consequence leaders are prepared to accept. If the operation encounters this failure—broad goals substitute for enforceable boundaries and rare cases create consequences outside the training or evaluation set.—the response for de-escalation is part of autonomy cannot compensate for the missing level 5 — act under preauthorized policy obligation. Reviewers should reconcile the evidence for both elements before they accept the next action: step down first and investigate from a safe boundary rather than preserving status while evidence is uncertain.

De-escalation is part of autonomy and Level 1 — Recommend protect different parts of the same commitment. Authority must contract when evidence weakens, context changes, or oversight capacity falls. The system proposes an answer or action; a person evaluates it before any downstream commitment. This is where a management ritual either becomes a real control or reveals itself as administrative theater. If the operation encounters this failure—autonomy rises through pilots but has no ordinary route downward short of a major incident.—the response for level 1 — recommend cannot compensate for the missing de-escalation is part of autonomy obligation. Reviewers should reconcile the evidence for both elements before they accept the next action: remain at recommendation when reviewers cannot reliably detect material errors.

Put the model into the management cadence

A framework becomes useful when it changes the normal cadence of work. It should shape intake, funding, design, delivery, and operating review without creating a separate governance industry around itself. The following sequence is deliberately small:

EXHIBIT 2

Start with one decision, then calibrate from evidence

1. **Move 1**
   Inventory actions by consequence and reversibility before assigning authority levels.
2. **Move 2**
   Set evidence gates for one workflow and one action class at a time.
3. **Move 3**
   Test intervention, rollback, and permission reduction under realistic load.
4. **Move 4**
   Review authority after changes to model, data, tools, policy, users, or oversight staffing.

- Inventory actions by consequence and reversibility before assigning authority levels.
- Set evidence gates for one workflow and one action class at a time.
- Test intervention, rollback, and permission reduction under realistic load.
- Review authority after changes to model, data, tools, policy, users, or oversight staffing.

The first cycle should be treated as calibration. Reviewers should note which questions changed a decision, which evidence was unavailable, where authority was unclear, and which controls cost out of proportion to the exposure. That record improves the five-level authority ladder without pretending the first version will fit every workflow or investment.

Ownership should remain close to the consequence. The PMO or AI-governance function can maintain the method, surface cross-portfolio patterns, and challenge weak evidence. It should not absorb the business owner's accountability for the result or the executive's accountability for the commitment.

Measures that reveal whether the control works

- Time from a material signal to a consequential decision.
- Commitments narrowed, redirected, paused, or stopped before avoidable exposure grew.
- Exceptions or assumptions discovered after the latest useful decision date.
- Scarce specialist and executive attention consumed by recurring unresolved conditions.
- Decisions reopened because the original evidence, boundary, or rationale was unclear.
- Operating outcomes and recovery performance after reliance expanded.

These measures are diagnostic, not a universal scorecard. Their purpose is to identify where the operating model fails repeatedly and whether governance is preserving options earlier. A lower count can indicate improvement, but it can also indicate weak detection. Leaders should read the measures with case evidence and frontline observation.

The strongest counterargument

The ladder does not claim autonomy is always desirable or that every workflow should climb. Some work should remain at recommendation or preparation indefinitely. The best level is the smallest authority that produces the required value within an operating boundary the organization can sustain.

A reasonable critic could also argue that experienced leaders already do this through judgment. Many do. The weakness is not the absence of judgment; it is the absence of a durable operating record when people, pressure, and conditions change. The five-level authority ladder makes the basis for judgment visible enough to challenge, execute, and revisit.

The shift

Autonomy is permission to affect work, systems, money, customers, or records—not a property a model possesses. Authority should expand through explicit levels only when boundary adherence, exception handling, recovery, and accountable intervention have produced sufficient evidence. The practical shift is from reporting activity to governing the conditions under which the organization relies on that activity.

That discipline does not make uncertainty disappear. It gives leaders a better place to stand: a named decision, evidence proportionate to consequence, an explicit boundary, a responsible owner, and a response when reality departs from the assumption. Those elements are enough to turn hidden operating risk into a choice the portfolio can make deliberately.

Notes and sources

[1] NIST, "AI Risk Management Framework Core," 2023. https://airc.nist.gov/airmf-resources/airmf/5-sec-core/

[2] NIST, "Artificial Intelligence Risk Management Framework: Generative Artificial Intelligence Profile," July 2024. https://nvlpubs.nist.gov/nistpubs/ai/nist.ai.600-1.pdf

About the author

Marco Policani is an enterprise portfolio, PMO, and AI operating-governance leader. He builds portfolio governance systems where AI carries the analytical load and named humans own every decision — an approach documented across case studies, walkthroughs, and working governance guides on his portfolio site.

Portfolio & governance library: policani.net/governance · Full portfolio: policani.net · LinkedIn: linkedin.com/in/marcpolicani