
WHITE PAPER · PORTFOLIO & AI GOVERNANCE

Enterprise AI Trust Is Built Through Operating Evidence, Not Tool Novelty

Why a capable AI pilot stalls at the safe edge of the business, and why a better model will not move it

Marco Policani

Enterprise Portfolio · PMO · AI Operating Governance

policani.net · [linkedin.com/in/marcpolicani](https://www.linkedin.com/in/marcpolicani)

July 2026

EXECUTIVE SUMMARY

Trust scales when a named owner can show what a system did, on what basis, within what limits, and who is accountable.

A capable AI pilot impresses everyone and then stalls before it reaches important work. Asking for a better model misreads the problem. The people accountable for using the system still lack proof that they can rely on it. This paper shows how a current evidence pack (one record of what the system did, on what basis, within what limits, and who answers for it) can provide that proof.

- 01 Usage dashboards and demo decks show that people tried the tool. They do not prove that it worked within known limits under a named owner.
- 02 Keep that record current as the system changes. The next team that wants to use the system can review what changed instead of restarting the trust debate.
- 03 The pack is not a certification or a substitute for judgment. It helps the owner decide whether to rely, retest, or hold back.

Why capable AI pilots stall

A leader runs a successful AI pilot. The demo lands, the trial gets funded, and then the system stalls. It stays on low-stakes tasks at the edge of the business. It never enters the core workflow that would justify the investment. The usual response is to ask for a better model.

That response misreads the problem. A stronger model will not move a stuck pilot when model quality is not the missing piece. The organization is holding the system outside important work because accountable leaders cannot answer one question: can we rely on it? Borrowing from safety-critical engineering, this paper treats the missing proof as an assurance case. It is a structured argument, backed by evidence, that a system is reliable enough for one use in one setting. Trust rests on the case the owner can show, not on the model alone. A better model will not move a pilot with an empty assurance case.

Pressure also comes from outside the program. Boards ask what AI delivers and what it risks. Buyers ask whether a provider governs the system, not only what it can do. BCG surveyed nearly 2,400 executives, including 640 CEOs, for its 2026 AI Radar. It found that 72 percent of CEOs say they are their company's main AI decision maker. Most also plan to keep investing even if AI does not pay off this year.[6] Deloitte frames governance, scale, and value as the issues that separate organizations gaining value from those still testing.[5] A pilot that cannot answer the reliance question stays parked.

The stakes rise as AI systems change. McKinsey says trust supports lasting adoption in core workflows. When systems begin to act, the risk grows from a wrong answer to a wrong action, tool misuse, or work beyond set limits.[1] Higher stakes require stronger proof before the business relies on the system.

What a business checks before it relies on a system

A demo proves that a system can produce a good result on a chosen example. Production asks a different question. Can the people accountable for the work see what the system did and decide, from evidence, whether to rely on it? That question is what moves a system into core work.

EXHIBIT 1

The four questions a named owner has to answer

What did it do?

A record of actions and outputs.

On what basis?

The sources, inputs, and reasoning trail behind them.

Within what limits?

Where the system has been tested and where it has not.

Who is accountable?

The named owner who answers for reliance on the output.

Trust becomes governable when a named owner can prove it on demand. This paper uses a concrete test. An AI system is trustworthy for a given use when its owner can show four things:

The owner can show what the system did through a record of its actions and outputs. The owner can show the basis for those results through sources, inputs, and a reasoning trail. The owner can state where the system has been tested and where it has not. The owner is also named and accountable for relying on the output.

If the owner cannot show all four, the system is not yet trustworthy for that use, even if the demo went well. This makes trust something built from evidence. It also lets a buyer, auditor, or board inspect the case without having attended the pilot.

The evidence pack

Research gives that evidence a useful shape. Before the current AI wave, researchers proposed FactSheets for AI services. They based them on supplier declarations used in other industries. A FactSheet records what a buyer needs to inspect a service instead of taking it on faith: purpose, performance, safety, security, and provenance. [2]

In business use, that becomes an evidence pack that travels with the system. It states the purpose and intended use. It records measured performance and the conditions of the test. It covers safety, security, data and model sources, tested limits, human review, escalation, and the named owner. Each item answers one of the four trust questions.

The pack also changes the cost of scaling. When trust remains unspoken, each new use restarts the same slow debate. With an evidence pack, a system that earned trust in one workflow arrives at the next with its proof attached. The review can focus on what changed. That lets the organization add trusted uses without lowering its standard.

Keeping the evidence current

A pack filed once is only a snapshot. Models drift, inputs change, and use can move beyond what was tested. A system trusted at launch can quietly stop earning that trust. NIST treats AI risk as a full-lifecycle concern, from design through retirement. Its Generative AI Profile suggests monitoring, feedback, red-team tests, independent

review, and incident response. It also asks leaders to examine how people and AI work together, vet suppliers, and trace training data.[3]

That work fits the NIST AI RMF functions: govern, map, measure, and manage. Together they form an ongoing loop. Teams study the setting, measure performance and risk, and decide how the system should be used as conditions change.[4]

The evidence pack therefore has to stay current. The named owner watches for drift and repeats tests as use expands. A stale pack cannot support a current trust claim.

This lifecycle approach does not have to slow the program. A current pack lets a team extend a trusted system into new work without starting from zero. It also helps the team catch a decline before it becomes an incident.

How an evidence pack fails

Evidence has failure modes of its own. Three deserve a name.

EXHIBIT 2

Three ways an evidence pack stops being evidence

Evidence theater

Pages that count activity but do not answer the four trust questions.

False precision

Numbers that imply more testing than the team performed.

Approval-only pack

A pack built to clear an approval but not kept current as the system changes.

The first is evidence theater. Teams produce usage dashboards and demo decks that count activity but do not prove reliable results. Usage shows that people tried the tool. It does not show that the tool did the right thing, for the right reason, within known limits, under a named owner. More pages do not make a stronger case.

The second is false precision. A performance number can imply more testing than the team performed. A number without its test conditions hides the gap it should reveal. So does a limits section that offers comfort but names no clear boundary.

The third is the approval-only pack. The team builds it to clear a gate and never updates it. Its evidence may have been true at launch, but the system keeps changing. Without an owner to maintain it, the pack becomes a history file and can no longer support trust.

Each failure creates the appearance of evidence without doing its job. The test is still simple. Can the named owner show what the system did, on what basis, within what limits, and answer for it today?

One pattern from governance work

A common pattern from governance work shows how this can operate. An AI system prepares analysis for a human decision. It shows its sources and the path behind its result. It states where it has and has not been tested. It sends low-confidence cases to a person instead of guessing. A reviewer can inspect each conclusion, and a named owner decides whether to rely on it.

The system earns its way into deeper work because each new use starts with that evidence. The team does not reopen the whole trust question. When something goes wrong, the record shows what happened and helps the team correct it. That is good governance in daily work.

The limits of an evidence pack

Two limits matter. First, no single enterprise standard governs every AI evidence pack. FactSheets provide a durable model, and NIST provides lifecycle guidance.[2,3] Each organization still has to design its own record. Treat the pack as a working pattern, not a certification.

Evidence reduces the trust problem but does not remove judgment. A complete pack still needs a skilled person to weigh it. A well-documented system can still be used beyond its tested limits. The named owner remains essential. No document can replace the judgment to rely on a system or hold it back.

Start with one high-consequence use

Start with one use, not the whole AI portfolio. Choose the use that would cause the most harm if it were wrong. Build one current evidence pack with a named owner for

that use. Make it easy to inspect. That first pack becomes the model for the next team because it shows what trust required in practice.

Return to the pilot parked at the edge of the business. A better model will not move it. Evidence will: what the system did, the basis for its result, its tested limits, and the owner who answers for its use. That evidence helps a capable tool enter important work and keep the confidence of buyers and boards. Without it, the tool stays at the edge.

Notes and sources

1. Gabriel Morgan Asaftei, Roger Roberts, Abby Sticha, and Cécile Prinsen, "State of AI trust in 2026: Shifting to the agentic era," McKinsey & Company, March 25, 2026. Verified August 3, 2026. <https://www.mckinsey.com/capabilities/tech-and-ai/our-insights/tech-forward/state-of-ai-trust-in-2026-shifting-to-the-agentic-era>

2. Matthew Arnold, Rachel K. E. Bellamy, Michael Hind, Stephanie Houde, Sameep Mehta, Aleksandra Mojsilović, Ravi Nair, Karthikeyan Natesan Ramamurthy, Darrell Reimer, Alexandra Olteanu, David Piorkowski, Jason Tsay, and Kush R. Varshney, "FactSheets: Increasing Trust in AI Services through Supplier's Declarations of Conformity," arXiv:1808.07261, 2018 (rev. 2019). Verified August 3, 2026. <https://arxiv.org/abs/1808.07261>

3. National Institute of Standards and Technology, "Artificial Intelligence Risk Management Framework: Generative Artificial Intelligence Profile," NIST AI 600-1, July 2024. Verified August 3, 2026. <https://doi.org/10.6028/NIST.AI.600-1>

4. National Institute of Standards and Technology, "AI Risk Management Framework Core," excerpt from AI RMF 1.0, 2023. Verified August 3, 2026. <https://airc.nist.gov/airmf-resources/airmf/5-sec-core/>

5. Deloitte, "The State of AI in the Enterprise," 2026 AI report, Deloitte AI Institute. Verified August 3, 2026. <https://www.deloitte.com/us/en/what-we-do/capabilities/applied-artificial-intelligence/content/state-of-ai-in-the-enterprise.html>

6. Jessica Apotheker, Sylvain Duranton, Vladimir Lukic, Nicolas de Bellefonds, and Christoph Schweizer, "As AI Investments Surge, CEOs Take the Lead," BCG AI Radar 2026, January 15, 2026. Verified August 3, 2026. <https://www.bcg.com/publications/2026/as-ai-investments-surge-ceos-take-the-lead>

About the author

Marco Policani is an enterprise portfolio, PMO, and AI operating-governance leader. He builds portfolio governance systems where AI carries the analytical load and named humans own every decision — an approach documented across case studies, walkthroughs, and working governance guides on his portfolio site.

Portfolio & governance library: policani.net/governance · Full portfolio: policani.net · LinkedIn: linkedin.com/in/marcpolicani

This white paper may be shared freely with attribution. © 2026 Marco Policani.