


Enterprise AI Trust Is Built Through Operating Evidence, Not Tool Novelty

Why a capable AI pilot stalls at the safe edge of the business, and why a better model will not move it



Marco Policani

Enterprise Portfolio · PMO · AI Operating Governance

policani.net · linkedin.com/in/marcpolicani

July 2026

EXECUTIVE SUMMARY

Trust scales when a named owner can show what a system did, on what basis, within what limits, and who is accountable.

A capable AI pilot impresses everyone and then stalls short of the work that would justify it. The reflex - ask for a better model - misreads the problem: what the system cannot yet do is prove itself to the people accountable for relying on it. This paper defines that proof as operating evidence, structures it as a pack that travels with the capability and stays current, and shows why buyers and boards ask to inspect it.

- 01** A capability is trustworthy when a named owner can show what it did, on what basis, within what limits, and who is accountable.
- 02** Package the proof as a living evidence pack, so a capability arrives at its next use with proof attached and the review covers only what changed.
- 03** Start with the single highest-consequence AI use and build one pack. The discipline spreads by demand, not mandate.

Why capable AI pilots stall

Picture a leader who has just run a successful AI pilot. The demo landed, the trial got funded, and then the capability stalled. It runs on low-stakes tasks at the edge of the business and never crosses into the core workflow that would justify the investment. The reflex, in the room and in the next budget cycle, is to ask for a better model.

That reflex misreads the problem. A more capable model does not move a pilot that is stuck, because capability is not what is missing. The organization has not let the system into the work that matters because it cannot answer, to the people accountable for the decision, one plain question: can we rely on this. That is a trust question, and no amount of novelty answers it. What is missing has a name in safety-critical engineering: an assurance case — a structured argument, backed by evidence, that a system is trustworthy enough to rely on for a specific use in a specific context. Trust attaches to the assurance case a system's owner can present, not to the model itself, which is why a more capable model does not move a pilot whose assurance case is empty.

The pressure to answer it also comes from outside. Boards ask management to report what AI is delivering and what it risks, in business terms, and enterprise buyers weigh whether a provider can show that its capability is governed, alongside what the capability does. BCG's 2026 AI Radar, a survey of nearly 2,400 executives including 640 CEOs, documents surging AI investment and a shift of AI

decisions to the CEO: 72 percent of CEOs say they are now their company's main decision maker on AI, and most plan to keep investing even if AI does not pay off this year.⁶ Deloitte's enterprise research frames governance, scaling, and value realization as the operating issues that separate the organizations capturing AI value from those still stuck in experimentation.⁵ The pilot that cannot answer the reliance question stays parked while a competitor that can answer it takes the deeper deployment.

The stakes also rise as the systems themselves change. McKinsey's account of the shift to the agentic era is that trust is what supports sustained adoption and integration into core workflows, and that as systems begin to act rather than only answer, the risk expands from saying the wrong thing to taking the wrong action, misusing a tool, or operating beyond its guardrails.¹ Higher stakes raise the evidence a business needs before it relies on a system.

What a business checks before it relies on a system

A demo shows one thing: that the system can produce a good result on a curated example. Reliance in production depends on something the demo never puts on screen, which is whether the people accountable for the work can see what the system did and decide, on evidence, whether to depend on it. Those are different questions, and the second one is the one that gets a capability into production work.

EXHIBIT 1

The four questions a named owner has to answer

What did it do? A record of actions and outputs.	On what basis? The sources, inputs, and reasoning trail behind them.
Within what limits? An honest statement of where the system is tested and where it is not.	Who is accountable? The named owner who answers for reliance on the output.

Trust becomes governable the moment it is defined as something a named human owner can demonstrate on demand. The working definition this paper uses is deliberately concrete. An AI capability is trustworthy, in operating terms, when a named owner can show:

What the system did, as a record of its actions and outputs; on what basis it did it, as the sources, inputs, and reasoning trail behind them; within what limits it is known to work, as an honest statement of where it is validated and where it is not; and who is accountable, as the named owner who answers for reliance on the output.

If any of those four cannot be shown, the capability is not yet trustworthy for that use, however well it performed in the demonstration. Defined this way, trust is something the owner assembles out of evidence, and that evidence is what lets an accountable decision-maker justify reliance once it exists. It also becomes inspectable by someone who was not in the room, which is exactly what a buyer, an auditor, or a board is asking for.

The evidence pack

The evidence that answers those four questions has a shape, and a research lineage that tells leaders they are not inventing a standard from scratch. Years before the current wave, researchers proposed AI-service FactSheets, modeled on the supplier's declarations of conformity used in other industries, to capture what a consumer of an AI service needs in order to examine it rather than take it on faith. The proposed contents were purpose, performance, safety, security, and provenance, completed by the provider for inspection by the user.²

Translated into enterprise practice, this becomes an evidence pack that travels with an AI capability into production. A workable pack states the capability's purpose and intended use; its measured performance and the conditions that measurement held under; its safety and security properties; the provenance of its data and models; the tested limits of its validation; the human oversight and escalation design; and the named owner. None of it is documentation for its own sake. Each item answers one of the four questions.

The pack also changes the economics of scaling. When trust is tacit, every new deployment reopens the same argument from the beginning, which is slow and inconsistent. When it is packaged as evidence, a capability that has earned reliance in one workflow arrives at the next with its proof attached, and the review covers only what changed. Evidence compounds, and that is what lets an organization go from one trusted use to many without lowering its standard each time.

Keeping the evidence current

A pack assembled once and filed is a snapshot, and AI systems do not hold still. Models drift, inputs shift, use expands past the envelope the system was validated for, and a capability that was trustworthy at launch can quietly stop being so. The public standards treat this as central. The NIST Generative AI Profile treats AI risk as a concern across the full lifecycle, from design and development through deployment, operation, and decommissioning. To keep that risk in view, it suggests continuous

monitoring, structured feedback, red-teaming, independent evaluation, and incident response, along with attention to human-AI configuration and to value-chain risks such as supplier vetting and training-data provenance.³

That maps onto the NIST AI Risk Management Framework's GOVERN, MAP, MEASURE, and MANAGE functions, which frame AI oversight as an ongoing loop for understanding context, assessing performance and risk, and deciding how a system should be used as it operates over time.⁴

The operating consequence is that the evidence pack has to be living. The named owner keeps it current, watches for drift, and refreshes the validation as use expands. A trust claim resting on a stale pack has expired even while the paperwork still sits in the folder.

Framing trust as a lifecycle also settles the usual worry that evidence discipline slows AI down. A living pack is what lets an organization move faster with less risk, because it can extend a trusted capability into new work without re-litigating its trustworthiness from zero, and it can catch a degradation before it becomes an incident.

How an evidence pack fails

Evidence discipline has its own failure modes, and naming them is what keeps a pack honest. Three are common enough to check for by name.

EXHIBIT 2

Three ways an evidence pack stops being evidence

Evidence theater

Volume that records activity instead of trustworthiness - usage dashboards, not the four answers.

False precision

Numbers that imply more validation than was performed, hiding the gap they should expose.

Launch artifact

A pack built to clear an approval and never maintained, decaying while the system changes.

The first is evidence theater: documentation produced for its own sake, thick with usage dashboards and demo decks that count activity. Usage data shows that people tried the tool; it says nothing about whether the tool did the right thing, for the right reason, within known limits, under someone's accountability. A pack that grows in volume without answering the four questions is theater, and a serious reviewer notices the substitution.

The second is false precision: numbers that imply more validation than was actually performed. A performance figure with no statement of the conditions it held under, or a limits section that reassures without marking an honest boundary, gives a reviewer confidence the evidence has not earned. Precision that outruns the testing behind it hides the gap it should expose.

The third is the launch artifact: a pack built to clear an approval and then never touched again. It is the lifecycle failure in a different form, evidence that was true at go-live and has been decaying since while the system it describes keeps changing. A pack that no owner maintains after launch has become a historical record, and it can no longer carry a trust claim.

The common thread is that each failure produces the appearance of evidence without its function. The test in every case is the same: can the named owner still show what the system did, on what basis, within what limits, and answer for it today, not at launch and not in the demo.

One pattern from governance work

The safest illustration is a described operating pattern, with no system named. In governance work, the durable pattern is an AI capability that prepares analysis feeding a human decision. What makes it trustworthy is the evidence it leaves behind: it shows its sources, exposes its reasoning trail, states where it is and is not validated, and routes low-confidence cases to a person instead of guessing. A reviewer can inspect how each conclusion was reached, and a named owner signs off on reliance.

The capability earns its way deeper because each new use inherits that evidence discipline instead of reopening the trust question. And when something does go wrong, as it eventually will, the organization can see it, explain it, and correct it, because the evidence was captured as the work happened. In operation, that is indistinguishable from good governance.

The limits of an evidence pack

Two limits are worth stating plainly. There is no single, universally adopted enterprise standard for an AI evidence pack; the FactSheets research provides a durable model and the NIST profiles provide authoritative lifecycle guidance, but organizations still assemble their own evidence structures.²³ Treat the pack as a pattern to implement; there is no certification to buy.

Evidence also reduces the trust problem without removing judgment. A complete pack still requires a competent human to weigh it, and a well-documented system can still be relied upon past its validated limits by someone who did not read the limits. The named owner remains the load-bearing element, and no amount of documentation substitutes for someone accountable actually exercising judgment about when to rely and when to hold back.

Start with one high-consequence use

The move is narrow on purpose. An organization does not earn trust across its whole AI portfolio at once, and trying to is how evidence discipline turns into theater. Pick the single AI use that carries the highest consequence if it is wrong, the one where reliance matters most and its absence is most dangerous, and build one living, owned, inspectable evidence pack for it. That first pack becomes the template, and the discipline spreads by demand, because the next team can see what reliance actually required.

Return to the pilot parked at the edge, and the reflex to ask for a better model. What actually moves it is evidence a named owner can show and keep current: what the system did, on what basis, within what limits, and who answers for relying on it. That evidence is what earns a capable tool its way into the work that matters and keeps the confidence of the buyers and boards that fund it. Without it, the most capable tool available stays exactly where the last one did, at the safe edge of the business.

Notes and sources

1. Gabriel Morgan Asaftei, Roger Roberts, Abby Sticha, and Cécile Prinsen, "State of AI trust in 2026: Shifting to the agentic era," McKinsey & Company, March 25, 2026. Verified August 3, 2026. <https://www.mckinsey.com/capabilities/tech-and-ai/our-insights/tech-forward/state-of-ai-trust-in-2026-shifting-to-the-agentic-era>
2. Matthew Arnold, Rachel K. E. Bellamy, Michael Hind, Stephanie Houde, Sameep Mehta, Aleksandra Mojsilovic, Ravi Nair, Karthikeyan Natesan Ramamurthy, Darrell Reimer, Alexandra Olteanu, David Piorkowski, Jason Tsay, and Kush R. Varshney, "FactSheets: Increasing Trust in AI Services through Supplier's Declarations of Conformity," arXiv:1808.07261, 2018 (rev. 2019). Verified August 3, 2026. <https://arxiv.org/abs/1808.07261>
3. National Institute of Standards and Technology, "Artificial Intelligence Risk Management Framework: Generative Artificial Intelligence Profile," NIST AI 600-1, July 2024. Verified August 3, 2026. <https://doi.org/10.6028/NIST.AI.600-1>
4. National Institute of Standards and Technology, "AI Risk Management Framework Core," excerpt from AI RMF 1.0, 2023. Verified August 3, 2026. <https://airc.nist.gov/airmf-resources/airmf/5-sec-core/>
5. Deloitte, "The State of AI in the Enterprise," 2026 AI report, Deloitte AI Institute. Verified August 3, 2026. <https://www.deloitte.com/us/en/what-we-do/capabilities/applied-artificial-intelligence/content/state-of-ai-in-the-enterprise.html>
6. Jessica Apotheker, Sylvain Duranton, Vladimir Lukic, Nicolas de Bellefonds, and Christoph Schweizer, "As AI Investments Surge, CEOs Take the Lead," BCG AI Radar 2026, January 15, 2026.

Verified August 3, 2026. <https://www.bcg.com/publications/2026/as-ai-investments-surge-ceos-take-the-lead>

About the author

Marco Policani is an enterprise portfolio, PMO, and AI operating-governance leader. He builds portfolio governance systems where AI carries the analytical load and named humans own every decision — an approach documented across case studies, walkthroughs, and working governance guides on his portfolio site.

Portfolio & governance library: policani.net/governance · Full portfolio: policani.net · LinkedIn: linkedin.com/in/marcpolicani

This white paper may be shared freely with attribution. © 2026 Marco Policani.