


Every Agent Needs a Human Operating Model

An agent belongs in production only after a named human can answer for what it decides



Marco Policani

Enterprise Portfolio · PMO · AI Operating Governance

policani.net · linkedin.com/in/marcpolicani

July 2026

EXECUTIVE SUMMARY

An agent without an operating model is not capability. It is unmanaged work, moving quickly.

Agents act like labor but are bought like technology, and that mismatch is where governance breaks. This paper gives the one-page fix: an operating-model canvas with seven elements a named human must be able to answer before any agent goes live - and the oversight-capacity ceiling that caps how many agents an organization can safely run.

- 01** Before an agent ships, one named person must be able to account for its decisions. If no name attaches, it is not ready.
- 02** Write the decision rights down. Silence defaults to the widest possible authority the agent can claim.
- 03** Oversight capacity, not model capability, caps safe scale. Fund it as a standing cost and make the canvas a funding gate.

CONTENTS

Executive premise

Why agents break the governance model you already have

The vendor is not the owner

An operating model, defined

The Agent Operating-Model Canvas

Reading the gap

Oversight capacity is the real ceiling

What it looks like in practice

What this changes for the portfolio

Evidence gaps this paper cannot close

The shift

Executive premise

An AI agent is not a tool you deploy, and it is not a hire you onboard. It is something organizations have little practice managing: a system that acts like labor but is bought and monitored like technology. That mismatch is where governance breaks. A business can approve an agent, assign it a license, point it at a workflow — and still have no answer for who owns its work when it makes a judgment call, escalates poorly, or produces output that looks finished and is wrong.

The claim of this paper is narrow and testable. An agent belongs in production only after it has an operating model: a role and scope, decision rights, an escalation and stop mechanism, a challenge protocol, acceptance criteria, an evidence duty, and one accountable human owner. Without those seven elements, the organization has not deployed capability. It has deployed unmanaged work.

The market is not short on agent enthusiasm. McKinsey reported in 2025 that 88 percent of respondents were using AI in at least one business function, while only 39 percent reported enterprise-level profit impact [1]. The gap is widest around agents. Gartner projects that more than 40 percent of agentic AI projects will be canceled by the end of 2027, citing rising costs, unclear value, and weak risk controls [2]. Those are not model-quality problems. They are operating-model problems: agents dropped into work that was never redesigned to hold them accountable.

This is a governance argument, not a caution against agents. The two get confused. Slowing an agent program down to satisfy a committee is not governance — it is friction with no owner. Governance is the opposite: what lets an organization move fast and still answer for the result. The winners will not be the companies with the most agents running. They will be the ones whose humans can still answer for what the agents did.

Why agents break the governance model you already have

Traditional operating models were built for control. They assume stable processes, clear handoffs, and predictable outcomes, with a human at each decision point. Agents do not behave that way. They make choices on their own, adapt as they run, and produce results nobody fully predicted. Bolting an autonomous agent onto a model designed for human workers, as Deloitte puts it, is like fitting a jet engine to a bicycle [3].

The deeper problem is one of category. Agents are neither capital nor labor. They act like workers but are funded like technology, and that creates gaps exactly where accountability is supposed to live: decision rights, risk and liability, quality assurance, and performance accountability [3]. An organization knows how to govern a software purchase. It knows how to govern an employee. It has almost no settled practice for a thing bought like the first that behaves like the second.

The evidence that this gap is unmanaged is direct. Deloitte's 2026 enterprise survey found that 84 percent of companies had not redesigned jobs to fit AI, even though automation expectations were high [3]. Redesign is exactly the work that produces decision rights, escalation paths, and ownership. Skipping it does not make the questions disappear. It defers them to the moment an agent does something no one authorized and no one can explain.

And layering agents onto broken processes does not fix those processes. It amplifies them. An agent handed an ambiguous workflow will resolve the ambiguity faster and more confidently than a person would — at a scale that makes the error harder to catch. Moving quickly is not the mistake. The

mistake is treating deployment as the finish line, when it is the moment the governance question becomes real.

The vendor is not the owner

Because agents are bought like technology, there is a standing temptation to believe accountability came with the purchase. It did not. A vendor can promise that a model performs to a spec. It cannot own the consequences of a decision the agent makes inside your workflow, against your data, on behalf of your customer. Accountability is not a feature of the product. It is a property of the operating model you build around the product.

That is why the accountable owner must be an internal, named person, not a contract clause. When an agent misclassifies a case or releases a wrong figure, the question a regulator, customer, or executive asks is not which vendor supplied the model. It is: who here was answerable for letting the agent act? If the honest answer is a procurement line item, the agent did not have an operating model. It had a license.

The practical test is simple, and worth applying before any agent goes live. Name the person who would stand up in a review and account for the agent's decisions. If that name does not exist — or if it is really a group that collectively means no one — the deployment is premature, however capable the model is.

An operating model, defined

An agent has an operating model when a named human can answer six questions before it goes live. What may it decide on its own? What must it escalate? Who owns its output? What evidence must it leave behind? How is it stopped? And what does a good result look like? If any answer is missing, the agent is not ready for production.

This is deliberately a work-design test, not a technical one. Model accuracy, speed, and cost matter, but they are not what fails in the field. What fails is the missing owner when a judgment goes wrong, the missing escalation path when the agent hits a case it should not decide, and the missing evidence trail when someone asks why. Those are governance objects, and they fit on a single page.

The Agent Operating-Model Canvas

The canvas has seven elements. None of them are new to anyone who has run a program or a portfolio. The discipline is applying them to a non-human actor before it starts working — not after it has caused a problem.

EXHIBIT 1

The human accountability map**Role and scope**

What the agent is for - and what it is not for.

Decision rights

Decide alone, get sign-off, or escalate. In writing.

Escalation and stop

A path for the cases it should not handle, and a halt that already exists.

Challenge and acceptance

Define how output is challenged and what good looks like before it arrives.

Evidence and human owner

Keep a reconstructable trail and name one person accountable for reliance.

1. Role and scope

State what the agent is for, and just as important, what it is not for. Scope is the first control because it bounds every other one. An agent scoped to draft a document and flag exceptions is a different governance object from one scoped to send the document to a customer. Scope creep is the most common way an agent ends up making decisions no one designed it to make. Naming the boundary in writing is what makes creep visible before it becomes an incident.

2. Decision rights

Decide in advance what the agent may decide on its own, what needs human sign-off, and what must be escalated. Deloitte frames the same three-way split and adds the crucial rider: set expectations for the evidence the agent must produce to support each decision [3]. Decision rights are where most agent programs are silent, and silence defaults to the widest possible authority. A written rule — the agent may reclassify but not approve, recommend but not release — turns an implicit risk into an explicit control.

3. Escalation and stop mechanism

An agent needs a path for the cases it should not handle, and a way to be halted when something is wrong. Deloitte's guidance is concrete: build stop mechanisms, decide who resolves exceptions, and choose where friction is valuable so errors cannot spread before a human notices [3]. Speed is the point of an agent. Unbounded speed is also how one bad assumption becomes a thousand bad outputs. The stop mechanism has to exist before it is needed, not be improvised during the failure.

4. Challenge protocol

An agent that only confirms is more dangerous than one that occasionally objects, because confident wrong answers are the expensive ones. In mature systems this becomes structural: McKinsey describes critic agents that challenge outputs, guardrail agents that enforce policy, and compliance agents that watch regulation, with every action logged [4]. Most organizations are not there yet, and do not need to be. What they need first is a rule that some outputs get challenged — by a second agent, a checklist, or a human — rather than accepted because they arrived quickly and looked complete. The protocol answers a simple question: before this output is trusted, what had to disagree with it and be overruled?

5. Acceptance criteria

Define what a good result looks like before the agent produces one. Acceptance criteria are the difference between reviewing an agent's work and merely receiving it. Without them, review collapses into a plausibility check — exactly the check a fluent model is best at defeating. Make the criteria specific to the workflow: the fields that must be correct, the sources that must be cited, the thresholds that must not be crossed, and the conditions where the output is rejected rather than edited. Written criteria also make review delegable.

6. Evidence and logging duty

An agent should leave a trail that lets a human reconstruct what it did and why. Deloitte's list is a workable minimum: log and review agent activity, audit behaviors, document rationales, and track disagreements [3]. Public standards help here. The NIST AI Risk Management Framework organizes oversight around govern, map, measure, and manage — an ongoing cycle, not a one-time approval [5]. Its Generative AI Profile treats risk as a lifecycle concern and calls for continuous monitoring, structured feedback, and incident response [6]. The evidence duty is what turns those frameworks from documents into daily practice.

7. Accountable human owner

Every agent needs one named person who owns its outcomes. Not a committee. Not the vendor. McKinsey's framing is that humans move above the loop: defining policies, monitoring outliers, adjusting involvement, rather than reviewing every action line by line [4]. Above the loop is not absent from the loop. The owner answers for the agent's decisions, can invoke the stop mechanism, and is accountable for the evidence trail being real. If no name can be attached to an agent, it should not be in production.

Reading the gap

Most agent failures show up as a symptom in the work long before anyone calls them a governance problem. The discipline is to read the symptom back to the missing canvas element, because each gap routes to a different fix. The table pairs the missing element with how the gap shows up and what it costs.

Missing element	How the gap shows up in the work	What it costs
Role and scope	The agent is quietly doing tasks no one assigned it	Unowned decisions pile up outside any review
Decision rights	No one can say whether the agent was allowed to make the call	Authority defaults to the widest reading; risk is invisible
Escalation and stop	A bad output is discovered only after it has spread	One wrong assumption scales into many wrong actions
Challenge protocol	Confident answers are accepted because they look complete	Fluent errors pass review; quality erodes silently
Acceptance criteria	Review becomes a plausibility check, not a real one	The organization cannot tell good output from convincing output
Evidence and logging	No one can reconstruct why a decision was made	Failures cannot be diagnosed, audited, or learned from
Accountable owner	When something breaks, the answer is "the system did it"	Accountability evaporates; the same failure repeats

The right column is the argument in miniature. Every entry is a governance cost, not a technology cost. None of them is fixed by a better model. All of them are fixed by deciding, in advance, the seven things a human owner must be able to answer.

Oversight capacity is the real ceiling

The canvas raises an operating question leaders underestimate: how many agents can one human actually own? Oversight is finite, and it is becoming the binding constraint. McKinsey's observation from early adopters is that a team of two to five people can already supervise 50 to 100 specialized agents running an end-to-end process [4]. A striking ratio — and a warning. The same report concludes that agentic adoption will be capped by how much oversight capacity humans can provide, making governance itself a potential bottleneck [4].

EXHIBIT 2

Scaling agents without scaling exposure

The instinct Add agents until the models stop improving.	The reality Human oversight runs out first. Unwatched agents scale exposure, not output.
The planning number Span of supervision: how many agents one owner can genuinely answer for.	The funding move Budget oversight as a standing cost. A completed canvas is the funding gate.

This reframes scaling. The instinct is to treat model capability as the limit and add agents until the models stop improving. The operating reality is that human oversight runs out first. An organization that adds agents without adding oversight will not scale its output. It will scale its exposure. This is an inference from the evidence, not a measured law — but the direction is consistent across the consulting and standards literature. Capacity to supervise, not capability to generate, is what caps safe scale.

There is a partial escape, and it is already visible: oversight itself can be scaled with technology — agents guarding other agents, guarded by humans [3]. Critic, guardrail, and compliance agents extend a small human team's reach, and the NIST profile supplies the actions that make such oversight credible: continuous monitoring, red-teaming, independent evaluation, incident loops [6]. But layered oversight does not remove the human ceiling. It raises it. Final accountability still rests with a named person, and the practical planning number — the span of supervision, how many agents one owner can genuinely answer for — belongs in the portfolio conversation next to cost and value.

What it looks like in practice

The safest illustrations are operating patterns, not named systems. In portfolio and PMO work, the recurring shape is the same. An agent is most useful when it is attached to a defined recurring workflow with a named owner and an evidence trail — and most dangerous when it is handed a workflow no one has redesigned.

One pattern from portfolio governance: an agent prepares decision evidence, surfacing inconsistencies and assembling the status a leadership cadence would otherwise spend its time reconstructing. The operating model matters more than the capability. The agent's scope is preparation, not decision. Its decision right is to recommend, not approve. Its acceptance criteria are the fields a reviewer must be able to trust. A named owner signs off before anything reaches the meeting. The value is real, but it is created by the surrounding governance, not the model alone.

A second pattern: an agent proposes reclassifications across a large initiative set. Left unsupervised, it would impose a taxonomy nobody agreed to. Inside an operating model, it recommends, a product owner confirms or amends, every change is logged, and the disagreements are tracked as signal. The agent widens coverage. The human keeps the judgment. That division of labor is the pattern that survives contact with real work.

A third pattern is intake triage: an agent reads incoming requests and routes them. It looks low-risk, which is exactly why it is worth governing. A mis-route is quiet, and a systematic mis-route is invisible until a backlog appears where nobody is watching. The safe version is unglamorous: the agent routes, low-confidence cases escalate instead of being guessed, routing decisions are logged for a weekly drift review, and one owner answers for routing quality. The agents most likely to escape governance are the ones that seem too minor to need it.

What this changes for the portfolio

Read together, the canvas and the oversight ceiling change how an agent program should be funded. An agent is not a one-time purchase whose cost ends at the license. It carries a standing oversight cost — the owner's attention, the review time, the evidence handling — for as long as it runs. A portfolio that funds agents without funding their oversight is understating cost and overstating capacity, and it will discover the gap at the worst time: when too many agents are live for too few owners to answer for.

The governance move is to make the operating model a funding gate. Before an agent is approved, the canvas is complete: owner, decision rights, stop mechanism, acceptance criteria, and a budgeted share of someone's oversight capacity. Agents that cannot clear the gate are not rejected on principle. They are sent back to have their operating model designed — which is usually where the real work was hiding.

Evidence gaps this paper cannot close

Two things here are not settled by current evidence, and naming them is more useful than papering over them. First, the span of supervision is not a known constant. McKinsey's 50-to-100 figure is an early-adopter observation, not a benchmark that transfers to every workflow or regulatory context. The safe number in a high-liability process may be far lower [4]. Treat span of supervision as a variable to measure in your own setting, not a target to hit.

Second, the return on oversight investment is not yet well quantified. The literature is clear that oversight capacity caps safe scale, but it does not tell a leader exactly how much oversight a given agent portfolio needs. Until that evidence matures, the defensible posture is conservative: budget oversight as a real cost, measure the span you can sustain, and let the number of live agents follow the oversight you have — not the other way around.

The shift

The temptation with agents is to measure progress in deployment: how many are running, how much they touch, how fast they move. That number is easy to grow and easy to mistake for value. The number that predicts whether the value survives is different. It is whether a named human can still answer, for each agent: what it may decide, what it must escalate, who owns its output, what evidence it leaves, how it is stopped, and what good looks like.

An agent that can act without an operating model is not autonomy. It is unmanaged work, moving quickly. Give every agent a role, an owner, decision rights, a challenge protocol, acceptance criteria, an evidence duty, and a way to be stopped, and the organization has deployed capability it can answer for. Skip that, and it has deployed something it will eventually have to explain. Build the operating model first. Then let the agent work.

Notes and sources

[1] Alex Singla, Alexander Sukharevsky, Bryce Hall, Lareina Yee, and Michael Chui, with Tara Balakrishnan, "The state of AI in 2025: Agents, innovation, and transformation," McKinsey & Company, November 5, 2025. <https://www.mckinsey.com/capabilities/quantumblack/our-insights/the-state-of-ai>

[2] Gartner, "Gartner Predicts Over 40% of Agentic AI Projects Will Be Canceled by End of 2027," press release, June 25, 2025. <https://www.gartner.com/en/newsroom/press-releases/2025-06-25-gartner-predicts-over-40-percent-of-agentic-ai-projects-will-be-canceled-by-end-of-2027>

[3] David Mallon, Brad Kreit, and Natasha Buckley, "Rethinking operating models for humans with agents," Deloitte Insights, April 2, 2026.

<https://www.deloitte.com/us/en/insights/topics/talent/operating-models-for-humans-ai-agents.html>

[4] Alexander Sukharevsky, Alexis Krivkovich, Arne Gast, Arsen Storozhev, Dana Maor, Deepak Mahadevan, Lari Hämäläinen, and Sandra Durth, "The agentic organization: Contours of the next paradigm for the AI era," McKinsey & Company, September 26, 2025. <https://www.mckinsey.com/capabilities/people-and-organizational-performance/our-insights/the-agentic-organization-contours-of-the-next-paradigm-for-the-ai-era>

[5] National Institute of Standards and Technology, "AI Risk Management Framework Core," excerpt from AI RMF 1.0, 2023. <https://airc.nist.gov/airmf-resources/airmf/5-sec-core/>

[6] National Institute of Standards and Technology, "Artificial Intelligence Risk Management Framework: Generative Artificial Intelligence Profile," NIST AI 600-1, July 2024. <https://doi.org/10.6028/NIST.AI.600-1>

About the author

Marco Policani is an enterprise portfolio, PMO, and AI operating-governance leader. He builds portfolio governance systems where AI carries the analytical load and named humans own every decision — an approach documented across case studies, walkthroughs, and working governance guides on his portfolio site.

Portfolio & governance library: policani.net/governance · Full portfolio: policani.net · LinkedIn: linkedin.com/in/marcpolicani

This white paper may be shared freely with attribution. © 2026 Marco Policani.