
WHITE PAPER · PORTFOLIO & AI GOVERNANCE

Every Agent Needs a Human Operating Model

What has to be true before a capable agent is allowed to act on real work in production

Marco Policani

Enterprise Portfolio · PMO · AI Operating Governance

policani.net · linkedin.com/in/marcpolicani

July 2026

EXECUTIVE SUMMARY

An agent should not act on real work until one named person can answer for what it decides, when it must stop, and what evidence it leaves.

A capable AI agent clears its pilot, and the go-live decision lands on one leader, who must answer for everything it does next. This paper takes a position: before the agent acts on real work, one named person, the owner, should be able to answer seven questions about it, set out on a one-page sheet called the Agent Operating-Model Canvas. Three of the seven follow guidance from Deloitte. The paper adds the single owner, a funding gate (no funding for real work until the canvas is complete), and the canvas as the written record. It also says what would show the position wrong.

- 01** Before an agent ships, one named person must be able to account for its decisions. If no name attaches, it is not ready.
- 02** Write down what the agent may decide, what needs sign-off, and what it must escalate. Silence gives it the widest authority anyone can read into it.
- 03** Fund an agent for real work only when its one-page canvas of seven answers is complete, and put the owner's time in the budget. A missing answer holds the funding. A weak answer narrows what the agent may do.

CONTENTS

The go-live decision	The Agent Operating-Model Canvas
Why your current controls do not fit agents	Three illustrations
The vendor is not the owner	Objections, and what the evidence does not show
Finding the owner, and paying for the owner's time	What to do next
	Notes and sources

The go-live decision

A capable AI agent, software that takes actions on its own, clears its pilot. The decision to put it into production lands on one leader's

desk. Approving it means assigning a license, pointing the agent at a live workflow, and letting it act on real customers and real data. At that moment the leader has no settled answer to a plain question. The agent makes a judgment call, hands the wrong case to the wrong person, or returns an output that looks finished and is wrong. Who in the organization owns that result?

This paper takes a position. An AI agent should not act on real work until one named person, the owner, can answer seven questions about it. What is it for? What may it decide on its own? How does it pass cases it should not decide up to a person, which this paper calls escalating, and how is it halted? What challenges its output? What does a good result look like? What evidence does it leave? Who owns the result?

An answer counts when it is written down and the owner can explain it. Together the answers are the agent's operating model, the people, rules, and records around it, and the one-page sheet that holds them is the Agent Operating-Model Canvas. Real work means work where the agent's output reaches customers, money, records, or decisions without a person redoing it first.

This is a position, not a finding. None of the sources used here compares agents that went live with these answers against agents that went live without them. Such a comparison would test it. Take agents doing similar work where a bad result could reach a customer or a record. Compare those that have a named owner, written rules on what they may decide, and a way to halt them against those that have none of the three. If the ungoverned agents made mistakes no more often, and at no higher cost, the position would be wrong for that kind of work. This paper expects the position to matter least where mistakes are cheap and easy to undo. There it asks for a name and a few sentences per answer. The funding approver, the person who releases the money, decides which work counts as cheap.

Use of AI is widespread. McKinsey's 2025 survey found that 88 percent of respondents used AI regularly in at least one business function, while 39 percent reported an effect on earnings at the enterprise level.[1] That survey covers AI broadly, not agents alone, so it is context for the position, not evidence for it. Gartner forecasts that more than 40 percent of agentic AI projects will be canceled by the end of 2027, citing escalating costs, unclear business value, or inadequate risk controls.[2] That is a prediction, not a count of cancellations. In the release, a Gartner analyst also says current models lack the maturity to reach complex business goals on their own. This paper reads the stated reasons of unclear value and inadequate

risk controls as questions about how agents are chosen and governed, which is the question it takes up.

Governing agents is easy to mistake for slowing them down. A committee that an agent program has to satisfy can be delay no one owns. The aim here is the opposite: let an organization move quickly and still answer for the result.

Why your current controls do not fit agents

Deloitte's authors argue that bolting autonomous agents onto operating models designed for human workers is like fitting a jet engine to a bicycle.[3] This paper's reading is that those models put a person at each decision point, while an agent takes some of those decisions itself and adapts as it runs.

The deeper problem is a category problem. Deloitte writes that agents act like workers but are funded like technology, which creates governance gaps, with ownership especially muddled over who decides what, who carries the risk and liability, who assures quality, and who answers for performance.[3] A company knows how to govern a software purchase and how to govern an employee. This paper's reading is that few companies have a settled practice for something they buy like the first and that acts like the second.

Deloitte's operating-models article also reports a finding from its State of AI in the Enterprise 2026 survey: 84 percent of companies have not redesigned jobs to fit AI. [3] The finding is about AI in general, not agents alone. Redesigning the job is where the rules on what a worker may decide, the escalation paths, and the ownership get written down. Skipping it does not make those questions go away. It leaves them for the moment an agent does something no one authorized and no one can explain.

Deloitte also warns that layering agents onto broken processes does not fix the processes and instead amplifies challenges.[3] A person handed an ambiguous workflow may stop and ask. An agent may resolve the ambiguity quickly and with confidence, in many more cases, which could make the error harder to catch. That is this paper's inference, not Deloitte's finding. Deployment is the point where the governance question becomes real, so it cannot be treated as the finish line.

The vendor is not the owner

Because companies buy agents like technology, they are tempted to treat accountability as part of the purchase. A vendor can warrant that a model performs to a specification. It cannot own the consequences of a decision the agent makes inside your workflow, against your data, on behalf of your customer. Accountability lives in the operating model you build around the product.

That is why the owner has to be a named person inside the company, not a clause in a contract. Say the agent misclassifies a case or releases a wrong figure. The regulator, the customer, or the executive is likely to ask who here was answerable for letting the agent act, more than which vendor supplied the model. If the honest answer is a line in a purchase order, the agent has a license and no operating model.

Finding the owner, and paying for the owner's time

The owner is the person who owns the business result of the workflow the agent works in, the one whose target or budget takes the loss when the agent is wrong. The owner is not the person who sponsored the purchase, and it is not the vendor. A technology team can be the owner only when it also owns the process the agent works in.

Test any candidate with four questions. Who takes the loss if the agent is wrong? Who can change the workflow's rules or narrow the agent's duties? Who can stop the agent today without asking anyone? Who has the time? A candidate who answers yes to all four is the owner. A no on one of the first three means the authority has to be granted in writing, or the owner is someone else. A no on the fourth means time has to be made. If the only name on offer is really a group that adds up to no one, the deployment is premature, however capable the model is.

If the workflow crosses several teams and no one owns the whole, the agent waits until someone is named, because the missing owner is itself the finding. The other route is to narrow the agent to the part of the workflow that one person does own. An owner may hand day-to-day review to a named reviewer. The answer to "what did it decide, and why?" stays with the owner.

The owner's time is a cost that is easy to leave out of a plan. As a planning estimate, not a measured figure, expect the owner and one colleague to spend a few hours writing the first canvas, then a standing part of every week reading the evidence trail, deciding the escalated cases, and staying reachable in case the agent has to be

stopped. Ask the candidate what they will stop doing to make room. If the answer is nothing, the agent has a name and no owner. How many agents one owner can answer for, and how to log that time against live volume, is covered in “The Limit on AI Agents May Be the People Who Supervise Them.”

The Agent Operating-Model Canvas

The canvas has one answer for each of seven elements. Accuracy, latency, and cost matter, but they do not settle who answers when a judgment goes wrong, what happens when the agent meets a case it should not decide, or how anyone learns why a decision was made. Those are work-design questions, and the canvas puts them in one place.

EXHIBIT 1

Who answers for what

1. Role and scope

What the agent is for, and what it is not for.

2. What it may decide

Decide alone, get sign-off, or escalate. In writing.

3. Escalation and stop

A path for cases it should not handle, and a halt that already exists.

4. What challenges it

Say what must challenge the output before anyone trusts it.

5. A good result

Define what good looks like before the output arrives.

6. Evidence it leaves

Keep a trail a human can follow.

7. Who owns the result

One named person answers for it. Not a committee, not the vendor.

Part of the canvas is not new. Deloitte’s article already lists three of the seven elements: define what agents can decide on their own, what needs human sign-off, and what must be escalated, with the evidence agents should produce; build stop mechanisms and clarify who resolves exceptions; and log and review agent activity,

audit behavior, document rationales and human interactions, and track disagreements.[3] The other four, the agent's role and scope, what challenges its output, what a good result looks like, and the single owner, are this paper's framing. McKinsey's authors describe critic and guardrail agents,[4] which relate to the fourth element, and NIST's frameworks call for clear roles and lifecycle oversight.[5][6] As far as this paper can tell, Deloitte's article does not propose a funding gate, a one-page record, or one named accountable person per agent.

This paper adds those three: one named owner who answers for all seven elements; a funding gate, a required checkpoint under which an agent is not funded for real work until its canvas is complete; and the canvas itself as the written record the gate checks. A missing answer holds the funding, and the agent can continue only as a limited test whose output a person checks or redoes before anyone uses it. A limited test is not a pilot: nothing it produces is used until a person has checked it. A weak answer narrows the approval to the work the answer covers. No source tests whether these additions improve outcomes.

The gate asks who answers for the agent and what it is allowed to decide. A different gate, covered in "The Demo-to-Production Gap," asks whether the tool has earned production use through tested hard cases, exception recovery, and funded support. The two can sit in one approval.

The seven elements are a checklist, not a chain, with one exception: the owner is last in the list but first in time. Find the owner before anything else, because the owner has to be able to answer the other six. The other six can be taken in any order.

1. Role and scope

Write down what the agent is for and, just as important, what it is not for. Scope is the first control because it sets the range of every other one. An agent that drafts a document and flags exceptions is a different governance problem from one that sends the document to a customer. Scope creep, where an agent's duties widen without anyone deciding they should, is one way an agent ends up making decisions no one designed it to make. Writing the scope down makes creep visible before it becomes an incident.

2. What it may decide, and what it may not

Decide in advance what the agent may decide on its own, what needs a human sign-off, and what it must escalate. Deloitte splits it the same three ways and adds a rider: set expectations for the evidence the agent must produce behind each decision.[3] If the rule is not written, the agent has the widest authority anyone can read into it. A

written rule, such as it may reclassify but not approve, or recommend but not release, turns a hidden risk into a control you can see.

3. How it escalates and stops

Set the path for the cases the agent should not handle, where it passes them up to a person who can decide, and set the stop switch that halts it when something is wrong. Deloitte's guidance here is concrete: build stop mechanisms, say who resolves exceptions, and decide where slowing the work down is worth it, so errors do not spread before a human notices.[3] Speed is the point of an agent. Speed with nothing to stop it is also how one bad assumption becomes many bad outputs. A stop switch limits how far an error spreads, and it has to exist before it is needed, because one built in the middle of a failure is hard to build well.

4. What has to challenge it

Decide which outputs have to be challenged before anyone trusts them. An agent that only confirms may be more dangerous than one that sometimes objects, because confident wrong answers can be the costly ones. McKinsey's authors describe an agentic organization in which critic agents will challenge outputs, guardrail agents will enforce policy, and compliance agents will monitor regulation, with every action able to be logged and explained in real time.[4] That is their forecast of where companies are heading, not a report of what companies run today. The rule can be light now. A second agent, a checklist, or a human reviewer all count, and accepting an output because it arrived fast and looked complete does not. The test is one question. Before anyone trusts this output, what checked it, and what did that check catch?

5. What a good result looks like

Define what a good result looks like before the agent produces one. Written standards are the difference between reviewing an agent's work and merely receiving it. Without them, review collapses into a check for plausibility, which a fluent model can pass without being right. Keep the standards specific to the workflow. Name the fields that must be correct, the sources that must be cited, the thresholds that must not be crossed, and the point at which the output is rejected rather than edited. Written standards also let someone else do the review, because a second person can apply them without guessing what the first reviewer had in mind.

6. What evidence it leaves

Keep a trail that lets a person work out what the agent did and why. Deloitte's list is a workable minimum: log and review what the agent does, audit its behavior, write down the reasons and the human interactions, and track disagreements.[3]

Public standards give that duty a concrete form. The NIST AI Risk Management Framework organizes risk work around four functions: govern, map, measure, and manage.[5] Its Generative AI Profile says risks can arise from design through decommissioning and suggests actions such as monitoring third-party systems in deployment, structured public feedback exercises such as red-teaming, and incident response.[6] Both are voluntary. Keeping the evidence is how an organization would turn them into daily practice.

7. Who owns the result

Name one person who owns what the agent produces. Not a committee, and not the vendor. McKinsey's authors describe people who set policy, watch for unusual cases, and adjust how involved they are, rather than checking every action line by line.[4] That is still active work, not a name on a chart. The owner answers for the agent's decisions, can use the stop switch, and is accountable for the evidence trail being real. The owner also exercises judgment in three places: deciding from the evidence trail whether to widen, narrow, or stop the agent; deciding which escalated cases need a person's call; and deciding when the written rules have to change. If no name can be attached to an agent, it has no operating model, and it should not be in production.

Three illustrations

Consider three agents. These illustrate how the canvas applies. They are not reports of agents anyone ran.

First, consider an agent that prepares the evidence for a portfolio decision by breaking down plan data and pointing out inconsistencies. Under the canvas, its job is preparation, not decision. It may recommend, not approve. The owner signs off before anything reaches the meeting.

Second, consider an agent that proposes reclassifications across a large set of initiatives. With no operating model, nothing stops it from imposing a filing scheme no one agreed to. Under the canvas, it recommends, a product manager confirms or

amends, every change is logged, and disagreements are tracked. The agent widens coverage, and the person keeps the judgment.

Third, consider an agent that reads incoming requests and routes them to the right queue. It looks low-risk, so it could easily be left out of governance. This paper's inference is that a single mis-route would be quiet, and a steady run of them could stay hidden until a backlog shows up somewhere no one is watching. Under the canvas, low-confidence cases are escalated, the routing is logged for a weekly review, and the owner answers for how well it works.

Objections, and what the evidence does not show

Nothing shows these seven elements are required. That is true, and the opening says so. The position rests on Deloitte's guidance,[3] McKinsey's description of where agentic organizations are heading,[4] two voluntary NIST frameworks,[5][6] and a forecast.[2] No source compares governed and ungoverned agents. Until someone reports that comparison, the paper's judgment is that writing the answers costs a few hours and that an unwritten answer costs more when a mistake arrives. That is a judgment, not a measurement.

Seven elements are too heavy for a small agent. For an agent whose mistakes are cheap and easy to undo, a sentence per answer is enough, and the owner can write it in one sitting. The depth should follow what a mistake would cost. The name stays.

A named owner just becomes the person who gets blamed. An owner with no time, no evidence trail, and no authority to stop the agent has the title without the job. That is why the canvas pairs the name with a stop switch, an evidence trail, and funded time. The point of the owner is a person who can decide, from the trail, to widen, narrow, or stop the agent.

A funding gate slows teams down. It can. A gate that takes weeks may lead teams to run agents around it, as a limited test that never ends. The owner writes the canvas, so a small agent can clear the gate quickly. This paper has no measure of how long a gate takes in practice, so the funding approver should track it.

Critic and guardrail agents will do the challenging. They may. McKinsey describes them,[4] and Deloitte describes agents guarding other agents that are then guarded by humans.[3] They still need a named person who can change or stop them, so the canvas counts them as the answer to element 4 only when someone owns them.

What the evidence does not show. No source measures how much owner time an agent needs, or what extra oversight returns. Gartner's forecast does not say which cancellations governance would have prevented.[2] The McKinsey passages describe a direction, not a result.[4]

What to do next

An agent's cost does not end at the license. The owner's attention, the review time, and the evidence handling continue for as long as it runs, and a portfolio that leaves them out understates the price.

Make a completed canvas a requirement for funding real work. The owner writes it, with a colleague from the team the agent works for: an answer to each of the seven elements, in a few sentences. It is written once before funding and revised when the agent's scope, model, or vendor changes. The funding approver, the person who releases the money, checks that all seven are answered and that the owner's time is in the budget. The rule from the canvas section applies: a missing answer holds the funding, and a weak one narrows the approval.

Agents already live get the same test. The owner completes the canvas now, and until it is complete the agent's scope narrows to work a person checks first. One useful count is how many live agents have an owner who can still answer all seven questions. For any that do not, the first spending is on writing the missing canvas and freeing the owner's time.

Return to the leader with the go-live decision on the desk. It comes down to one question. Can a named person say what this agent may decide, when it must stop, and what evidence it leaves, and is that person's time funded? If so, the agent can go live. If not, it has a license and no operating model. Build the operating model first. Then let the agent work.

Notes and sources

[1] Alex Singla, Alexander Sukharevsky, Bryce Hall, Lareina Yee, and Michael Chui, with Tara Balakrishnan, "The state of AI in 2025: Agents, innovation, and transformation," McKinsey & Company, November 5, 2025.

<https://www.mckinsey.com/capabilities/quantumblack/our-insights/the-state-of-ai-2025>

[2] Gartner, "Gartner Predicts Over 40% of Agentic AI Projects Will Be Canceled by End of 2027," press release, June 25, 2025.

<https://www.gartner.com/en/newsroom/press-releases/2025-06-25-gartner-predicts-over-40-percent-of-agentic-ai-projects-will-be-canceled-by-end-of-2027>

[3] David Mallon, Brad Kreit, and Natasha Buckley, "Rethinking operating models for humans with agents," Deloitte Insights, April 2, 2026.

<https://www.deloitte.com/us/en/insights/topics/talent/operating-models-for-humans-ai-agents.html>

[4] Alexander Sukharevsky, Alexis Krivkovich, Arne Gast, Arsen Storozhev, Dana Maor, Deepak Mahadevan, Lari Hämäläinen, and Sandra Durth, "The agentic organization: Contours of the next paradigm for the AI era," McKinsey & Company, September 26, 2025. <https://www.mckinsey.com/capabilities/people-and-organizational-performance/our-insights/the-agentic-organization-contours-of-the-next-paradigm-for-the-ai-era>

[5] National Institute of Standards and Technology, "AI Risk Management Framework Core," excerpt from AI RMF 1.0, 2023. <https://airc.nist.gov/airmf-resources/airmf/5-sec-core/>

[6] National Institute of Standards and Technology, "Artificial Intelligence Risk Management Framework: Generative Artificial Intelligence Profile," NIST AI 600-1, July 2024. <https://doi.org/10.6028/NIST.AI.600-1>

About the author

Marco Policani is an enterprise portfolio, PMO, and AI operating-governance leader. He builds portfolio governance systems where AI carries the analytical load and named humans own every decision — an approach documented across case studies, walkthroughs, and working governance guides on his portfolio site.

Portfolio & governance library: policani.net/governance · Full portfolio: policani.net · LinkedIn: linkedin.com/in/marcpolicani

This white paper may be shared freely with attribution. © 2026 Marco Policani.