

Governing AI as Operational Change

A practical standard for measuring what enterprise AI really delivers



Marco Policani

Enterprise Portfolio · PMO · AI Operating Governance

policani.net · linkedin.com/in/marcpolicani

July 2026

EXECUTIVE SUMMARY

Meaningful adoption is a condition, not a utilization rate: AI in the critical path of recurring work, changing results you can measure.

Adoption dashboards measure software activity. Leadership funds AI to change work. This paper gives the standard and the three-layer evidence model that close the gap - platform telemetry for breadth, workflow metrics for value, and disciplined comparison for interpretation - plus the governance that turns evidence into scale, stabilize, redesign, or stop decisions.

- 01** Active users and meaningful users are different populations. Averages hide concentration; distributions reveal it.
- 02** Set the evidence threshold before measuring. Programs that set it after define it to fit the answer they already have.
- 03** Skip intrusive monitoring. In experiments, algorithmic surveillance lowered perceived autonomy and raised resistance.

The measurement gaps most programs ignore

Enterprise AI programs have a measurement problem, and it is not the one most teams think it is.

The data exists. Microsoft Copilot, OpenAI Enterprise, and similar platforms expose seat-level usage data by default. Active users. Session counts. Feature engagement. Adoption curves. Most companies can produce these reports within days of deployment. Many treat them as evidence of progress.

They are not. They are evidence of software activity.

Leadership funds AI to change recurring work. Faster cycle times. More throughput. Less rework. Better decisions at the point of action. None of those outcomes appear in an active-user dashboard. The gap between what gets reported and what actually matters is where most AI programs quietly build up risk.

The scale of the problem is not theoretical. In a survey of 644 respondents in the U.S., Germany, and the U.K., Gartner found that estimating and demonstrating AI value is the primary obstacle to AI adoption, cited by 49 percent of respondents — ahead of talent shortages, technical difficulties, and data-related problems.¹

Broad adoption has not produced matching results. MIT NANDA's preliminary 2025 study, as reported by Fortune, found generative AI falling short for 95 percent of companies in its dataset: about 5

percent of AI pilot programs achieved rapid revenue acceleration, while the vast majority stalled with little to no measurable impact on P&L.² A Gartner survey of more than 5,000 digital workers found employees self-report saving 3.6 hours per week — and found the benefit is uneven.³ Turning personal time savings into company-level results remains the central challenge. Gartner separately predicted that at least 30 percent of generative AI projects would be abandoned after proof of concept by the end of 2025, citing poor data quality, inadequate risk controls, escalating costs, or unclear business value.⁴

These numbers describe one problem from different angles: AI is producing activity, and sometimes personal productivity, without reliably producing the business results that justify the spend. AI use is a workflow question, not a software-activity question, and what follows is a practical model for governing it that way. The failure has an old name: Goodhart's Law, the rule that once a proxy becomes the number people manage to, it stops measuring the thing it once stood for. Usage was only ever a proxy for value; managed as the target, it stops tracking value at all.

Why current approaches fall short

Software-level metrics fail for five reasons, and they fail the same way every time.

There is no denominator for real adoption. A user who opened the tool twice looks identical, in most reports, to a user whose weekly workflow now depends on it. Active users and meaningful users are not the same population.

Power users distort the picture. A small group of enthusiasts can make an organization look broadly deployed when the real footprint is narrow. OpenAI's enterprise research found frontier firms generate roughly seven times more messages to specialized GPTs than median enterprises.⁵ Averages hide this. Distributions reveal it.

There is no baseline. If workflow performance was not measured before enablement, post-launch claims are retrospective and anecdotal. Companies treat this step as optional overhead. It is not.

Activity data and outcome data live in different systems. AI usage is logged by user and date. Business outcomes live in Jira, ServiceNow, Salesforce, or a BI layer, logged by ticket, case, or document. Without a deliberate bridge, the two signals cannot be read together.

And there is no accepted evidence threshold. Self-reported time savings are useful as a directional signal, but they are not enough to support an investment decision. Pre/post comparison, matched cohorts, manager validation, and quality sampling are practical replacements. Most AI programs never formally adopt one.

The result: plenty of visible reporting, and no signal a leader can decide on.

A better standard for meaningful adoption

Meaningful adoption is a condition, not a utilization rate:

AI sits in the critical path of recurring work, at enough scale to change workflow results or decision inputs in a way you can measure — across more than a handful of power users.

This is the standard AI investment should be evaluated against. It aligns measurement to what leadership actually funds. It uses data companies already have. And it avoids the trust costs of intrusive monitoring.

Programs that last tend to share this discipline. In Gartner's maturity survey, 45 percent of leaders at high-maturity organizations said their AI initiatives stay in production three years or more, against 20 percent at low-maturity ones. Trust separates them too: in 57 percent of high-maturity organizations, business units trust and are ready to use new AI solutions, against 14 percent at low maturity. Gartner also reports that most high-maturity organizations use defined metrics to evaluate AI project success.⁶

The three-layer evidence model

Meeting the standard takes three layers, read in order — each answers a question the one before it cannot. Platform telemetry shows whether the tool has reached the work. Workflow performance metrics show whether the work's results moved. Explanatory validation shows whether AI is a credible cause of that movement, well enough to act on. Telemetry alone reports activity; workflow metrics without validation report a change no one can yet attribute; only the three together turn data the company already owns into evidence a leader can scale, stabilize, redesign, or stop on.

EXHIBIT 1

Three layers make AI value decision-grade

- 1 Platform telemetry**
Who is active, how often, and is adoption broad or concentrated?
- 2 Workflow performance**
Cycle time, throughput, quality, rework - from the systems that run the operation.
- 3 Explanatory validation**
Pre/post, matched cohorts, manager checks. Decision-grade, not perfect.

Layer one: Platform telemetry

The breadth layer. Use platform analytics to confirm who is active, how often, and whether adoption is broad or concentrated. Microsoft's Copilot reporting architecture separates readiness and adoption reporting from business-value and ROI reporting — a distinction that matters, because mixing the layers produces neither.⁷ This data is easy to get and useful as a consistency check. It is not, by itself, evidence of business value.

Layer two: Workflow performance metrics

The value layer. The questions here are operational. Did cycle time improve? Did throughput rise? Did quality hold? Did rework and escalations fall? The answers already exist in the systems that run the operation. The data does not need to be invented. It needs to be reviewed alongside the telemetry, not separately from it. Microsoft's Copilot Business Impact report makes the link explicit, requiring organizations to upload the operating metrics they want to analyze alongside usage.⁸ These metrics establish that results changed, not what caused the change: a cycle time can fall because volume dropped, a process was redesigned, or the team changed. Reading the outcome beside the telemetry narrows the candidates; confirming AI as the cause is the third layer's job.

Layer three: Explanatory validation

The interpretation layer. When workflow metrics move, something caused the movement. Disciplined comparisons give the causal argument its structure: pre/post analysis on the same workflow, matched cohorts, manager pulse checks, quality sampling. All standard, all practical in enterprise settings. Perfect attribution is not the goal. Decision-grade evidence is.

NIST's AI Risk Management Framework applies a parallel discipline to AI risk in its MEASURE function. It calls for methods to analyze, assess, benchmark, and monitor AI risk and related impacts, and for documenting risks that will not or cannot be measured, a practice that carries over to the attribution problem.⁹

Making it operational

The model asks for four commitments before any workflow enters measurement.

EXHIBIT 2**Every measured workflow ends in one of four decisions****Scale**

Evidence supports further rollout.

Stabilize

Value exists but adoption consistency is uneven.

Redesign

The workflow or enablement approach is flawed.

Stop

No meaningful improvement against agreed metrics.

Governance ownership

Three roles need to be explicit. A Workflow Owner confirms the metrics reflect real operating performance. An Analytics Owner validates source-system integrity. An AI Program Owner manages enablement and coordinates across teams. An executive forum — CIO staff, transformation office, or AI steering council — reviews evidence and makes continuation decisions. Adapt the roles to existing governance; the minimum is that they exist.

Workflow qualification

Not every AI use case warrants formal measurement. A workflow enters the program only if it has enough volume to produce usable data, repeats often enough to observe change, is already tracked in a system of record, and matters enough that leadership will act on the result. OpenAI's enterprise report describes leading firms in similar terms: they run continuous evaluations against real-world outcomes, pair centralized governance with distributed enablement, and embed AI as a core capability rather than a peripheral productivity tool.⁵

Evidence thresholds, set in advance

No workflow should be called value-generating without clearing a minimum standard:

Baseline observation period: 4-8 weeks minimum, adjusted to the workflow's cycle time.

Post-enablement window: 4-8 weeks minimum, matched to the baseline.

Primary operating metrics: at least two per workflow.

Comparison method: at least one — pre/post, matched cohort, or quality sampling.

Telemetry check confirming breadth of use, not just aggregate activity.

Sign-off from both the Workflow Owner and the Analytics Owner.

Define the threshold before measurement begins. Programs that define it after almost always define it to fit the answer they already have.

A review cadence tied to decisions

Reporting that ends in a dashboard is not governance. Monthly operating reviews, plus quarterly executive reviews, should drive one of four explicit outcomes for each workflow:

Status	Meaning
Scale	Evidence supports further rollout
Stabilize	Value exists but adoption consistency is uneven
Redesign	Workflow or enablement approach is flawed
Stop	No meaningful improvement against agreed metrics

These four outcomes are what separate real management from reporting theater.

A note on monitoring and trust

The model deliberately avoids intrusive measurement. No keystroke logging. No screen monitoring. No per-task AI tracking. Not mainly for ethical reasons — for practical ones.

Research by Rachel Schlund and Emily M. Zitek found that participants under algorithmic surveillance perceived less autonomy than those surveilled by humans. Across four experiments with 1,195 students and online-panel participants, algorithmic surveillance also increased intentions to resist and, in two lab idea-generation tasks, reduced performance. Framing mattered: when monitoring was presented as developmental rather than evaluative, the differences in autonomy and resistance diminished.¹⁰

The track record is instructive. In December 2020, Microsoft said it was removing user names and associated actions from Productivity Score after criticism of the preview feature. The company shifted the product to organization-level adoption measures and stated that it should not expose how an individual used Microsoft 365.¹¹

Heavy-handed measurement does not just create legal exposure. It damages the adoption it is supposed to track.

What this looks like in practice

Three workflow types show how the model applies. They are illustrations, not use-case recommendations — proof that the means of measurement are practical and mostly already available.

Service and case workflows

Metrics: case cycle time, throughput per agent per week, reopen rate within 7 days, escalation rate.

Sources: ServiceNow, Zendesk, Salesforce Service Cloud, internal BI.

Evidence standard: a 6-week baseline against a 6-week post-launch window. Confirm usage is not concentrated in a few agents before drawing team-level conclusions.

Document and proposal workflows

Metrics: time to first usable draft, revision cycles, approval cycle time, rework rate.

Sources: document management systems, approval logs, SharePoint, PMO trackers, BI extracts.

Evidence standard: compare AI-assisted teams against similar teams or a prior period. Have managers confirm whether the saved time became volume, quality, or speed.

Requirements and internal analysis workflows

Metrics: draft turnaround time, revision count, decision latency, defect leakage into downstream work.

Sources: Jira, Azure DevOps, PMO artifacts, governance logs.

Evidence standard: pair timing metrics with at least one quality metric, so speed gains are not mistaken for improvement if rework rises later.

From usage reports to workflow evidence

Governed as a software rollout, an AI program produces adoption dashboards, active-user reports, and usage curves. Governed as an operating change, it produces workflow baselines, evidence thresholds, named owners, and explicit decisions about where investment grows or stops. The change is control over what claims the existing data can support, not more data collected.

The central management error in most AI programs is evidentiary, not technical. Companies ask the easiest question first — are people using the tool? — when the question that justifies continued investment is whether AI has changed a recurring workflow enough to warrant scaling.

The few programs producing measurable returns evaluate tools on business outcomes rather than software benchmarks, as the successful buyers in MIT NANDA's study did,² and hold vendors and internal teams to operating metrics. The discipline is uncommon rather than exotic. If a program has delivered, it should be able to demonstrate it; if it cannot, the honest answer is that it is still experimenting — a legitimate stage, but a different claim than operational impact.

AI that does not change a workflow metric leadership already cares about may be promising. It has not yet reached meaningful adoption.

Notes and sources

1. Gartner, "Gartner Survey Finds Generative AI Is Now the Most Frequently Deployed AI Solution in Organizations," May 7, 2024. Verified August 3, 2026. <https://www.gartner.com/en/newsroom/press-releases/2024-05-07-gartner-survey-finds-generative-ai-is-now-the-most-frequently-deployed-ai-solution-in-organizations>
2. Sheryl Estrada, "MIT report: 95% of generative AI pilots at companies are failing," Fortune, August 18, 2025, reporting findings from Aditya Challapally, Chris Pease, Ramesh Raskar, and Pradyumna Chari, "The GenAI Divide: State of AI in Business 2025," MIT NANDA. Verified August 3, 2026. <https://fortune.com/2025/08/18/mit-report-95-percent-generative-ai-pilots-at-companies-failing-cfo/>
3. Gartner, "Gartner Identifies Four Emerging Challenges to Delivering Value from AI Safely and at Scale," October 21, 2024. Verified August 3, 2026. <https://www.gartner.com/en/newsroom/press-releases/2024-10-21-gartner-identifies-four-emerging-challenges-to-delivering-value-from-ai-safely-and-at-scale>
4. Gartner, "Gartner Predicts 30% of Generative AI Projects Will Be Abandoned After Proof of Concept by End of 2025," July 29, 2024. Verified August 3, 2026. <https://www.gartner.com/en/newsroom/press-releases/2024-07-29-gartner-predicts-30-percent-of-generative-ai-projects-will-be-abandoned-after-proof-of-concept-by-end-of-2025>
5. OpenAI, "The State of Enterprise AI: 2025 Report," 2025. Verified August 3, 2026. <https://openai.com/business/guides-and-resources/the-state-of-enterprise-ai-2025-report/>
6. Gartner, "Gartner Survey Finds 45% of Organizations With High AI Maturity Keep AI Projects Operational for at Least Three Years," June 30, 2025. Verified August 3, 2026. <https://www.gartner.com/en/newsroom/press-releases/2025-06-30-gartner-survey-finds-forty-five-percent-of-organizations-with-high-artificial-intelligence-maturity-keep-artificial-intelligence-projects-operational-for-at-least-three-years>
7. Microsoft, "Copilot controls measurement and reporting," Microsoft Learn, updated February 25, 2026. Verified September 15, 2026. <https://learn.microsoft.com/en-us/microsoft-365/copilot/copilot-controls/measurement-reporting>

8. Microsoft, "Copilot Business Impact Report," Microsoft Learn, updated February 17, 2026. Verified August 3, 2026. <https://learn.microsoft.com/en-us/viva/insights/advanced/analyst/templates/copilot-business-impact>
9. National Institute of Standards and Technology, "AI Risk Management Framework Core," excerpt from Artificial Intelligence Risk Management Framework (AI RMF 1.0), January 2023. Verified August 3, 2026. <https://airc.nist.gov/airmf-resources/airmf/5-sec-core/>
10. Rachel Schlund and Emily M. Zitek, "Algorithmic versus human surveillance leads to lower perceptions of autonomy and increased resistance," Communications Psychology 2, article 53, June 6, 2024. Verified August 3, 2026. <https://www.nature.com/articles/s44271-024-00102-8>
11. Jared Spataro, "Our commitment to privacy in Microsoft Productivity Score," Microsoft 365 Blog, December 1, 2020. Verified August 3, 2026. <https://www.microsoft.com/en-us/microsoft-365/blog/2020/12/01/our-commitment-to-privacy-in-microsoft-productivity-score/>

About the author

Marco Policani is an enterprise portfolio, PMO, and AI operating-governance leader. He builds portfolio governance systems where AI carries the analytical load and named humans own every decision — an approach documented across case studies, walkthroughs, and working governance guides on his portfolio site.

Portfolio & governance library: policani.net/governance · Full portfolio: policani.net · LinkedIn: linkedin.com/in/marcpolicani

This white paper may be shared freely with attribution. © 2026 Marco Policani.