
WHITE PAPER · PORTFOLIO & AI GOVERNANCE

Measure AI by the Workflow It Changes

Three evidence checks turn AI usage into a choice

Marco Policani

Enterprise Portfolio · PMO · AI Operating Governance

policani.net · linkedin.com/in/marcpolicani

July 2026

EXECUTIVE SUMMARY

Measure one recurring workflow in three layers: tool use, workflow results, and evidence that AI contributed.

An AI dashboard can be green while the workflow stays the same. This paper measures one recurring workflow in three layers: who used the tool, what result changed, and whether AI is a credible reason for the change.

- 01** Treat tool use, workflow results, and the reason for a change as three separate questions.
- 02** Set the baseline, quality measure, comparison, and owners before the team sees the answer.
- 03** End each review with a choice to continue, hold, change, or stop. Measure the workflow, using the least personal data that can answer the question.

The dashboard can be green while the work stays the same

Most AI tools can report active users, sessions, and feature use within days. Those numbers answer a useful question: Did people open the tool? They do not answer the question leaders funded the program to solve: Did recurring work get faster, better, or less costly?

That gap is common. In a Gartner survey of 644 respondents, 49 percent named estimating and showing AI value as the main barrier to adoption.[1] Another Gartner survey found that workers reported saving 3.6 hours per week, but the gains were uneven.[2] Self-reported time saved may be real. It still does not show what the organization gained from that time.

A preliminary MIT NANDA study, reported by Fortune, found that about 5 percent of the AI pilots in its dataset produced rapid revenue gains. Most showed little or no measurable effect on profit and loss.[3] The finding comes through a secondary report and does not cover every kind of AI work. Gartner also predicted that at least 30 percent of generative AI projects would be dropped after proof of concept by the end of 2025.[4] That was a forecast, not an observed result. Together, these findings

explain why leaders should require evidence beyond activity reports before approving another investment.

The unit of measurement should be the recurring workflow: a service case, proposal, requirement, analysis, approval, or other piece of work that already has an owner and a result.

Start with the result leaders care about

Choose a workflow where leaders will act on the evidence. Write down the result the organization already cares about. It may be cycle time, throughput, error rate, rework, escalation, approval time, or a decision-quality check.

EXHIBIT 1

Three layers connect AI use to a workflow result

- ① **Tool use**
Did AI reach the people and tasks in the workflow?
- ② **Workflow result**
Did cycle time, quality, rework, or another agreed result change?
- ③ **Reason for change**
Does the comparison show that AI likely helped cause the movement?

Then record a baseline before changing the workflow. Without a baseline, a later claim rests on memory and opinion. Four to eight weeks may be enough for frequent work. A slower workflow may need a longer period. The point is to cover enough normal variation to make the comparison useful.

Tool use still matters. If almost no one uses the AI, the program cannot reasonably claim that AI changed the workflow. Yet averages can hide a narrow footprint. OpenAI reports that firms it calls frontier firms send about seven times more messages to specialized GPTs than the median enterprise.[5] That vendor report describes use patterns, not business results. It is a reminder to look at the spread of use instead of one average.

Ask three questions in order:

- Did the tool reach the people and tasks in the workflow?

- Did the workflow result change?
- Is AI a credible reason for that change?

Each question needs different evidence.

AI value needs evidence on all three questions: the tool reached the workflow, the workflow result changed, and AI credibly contributed. Leaders should decide in advance what evidence will trigger the next choice.

Layer one: Did AI reach the workflow?

Platform data shows who used the tool, how often, and whether use was broad or limited to a few enthusiasts. Microsoft separates Copilot readiness and adoption reports from reports about business value and return on investment.[6] That separation matters. If AI never reaches the workflow, it cannot change the result. Measure reach first, then measure the result separately.

Read the distribution, not only the total. Check whether the people who perform the target workflow are active. Look for teams with little or no use, sudden drops, and use concentrated in a handful of people. These findings tell the program where to investigate. They do not prove value.

Protect trust while measuring

The program does not need keystroke logs, screen recordings, or a record of every prompt. Those methods can damage the behavior the program is trying to understand.

In four studies with 1,195 students and online participants, Schlund and Zitek found that algorithmic monitoring reduced perceived autonomy and raised resistance compared with human monitoring. Two lab idea-generation tasks also showed lower performance. The setting limits how far the result can be applied to employees, but it shows a risk worth testing in workplace settings. The negative effects were smaller when monitoring was framed as developmental rather than evaluative.[7]

Microsoft faced a similar trust problem with its Productivity Score preview. In 2020, it removed user names and their associated actions and shifted toward organization-level measures.[8]

Use the least personal data that can answer the workflow question. Tell people what is measured, why it matters, who can see it, and which decision it supports.

Layer two: Did the workflow result change?

The second layer comes from the system that already records the work. A service team may use case cycle time, reopen rate, and escalations. A proposal team may use time to first usable draft, revision rounds, and approval time. An analysis team may use turnaround time, defects found later, and the time a decision remained open.

Microsoft's Copilot Business Impact report asks organizations to upload business outcome data beside tool-use data.[9] The method creates the needed bridge between the tool and the workflow. It does not prove that AI caused the change.

Other changes can move a result. Faster work can hide errors. Pair speed or volume with a quality measure so gains do not hide rework.

Layer three: Why did the result change?

The third layer tests whether AI is a credible reason for the movement. Use the strongest practical comparison the workflow allows:

- compare the same workflow before and after the change;
- compare similar teams with and without the AI support;
- sample work for quality before and after;
- ask managers what changed and check their account against the records.

Most teams cannot run a controlled experiment while serving customers. The goal is evidence strong enough for the next decision. State what the comparison can and cannot support.

The NIST AI RMF Core uses a similar discipline for risk. Its MEASURE function calls for analysis, tests, comparison, monitoring, and a record of risks that cannot be measured.[10] The same habit helps here: document uncertainty instead of hiding it behind a precise chart.

Set the rules before seeing the result

For each workflow, name three owners. The workflow owner confirms that the measures reflect real work. The data owner checks that the source records are complete and reliable. The AI program owner manages access, support, and changes to the tool. One named leader or review group then decides what happens next.

Write the evidence rule before measurement begins. It should state:

- the baseline and follow-up periods;
- two or more workflow measures, including quality;
- the comparison method;
- the level and spread of tool use needed for a fair test;
- the people who must confirm the result.

Do not set a universal improvement target. A useful change depends on the cost, risk, and volume of the workflow. Set the threshold for the specific decision before the team knows the answer.

Gartner reports that high-maturity organizations more often keep AI work live for at least three years, report stronger business trust, and use defined success measures. [11] The survey shows an association, not proof that the measures caused better results. It still supports the practice of agreeing on success before judging it.

End every review with a choice

A dashboard is useful only when it leads to a decision. Each workflow review should end with one of four choices:

EXHIBIT 2

Every review ends with one of four choices

Continue

The evidence supports more use in the same workflow.

Hold

Keep current use while more evidence develops.

Change

Adjust the workflow, tool, training, or review rule and test again.

Stop

The result or risk does not justify continued use.

- Continue: the evidence supports more use in the same workflow.
- Hold: keep the current use while more evidence develops.
- Change: adjust the workflow, tool, training, or review rule and test again.
- Stop: end the use because the result or risk does not justify it.

The record should name the choice, owner, reason, next check, and any condition that would reopen the decision. That record lets the decision owner explain why the team continued, held, changed, or stopped the use.

What to do next

Choose one recurring workflow with enough volume to measure. Before the next AI review:

- name the workflow, data, and program owners;
- write the baseline and follow-up periods;
- choose one speed or volume measure and one quality measure;
- decide how tool use will be matched to workflow records;
- choose the comparison method;

- write the four possible decisions.

After the follow-up period, ask one question: Which workflow result changed, by how much, against what baseline, and with what evidence that AI contributed?

Judge the result against the evidence rule written before the test. Then continue, hold, change, or stop.

If the program can answer, leaders have evidence for the next decision. If it cannot, the work may still be a useful experiment. It is not yet proof that the operation improved.

Notes and sources

[1] Gartner, "Gartner Survey Finds Generative AI Is Now the Most Frequently Deployed AI Solution in Organizations," May 7, 2024.

<https://www.gartner.com/en/newsroom/press-releases/2024-05-07-gartner-survey-finds-generative-ai-is-now-the-most-frequently-deployed-ai-solution-in-organizations>

[2] Gartner, "Gartner Identifies Four Emerging Challenges to Delivering Value from AI Safely and at Scale," October 21, 2024.

<https://www.gartner.com/en/newsroom/press-releases/2024-10-21-gartner-identifies-four-emerging-challenges-to-delivering-value-from-ai-safely-and-at-scale>

[3] Sheryl Estrada, "MIT report: 95% of generative AI pilots at companies are failing," Fortune, August 18, 2025, reporting findings from Aditya Challapally, Chris Pease, Ramesh Raskar, and Pradyumna Chari, "The GenAI Divide: State of AI in Business 2025," MIT NANDA. <https://fortune.com/2025/08/18/mit-report-95-percent-generative-ai-pilots-at-companies-failing-cfo/>

[4] Gartner, "Gartner Predicts 30% of Generative AI Projects Will Be Abandoned After Proof of Concept by End of 2025," July 29, 2024.

<https://www.gartner.com/en/newsroom/press-releases/2024-07-29-gartner-predicts-30-percent-of-generative-ai-projects-will-be-abandoned-after-proof-of-concept-by-end-of-2025>

[5] OpenAI, "The State of Enterprise AI: 2025 Report," 2025.

<https://openai.com/business/guides-and-resources/the-state-of-enterprise-ai-2025-report/>

- [6] Microsoft, "Copilot controls measurement and reporting," Microsoft Learn, updated February 25, 2026. <https://learn.microsoft.com/en-us/microsoft-365/copilot/copilot-controls/measurement-reporting>
- [7] Rachel Schlund and Emily M. Zitek, "Algorithmic versus human surveillance leads to lower perceptions of autonomy and increased resistance," Communications Psychology 2, article 53, June 6, 2024. <https://www.nature.com/articles/s44271-024-00102-8>
- [8] Jared Spataro, "Our commitment to privacy in Microsoft Productivity Score," Microsoft 365 Blog, December 1, 2020. <https://www.microsoft.com/en-us/microsoft-365/blog/2020/12/01/our-commitment-to-privacy-in-microsoft-productivity-score/>
- [9] Microsoft, "Copilot Business Impact Report," Microsoft Learn, updated February 17, 2026. <https://learn.microsoft.com/en-us/viva/insights/advanced/analyst/templates/copilot-business-impact>
- [10] National Institute of Standards and Technology, "AI Risk Management Framework Core," excerpt from Artificial Intelligence Risk Management Framework (AI RMF 1.0), January 2023. <https://airc.nist.gov/airmf-resources/airmf/5-sec-core/>
- [11] Gartner, "Gartner Survey Finds 45% of Organizations With High AI Maturity Keep AI Projects Operational for at Least Three Years," June 30, 2025. <https://www.gartner.com/en/newsroom/press-releases/2025-06-30-gartner-survey-finds-forty-five-percent-of-organizations-with-high-artificial-intelligence-maturity-keep-artificial-intelligence-projects-operational-for-at-least-three-years>

About the author

Marco Policani is an enterprise portfolio, PMO, and AI operating-governance leader. He builds portfolio governance systems where AI carries the analytical load and named humans own every decision — an approach documented across case studies, walkthroughs, and working governance guides on his portfolio site.

Portfolio & governance library: policani.net/governance · Full portfolio: policani.net · LinkedIn: linkedin.com/in/marcpolicani

This white paper may be shared freely with attribution. © 2026 Marco Policani.