

9

WHITE PAPER · PORTFOLIO & AI GOVERNANCE

Governing AI as Operational Change

A practical standard for measuring what enterprise AI really delivers

Marco Policani

Enterprise Portfolio · PMO · AI Operating Governance

policani.net · linkedin.com/in/marcpolicani

July 2026

EXECUTIVE SUMMARY**Meaningful adoption is a condition, not a utilization rate: AI in the critical path of recurring work, changing results you can measure.**

Adoption dashboards measure software activity. Leadership funds AI to change work. This paper gives the standard and the three-layer evidence model that close the gap - platform telemetry for breadth, workflow metrics for value, and disciplined comparison for interpretation - plus the governance that turns evidence into scale, stabilize, redesign, or stop decisions.

- 01** Active users and meaningful users are different populations. Averages hide concentration; distributions reveal it.
- 02** Set the evidence threshold before measuring. Programs that set it after define it to fit the answer they already have.
- 03** Skip intrusive monitoring. AI-driven surveillance measurably damages the adoption it claims to track.

The measurement gaps most programs ignore

Enterprise AI programs have a measurement problem, and it is not the one most teams think it is.

The data exists. Microsoft Copilot, OpenAI Enterprise, and similar platforms expose seat-level usage data by default. Active users. Session counts. Feature engagement. Adoption curves. Most companies can produce these reports within days of deployment. Many treat them as evidence of progress.

They are not. They are evidence of software activity.

Leadership funds AI to change recurring work. Faster cycle times. More throughput. Less rework. Better decisions at the point of action. None of those outcomes appear in an active-user dashboard. The gap between what gets reported and what actually matters is where most AI programs quietly build up risk.

The scale of the problem is not theoretical. In a survey of 644 companies, Gartner found that demonstrating AI value is the single biggest adoption barrier, cited by 49 percent of respondents — ahead of talent shortages, technical challenges, and data problems.

Over 80 percent of companies have explored or piloted generative AI. Nearly 40 percent report deployment. But deployment has not produced results. MIT's NANDA initiative found that 95 percent of enterprise AI programs deliver no measurable profit impact, despite \$30-40 billion in annual spend.

The same Gartner survey found employees self-report saving 3.6 hours per week — and found the benefit is uneven. Turning personal time savings into company-level results remains the central challenge. Gartner separately projects at least 30 percent of generative AI projects will be dropped after proof of concept, mostly over cost and unclear value.

These numbers describe one problem from different angles. AI is producing activity, and sometimes personal productivity, without reliably producing the business results that justify the spend. The argument here is simple. AI use is a workflow question, not a software-activity question. What follows is a practical model for governing it that way.

Why current approaches fall short

Software-level metrics fail for five reasons, and they fail the same way every time.

There is no denominator for real adoption. A user who opened the tool twice looks identical, in most reports, to a user whose weekly workflow now depends on it. Active users and meaningful users are not the same population.

Power users distort the picture. A small group of enthusiasts can make an organization look broadly deployed when the real footprint is narrow. OpenAI's enterprise research found frontier firms generate roughly seven times more messages to specialized AI tools than median enterprises. Averages hide this. Distributions reveal it.

There is no baseline. If workflow performance was not measured before enablement, post-launch claims are retrospective and anecdotal. Companies treat this step as optional overhead. It is not.

Activity data and outcome data live in different systems. AI usage is logged by user and date. Business outcomes live in Jira, ServiceNow, Salesforce, or a BI layer, logged by ticket, case, or document. Without a deliberate bridge, the two signals cannot be read together.

And there is no accepted evidence threshold. Companies reject weak self-reporting but fail to replace it with anything better. The research on self-reported performance is clear: people present their own work favorably, and the effect worsens under evaluation. Pre/post comparison, matched cohorts, and manager validation are practical replacements. Most AI programs never formally adopt one.

The result: plenty of visible reporting, and no signal a leader can decide on.

A better standard for meaningful adoption

Meaningful adoption is not a utilization rate. It is a condition:

AI sits in the critical path of recurring work, at enough scale to change workflow results or decision inputs in a way you can measure — across more than a handful of power users.

This is the standard AI investment should be evaluated against. It aligns measurement to what leadership actually funds. It uses data companies already have. And it avoids the trust costs of intrusive monitoring.

The programs that last are the ones that set this discipline early. Gartner's maturity research found 45 percent of high-maturity companies keep AI in production three years or more, against 20 percent of low-maturity ones. Trust separates them too: 57 percent of high-maturity companies say business units trust and will use new AI, against 14 percent at low maturity. Explicit metrics, the survey found, feed directly into program success.

The three-layer evidence model

Meeting the standard takes three layers working together — not any one of them alone.

EXHIBIT 1

Three layers make AI value decision-grade

- 1 **Platform telemetry**
Who is active, how often, and is adoption broad or concentrated?
- 2 **Workflow performance metrics**
Cycle time, throughput, quality, rework - from the systems that run the operation.
- 3 **Explanatory validation**
Pre/post, matched cohorts, manager checks. Decision-grade, not perfect.

Layer one: Platform telemetry

The breadth layer. Use platform analytics to confirm who is active, how often, and whether adoption is broad or concentrated. Microsoft's Copilot analytics architecture separates operational reporting (license assignment, adoption tracking) from strategic reporting (workflow impact, business outcomes) — a distinction that matters, because mixing the layers produces neither. This data is easy to get and useful as a consistency check. It is not, by itself, evidence of business value.

Layer two: Workflow performance metrics

The value layer. The questions here are operational. Did cycle time improve? Did throughput rise? Did quality hold? Did rework and escalations fall? The answers already exist in the systems that run the operation. The data does not need to be invented. It needs to be reviewed alongside the telemetry, not

separately from it. Microsoft's Copilot Business Impact framework makes the link explicit, recommending companies upload key operating metrics into the analytics layer so usage and outcomes can be compared.

Layer three: Explanatory validation

The interpretation layer. When workflow metrics move, something caused the movement. Disciplined comparisons give the causal argument its structure: pre/post analysis on the same workflow, matched cohorts, manager pulse checks, quality sampling. All standard, all practical in enterprise settings. Perfect attribution is not the goal. Decision-grade evidence is.

NIST's AI Risk Management Framework backs this approach in its MEASURE function. It calls for tools to analyze, benchmark, and monitor AI impacts — and for writing down what cannot yet be measured reliably, which speaks directly to the attribution problem.

Making it operational

The model asks for four commitments before any workflow enters measurement.

EXHIBIT 2

Every measured workflow ends in one of four decisions

Scale

Evidence supports further rollout.

Stabilize

Value exists but adoption consistency is uneven.

Redesign

The workflow or enablement approach is flawed.

Stop

No meaningful improvement against agreed metrics.

Governance ownership

Three roles need to be explicit. A Workflow Owner confirms the metrics reflect real operating performance. An Analytics Owner validates source-system integrity. An AI Program Owner manages enablement and coordinates across teams. An executive forum — CIO staff, transformation office, or AI steering council — reviews evidence and makes continuation decisions. Adapt the roles to existing governance; the minimum is that they exist.

Workflow qualification

Not every AI use case warrants formal measurement. A workflow enters the program only if it has enough volume to produce usable data, repeats often enough to observe change, is already tracked in a system of record, and matters enough that leadership will act on the result. OpenAI's enterprise guidance points the same way: measure where there is a measurable change, a defined workflow, and alignment to priorities.

Evidence thresholds, set in advance

No workflow should be called value-generating without clearing a minimum standard:

Baseline observation period: 4-8 weeks minimum, adjusted to the workflow's cycle time.

Post-enablement window: 4-8 weeks minimum, matched to the baseline.

Primary operating metrics: at least two per workflow.

Comparison method: at least one — pre/post, matched cohort, or quality sampling.

Telemetry check confirming breadth of use, not just aggregate activity.

Sign-off from both the Workflow Owner and the Analytics Owner.

Define the threshold before measurement begins. Programs that define it after almost always define it to fit the answer they already have.

A review cadence tied to decisions

Reporting that ends in a dashboard is not governance. Monthly operating reviews, plus quarterly executive reviews, should drive one of four explicit outcomes for each workflow:

Status	Meaning
Scale	Evidence supports further rollout
Stabilize	Value exists but adoption consistency is uneven
Redesign	Workflow or enablement approach is flawed
Stop	No meaningful improvement against agreed metrics

These four outcomes are what separate real management from reporting theater.

A note on monitoring and trust

The model deliberately avoids intrusive measurement. No keystroke logging. No screen monitoring. No per-task AI tracking. Not mainly for ethical reasons — for practical ones.

Research from Cornell's ILR School found that employees under AI-driven monitoring report a much greater loss of autonomy than those monitored by humans. They are also measurably more likely to resist, complain, perform worse, and consider quitting. Four controlled experiments, 1,195 participants, same result. Framing mattered: positioned as support for growth, the negative effects largely disappeared. Framed as performance evaluation, they did not.

The track record is instructive. In November 2020, Microsoft's Productivity Score feature exposed per-person usage data across Microsoft 365. It showed how much each named employee emailed and communicated. Within days of public criticism, Microsoft removed all individual names and per-person metrics from the product.

Heavy-handed measurement does not just create legal exposure. It damages the adoption it is supposed to track.

What this looks like in practice

Three workflow types show how the model applies. These are not use-case recommendations. They are proof that the means of measurement are practical and mostly already available.

Service and case workflows

Metrics: case cycle time, throughput per agent per week, reopen rate within 7 days, escalation rate.

Sources: ServiceNow, Zendesk, Salesforce Service Cloud, internal BI.

Evidence standard: a 6-week baseline against a 6-week post-launch window. Confirm usage is not concentrated in a few agents before drawing team-level conclusions.

Document and proposal workflows

Metrics: time to first usable draft, revision cycles, approval cycle time, rework rate.

Sources: document management systems, approval logs, SharePoint, PMO trackers, BI extracts.

Evidence standard: compare AI-assisted teams against similar teams or a prior period. Have managers confirm whether the saved time became volume, quality, or speed.

Requirements and internal analysis workflows

Metrics: draft turnaround time, revision count, decision latency, defect leakage into downstream work.

Sources: Jira, Azure DevOps, PMO artifacts, governance logs.

Evidence standard: pair timing metrics with at least one quality metric, so speed gains are not mistaken for improvement if rework rises later.

The shift this produces

Before this model, AI is governed as a software rollout. Adoption dashboards. Active-user reports. Usage curves. General claims about productivity.

After it, AI is governed as an operating change. Workflow baselines. Evidence thresholds. Named owners. Explicit decisions about where investment grows or stops.

The shift is not about collecting more data. It is about better control over what claims can be made from the data that already exists.

The underlying argument

The central management error in most AI programs is not technical. It is a problem of evidence.

Companies ask the easiest question first: are people using the tool? That question has a place, but only at the surface. The question that justifies continued investment is different: has AI changed a recurring workflow in a way that warrants scaling?

The few programs producing measurable returns share one trait. They evaluate tools on business outcomes, not software benchmarks, and they hold vendors and internal teams to operational metrics. The discipline is not exotic. It is just uncommon.

If a program has delivered, it should be able to demonstrate it. If it cannot, the honest answer is that it is still experimenting — a legitimate stage, but a different claim than operational impact.

AI that does not change a workflow metric leadership already cares about may be promising. It is not yet meaningful adoption.

About the author

Marco Policani is an enterprise portfolio, PMO, and AI operating-governance leader. He builds portfolio governance systems where AI carries the analytical load and named humans own every decision — an approach documented across case studies, walkthroughs, and working governance guides on his portfolio site.

Portfolio & governance library: policani.net/governance · Full portfolio: policani.net · LinkedIn: linkedin.com/in/marcpolicani

This white paper may be shared freely with attribution. © 2026 Marco Policani.