
WHITE PAPER · PORTFOLIO & DELIVERY GOVERNANCE

A PMO Dashboard Should Show Investment Health

A decision interface, not a longer status deck

Marco Policani

Enterprise Portfolio · PMO · AI Operating Governance

policani.net · linkedin.com/in/marcpolicani

July 2026

EXECUTIVE SUMMARY

A dashboard earns attention when it changes an investment while options remain.

A portfolio dashboard built from project-tool exports can describe delivery activity honestly and still support funding decisions poorly. This paper replaces the status deck with one page per investment and six lenses, which are six questions asked of every investment: outcome, confidence, capacity, adoption, dependency, and decision readiness. Each lens stays only if its answer changes a decision a leader can still make.

- 01** Status asks if work is proceeding as planned. Health asks if this is still a good use of money, people, and attention.
- 02** Do not combine the lenses into one score. When they disagree, such as strong delivery with falling confidence, the disagreement is the information, and a blended number would hide it.
- 03** Start by pruning your current dashboard: retire any field that has not changed a decision in two review cycles, then build the page for your largest investments.

The argument

A dashboard earns attention when it helps leaders change an investment while options remain. Schedule and spend matter, but neither can show whether the value is still plausible, the people are available, the receiving operation is ready, the outside commitments will hold, or a decision is overdue.

Here is a scene I have watched play out in a portfolio review. The deck is long and accurate in every cell. Near the end, a CFO asks a simple question about the largest investment on the page: if we stopped this today, what would we actually lose? Nobody can answer from the deck. The deck knows the burn rate, the milestone variance, and the issue count. It does not know what the money is still expected to buy, how confident anyone is, or what decision would change the outcome. The one question that justified the meeting has no page.

A deck built from whatever the project tools can export can look like that one. It describes delivery activity honestly and supports investment decisions poorly. More pages do not fix the gap. Changing what the page is for does: from describing the work to supporting the next real choice about each investment.

This paper is written for portfolio executives, sponsors, PMO leaders, and investment owners. It proposes one page per investment, built from six lenses. A lens here is a question asked of every investment, kept only because its answer changes a decision a leader can still make. The six are outcome, confidence, capacity, adoption, dependency, and decision readiness. They are this paper's own design. The sources it cites show that status reporting can miss value and that dashboard ratings depend on upkeep. None tests this set or shows that a dashboard built this way leads to better decisions. The paper offers the page as a management instrument, not a scoring formula, and ends with signals for testing it.

The page is also different from the one filed at each funding decision, which "Portfolio Governance Is a Funding-Discipline Problem, Not a Reporting Problem" covers. That page already carries outcome evidence and the capacity a project holds. This is the standing page a forum reads between funding decisions: it reads those two again each month and adds confidence, adoption, dependency, and open decisions. Earlier papers cover other pieces it reads, and each is named where it applies.

Why status is not health

Status answers one question: is the work proceeding as planned? Health answers a different one: is this still a good use of the organization's money, people, and attention? The two can diverge. A project can be green on plan and dead on value, for example because the market moved, the sponsor left, or the process it automates is being retired. It can be red on plan and healthy on value: late, over budget, and still the best available path to an outcome that matters more than it did at approval.

Studies of project outcomes report value losses alongside delivery problems, though none was built to test a dashboard. In the Project Management Institute's 2026 survey, 61 percent of project professionals reported some form of value loss on complex projects.[1] The report describes these as projects that deliver outputs while losing sight of outcomes. The figure is self-reported and concerns complexity. PMI's category figures count respondents who had any of several problems, so they do not show that value loss is more common than delivery disruption.

McKinsey and the University of Oxford compared predicted and actual budgets, schedules, and benefits for more than 5,400 IT projects. For large projects, those with initial price tags above \$15 million, the averages were 45 percent over budget, 7 percent over time, and 56 percent less value than predicted.[2] The three figures measure different things, a cost overrun, a schedule overrun, and a shortfall against predicted benefits, so the study does not rank them. It shows that value can be measured against the original prediction and that, for large IT projects, it fell well short on average. It does not say whether projects that stayed on budget and schedule delivered the value, so it cannot show that a project that is green on plan delivers little. In their survey-based discussion of causes, the authors write that IT initiatives too often manage projects purely to budget and schedule targets.[2] A page limited to those targets has no field where a value shortfall could appear.

Reporting tools lean toward status for structural reasons. Dates, spend, and issue counts come out of delivery tools automatically. Health takes a judgment about outcomes, a fresh look at assumptions, and a conversation with the receiving operation. That difference in cost is this paper's explanation for why status fills the page. No source here measures it. A rating that depends on judgment also depends on upkeep. The U.S. Government Accountability Office (GAO) checked 80 investments on the federal IT Dashboard against their supporting documents and found 53 of the chief information officers' risk ratings consistent, 20 partly consistent, and 7 inconsistent.[3] GAO named late updates and differences in how performance was tracked among the causes, and warned that one agency's yearly reset of its cost and schedule baselines raises the risk of undetected variances.[3] The sample is federal IT, and GAO does not say the ratings were shaded on purpose. It does show that a color on a dashboard is only as current and as consistent as the work behind it, and the lenses below have the same problem.

A second reason is the difference between lagging and leading measures. The Occupational Safety and Health Administration (OSHA) uses the terms for workplace safety: leading indicators are proactive, preventive measures that show whether activities are working, and lagging indicators count events that already happened. OSHA says a good program uses leading indicators to drive change and lagging indicators to measure effectiveness.[4] This paper borrows the terms for investments. Spend to date and milestones met are records of the past, so they are lagging. By the time they move, the options may have narrowed. The six lenses are meant to be leading measures, signals that move while a leader can still change the investment. Whether they lead is the paper's claim to test, not a finding.

If leaders act on the signal in front of them, their attention goes where the page shows trouble. That is likely to be the projects that are red on plan, which are often

merely late, while projects that are green on plan and buying little get none. This is the paper's inference about how a forum behaves. No source here measures it.

The six-lens investment health view

The view puts six questions on one page per investment. Each lens earns its place by changing a decision a leader can still make: continue, reshape, or stop; how much further to commit; where to put scarce people; how to roll the change out; when to trigger a fallback; and whether a choice is made now or passed along. A question that changes none of those belongs in the reference material. The same test applies to the set over time. A seventh lens joins when it changes a decision the six cannot, and a lens that has changed no decision in two review cycles is dropped, as a field is.

EXHIBIT 1

Five evidence lenses feed one decision queue

- 
Outcome health
 Remaining opportunity and whether the path from delivery to value still holds.
- 
Confidence health
 Premises rising, stable, falling, or still untested as evidence lands.
- 
Capacity health
 Named scarce roles, overload periods, substitutions, and displaced work.
- 
Adoption health
 Observable readiness in the receiving operation before go-live.
- 
Dependency health
 Owner, requirement, latest useful date, confidence, and consequence.

The six lenses are a checklist by design. Each is a separate question, an investment can fail on any one of them, and none has to be answered before another, so they can be read in any order. The sixth differs in one way: it collects the decisions the other five raise, so it is where their answers are acted on. The sections around the lenses do depend on each other. The case that a status deck cannot show whether

an investment still deserves its funding comes first, because the lenses only earn their place against what that deck cannot say.

A lens left blank is not green. Show it as blank and put it on the decision queue with an owner and a date. A blank in any lens holds that investment's next funding step until the owner fills it or the forum records why it proceeds without it. A blank outcome line at a renewal is the plainest case, because the review cannot say what the money is buying. A blank differs from an untested premise: untested is a state someone has chosen to report, and blank is a missing report. A weak answer is different again. It is information, and it goes to the decision the lens supports.

Outcome health

The first lens states what operating result the money is buying and whether the path from delivery to that result is still believable. The result is the change in how the business runs, such as cycle time down, errors down, capacity released, or exposure closed, not the list of things the project hands over. Beside it sit the assumption most likely to break that path and the latest evidence that the path is holding or fraying.

Writing the outcome line can force a negotiation some investments never had. Suppose the business case said cost avoidance, the sponsor meant capability, and finance booked headcount. An investment whose outcome cannot be stated in one operating sentence may be carrying an unresolved disagreement rather than a reporting gap.

"Benefits Realization Is Where Transformations Quietly Fail" covers naming the benefit owner, the measure, and the baseline at approval. This lens reads them while the investment is still running: the trend on the operating measure, taken from the receiving operation, against its baseline; the benefit owner's current estimate of the remaining opportunity; and what changed since approval. Here delivery progress stops standing in for value progress. A build can be eighty percent complete while the need has changed, for example because the volumes that justified the automation fell. Completion percentage says nothing about that.

The decision it supports is the CFO's question from the opening: continue, reshape, or stop. Judge it on the remaining opportunity against the remaining cost, including the cost of stopping, such as contract exit terms and work that cannot be reused. Never judge it on the money already spent.

Confidence health

The second lens tracks whether confidence in the investment's premises is rising or falling as evidence lands. Confidence is a statement about tested assumptions, not optimism and not the status color: which premises have been examined since the last review, what the tests showed, and what remains untested.

Each premise gets one of four states: rising, stable, falling, or untested. Four states instead of probabilities is a design choice. A probability asks for precision the evidence may not support, and a state is simple enough to keep current. The paper has not tested it against the alternative.

The reason to track movement is that the approval-time case is weak evidence by itself. Bent Flyvbjerg, whose evidence comes mostly from large transportation projects, reports that front-end estimates of costs and benefits typically differ significantly from actual results, so they predict value and viability poorly. He proposes checking a proposal against the record of similar projects.[5] A benefits forecast quoted to the same decimal through six review cycles has not been tested again. Real confidence moves: up when an integration proves out or a pilot converts, down when a vendor slips or pilot numbers come back soft.

The format is a short assumptions register: usually three to five premises the case leans on, the evidence that last touched each, the direction of movement, and the next scheduled test. "Give Stage Gates Teeth by Shrinking What They Decide" covers the facts an approval rests on and the triggers that send it back to the board. The register records how those premises are holding between approvals. The decision this lens supports is sizing: expand the commitment where confidence has been earned, and bound the next round of funding where it has not.

Capacity health

The third lens asks whether the people the investment assumes actually exist for it. An investment is unhealthy when it depends on scarce roles that are not available, even if its schedule still looks recoverable, because the recovery plan uses the same hours as the original plan. "The EPMO's Job Is Capacity, Not Calendars" covers how to find the scarce roles and keep one view of demand against supply. This lens reads that view for one investment: named allocations for the constrained roles, overload periods, substitutions in effect, and the other work this investment's staffing displaces.

Overload may show up first where a status report does not look, such as reviews done in less depth or testing squeezed, before it shows as slower delivery or more defects. That is this paper's expectation, not a measured sequence. The decision this lens supports is a portfolio decision, not a team one: sequence the work, narrow it, or release capacity on purpose, instead of leaving the constrained people to absorb the overload.

Adoption and operating health

The fourth lens asks whether the receiving operation is becoming able to use, support, and sustain what is being delivered. Delivery is incomplete until it is, and the paper treats the evidence as something to collect before go-live, not after: process owners engaged in design, supervisors trained before launch, a staffed support model, controls assigned to named owners, and early users giving real feedback. "A Project Is Not Ready Because a Plan Exists" covers getting the right people named, with their time confirmed, before kickoff. This lens checks, through delivery, that they did what they committed to.

A likely failure is deferral. Training, process ownership, and support readiness get scheduled for closer to launch while the dashboard tracks build milestones, and launch arrives at an operation that first engaged with the change the week it landed. The go-live can succeed technically and fail operationally. That is the paper's reasoning, not a measured rate.

Put leading indicators on the page from the first month: participation in design sessions, process-owner acceptance, support-readiness checkpoints, pilot adoption, and what pilot users actually did. The paper prefers what the receiving operation did to what it said. A process owner who changed a standing procedure, pilot users who came back in week three without being chased, and an operation that staffed its own support rota are harder to produce for the dashboard's sake than a survey score or an attendance count. That is a judgment. No source here compares behavior with survey results as predictors of adoption, and a survey can still catch a problem early. The decision this lens supports is how to roll out: slow it, stage it, or reinforce the receiving side while the plan is still cheap to change.

Dependency and risk health

The fifth lens asks whether the outside commitments this investment depends on will hold when it needs them. It treats each dependency as a commitment with an owner and a date, not a list of outside concerns. Every material dependency appears with the same five facts: who owns it, what must be in place, the latest useful date, current confidence, and the consequence if it fails. The latest useful date is the last point at which the commitment can still change the work. "Decision Friction Is a Portfolio Risk" uses the same term for decisions.

Consider a platform team that misses three integration windows. Each slip is logged as an issue, and the investment's completion date does not move. The dependency lives in a risk log, the forecast lives in a schedule, and no page puts them side by side. This lens does.

Beside the dependencies sit the conventional risks, each with its contingency option, what that option costs, and the date by which triggering it is still possible. The decision this lens supports is timing: invoke a contingency or re-sequence while the alternative still exists, instead of after the dependency's failure has chosen for everyone.

Decision readiness

The sixth lens collects the decisions the other five raise: what choice does this investment need, from whom, by when, and what happens if it waits? Pages of accurate description that end with no decision request are the sign of a dashboard that only mirrors the work.

Each open decision on the page is an entry from the decision inventory in "Decision Friction Is a Portfolio Risk," which sets out the fields each entry carries and how to fill them. This paper adds only where the entries sit: a visible queue on each investment's page, with a record below it of decisions made, deferred, escalated, or allowed to expire, so the forum can see its own performance as a deciding body.

The paper expects that a forum that opens with the queue spends its time deciding, while a forum that opens with status spends it being briefed and runs out of clock when the decisions arrive. Once latest useful dates are recorded, some forums may find their real constraint is throughput, not information. The inventory paper covers the queue measures that show this, and "The EPMO's Job Is Capacity, Not Calendars" counts decision capacity alongside people.

How the lenses read together

Read the lenses as a set, because their disagreements carry the information. Strong outcome and weak adoption is a rollout problem with a known fix. Strong delivery and falling confidence is a case for narrowing the next round of funding. Weak capacity under strong everything else is a collision with other work that will show up later as delivery problems. A single blended score would average these into silence, which is why the view computes none.

The set should also steer intervention. A leader reading one signal, usually schedule, is likely to intervene on schedule, which tends to mean re-planning. A leader reading six lenses can intervene on the weak lens, with the benefit owner, the capacity owner, the platform team, or the forum itself. Both are the paper's expectations, and the signals below are meant to check them.

One page per investment is a constraint worth defending. An investment that needs three pages is back to being a report. Everything below the six lenses is reference material: available on request, absent from the meeting.

One investment, read through the view

Consider a customer-onboarding automation, fourteen months in, seventy percent delivered, slightly under budget. The figures are invented to show how the page reads. The conventional dashboard shows green with a minor schedule note. The health view shows more.

Outcome: the case promised a forty percent cut in onboarding cycle time. Volumes in the receiving unit have fallen by a quarter since approval, and the benefit owner's refreshed estimate cuts the opportunity roughly in half. Confidence: the integration premise proved out cleanly, but the pilot's completion rates came back soft, and the adoption premise has gone untested for two quarters. Capacity: the constrained data engineer is committed to incident remediation through next quarter, and her hours are covered by a contractor who has never touched the legacy schema.

Adoption: process-owner acceptance is signed, but supervisor training has slipped twice and the support model is unstaffed. Dependency: the identity-platform upgrade this rollout requires has slipped once, with the next window in eleven weeks and no contingency chosen. Decision readiness: one item, open seven weeks, on whether to narrow the rollout to the two highest-volume segments, with a recommendation attached and a latest useful date in twenty days.

Read as a set, the page says the investment can still be delivered and the case has shrunk. It points to the narrowing decision now, a staged rollout matched to the platform window, and a renegotiated benefit line. None of that shows in green. All of it could be decided in one meeting, from one page.

Where the model stops

This section comes before the conversion steps because it decides how far to take them.

The lenses are one author's design. The sources behind this paper support the problem and some of the parts, such as measuring value separately from cost and treating ratings as dependent on upkeep. They do not support this particular set of six, the four-state confidence scale, or the one-page limit. A reader whose investments are governed differently should keep the lenses that change a decision and drop the rest.

Reading a few measures together is not new. The balanced scorecard is the best-known earlier example, and this paper does not claim to improve on it. The differences are the unit, one investment and its next decision, and the test each lens must pass.

Measures must fit the investment's mechanism and consequence, so no six lenses form a universal set of key performance indicators. The standard is decision usefulness: each indicator should reveal something leaders can influence. For the same reason the view does not rank investments. A ranking from six lenses would need weights, and the weights would be a judgment this paper does not try to settle. Comparing requests for funding is the job of the funding page.

Any judgment field can be performed. Confidence entries can drift optimistic, and adoption checkpoints can be marked complete by the team that owns the milestone. The GAO findings above show that ratings fall out of date even without intent. The countermeasures are ordinary but must exist. The PMO spot-checks lens entries against observable events such as a signed acceptance, a staffing record, or a test result, and the forum treats a lens that never moved before a surprise as a question for the post-incident review.

The page costs time. As a planning estimate, not a measured cost, the first page for a large investment takes a PMO analyst and the owners of the lenses a few hours, and later monthly refreshes take less because four of the six lenses are short judgments, not data pulls. The cost is justified only if the page replaces that investment's status

slides. If it is added on top of them, the deck grows and the forum has more to read, not less.

The view does not replace financial reporting, program management, or delivery tooling. It sits above them. It will not survive a forum that does not want to decide: an executive team determined to treat portfolio review as a briefing will defeat any instrument. The six lenses make that choice visible, which is as much as an instrument can do.

Converting the deck you already have

The conversion is a pruning exercise before it is a building exercise:

- Retire any field that has not changed a decision in two review cycles.
- Put outcome, confidence, capacity, adoption, dependency, and decision readiness on one page per investment.
- Allow narrative only where it explains a change, an uncertainty, or a requested choice.
- Track actions from the dashboard and test whether the forum resolved them before their latest useful dates.

The retirement rule does the heavy lifting. A field can persist because of a past incident, a reporting requirement nobody remembers, or a leader who has left, and each costs collection effort and attention. Count what you retire, and ask the people who received the old deck whether they noticed. This paper has no figure for how many fields that removes, and the number will vary by dashboard. A field required by regulation or contract moves to the reference material, off the page.

Start with the largest investments, where the gap between health and status costs most. An investment is whatever your forum funds or approves: a project, a program, or a product. When one cannot be separated, such as a shared platform serving several programs, give the platform its own page and let each program's dependency lens point to it. Smaller investments stay on the current status report until the first pages have changed a decision. The first cycle is likely to be rough: outcome statements may need renegotiation, assumptions registers may be thin, and decision queues may show items open for months. That roughness shows what was never settled.

Expect resistance in places. Benefit owners may resist refreshed estimates because a smaller case can read as their failure; a case revised early is stewardship, and one

defended to the end is the failure. Project managers may resist confidence entries that move down if the dashboard has been treated as a performance document. The forum can visibly reward the first team that downgrades its own confidence and gets a faster decision because of it. The paper expects the culture of the page to be set by the first few such moments, not by its template. That is its judgment, not a finding.

Data plumbing matters less than it appears. Four of the six lenses run on judgments refreshed monthly: the assumptions register, the benefit owner's estimate, the adoption checkpoints, and the decision queue. Capacity and dependency can start from the EPMO's registers where they exist. Waiting for a data warehouse delays the first page.

The PMO curates the page and does not adjudicate. Its jobs are the honesty of the page, such as challenging a confidence entry that never moves or a dependency marked green by hope, and the quality of the queue. The judgments inside the lenses belong to benefit owners, capacity owners, and sponsors. The decisions belong to the forum.

Signals the instrument is working

The page's own health is measurable:

EXHIBIT 2

How to tell the page is steering

Decisions from the page

Made before their latest useful dates.

Targeted interventions

Aimed at the weak lens, not generic schedule pressure.

Confidence that moves

Entries changing during the quarter, in either direction.

A shrinking deck

Fields retired over time, not added back until the deck is long again.

- Decisions made from the page, and whether they beat their latest useful dates.
- Interventions that targeted the weak lens, versus generic schedule pressure.
- Confidence entries that moved during the quarter, in either direction.

- Adoption signals visible before go-live on every investment past design.
- Dependency slips that changed a forecast within one cycle, and fields retired versus fields added over time.

Because no source shows that a page like this improves decisions, test it on the first few investments. Record when in each project reshaping or stopping decisions were made, and the share of decisions made before their latest useful dates, and compare them with investments still on the old deck. If a lens changes nothing, drop it, as the rule above directs.

The last measure guards the page's future. The paper expects dashboards to accrete fields: every incident proposes one, and none volunteers to leave. Net field count over a year shows whether the page is surviving the organization's habits.

The timing completes the design. The page refreshes monthly, the forum meets monthly, and the decision queue is the agenda. An investment whose page shows no weak lens and no open decision gets thirty seconds and a consent motion. The forum's clock belongs to weak lenses and due decisions, in order of the cost of waiting.

The instrument, not the mirror

A dashboard earns attention when it changes an investment while options remain. The test of the page is the CFO's question: if we stopped this today, what would we lose? A health view answers from the page with the remaining opportunity, the confidence earned, the capacity trapped, the adoption forming, the dependencies surviving, and the decision now due.

The paper expects a forum that adopts the page to spend more of its time deciding and less being briefed, and to reshape or stop investments earlier, while the narrowing is still a choice and the capacity still has somewhere to go. No source tests that expectation, which is why the signals above exist.

Notes and sources

[1] Project Management Institute, "Pulse of the Profession® 2026: Driving Success in Complex Projects," May 2026. <https://www.pmi.org/learning/thought-leadership/driving-success-in-complex-projects>

- [2] Michael Bloch, Sven Blumberg, and Jürgen Laartz, "Delivering large-scale IT projects on time, on budget, and on value," McKinsey & Company, October 1, 2012. <https://www.mckinsey.com/capabilities/tech-and-ai/our-insights/delivering-large-scale-it-projects-on-time-on-budget-and-on-value>
- [3] U.S. Government Accountability Office, "IT Dashboard: Agencies Are Managing Investment Risk, but Related Ratings Need to Be More Accurate and Available," GAO-14-64, December 12, 2013. <https://www.gao.gov/products/gao-14-64>
- [4] Occupational Safety and Health Administration, "Using Leading Indicators to Improve Safety and Health Outcomes," U.S. Department of Labor. <https://www.osha.gov/leading-indicators>
- [5] Bent Flyvbjerg, "Quality Control and Due Diligence in Project Management: Getting Decisions Right by Taking the Outside View," 2013. <https://arxiv.org/abs/1302.2544>

About the author

Marco Policani is an enterprise portfolio, PMO, and AI operating-governance leader. He builds portfolio governance systems where AI carries the analytical load and named humans own every decision — an approach documented across case studies, walkthroughs, and working governance guides on his portfolio site.

Portfolio & governance library: policani.net/governance · Full portfolio: policani.net · LinkedIn: linkedin.com/in/marcpolicani

This white paper may be shared freely with attribution. © 2026 Marco Policani.