
WHITE PAPER · PORTFOLIO & DELIVERY GOVERNANCE

A PMO Dashboard Should Show Investment Health

A decision interface, not a longer status deck

Marco Policani

Enterprise Portfolio · PMO · AI Operating Governance

policani.net · linkedin.com/in/marcpolicani

July 2026

EXECUTIVE SUMMARY

A portfolio dashboard should make the next investment decision easier, not make activity look more complete.

A dashboard earns attention when it helps leaders alter an investment while options remain. Schedule and spend matter, but neither can show whether value is still plausible, capacity is credible, adoption is forming, dependencies are survivable, or a decision is overdue. This paper gives portfolio executives, sponsors, PMO leaders, and investment owners a practical the six-lens investment health view for making the next commitment explicit.

- 01** Outcome health: The dashboard should state the operating result being purchased and whether the causal path from outputs to that result remains credible.
- 02** Confidence health: Confidence should rise or fall as premises are tested; it is different from optimistic status commentary.
- 03** Capacity health: An investment is unhealthy when it depends on unavailable scarce roles even if its schedule still appears recoverable.
- 04** Adoption and operating health: Delivery is incomplete until the receiving operation can use, support, and sustain the change.

CONTENTS

The argument	Dependency and risk health
Why the conventional approach breaks	Decision readiness
The six-lens investment health view	How the elements interact
Outcome health	Put the model into the management cadence
Confidence health	Measures that reveal whether the control works
Capacity health	The strongest counterargument
Adoption and operating health	The shift

The argument

A dashboard earns attention when it helps leaders alter an investment while options remain. Schedule and spend matter, but neither can show whether value is still plausible, capacity is credible, adoption is forming, dependencies are survivable, or a decision is overdue.

Most portfolio dashboards are accurate descriptions of delivery activity and weak instruments for investment governance. They show milestones, budget variance, issue counts, and status colors. Leaders can read every page and still be unable to answer the questions that justify continued funding: Is the expected outcome still worth pursuing? Has confidence improved? What scarce capacity is trapped? Which decision would materially change the result?

This paper is written for portfolio executives, sponsors, PMO leaders, and investment owners. It offers the six-lens investment health view as a management instrument: a way to connect operating evidence to the next consequential choice. The cited sources provide external context; the framework and its application are Marco Policani's synthesis. It is not a predictive score and does not replace professional, legal, security, financial, or domain judgment.

Why the conventional approach breaks

This is where a management ritual either becomes a real control or reveals itself as administrative theater. Conventional governance tends to separate planning, approval, delivery reporting, and operating ownership. That separation allows a premise to pass through several forums without any one forum owning the decision it implies. The project or workflow then carries an assumption that was acknowledged socially but never accepted operationally.

A polished status can coexist with this failure because status describes activity; it does not prove the premise that authorizes reliance. The answer is not heavier control at every step. It is to place a clear, proportionate test at the point where the organization increases money, capacity, access, dependency, or action authority.

The six-lens investment health view

The model has 6 connected elements. Each makes a different condition visible, but they should be reviewed together because strength in one area does not cancel a missing obligation in another. The model is designed for a short executive conversation supported by inspectable evidence, not a long maturity assessment.

EXHIBIT 1

The six-lens investment health view connects evidence to action

- 1 **Outcome health**
The dashboard should state the operating result being purchased and whether the causal path from outputs to that result remains credible.
- 2 **Confidence health**
Confidence should rise or fall as premises are tested; it is different from optimistic status commentary.
- 3 **Capacity health**
An investment is unhealthy when it depends on unavailable scarce roles even if its schedule still appears recoverable.
- 4 **Adoption and operating health**
Delivery is incomplete until the receiving operation can use, support, and sustain the change.

5

Dependency and risk health

Dependencies deserve treatment as commitments with owners and dates, not a list of external concerns.

6

Decision readiness

The dashboard should make the next consequential choice unmistakable.

Outcome health

The dashboard should state the operating result being purchased and whether the causal path from outputs to that result remains credible. Viewed from the portfolio, this is a resource-allocation problem before it becomes a delivery problem.

Delivery progress is green while the original need, user behavior, policy, or business process has changed. Reporting usually captures the symptom after it has been translated into schedule, cost, quality, or risk. It rarely captures the original unresolved choice.

Show the baseline, intended outcome, leading evidence, benefit owner, and the assumption most likely to break the value path. Good governance preserves options. It creates a legitimate route to narrow, stage, redirect, pause, or stop without labeling learning as failure.

Trend evidence from the receiving operation alongside completion percentages from the project plan. A lack of evidence is itself useful information if the next decision allows a smaller experiment, tighter boundary, or deliberate wait.

Continue, reshape, or stop according to the remaining outcome opportunity rather than the amount already spent. The response should occur while it can still change the outcome, not after the issue has been translated into an unrecoverable variance.

A tabletop review of outcome health should use the observed break as its scenario: Delivery progress is green while the original need, user behavior, policy, or business process has changed. Participants then apply the proposed response. They show the baseline, intended outcome, leading evidence, benefit owner, and the assumption most likely to break the value path. The review identifies where authority, information, time, or recovery would run out. It is useful only if it changes the boundary or confirms who will act when the condition appears.

The minimum record for outcome health is anchored in actual proof. It includes trend evidence from the receiving operation alongside completion percentages from the project plan. The record also carries the choice supported by that proof: Continue, reshape, or stop according to the remaining outcome opportunity rather than the amount already spent. Keeping evidence and decision together lets a later owner distinguish a deliberate residual risk from a premise that nobody examined. It also gives the team a clear reason to reopen the choice when the underlying facts change.

Across the work governed by portfolio executives, sponsors, PMO leaders, and investment owners, occurrences of this pattern should be compared rather than treated as unrelated project stories. The positive standard is: The dashboard should state the operating result being purchased and whether the causal path from outputs to that result remains credible. If the same failure recurs, leaders can change delegation, capacity, intake, policy, or the intervention itself. That portfolio response is usually more valuable than asking each team to absorb another local workaround.

Two questions test outcome health. What evidence would show that the condition is no longer adequate? Would that evidence arrive early enough to preserve a feasible choice? The intended choice is clear: Continue, reshape, or stop according to the remaining outcome opportunity rather than the amount already spent. A weak answer to either question calls for a smaller commitment, an earlier signal, or a safer fallback. A strong answer earns movement without claiming that uncertainty has disappeared.

Confidence health

Confidence should rise or fall as premises are tested; it is different from optimistic status commentary. A useful governance design therefore begins with the consequence leaders are prepared to accept.

Forecasts remain precise despite unresolved technical, adoption, dependency, or operating questions. A local fix can keep the work moving while worsening the portfolio: another specialist is borrowed, another exception is tolerated, and another dependency is promised twice.

Name the material assumptions, the evidence obtained since the last review, and what remains unknown. The control is complete only when the response is known: who acts, how quickly, within what authority, and with which safe fallback.

Test results, ranges, confidence changes, contradictory observations, and the next evidence event. The receiving operation must recognize the evidence as relevant. Technical success alone cannot prove that a business workflow can carry the change.

Bound the next commitment when confidence has not earned expansion. That choice should be recorded in operational terms: what may proceed, what remains outside the boundary, who owns the next move, and what would reopen the decision.

A tabletop review of confidence health should use the observed break as its scenario: Forecasts remain precise despite unresolved technical, adoption, dependency, or operating questions. Participants then apply the proposed response. They name the material assumptions, the evidence obtained since the last review, and what remains unknown. The review identifies where authority, information, time, or recovery would run out. It is useful only if it changes the boundary or confirms who will act when the condition appears.

The minimum record for confidence health is anchored in actual proof. It includes test results, ranges, confidence changes, contradictory observations, and the next evidence event. The record also carries the choice supported by that proof: Bound the next commitment when confidence has not earned expansion. Keeping evidence and decision together lets a later owner distinguish a deliberate residual risk from a premise that nobody examined. It also gives the team a clear reason to reopen the choice when the underlying facts change.

Across the work governed by portfolio executives, sponsors, PMO leaders, and investment owners, occurrences of this pattern should be compared rather than treated as unrelated project stories. The positive standard is: Confidence should rise or fall as premises are tested; it is different from optimistic status commentary. If the same failure recurs, leaders can change delegation, capacity, intake, policy, or the intervention itself. That portfolio response is usually more valuable than asking each team to absorb another local workaround.

Two questions test confidence health. What evidence would show that the condition is no longer adequate? Would that evidence arrive early enough to preserve a feasible choice? The intended choice is clear: Bound the next commitment when confidence has not earned expansion. A weak answer to either question calls for a smaller

commitment, an earlier signal, or a safer fallback. A strong answer earns movement without claiming that uncertainty has disappeared.

Capacity health

An investment is unhealthy when it depends on unavailable scarce roles even if its schedule still appears recoverable. A strong operating model makes the hidden negotiation explicit while options are still available.

Teams compensate for missing capacity by multitasking critical people, accepting quality debt, or shifting work into later phases. A polished status can coexist with this failure because status describes activity; it does not prove the premise that authorizes reliance.

Report constrained roles, decision load, dependency capacity, and the impact on other investments. The aim is not certainty. It is a justified next commitment with an observable basis and a viable way back.

Named allocations, overload periods, vacancies, substitution limits, and work displaced by the current commitment. Measures should expose changes in the operation rather than reward production of artifacts. Activity can support the story; it cannot substitute for the result.

Sequence, narrow, or release capacity instead of treating overload as a delivery-team problem. When the premise is weak, narrowing is often a better decision than forcing a binary choice between full approval and cancellation.

A tabletop review of capacity health should use the observed break as its scenario: Teams compensate for missing capacity by multitasking critical people, accepting quality debt, or shifting work into later phases. Participants then apply the proposed response. They report constrained roles, decision load, dependency capacity, and the impact on other investments. The review identifies where authority, information, time, or recovery would run out. It is useful only if it changes the boundary or confirms who will act when the condition appears.

The minimum record for capacity health is anchored in actual proof. It includes named allocations, overload periods, vacancies, substitution limits, and work displaced by the current commitment. The record also carries the choice supported by that proof: Sequence, narrow, or release capacity instead of treating overload as a delivery-team problem. Keeping evidence and decision together lets a later owner distinguish a deliberate residual risk from a premise that nobody examined. It also gives the team a clear reason to reopen the choice when the underlying facts change.

Across the work governed by portfolio executives, sponsors, PMO leaders, and investment owners, occurrences of this pattern should be compared rather than treated as unrelated project stories. The positive standard is: An investment is unhealthy when it depends on unavailable scarce roles even if its schedule still appears recoverable. If the same failure recurs, leaders can change delegation, capacity, intake, policy, or the intervention itself. That portfolio response is usually more valuable than asking each team to absorb another local workaround.

Two questions test capacity health. What evidence would show that the condition is no longer adequate? Would that evidence arrive early enough to preserve a feasible choice? The intended choice is clear: Sequence, narrow, or release capacity instead of treating overload as a delivery-team problem. A weak answer to either question calls for a smaller commitment, an earlier signal, or a safer fallback. A strong answer earns movement without claiming that uncertainty has disappeared.

Adoption and operating health

Delivery is incomplete until the receiving operation can use, support, and sustain the change. Teams move faster when uncertainty is bounded honestly than when it is concealed inside an optimistic plan.

Training, process ownership, support, controls, and local behavior are deferred until launch while the dashboard stays focused on build milestones. Under pressure, the workaround is predictable: capable people compensate, local decisions proliferate, and the formal record stops describing how the operation actually runs.

Add adoption signals, operational readiness, support load, control ownership, and transition risks early. The design should state what remains unknown. Acknowledged uncertainty can be bounded; disguised uncertainty is inherited by the operation.

Representative user evidence, process-owner acceptance, support readiness, exception volumes, and transition commitments. Where consequence is high, independent checking or cross-functional challenge reduces the risk that advocacy becomes its own proof.

Change the rollout or investment shape when operating absorption is weaker than the delivery plan assumes. Escalation is appropriate when authority is missing, not as a routine substitute for clear delegation.

A tabletop review of adoption and operating health should use the observed break as its scenario: Training, process ownership, support, controls, and local behavior are deferred until launch while the dashboard stays focused on build milestones. Participants then apply the proposed response. They add adoption signals, operational readiness, support load, control ownership, and transition risks early. The review identifies where authority, information, time, or recovery would run out. It is useful only if it changes the boundary or confirms who will act when the condition appears.

The minimum record for adoption and operating health is anchored in actual proof. It includes representative user evidence, process-owner acceptance, support readiness, exception volumes, and transition commitments. The record also carries the choice supported by that proof: Change the rollout or investment shape when operating absorption is weaker than the delivery plan assumes. Keeping evidence and decision together lets a later owner distinguish a deliberate residual risk from a premise that nobody examined. It also gives the team a clear reason to reopen the choice when the underlying facts change.

Across the work governed by portfolio executives, sponsors, PMO leaders, and investment owners, occurrences of this pattern should be compared rather than treated as unrelated project stories. The positive standard is: Delivery is incomplete until the receiving operation can use, support, and sustain the change. If the same failure recurs, leaders can change delegation, capacity, intake, policy, or the intervention itself. That portfolio response is usually more valuable than asking each team to absorb another local workaround.

Two questions test adoption and operating health. What evidence would show that the condition is no longer adequate? Would that evidence arrive early enough to preserve a feasible choice? The intended choice is clear: Change the rollout or investment shape when operating absorption is weaker than the delivery plan assumes. A weak answer to either question calls for a smaller commitment, an earlier signal, or a safer fallback. A strong answer earns movement without claiming that uncertainty has disappeared.

Dependency and risk health

Dependencies deserve treatment as commitments with owners and dates, not a list of external concerns. The practical failure occurs between the formal approval and the ordinary pressure of running the work.

A shared platform, vendor, policy decision, or upstream program slips repeatedly without changing the investment forecast. The happy path can survive this weakness. Ordinary variation exposes it: a disputed input, absent owner, competing commitment, or case that falls outside the assumed rule.

Show dependency owner, required condition, latest useful date, confidence, and consequence of failure beside conventional risks. The smallest credible control is the one that changes behavior when the evidence is weak. Anything else is commentary.

Accepted commitments, interface evidence, contingency options, and the decision date before delay becomes irreversible. Qualitative observation belongs beside quantitative measures when workflow behavior, adoption, or judgment explains why the number moved.

Trigger a contingency or re-sequence when dependency confidence falls below what the next commitment needs. A stop or rollback is evidence that the control works when it protects greater value elsewhere in the portfolio.

A tabletop review of dependency and risk health should use the observed break as its scenario: A shared platform, vendor, policy decision, or upstream program slips repeatedly without changing the investment forecast. Participants then apply the proposed response. They show dependency owner, required condition, latest useful date, confidence, and consequence of failure beside conventional risks. The review identifies where authority, information, time, or recovery would run out. It is useful only if it changes the boundary or confirms who will act when the condition appears.

The minimum record for dependency and risk health is anchored in actual proof. It includes accepted commitments, interface evidence, contingency options, and the decision date before delay becomes irreversible. The record also carries the choice supported by that proof: Trigger a contingency or re-sequence when dependency confidence falls below what the next commitment needs. Keeping evidence and decision together lets a later owner distinguish a deliberate residual risk from a premise that nobody examined. It also gives the team a clear reason to reopen the choice when the underlying facts change.

Across the work governed by portfolio executives, sponsors, PMO leaders, and investment owners, occurrences of this pattern should be compared rather than treated as unrelated project stories. The positive standard is: Dependencies deserve treatment as commitments with owners and dates, not a list of external concerns. If the same failure recurs, leaders can change delegation, capacity, intake, policy, or the intervention itself. That portfolio response is usually more valuable than asking each team to absorb another local workaround.

Two questions test dependency and risk health. What evidence would show that the condition is no longer adequate? Would that evidence arrive early enough to preserve a feasible choice? The intended choice is clear: Trigger a contingency or re-sequence when dependency confidence falls below what the next commitment needs. A weak answer to either question calls for a smaller commitment, an earlier signal, or a safer fallback. A strong answer earns movement without claiming that uncertainty has disappeared.

Decision readiness

The dashboard should make the next consequential choice unmistakable. The visible artifact is often complete before the operating condition behind it is credible.

Pages of status end with no decision request, so uncertainty ages between forums and teams improvise. The organization may call this execution risk, yet the failure began when authority and evidence were separated at the decision point.

Name the decision, owner, evidence, options, recommendation, latest useful date, and consequence of waiting. A useful test asks what ordinary operating pressure would break the premise and how that break would be detected before consequence spreads.

A visible decision queue and a record of choices made, deferred, escalated, or allowed to expire. Evidence should arrive before the latest useful decision date; perfect analysis delivered after options expire has little governance value.

Use the review to change the portfolio, not simply acknowledge its condition. The operating owner should confirm that the decision can be executed with the people, systems, and time actually available.

A tabletop review of decision readiness should use the observed break as its scenario: Pages of status end with no decision request, so uncertainty ages between forums and teams improvise. Participants then apply the proposed response. They name the decision, owner, evidence, options, recommendation, latest useful date, and consequence of waiting. The review identifies where authority, information, time, or recovery would run out. It is useful only if it changes the boundary or confirms who will act when the condition appears.

The minimum record for decision readiness is anchored in actual proof. It includes a visible decision queue and a record of choices made, deferred, escalated, or allowed to expire. The record also carries the choice supported by that proof: Use the review to change the portfolio, not simply acknowledge its condition. Keeping evidence and decision together lets a later owner distinguish a deliberate residual risk from a premise that nobody examined. It also gives the team a clear reason to reopen the choice when the underlying facts change.

Across the work governed by portfolio executives, sponsors, PMO leaders, and investment owners, occurrences of this pattern should be compared rather than treated as unrelated project stories. The positive standard is: The dashboard should make the next consequential choice unmistakable. If the same failure recurs, leaders can change delegation, capacity, intake, policy, or the intervention itself. That portfolio response is usually more valuable than asking each team to absorb another local workaround.

Two questions test decision readiness. What evidence would show that the condition is no longer adequate? Would that evidence arrive early enough to preserve a feasible choice? The intended choice is clear: Use the review to change the portfolio, not simply acknowledge its condition. A weak answer to either question calls for a smaller commitment, an earlier signal, or a safer fallback. A strong answer earns movement without claiming that uncertainty has disappeared.

How the elements interact

Outcome health and Confidence health protect different parts of the same commitment. The dashboard should state the operating result being purchased and whether the causal path from outputs to that result remains credible. Confidence should rise or fall as premises are tested; it is different from optimistic status commentary. A strong operating model makes the hidden negotiation explicit while options are still available. If the operation encounters this failure—delivery progress is green while the original need, user behavior, policy, or business process has changed.—the response for confidence health cannot compensate for the missing outcome health obligation. Reviewers should reconcile the evidence for both elements before they accept the next action: bound the next commitment when confidence has not earned expansion.

Confidence health and Capacity health protect different parts of the same commitment. Confidence should rise or fall as premises are tested; it is different from optimistic status commentary. An investment is unhealthy when it depends on unavailable scarce roles even if its schedule still appears recoverable. The work becomes governable when ownership, evidence, boundaries, and response are connected to a real choice. If the operation encounters this failure—forecasts remain precise despite unresolved technical, adoption, dependency, or operating questions.—the response for capacity health cannot compensate for the missing confidence health obligation. Reviewers should reconcile the evidence for both elements before they accept the next action: sequence, narrow, or release capacity instead of treating overload as a delivery-team problem.

Capacity health and Adoption and operating health protect different parts of the same commitment. An investment is unhealthy when it depends on unavailable scarce roles even if its schedule still appears recoverable. Delivery is incomplete until the receiving operation can use, support, and sustain the change. The cost appears downstream, but the missing decision sits upstream. If the operation encounters this failure—teams compensate for missing capacity by multitasking critical people, accepting quality debt, or shifting work into later phases.—the response for adoption and operating health cannot compensate for the missing capacity health obligation. Reviewers should reconcile the evidence for both elements before they accept the next action: change the rollout or investment shape when operating absorption is weaker than the delivery plan assumes.

Adoption and operating health and Dependency and risk health protect different parts of the same commitment. Delivery is incomplete until the receiving operation can use, support, and sustain the change. Dependencies deserve treatment as commitments with owners and dates, not a list of external concerns. Teams move faster when uncertainty is bounded honestly than when it is concealed inside an optimistic plan. If the operation encounters this failure—training, process ownership, support, controls, and local behavior are deferred until launch while the dashboard stays focused on build milestones.—the response for dependency and risk health cannot compensate for the missing adoption and operating health obligation. Reviewers should reconcile the evidence for both elements before they accept the next action: trigger a contingency or re-sequence when dependency confidence falls below what the next commitment needs.

Dependency and risk health and Decision readiness protect different parts of the same commitment. Dependencies deserve treatment as commitments with owners and dates, not a list of external concerns. The dashboard should make the next consequential choice unmistakable. The framework in this paper is designed to make that boundary usable in ordinary management cadence. If the operation encounters this failure—a shared platform, vendor, policy decision, or upstream program slips repeatedly without changing the investment forecast.—the response for decision readiness cannot compensate for the missing dependency and risk health obligation. Reviewers should reconcile the evidence for both elements before they accept the next action: use the review to change the portfolio, not simply acknowledge its condition.

Decision readiness and Outcome health protect different parts of the same commitment. The dashboard should make the next consequential choice unmistakable. The dashboard should state the operating result being purchased and whether the causal path from outputs to that result remains credible. The distinction matters because leaders authorize commitments, while teams inherit whatever the authorization left unresolved. If the operation encounters this failure—pages of status end with no decision request, so uncertainty ages between forums and teams improvise.—the response for outcome health cannot compensate for the missing decision readiness obligation. Reviewers should reconcile the evidence for both elements before they accept the next action: continue, reshape, or stop according to the remaining outcome opportunity rather than the amount already spent.

Put the model into the management cadence

A framework becomes useful when it changes the normal cadence of work. It should shape intake, funding, design, delivery, and operating review without creating a separate governance industry around itself. The following sequence is deliberately small:

EXHIBIT 2

Start with one decision, then calibrate from evidence

- 1 Move 1**
Retire any field that has not changed a decision in two review cycles.
- 2 Move 2**
Put outcome, confidence, capacity, adoption, dependency, and decision readiness on one page per investment.
- 3 Move 3**
Allow narrative only where it explains a change, uncertainty, or requested choice.
- 4 Move 4**
Track actions from the dashboard and test whether the forum resolved them before their latest useful dates.

- Retire any field that has not changed a decision in two review cycles.
- Put outcome, confidence, capacity, adoption, dependency, and decision readiness on one page per investment.
- Allow narrative only where it explains a change, uncertainty, or requested choice.
- Track actions from the dashboard and test whether the forum resolved them before their latest useful dates.

The first cycle should be treated as calibration. Reviewers should note which questions changed a decision, which evidence was unavailable, where authority was unclear, and which controls cost out of proportion to the exposure. That record improves the six-lens investment health view without pretending the first version will fit every workflow or investment.

Ownership should remain close to the consequence. The PMO or AI-governance function can maintain the method, surface cross-portfolio patterns, and challenge weak evidence. It should not absorb the business owner's accountability for the result or the executive's accountability for the commitment.

Measures that reveal whether the control works

- Time from a material signal to a consequential decision.
- Commitments narrowed, redirected, paused, or stopped before avoidable exposure grew.
- Exceptions or assumptions discovered after the latest useful decision date.
- Scarce specialist and executive attention consumed by recurring unresolved conditions.
- Decisions reopened because the original evidence, boundary, or rationale was unclear.
- Operating outcomes and recovery performance after reliance expanded.

These measures are diagnostic, not a universal scorecard. Their purpose is to identify where the operating model fails repeatedly and whether governance is preserving options earlier. A lower count can indicate improvement, but it can also indicate weak detection. Leaders should read the measures with case evidence and frontline observation.

The strongest counterargument

No six lenses form a universal KPI set. Measures must fit the investment's mechanism and consequence. The standard is decision usefulness: each indicator should reveal a condition leaders can influence, not manufacture a comparable score where none exists.

A reasonable critic could also argue that experienced leaders already do this through judgment. Many do. The weakness is not the absence of judgment; it is the absence of a durable operating record when people, pressure, and conditions change. The six-lens investment health view makes the basis for judgment visible enough to challenge, execute, and revisit.

The shift

A dashboard earns attention when it helps leaders alter an investment while options remain. Schedule and spend matter, but neither can show whether value is still plausible, capacity is credible, adoption is forming, dependencies are survivable, or a decision is overdue. The practical shift is from reporting activity to governing the conditions under which the organization relies on that activity.

That discipline does not make uncertainty disappear. It gives leaders a better place to stand: a named decision, evidence proportionate to consequence, an explicit boundary, a responsible owner, and a response when reality departs from the assumption. Those elements are enough to turn hidden operating risk into a choice the portfolio can make deliberately.

Notes and sources

[1] Project Management Institute, "Pulse of the Profession 2026: Driving Success in Complex Projects," May 2026.
<https://www.pmi.org/learning/thought-leadership/driving-success-in-complex-projects>

[2] CIO, "IT Governance: Best Practices for Aligning IT and Business Strategy."
<https://www.cio.com/article/272051/it-governance-a-formal-way-to-align-it-and-business-strategy.html>

About the author

Marco Policani is an enterprise portfolio, PMO, and AI operating-governance leader. He builds portfolio governance systems where AI carries the analytical load and named humans own every decision — an approach documented across case studies, walkthroughs, and working governance guides on his portfolio site.

Portfolio & governance library: policani.net/governance · Full portfolio: policani.net · LinkedIn: linkedin.com/in/marcpolicani

This white paper may be shared freely with attribution. © 2026 Marco Policani.