
WHITE PAPER · AI OPERATING GOVERNANCE

Local AI Versus Cloud AI Is a Governance Decision

Choose the operating boundary before the model

Marco Policani

Enterprise Portfolio · PMO · AI Operating Governance

policani.net · linkedin.com/in/marcpolicani

July 2026

EXECUTIVE SUMMARY

The operating boundary should be chosen around data, access, action authority, resilience, and exit—not model fashion.

Local, cloud, hybrid, or no deployment should follow an operating-boundary decision about data, access, action authority, evaluation, resilience, support, suppliers, and exit. Model fashion is not a governance criterion. This paper gives CIOs, architecture leaders, security leaders, AI program owners, and workflow executives a practical the seven-boundary deployment decision for making the next commitment explicit.

- 01** Data boundary: Deployment must fit the sensitivity, residency, lineage, retention, and secondary-use rules of the actual workflow data.
- 02** Access and identity boundary: The decision includes who can reach models, data, tools, administrative functions, and support channels.
- 03** Action-authority boundary: Inference location does not decide what the system is permitted to do inside a workflow.
- 04** Evaluation and change boundary: Every option needs a way to evaluate behavior as models, prompts, data, and users change.

CONTENTS

The argument	Resilience and support boundary
Why the conventional approach breaks	Supplier and exit boundary
The seven-boundary deployment decision	How the elements interact
Data boundary	Put the model into the management cadence
Access and identity boundary	Measures that reveal whether the control works
Action-authority boundary	The strongest counterargument
Evaluation and change boundary	The shift

The argument

Local, cloud, hybrid, or no deployment should follow an operating-boundary decision about data, access, action authority, evaluation, resilience, support, suppliers, and exit. Model fashion is not a governance criterion.

Architecture debates often begin with model performance and end with governance treated as a constraint on the preferred answer. The order should be reversed. Leaders first need to decide where data may move, who can

access the service, what actions it may influence, how failure is contained, who operates it, and how the organization exits the arrangement.

NIST's generative-AI risk profile and CISA's cloud security architecture both reinforce the need to connect technology choices to managed risk, shared responsibility, and an inspectable operating design [1][2].

This paper is written for CIOs, architecture leaders, security leaders, AI program owners, and workflow executives. It offers the seven-boundary deployment decision as a management instrument: a way to connect operating evidence to the next consequential choice. The cited sources provide external context; the framework and its application are Marco Policani's synthesis. It is not a predictive score and does not replace professional, legal, security, financial, or domain judgment.

Why the conventional approach breaks

Teams move faster when uncertainty is bounded honestly than when it is concealed inside an optimistic plan. Conventional governance tends to separate planning, approval, delivery reporting, and operating ownership. That separation allows a premise to pass through several forums without any one forum owning the decision it implies. The project or workflow then carries an assumption that was acknowledged socially but never accepted operationally.

What looks like a team-level miss is frequently a portfolio choice that nobody made explicitly. Scarce attention is then spent reconciling consequences rather than choosing among alternatives. The answer is not heavier control at every step. It is to place a clear, proportionate test at the point where the organization increases money, capacity, access, dependency, or action authority.

The seven-boundary deployment decision

The model has 6 connected elements. Each makes a different condition visible, but they should be reviewed together because strength in one area does not cancel a missing obligation in another. The model is designed for a short executive conversation supported by inspectable evidence, not a long maturity assessment.

EXHIBIT 1

The seven-boundary deployment decision connects evidence to action

- 1 **Data boundary**
Deployment must fit the sensitivity, residency, lineage, retention, and secondary-use rules of the actual workflow data.
- 2 **Access and identity boundary**
The decision includes who can reach models, data, tools, administrative functions, and support channels.
- 3 **Action-authority boundary**
Inference location does not decide what the system is permitted to do inside a workflow.
- 4 **Evaluation and change boundary**
Every option needs a way to evaluate behavior as models, prompts, data, and users change.

5

Resilience and support boundary

Availability, latency, incident response, continuity, and recovery obligations differ by deployment pattern.

6

Supplier and exit boundary

A deployment choice creates dependencies on models, hardware, platforms, contracts, skills, and data formats.

Data boundary

Deployment must fit the sensitivity, residency, lineage, retention, and secondary-use rules of the actual workflow data. The practical failure occurs between the formal approval and the ordinary pressure of running the work.

Teams classify the model in abstract while prompts, retrieved content, logs, embeddings, and support artifacts cross different boundaries. The happy path can survive this weakness. Ordinary variation exposes it: a disputed input, absent owner, competing commitment, or case that falls outside the assumed rule.

Trace data through input, processing, retrieval, telemetry, administration, backup, and deletion for each option. The smallest credible control is the one that changes behavior when the evidence is weak. Anything else is commentary.

Data-flow diagrams, classifications, contractual terms, retention controls, access records, and verified deletion paths. Qualitative observation belongs beside quantitative measures when workflow behavior, adoption, or judgment explains why the number moved.

Choose a bounded or local pattern when required controls cannot be demonstrated in the cloud; choose cloud when its managed controls better fit the need. A stop or rollback is evidence that the control works when it protects greater value elsewhere in the portfolio.

A tabletop review of data boundary should use the observed break as its scenario: Teams classify the model in abstract while prompts, retrieved content, logs, embeddings, and support artifacts cross different boundaries. Participants then apply the proposed response. They trace data through input, processing, retrieval, telemetry, administration, backup, and deletion for each option. The review identifies where authority, information, time, or recovery would run out. It is useful only if it changes the boundary or confirms who will act when the condition appears.

The minimum record for data boundary is anchored in actual proof. It includes data-flow diagrams, classifications, contractual terms, retention controls, access records, and verified deletion paths. The record also carries the choice supported by that proof: Choose a bounded or local pattern when required controls cannot be demonstrated in the cloud; choose cloud when its managed controls better fit the need. Keeping evidence and decision together lets a later owner distinguish a deliberate residual risk from a premise that nobody examined. It also gives the team a clear reason to reopen the choice when the underlying facts change.

Across the work governed by CIOs, architecture leaders, security leaders, AI program owners, and workflow executives, occurrences of this pattern should be compared rather than treated as unrelated project stories. The positive standard is: Deployment must fit the sensitivity, residency, lineage, retention, and secondary-use rules of the actual workflow data. If the same failure recurs, leaders can change delegation, capacity, intake, policy, or the

intervention itself. That portfolio response is usually more valuable than asking each team to absorb another local workaround.

Two questions test data boundary. What evidence would show that the condition is no longer adequate? Would that evidence arrive early enough to preserve a feasible choice? The intended choice is clear: Choose a bounded or local pattern when required controls cannot be demonstrated in the cloud; choose cloud when its managed controls better fit the need. A weak answer to either question calls for a smaller commitment, an earlier signal, or a safer fallback. A strong answer earns movement without claiming that uncertainty has disappeared.

Access and identity boundary

The decision includes who can reach models, data, tools, administrative functions, and support channels. The visible artifact is often complete before the operating condition behind it is credible.

A local installation is assumed private despite weak endpoint controls, or a cloud service is assumed controlled because corporate login exists. The organization may call this execution risk, yet the failure began when authority and evidence were separated at the decision point.

Map identities, privileges, service accounts, tenant boundaries, administrative access, and separation of duties. A useful test asks what ordinary operating pressure would break the premise and how that break would be detected before consequence spreads.

Provisioning records, least-privilege tests, privileged access monitoring, and offboarding evidence. Evidence should arrive before the latest useful decision date; perfect analysis delivered after options expire has little governance value.

Reject any option whose identity model cannot support the required accountability. The operating owner should confirm that the decision can be executed with the people, systems, and time actually available.

A tabletop review of access and identity boundary should use the observed break as its scenario: A local installation is assumed private despite weak endpoint controls, or a cloud service is assumed controlled because corporate login exists. Participants then apply the proposed response. They map identities, privileges, service accounts, tenant boundaries, administrative access, and separation of duties. The review identifies where authority, information, time, or recovery would run out. It is useful only if it changes the boundary or confirms who will act when the condition appears.

The minimum record for access and identity boundary is anchored in actual proof. It includes provisioning records, least-privilege tests, privileged access monitoring, and offboarding evidence. The record also carries the choice supported by that proof: Reject any option whose identity model cannot support the required accountability. Keeping evidence and decision together lets a later owner distinguish a deliberate residual risk from a premise that nobody examined. It also gives the team a clear reason to reopen the choice when the underlying facts change.

Across the work governed by CIOs, architecture leaders, security leaders, AI program owners, and workflow executives, occurrences of this pattern should be compared rather than treated as unrelated project stories. The positive standard is: The decision includes who can reach models, data, tools, administrative functions, and support channels. If the same failure recurs, leaders can change delegation, capacity, intake, policy, or the intervention itself. That portfolio response is usually more valuable than asking each team to absorb another local workaround.

Two questions test access and identity boundary. What evidence would show that the condition is no longer adequate? Would that evidence arrive early enough to preserve a feasible choice? The intended choice is clear: Reject any option whose identity model cannot support the required accountability. A weak answer to either question calls for a smaller commitment, an earlier signal, or a safer fallback. A strong answer earns movement without claiming that uncertainty has disappeared.

Action-authority boundary

Inference location does not decide what the system is permitted to do inside a workflow. The issue is not a shortage of activity; it is a shortage of inspectable commitments.

Leaders treat local deployment as inherently safer while granting broad tool or system access, or restrict a well-controlled cloud service unnecessarily. The missing condition is often treated as something the team will resolve later, even though the team lacks the authority or capacity to resolve it.

Specify recommend, prepare, approve, execute, and recover permissions separately from hosting. The forum should be able to distinguish an acceptable residual risk from a missing obligation that has merely been renamed.

Tool scopes, approval records, transaction limits, reversible actions, and tested kill paths. The evidence may be compact, but it must be inspectable. A future reviewer should be able to see what was examined, what conclusion was drawn, and which limitation remained.

Select architecture only after the permitted action envelope is clear. Leaders should prefer the smallest commitment that preserves learning and limits irreversible exposure.

A tabletop review of action-authority boundary should use the observed break as its scenario: Leaders treat local deployment as inherently safer while granting broad tool or system access, or restrict a well-controlled cloud service unnecessarily. Participants then apply the proposed response. They specify recommend, prepare, approve, execute, and recover permissions separately from hosting. The review identifies where authority, information, time, or recovery would run out. It is useful only if it changes the boundary or confirms who will act when the condition appears.

The minimum record for action-authority boundary is anchored in actual proof. It includes tool scopes, approval records, transaction limits, reversible actions, and tested kill paths. The record also carries the choice supported by that proof: Select architecture only after the permitted action envelope is clear. Keeping evidence and decision together lets a later owner distinguish a deliberate residual risk from a premise that nobody examined. It also gives the team a clear reason to reopen the choice when the underlying facts change.

Across the work governed by CIOs, architecture leaders, security leaders, AI program owners, and workflow executives, occurrences of this pattern should be compared rather than treated as unrelated project stories. The positive standard is: Inference location does not decide what the system is permitted to do inside a workflow. If the same failure recurs, leaders can change delegation, capacity, intake, policy, or the intervention itself. That portfolio response is usually more valuable than asking each team to absorb another local workaround.

Two questions test action-authority boundary. What evidence would show that the condition is no longer adequate? Would that evidence arrive early enough to preserve a feasible choice? The intended choice is clear: Select architecture only after the permitted action envelope is clear. A weak answer to either question calls for a

smaller commitment, an earlier signal, or a safer fallback. A strong answer earns movement without claiming that uncertainty has disappeared.

Evaluation and change boundary

Every option needs a way to evaluate behavior as models, prompts, data, and users change. The cost appears downstream, but the missing decision sits upstream.

A local model becomes stale and unsupported, or a cloud model changes beneath a static approval. The recurring signal is translation work: people spend time discovering what was meant, who may decide, and how to repair the unsupported assumption.

Define update authority, regression testing, release evidence, monitoring, and conditions requiring reapproval. This does not require a larger pack. It requires a traceable connection among the premise, the evidence, the boundary, and the action that follows.

Version records, representative evaluations, provider change notices, rollback capability, and unresolved regressions. Evidence has a shelf life. The record needs the conditions or date that would make the conclusion stale.

Prefer the option whose change model the organization can actually govern. A hold is governable when the return conditions are explicit; otherwise it becomes an unmanaged queue.

A tabletop review of evaluation and change boundary should use the observed break as its scenario: A local model becomes stale and unsupported, or a cloud model changes beneath a static approval. Participants then apply the proposed response. They define update authority, regression testing, release evidence, monitoring, and conditions requiring reapproval. The review identifies where authority, information, time, or recovery would run out. It is useful only if it changes the boundary or confirms who will act when the condition appears.

The minimum record for evaluation and change boundary is anchored in actual proof. It includes version records, representative evaluations, provider change notices, rollback capability, and unresolved regressions. The record also carries the choice supported by that proof: Prefer the option whose change model the organization can actually govern. Keeping evidence and decision together lets a later owner distinguish a deliberate residual risk from a premise that nobody examined. It also gives the team a clear reason to reopen the choice when the underlying facts change.

Across the work governed by CIOs, architecture leaders, security leaders, AI program owners, and workflow executives, occurrences of this pattern should be compared rather than treated as unrelated project stories. The positive standard is: Every option needs a way to evaluate behavior as models, prompts, data, and users change. If the same failure recurs, leaders can change delegation, capacity, intake, policy, or the intervention itself. That portfolio response is usually more valuable than asking each team to absorb another local workaround.

Two questions test evaluation and change boundary. What evidence would show that the condition is no longer adequate? Would that evidence arrive early enough to preserve a feasible choice? The intended choice is clear: Prefer the option whose change model the organization can actually govern. A weak answer to either question calls for a smaller commitment, an earlier signal, or a safer fallback. A strong answer earns movement without claiming that uncertainty has disappeared.

Resilience and support boundary

Availability, latency, incident response, continuity, and recovery obligations differ by deployment pattern. The distinction matters because leaders authorize commitments, while teams inherit whatever the authorization left unresolved.

Local capacity is under-engineered because cloud costs are visible, or cloud dependence is accepted without a viable degraded mode. What looks like a team-level miss is frequently a portfolio choice that nobody made explicitly. Scarce attention is then spent reconciling consequences rather than choosing among alternatives.

Model failure modes, operating hours, connectivity, capacity peaks, support skills, recovery objectives, and fallback work. A named owner must be able to change the work; receiving a notification after the boundary is crossed is insufficient.

Load tests, continuity exercises, support rosters, dependency maps, recovery results, and backlog trends. The evidence owner is accountable for integrity, while the decision owner remains accountable for the commitment. Keeping those roles distinct improves challenge.

Use hybrid, bounded functionality, or non-adoption when neither option can meet the operational need responsibly. The residual uncertainty belongs in the decision record so later outcomes are compared with what leaders actually knew at the time.

A tabletop review of resilience and support boundary should use the observed break as its scenario: Local capacity is under-engineered because cloud costs are visible, or cloud dependence is accepted without a viable degraded mode. Participants then apply the proposed response. They model failure modes, operating hours, connectivity, capacity peaks, support skills, recovery objectives, and fallback work. The review identifies where authority, information, time, or recovery would run out. It is useful only if it changes the boundary or confirms who will act when the condition appears.

The minimum record for resilience and support boundary is anchored in actual proof. It includes load tests, continuity exercises, support rosters, dependency maps, recovery results, and backlog trends. The record also carries the choice supported by that proof: Use hybrid, bounded functionality, or non-adoption when neither option can meet the operational need responsibly. Keeping evidence and decision together lets a later owner distinguish a deliberate residual risk from a premise that nobody examined. It also gives the team a clear reason to reopen the choice when the underlying facts change.

Across the work governed by CIOs, architecture leaders, security leaders, AI program owners, and workflow executives, occurrences of this pattern should be compared rather than treated as unrelated project stories. The positive standard is: Availability, latency, incident response, continuity, and recovery obligations differ by deployment pattern. If the same failure recurs, leaders can change delegation, capacity, intake, policy, or the intervention itself. That portfolio response is usually more valuable than asking each team to absorb another local workaround.

Two questions test resilience and support boundary. What evidence would show that the condition is no longer adequate? Would that evidence arrive early enough to preserve a feasible choice? The intended choice is clear: Use hybrid, bounded functionality, or non-adoption when neither option can meet the operational need responsibly. A weak answer to either question calls for a smaller commitment, an earlier signal, or a safer fallback. A strong answer earns movement without claiming that uncertainty has disappeared.

Supplier and exit boundary

A deployment choice creates dependencies on models, hardware, platforms, contracts, skills, and data formats. The control belongs at the point where reliance expands, not in a retrospective explanation after the outcome is known.

The organization compares current run cost while ignoring migration, retraining, data extraction, replacement controls, and concentration. Momentum makes the weakness self-reinforcing. Each additional commitment raises the political and economic cost of revisiting the premise.

Describe how service ends, data returns or is destroyed, workflows continue, and controls transfer to a successor. The record should travel with the work so that later participants can see the basis for reliance instead of reconstructing it from memory.

Exit clauses, export tests, replacement lead times, asset inventory, documentation, and named transition ownership. The strongest record shows both what supports the choice and what could overturn it.

Treat an untested exit as a material commitment, not a future procurement detail. Decision quality improves when dissent and competing evidence are preserved instead of flattened into a unanimous-looking status.

A tabletop review of supplier and exit boundary should use the observed break as its scenario: The organization compares current run cost while ignoring migration, retraining, data extraction, replacement controls, and concentration. Participants then apply the proposed response. They describe how service ends, data returns or is destroyed, workflows continue, and controls transfer to a successor. The review identifies where authority, information, time, or recovery would run out. It is useful only if it changes the boundary or confirms who will act when the condition appears.

The minimum record for supplier and exit boundary is anchored in actual proof. It includes exit clauses, export tests, replacement lead times, asset inventory, documentation, and named transition ownership. The record also carries the choice supported by that proof: Treat an untested exit as a material commitment, not a future procurement detail. Keeping evidence and decision together lets a later owner distinguish a deliberate residual risk from a premise that nobody examined. It also gives the team a clear reason to reopen the choice when the underlying facts change.

Across the work governed by CIOs, architecture leaders, security leaders, AI program owners, and workflow executives, occurrences of this pattern should be compared rather than treated as unrelated project stories. The positive standard is: A deployment choice creates dependencies on models, hardware, platforms, contracts, skills, and data formats. If the same failure recurs, leaders can change delegation, capacity, intake, policy, or the intervention itself. That portfolio response is usually more valuable than asking each team to absorb another local workaround.

Two questions test supplier and exit boundary. What evidence would show that the condition is no longer adequate? Would that evidence arrive early enough to preserve a feasible choice? The intended choice is clear: Treat an untested exit as a material commitment, not a future procurement detail. A weak answer to either question calls for a smaller commitment, an earlier signal, or a safer fallback. A strong answer earns movement without claiming that uncertainty has disappeared.

How the elements interact

Data boundary and Access and identity boundary protect different parts of the same commitment. Deployment must fit the sensitivity, residency, lineage, retention, and secondary-use rules of the actual workflow data. The decision includes who can reach models, data, tools, administrative functions, and support channels. The distinction matters because leaders authorize commitments, while teams inherit whatever the authorization left unresolved. If the operation encounters this failure—teams classify the model in abstract while prompts, retrieved content, logs, embeddings, and support artifacts cross different boundaries.—the response for access and identity boundary cannot compensate for the missing data boundary obligation. Reviewers should reconcile the evidence for both elements before they accept the next action: reject any option whose identity model cannot support the required accountability.

Access and identity boundary and Action-authority boundary protect different parts of the same commitment. The decision includes who can reach models, data, tools, administrative functions, and support channels. Inference location does not decide what the system is permitted to do inside a workflow. The practical failure occurs between the formal approval and the ordinary pressure of running the work. If the operation encounters this failure—a local installation is assumed private despite weak endpoint controls, or a cloud service is assumed controlled because corporate login exists.—the response for action-authority boundary cannot compensate for the missing access and identity boundary obligation. Reviewers should reconcile the evidence for both elements before they accept the next action: select architecture only after the permitted action envelope is clear.

Action-authority boundary and Evaluation and change boundary protect different parts of the same commitment. Inference location does not decide what the system is permitted to do inside a workflow. Every option needs a way to evaluate behavior as models, prompts, data, and users change. Viewed from the portfolio, this is a resource-allocation problem before it becomes a delivery problem. If the operation encounters this failure—leaders treat local deployment as inherently safer while granting broad tool or system access, or restrict a well-controlled cloud service unnecessarily.—the response for evaluation and change boundary cannot compensate for the missing action-authority boundary obligation. Reviewers should reconcile the evidence for both elements before they accept the next action: prefer the option whose change model the organization can actually govern.

Evaluation and change boundary and Resilience and support boundary protect different parts of the same commitment. Every option needs a way to evaluate behavior as models, prompts, data, and users change. Availability, latency, incident response, continuity, and recovery obligations differ by deployment pattern. The control belongs at the point where reliance expands, not in a retrospective explanation after the outcome is known. If the operation encounters this failure—a local model becomes stale and unsupported, or a cloud model changes beneath a static approval.—the response for resilience and support boundary cannot compensate for the missing evaluation and change boundary obligation. Reviewers should reconcile the evidence for both elements before they accept the next action: use hybrid, bounded functionality, or non-adoption when neither option can meet the operational need responsibly.

Resilience and support boundary and Supplier and exit boundary protect different parts of the same commitment. Availability, latency, incident response, continuity, and recovery obligations differ by deployment pattern. A deployment choice creates dependencies on models, hardware, platforms, contracts, skills, and data formats. The visible artifact is often complete before the operating condition behind it is credible. If the operation encounters this failure—local capacity is under-engineered because cloud costs are visible, or cloud dependence is accepted without a viable degraded mode.—the response for supplier and exit boundary cannot compensate for the missing

resilience and support boundary obligation. Reviewers should reconcile the evidence for both elements before they accept the next action: treat an untested exit as a material commitment, not a future procurement detail.

Supplier and exit boundary and Data boundary protect different parts of the same commitment. A deployment choice creates dependencies on models, hardware, platforms, contracts, skills, and data formats. Deployment must fit the sensitivity, residency, lineage, retention, and secondary-use rules of the actual workflow data. A useful governance design therefore begins with the consequence leaders are prepared to accept. If the operation encounters this failure—the organization compares current run cost while ignoring migration, retraining, data extraction, replacement controls, and concentration.—the response for data boundary cannot compensate for the missing supplier and exit boundary obligation. Reviewers should reconcile the evidence for both elements before they accept the next action: choose a bounded or local pattern when required controls cannot be demonstrated in the cloud; choose cloud when its managed controls better fit the need.

Put the model into the management cadence

A framework becomes useful when it changes the normal cadence of work. It should shape intake, funding, design, delivery, and operating review without creating a separate governance industry around itself. The following sequence is deliberately small:

EXHIBIT 2

Start with one decision, then calibrate from evidence

- 1 **Move 1**
Begin with two real workflows; abstract enterprise-wide scoring hides important differences.
- 2 **Move 2**
Complete the data and action maps before comparing models.
- 3 **Move 3**
Estimate standing support, evaluation, security, and recovery capacity for each option.
- 4 **Move 4**
Record the decision, residual risk, triggers for reconsideration, and exit obligations.

- Begin with two real workflows; abstract enterprise-wide scoring hides important differences.
- Complete the data and action maps before comparing models.
- Estimate standing support, evaluation, security, and recovery capacity for each option.
- Record the decision, residual risk, triggers for reconsideration, and exit obligations.

The first cycle should be treated as calibration. Reviewers should note which questions changed a decision, which evidence was unavailable, where authority was unclear, and which controls cost out of proportion to the exposure. That record improves the seven-boundary deployment decision without pretending the first version will fit every workflow or investment.

Ownership should remain close to the consequence. The PMO or AI-governance function can maintain the method, surface cross-portfolio patterns, and challenge weak evidence. It should not absorb the business owner's accountability for the result or the executive's accountability for the commitment.

Measures that reveal whether the control works

- Exceptions or assumptions discovered after the latest useful decision date.
- Scarce specialist and executive attention consumed by recurring unresolved conditions.
- Decisions reopened because the original evidence, boundary, or rationale was unclear.
- Operating outcomes and recovery performance after reliance expanded.
- Time from a material signal to a consequential decision.
- Commitments narrowed, redirected, paused, or stopped before avoidable exposure grew.

These measures are diagnostic, not a universal scorecard. Their purpose is to identify where the operating model fails repeatedly and whether governance is preserving options earlier. A lower count can indicate improvement, but it can also indicate weak detection. Leaders should read the measures with case evidence and frontline observation.

The strongest counterargument

Local is not automatically private, cloud is not automatically resilient, and hybrid is not automatically balanced. Outcomes depend on implementation and operating discipline. The framework supports a reasoned choice; it does not declare one deployment pattern universally superior.

A reasonable critic could also argue that experienced leaders already do this through judgment. Many do. The weakness is not the absence of judgment; it is the absence of a durable operating record when people, pressure, and conditions change. The seven-boundary deployment decision makes the basis for judgment visible enough to challenge, execute, and revisit.

The shift

Local, cloud, hybrid, or no deployment should follow an operating-boundary decision about data, access, action authority, evaluation, resilience, support, suppliers, and exit. Model fashion is not a governance criterion. The practical shift is from reporting activity to governing the conditions under which the organization relies on that activity.

That discipline does not make uncertainty disappear. It gives leaders a better place to stand: a named decision, evidence proportionate to consequence, an explicit boundary, a responsible owner, and a response when reality departs from the assumption. Those elements are enough to turn hidden operating risk into a choice the portfolio can make deliberately.

Notes and sources

[1] NIST, "Artificial Intelligence Risk Management Framework: Generative Artificial Intelligence Profile," July 2024. <https://nvlpubs.nist.gov/nistpubs/ai/nist.ai.600-1.pdf>

[2] Cybersecurity and Infrastructure Security Agency, "Cloud Security Technical Reference Architecture, Version 2," June 2022. https://www.cisa.gov/sites/default/files/2023-02/cloud_security_technical_reference_architecture_2.pdf

About the author

Marco Policani is an enterprise portfolio, PMO, and AI operating-governance leader. He builds portfolio governance systems where AI carries the analytical load and named humans own every decision — an approach documented across case studies, walkthroughs, and working governance guides on his portfolio site.

Portfolio & governance library: policani.net/governance · Full portfolio: policani.net · LinkedIn: linkedin.com/in/marcpolicani

This white paper may be shared freely with attribution. © 2026 Marco Policani.