
WHITE PAPER · AI OPERATING GOVERNANCE

Where Should This AI Run?

Seven checks for local, cloud, and hybrid AI

Marco Policani

Enterprise Portfolio · PMO · AI Operating Governance

policani.net · linkedin.com/in/marcpolicani

September 2026

EXECUTIVE SUMMARY

A claims team has two offers. Neither yet shows who can use the data, what AI may do, or who restores service.

A claims team has a local model and a cloud offer. Neither says who may read the records, what AI may do, or who restores service. The workflow's needs, not the hosting label, decide which option can work.

- 01** Check data, access, actions, changes, failures, operations, and supplier exit before choosing where AI runs.
- 02** A local system is not automatically private, and a cloud service is not automatically reliable.
- 03** Have the workflow owner recommend an option and a named leader approve it and accept the remaining risk.

A model comparison leaves out the hard choice

A business leader must approve an AI tool for claims staff. The team offers a local model because the records are sensitive. It also offers a cloud service with better support. Yet neither proposal says who may read those records, what the tool may do, or who restores service when it fails. A model comparison cannot settle those questions.

The leader must decide before staff depend on the tool. A weak choice can expose records, interrupt claims work, or leave no team able to fix the AI system, meaning the model and the setup that runs it. The same leader may also face a low-risk writing aid, where a costly local setup solves no clear problem.

NIST's AI guidance asks organizations to manage risk across the system lifecycle, including third-party components.[1] CISA's cloud guidance explains that the provider and customer each have duties that depend on the service.[2] NIST describes cloud computing as a set of service and deployment choices, not one fixed design.[3]

Two workflows can need different answers

The claims tool may need strict limits on data, access, and actions. It might use a limited local step or a tightly controlled cloud service. The writing aid may fit a standard cloud service because its team does not have the skills to run a local system. One company rule based on location would miss those differences.

Hybrid also needs a clear description. A team may keep one sensitive step local and use cloud AI for the rest. It may use a stronger cloud model for difficult cases and a smaller local model for routine work. Or it may use a cloud service normally and a limited local option during an outage. These are different designs with different risks.

The choice turns on the workflow and the controls it needs. A location does not make AI private, safe, reliable, or easy to leave. A local service can have weak access controls. A cloud service can have stronger controls than a small team can operate alone. Compare only options that meet the work's needs.

Make seven checks

The two workflows show why the team must check more than model quality. Start with the data and who can reach it. Then set what AI may do. Test how the option will change, fail, run over time, and end.

EXHIBIT 1

Seven checks before choosing where AI runs

- 1 Data**
Trace prompts, files, logs, support access, backups, and deletion.
- 2 Access**
Name users, administrators, service accounts, and vendor support.
- 3 Actions**
Set the strongest action AI may take and the human decision it cannot make.
- 4 Changes**
Name the approver, real-case test, and rollback for each change.
- 5 Failures**
Test the busy day, the bad day, and the lower-service fallback.
- 6 Operations**
Choose the option the organization will staff and fund in year three.
- 7 Supplier and exit**
Check vendor duties, data export, replacement time, and who would move the workflow.

1. Data

Trace the data that enters, leaves, and stays behind: prompts, source files, logs, support access, backups, and deletion. Say which data is allowed and where it may go. If a cloud service cannot show the required controls, do not use it for that workflow. If a local service cannot protect the data well enough, do not treat its location as proof.

2. Access

Name who can use the system and who can administer it. Include service accounts, vendor support, and people who can change settings. A system is not ready if its owner cannot show who did what. Strong identity controls matter in either location.

3. Actions

Write the strongest action the system may take. It may draft, recommend, prepare, or act. Name the decision a person must still make and the way to stop the system. Where a model runs does not decide how much authority it should have.

4. Changes

Cloud providers can change models and services. Local teams must patch, update, and retire their own models. For each option, name who approves a change, how the team tests it with real cases, and how it rolls back when results get worse.

5. Failures

Test the busy day and the bad day. Check response time, service limits, recovery, and a lower-service way to keep the workflow moving. A local system without enough capacity and a cloud system without a workable fallback can both fail when the work needs them most.

6. Operations

Ask who will run the system in year three. A local system needs people who can operate and patch it. A cloud service needs people who can manage the vendor, cost, access, and changes. Choose the option whose ongoing work the organization will actually staff and fund.

7. Supplier and exit

Every choice creates a dependency. Test data export at a real size, check how long a replacement takes, and name the person who would move the workflow. A contract does not remove the customer's responsibility for its data, users, and workflow.

Keep one decision record

The workflow owner brings evidence and recommends an option. A named leader with approval authority chooses where the AI runs and accepts the risks within that authority. Security, privacy, and operations leads set limits the option must meet. If a risk exceeds the approver's authority, take it to a leader who can accept it.

EXHIBIT 2

Record the recommendation, approver, remaining risk, and what sends the choice back for review.

Keep one page for each choice. Record the workflow, seven checks, evidence, recommendation, final decision, approver, risks that remain, and events that require a

new review. A new data class, a stronger action, a provider change, or a serious incident should bring the choice back for review.

NIST's public-cloud guidance explains that security and privacy depend on the controls used, not simply on the deployment label.[4] Its cloud recommendations also help teams compare service choices and responsibilities.[5] Use that evidence to make a choice that another person can inspect and challenge.

What to do next

Before the next model comparison, choose two real workflows. Write the data, access, action, change, failure, operations, and supplier-and-exit checks for each. Have the workflow owner recommend an option, then name who can approve it and accept the remaining risk.

Notes and sources

[1] Chloe Autio, Reva Schwartz, Jesse Dunietz, Shomik Jain, Martin Stanley, Elham Tabassi, Patrick Hall, and Kamie Roberts, "Artificial Intelligence Risk Management Framework: Generative Artificial Intelligence Profile," NIST AI 600-1, July 26, 2024. <https://doi.org/10.6028/NIST.AI.600-1>

[2] Cybersecurity and Infrastructure Security Agency, United States Digital Service, and Federal Risk and Authorization Management Program, "Cloud Security Technical Reference Architecture," Version 2.0, June 2022. https://www.cisa.gov/sites/default/files/2023-02/cloud_security_technical_reference_architecture_2.pdf

[3] Peter Mell and Timothy Grance, "The NIST Definition of Cloud Computing," NIST Special Publication 800-145, September 2011. <https://nvlpubs.nist.gov/nistpubs/Legacy/SP/nistspecialpublication800-145.pdf>

[4] Wayne Jansen and Timothy Grance, "Guidelines on Security and Privacy in Public Cloud Computing," NIST Special Publication 800-144, December 2011. <https://nvlpubs.nist.gov/nistpubs/Legacy/SP/nistspecialpublication800-144.pdf>

[5] Lee Badger, Tim Grance, Robert Patt-Corner, and Jeff Voas, "Cloud Computing Synopsis and Recommendations," NIST Special Publication 800-146, May 2012. <https://nvlpubs.nist.gov/nistpubs/Legacy/SP/nistspecialpublication800-146.pdf>

About the author

Marco Policani is an enterprise portfolio, PMO, and AI operating-governance leader. He builds portfolio governance systems where AI carries the analytical load and named humans own every decision — an approach documented across case studies, walkthroughs, and working governance guides on his portfolio site.

Portfolio & governance library: policani.net/governance · Full portfolio: policani.net · LinkedIn: linkedin.com/in/marcpolicani

This white paper may be shared freely with attribution. © 2026 Marco Policani.