


Oversight Capacity, Not Model Capability, Is the Ceiling on AI Scale

You were told to scale the AI by adding agents. The first limit you hit is one you have never measured.



Marco Policani

Enterprise Portfolio · PMO · AI Operating Governance

policani.net · linkedin.com/in/marcpolicani

July 2026

EXECUTIVE SUMMARY

Capability builds the output. Oversight decides how much of it the organization can safely own.

You were told to scale the AI by adding agents, and the plan assumes the model is the limit. But every agent draws on the same finite pool of people who have to review, correct, and answer for its output, and there is a point past which no one on the team can still answer for the work. Most companies cross it without noticing and find out where it was from an incident nobody caught. This paper shows how to find that limit before it finds you, spend it deliberately, and raise it.

- 01** For each agent, one test: can a named human answer for what it did this week? When the honest answer is no, the span is exceeded.
- 02** Match oversight to risk. Sample the low-stakes, concentrate judgment where errors are costly or hidden, escalate the uncertain.
- 03** Charge every new AI deployment against the shared oversight pool - and be willing to say 'not yet' on oversight grounds.

The limit you never counted

You have been handed the mandate to scale the AI: add agents, add use cases, get more done with what the models can now do. The plan behind that mandate assumes the limit is the model — better models, more of them, more output. For a widening set of tasks that assumption is now wrong, and the miss is expensive.

What actually caps safe scale is oversight capacity: the finite human attention and judgment available to review, correct, and answer for what the AI does. McKinsey states it directly — the scale of agentic adoption will be capped by how much oversight capacity humans can provide, which makes governance itself a potential bottleneck ^[1]. Its figure, two to five people who can already supervise 50 to 100 specialized agents, reads as impressive until you read it as a ceiling: there is a number of agents past which a given team can no longer answer for the work ^[1].

This limit behaves differently from a model limit. It can be measured and budgeted like any other scarce resource, and how high it sits depends on how a company designs its work, not on how good its models get. The next question is the practical one: what, exactly, does a company run out of when it runs out of oversight?

What you are actually running out of

For most of enterprise AI's short history the model really was the limit. It could not do the task reliably, so adoption meant making it good enough. That work is now finished for a widening set of tasks, and the bottleneck moves to the human who has to check the output, own the decision, and answer when it is wrong. Generating output scales at near-zero cost; supervising it does not, because it runs through a person whose attention is fixed and whose judgment cannot be run in parallel. When one side of a system scales and the other side holds still, the side that holds still sets the limit. In enterprise AI, that side is human oversight.

The limit you are about to hit is not in the model; it is the number of decisions your people can still stand behind — a pool every new agent draws down, and one you have probably never counted.

Three properties make oversight capacity the binding constraint. It is finite: a person can genuinely own only so many decisions before ownership becomes a signature on work they did not really inspect. It is shared: every AI system draws on the same pool, so an agent added in one function quietly lowers the oversight available everywhere else. And it is spent whether or not anyone measures it.

That last property is why the limit is crossed silently. A company that never measures its oversight has no reading of how much it has left, so it keeps adding agents until the work outruns the pool. The first proof the limit was there is an incident nobody had the capacity to catch.

The span of supervision

The number that makes this manageable is the span of supervision: how many agents, decisions, or units of AI output one accountable human can genuinely own. It is the AI-era version of span of control, and it belongs in operating design the same way. McKinsey's two-to-five people per 50-to-100 agents is one reported span^[1]. The ratio is an early figure rather than a law, though the shape is durable: the number exists, and it is not infinite. Human-supervisory-control research shows what happens when a supervisor's attention runs short. In a simulation where one operator supervised multiple semi-autonomous aircraft, operators who reached cognitive saturation fell back on two very different management strategies with uneven results, and those given management-by-exception support showed automation bias [6]. A second simulation found where the limit fell in one task: when 42 U.S. Navy personnel supervised 8, 12, or 16 highly autonomous missiles, performance and situation awareness were significantly degraded at 16 [7]. AI agents raise a similar problem in a new domain: the more that run, the sooner a single accountable human runs out of the attention to own them.

The span varies with the work rather than holding at one figure. It shrinks when stakes are high and errors are hard to spot, and grows when stakes are low and errors are cheap and obvious, so it has to be estimated per workflow rather than set once for all of AI.

That estimate needs a test a leader can actually run. For each agent, can a named human say what it did this week and answer for it? When the honest answer becomes no, the span has already been exceeded, whatever the plan on paper says.

Calibrate review to the risk

One misreading has to be cleared first, because it quietly caps AI at human speed: the idea that oversight means a human reviews every action. Reviewing everything only recreates the bottleneck the AI was meant to relieve. Effective oversight samples and escalates by risk. In McKinsey's framing, humans move above the loop — setting policy, watching outliers, and adjusting how involved they are, rather than checking every action ^[1].

This calibration is where most of the real capacity gain lives. Uniform heavy review spends the scarce resource on low-risk output that never needed it, leaving too little for the high-risk output that did. Matching oversight to risk — sampling the low-stakes, concentrating judgment where errors are costly or hidden, and escalating the genuinely uncertain — stretches a fixed pool of attention across far more AI. McKinsey expects the companies that get this balance right to capture more of the agentic advantage ^[1].

Raising the ceiling with layered oversight

The span is not fixed, and raising it without lowering the standard is the central move in scaling AI. The way to raise it is to use technology to extend human oversight rather than replace it. McKinsey describes embedding control agents into the workflow: critic agents that challenge outputs, guardrail agents that enforce policy, and compliance agents that watch regulation, with every action able to be logged and explained ^[1]. Deloitte frames the same design as agents guarding other agents, with humans guarding them in turn ^[2]. Layered this way, a small team can answer for a much larger agent population, because routine supervision is itself automated and the people concentrate on the exceptions and the accountability.

Public standards supply the mechanisms that give layered oversight teeth. The NIST Generative AI Profile suggests actions for continuous monitoring, red-teaming, independent evaluation, incident response, and structured feedback across the lifecycle ^[3]. Each one lets oversight scale without a human watching every action: monitoring catches drift, red-teaming surfaces failure modes before production, and incident response bounds the damage when something slips. The NIST AI Risk Management Framework's govern, map, measure, and manage cycle keeps this a standing discipline rather than a launch-day checklist ^[4].

Layered oversight raises the ceiling, but a ceiling remains. Every oversight agent still answers to a human, and the critics and guardrails are themselves systems that can fail and that take some

oversight to maintain. Treat automation as a multiplier on the reach of the human team, push the ceiling as high as the work safely allows, and keep in view that it is still there.

Oversight is a portfolio decision

Because oversight is finite and shared, where to spend it is a portfolio decision. Every AI system competes for the same pool of judgment, so the test for any deployment is whether it is worth the oversight it will consume, given everything else that oversight could cover. A company that approves AI one system at a time, each justified on its own merits, will over-commit its oversight, because no single approval ever sees the shared constraint it is drawing down.

Leaders can read which constraint they are actually managing from what they reach for when they want to scale. The table pairs the capability lever, which is largely exhausted, with the oversight lever that still moves safe scale.

The capability lever	What it assumes	The oversight lever that actually scales
Deploy a better model	The limit is model quality	Raise the span of supervision per human
Add more agents	More agents means more output	Add oversight capacity as you add agents
Trust the model more	Better models need less review	Match oversight to risk, not to trust
Review every output	Diligence means checking everything	Move above the loop; sample and escalate by risk
Buy more capability	Capability is the bottleneck	Layer oversight agents under human accountability

Every left-column move adds output the company cannot answer for; every right-column move adds the capacity to answer for it, which is what actually raises safe scale.

This connects oversight to funding discipline. The highest-value move is often better allocation rather than more review. Shifting oversight off a low-risk system that was over-watched and onto a high-risk one that was under-watched raises safe scale without adding a single reviewer, the same way reallocating capacity in a project portfolio raises throughput without new hiring. Leaders who feel out of oversight capacity are frequently spending what they have on the wrong AI, so the first move is usually to reallocate rather than to add.

The same logic sets where the marginal dollar goes. Once oversight is the limit, more capability the company cannot supervise buys nothing; the return is higher on raising the ceiling — monitoring, layered control, and workflow redesign that lowers the oversight each unit of output needs — than on more model underneath it.

The failure signature of exceeded oversight

Exceeded oversight has a recognizable signature, and naming it lets a leader catch the problem before the incident that usually announces it. The first sign is ownership in name only: a system with an owner on paper who cannot actually account for what it has been doing. When ownership is only a name in a document, the oversight behind it has already thinned to nothing real.

EXHIBIT 2

Three signs the ceiling is already exceeded

Ownership in name only

The owner on paper cannot say what the system did this week.

Rubber-stamp review

Approval without inspection - a false record of governance.

Surprise incidents

Failures surfacing from corners nobody was really watching.

The corrective

Shrink the AI footprint, or add oversight before adding AI.

The second sign is review that has quietly become rubber-stamping. When output outruns the capacity to examine it, review turns into approval: a human signs off on work they did not truly inspect, because inspecting all of it is no longer possible. That is more dangerous than no review at all, because it produces a false record that the work was governed.

The third sign is surprise: incidents that emerge from corners of the AI estate nobody was really watching, surfacing downstream after the fact. It is the direct operational symptom of oversight spread too thin. The corrective for all three is the same, and it is uncomfortable because it means slowing something down — reduce the live AI footprint to what current oversight can genuinely cover, or add oversight before adding more AI. Responding to these signs by demanding a better model treats the symptom in the wrong system.

The capacity is your experienced people

That oversight pool is specific, experienced people. Oversight is the judgment of humans who know the work well enough to tell when the AI is wrong — a capability that takes years to build and cannot be summoned on demand. That makes the capacity finite and slow to grow, and it ties the ceiling on AI scale to the depth of experienced judgment a company keeps.

This carries an uncomfortable warning. Companies under pressure to show AI savings sometimes cut the experienced staff who are their oversight capacity, on the theory that AI now covers the work — but that work included the judgment that supervises the AI. The cut lowers the oversight ceiling at the very moment the AI is being scaled against it, and the shortage shows up the way oversight shortages always do: as an error nobody had the capacity to catch. Building oversight capacity is therefore partly a workforce problem, and the people who can supervise AI are an asset to grow and keep deliberately.

Designing for oversight capacity

Pulling the discipline together, any AI program can run oversight capacity as six standing moves.

EXHIBIT 1

The oversight-capacity map

Estimate and budget

Size the span by stakes and ambiguity; charge each deployment to the shared pool.

Hold the line

Live agents track the humans who can answer for them.

Match oversight to risk

Sample low-stakes output. Concentrate judgment where it counts.

Layer the oversight

Critic and guardrail agents under named human accountability.

Wire in the instruments

Monitoring, red-teaming, incident response.

- Estimate the span of supervision per workflow, driven by stakes, ambiguity, and how detectable errors are, since one number cannot cover all of AI.

- Budget oversight as a real, finite resource, and charge every new AI deployment against the shared attention pool it will consume.
- Hold the number of live agents within the oversight available, so live AI tracks the humans who can answer for it.
- Match oversight to risk: sample low-stakes output, concentrate judgment where errors are costly or hidden, and escalate the uncertain.
- Layer automated oversight — critic, guardrail, and compliance agents — under a named human who answers for it.
- Wire in monitoring, red-teaming, and incident response, so decay is caught by the system rather than by a person watching everything.

Two limits keep this honest. The published ratio — two to five people supervising 50 to 100 agents — is an early figure reported by McKinsey, not a benchmark that transfers across industries or risk levels ^[1]; it indicates that a ceiling exists, without telling you the safe number in your setting, so measure your own span rather than import a headline. And the return on oversight investment is not yet well quantified: the sources argue that oversight caps safe scale, but not yet how much a given portfolio needs to stay safe. Until that matures, the conservative posture is the defensible one — budget oversight generously, measure the span you can actually sustain, and let live AI follow it. The specific ratios are early; the principle beneath them is not. Human-factors research shows how that oversight fails: when monitoring competes with other tasks for an operator's attention, and when automation has been consistently reliable, people check it less and miss its failures, and simple practice does not overcome the effect [5].

Back to the scaling mandate

Go back to the mandate to scale the AI. More agents and better models raise output for a while, and then the ceiling arrives — when the work outruns the people who can answer for it. How high that ceiling sits is the part a leader actually controls, and it is set by the span of supervision, the layered oversight around it, and the experienced judgment the whole thing is made of. Count that human capacity, spend it deliberately, and the ceiling rises as far as the work safely allows. Ignore it, and the ceiling still arrives, on its own schedule — as the day the company can no longer answer for what its AI just did.

Notes and sources

[1] Alexander Sukharevsky, Alexis Krivkovich, Arne Gast, Arsen Storozhev, Dana Maor, Deepak Mahadevan, Lari Hämäläinen, and Sandra Durth, "The agentic organization: Contours of the next paradigm for the AI era," McKinsey & Company, September 26, 2025.

<https://www.mckinsey.com/capabilities/people-and-organizational-performance/our-insights/the-agentic-organization-contours-of-the-next-paradigm-for-the-ai-era>

[2] David Mallon, Brad Kreit, and Natasha Buckley, "Rethinking operating models for humans with agents," Deloitte Insights, April 2, 2026.

<https://www.deloitte.com/us/en/insights/topics/talent/operating-models-for-humans-ai-agents.html>

[3] Chloe Autio, Reva Schwartz, Jesse Dunietz, Shomik Jain, Martin Stanley, Elham Tabassi, Patrick Hall, and Kamie Roberts, "Artificial Intelligence Risk Management Framework: Generative Artificial Intelligence Profile," NIST AI 600-1, July 26, 2024. <https://doi.org/10.6028/NIST.AI.600-1>

[4] National Institute of Standards and Technology, "AI Risk Management Framework Core," excerpt from AI RMF 1.0, 2023. <https://airc.nist.gov/airmf-resources/airmf/5-sec-core/>

[5] Raja Parasuraman and Dietrich H. Manzey, "Complacency and Bias in Human Use of Automation: An Attentional Integration," Human Factors 52, no. 3 (2010): 381-410. <https://depositonce.tu-berlin.de/bitstreams/cafd2873-814b-4c59-bab1-add42e249d2/download>

[6] M. L. Cummings and P. J. Mitchell, "Operator Scheduling Strategies in Supervisory Control of Multiple UAVs," Aerospace Science and Technology 11 (2007): 339-348. https://dspace.mit.edu/bitstream/handle/1721.1/90289/CummingsMitchell_AST06.pdf

[7] M. L. Cummings and Stephanie Guerlain, "Developing Operator Capacity Estimates for Supervisory Control of Autonomous Vehicles," Human Factors 49, no. 1 (2007): 1-15. <https://europepmc.org/article/MED/17315838>

About the author

Marco Policani is an enterprise portfolio, PMO, and AI operating-governance leader. He builds portfolio governance systems where AI carries the analytical load and named humans own every decision — an approach documented across case studies, walkthroughs, and working governance guides on his portfolio site.

Portfolio & governance library: policani.net/governance · Full portfolio: policani.net · LinkedIn: linkedin.com/in/marcpolicani

This white paper may be shared freely with attribution. © 2026 Marco Policani.