
WHITE PAPER · PORTFOLIO & DELIVERY GOVERNANCE

A Portfolio Needs Scenarios, Not a Single Annual Plan

Pre-agree contingent choices before surprise forces them

Marco Policani

Enterprise Portfolio · PMO · AI Operating Governance

policani.net · linkedin.com/in/marcpolicani

July 2026

EXECUTIVE SUMMARY

A plan alone has no answer when its assumptions break. Scenarios kept beside it let leaders choose from alternatives prepared in advance.

When one assumption in an annual plan breaks, the plan has no answer prepared, and leaders reorder priorities in an emergency meeting. This paper keeps the plan as the version that budgets, hiring, and contracts rest on, and adds scenarios beside it. Each scenario is an alternative portfolio you could actually run, with a trigger someone can observe, one named owner, and first moves agreed in advance that you can reverse. It adapts established scenario planning to the portfolio, and no source reviewed here tests the whole cycle.

- 01** Build two genuinely different alternatives you could actually run. Expect each one to delay or drop some initiatives rather than keep them all.
- 02** Set each trigger by lead time and give it one named owner, so someone is responsible for noticing.
- 03** Agree a first move you can reverse, and the money behind it. Count emergency reprioritizations, the hurried meetings that reorder priorities after a shock, for a year before you start, so you have something to compare with.

When the plan's premise breaks

A hiring freeze lands after the plan is approved. The plan assumed the people it needs, and it says nothing about what to do without them. Leaders call a meeting to reorder the portfolio's priorities. This paper calls such a meeting an emergency reprioritization. It is run in a hurry by whoever can attend, with thin information, while each sponsor defends an initiative. I have supported portfolios that went through this. A planning session could have imagined the shock. The organization had one approved plan and nothing prepared beside it. That is an observation from practice, not a measured finding.

This paper argues that a portfolio needs more than the plan: a small set of scenarios kept beside it. A scenario here is an alternative version of the portfolio, tested in advance so leaders know they could run it. The paper also calls it an alternative. Each

one comes with a trigger, a signal someone can observe that tells leaders it is time to choose, and with the first moves leaders agree to make if it fires. When a surprise arrives, leaders meet to choose among alternatives already prepared instead of inventing them in the meeting.

The title's "not a single plan" means not a plan alone. The approved plan stays as the baseline, the version everyone commits to, because budgets, hiring, vendor contracts, and promises across teams all need one agreed version. The argument is against having only the plan. The title's "needs" is limited too. The paper claims that a portfolio with scarce people, shared dependencies, and slow decisions is better placed to choose well when it has prepared its choices. No source reviewed here compares portfolios that prepared alternatives with portfolios that re-planned after the surprise. The claim rests on the reasoning below and on the sources it uses, not on a measured result.

Revising a portfolio during the year is not new. Robert Cooper and Scott Edgett describe portfolio management for new products as a dynamic decision process in which the list of active projects is constantly updated and revised.[1] This paper's point is what that revising leaves out. A re-forecast answers "what changed?" It does not answer "what do we do by Friday?" Only an alternative composed beforehand does.

Scenario planning is an established method for preparing such answers. Paul Schoemaker describes it as a method for imagining possible futures that compensates for two common errors in decision making, underpredicting and overpredicting change.[2] Hugh Courtney and his McKinsey coauthors apply the idea to strategy under uncertainty: develop a limited number of scenarios, and watch for the trigger events that show a market moving toward one of them.[3] What these sources leave open is the portfolio side: what the organization's own work would look like under each future, who has agreed to the trade-offs, and who watches for the signal. This paper's scenarios are alternative portfolios, composed against scarce capacity, with a named owner for each signal and first moves agreed in advance. That is this paper's adaptation of existing methods, untested as a whole.

What one plan costs when it breaks

This section is the paper's reading of what a plan costs when it breaks, drawn from practice and from one study. It is not a measured result.

The damage is mostly about timing. When a surprise arrives, the organization has to do four things at once: notice the change, invent a response, negotiate it across everyone affected, and execute it. Meanwhile delivery keeps using capacity against

the old assumptions. Each step adds delay, and the delays tend to stack. By the time the new priorities land, scarce people may have kept working on what the new facts had already demoted.

A plan also resists change. In a McKinsey study of 1,616 US-listed companies, the capital each business unit received correlated 0.92, on average, with what it received the year before.[4] Over periods shorter than three years, the same study found that companies that moved more capital earned lower shareholder returns, and one explanation the authors offer is that investors recognize value only once results become visible.[4] This paper reads that as a reason to confirm a move before committing to it. The study measures business units, not initiatives, so applying it to a portfolio's initiatives is this paper's inference. "Capital Reallocation Is the Real Portfolio Decision" covers the evidence and the review that can act on it.

Negotiating in an emergency can also pick the wrong winners. The study's authors describe the politics: a leader whose unit wins less than last year is seen as weak, so leaders defend their share. They report no figures for it.[4] In this paper's reading, that pull is strongest in an emergency meeting. The initiatives that survive are likely to be the ones with the best advocates in the room that week, not the ones the portfolio would pick with a clear head. Money already spent pulls leaders toward spending more, and rarely does anyone speak up for keeping choices open. The cut may land where resistance is weakest.

There is a quieter cost. Some threats and openings may never be recognized as portfolio events. Without named uncertainties and signals to watch, a slow slide, such as a dependency slipping a little each month or demand softening each quarter, may never cross a line anyone owns. Nobody has to decide to ignore it. The portfolio has no routine way to notice.

The cycle in six parts

The cycle runs beside the baseline plan. The scenarios answer a question the plan cannot ask: what do we do if this stops being true? The six parts depend on each other, though not in a straight line. Choosing the uncertainties decides what the alternatives vary, and building the alternatives comes next. The one-page view, the triggers, and the agreed moves are each about an alternative, so none can be written until the alternatives exist, and the page and the triggers can be written in either order. An uncertainty is also usable only if its signal can be observed, so the first part is checked against the triggers. The last part checks all the others. "When it fails, work backward" uses these dependencies to find where a cycle broke.

EXHIBIT 1

Six parts turn a plan into prepared alternatives

- 1 **Pick uncertainties that change a choice**
A scenario should vary something that would change what the portfolio does, not catalog every imaginable event.
- 2 **Build alternatives you could run**
Each scenario has to fit the capacity and dependency circumstances it assumes.
- 3 **Show what each alternative protects**
Leaders need to see what choosing a scenario protects and what it gives up.
- 4 **Set triggers someone can see**
A scenario becomes usable when someone can spot movement toward it early enough to act.
- 5 **Agree the moves in advance**
The portfolio should already know which actions become candidates under each trigger.
- 6 **Review the assumptions, briefly**
Ask what changed and whether the current alternatives still fit, rather than repeating the annual ranking ritual.

Pick the uncertainties that change a choice

This step is easy to get wrong. This paper's reading is that teams often pick uncertainties that are interesting rather than decisive. A scenario should vary something that would change what the portfolio does: the order of the work, the funding, the scope, the staffing, how a dependency is handled. If leaders would do the same things either way, that uncertainty is background noise.

To find candidates, ask each initiative lead which assumption, if it broke, would change their plan. The assumptions that appear in several answers, such as a shared vendor, a scarce specialist, or a regulatory date, are the candidates, because a break there reaches the most work. "The EPMO's Job Is Capacity, Not Calendars" (an EPMO is an enterprise project management office) describes how to find the scarce people. Then ask two questions of each candidate. Could this move enough to matter inside the planning horizon? If it did, which commitments would we want to change? The candidates that pass both are the uncertainties the scenarios vary. Start with two or three. More multiplies the combinations, and this paper's judgment is that the extra work adds few decisions.

An uncertainty is only usable if you can see data that speaks to it. Capacity shows up in attrition and in how long open positions sit unfilled. Dependencies show up in delivery against commitments. Demand shows up in pipeline, volumes, or a regulator's published calendar. If nothing observable stands in for the uncertainty, it can still shape your thinking, but it cannot carry a trigger, so review it on the calendar only.

Keeping the set small also forces an honest conversation about what would really move the portfolio. I have watched leadership teams discover, while ranking uncertainties, that they had been running the portfolio on assumptions that clashed and that nobody had examined. Getting that into the open before anything breaks is part of the work.

Build alternatives you could actually run

Each scenario has to be a portfolio you could actually run if its assumptions came true. This step separates a real operating alternative from a slide. A downside case that lowers the benefit forecast while leaving the same long list of initiatives staffed by the same overcommitted people is the baseline plan with a lower forecast.

Schoemaker's test for scenarios is that they describe generically different futures, not variations on one theme.[2] Courtney and colleagues advise avoiding redundant scenarios that have no unique implications for decisions, and keeping the number limited, because more than four or five tends to hinder decision making.[3] This paper applies the same tests to portfolios and starts with two alternatives beside the baseline, which is its own choice for keeping the work small.

The real work is composing them. Under this scenario's constraints, which work continues at full size, which narrows, which moves later, which stops, and what new work comes in? The scarce people, the shared environments, and the decision forums all have to clear. If your reduced-capacity scenario still assumes output from the architect who left, it fails its own premise. Expect every alternative to need real subtraction. A scenario in which nothing stops may be a sign the group negotiated it to avoid offense. If contracts or law lock the work, say so on the page. The alternative can then change only timing or add money, and that limit is part of what leaders learn.

Composing the alternatives is useful even if no trigger ever fires. It shows which initiatives carry weight in every future and which do not.

Show what each alternative protects

Leaders can only choose a scenario when they can see what it protects and what it gives up. A single financial ranking can hide exactly that. Resilience, regulatory

exposure, customer commitments, the skills you are building, and the freedom to act later can all disappear into one number.

Give each alternative one page. On it: outcomes protected, choices kept open, risks accepted, skills built or given up, commitments sacrificed, and the assumptions it leans on hardest. Keep the uncertainty visible. Use ranges instead of single numbers, and note where the team disagreed about a consequence. The aim is a choice whose costs the choosers have looked at directly, not a winner the spreadsheet computed.

Set triggers someone can see

A scenario without a trigger is a document. Schoemaker says that developing early indicators for each scenario helps an organization recognize which way things are heading.[2] Courtney and colleagues point to the trigger events that show a market moving toward one scenario or another.[3] Both write about business strategy, so what follows about owners, lead time, and confirmation windows is this paper's design.

Give each uncertainty a few observable signals, a threshold that means "convene the choice," a named person who watches, and a schedule for looking. Observable is the word that matters. "Market deterioration" is a mood. "Two straight quarters of pipeline below the planning level" is a fact someone can own.

One person owns each trigger. A threshold watched by a team may be watched by nobody, because every member can reasonably assume someone else is watching. The owner's job is small: check on schedule, log the value, raise the flag at the threshold. A missed check is reported at the next review, and a trigger nobody checks counts as not set.

Setting the threshold trades false alarms against late ones. Settle it by lead time: put the threshold where the move you would make still has room to run well. A trigger that fires once the response is already forced was set for observers, not deciders. When the signal is noisy, add a confirmation window and react to movement that holds, not to spikes, while reacting is still cheap. The same triggers can cover an opening, such as an acquisition target becoming available.

A trigger convenes the decision. It does not make it. The decision owner, the one person with authority to move money and people across initiatives, and the record of each move or hold are set out in "Capital Reallocation Is the Real Portfolio Decision."

Agree the moves before you need them

When a trigger fires, a prepared portfolio starts its next meeting from alternatives. An unprepared one starts from scratch. The agreed moves are what a fired trigger starts.

For each scenario, write down the candidate actions: delay named work, narrow a scope, protect a scarce person, change the order, let a dependency go, bring forward work that builds a needed skill. Name the person who decides, the commitments each move affects, the first step you could reverse, and what you will have to tell people.

The first reversible step matters most, because it lets you respond before you are certain. Robert Pindyck's analysis of investment under uncertainty says that an irreversible commitment gives up the chance to wait for new information, and that the lost chance is part of the cost.[5] His work is theory, not a study of portfolios. It supports the design here: make the smallest move that keeps your choices open while the evidence builds. Pause the next round of funding instead of canceling it. Hold the architect instead of reassigning her. Open the vendor conversation instead of signing. Courtney and colleagues call moves that secure the gains of the best case while limiting losses in the worst "options," and moves that pay off whatever happens "no-regrets moves." [3]

Agree the money at the same time. Contingent money is money set aside, or authority to move money, that a named trigger releases. A move that needs a fresh business case on the day the trigger fires was only ever discussed, not agreed. Two arrangements work: a small contingency fund at portfolio level, with the right to draw on it tied to named triggers, or pre-approved authority to move up to a stated amount between named initiatives, with a review afterward. Money held back is money not spent on the plan, and that is a real cost.

Design the way back too. Write down when to return to the earlier scenario. Without that, every contingent move becomes irreversible in practice, and the portfolio may get more conservative with each firing.

Review the assumptions, and keep it short

A short recurring review keeps the cycle alive. Its agenda is narrow: which signals moved, which assumptions changed, whether the current alternatives are still workable, and which scenarios have expired. A scenario built around an uncertainty that has since resolved has done its job. Retire it and bring in whichever uncertainty has since earned the seat. "Give Stage Gates Teeth by Shrinking What They Decide" describes the same idea for approvals: an end date and a named trigger send the decision back to the board. A scenario expires in the same way, when its uncertainty resolves or its end date arrives.

How often the review meets, and how to avoid moving so often that the portfolio churns, are covered in "Capital Reallocation Is the Real Portfolio Decision." This paper adds only the scope. The review looks at assumptions and feasibility, and only a fired

trigger or an expired assumption reopens what the portfolio contains. Left alone, a review like this can drift back toward annual-planning theater: a full re-prioritization, fresh narratives, and decks that grow each cycle.

A vendor slip, run through the cycle

Consider an illustration built to show the cycle. It is not a case, and it reports no result. A portfolio's largest program depends on a vendor platform whose recent releases slipped, and the vendor has just announced a restructuring. The uncertainty passes the two-question test: the slip could happen this year and would change the order of the work, the staffing, and a contract decision.

The alternatives: the baseline plan assumes the vendor delivers roughly as promised. Scenario B is the portfolio you would run if the vendor slips for a long time. The dependent program moves its integration phase behind a workstream that needs nothing from the vendor, the team scopes a stopgap for the interim, and the contract's renewal option converts to a shorter term. Scenario C is the portfolio you would run if the vendor fails. Migration to another platform starts, some mid-tier initiatives delay to fund it, and the whole integration team moves to the migration. Both alternatives clear the test for scarce people and dependencies.

The triggers and moves: the program's vendor manager tracks release dates against what the vendor committed to, with a threshold of one more slip past an agreed margin. A second, softer signal, the vendor's named engineering lead leaving, convenes the review straight away. Scenario B's first step is reversible: hold the integration start and open the contract conversation. Scenario C's first step is a small paid discovery engagement with the other platform, funded from a small contingency set aside for it.

If the vendor then slipped again, the review would meet, take scenario B's first step, and set a confirmation window for scenario C's discovery work. The illustration does not show that this would work. It shows the starting position: leaders open the meeting with alternatives already composed, an owner watching the signal, and money that needs no new business case.

When it fails, work backward

Most of the cycle hangs on the alternatives, so in this paper's reading a failure tends to trace back to them. A trigger nobody acted on usually points to a move nobody really agreed. That points to consequences nobody faced. That most often points to alternatives nobody composed honestly.

EXHIBIT 2

What each part prepares you to decide

Pick uncertainties that change a choice

Leave out any uncertainty that would not change a portfolio commitment.

Build alternatives you could run

Reject a scenario that keeps every priority alive by hiding the overload.

Show what each alternative protects

Choose the portfolio whose tradeoffs leaders are prepared to own.

Set triggers someone can see

Convene the choice when the trigger enters its action range, not after the old plan has failed.

Agree the moves in advance

Make the smallest move that keeps choices open while you confirm the scenario.

Review the assumptions, briefly

Update the set when the decisions change, rather than keeping scenarios alive for the deck.

The honest test is behavioral, and you run it the first time a trigger really fires. Did the response start from the prepared alternatives? If leaders met, opened the scenario file, and negotiated from its moves, the cycle works, whatever else is wrong with it. If the old emergency ritual came back and nobody opened the file, the organization has scenario documents rather than scenario governance. Fix the agreement step or stop paying for the documents.

A portfolio may go a year without a trigger firing, and a quiet year does not show the cycle works. In that case run the test on paper. Take the last real surprise and ask whether the alternatives would have covered it, who would have been watching, and what the first move would have been.

Objections and costs

Rolling forecasts and quarterly re-planning already do this. They do part of it. A re-forecast tells you the premise broke. A composed alternative tells you what to do about it by Friday. Courtney and colleagues make a related point: when the uncertainty that remains is small enough, a single forecast is a sufficient basis for strategy, and scenarios become useful when a few distinct outcomes would change the strategy.[3] No source reviewed here shows that prepared alternatives beat a

good re-plan. Where surprises are slow and the portfolio has slack, a good re-plan may do as well, and the cycle is not worth its cost. It earns its place where people are scarce, dependencies are shared, and decisions are slow.

Scenarios do not make forecasts certain, and the surprise that arrives may not be one you prepared. Scenarios should not turn into predictions in disguise. What they give you is preparation. They put the consequences in front of you, keep choices open, and shorten the time between noticing a change and making a coherent portfolio choice. Even when the surprise differs, the composition work shows which initiatives carry weight in every future and who has already agreed to what. It cannot cover the surprise nobody imagined.

Agreeing sacrifices in advance can be as political as arguing them in an emergency, and less informed. Sponsors may bargain hard over a hypothetical, or agree to a cut that they would resist once it is real. The defense is partial. The one-page view puts each sponsor's trade-off in front of them while there is time to change the composition, and the agreement is reopened only when a trigger fires or an assumption expires. A sponsor who still disagrees when the trigger fires has not been overruled by a document. The decision owner decides, with the earlier agreement as the starting point.

Scenario work costs time and money, and it can become its own industry. No source prices the cost. It is the time to compose each alternative, which falls on the EPMO and the leads of the initiatives involved, a recurring check by each trigger owner, a short review, and the contingent money held back from other uses. The work can also eat planning capacity that delivery needed. The defense is to keep it deliberately small: two or three uncertainties, two alternatives, triggers on one page, a short review. If your organization cannot keep it that small, it should suspect the exercise has become intellectually interesting, which is not the same as operationally useful. A small portfolio with few shared dependencies may not need the cycle. Counting emergency reprioritizations before and after is a rough check on whether it pays for itself, not proof, because other things change over the same period.

What to do next

Start with one portfolio, or one part of it, where people and dependencies are scarce. Install the cycle in small pieces.

- Build a baseline plus two genuinely different alternatives you could run, and test them against scarce people and dependencies before anyone argues about which they prefer.
- Limit triggers to signals leaders can see and act on, each with one named owner.

- Agree each alternative's first reversible move and the contingent money behind it.
- Run a short scenario review before each portfolio funding decision, and after any real shock.

Order matters. Compose the alternatives first and write the triggers afterward, not the other way round. Working out what you would run if capacity fell shows you which capacity signal matters and where its threshold sits.

Handle sponsor politics in advance. The owner of the initiative that gets delayed under scenario B should hear it from the one-page view, argue it, and either change the composition or accept it. Imposed on the day a trigger fires, the same cut may feel like an ambush.

Ownership follows the usual portfolio split. The EPMO or portfolio office runs the cycle: it keeps the alternatives workable, makes sure the trigger owners are watching, and keeps the review short. The portfolio's decision owner, one person with authority to move money and people across initiatives, makes the choices: which alternative, when to move, when to come back. "Capital Reallocation Is the Real Portfolio Decision" describes that role.

A working cycle leaves a record:

- Uncertainties with named signals, owners, and current values, reviewed on schedule.
- Alternatives that still work against current capacity and dependencies.
- Triggers that fired and convened a decision in time to act.
- First moves made reversibly, then confirmed, extended, or undone.
- The number of emergency reprioritizations per year, which the cycle exists to reduce.

Count the emergency reprioritizations for a year before you start, or the number has nothing to compare with. Read the record once a year to calibrate: raise thresholds for triggers that moved nothing, add signals for deteriorations nobody saw coming, and drop moves no firing ever used.

A single plan asks one future to show up on schedule and treats every departure as an exception to absorb. A scenario cycle assumes the future will argue back and prepares your side in advance. The change is in what a surprise means. It stops being a planning crisis and becomes a trigger doing its job, noticed early and met with alternatives already prepared. The two approaches look similar on paper. The difference shows up the week something breaks.

Notes and sources

- [1] Robert G. Cooper and Scott J. Edgett, Portfolio Management: Fundamental for New Product Success, Stage-Gate International. <https://www.stage-gate.com/blog/portfolio-management-fundamental-for-new-product-success/>
- [2] Paul J. H. Schoemaker, Scenario Planning: A Tool for Strategic Thinking, Sloan Management Review 36(2), Winter 1995, 25-40.
https://www.ftms.edu.my/images/Document/MOD001074%20-%20Strategic%20Management%20Analysis/WK4_SR_MOD001074_Schoemaker_1995.pdf
- [3] Hugh Courtney, Jane Kirkland, and Patrick Viguier, Strategy Under Uncertainty, McKinsey Quarterly (adapted from Harvard Business Review, November-December 1997).
<https://www.mckinsey.com/~media/McKinsey/Featured%20Insights/Managing%20In%20Uncertainty/Strategy%20under%20uncertainty/Strategy%20under%20uncertainty.pdf>
- [4] Stephen Hall, Dan Lovallo, and Reinier Musters, How to Put Your Money Where Your Strategy Is, McKinsey & Company, March 1, 2012.
<https://www.mckinsey.com/capabilities/strategy-and-corporate-finance/our-insights/how-to-put-your-money-where-your-strategy-is>
- [5] Robert S. Pindyck, Irreversibility, Uncertainty, and Investment, Journal of Economic Literature 29(3), 1991, 1110-1148; NBER Working Paper 3307, 1990.
<https://www.nber.org/papers/w3307>

About the author

Marco Policani is an enterprise portfolio, PMO, and AI operating-governance leader. He builds portfolio governance systems where AI carries the analytical load and named humans own every decision — an approach documented across case studies, walkthroughs, and working governance guides on his portfolio site.

Portfolio & governance library: policani.net/governance · Full portfolio: policani.net · LinkedIn: linkedin.com/in/marcpolicani

This white paper may be shared freely with attribution. © 2026 Marco Policani.