
WHITE PAPER · PORTFOLIO & DELIVERY GOVERNANCE

A Portfolio Needs Scenarios, Not a Single Annual Plan

Pre-agree contingent choices before surprise forces them

Marco Policani

Enterprise Portfolio · PMO · AI Operating Governance

policani.net · linkedin.com/in/marcpolicani

July 2026

EXECUTIVE SUMMARY

Scenarios create options before a capacity change, dependency slip, or new demand forces an unprepared choice.

A single annual plan tends to fail quietly. Nothing gets canceled, the delay spreads evenly across every initiative, and the portfolio only reprioritizes in an emergency, under sunk-cost pressure. A scenario cycle prepares the alternatives first. You build the current plan plus two alternative portfolios you could actually run. You set triggers: signals someone can observe that tell leaders it is time to make a decision, each with one named owner. You agree the first moves in advance, and you can reverse them.

- 01** Build two genuinely different alternatives you could actually run. Make at least one of them delay or drop initiatives rather than keep them all.
- 02** Set each trigger by lead time and give it one named owner, so noticing is never the bottleneck.
- 03** Measure the cycle by how many emergency reprioritizations it avoids. When a surprise arrives, leaders should meet to choose among options already prepared.

Three surprises before June

An annual plan is a good commitment baseline. It is a poor description of every possible future. A short set of workable alternatives lets leaders agree in advance what they will protect, delay, reshape, or stop when capacity, demand, dependencies, or assumptions change. Moving money and people on that evidence is what separates governing a portfolio from administering a plan. Research on capital reallocation found that companies that shifted more capital across their business units over 15 years earned higher shareholder returns, on average, than companies whose allocations barely moved.[3] Treating the portfolio as a decision you keep making, not one you set once a year, is a long-standing principle in new-product portfolio management.[5]

I learned to distrust single plans the year a portfolio I supported took three surprises before June. A hiring freeze in February. An acquired product line in April. A platform vendor pushing a migration nobody had scheduled. None of them was exotic, and any planning session could have imagined all three. What the organization had instead was one approved plan. So each surprise set off the same ritual: an emergency reprioritization, run in a hurry, by whoever could attend. Those meetings happen in the circumstances that produce bad portfolio decisions — thin information, sunk-cost pressure, and a room of people defending their own initiatives.

Every plan eventually meets a surprise. What matters is what happens next. When a planning premise breaks, most portfolios keep up the plan's appearance: every initiative stays alive, delay spreads evenly, and leaders share out the overload as a kind of fairness. The plan document stays formally in force long after it stopped describing anything. The other response is to choose again, on purpose, from options prepared earlier. Preparing them is what scenario work buys. The aim is better decisions, not more forecasting — judged on quality, speed, execution, and the effort each choice consumes.[4]

A scenario here is an operating option, not a forecast or a strategy story. It is an alternative version of the portfolio, tested in advance so you know you could run it, with the signals that would trigger it and the moves it implies. This paper is for portfolio executives, EPMO leaders, strategy leaders, and functional executives. It lays out a six-part cycle for building and running those options. Flyvbjerg's due-diligence work argues for testing front-end estimates against the outside view before you commit.[2] PMI's complex-project research finds that high-performing organizations prioritize sponsor alignment at project initiation.[1] Scenario planning has a long history in strategy: for decades, planners have used it to rehearse decisions against a few different but plausible futures rather than one forecast.[6] What this paper adds is the portfolio side — agreed-in-advance alternatives with named triggers and reallocation rules, checked for feasibility before the surprise arrives. I assembled the cycle the hard way, one absorbed surprise at a time.

What one plan costs when it breaks

The damage is mostly about timing. When a surprise arrives, the organization has to do four things at once: notice the change, invent a response, negotiate it across everyone affected, and execute it. Meanwhile delivery keeps burning capacity against the old assumptions. Each step carries its own delay, and the delays stack. By the time the reprioritization lands, months of scarce capacity have gone into work the new reality had already demoted.

Negotiating under fire also picks the wrong winners. In an emergency reprioritization, the initiatives that survive tend to be the ones with the best advocates in the room that week, not the ones the portfolio would pick with a clear head. Sunk cost argues loudly. Nobody speaks up for keeping options open. The cut lands wherever resistance is weakest. Each of those calls can be defended on its own, and together they add up to no direction at all. That is why portfolios that live this way feel permanently busy and oddly aimless.

People who defend the annual plan are right about one thing, and this cycle should say so plainly: organizations need a commitment baseline. Budgets, hiring, vendor contracts, and promises across functions all need one stated format. The argument is not against having a plan; it is against having only a plan. One composed future, with no prepared alternatives beside it, gets defended long past its premises — because the only other option is improvising in public.

There is a quieter cost. Some threats and opportunities never get recognized as portfolio events at all. Without named uncertainties and signals to watch, a slow slide — a dependency slipping a little each month, demand softening quarter by quarter — never crosses a line anyone owns. Nobody decides to ignore it. The portfolio simply has no way to notice.

The cycle in six parts

The cycle runs continuously beside the baseline plan. The plan stays the commitment of record. The scenarios are the prepared answers to a question the plan cannot ask: what do we do if this stops being true?

EXHIBIT 1

Six moves turn a plan into prepared options

- 1 **Pick uncertainties that change a choice**
A scenario should vary something that would change what the portfolio does, not catalog every imaginable event.
- 2 **Build alternatives you could run**
Each scenario has to fit the capacity and dependency circumstances it assumes.
- 3 **Show what each option protects**
Leaders need to see what choosing a scenario protects and what it gives up.
- 4 **Set triggers someone can see**
A scenario becomes usable when someone can spot movement toward it early enough to act.
- 5 **Agree the moves in advance**
The portfolio should already know which actions become candidates under each trigger.
- 6 **Review the assumptions, briefly**
Ask what changed and whether the current alternatives still fit, rather than repeating the annual ranking ritual.

Pick the uncertainties that change a choice

Scenario work fails most often at this first step, because teams pick uncertainties that are interesting rather than decisive. A scenario should vary something that would change what the portfolio actually does — the order of the work, the funding, the scope, the staffing, how you handle a dependency. If leaders would do the same things either way, that uncertainty is background noise.

Find the signal at this step too. An uncertainty is only usable if you can see data that speaks to it. Capacity shows up in attrition and in how long requisitions sit open. Dependencies show up in delivery against commitments. Demand shows up in pipeline, volumes, or a regulator's published calendar. If nothing observable stands in for the uncertainty, it can still shape your thinking, but it cannot carry a trigger. Favor the ones you can watch.

Ask two questions of each candidate. Could this move enough to matter inside the planning horizon? If it did, which commitments would we want to change? The ones that pass both become scenario axes: a regulatory date that could shift, whether you keep a critical skill, how confident you are in a dependency's dates, a demand

assumption with real spread. Two or three are enough. Past that, the combinations make work without adding decisions.

Keeping the set small also forces an honest ranking conversation: of everything that could change, what would really move this portfolio? I have watched leadership teams discover mid-argument that they disagreed about the answer. They had been running the portfolio on assumptions that clashed and that nobody had examined. Getting that disagreement into the open before anything breaks is the work paying for itself.

Build alternatives you could actually run

Each scenario has to be a portfolio you could actually run under the circumstances it assumes. This step is what separates a real operating option from a slide. A downside case that lowers the benefit forecast while leaving the same forty initiatives staffed by the same overcommitted people is just the baseline plan wearing a worried expression.

The real work is composing it. Under this scenario's constraints, which work continues at full size, which narrows, which moves later, which stops, and what new work comes in? The scarce roles, the shared environments, and the decision forums all have to clear. If your reduced-capacity scenario still assumes output from the architect who left, it fails its own premise. Expect every alternative to require real subtraction. A scenario in which nothing stops is usually a sign the group negotiated it to avoid offense rather than built it to fit the constraint.

This testing pays for itself even if no trigger ever fires. It shows which initiatives carry weight in every future and which are along for the ride. That alone sharpens the next intake round: a proposal that would survive no adverse scenario is one the portfolio is carrying on optimism.

Show what each option protects

Leaders can only choose a scenario when they can see what it protects and what it gives up. A financial ranking flattens exactly the wrong things. Resilience, regulatory exposure, customer commitments, the skills you are building, and the freedom to act later all disappear into a net-present-value ordering that no executive really believes.

Give each alternative one page. On it: outcomes protected, options preserved, risks accepted, skills built or given up, commitments sacrificed, and the assumptions it leans on hardest. Keep the uncertainty visible. Use ranges instead of single numbers,

and note where the team disagreed about a consequence. You want a choice whose costs the choosers have looked at directly, not a winner the spreadsheet computed.

Set triggers someone can see

A scenario without a trigger is a document. With one, it becomes a standing decision waiting for its trigger. Give each uncertainty a few observable signals, a threshold range that means “convene the choice,” a named person who watches, and a stated schedule for looking. Observable is the word that matters. “Market deterioration” is a mood. “Two straight quarters of pipeline below X” is a fact someone can own.

One person owns each trigger, by design. A threshold watched by a team is watched by nobody, because every member reasonably assumes someone else is watching. The owner’s job is small: check on schedule, log the value, raise the flag at the threshold. Because it is small, it survives busy quarters that would kill a heavier duty.

You are trading false alarms against late ones. Settle it by lead time: put the threshold wherever the move you would make still has room to run well. A trigger that fires once the response is already forced was set for observers, not deciders. When the signal is noisy, add a confirmation window instead of raising the threshold. React to movement that holds, not to spikes, while reacting is still cheap.

A trigger convenes the decision. It does not make it. Leaders still own the judgment call. All it takes away is the chance of nobody noticing, and the political cost of being the person who raises the alarm.

Agree the moves before you need them

When a trigger fires, a prepared portfolio starts its next meeting from options. An unprepared one starts from scratch. The agreed moves are those options. For each scenario, write down the candidate actions — delay named work, narrow a scope, protect a scarce resource, change the order, let a dependency go, speed up a capability. Name the person who decides, the commitments each move affects, the first step you could reverse, and what each move means you have to tell people.

That first reversible step matters most, because it lets you respond before you are certain. Commit fully the first time a trigger fires and you over-steer. Wait until the scenario is undeniable and you wasted the preparation. Take the middle: make the smallest move that keeps your options open while the evidence builds. Pause the next round of funding instead of canceling. Hold the architect instead of reassigning

her. Open the vendor conversation instead of signing. Each one buys information cheaply.

Design the way back at the same time. Write down when to return to the earlier scenario. Skip that and every contingent move becomes a one-way door, and the portfolio gets more conservative with each firing. A move you can enter, evaluate, and exit steers the portfolio. A move you can only enter is a slow-motion emergency.

Review the assumptions, and keep it short

A short recurring review keeps the cycle alive, with a narrow agenda: which signals moved, which assumptions changed, whether the current alternatives are still workable, and which scenarios have expired. Expiry is a sign of health. A scenario built around an uncertainty that has now resolved has done its job. Retire it and bring in whichever uncertainty has since earned the seat.

The review's enemy is institutional gravity. Left alone, it drifts back toward annual-planning theater: full re-prioritization, fresh narratives, decks that grow every cycle. Hold the scope. The review looks at assumptions and feasibility. Only a fired trigger or an expired premise reopens what the portfolio contains. Twenty minutes on the quarterly portfolio agenda keeps the habit alive. A standalone scenario summit kills it within a year.

A vendor slip, run through the cycle

Here is the chain end to end. The uncertainty: a portfolio's largest program depends on a vendor platform whose last two releases slipped, and the vendor has just announced a restructuring. It passes the two-question test. The risk could come true this year, and if it did, it would change the order of the work, the staffing, and a contract decision.

The alternatives: the baseline assumes the vendor delivers within a quarter of what it promised. Scenario B assumes a two-quarter slip. The dependent program moves its integration phase behind a workstream that needs nothing from the vendor, the team scopes a stopgap capability for the interim, and the contract's renewal option converts to a shorter term. Scenario C assumes the vendor fails. Migration to the other platform starts, two mid-tier initiatives delay to fund it, and the whole constrained integration team moves to the migration. Both alternatives clear the constraint model. The team composed both calmly, in an afternoon workshop, months before anything happened.

The triggers: the program's vendor manager tracks release dates against what the vendor committed to, with a threshold of one more slip past four weeks. A second, softer signal — the vendor's named engineering lead leaves — convenes the review straight away. The moves: scenario B's first step is reversible, holding the integration start and opening the contract conversation. Scenario C's first step is a small paid discovery engagement with the other platform.

When the vendor slipped again in August, the review met within a week, took scenario B's first step, and set a confirmation window for scenario C's discovery work. Total decision time: nine days, most of it scheduling. I have lived the unprepared version of that same event. It eats a quarter, and the migration starts from nothing with the integration team already burned out.

When it fails, work backward

The parts depend on each other in one direction. Your triggers are only as good as the uncertainties you chose. Your moves are only executable if the alternatives were workable. The review means nothing unless the triggers are observable. So when the cycle disappoints, trace backward. A trigger nobody acted on usually points to a move nobody really agreed. That points to consequences nobody faced. That, most often, points to alternatives nobody composed honestly.

EXHIBIT 2

What each part prepares you to decide**Pick uncertainties that change a choice**

Leave out any uncertainty that would not change a portfolio commitment.

Build alternatives you could run

Reject a scenario that keeps every priority alive by hiding the overload.

Show what each option protects

Choose the portfolio whose tradeoffs leaders are prepared to own.

Set triggers someone can see

Convene the choice when the trigger enters its action range, not after the old plan has failed.

Agree the moves in advance

Make the smallest move that keeps options open while you confirm the scenario.

Review the assumptions, briefly

Update the set when the decisions change, rather than keeping scenarios alive for the deck.

The honest test is behavioral, and you run it the first time a trigger really fires. Did the response start from the prepared options? If leaders met, opened the scenario file, and negotiated from its moves, the cycle works, whatever else is wrong with it. If the old emergency ritual came back and nobody opened the file, the organization has scenario documents rather than scenario governance. Fix the agreement step or stop paying for the documents.

Starting small beside the annual plan

You can install the cycle in small pieces:

- Build a baseline plus two genuinely different alternatives you could run, and test them against scarce roles and dependencies before anyone argues about which they prefer.
- Limit triggers to signals leaders can see and act on.
- Run a short scenario review before quarterly portfolio decisions, and after any real shock.

Order matters. Compose the alternatives first and let those one-page views drive the trigger discussion, not the other way round. Teams that start by brainstorming triggers build elaborate sensing for options they never compose. Teams that start with two honest alternatives find their triggers on the way. Working out what we would run if capacity fell shows you exactly which capacity signal matters and where its threshold sits.

Sort out the money early, because contingent moves need contingent money. A move that needs a fresh business case the day the trigger fires was only ever discussed, not agreed. Two approaches work. Hold a small contingency fund at portfolio level and tie the right to draw on it to named triggers. Or pre-approve reallocation authority, so the forum can move up to a stated amount between named initiatives when a trigger fires, reviewed afterward. Either one turns that first reversible step from a budgeting exercise into an operating one.

Handle sponsor politics in advance too. The owner of the initiative that gets delayed under scenario B should not find out on the day the trigger fires. The one-page conversation is where that sponsor hears it, argues it, and either changes the composition or accepts it. Executives who would fight that cut in an emergency can often agree calmly, months earlier, that under scenario B this is the program that waits. A sacrifice negotiated in advance is a governance agreement. Imposed during a crisis, it is an ambush, and ambushed sponsors fight back. Much of the cycle's value is moving that negotiation to the cheap end of the timeline.

Ownership follows the usual portfolio split. The EPMO or portfolio function runs the cycle: it keeps the alternatives workable, makes sure the signal owners are actually watching, and keeps the review short. The choices — which alternative, when to move, when to come back — belong to the executive forum. The only new role is the trigger owner, one named person per signal, for the reasons already given.

If your organization runs rolling forecasts or quarterly funding, the cycle sits alongside them. Rolling processes re-examine the plan on a calendar. The scenario cycle re-examines it on evidence. The quarterly re-forecast is a scheduled review, a fired trigger is an unscheduled one, and both go to the same forum. What the scenario cycle adds is the prepared option. Re-forecasting tells you the premise broke. Only a composed alternative tells you what to do about it by Friday.

What the record should show

A working cycle leaves evidence:

- Uncertainties with named signals, owners, and current values, reviewed on schedule.

- Alternatives that still work against the current constraint model.
- Triggers that fired and convened a decision inside their intended lead time.
- First moves made reversibly, then confirmed, extended, or undone.
- Scenarios retired once their uncertainty resolved, and replaced by live ones.
- Emergency reprioritizations per year — the number the cycle exists to reduce.

Read the record once a year to calibrate, not just to hold people to account. Triggers that fired but moved nothing argue for higher thresholds. Deteriorations that nobody saw coming argue for new signals. Moves that no firing ever used may be decoration. The cycle rests on assumptions too, and it deserves the same regular challenge it imposes on the plan.

One measure matters most: does the organization still hold emergency reprioritizations? When the cycle works, events that would have been emergencies arrive as convened decisions instead. One avoided emergency typically repays a year of review minutes.

Two objections worth taking seriously

First, scenarios do not make forecasts any more certain, and they should not turn into predictions in disguise. What they give you is preparation. They put the consequences in front of you, keep options open, and shorten the time between noticing a changed circumstance and making a coherent portfolio choice.

Second, and this objection has more force than it usually gets: scenario work can become its own industry, eating planning capacity that delivery needed. The defense is keeping it deliberately small — two or three uncertainties, two alternatives, triggers on one page, a twenty-minute review. If your organization cannot keep it that small, it should suspect the exercise has become intellectually interesting, which is a different thing from operationally useful. It deserves exactly as much investment as the emergencies it prevents, and no more.

From one plan to prepared options

A single plan asks one future to show up on schedule and treats every departure from it as an exception to absorb. A scenario cycle assumes the future will argue back, and prepares your side of the argument in advance: uncertainties named, alternatives composed, consequences faced, triggers owned, first moves reversible, return paths designed.

Upside deserves the same machinery, and that gets lost because scenario language leans defensive. An acquisition target coming available, a competitor stumbling, a regulation opening a market — those are triggers too. A portfolio with an agreed upside alternative can commit capacity to a sudden opening while competitors are still scheduling their emergency meeting. Preparation cuts both ways. Only the habit of treating it as insurance makes it feel otherwise.

The deeper change is what surprise means. Most surprise stops being a planning crisis and becomes a trigger doing its job: noticed early, met with prepared options, decided while deciding is still cheap. The two plans look similar on paper. The difference shows up the week something breaks.

Notes and sources

[1] Project Management Institute, "Pulse of the Profession 2026: Driving Success in Complex Projects," May 2026. <https://www.pmi.org/learning/thought-leadership/driving-success-in-complex-projects>

[2] Bent Flyvbjerg, "Quality Control and Due Diligence in Project Management," 2013. <https://arxiv.org/abs/1302.2544>

[3] Stephen Hall, Dan Lovallo, and Reinier Musters, "How to Put Your Money Where Your Strategy Is," McKinsey & Company, March 1, 2012. <https://www.mckinsey.com/capabilities/strategy-and-corporate-finance/our-insights/how-to-put-your-money-where-your-strategy-is>

[4] Marcia Blenko, Michael Mankins, and Paul Rogers, "The Five Steps to Better Decisions," Bain & Company, August 6, 2013. <https://www.bain.com/insights/the-five-steps-to-better-decisions/>

[5] Robert G. Cooper and Scott J. Edgett, "Portfolio Management: Fundamental for New Product Success," Stage-Gate International. <https://www.stage-gate.com/blog/portfolio-management-fundamental-for-new-product-success/>

[6] Paul J. H. Schoemaker, "Scenario Planning: A Tool for Strategic Thinking," Sloan Management Review 36(2), Winter 1995, 25–40. https://www.ftms.edu.my/images/Document/MOD001074%20-%20Strategic%20Management%20Analysis/WK4_SR_MOD001074_Schoemaker_1995.pdf

About the author

Marco Policani is an enterprise portfolio, PMO, and AI operating-governance leader. He builds portfolio governance systems where AI carries the analytical load and named humans own every decision — an approach documented across case studies, walkthroughs, and working governance guides on his portfolio site.

Portfolio & governance library: policani.net/governance · Full portfolio: policani.net · LinkedIn: linkedin.com/in/marcpolicani

This white paper may be shared freely with attribution. © 2026 Marco Policani.