
WHITE PAPER · AI OPERATING GOVERNANCE

An AI Evaluation Is a Production Operating Control

Govern changing workflow behavior, not a one-time test

Marco Policani

Enterprise Portfolio · PMO · AI Operating Governance

policani.net · linkedin.com/in/marcpolicani

July 2026

DECISION SIGNAL

Evaluation continues in production because workflow behavior, context, and reliance continue to change.

The operating problem

Scenario set is not a label to add to a project charter. It is a commitment with an owner, a boundary, and evidence that a sponsor can examine. The point is to turn an assumption into an explicit condition of work.

In practice, teams should ask what would make scenario set fail under ordinary operating pressure: a competing priority, a missing input, an ambiguous decision, a changed dependency, or an exception that no one owns. The answer identifies the smallest control required before reliance expands.

The standard is deliberately practical. A team can narrow the commitment, stage discovery, add a named review, secure a decision, or stop the work. Each is better than continuing under a premise that the organization has not actually accepted.

The operating standard

The first design question is ownership. Someone must be able to say whether scenario set is adequate for the next commitment, and that person needs information early enough to act. Naming an owner is not bureaucracy; it prevents a consequential judgment from being distributed so widely that nobody can make it.

The second design question is evidence. Evidence need not be a large report. It may be a decision record, a sample of real workflow cases, a capacity confirmation, a stakeholder commitment, or a documented test. The useful standard is whether a skeptical sponsor can understand why scenario set is credible and what remains uncertain.

The third design question is boundary. Teams need to know what scenario set does not authorize, what condition would cause it to be revisited, and which escalation path applies when the boundary is crossed. A visible boundary turns surprise into a managed exception rather than an informal renegotiation.

What must become explicit

The fourth design question is cadence. Operating conditions change after a decision. A lightweight recurring review asks whether scenario set still fits the work, whether the original evidence still applies, and whether reliance should expand, remain bounded, or be reduced. That is how a framework becomes an operating control.

The final design question is learning. Recurring problems around scenario set should improve the next decision rather than become folklore about a difficult project. Capturing the pattern gives portfolio leaders better questions, more realistic commitments, and a way to distinguish a one-off event from a weakness in the operating model.

This is also where proportionality matters. A low-consequence decision can use a short record and a rapid review. A decision that changes a recurring workflow, exposes sensitive information, consumes scarce capacity, or creates a difficult-to-reverse dependency deserves more evidence and a clearer escalation path. The framework scales the control to the commitment rather than applying one ritual to every kind of work.

Applying the test

A useful operating illustration is a team that appears ready until an ordinary case falls outside the happy path. The work then pauses while people discover who may decide, what outcome is actually protected, whose participation is needed, or what capacity is available. Designing scenario set in advance makes that ordinary case visible before it becomes expensive.

The portfolio implication is cumulative. One unresolved condition may be tolerable; many unresolved conditions across active initiatives create a hidden demand on senior attention and specialist capacity. Making scenario set explicit allows leaders to decide which uncertainty deserves investment, which should remain bounded, and which initiative should wait.

The practical record should be concise enough to be used. It should identify the next commitment, the owner, the evidence examined, the boundary of reliance, and the next review date. If the record cannot make those elements plain, it is not yet helping the team govern scenario set; it is simply producing more documentation.

A proportionate control

A useful review also names the decision that follows from the evidence. If scenario set is adequate, the organization may extend the commitment. If it is not, the team should be able to narrow the scope, add discovery, change the intervention, or hold the work without treating any of those choices as failure. This keeps governance attached to a real choice instead of turning it into commentary on work that will continue regardless.

Leaders should resist the temptation to convert this test into a universal score. The relevant question is whether the available evidence is sufficient for this particular next step, given the consequence of error and the reversibility of the commitment. A small, reversible trial can proceed with less evidence than an enterprise-wide rollout, a material dependency, or an automated action that is difficult to undo.

This distinction matters because a portfolio is not improved by making every initiative look equally governed. It is improved when leaders can see which commitments have earned confidence, which still need evidence, and where limited attention should be applied next. The record makes that judgment repeatable without pretending that the judgment can be automated.

The shift

The habit can begin with a short standing question in the normal management cadence: what would we need to see before we extend reliance on scenario set? The answer should be recorded where the people doing the work can find it. That modest discipline reduces retrospective debate because it makes the original basis for the decision available when conditions change.

Acceptance boundary is not a label to add to a project charter. It is a commitment with an owner, a boundary, and evidence that a sponsor can examine. The point is to turn an assumption into an explicit condition of work.

In practice, teams should ask what would make acceptance boundary fail under ordinary operating pressure: a competing priority, a missing input, an ambiguous decision, a changed dependency, or an exception that no one owns. The answer identifies the smallest control required before reliance expands.

The standard is deliberately practical. A team can narrow the commitment, stage discovery, add a named review, secure a decision, or stop the work. Each is better than continuing under a premise that the organization has not actually accepted.

The first design question is ownership. Someone must be able to say whether acceptance boundary is adequate for the next commitment, and that person needs information early enough to act. Naming an owner is not bureaucracy; it prevents a consequential judgment from being distributed so widely that nobody can make it.

The second design question is evidence. Evidence need not be a large report. It may be a decision record, a sample of real workflow cases, a capacity confirmation, a stakeholder commitment, or a documented test. The useful standard is whether a skeptical sponsor can understand why acceptance boundary is credible and what remains uncertain.

The third design question is boundary. Teams need to know what acceptance boundary does not authorize, what condition would cause it to be revisited, and which escalation path applies when the boundary is crossed. A visible boundary turns surprise into a managed exception rather than an informal renegotiation.

The fourth design question is cadence. Operating conditions change after a decision. A lightweight recurring review asks whether acceptance boundary still fits the work, whether the original evidence still applies, and whether reliance should expand, remain bounded, or be reduced. That is how a framework becomes an operating control.

The final design question is learning. Recurring problems around acceptance boundary should improve the next decision rather than become folklore about a difficult project. Capturing the pattern gives portfolio leaders better questions, more realistic commitments, and a way to distinguish a one-off event from a weakness in the operating model.

This is also where proportionality matters. A low-consequence decision can use a short record and a rapid review. A decision that changes a recurring workflow, exposes sensitive information, consumes scarce capacity, or creates a difficult-to-reverse dependency deserves more evidence and a clearer escalation path. The framework scales the control to the commitment rather than applying one ritual to every kind of work.

A useful operating illustration is a team that appears ready until an ordinary case falls outside the happy path. The work then pauses while people discover who may decide, what outcome is actually protected, whose participation is needed, or what capacity is available. Designing acceptance boundary in advance makes that ordinary case visible before it becomes expensive.

The portfolio implication is cumulative. One unresolved condition may be tolerable; many unresolved conditions across active initiatives create a hidden demand on senior attention and specialist capacity. Making acceptance boundary explicit allows leaders to decide which uncertainty deserves investment, which should remain bounded, and which initiative should wait.

The practical record should be concise enough to be used. It should identify the next commitment, the owner, the evidence examined, the boundary of reliance, and the next review date. If the record cannot make those elements plain, it is not yet helping the team govern acceptance boundary; it is simply producing more documentation.

A useful review also names the decision that follows from the evidence. If acceptance boundary is adequate, the organization may extend the commitment. If it is not, the team should be able to narrow the scope, add discovery, change the intervention, or hold the work without treating any of those choices as failure. This keeps governance attached to a real choice instead of turning it into commentary on work that will continue regardless.

Leaders should resist the temptation to convert this test into a universal score. The relevant question is whether the available evidence is sufficient for this particular next step, given the consequence of error and the reversibility of the commitment. A small, reversible trial can proceed with less evidence than an enterprise-wide rollout, a material dependency, or an automated action that is difficult to undo.

This distinction matters because a portfolio is not improved by making every initiative look equally governed. It is improved when leaders can see which commitments have earned confidence, which still need evidence,

and where limited attention should be applied next. The record makes that judgment repeatable without pretending that the judgment can be automated.

The habit can begin with a short standing question in the normal management cadence: what would we need to see before we extend reliance on acceptance boundary? The answer should be recorded where the people doing the work can find it. That modest discipline reduces retrospective debate because it makes the original basis for the decision available when conditions change.

Production signals is not a label to add to a project charter. It is a commitment with an owner, a boundary, and evidence that a sponsor can examine. The point is to turn an assumption into an explicit condition of work.

In practice, teams should ask what would make production signals fail under ordinary operating pressure: a competing priority, a missing input, an ambiguous decision, a changed dependency, or an exception that no one owns. The answer identifies the smallest control required before reliance expands.

The standard is deliberately practical. A team can narrow the commitment, stage discovery, add a named review, secure a decision, or stop the work. Each is better than continuing under a premise that the organization has not actually accepted.

The first design question is ownership. Someone must be able to say whether production signals is adequate for the next commitment, and that person needs information early enough to act. Naming an owner is not bureaucracy; it prevents a consequential judgment from being distributed so widely that nobody can make it.

The second design question is evidence. Evidence need not be a large report. It may be a decision record, a sample of real workflow cases, a capacity confirmation, a stakeholder commitment, or a documented test. The useful standard is whether a skeptical sponsor can understand why production signals is credible and what remains uncertain.

The third design question is boundary. Teams need to know what production signals does not authorize, what condition would cause it to be revisited, and which escalation path applies when the boundary is crossed. A visible boundary turns surprise into a managed exception rather than an informal renegotiation.

The fourth design question is cadence. Operating conditions change after a decision. A lightweight recurring review asks whether production signals still fits the work, whether the original evidence still applies, and whether reliance should expand, remain bounded, or be reduced. That is how a framework becomes an operating control.

The final design question is learning. Recurring problems around production signals should improve the next decision rather than become folklore about a difficult project. Capturing the pattern gives portfolio leaders better questions, more realistic commitments, and a way to distinguish a one-off event from a weakness in the operating model.

This is also where proportionality matters. A low-consequence decision can use a short record and a rapid review. A decision that changes a recurring workflow, exposes sensitive information, consumes scarce

capacity, or creates a difficult-to-reverse dependency deserves more evidence and a clearer escalation path. The framework scales the control to the commitment rather than applying one ritual to every kind of work.

A useful operating illustration is a team that appears ready until an ordinary case falls outside the happy path. The work then pauses while people discover who may decide, what outcome is actually protected, whose participation is needed, or what capacity is available. Designing production signals in advance makes that ordinary case visible before it becomes expensive.

The portfolio implication is cumulative. One unresolved condition may be tolerable; many unresolved conditions across active initiatives create a hidden demand on senior attention and specialist capacity. Making production signals explicit allows leaders to decide which uncertainty deserves investment, which should remain bounded, and which initiative should wait.

The practical record should be concise enough to be used. It should identify the next commitment, the owner, the evidence examined, the boundary of reliance, and the next review date. If the record cannot make those elements plain, it is not yet helping the team govern production signals; it is simply producing more documentation.

A useful review also names the decision that follows from the evidence. If production signals is adequate, the organization may extend the commitment. If it is not, the team should be able to narrow the scope, add discovery, change the intervention, or hold the work without treating any of those choices as failure. This keeps governance attached to a real choice instead of turning it into commentary on work that will continue regardless.

Leaders should resist the temptation to convert this test into a universal score. The relevant question is whether the available evidence is sufficient for this particular next step, given the consequence of error and the reversibility of the commitment. A small, reversible trial can proceed with less evidence than an enterprise-wide rollout, a material dependency, or an automated action that is difficult to undo.

This distinction matters because a portfolio is not improved by making every initiative look equally governed. It is improved when leaders can see which commitments have earned confidence, which still need evidence, and where limited attention should be applied next. The record makes that judgment repeatable without pretending that the judgment can be automated.

The habit can begin with a short standing question in the normal management cadence: what would we need to see before we extend reliance on production signals? The answer should be recorded where the people doing the work can find it. That modest discipline reduces retrospective debate because it makes the original basis for the decision available when conditions change.

Review and deactivation is not a label to add to a project charter. It is a commitment with an owner, a boundary, and evidence that a sponsor can examine. The point is to turn an assumption into an explicit condition of work.

In practice, teams should ask what would make review and deactivation fail under ordinary operating pressure: a competing priority, a missing input, an ambiguous decision, a changed dependency, or an

exception that no one owns. The answer identifies the smallest control required before reliance expands.

The standard is deliberately practical. A team can narrow the commitment, stage discovery, add a named review, secure a decision, or stop the work. Each is better than continuing under a premise that the organization has not actually accepted.

The first design question is ownership. Someone must be able to say whether review and deactivation is adequate for the next commitment, and that person needs information early enough to act. Naming an owner is not bureaucracy; it prevents a consequential judgment from being distributed so widely that nobody can make it.

The second design question is evidence. Evidence need not be a large report. It may be a decision record, a sample of real workflow cases, a capacity confirmation, a stakeholder commitment, or a documented test. The useful standard is whether a skeptical sponsor can understand why review and deactivation is credible and what remains uncertain.

The third design question is boundary. Teams need to know what review and deactivation does not authorize, what condition would cause it to be revisited, and which escalation path applies when the boundary is crossed. A visible boundary turns surprise into a managed exception rather than an informal renegotiation.

The fourth design question is cadence. Operating conditions change after a decision. A lightweight recurring review asks whether review and deactivation still fits the work, whether the original evidence still applies, and whether reliance should expand, remain bounded, or be reduced. That is how a framework becomes an operating control.

The final design question is learning. Recurring problems around review and deactivation should improve the next decision rather than become folklore about a difficult project. Capturing the pattern gives portfolio leaders better questions, more realistic commitments, and a way to distinguish a one-off event from a weakness in the operating model.

This is also where proportionality matters. A low-consequence decision can use a short record and a rapid review. A decision that changes a recurring workflow, exposes sensitive information, consumes scarce capacity, or creates a difficult-to-reverse dependency deserves more evidence and a clearer escalation path. The framework scales the control to the commitment rather than applying one ritual to every kind of work.

A useful operating illustration is a team that appears ready until an ordinary case falls outside the happy path. The work then pauses while people discover who may decide, what outcome is actually protected, whose participation is needed, or what capacity is available. Designing review and deactivation in advance makes that ordinary case visible before it becomes expensive.

The portfolio implication is cumulative. One unresolved condition may be tolerable; many unresolved conditions across active initiatives create a hidden demand on senior attention and specialist capacity. Making review and deactivation explicit allows leaders to decide which uncertainty deserves investment, which should remain bounded, and which initiative should wait.

The practical record should be concise enough to be used. It should identify the next commitment, the owner, the evidence examined, the boundary of reliance, and the next review date. If the record cannot make those elements plain, it is not yet helping the team govern review and deactivation; it is simply producing more documentation.

A useful review also names the decision that follows from the evidence. If review and deactivation is adequate, the organization may extend the commitment. If it is not, the team should be able to narrow the scope, add discovery, change the intervention, or hold the work without treating any of those choices as failure. This keeps governance attached to a real choice instead of turning it into commentary on work that will continue regardless.

Leaders should resist the temptation to convert this test into a universal score. The relevant question is whether the available evidence is sufficient for this particular next step, given the consequence of error and the reversibility of the commitment. A small, reversible trial can proceed with less evidence than an enterprise-wide rollout, a material dependency, or an automated action that is difficult to undo.

This distinction matters because a portfolio is not improved by making every initiative look equally governed. It is improved when leaders can see which commitments have earned confidence, which still need evidence, and where limited attention should be applied next. The record makes that judgment repeatable without pretending that the judgment can be automated.

The habit can begin with a short standing question in the normal management cadence: what would we need to see before we extend reliance on review and deactivation? The answer should be recorded where the people doing the work can find it. That modest discipline reduces retrospective debate because it makes the original basis for the decision available when conditions change.

The framework becomes useful when it changes a real decision. A recurring review should ask what has changed, which assumption is now visible, who owns the next move, and what evidence would justify a larger commitment. That is governance as operating work rather than retrospective reporting.

There is a counterargument worth taking seriously: additional controls can become ceremony. The answer is not to remove every control. It is to use the lightest control that makes the relevant decision, boundary, and accountability clear. Reversible work deserves a lighter standard than consequential work.

The shift is straightforward. An AI Evaluation Is a Production Operating Control asks leaders to govern the conditions of reliance, not merely the activity around an initiative. When those conditions are explicit, teams can move faster with less hidden risk. When they are not, apparent momentum is only a more organized form of uncertainty.

The discipline is deliberately modest. It does not promise that uncertainty disappears or that a single framework will resolve competing interests. It gives the people responsible for the next commitment a common way to say what is known, what is not yet known, and what decision is justified now. That is enough to turn a hidden operating risk into work leaders can actually govern.

Notes and sources

[1] NIST, "Artificial Intelligence Risk Management Framework: Generative Artificial Intelligence Profile," July 2024. <https://nvlpubs.nist.gov/nistpubs/ai/nist.ai.600-1.pdf>

[2] Blake Bullwinkel et al., "Lessons From Red Teaming 100 Generative AI Products," Microsoft Research, January 2025. <https://www.microsoft.com/en-us/research/publication/lessons-from-red-teaming-100-generative-ai-products/>

About the author

Marco Policani is an enterprise portfolio, PMO, and AI operating-governance leader. He builds portfolio governance systems where AI carries the analytical load and named humans own every decision — an approach documented across case studies, walkthroughs, and working governance guides on his portfolio site.

Portfolio & governance library: policani.net/governance · Full portfolio: policani.net · LinkedIn: linkedin.com/in/marcpolicani

This white paper may be shared freely with attribution. © 2026 Marco Policani.