14

WHITE PAPER · AI OPERATING GOVERNANCE

An AI Evaluation Is a Production Operating Control

Govern changing workflow behavior, not a one-time test

Marco Policani

Enterprise Portfolio · PMO · AI Operating Governance

policani.net · linkedin.com/in/marcpolicani

July 2026

EXECUTIVE SUMMARY

Evaluation continues in production because workflow behavior, context, and reliance continue to change.

Benchmarking, red teaming, and production evaluation answer different questions. Production evaluation is a living control that tests representative workflow cases, monitors change, converts incidents into cases, and changes permitted reliance when evidence changes. This paper gives AI product leaders, workflow owners, risk teams, and technology executives a practical the production evaluation loop for making the next commitment explicit.

01 Separate three evaluation jobs: Benchmarks compare capability, red teams search for weaknesses, and operating evaluations test a specific workflow under its own consequences.

02 Build representative workflow cases: The case set should reflect ordinary work, important subgroups, exceptions, difficult inputs, and foreseeable misuse.

03 Define thresholds by consequence: Average quality can hide severe failures; thresholds should combine performance, severity, and control response.

04 Monitor production signals: Evaluation continues through drift, exceptions, overrides, complaints, incidents, and changes in user behavior.

CONTENTS

The argument
Why the conventional approach breaks
The production evaluation loop
Separate three evaluation jobs
Build representative workflow cases
Define thresholds by consequence
Monitor production signals
Turn change into a test event
Learn from incidents and near misses
Assign authority to change reliance
How the elements interact
Put the model into the management cadence
Measures that reveal whether the control works
The strongest counterargument
The shift

The argument

Benchmarking, red teaming, and production evaluation answer different questions. Production evaluation is a living control that tests representative workflow cases, monitors change, converts incidents into cases, and changes permitted reliance when evidence changes.

A model can pass an evaluation and still fail the operation that adopts it. The test may use the wrong cases, the workflow may change, users may rely on output differently, retrieval data may drift, or a vendor update may alter behavior. The evaluation was not necessarily useless; it was insufficiently connected to production decisions.

NIST's generative-AI profile and Microsoft's red-team lessons support systematic testing and learning [1][2]. The operating challenge is to keep that work alive after launch and tie it to authority.

This paper is written for AI product leaders, workflow owners, risk teams, and technology executives. It offers the production evaluation loop as a management instrument: a way to connect operating evidence to the next consequential choice. The cited sources provide external context; the framework and its application are Marco Policani's synthesis. It is not a predictive score and does not replace professional, legal, security, financial, or domain judgment.

Why the conventional approach breaks

The distinction matters because leaders authorize commitments, while teams inherit whatever the authorization left unresolved. Conventional governance tends to separate planning, approval, delivery reporting, and operating ownership. That separation allows a premise to pass through several forums without any one forum owning the decision it implies. The project or workflow then carries an assumption that was acknowledged socially but never accepted operationally.

Reporting usually captures the symptom after it has been translated into schedule, cost, quality, or risk. It rarely captures the original unresolved choice. The answer is not heavier control at every step. It is to place a clear, proportionate test at the point where the organization increases money, capacity, access, dependency, or action authority.

The production evaluation loop

The model has 7 connected elements. Each makes a different condition visible, but they should be reviewed together because strength in one area does not cancel a missing obligation in another. The model is designed for a short executive conversation supported by inspectable evidence, not a long maturity assessment.

EXHIBIT 1

The production evaluation loop connects evidence to action

1. **Separate three evaluation jobs**
   Benchmarks compare capability, red teams search for weaknesses, and operating evaluations test a specific workflow under its own consequences.
2. **Build representative workflow cases**
   The case set should reflect ordinary work, important subgroups, exceptions, difficult inputs, and foreseeable misuse.
3. **Define thresholds by consequence**
   Average quality can hide severe failures; thresholds should combine performance, severity, and control response.

4. **Monitor production signals**
   Evaluation continues through drift, exceptions, overrides, complaints, incidents, and changes in user behavior.
5. **Turn change into a test event**
   Model, prompt, retrieval, policy, tool, user, and workflow changes can invalidate prior evidence.
6. **Learn from incidents and near misses**
   Production failures should expand the case set and improve boundaries, not end with a local correction.
7. **Assign authority to change reliance**
   Evaluation matters only when someone can narrow, pause, restore, or retire use.

Separate three evaluation jobs

Benchmarks compare capability, red teams search for weaknesses, and operating evaluations test a specific workflow under its own consequences. Teams move faster when uncertainty is bounded honestly than when it is concealed inside an optimistic plan.

A strong benchmark or pre-launch red team is cited as proof that every production use is acceptable. Under pressure, the workaround is predictable: capable people compensate, local decisions proliferate, and the formal record stops describing how the operation actually runs.

Document which question each test answers and what decision it is allowed to support. The design should state what remains unknown. Acknowledged uncertainty can be bounded; disguised uncertainty is inherited by the operation.

Test purpose, population, environment, limitations, owner, and explicit link to a reliance decision. Where consequence is high, independent checking or cross-functional challenge reduces the risk that advocacy becomes its own proof.

Commission the missing test rather than stretching one result beyond its scope. Escalation is appropriate when authority is missing, not as a routine substitute for clear delegation.

A tabletop review of separate three evaluation jobs should use the observed break as its scenario: A strong benchmark or pre-launch red team is cited as proof that every production use is acceptable. Participants then apply the proposed response. They document which question each test answers and what decision it is allowed to support. The review identifies where authority, information, time, or recovery would run out. It is useful only if it changes the boundary or confirms who will act when the condition appears.

The minimum record for separate three evaluation jobs is anchored in actual proof. It includes test purpose, population, environment, limitations, owner, and explicit link to a reliance decision. The record also carries the choice supported by that proof: Commission the missing test rather than stretching one result beyond its scope. Keeping evidence and decision together lets a later owner distinguish a deliberate residual risk from a premise that nobody examined. It also gives the team a clear reason to reopen the choice when the underlying facts change.

Across the work governed by AI product leaders, workflow owners, risk teams, and technology executives, occurrences of this pattern should be compared rather than treated as unrelated project stories. The positive standard is: Benchmarks compare capability, red teams search for weaknesses, and operating evaluations test a specific workflow under its own consequences. If the same failure recurs, leaders can change delegation, capacity, intake, policy, or the intervention itself. That portfolio response is usually more valuable than asking each team to absorb another local workaround.

Two questions test separate three evaluation jobs. What evidence would show that the condition is no longer adequate? Would that evidence arrive early enough to preserve a feasible choice? The intended choice is clear: Commission the missing test rather than stretching one result beyond its scope. A weak answer to either question calls for a smaller commitment, an earlier signal, or a safer fallback. A strong answer earns movement without claiming that uncertainty has disappeared.

Build representative workflow cases

The case set should reflect ordinary work, important subgroups, exceptions, difficult inputs, and foreseeable misuse. The practical failure occurs between the formal approval and the ordinary pressure of running the work.

Teams evaluate on convenient examples supplied by builders, excluding rare cases that create most of the risk or effort. The happy path can survive this weakness. Ordinary variation exposes it: a disputed input, absent owner, competing commitment, or case that falls outside the assumed rule.

Sample production history, interview domain experts, include edge and adversarial conditions, and record why coverage is sufficient. The smallest credible control is the one that changes behavior when the evidence is weak. Anything else is commentary.

Case provenance, frequency, consequence, subgroup coverage, expected response, and uncertainty about the reference answer. Qualitative observation belongs beside quantitative measures when workflow behavior, adoption, or judgment explains why the number moved.

Limit permitted use to the population the case set can reasonably represent. A stop or rollback is evidence that the control works when it protects greater value elsewhere in the portfolio.

A tabletop review of build representative workflow cases should use the observed break as its scenario: Teams evaluate on convenient examples supplied by builders, excluding rare cases that create most of the risk or effort. Participants then apply the proposed response. They sample production history, interview domain experts, include edge and adversarial conditions, and record why coverage is sufficient. The review identifies where authority, information, time, or recovery would run out. It is useful only if it changes the boundary or confirms who will act when the condition appears.

The minimum record for build representative workflow cases is anchored in actual proof. It includes case provenance, frequency, consequence, subgroup coverage, expected response, and uncertainty about the reference answer. The record also carries the choice supported by that proof: Limit permitted use to the population the case set can reasonably represent. Keeping evidence and decision together lets a later owner distinguish a deliberate residual risk from a premise that nobody examined. It also gives the team a clear reason to reopen the choice when the underlying facts change.

Across the work governed by AI product leaders, workflow owners, risk teams, and technology executives, occurrences of this pattern should be compared rather than treated as unrelated project stories. The positive standard is: The case set should reflect ordinary work, important subgroups, exceptions, difficult inputs, and foreseeable misuse. If the same failure recurs, leaders can change delegation, capacity, intake, policy, or the intervention itself. That portfolio response is usually more valuable than asking each team to absorb another local workaround.

Two questions test build representative workflow cases. What evidence would show that the condition is no longer adequate? Would that evidence arrive early enough to preserve a feasible choice? The intended choice is clear: Limit permitted use to the population the case set can reasonably represent. A weak answer to either question calls for a smaller commitment, an earlier signal, or a safer fallback. A strong answer earns movement without claiming that uncertainty has disappeared.

Define thresholds by consequence

Average quality can hide severe failures; thresholds should combine performance, severity, and control response. The visible artifact is often complete before the operating condition behind it is credible.

A single accuracy number allows material harms in a small class to disappear inside aggregate success. The organization may call this execution risk, yet the failure began when authority and evidence were separated at the decision point.

Set class-specific acceptance, abstention, escalation, latency, and recovery requirements before results are known. A useful test asks what ordinary operating pressure would break the premise and how that break would be detected before consequence spreads.

Pass rates by case class, severe failure counts, false acceptance and rejection, review load, and residual exposure. Evidence should arrive before the latest useful decision date; perfect analysis delivered after options expire has little governance value.

Pause or narrow reliance when a critical threshold fails even if the aggregate average looks strong. The operating owner should confirm that the decision can be executed with the people, systems, and time actually available.

A tabletop review of define thresholds by consequence should use the observed break as its scenario: A single accuracy number allows material harms in a small class to disappear inside aggregate success. Participants then apply the proposed response. They set class-specific acceptance, abstention, escalation, latency, and recovery requirements before results are known. The review identifies where authority, information, time, or recovery would run out. It is useful only if it changes the boundary or confirms who will act when the condition appears.

The minimum record for define thresholds by consequence is anchored in actual proof. It includes pass rates by case class, severe failure counts, false acceptance and rejection, review load, and residual exposure. The record also carries the choice supported by that proof: Pause or narrow reliance when a critical threshold fails even if the aggregate average looks strong. Keeping evidence and decision together lets a later owner distinguish a deliberate residual risk from a premise that nobody examined. It also gives the team a clear reason to reopen the choice when the underlying facts change.

Across the work governed by AI product leaders, workflow owners, risk teams, and technology executives, occurrences of this pattern should be compared rather than treated as unrelated project stories. The positive standard is: Average quality can hide severe failures; thresholds should combine performance, severity, and control response. If the same failure recurs, leaders can change delegation, capacity, intake, policy, or the intervention itself. That portfolio response is usually more valuable than asking each team to absorb another local workaround.

Two questions test define thresholds by consequence. What evidence would show that the condition is no longer adequate? Would that evidence arrive early enough to preserve a feasible choice? The intended choice is clear: Pause or narrow reliance when a critical threshold fails even if the aggregate average looks strong. A weak answer to either question calls for a smaller commitment, an earlier signal, or a safer fallback. A strong answer earns movement without claiming that uncertainty has disappeared.

Monitor production signals

Evaluation continues through drift, exceptions, overrides, complaints, incidents, and changes in user behavior. The issue is not a shortage of activity; it is a shortage of inspectable commitments.

The organization monitors uptime and cost but cannot tell whether answers or actions remain fit for the workflow. The missing condition is often treated as something the team will resolve later, even though the team lacks the authority or capacity to resolve it.

Instrument inputs, outputs, review dispositions, overrides, failures, data shifts, and operational outcomes with appropriate privacy controls. The forum should be able to distinguish an acceptable residual risk from a missing obligation that has merely been renamed.

Trend and distribution changes, sampled case reviews, unresolved exceptions, reviewer disagreement, and outcome movement. The evidence may be compact, but it must be inspectable. A future reviewer should be able to see what was examined, what conclusion was drawn, and which limitation remained.

Trigger focused evaluation when signals move beyond expected variation. Leaders should prefer the smallest commitment that preserves learning and limits irreversible exposure.

A tabletop review of monitor production signals should use the observed break as its scenario: The organization monitors uptime and cost but cannot tell whether answers or actions remain fit for the workflow. Participants then apply the proposed response. They instrument inputs, outputs, review dispositions, overrides, failures, data shifts, and operational outcomes with appropriate privacy controls. The review identifies where authority, information, time, or recovery would run out. It is useful only if it changes the boundary or confirms who will act when the condition appears.

The minimum record for monitor production signals is anchored in actual proof. It includes trend and distribution changes, sampled case reviews, unresolved exceptions, reviewer disagreement, and outcome movement. The record also carries the choice supported by that proof: Trigger focused evaluation when signals move beyond expected variation. Keeping evidence and decision together lets a later owner distinguish a deliberate residual risk from a premise that nobody examined. It also gives the team a clear reason to reopen the choice when the underlying facts change.

Across the work governed by AI product leaders, workflow owners, risk teams, and technology executives, occurrences of this pattern should be compared rather than treated as unrelated project stories. The positive standard is: Evaluation continues through drift, exceptions, overrides, complaints, incidents, and changes in user behavior. If the same failure recurs, leaders can change delegation, capacity, intake, policy, or the intervention itself. That portfolio response is usually more valuable than asking each team to absorb another local workaround.

Two questions test monitor production signals. What evidence would show that the condition is no longer adequate? Would that evidence arrive early enough to preserve a feasible choice? The intended choice is clear: Trigger focused evaluation when signals move beyond expected variation. A weak answer to either question calls for a smaller commitment, an earlier signal, or a safer fallback. A strong answer earns movement without claiming that uncertainty has disappeared.

Turn change into a test event

Model, prompt, retrieval, policy, tool, user, and workflow changes can invalidate prior evidence. The cost appears downstream, but the missing decision sits upstream.

Updates are treated as technical maintenance and deployed without checking the cases that justified reliance. The recurring signal is translation work: people spend time discovering what was meant, who may decide, and how to repair the unsupported assumption.

Classify changes, define regression scope, preserve version history, and require approval proportional to affected consequence. This does not require a larger pack. It requires a traceable connection among the premise, the evidence, the boundary, and the action that follows.

Before-and-after results, changed cases, release owner, rollback readiness, and post-change observation. Evidence has a shelf life. The record needs the conditions or date that would make the conclusion stale.

Withhold or reduce deployment when regression evidence is incomplete. A hold is governable when the return conditions are explicit; otherwise it becomes an unmanaged queue.

A tabletop review of turn change into a test event should use the observed break as its scenario: Updates are treated as technical maintenance and deployed without checking the cases that justified reliance. Participants then apply the proposed response. They classify changes, define regression scope, preserve version history, and require approval proportional to affected consequence. The review identifies where authority, information, time, or recovery would run out. It is useful only if it changes the boundary or confirms who will act when the condition appears.

The minimum record for turn change into a test event is anchored in actual proof. It includes before-and-after results, changed cases, release owner, rollback readiness, and post-change observation. The record also carries the choice supported by that proof: Withhold or reduce deployment when regression evidence is incomplete. Keeping evidence and decision together lets a later owner distinguish a deliberate residual risk from a premise that nobody examined. It also gives the team a clear reason to reopen the choice when the underlying facts change.

Across the work governed by AI product leaders, workflow owners, risk teams, and technology executives, occurrences of this pattern should be compared rather than treated as unrelated project stories. The positive standard is: Model, prompt, retrieval, policy, tool, user, and workflow changes can invalidate prior evidence. If the same failure recurs, leaders can change delegation, capacity, intake, policy, or the intervention itself. That portfolio response is usually more valuable than asking each team to absorb another local workaround.

Two questions test turn change into a test event. What evidence would show that the condition is no longer adequate? Would that evidence arrive early enough to preserve a feasible choice? The intended choice is clear: Withhold or reduce deployment when regression evidence is incomplete. A weak answer to either question calls for a smaller commitment, an earlier signal, or a safer fallback. A strong answer earns movement without claiming that uncertainty has disappeared.

Learn from incidents and near misses

Production failures should expand the case set and improve boundaries, not end with a local correction. The distinction matters because leaders authorize commitments, while teams inherit whatever the authorization left unresolved.

Teams fix one prompt or record while the failure mechanism remains possible elsewhere. What looks like a team-level miss is frequently a portfolio choice that nobody made explicitly. Scarce attention is then spent reconciling consequences rather than choosing among alternatives.

Perform a short causal review across model, data, interface, instructions, human review, permissions, and operating pressure. A named owner must be able to change the work; receiving a notification after the boundary is crossed is insufficient.

New regression cases, control changes, recurrence monitoring, responsible owner, and evidence of effectiveness. The evidence owner is accountable for integrity, while the decision owner remains accountable for the commitment. Keeping those roles distinct improves challenge.

Change permitted use when the incident reveals a class of exposure larger than the original case. The residual uncertainty belongs in the decision record so later outcomes are compared with what leaders actually knew at the time.

A tabletop review of learn from incidents and near misses should use the observed break as its scenario: Teams fix one prompt or record while the failure mechanism remains possible elsewhere. Participants then apply the proposed response. They perform a short causal review across model, data, interface, instructions, human review, permissions, and operating pressure. The review identifies where authority, information, time, or recovery would run out. It is useful only if it changes the boundary or confirms who will act when the condition appears.

The minimum record for learn from incidents and near misses is anchored in actual proof. It includes new regression cases, control changes, recurrence monitoring, responsible owner, and evidence of effectiveness. The record also carries the choice supported by that proof: Change permitted use when the incident reveals a class of exposure larger than the original case. Keeping evidence and decision together lets a later owner distinguish a deliberate residual risk from a premise that nobody examined. It also gives the team a clear reason to reopen the choice when the underlying facts change.

Across the work governed by AI product leaders, workflow owners, risk teams, and technology executives, occurrences of this pattern should be compared rather than treated as unrelated project stories. The positive standard is: Production failures should expand the case set and improve boundaries, not end with a local correction. If the same failure recurs, leaders can change delegation, capacity, intake, policy, or the intervention itself. That portfolio response is usually more valuable than asking each team to absorb another local workaround.

Two questions test learn from incidents and near misses. What evidence would show that the condition is no longer adequate? Would that evidence arrive early enough to preserve a feasible choice? The intended choice is clear: Change permitted use when the incident reveals a class of exposure larger than the original case. A weak answer to either question calls for a smaller commitment, an earlier signal, or a safer fallback. A strong answer earns movement without claiming that uncertainty has disappeared.

Assign authority to change reliance

Evaluation matters only when someone can narrow, pause, restore, or retire use. The control belongs at the point where reliance expands, not in a retrospective explanation after the outcome is known.

Analysts publish results while product momentum and adoption continue independently. Momentum makes the weakness self-reinforcing. Each additional commitment raises the political and economic cost of revisiting the premise.

Name the workflow owner, technical owner, risk challenger, forum, thresholds, and emergency stop route. The record should travel with the work so that later participants can see the basis for reliance instead of reconstructing it from memory.

Recorded decisions, time from trigger to action, exceptions to threshold policy, and evidence required for restoration. The strongest record shows both what supports the choice and what could overturn it.

Treat evaluation as an operating control, not a research report. Decision quality improves when dissent and competing evidence are preserved instead of flattened into a unanimous-looking status.

A tabletop review of assign authority to change reliance should use the observed break as its scenario: Analysts publish results while product momentum and adoption continue independently. Participants then apply the proposed response. They name the workflow owner, technical owner, risk challenger, forum, thresholds, and emergency stop route. The review identifies where authority, information, time, or recovery would run out. It is useful only if it changes the boundary or confirms who will act when the condition appears.

The minimum record for assign authority to change reliance is anchored in actual proof. It includes recorded decisions, time from trigger to action, exceptions to threshold policy, and evidence required for restoration. The record also carries the choice supported by that proof: Treat evaluation as an operating control, not a research report. Keeping evidence and decision together lets a later owner distinguish a deliberate residual risk from a premise that nobody examined. It also gives the team a clear reason to reopen the choice when the underlying facts change.

Across the work governed by AI product leaders, workflow owners, risk teams, and technology executives, occurrences of this pattern should be compared rather than treated as unrelated project stories. The positive standard is: Evaluation matters only when someone can narrow, pause, restore, or retire use. If the same failure recurs, leaders can change delegation, capacity, intake, policy, or the intervention itself. That portfolio response is usually more valuable than asking each team to absorb another local workaround.

Two questions test assign authority to change reliance. What evidence would show that the condition is no longer adequate? Would that evidence arrive early enough to preserve a feasible choice? The intended choice is clear: Treat evaluation as an operating control, not a research report. A weak answer to either question calls for a smaller commitment, an earlier signal, or a safer fallback. A strong answer earns movement without claiming that uncertainty has disappeared.

How the elements interact

Separate three evaluation jobs and Build representative workflow cases protect different parts of the same commitment. Benchmarks compare capability, red teams search for weaknesses, and operating evaluations test a specific workflow under its own consequences. The case set should reflect ordinary work, important subgroups, exceptions, difficult inputs, and foreseeable misuse. Viewed from the portfolio, this is a resource-allocation problem before it becomes a delivery problem. If the operation encounters this failure—a strong benchmark or pre-launch red team is cited as proof that every production use is acceptable.—the response for build representative workflow cases cannot compensate for the missing separate three evaluation jobs obligation. Reviewers should reconcile the evidence for both elements before they accept the next action: limit permitted use to the population the case set can reasonably represent.

Build representative workflow cases and Define thresholds by consequence protect different parts of the same commitment. The case set should reflect ordinary work, important subgroups, exceptions, difficult inputs, and foreseeable misuse. Average quality can hide severe failures; thresholds should combine performance, severity, and control response. The control belongs at the point where reliance expands, not in a retrospective explanation after the outcome is known. If the operation encounters this failure—teams evaluate on convenient examples supplied by builders, excluding rare cases that create most of the risk or effort.—the response for define thresholds by consequence cannot compensate for the missing build representative workflow cases obligation. Reviewers should reconcile the evidence for both elements before they accept the next action: pause or narrow reliance when a critical threshold fails even if the aggregate average looks strong.

Define thresholds by consequence and Monitor production signals protect different parts of the same commitment. Average quality can hide severe failures; thresholds should combine performance, severity, and control response. Evaluation continues through drift, exceptions, overrides, complaints, incidents, and changes in user behavior. The visible artifact is often complete before the operating condition behind it is credible. If the operation encounters this failure—a single accuracy number allows material harms in a small class to disappear inside aggregate success.—the response for monitor production signals cannot compensate for the missing define thresholds by consequence obligation. Reviewers should reconcile the evidence for both elements before they accept the next action: trigger focused evaluation when signals move beyond expected variation.

Monitor production signals and Turn change into a test event protect different parts of the same commitment. Evaluation continues through drift, exceptions, overrides, complaints, incidents, and changes in user behavior. Model, prompt, retrieval, policy, tool, user, and workflow changes can invalidate prior evidence. A useful governance design therefore begins with the consequence leaders are prepared to accept. If the operation encounters this failure—the organization monitors uptime and cost but cannot tell whether answers or actions remain fit for the workflow.—the response for turn change into a test event cannot compensate for the missing monitor production signals obligation. Reviewers should reconcile the evidence for both elements before they accept the next action: withhold or reduce deployment when regression evidence is incomplete.

Turn change into a test event and Learn from incidents and near misses protect different parts of the same commitment. Model, prompt, retrieval, policy, tool, user, and workflow changes can invalidate prior evidence. Production failures should expand the case set and improve boundaries, not end with a local correction. This is where a management ritual either becomes a real control or reveals itself as administrative theater. If the operation encounters this failure—updates are treated as technical maintenance and deployed without checking the cases that justified reliance.—the response for learn from incidents and near misses cannot compensate for the missing turn change into a test event obligation. Reviewers should reconcile the evidence for both elements before they accept the next action: change permitted use when the incident reveals a class of exposure larger than the original case.

Learn from incidents and near misses and Assign authority to change reliance protect different parts of the same commitment. Production failures should expand the case set and improve boundaries, not end with a local correction. Evaluation matters only when someone can narrow, pause, restore, or retire use. The issue is not a shortage of activity; it is a shortage of inspectable commitments. If the operation encounters this failure—teams fix one prompt or record while the failure mechanism remains possible elsewhere.—the response for assign authority to change reliance cannot compensate for the missing learn from incidents and near misses obligation. Reviewers should reconcile the evidence for both elements before they accept the next action: treat evaluation as an operating control, not a research report.

Assign authority to change reliance and Separate three evaluation jobs protect different parts of the same commitment. Evaluation matters only when someone can narrow, pause, restore, or retire use. Benchmarks compare capability, red teams search for weaknesses, and operating evaluations test a specific workflow under its own consequences. A strong operating model makes the hidden negotiation explicit while options are still available. If the operation encounters this failure—analysts publish results while product momentum and adoption continue independently.—the response for separate three evaluation jobs cannot compensate for the missing assign authority to change reliance obligation. Reviewers should reconcile the evidence for both elements before they accept the next action: commission the missing test rather than stretching one result beyond its scope.

Put the model into the management cadence

A framework becomes useful when it changes the normal cadence of work. It should shape intake, funding, design, delivery, and operating review without creating a separate governance industry around itself. The following sequence is deliberately small:

EXHIBIT 2

Start with one decision, then calibrate from evidence

1. **Move 1**
   Create the first case set from real workflow history, not a generic benchmark.
2. **Move 2**
   Write acceptance and stop thresholds before running the evaluation.
3. **Move 3**
   Connect production telemetry and reviewer decisions to a maintained evaluation backlog.
4. **Move 4**
   Run regression after material change and a scheduled review even when no incident occurs.

- Create the first case set from real workflow history, not a generic benchmark.
- Write acceptance and stop thresholds before running the evaluation.
- Connect production telemetry and reviewer decisions to a maintained evaluation backlog.

- Run regression after material change and a scheduled review even when no incident occurs.

The first cycle should be treated as calibration. Reviewers should note which questions changed a decision, which evidence was unavailable, where authority was unclear, and which controls cost out of proportion to the exposure. That record improves the production evaluation loop without pretending the first version will fit every workflow or investment.

Ownership should remain close to the consequence. The PMO or AI-governance function can maintain the method, surface cross-portfolio patterns, and challenge weak evidence. It should not absorb the business owner's accountability for the result or the executive's accountability for the commitment.

Measures that reveal whether the control works

- Commitments narrowed, redirected, paused, or stopped before avoidable exposure grew.
- Exceptions or assumptions discovered after the latest useful decision date.
- Scarce specialist and executive attention consumed by recurring unresolved conditions.
- Decisions reopened because the original evidence, boundary, or rationale was unclear.
- Operating outcomes and recovery performance after reliance expanded.
- Time from a material signal to a consequential decision.

These measures are diagnostic, not a universal scorecard. Their purpose is to identify where the operating model fails repeatedly and whether governance is preserving options earlier. A lower count can indicate improvement, but it can also indicate weak detection. Leaders should read the measures with case evidence and frontline observation.

The strongest counterargument

No evaluation can prove universal safety or future performance. It provides bounded evidence for a named use under observed conditions. Honest limits, maintained cases, and authority to change reliance are more valuable than a large one-time score.

A reasonable critic could also argue that experienced leaders already do this through judgment. Many do. The weakness is not the absence of judgment; it is the absence of a durable operating record when people, pressure, and conditions change. The production evaluation loop makes the basis for judgment visible enough to challenge, execute, and revisit.

The shift

Benchmarking, red teaming, and production evaluation answer different questions. Production evaluation is a living control that tests representative workflow cases, monitors change, converts incidents into cases, and changes permitted reliance when evidence changes. The practical shift is from reporting activity to governing the conditions under which the organization relies on that activity.

That discipline does not make uncertainty disappear. It gives leaders a better place to stand: a named decision, evidence proportionate to consequence, an explicit boundary, a responsible owner, and a response when reality departs from the assumption. Those elements are enough to turn hidden operating risk into a choice the portfolio can make deliberately.

Notes and sources

[1] NIST, "Artificial Intelligence Risk Management Framework: Generative Artificial Intelligence Profile," July 2024. https://nvlpubs.nist.gov/nistpubs/ai/nist.ai.600-1.pdf

[2] Blake Bullwinkel et al., "Lessons From Red Teaming 100 Generative AI Products," Microsoft Research, January 2025. https://www.microsoft.com/en-us/research/publication/lessons-from-red-teaming-100-generative-ai-products/

About the author

Marco Policani is an enterprise portfolio, PMO, and AI operating-governance leader. He builds portfolio governance systems where AI carries the analytical load and named humans own every decision — an approach documented across case studies, walkthroughs, and working governance guides on his portfolio site.

Portfolio & governance library: policani.net/governance · Full portfolio: policani.net · LinkedIn: linkedin.com/in/marcpolicani

This white paper may be shared freely with attribution. © 2026 Marco Policani.