


A Project Is Not Ready Because a Plan Exists

Four operating commitments before work earns a start



Marco Policani

Enterprise Portfolio · PMO · AI Operating Governance

policani.net · linkedin.com/in/marcpolicani

July 2026

EXECUTIVE SUMMARY

A plan describes intent. Readiness is whether the organization committed to the conditions the intent depends on.

Projects arrive at kickoff with credible schedules and polished charters and are still unready to start, because documents stand in for commitments nobody made. This paper proposes a four-commitment readiness test – decision authority, outcome and boundary, scheduled participation, named capacity – applied before a start date is granted.

- 01** List the first three or four hard decisions before kickoff and confirm in writing who can make each one.
- 02** Convert stakeholder maps into dated commitments acknowledged by managers. Six names with hours beats forty names on a slide.
- 03** Resist partial credit. The weak commitment decides where the project breaks first; the strong ones only set the momentum.

The argument

Readiness is not the presence of a populated plan. It is an earned decision that the organization can sustain the next commitment without hiding unresolved sponsorship, scope, participation, or capacity inside the delivery team.

I have watched projects arrive at kickoff with a credible schedule, a polished charter, and a room full of supportive leaders — and still be unready to start. What was missing was never another planning artifact. It was a set of operating commitments nobody had actually made. Who would settle the first hard decision? Which outcome was protected when tradeoffs arrived? Which business experts would show up for design and testing? And did the named architect have hours left after the meeting ended?

A kickoff date generates its own momentum. Once invitations go out and a team is announced, asking whether the work should start at all reads as obstruction. That social pressure is exactly why readiness needs a decision standard that operates before the calendar takes over. Once the date exists, the burden of proof flips. The skeptic must now justify delay, when the sponsor should still be justifying the start.

This paper is written for sponsors, portfolio leaders, PMO leaders, and delivery executives. It proposes a four-commitment readiness test: a short set of operating questions a sponsor must answer with evidence before a start date is granted. The test distills my own portfolio and PMO work. The cited research provides outside context on why complex projects miss and how due diligence before commitment pays for itself [1] [2]. It is a management instrument, not a predictive score.

Two boundaries on scope. The test is not a maturity model, and it does not replace stage gates. It governs one moment — the grant of a start date — and assumes a portfolio process on either side: intake that

shapes demand before it, gates that re-examine the commitment after it. Organizations without those structures can still use the test. It simply becomes the first governance the work meets, which is often where its findings are loudest.

How a populated plan passes for readiness

Planning artifacts are cheap to complete relative to the commitments they describe. A schedule can be built without confirming anyone will staff it. A charter can name a sponsor who never accepted the decisions the role implies. A stakeholder map can list departments whose managers have promised nothing. Each document is individually defensible, and together they feel like rigor.

Approval forums reinforce the swap. Intake reviews are built to inspect documents, because documents are what can be attached to an agenda. The four conditions that actually decide the first ninety days — decision authority, outcome boundaries, participation, and named capacity — live in other people's calendars and intentions, where a checklist cannot see them. So the forum approves the paperwork, the paperwork stands in for the commitments, and the delivery team inherits the gap between them.

The inheritance is quiet. A team missing real sponsorship does not announce it. It shops decisions through informal networks and absorbs delay. A team missing business participation does not stop. It guesses, and the guesses surface at acceptance. By the time the gap becomes schedule variance, the cheap responses — start later, start smaller, staff differently — have expired. Flyvbjerg's argument for front-end quality control makes the general case: the inexpensive moment to challenge a commitment is before it is made, and organizations systematically underinvest in that moment [2].

Incentives finish the job. A sponsor is rewarded for securing the start. The planning consultancy is paid for the plan. The PMO is often measured on approval throughput. Nobody in that chain is paid to say "not yet" — which is why the test has to exist as structure. It puts the unpriced commitments into the price of a start date.

The four-commitment readiness test

The test asks whether four commitments exist in operating form, on the specific work being approved, with evidence a skeptical reviewer could inspect. It is designed for a short conversation between a sponsor and a portfolio forum, and it deliberately ignores most of what a planning pack contains.

EXHIBIT 1

Four commitments, evidenced before a start date

01	Decision authority The sponsor owns forced choices or an accepted escalation route.
02	Outcome with a boundary Name one testable operating result and what this tranche will not attempt.
03	Scheduled participation Managers acknowledge dated commitments for design, validation, and acceptance.
04	Named capacity Confirm constrained roles by name and test allocations against other work.

The sponsor can decide

The first commitment is decision authority. The sponsor must hold authority over the choices this work will force — scope disputes, policy calls, funding moves, conflicts that cross functions — or hold a named, accepted route to whoever does. Visibility into status is not the job. A sponsor who can attend reviews but cannot settle a contested boundary is an audience, and the project will discover that at the first real disagreement.

The failure pattern is familiar in large organizations: a nominal sponsor several levels below the decisions the work implies. When the contested question finally lands — which regulatory reading applies, whose budget covers the integration, which department changes its process — the sponsor becomes a messenger, and the question begins a tour of steering groups. Delivery slows to the speed of that tour. The team learns to route around formal channels, which works until the informal answer contradicts a formal one.

The working control is concrete. Before kickoff, list the first three or four consequential decisions the work is likely to raise. For each, confirm in writing who can make it, and get the escalation route accepted by the people on it. This takes a sponsor an afternoon. Resistance to the exercise is itself diagnostic: a sponsor unwilling to list hard decisions is usually one who suspects the authority is not theirs.

The record is half a page. The likely decisions, the name against each, the route when the name cannot bind the organization, and the sponsor's calendar hold for the first gate. In my experience, producing this record has a way of surfacing that the most consequential decisions belong to an executive who has never heard of the project. That conversation happens before kickoff instead of in month four.

The decision the test enables is equally concrete: start, narrow the tranche to what the confirmed authority covers, or hold until authority matches the commitment being requested.

Outcome and boundary are explicit

The second commitment is a testable outcome with a visible edge. One operational result, expressed as a change in how the business runs — cycle time, error rate, capacity released, exposure closed — plus an honest statement of what this tranche will not attempt. Deliverable language does not qualify. "Implement the platform" names an output. It obligates no one to a result.

Without the boundary, scope becomes a negotiation conducted in fragments across the delivery period, and every fragment lands on the team. Success drifts toward whatever was built. Exclusions discovered late arrive as betrayals. The process owner assumed the workflow was included. The compliance officer assumed the report was. Each assumption seemed reasonable, because nothing visible contradicted it.

The control: write the outcome the way the receiving operation would measure it, identify the baseline or a credible proxy now, and publish the exclusions with the same prominence as the deliverables. The process owner who will live with the result should accept the outcome statement before kickoff. Acceptance in writing — not attendance at the meeting where it was shown.

Enabling and infrastructure work deserves an honest note, because "express the outcome as an operating change" is hardest where the benefit is indirect. A data-platform tranche does not move a customer metric by itself. The discipline still applies, one level down. Name the operating condition the tranche creates — which consuming teams can do what, by when, that they cannot do today — and the first dependent use that will prove it. An enabling project with no named first consumer is the purest form of activity funding, and the test should say so.

The forum's question at the next tranche is then answerable: what result would justify continuing, and what result forces redesign? If that cannot be stated at approval time, the organization is funding activity and hoping.

Stakeholders will participate

The third commitment converts stakeholder sentiment into scheduled behavior. Support expressed in a steering meeting costs nothing. The work needs specific people: the operations expert for design workshops, the data steward for validation, the supervisors who will accept the change, the manager who owns adoption after launch. Every one of those people has a full-time job that predates this project.

The predictable break: names on a stakeholder map, nobody in the room. Design proceeds on secondhand knowledge. Testing samples what the builders imagined instead of what the operation does. The go-live lands on a workforce that first engaged with the change the week it arrived. None of this appears on a status report as "stakeholder absence." It appears later, as rework and slow adoption, attributed to the team.

The control: convert each label on the map into dated commitments tied to concrete work — these workshops, this acceptance decision, this data validation, this transition ownership. Then the harder half: the participants' managers acknowledge the time, because participation unfunded by management evaporates under quarter-end pressure. A fallback should exist for the one expert everyone names. That person is the portfolio's real constraint, and pretending otherwise readies nothing.

A long stakeholder map, run through this discipline, tends to reduce to a handful of people whose time actually matters in the first quarter — and often one or two of them cannot be freed. The map implied broad readiness. The dated commitments reveal a narrow, specific constraint the sponsor can then negotiate. The readiness fact is the six names with acknowledged hours, not the forty on the stakeholder map.

Where participation cannot be secured, the honest responses are to shrink the scope to what secured participation can carry, move the date to when people are available, or stop claiming the dependency is covered. All three beat starting on schedule with a hollow map.

Named capacity and capability exist

The fourth commitment is that the staffing plan survives contact with names. Generic full-time equivalents are an accounting fiction at the moment of readiness. The plan needs this architect, that product owner, a data engineer with the relevant domain, and a subject expert who is not already promised to two other initiatives. Scarce roles decide feasibility, and headcount totals hide them.

A pattern I have seen more than once: the same senior architect appears, in good faith, in the staffing assumptions of several concurrently approved projects. Each plan is individually plausible. The combination is arithmetically impossible. Nobody lied — the portfolio simply had no view that added the names up. The collision surfaces months later as delayed projects and one exhausted architect, resolved by improvisation instead of choice.

The control: plan the constrained roles by name and by period, test the allocation against other active work and expected peaks rather than an annual average, and include the decision load. Executives who must choose are a capacity pool too, and a portfolio can exhaust them as surely as it exhausts engineers. A capacity owner — someone with authority to resolve collisions — should sign, because a collision without an owner resolves itself downward, onto the teams.

Two capacity details are chronically underweighted. Onboarding lead time: a role filled by requisition is not filled at kickoff, and the plan should carry the realistic weeks between offer and productive contribution. And decision load: if the design phase needs six consequential executive choices in eight weeks, the executives' calendars are part of the staffing plan. A portfolio that tracks engineer hours to the half-day while leaving executive attention unmodeled will starve its projects at the top — politely, and invisibly.

When the portfolio cannot supply the scarce capability without damaging work already in flight, the readiness decision is sequencing: reshape, wait, or consciously trade against the in-flight work. Starting on capacity that does not exist is the failure to avoid; waiting is not.

Where one weak commitment defeats the other three

The four commitments protect one another, and a gap in any of them drains the others. An empowered sponsor with a vague outcome makes fast decisions about the wrong question. A crisp outcome without participation is measured precisely and delivered blind. Committed participants without named capacity attend workshops for a build team that keeps dissolving. Real capacity under an absent sponsor executes competently into every unresolved conflict.

This is why the test resists partial credit. Three strong answers and one hopeful one is how false starts happen. The weak commitment determines where the project breaks first. The strong ones determine only how much momentum exists when it does. A reviewer should probe the weakest of the four, on the theory that the sponsor has already rehearsed the strongest.

A worked chain shows the mechanics. A sponsor without authority defers a scope dispute. The outcome statement quietly stretches to cover both readings. The operations experts, seeing the target move, deprioritize workshops for work with firmer edges. The delivery team backfills the missing knowledge with contractor hours the capacity plan never carried. Each step is a rational local response to the previous one. The audit trail will show a staffing problem. The cause was a decision-rights problem, four moves earlier.

The interaction also sets the repair order. Authority first, because every later repair needs a decider. Outcome second, because participation and capacity cannot be scoped against an undefined result. Participation and capacity follow — and they are usually where the test surfaces the portfolio-level truth: the organization has approved more concurrent change than its scarce people can absorb.

Installing the test without growing a bureaucracy

The test earns its keep only if it changes when starts happen. That means it lives inside intake and portfolio review. The sequence is deliberately small:

EXHIBIT 2

Five decisions the test can produce

Approve

All four commitments evidenced. The work has earned its date.

Approve with a boundary

Start narrowed to what the confirmed authority and capacity cover.

Stage discovery

A time-boxed tranche to establish the baseline and confirm the names.

Narrow or hold

Not a defeat - a portfolio truth surfaced before the burn rate begins.

- Use the test before a kickoff is scheduled, not inside the kickoff deck.
- Record one decision for the next commitment: approve, approve with a boundary, stage discovery, narrow, or hold.
- Re-test readiness when a material dependency, sponsor, outcome, or capacity assumption changes.
- Review false starts quarterly, and change the intake questions that let them through.

The re-test trigger is the piece most organizations skip. A readiness decision is evidence with a shelf life. When the sponsor changes, a constrained role leaves, the outcome is renegotiated, or a dependency slips past its assumed date, the original grant no longer describes the work. Any of those events should reopen

the four questions — briefly, against what changed — rather than letting the start date carry authority it no longer earns.

Expect the first cycle to be rough. Sponsors will experience the questions as friction, because the questions move work from the delivery period — where it was invisible and expensive — into the approval period, where it is visible and cheap. A PMO can hold the line by keeping the test short and the consequence real: four commitments, evidence for each, one recorded decision. The moment the test grows into a readiness workbook, sponsors will satisfy the workbook, and the substitution of documents for commitments will have returned.

Staged discovery deserves rehabilitation as an outcome. When the test fails on evidence rather than substance — the outcome is plausible but unbaselined, the capacity picture is unclear rather than bad — the right decision is often a small, time-boxed discovery tranche with its own narrow test: establish the baseline, confirm the names, return in six weeks. That preserves momentum without granting the full commitment, and it converts "hold" from a defeat into a defined path back. The record for any of these decisions is one page.

Ownership belongs where consequence lives. The PMO maintains the method, spots patterns across projects, and challenges thin evidence. It does not absorb the sponsor's accountability for the start decision or the business owner's accountability for the result.

Running the review

The review itself is thirty minutes if the preparation is real. In the room: the sponsor, the process owner who accepted the outcome, the capacity owner, and whoever chairs intake. The sequence follows the repair order — authority, outcome, participation, capacity — and the chair's discipline is to ask for the evidence, not the narrative. "Who accepted the escalation route?" is a better question than "Is sponsorship strong?" because it has a checkable answer.

Tone matters more than format. Run as an audit, the review gets curated evidence from defended sponsors. Run as a service — the portfolio helping a sponsor find the collisions while they are still cheap — the same questions produce candor. That distinction decides whether the test learns the truth at intake or the delivery team learns it in month four. The chair should close every review by naming the single weakest commitment aloud, whatever the decision, so the record carries a warning as well as a verdict.

Signals that the test is working

The point is fewer false starts and earlier honest holds, so measure exactly that:

- Projects held, narrowed, or staged at intake — with the reason recorded — rather than started on hope.
- Time from kickoff to first consequential sponsor decision, and whether it was made by the named route.
- Participation commitments honored in the first ninety days versus promised at approval, and scarce-role collisions found at intake versus in flight.
- Re-tests triggered by changed assumptions, and what they changed.
- False starts per quarter, and which intake question each one evaded.

Read these as diagnostics, in context. A quarter with several holds may be the test working or demand quality falling. The recorded reasons distinguish the two. The number that matters most over a year is the last one — false starts — because it prices what the test exists to prevent.

Watch for gaming, because any gate teaches people to pass it. The tell is evidence that satisfies the letter while the commitment stays hollow: an escalation route "accepted" by an executive assistant, participation hours acknowledged by a manager who oversees none of the named people, a capacity sign-off from someone with no authority to resolve a collision. Spot checks against the first ninety days of actual behavior — did the route get used, did the hours materialize — keep the evidence honest.

The fair objection

The test is not a universal score and cannot eliminate uncertainty. A small reversible discovery effort needs less proof than a large, dependent, hard-to-reverse commitment. The discipline is matching evidence to consequence: require enough proof for the stakes of the next step.

A second objection deserves a direct answer: experienced sponsors already weigh these things by judgment. Many do. But judgment held in one person's head does not survive that person's calendar, and it cannot be challenged, compared, or taught. Writing four commitments down gives judgment a record rather than replacing it — so the next reviewer can see what was actually promised when the start was granted, and so a false start teaches the intake process instead of embarrassing one sponsor.

The objection that the test slows starts is true and incomplete. It slows starts by days and weeks — far less than the quarters a false start costs. An organization that wants speed should measure it from approval to first operating result, not from approval to first status meeting — and on that clock, the test is one of the few governance instruments that reliably buys time back.

What changes

A populated plan stops functioning as a permission slip. Kickoff becomes something the work earns by evidence: a sponsor who can decide, an outcome with an edge, people who will show up, and names that add up. Projects that pass start faster, because the first collisions were resolved while resolving them was cheap. Projects that fail the test surface a portfolio truth at the only moment the portfolio can still act on it — before the burn rate begins.

Over a year, the effect compounds. Fewer false starts release the scarce roles that false starts consume, which raises the readiness of everything else in the queue. The test does not add capacity. It stops the portfolio from spending capacity on work that was never able to use it.

Notes and sources

[1] Project Management Institute, "Pulse of the Profession 2026: Driving Success in Complex Projects," May 2026. <https://www.pmi.org/learning/thought-leadership/driving-success-in-complex-projects>

[2] Bent Flyvbjerg, "Quality Control and Due Diligence in Project Management," 2013.
<https://arxiv.org/abs/1302.2544>

About the author

Marco Policani is an enterprise portfolio, PMO, and AI operating-governance leader. He builds portfolio governance systems where AI carries the analytical load and named humans own every decision — an approach documented across case studies, walkthroughs, and working governance guides on his portfolio site.

Portfolio & governance library: policani.net/governance · Full portfolio: policani.net · LinkedIn: linkedin.com/in/marcpolicani

This white paper may be shared freely with attribution. © 2026 Marco Policani.