
WHITE PAPER · PORTFOLIO & DELIVERY GOVERNANCE

A Project Is Not Ready Because a Plan Exists

Four promises to check before work earns a start

Marco Policani

Enterprise Portfolio · PMO · AI Operating Governance

policani.net · linkedin.com/in/marcpolicani

July 2026

EXECUTIVE SUMMARY

A finished plan does not make a project ready. It is ready when the people it depends on have agreed to what the plan assumes.

A project can reach kickoff with a sound schedule and still lack four promises from people outside the delivery team: someone who can settle hard choices, a result the business can test, people who will take part, and staff who have time. Check all four before setting a start date.

- 01** List the first three or four hard decisions the project will face, and confirm that the person who can make each one will take the question.
- 02** For each person the project needs from the business, name the task and the date, and have that person's manager confirm the time. Six people with confirmed hours tell you more than forty names on a slide.
- 03** If any of the four promises lacks proof, change the start decision: narrow the project, fund a short discovery stage, or wait. Press hardest on the weakest one.

The argument

I have watched projects arrive at kickoff with a sound schedule, a signed charter, and supportive leaders in the room, yet still be unready to start. The missing pieces were promises from people outside the delivery team. Who would settle the first hard decision, and which result would guide that choice? Which business experts would join design and testing, and did the named architect have hours left after the meeting ended?

A complete plan cannot prove that a project is ready, because the people it depends on may not have made the promises the plan assumes. Leaders still need to settle who can decide, what result they want, who will take part, and who has time to do the work.

In this paper's reading, a kickoff date creates its own momentum. Once invitations go out and a team is named, a request to delay can sound like obstruction. Before the

date is set, the sponsor has to explain why the project can start; after it, the person who raises a concern has to defend a delay. Leaders should test readiness first.

Large IT projects can miss badly. A McKinsey-Oxford study of large IT projects found that they ran 45 percent over budget on average and delivered 56 percent less value than expected.[1] The study covers large IT projects and did not test these four promises, so it shows the stakes there, not the cause.

Cooper and Edgett describe new-product portfolio management as choices to select, rank, speed up, or stop work as facts change,[2] and Flyvbjerg warns that a plan's inside view can give poor forecasts.[3] This paper applies that decision logic and caution to a project's start date. PMI's 2026 survey describes how complexity affects projects,[4] and Bain's research links how well firms make and carry out decisions to their financial results;[5] PMI, Flyvbjerg, Cooper and Edgett, and Bain serve as context. None of these sources tests the four questions that follow. They come from my portfolio and PMO work, and they are a decision test, not a score that predicts outcomes.

This paper is for sponsors, portfolio leaders, PMO leaders, and delivery executives who must decide whether to grant a start date. It offers four questions the sponsor must answer with evidence, one for each promise, and it covers only that one decision.

The test fits between the review that shapes a request and the later reviews that check whether the project still deserves support. A team without those steps can still use the four questions to find gaps before work begins.

How a populated plan passes for readiness

A team can finish a schedule before anyone agrees to staff it, and a charter can name a sponsor who has not agreed to make the hard calls. A list of people affected by the project can name departments whose managers have promised no time. Each document is accurate on its own. Together, they give leaders a false sense of readiness.

Review groups often inspect documents because they fit an agenda, but the harder facts sit in other people's authority and schedules. Who can make a disputed choice, what result will count, and which people will take part? Do those people have time in the first ninety days? A checklist cannot answer for them, so the delivery team deals with the missing promises after leaders approve the paperwork.

These gaps stay hidden while a team gets decisions through informal channels, or guesses because business experts are absent. The guesses come to light when the

business tests the result, after the cheaper choices to start later, start smaller, or change staff have passed. Flyvbjerg argues that early cost and benefit estimates often predict poorly and that teams must work to remove bias before the decision.[3]

In this paper's reading, three parties have a reason to avoid saying "not yet": a sponsor may be rewarded for getting work started, a planning firm is paid for its plan, and a PMO may be measured by how many requests it clears. The short test makes the missing promises part of the decision to set a date.

The four-promise readiness test

For this project, the sponsor must show evidence for four promises that a reviewer can check. The sponsor and review group test them in a short meeting, using only the parts of the plan that answer those questions.

EXHIBIT 1

Four promises to check before kickoff

01	Who can decide Name who will settle hard choices.
02	What result will count Name one result to test and what this stage leaves out.
03	Who will take part Managers confirm time for design, tests, and handover.
04	Who has time to build Check named staff against their other projects.

The sponsor can decide

First, the sponsor must be able to settle choices about scope, policy, funding, or a dispute between departments. If another leader must decide, that leader must agree to receive and answer the question. A sponsor who attends meetings but cannot settle a dispute leaves the team waiting at the first hard choice.

In a large firm, the named sponsor may sit several levels below the person who can decide who pays for an integration or which department changes its work. The sponsor carries each question from group to group while the team waits. To keep moving, the team may seek an informal answer that later conflicts with the formal decision.

Before kickoff, list the first three or four hard decisions the project may require and write down who can make each one. If the sponsor cannot decide, confirm that the person who can will take the question. This takes about an afternoon, and resistance to naming the choices may show that the sponsor lacks authority.

The record fits on half a page: each likely decision, who can make it, where it goes if that person cannot, and time held for the first review. A list like this can reveal a senior leader who owns a key choice but has not heard of the project. The sponsor can then reach that leader before kickoff instead of finding the gap after work has started.

Leaders can then start, narrow the project to choices the named people can make, or wait until the right leader agrees to decide.

Name the result and its limits

Second, name one result the business can test, such as shorter cycle time, fewer errors, more staff time, or less risk. State what this round of funding will leave out as clearly as what it will include. "Implement the platform" names a product to deliver, but says nothing about what should improve after people use it.

Without a clear limit, people add requests throughout the project, and the team must settle each one while it builds. Success becomes whatever the team happened to finish. The person who runs the process may expect a new workflow, while the compliance officer expects a report. If no one wrote down what was excluded, both have reason to feel misled when the project closes.

Write the result using measures the team that will use it already tracks, and record today's level or a sound estimate if no measure exists. The person who will live with the result must agree to that statement in writing before kickoff; attendance at a meeting is not agreement.

This is harder for work that supports other teams, because a data platform may not change a customer measure on its own. Instead, name what the funded work will let another team do, by when, that it cannot do today. Identify the first team that will use it and the first use that will show whether it works. Without a named first user, the sponsor has described activity without showing how it will help.

At the next funding decision, leaders can ask what result would justify more work and what result would require a new plan. If no one can answer when the project is approved, the firm is paying for activity while hoping it helps.

The right people will take part

Third, the people the project needs must agree to take part, because support in a steering meeting does not reserve their time. The team may need an operations expert to design the work, a data steward to check records, and supervisors to accept the change. Their managers must also plan for the work after launch, while each person still has a full-time job.

A long list of affected people can hide an empty workshop, leaving designers to depend on secondhand knowledge. Tests then cover what builders expect rather than what staff actually do. Some people first see the change when it goes live. A status report can miss their absence until leaders see rework and slow use and blame the delivery team.

For each person the team needs, name the task and date: a workshop, test, data check, or handover. Ask that person's manager to confirm the time, since routine work can take priority when a deadline arrives. Name a backup for the expert every project depends on; if that expert has no time, the sponsor needs to face the conflict before setting a date.

A list of forty affected people may shrink to six whose time matters in the first quarter, and two of them may not be free. That gives the sponsor a specific problem to solve. Six names with confirmed hours tell leaders more about readiness than forty names on a slide.

If those people cannot join, reduce the work, move the date, or wait; a list of names earns no start.

Name the staff and their time

Fourth, name the people who can do the work and the periods when they are free, because a total staff count hides roles that are hard to fill. The plan may need one architect, a product owner, a data engineer who knows this field, and an expert already named on other projects. The availability of those people, rather than total headcount, decides whether the plan can work.

Consider a test lead named in several approved plans, each of which looks sound on its own. Together, the plans may ask for more hours than one person has, and no one

has added the assignments across projects. If so, the teams face delays and the test lead faces too much work, and leaders then have to make a choice they could have made before kickoff.

Plan scarce roles by name and by month or phase, then check each person's other work and the weeks when demand peaks. Count leaders' decisions too, since their time runs out just as an engineer's does. A leader who can settle staff conflicts should sign the plan; otherwise, teams will fight over the same people after work starts.

Allow time for new staff to join and learn the work; a posted job does not fill a role at kickoff. Reserve time for leaders to decide as well. If design needs six major choices in eight weeks, those meetings belong in the staffing plan. Tracking engineers' hours in detail will not help when leaders cannot make the choices those engineers need.

If the same expert is needed on work already under way, leaders must choose whether to change the new project, wait, or move time from existing work and accept the cost. Starting with hours that do not exist puts both projects at risk.

Where one weak promise defeats the other three

The four promises depend on one another. A sponsor with authority but no clear result can decide fast and still choose poorly, while a clear result will not help if business experts never join the work. Those experts still cannot get a solution built when the delivery team lacks people. A fully staffed team will wait when no sponsor can settle a dispute.

Three firm promises cannot make up for one hopeful answer, because a false start often shows first at the weak promise. The other three keep the work moving until it reaches that gap. Reviewers should press hardest on the weakest answer, which the sponsor has likely had less time to prepare.

Consider a sponsor who cannot settle a scope dispute: the stated result grows to cover both views. Operations experts see the target move and skip workshops for work with clearer limits. The team hires contractors to fill gaps in knowledge, though the staffing plan had no money for them. Each step responds to the last, so the record may point to a staffing problem even though the first gap was the sponsor's lack of authority.

That chain suggests an order for repair: first settle who can decide, then agree on the result so leaders know which people and skills it needs. Check who can take part and who can build it after that. Those last two checks may show that the firm has started more projects than its key people can handle.

Installing the test without growing a bureaucracy

Use the test when leaders review a request for new work, while they can still change the start decision if a promise is missing. The review ends with one recorded choice.

EXHIBIT 2

Choices after the readiness review

Start

All four promises have proof. Set the date.

Start with limits

Begin only the work that named leaders and staff can support.

Fund discovery

Measure the starting point and confirm the people before a full start.

Reduce the work

Cut the project to what the proven promises can support.

Wait

Move the date while leaders still have that choice.

The steps are short:

- Use the test before a kickoff is scheduled, not inside the kickoff deck.
- Record one choice for the next stage: start, start with limits, fund discovery, reduce the work, or wait.
- Review readiness again when a key person, result, staff plan, or another team's promised date changes.
- Review false starts quarterly, and change the intake questions that let them through.

A start decision depends on facts that can change. If the sponsor leaves, a key worker moves, the result changes, or another team misses a date, review the four promises again with a focus on what changed. The original approval cannot cover a project that now needs different people or decisions.

Sponsors may resist the questions at first because the test brings hard choices into the review before teams spend time on a weak start. The PMO can keep it useful by asking for four promises, evidence for each, and one recorded decision. A long

workbook can recreate the problem: people fill it out while the promises remain missing.

When evidence is missing, a short discovery stage can be the right decision, even if the result looks useful. No one may know today's level, or staff may be free but their schedule unclear. Leaders can fund a small first stage to measure the baseline, confirm the people, and return in six weeks. The team learns without promising the full project; record the choice and the next test on one page.

The PMO keeps the questions, spots repeated problems, and challenges weak proof. The sponsor decides when to start; the business owner owns the result.

Running the review

With the evidence ready, the review can take thirty minutes. Invite the sponsor, the person who owns the result, the person who can settle staffing conflicts, and the review chair. Ask about authority, result, participation, and staff time, in that order. The chair should ask for proof: "Who agreed to take a disputed decision?" has a checkable answer, while "Is sponsorship strong?" does not.

How the chair asks matters: sponsors who feel accused may bring only proof that helps their case. A review that helps them find conflicts early is more likely to get a candid answer. Before the meeting ends, the chair should name the weakest promise even when the work can start, so the team knows what to check.

Signals that the test is working

The test aims to catch weak starts while leaders can still change course. Track the decisions it changes and what happens after those decisions:

- Projects held, reduced, or given a short discovery stage, with the reason recorded.
- Time from kickoff to the first major sponsor decision, and whether the named person made it.
- People who took part in the first ninety days compared with those promised at approval; staff conflicts found before or after kickoff.
- Reviews triggered by a change in the plan, and the decisions those reviews changed.
- False starts each quarter, and which question failed to reveal the gap.

Read the numbers with their reasons, because more projects held in one quarter could mean the test found risks earlier or that weaker requests came in. The recorded reasons help leaders tell the difference. Over a year, watch false starts most closely, since preventing them is the purpose of the test.

People can learn to pass a test on paper without making the underlying promise. A route for hard choices may have an assistant's approval, while the leader who must decide has not agreed. A manager may promise hours for staff they do not supervise, or someone may sign a staffing plan without power to settle conflicts. Check what happens in the first ninety days: did the named leader decide, and did the promised people take part?

The fair objection

The four questions do not remove uncertainty or produce a score that fits every project. A small test that can be stopped needs less proof than a large project that depends on many other teams. Ask for enough evidence to match the cost and risk of the next step, rather than demanding the same packet for every request.

Experienced sponsors may already ask these questions. Record the answers anyway. If the project stalls, the review group can see what leaders promised and which question missed the gap.

A readiness review takes time, but an unsupported start can create more delay. This is a risk to test, not a measured payoff for these four questions. To judge speed, measure from approval to the first useful result rather than the first status meeting. Compare projects that passed the review with projects that needed a new start; keep the review only if it saves time and waste in practice.

What changes

A complete plan no longer grants a start by itself. Leaders ask for a sponsor who can decide, a clear result and limit, people who will take part, and staff who have time. If one promise is missing, they can change the project while choices are still open; when all four hold, the team can begin with fewer conflicts left for later.

Over time, fewer false starts should leave more time for the skilled people every project needs. The test creates no new staff hours, but it helps leaders stop promising the same hours to projects that cannot use them well.

Notes and sources

- [1] Michael Bloch, Sven Blumberg, and Jürgen Laartz, "Delivering large-scale IT projects on time, on budget, and on value," McKinsey & Company, 2012. <https://www.mckinsey.com/capabilities/tech-and-ai/our-insights/delivering-large-scale-it-projects-on-time-on-budget-and-on-value>
- [2] Robert G. Cooper and Scott J. Edgett, "Portfolio Management: Fundamental for New Product Success," Stage-Gate International. <https://www.stage-gate.com/blog/portfolio-management-fundamental-for-new-product-success/>
- [3] Bent Flyvbjerg, "Quality Control and Due Diligence in Project Management: Getting Decisions Right by Taking the Outside View," 2013. <https://arxiv.org/abs/1302.2544>
- [4] Project Management Institute, "Pulse of the Profession 2026: Driving Success in Complex Projects," May 2026. <https://www.pmi.org/learning/thought-leadership/driving-success-in-complex-projects>
- [5] Marcia W. Blenko and Michael C. Mankins, "Measuring Decision Effectiveness," Bain & Company, June 5, 2012. <https://www.bain.com/insights/measuring-decision-effectiveness/>

About the author

Marco Policani is an enterprise portfolio, PMO, and AI operating-governance leader. He builds portfolio governance systems where AI carries the analytical load and named humans own every decision — an approach documented across case studies, walkthroughs, and working governance guides on his portfolio site.

Portfolio & governance library: policani.net/governance · Full portfolio: policani.net · LinkedIn: linkedin.com/in/marcpolicani

This white paper may be shared freely with attribution. © 2026 Marco Policani.