
WHITE PAPER · PORTFOLIO & AI GOVERNANCE

Build One Evidence Record for the AI Decision and Your Part in It

One record, written while the program runs

Marco Policani

Enterprise Portfolio · PMO · AI Operating Governance

policani.net · linkedin.com/in/marcpolicani

July 2026



EXECUTIVE SUMMARY

One record, written while the program runs and built around a single claim, supports the funding decision now and your own account of your part later.

A sponsor, the executive who approves the money, asks what can be inspected beyond the demonstration of an AI program. A later employer or promotion panel asks what you personally did. People remember their own share of joint work as larger than the others involved credit them, so an account rebuilt later from memory may be the weaker one.

- 01** Pick the one claim the next decision depends on, and write a record of it as the program runs: the evidence, the limits, and who made each choice.
- 02** Someone outside the program checks the evidence, and the people named confirm their entries. Weak evidence narrows the claim, and missing evidence keeps it out of the decision.
- 03** Confidential work stays with the employer or client. Carry forward a short description of your own part, not the document.

Two readers will ask what you can show

Suppose you lead an AI program and your next funding decision is close. The demonstration works, and the team can explain the workflow and the benefit. The sponsor, the executive who approves the money, asks: What can we inspect that shows this works outside the demonstration, for which uses, and who answers for it?

A later reader will ask a related question. Another employer, a promotion panel, or a client will want to know what you personally did on the program and what came of it. AI programs are new enough that a leader's track record on them may be short, which leaves both readers with less to go on.

This paper expects that two beliefs carry a leader through both questions. One is that the program's own documents will answer them. The other is that you will remember what you did. The next two sections test each belief against what published guidance covers and what memory studies show, and neither holds up well. The usual documents describe the AI service and how it is managed, not who made each decision. Memory favors the person remembering. The cost falls on the sponsor, who decides on thin ground, and on you, whose later account is one nobody can check.

This paper recommends one record, written while the program runs and built around a single claim the decision depends on. The sponsor reads it to decide about the program. Later, you draw a short statement of your own part from it. Many leaders keep a personal file of achievements, and the difference is checking: this record is read and tested by other people while the work is fresh, which a personal file may not be. No study tests whether such a record changes a funding decision or a hiring outcome, so it is this paper's design, not a proven method.

The program's documents describe the service, not your choices

The program half of this problem has its own paper. "The Demo-to-Production Gap" covers the written approval for the AI tool itself, which the business owner signs. The record here is the program leader's working file that supplies the evidence behind such an approval, and it adds one thing that approval does not need: a trail of who made which choice.

Four published documents suggest what to write down about an AI service and its management. A research team proposed FactSheets, a document the supplier fills in to describe a service's purpose, performance, safety, security, and provenance, meaning where its data and models came from.[1] NIST's Generative AI Profile suggests voluntary actions for monitoring, evaluation, and incident response, meaning what to do when something goes wrong.[2] NIST's AI Risk Management Framework asks organizations to document who is responsible for managing AI risk and to plan periodic review.[3] The GAO accountability framework lists governance, data, performance, and monitoring practices for managers and auditors.[4] They offer categories for a record. None of them shows that buyers or boards use the categories, and none records which person made which choice.

The six fields below borrow their evidence, limits, and owner categories from these documents. They add the field the documents lack, which is who made each choice.

Memory favors the person remembering

Long after the decision, you will describe your part from memory and from whatever files remain. Those files may show what the program did without showing who decided what.

Memory fills that gap unevenly. Ross and Sicoly ran five experiments with married couples, student discussion groups, basketball players, and graduate students. People recalled more of their own contributions to joint work than of the others', and they claimed more responsibility for a joint product than the other participants credited to them.[5] In the couples study, each partner rated who was responsible for 20 shared activities, from shopping to causing arguments. Across them, 27 of 37 couples together claimed more than 100 percent, although the average overestimate was small. The authors say the bias can arise without any wish to deceive. They did not study workplaces, programs, or leaders. This paper treats the result as a warning about memory, not as a measurement of any leader.

Two other findings from the same paper shape the design. First, notes did not clearly help. Pairs kept notes during a 45-minute discussion, each person recording either their own contributions or their partner's, and that choice did not change who people credited. The study had no group without notes. The authors suspected the notes, about a page long, were too brief, and that much of what was written seemed irrelevant to the final decisions.[5] Second, graduate students who were first asked about their supervisors' contributions to their theses credited the supervisors with about twice the overall responsibility, 33 percent against 17 percent, of those first asked about their own.[5] This paper infers from both that a record should tie each choice to the decision it shaped and should name the other contributors.

What the later reader looks for is only partly known. In LinkedIn's own survey of 1,271 recruiting professionals, 93 percent said that accurately assessing a candidate's skills is crucial for improving the quality of a hire.[6] That is LinkedIn's research on its own members, and it shows only that recruiters rate skills assessment as important.

The record: one claim, six fields

The record is a short document, a few pages at most. You write it, the sponsor reads it, and you update it at each decision. The sponsor's decision comes first because it is the nearer deadline and it forces the record to be checkable. The statement for the later reader is drawn from the same record afterward.

Pick the claim by asking: If this claim were false, would the sponsor decide differently? Suppose the program is an AI assistant that answers customer requests. "The program is a success" fails the test because nothing can be checked against it. "The assistant answers these kinds of request correctly in these settings, with these exceptions" passes. If you cannot narrow the claim that far, the decision itself may be too vague, and the first move is to ask the sponsor which question it turns on.

Then fill in six fields:

- **Claim.** State exactly what the program does or has achieved, in terms the evidence can test. General labels such as "successful," "transformative," or "trusted" are too broad to test. Your own part goes in the next field, not here.
- **Context and work.** Record the circumstances, the work performed, and the choices that shaped the result. For each choice, say who made it and who else contributed. Ask each person you name to read the entry about them and correct it. If someone disputes an entry, the record shows both versions instead of dropping it. This is the field the later reader depends on.
- **Evidence produced.** Attach the file, measurement, decision record (the note, ticket, or approval written when the choice was made), test result, or source that supports the claim. Put it close enough that a reader can inspect it without reconstructing the story.
- **Result or learning.** Say what changed or failed. Include what the work revealed when that learning supports a decision to continue, change, or stop.
- **Limits.** State what the evidence does not establish: the circumstances it does not cover, the comparisons nobody made, and the uncertainty that would otherwise invite a broader conclusion. A visible limit shows where the evidence stops and makes the rest of the record easier to believe.
- **Owner and next decision.** Name the person accountable for keeping the record current, set the date or event that triggers a refresh, and say what decision comes next. If you move on, the owner changes and the record says so.

A weak field and a missing field mean different things. This paper applies the rule from "The Demo-to-Production Gap" to one claim instead of a whole approval. A weak field means evidence exists but covers little, such as a test on one office's requests. Narrow the claim to what the evidence covers. A missing field means no evidence exists yet. The claim then stays out of the sponsor's decision and out of any later account until the evidence exists, although the program can continue as a small trial on low-risk work. A limits field left empty counts as missing, not as "no limits."

Two checks protect the record. Before the review, someone outside the program, such as an analyst the sponsor names, checks the claim against its evidence and the choices against the decision records. The people named in the context field confirm their entries.

What can go wrong

The record gives the sponsor material to use against you. It shows where your evidence stops and what went wrong, and a sponsor could use both. This is a real cost, and the paper cannot remove it. This paper's judgment, which no source tests, is that a limit you state in your own words reads better than one a reviewer finds later. Write what the evidence shows and leave out opinions about colleagues. If a sponsor penalizes plainly stated limits, the record will show that early.

Much of the work is confidential. Confidential work can remain private. Evidence can be anonymized, aggregated, or described through a method without exposing a client or an internal program. Keep only what the later reader needs and what confidentiality permits. The record belongs to the program and the employer. Before keeping a copy or describing it elsewhere, check the confidentiality terms with a manager or legal contact for the employer or client. What you carry forward is a short statement of your own part, not the document: two or three sentences saying which claim you worked on, what you decided or did, who else contributed, and what evidence exists.

The record can be gamed or turn into paperwork. The risk is that people write the evidence they have rather than the evidence the decision needs, that limits become boilerplate, and that nobody reads the result. Four checks help. Attach the evidence rather than a summary of it. Treat an empty limits field as missing. Have the outside reader and the named people check as described above. Stay with one claim so the record is short enough to read. None of these removes the risk, and a record that nobody reads is paperwork.

It costs time and slows ordinary decisions. As planning estimates, not measurements, expect about two hours to set up a record around one claim, about fifteen minutes to update it at each decision, a few minutes for each named person to confirm an entry, and about an hour for the outside check. A known expert in a trusted relationship can move faster when the stakes and the cost of error are low. The record earns its cost when the stakes justify checking or when the reader does not know you.

Nothing here shows the record works. The published documents offer categories, not proof of adoption. The memory studies concern couples, students, and teams, not leaders. LinkedIn's figure is a belief. No source says how often a leader is asked what they personally did, or whether a record changes the outcome when they are. Credentials, reputation, and trusted relationships continue to matter, and the record adds to them without replacing them.

What to do next

Before the next decision that matters, pick one claim and fill in the six fields. Give the sponsor the record before the meeting, and have someone outside the program check the evidence for that claim. If a field is weak, narrow the claim. If a field is missing, keep the claim out of the decision until the evidence exists.

Decide in advance, under the confidentiality terms, which parts of the context and work field you may keep for later. Update the record at each decision. When you leave the program, write the short statement of your own part from the record while it is still in front of you, and show it to the same manager or legal contact before you share it.

The sponsor's question then has an answer on paper: what can be inspected, for which uses, and who answers for it. The later reader gets the same answer, narrowed to your part and checked by the people who were there.

Notes and sources

[1] Matthew Arnold, Rachel K. E. Bellamy, Michael Hind, Stephanie Houde, Sameep Mehta, Aleksandra Mojsilović, Ravi Nair, Karthikeyan Natesan Ramamurthy, Darrell Reimer, Alexandra Olteanu, David Piorkowski, Jason Tsay, and Kush R. Varshney, "FactSheets: Increasing Trust in AI Services through Supplier's Declarations of Conformity," arXiv:1808.07261, submitted August 22, 2018, revised February 7, 2019. <https://arxiv.org/abs/1808.07261>

[2] National Institute of Standards and Technology, "Artificial Intelligence Risk Management Framework: Generative Artificial Intelligence Profile," NIST AI 600-1, July 2024. <https://doi.org/10.6028/NIST.AI.600-1>

[3] National Institute of Standards and Technology, "AI Risk Management Framework Core," excerpt from AI RMF 1.0, 2023. <https://airc.nist.gov/airmf-resources/airmf/5-sec-core/> <https://airc.nist.gov/airmf-resources/airmf/5-sec-core/>

[4] U.S. Government Accountability Office, "Artificial Intelligence: An Accountability Framework for Federal Agencies and Other Entities," GAO-21-519SP, June 30, 2021. <https://www.gao.gov/products/gao-21-519sp>

[5] Michael Ross and Fiore Sicoly, "Egocentric Biases in Availability and Attribution," Journal of Personality and Social Psychology 37(3), 322-336, 1979.
http://web.mit.edu/curhan/www/docs/Articles/biases/37_J_Personality_Social_Psychology_322_%28Ross%29.pdf
http://web.mit.edu/curhan/www/docs/Articles/biases/37_J_Personality_Social_Psychology_322_%28Ross%29.pdf

[6] Greg Lewis, "LinkedIn Report: How AI Will Redefine Recruiting in 2025," LinkedIn Talent Blog, February 13, 2025. <https://www.linkedin.com/business/talent/blog/talent-acquisition/future-of-recruiting-2025>

About the author

Marco Policani is an enterprise portfolio, PMO, and AI operating-governance leader. He builds portfolio governance systems where AI carries the analytical load and named humans own every decision — an approach documented across case studies, walkthroughs, and working governance guides on his portfolio site.

Portfolio & governance library: policani.net/governance · Full portfolio: policani.net · LinkedIn: [linkedin.com/in/marcpolicani](https://www.linkedin.com/in/marcpolicani)

This white paper may be shared freely with attribution. © 2026 Marco Policani.