

The Prove-It Economy Is Coming for Careers and AI Programs

How one record can back an AI funding decision and a leader's account of their own part

Marco Policani

Enterprise Portfolio · PMO · AI Operating Governance

policani.net · linkedin.com/in/marcpolicani

July 2026

EXECUTIVE SUMMARY

For an AI-program leader, the evidence needed to fund the program and support a career claim comes from the same operating record.

An AI-program leader can enter a review with a working demonstration and still lack the record needed for the next decision. The same leader may later have to show a hiring committee what they actually governed and delivered, and if the record is thin, confidence has to carry more weight than either decision warrants.

- 01** A demonstration shows capability under selected conditions. A consequential decision also needs results, limits, and accountable ownership.
- 02** Start with the claim the next decision depends on and record it as the program runs: what was done and why, the evidence, the result or lesson, its limits, the owner, and the next decision.
- 03** Let the résumé point to the same underlying records, and claim only your own part in what those records show.

The sponsor's harder question

You lead an enterprise AI program and are preparing the case for its next funding decision. The demonstration works. The team can explain the model, the workflow, and the intended benefit. Your own professional record says you have led difficult transformation work. Then the sponsor asks a harder question: What can we inspect that shows this program works outside the demonstration, within known limits, under accountable ownership?

The usual signals still have value. A demonstration shows capability under selected conditions. A respected title or brand earns initial attention. A clear explanation helps an audience understand the case. The sponsor still needs evidence for a decision made under real operating conditions.

The missing evidence matters beyond the review meeting. An AI system may mediate how an evaluator finds and compares claims. A true claim with a thin record may then lose to a narrower claim backed by sources or artifacts the evaluator can inspect.

The program leader faces two exposed claims. The sponsor must decide whether the program deserves more funding or reliance. A future hiring committee may assess what the leader actually governed and delivered. If the available evidence is thin, confidence must carry more weight than either decision warrants.

A demonstration leaves the board's question open

AI programs face a related test inside the enterprise. A persuasive demonstration establishes that a system can perform under the conditions shown. A funding or reliance decision also needs evidence from normal use. The sponsor needs to understand the result and its conditions. The review must also cover known limits, the response to failure, and who owns the decision to continue.

The evidence base supports that distinction. AI-service FactSheets propose supplier disclosures about an AI service.[2] The NIST Generative AI Profile suggests voluntary actions for monitoring, evaluation, and incident response across the lifecycle.[3] The GAO accountability framework adds governance and management considerations.[4] Together, these sources offer categories for a record; they do not establish adoption by buyers or boards.

The program leader therefore has a decision problem. An honest demonstration and a capable team can still leave the sponsor unable to reconstruct what happened under operating conditions. The funding case remains incomplete until the record covers the result, its limits, and accountable ownership.

Evaluators can inspect the record

The sponsor will read that record directly. A later evaluator may encounter only what a search or answer system can retrieve. The career side of the problem therefore depends first on whether the relevant evidence exists in accessible form.

Research on generative engine optimization tested what happens after retrieval: how visible an already-retrieved source becomes in a generated answer. On its benchmark, adding citations, quotations from credible sources, and statistics improved a source's visibility in generated answers.[1] The finding is bounded to the benchmark and one deployed engine. It shows that evidence-bearing text can give a generative system material to use in an answer.

That result does not predict any individual's hiring outcome. It identifies a practical disadvantage for an accurate claim that is poorly documented: neither a person nor a system can inspect evidence that was never retained.

Use the standard where scrutiny earns its cost

The record should be proportional to the decision. A known expert can move quickly in a trusted relationship when the stakes and cost of error are low. Formal proof for every conversation would add delay without improving ordinary decisions. The record earns its cost when the stakes justify checking, the evaluator is unfamiliar with the claimant, or a machine mediates the assessment.

Evidence can also be staged. Sources may fail to support the sentence beside them. Statistics may be stripped of context. An artifact may look rigorous while proving very little. The evaluator must therefore check support at the level of the claim and keep the boundary visible.

This leaves the leader with two apparently separate jobs. The program team can assemble material for the board. Later, the leader can reconstruct a career story from memory and whatever artifacts remain. The audiences and decisions differ, so separate preparation can seem reasonable. The unresolved question is whether the evidence must be created twice.

One record carries both kinds of credibility

The board review and the later career claim should draw from one operating record created during delivery. That record connects a consequential claim to the decisions and results that produced it. At that point, the conditions and limits are still checkable.

EXHIBIT 1

The board review and the later career claim should draw from one operating record created during delivery.

The presentations remain different. A board decides whether to fund, rely on, redirect, or stop. A career reader assesses the leader's judgment, contribution, and result. Adapting the record for the second decision means selecting only evidence relevant to the leader's part in the work.

Careful leaders can benefit from this discipline. A visible limit shows where the evidence stops and helps an evaluator judge the result. The record becomes more credible when it includes uncertainty that would otherwise make the claim look broader than the work supports.

Build one proof record

Before the next consequential review, select one claim the decision depends on. Build one record around that claim with six fields:

EXHIBIT 2

One record supports the next consequential decision

1

Claim

State exactly what the program or leader is said to have achieved.

- 2 **Context and work**
Record the conditions, work, and choices that shaped the result.
- 3 **Evidence produced**
Attach the artifact, measurement, decision record, test, or source.
- 4 **Result or learning**
Say what changed, failed, or became clear enough to decide.
- 5 **Limits**
State what the evidence does not establish.
- 6 **Owner and next decision**
Name accountability, the refresh trigger, and the decision.

The paper's inference from the GEO benchmark is narrow: direct sources and explicit claim language may make evidence easier for a generative system to use. Stable artifacts also help a person reconstruct the decision. Write for people and retain enough context for either evaluator to inspect the claim.

The owner matters because conditions change. Models and workflows change too. A result can outlive the context that made it true, so an unowned record eventually becomes unreliable.

Let the resume point to the same evidence

No source here measures an individual's hiring outcome. LinkedIn's 2025 recruiting research reported that 93 percent of surveyed talent-acquisition professionals considered accurate skills assessment crucial to quality of hire. In LinkedIn's platform data from the same report, organizations with the most skills-based searches were 12 percent more likely to make a quality hire.[5] These findings support attention to demonstrable skills. The effect of a public portfolio remains unmeasured here.

For the accountable program leader, the practical implication is narrow. Treat the resume as an index to claims that have evidence behind them. "Led an AI transformation" asks the evaluator to trust a label. A smaller claim can point to an artifact or measured result. It can also show a decision pattern or documented learning, with a clear boundary around the leader's contribution.

Confidential work can remain private. Evidence may be anonymized, aggregated, or described through a method without exposing a client or internal program. Retain only what the career claim needs and what confidentiality permits. A credential can earn attention; the evidence then gives the evaluator something to assess.

Two evidence gaps remain. The GEO research demonstrates a benchmark mechanism; it supplies no universal rule or timetable for AI-mediated discovery.[1] The recruiting research supports skills assessment, while the effect of inspectable public work on an individual's hiring outcome remains unmeasured here.[5] Credentials, reputation, and trusted relationships continue to matter.

Put proof into the work

Return to the program review. Create one claim-level record before expanding the deck. Use its six fields to show the sponsor what happened and what remains uncertain. If the record cannot support the claim, narrow the claim or delay the decision until the missing evidence exists.

Make the record a normal output of delivery for the claims with the highest stakes. Start with one program decision and test whether the six fields help the sponsor decide. Reuse the structure after that test.

Build the record during the work, while its context can still be checked. Later adaptation then becomes a controlled use of retained evidence rather than a claim reconstructed from memory.

Notes and sources

[1] Pranjal Aggarwal, Vishvak Murahari, Tanmay Rajpurohit, Ashwin Kalyan, Karthik Narasimhan, and Ameet Deshpande, "GEO: Generative Engine Optimization," accepted to KDD 2024; arXiv:2311.09735, submitted November 16, 2023, revised June 28, 2024. <https://arxiv.org/html/2311.09735>

[2] Matthew Arnold, Rachel K. E. Bellamy, Michael Hind, Stephanie Houde, Sameep Mehta, Aleksandra Mojsilović, Ravi Nair, Karthikeyan Natesan Ramamurthy, Darrell Reimer, Alexandra Olteanu, David Piorkowski, Jason Tsay, and Kush R. Varshney, "FactSheets: Increasing Trust in AI Services through Supplier's Declarations of Conformity," arXiv:1808.07261, submitted August 22, 2018, revised February 7, 2019. <https://arxiv.org/abs/1808.07261>

[3] National Institute of Standards and Technology, "Artificial Intelligence Risk Management Framework: Generative Artificial Intelligence Profile," NIST AI 600-1, July 2024. <https://doi.org/10.6028/NIST.AI.600-1>

[4] U.S. Government Accountability Office, "Artificial Intelligence: An Accountability Framework for Federal Agencies and Other Entities," GAO-21-519SP, June 30, 2021. <https://www.gao.gov/products/gao-21-519sp>

[5] Greg Lewis, "LinkedIn Report: How AI Will Redefine Recruiting in 2025," LinkedIn Talent Blog, February 13, 2025. <https://www.linkedin.com/business/talent/blog/talent-acquisition/future-of-recruiting-2025>

About the author

Marco Policani is an enterprise portfolio, PMO, and AI operating-governance leader. He builds portfolio governance systems where AI carries the analytical load and named humans own every decision — an approach documented across case studies, walkthroughs, and working governance guides on his portfolio site.

Portfolio & governance library: policani.net/governance · Full portfolio: policani.net · LinkedIn: linkedin.com/in/marcpolicani

This white paper may be shared freely with attribution. © 2026 Marco Policani.