
WHITE PAPER · PORTFOLIO & DELIVERY GOVERNANCE

Separate the Decision From the Outcome

Why reaching the funded target does not settle whether the outcome—or the decision—was sound

Marco Policani

Enterprise Portfolio · PMO · AI Operating Governance

policani.net · [linkedin.com/in/marcpolicani](https://www.linkedin.com/in/marcpolicani)

August 2026

EXECUTIVE SUMMARY

Judge a finished program on its full material effects, not just the funded target; record how it got there; and grade the approval on what was knowable when the money was committed.

A green target does not prove the outcome was good or the approval was sound. Review three things separately: the full outcome, the route the team took compared with the approved plan, and the original decision.

- 01** Score the full outcome, not just the funded target.
- 02** Grade the approval using only what leaders knew at the time.
- 03** Record whether a gap came from uncertainty the case declared, a change in delivery, or a repeated error in the way estimates are made.

CONTENTS

A target is not the whole outcome	Score the whole outcome
Why results distort judgment	Attribute the gap before assigning credit
What counts as a decision	Change the review, not just the form
Use two grades, not one verdict	What to do next
Reconstruct the approval before opening the result	Notes and sources

A target is not the whole outcome

Suppose a drunk driver arrives home uninjured and in record time. If the scorecard records only his arrival time and whether he was hurt, the trip looks successful. In fact he hit a cat, broke laws, and drove through someone's yard. He reached the destination. His choice was reckless. His route caused harm. The full outcome was not a success.

Portfolio reviews can hide the same mistake. A program can hit its benefit number while bypassing controls, exhausting a key team, shifting risk into operations, or harming customers. One green target does not describe everything the program caused.

Reviews can also make the opposite mistake. A program can miss a target after an outside event that no one could reasonably predict. If the original case named its limits, compared real options, and made the trade-offs clear, a bad result does not prove the approval was poor.

Keep three questions apart, and answer them in this order:

- Was the original decision sound, based only on what leaders knew at approval?
- What was achieved when the full consequence is counted?
- How was it achieved compared with the approved route?

The review needs two grades and one attribution record. Before seeing the outcome score, grade the approval from its original evidence. Then score the full outcome and compare the approved route with the route the team actually took. That comparison helps explain the gap without allowing a later result to rewrite what was knowable at the time.

Why results distort judgment

Outcome bias is the habit of judging a decision by the result that followed it. Baron and Hershey gave people the same decision, the same odds, and the same information. They changed only the outcome. Participants rated the decision and the decision maker more highly when the result was good.[1]

A 2023 replication with 692 participants found the effect again in both versions it tested.[2] The result matters to the people affected by it. It does not prove that the earlier choice was better or worse than the same choice would have been before anyone knew the result.

Jan van Ours studied manager replacement in European football, where clubs had ongoing performance indicators. Recent results still drove many replacement decisions. On average, replacement did no better than regression to the mean.[3] The narrow lesson is that leaders can see better indicators and still let the latest result dominate their judgment.

Michael Mauboussin and Dan Callahan make the same practical point: when luck affects a result, the outcome does not fully explain the process that produced it. The part leaders can improve is the process.[4]

Portfolio leaders face a harder version of the problem. Their work has fewer comparable cases and slower feedback. A single finished program is thin evidence. Flyvbjerg and Budzier studied 1,471 IT projects. The average cost overrun was 27 percent, but one in six projects had much larger overruns and delays.[5] One win or loss cannot explain the quality of the decision that came before it.

What counts as a decision

A portfolio decision is a commitment made at a point in time. It uses a defined set of information, chooses among real alternatives, accepts trade-offs, and names an authority who can act.

The Society of Decision Professionals describes six parts: the frame, alternatives, information, values and trade-offs, reasoning, and commitment to action.[6] This is not a promise of success. It is an inspection list. A reviewer can see whether the record supports each part or whether something was missing.

The information test needs special care. Hindsight can quietly enter the review. A reference class of comparable past work was available at approval, whether or not the team used it. Flyvbjerg calls this the outside view.[7] If a useful base rate existed and the team ignored it, that is a weakness in the decision. A named vendor being acquired in a particular quarter may not be.

Use two grades, not one verdict

Put decision quality on one axis and full outcome on the other. That creates four cases.

EXHIBIT 1

Four cases need four responses

- ① **Sound decision, good outcome**
Keep the reasoning, not just the tactics.
- ② **Sound decision, bad outcome**
Update the risk or assumption; protect disciplined risk-taking.
- ③ **Weak decision, good outcome**
Review the method before luck becomes policy.
- ④ **Weak decision, bad outcome**
Name the missing input and repair the process.

Sound decision, good outcome. Celebrate the result, but record the reasoning rather than copying every tactic. A method should become standard because its logic holds across cases, not because it happened to be present for one win.

Sound decision, bad outcome. Protect disciplined risk-taking. If the case named a real dependency, set a reasonable range, and made its trade-offs clear, update the assumption or risk price. Do not punish the sponsor for an event the record could not have predicted.

Weak decision, good outcome. This is the case leaders miss most often. The program lands, the sponsor is rewarded, and the method becomes precedent. Review a sample of successful programs with the same care as failures. Otherwise luck can become policy.

Weak decision, bad outcome. Name the missing input, not the person. "The sponsor overcommitted" explains little. "The case had one real alternative and no estimate range" identifies a defect that a team can repair.

The two grades do not settle why the program landed in a particular case. The attribution record does that work. It says whether the gap came from declared uncertainty, a change in execution, or a repeated flaw in the model.

Reconstruct the approval before opening the result

The cleanest review starts with a record created at approval. Keep it short. It should include the problem and decision date; the person who decided; the options

considered and why the others lost; the main assumptions, risks, and trade-offs; a reference class when one exists; ranges for cost, schedule, and benefit; and the evidence that would cause leaders to revisit the choice.

At review, use that record to grade the decision. Keep later facts in a separate column. They are useful for learning, but they cannot change the evidence that was available at the time.

The grade attaches to the decision, not to the decision maker. Sponsors will provide an honest record only if a sound decision can still receive a bad outcome without becoming a personal failure.

Two defects are common. Some cases present one serious option and two options written to lose. Others offer a range so wide that no outcome could fall outside it. Both defects are visible from the approval record. Neither requires hindsight.

Score the whole outcome

Start with the funded target, then go wider. Check benefit, cost, schedule, and risk against what the case promised. Count material effects the case did not price: platform debt, customer harm, work the organization can no longer do, or risk shifted to another part of the business.

Next compare the approved route with the route the team took. Were controls bypassed? Did scope change without a new decision? Did a named dependency go unmanaged? The route matters even when the destination is reached.

The result is strong evidence about what happened. Money was spent. A customer was harmed or helped. A benefit arrived or did not. It is weaker evidence about whether a decision was sound when the case contained a real range and uncertainty was material. The outcome score should never edit the approval-time grade.

Attribute the gap before assigning credit

Use three tests in order.

EXHIBIT 2

Three tests explain a gap

- ① **Declared uncertainty**
Check whether the result landed inside the stated range.
- ② **Delivery change**
Compare the approved route with the route taken.
- ③ **Model error**
Look for repeated error in the same direction across comparable decisions.

First, ask whether the result landed inside the range the decision declared. A result inside an honest range may be uncertainty resolving. If the case declared no range, the review has already found a decision defect: the case claimed more certainty than the work could support.

Second, ask whether the organization did what it approved. Added scope, changed sequencing, withdrawn staffing, bypassed controls, or ignored dependencies may belong to delivery governance. The comparison only works when the approval record is specific enough to compare with actual work.

Third, look across comparable decisions. One miss can be noise. A repeated error in the same direction is a reason to inspect the estimating model, gate criteria, or delivery process. Errors scattered around an estimate may show an imprecise but roughly centered process. Repeated optimism points to a systematic problem.

These categories can overlap. Record the evidence and the judgment rather than pretending the answer is certain. Early judgments can change as a portfolio collects more cases. A written attribution makes that revision possible.

Change the review, not just the form

Start with one decision above a sensible commitment threshold. At approval, lock the short decision record. Do not try to recreate every old decision after results are known. A retrospective record is useful only when it is clearly labeled as such.

At the next portfolio review, start with the approval record. Grade the original decision before the reviewer reads the outcome score. Then score what was achieved, compare the route taken with the approved route, and record the attribution.

Review successful programs as well as failed ones. Measure calibration across comparable decisions, not sponsor rankings from isolated outcomes. A sponsor with a strong approval record and a bad result may have priced a risk too low. A sponsor with a weak record and a good result may have been lucky. Both findings help leaders improve the next choice.

What to do next

At the next stage gate above your commitment threshold, require a one-page decision record. At the next review, grade the original decision from the approval-time record before opening the outcome score. Then score the full outcome and compare the approved route with the route taken.

EXHIBIT 3

| Separate the full outcome, route, and decision.

Then review one successful program at the same depth as a failure. If leaders can separate the result, the route, and the decision, they can learn what to repeat without turning luck into a standard or treating every miss as poor judgment.

Notes and sources

[1] Jonathan Baron and John C. Hershey, Outcome bias in decision evaluation, *Journal of Personality and Social Psychology* 54(4), April 1988, 569–579.

<https://pubmed.ncbi.nlm.nih.gov/3367280/>

[2] Sriraj Aiyer, Hoi Ching Kam, Ka Yuk Ng, Nathaniel A. Young, Jiaxin Shi, and Gilad Feldman, Outcomes Affect Evaluations of Decision Quality, *International Review of Social Psychology* 36(1), article 12, July 28, 2023. <https://rips-irsp.com/articles/10.5334/irsp.751>

[3] Jan C. van Ours, Outcome bias in managerial decisions, *Journal of Economic Psychology* 112, January 2026, article 102872.

<https://www.sciencedirect.com/science/article/pii/S0167487025000844>

[4] Michael J. Mauboussin and Dan Callahan, Outcome Bias and the Interpreter: How Our Minds Confuse Skill and Luck, *Credit Suisse Global Financial Strategies*, October

15, 2013. <https://www.newedgewealth.com/wp-content/uploads/document-805810220.pdf>

[5] Bent Flyvbjerg and Alexander Budzier, Why Your IT Project Might Be Riskier Than You Think, Harvard Business Review 89(9), September 2011; open-access version arXiv:1304.0265. <https://arxiv.org/abs/1304.0265>

[6] Society of Decision Professionals, "Decision Quality."
<https://www.decisionprofessionals.com/who-we-are/decision-quality>

[7] Bent Flyvbjerg, Quality Control and Due Diligence in Project Management: Getting Decisions Right by Taking the Outside View, arXiv:1302.2544, 2013.
<https://arxiv.org/abs/1302.2544>

About the author

Marco Policani is an enterprise portfolio, PMO, and AI operating-governance leader. He builds portfolio governance systems where AI carries the analytical load and named humans own every decision — an approach documented across case studies, walkthroughs, and working governance guides on his portfolio site.

Portfolio & governance library: policani.net/governance · Full portfolio: policani.net · LinkedIn: linkedin.com/in/marcpolicani

This white paper may be shared freely with attribution. © 2026 Marco Policani.