

Stage Gates Work Only When They Can Kill Projects

Decision rights that protect scarce capacity

Marco Policani

Enterprise Portfolio · PMO · AI Operating Governance

policani.net · linkedin.com/in/marcpolicani

July 2026

EXECUTIVE SUMMARY

A gate earns its place only when it can continue, redirect, pause, or stop the next commitment.

A stage gate governs only when it connects evidence to a consequential choice. If every project advances regardless of what the review finds, the gate is a status meeting wearing the language of control. This paper gives portfolio boards, sponsors, PMO leaders, and investment owners a practical a gate that earns the next commitment for making the next commitment explicit.

- 01** Delegate real decision authority: The gate forum needs explicit authority to continue, redirect, narrow, pause, or stop the next commitment.
- 02** Ask for evidence proportionate to commitment: Evidence should answer the risk created by the next tranche, not satisfy the same checklist at every stage.
- 03** Keep all five outcomes credible: Advance is one valid outcome among several; redirection, narrowing, pausing, and stopping are evidence of governance working.
- 04** Separate evidence owners from advocates: The sponsor may advocate for the investment, but critical evidence needs an owner accountable for its integrity.

CONTENTS

The argument	Make conditions expire
Why the conventional approach breaks	How the elements interact
A gate that earns the next commitment	Put the model into the management cadence
Delegate real decision authority	Measures that reveal whether the control works
Ask for evidence proportionate to commitment	The strongest counterargument
Keep all five outcomes credible	The shift
Separate evidence owners from advocates	

The argument

A stage gate governs only when it connects evidence to a consequential choice. If every project advances regardless of what the review finds, the gate is a status meeting wearing the language of control.

The most revealing question at a gate review is not what color the project is. It is what the forum is allowed to decide. In weak systems the answer is effectively nothing: leaders can ask for clarification, express concern, and

request another slide, but the work continues because staffing, vendor commitments, and political expectations already assume approval.

That arrangement rewards momentum over evidence. It also turns early uncertainty into later sunk cost, when stopping becomes harder and the organization has fewer options.

This paper is written for portfolio boards, sponsors, PMO leaders, and investment owners. It offers a gate that earns the next commitment as a management instrument: a way to connect operating evidence to the next consequential choice. The cited sources provide external context; the framework and its application are Marco Policani's synthesis. It is not a predictive score and does not replace professional, legal, security, financial, or domain judgment.

Why the conventional approach breaks

The visible artifact is often complete before the operating condition behind it is credible. Conventional governance tends to separate planning, approval, delivery reporting, and operating ownership. That separation allows a premise to pass through several forums without any one forum owning the decision it implies. The project or workflow then carries an assumption that was acknowledged socially but never accepted operationally.

Because responsibility is distributed, no single event appears large enough to stop the work. The accumulated operating load is the real exposure. The answer is not heavier control at every step. It is to place a clear, proportionate test at the point where the organization increases money, capacity, access, dependency, or action authority.

A gate that earns the next commitment

The model has 5 connected elements. Each makes a different condition visible, but they should be reviewed together because strength in one area does not cancel a missing obligation in another. The model is designed for a short executive conversation supported by inspectable evidence, not a long maturity assessment.

EXHIBIT 1

A gate that earns the next commitment connects evidence to action

- 1 **Delegate real decision authority**
The gate forum needs explicit authority to continue, redirect, narrow, pause, or stop the next commitment.
- 2 **Ask for evidence proportionate to commitment**
Evidence should answer the risk created by the next tranche, not satisfy the same checklist at every stage.
- 3 **Keep all five outcomes credible**
Advance is one valid outcome among several; redirection, narrowing, pausing, and stopping are evidence of governance working.
- 4 **Separate evidence owners from advocates**
The sponsor may advocate for the investment, but critical evidence needs an owner accountable for its integrity.

5

Make conditions expire

A gate decision is based on evidence that can age; approval should identify the assumptions that would reopen it.

Delegate real decision authority

The gate forum needs explicit authority to continue, redirect, narrow, pause, or stop the next commitment. A useful governance design therefore begins with the consequence leaders are prepared to accept.

Participants advise while the actual decision sits elsewhere, so the project interprets discussion as approval and moves on. A local fix can keep the work moving while worsening the portfolio: another specialist is borrowed, another exception is tolerated, and another dependency is promised twice.

Publish the forum's decision rights, quorum, reserved matters, and escalation path before evidence is prepared. The control is complete only when the response is known: who acts, how quickly, within what authority, and with which safe fallback.

A charter for the gate, named decision owner, recorded decision, and proof that resource commitments follow the outcome. The receiving operation must recognize the evidence as relevant. Technical success alone cannot prove that a business workflow can carry the change.

If the forum cannot change money, capacity, scope, or reliance, redesign the forum rather than adding ceremony. That choice should be recorded in operational terms: what may proceed, what remains outside the boundary, who owns the next move, and what would reopen the decision.

A tabletop review of delegate real decision authority should use the observed break as its scenario: Participants advise while the actual decision sits elsewhere, so the project interprets discussion as approval and moves on. Participants then apply the proposed response. They publish the forum's decision rights, quorum, reserved matters, and escalation path before evidence is prepared. The review identifies where authority, information, time, or recovery would run out. It is useful only if it changes the boundary or confirms who will act when the condition appears.

The minimum record for delegate real decision authority is anchored in actual proof. It includes a charter for the gate, named decision owner, recorded decision, and proof that resource commitments follow the outcome. The record also carries the choice supported by that proof: If the forum cannot change money, capacity, scope, or reliance, redesign the forum rather than adding ceremony. Keeping evidence and decision together lets a later owner distinguish a deliberate residual risk from a premise that nobody examined. It also gives the team a clear reason to reopen the choice when the underlying facts change.

Across the work governed by portfolio boards, sponsors, PMO leaders, and investment owners, occurrences of this pattern should be compared rather than treated as unrelated project stories. The positive standard is: The gate forum needs explicit authority to continue, redirect, narrow, pause, or stop the next commitment. If the same failure recurs, leaders can change delegation, capacity, intake, policy, or the intervention itself. That portfolio response is usually more valuable than asking each team to absorb another local workaround.

Two questions test delegate real decision authority. What evidence would show that the condition is no longer adequate? Would that evidence arrive early enough to preserve a feasible choice? The intended choice is clear: If

the forum cannot change money, capacity, scope, or reliance, redesign the forum rather than adding ceremony. A weak answer to either question calls for a smaller commitment, an earlier signal, or a safer fallback. A strong answer earns movement without claiming that uncertainty has disappeared.

Ask for evidence proportionate to commitment

Evidence should answer the risk created by the next tranche, not satisfy the same checklist at every stage. A strong operating model makes the hidden negotiation explicit while options are still available.

Teams produce large packs of low-value material while the decisive uncertainty—customer adoption, integration feasibility, capacity, or operating ownership—remains unresolved. A polished status can coexist with this failure because status describes activity; it does not prove the premise that authorizes reliance.

State the next commitment first, then ask what must be true and what evidence would make a skeptical decision maker comfortable extending it. The aim is not certainty. It is a justified next commitment with an observable basis and a viable way back.

A concise evidence map linking each material premise to a test, result, owner, and unresolved uncertainty. Measures should expose changes in the operation rather than reward production of artifacts. Activity can support the story; it cannot substitute for the result.

Advance reversible exploration with modest evidence; demand stronger proof before irreversible spend, dependency, or enterprise reliance. When the premise is weak, narrowing is often a better decision than forcing a binary choice between full approval and cancellation.

A tabletop review of ask for evidence proportionate to commitment should use the observed break as its scenario: Teams produce large packs of low-value material while the decisive uncertainty—customer adoption, integration feasibility, capacity, or operating ownership—remains unresolved. Participants then apply the proposed response. They state the next commitment first, then ask what must be true and what evidence would make a skeptical decision maker comfortable extending it. The review identifies where authority, information, time, or recovery would run out. It is useful only if it changes the boundary or confirms who will act when the condition appears.

The minimum record for ask for evidence proportionate to commitment is anchored in actual proof. It includes a concise evidence map linking each material premise to a test, result, owner, and unresolved uncertainty. The record also carries the choice supported by that proof: Advance reversible exploration with modest evidence; demand stronger proof before irreversible spend, dependency, or enterprise reliance. Keeping evidence and decision together lets a later owner distinguish a deliberate residual risk from a premise that nobody examined. It also gives the team a clear reason to reopen the choice when the underlying facts change.

Across the work governed by portfolio boards, sponsors, PMO leaders, and investment owners, occurrences of this pattern should be compared rather than treated as unrelated project stories. The positive standard is: Evidence should answer the risk created by the next tranche, not satisfy the same checklist at every stage. If the same failure recurs, leaders can change delegation, capacity, intake, policy, or the intervention itself. That portfolio response is usually more valuable than asking each team to absorb another local workaround.

Two questions test ask for evidence proportionate to commitment. What evidence would show that the condition is no longer adequate? Would that evidence arrive early enough to preserve a feasible choice? The intended choice is clear: Advance reversible exploration with modest evidence; demand stronger proof before irreversible spend,

dependency, or enterprise reliance. A weak answer to either question calls for a smaller commitment, an earlier signal, or a safer fallback. A strong answer earns movement without claiming that uncertainty has disappeared.

Keep all five outcomes credible

Advance is one valid outcome among several; redirection, narrowing, pausing, and stopping are evidence of governance working. Teams move faster when uncertainty is bounded honestly than when it is concealed inside an optimistic plan.

Organizations call every gate successful only when it passes, teaching teams to defend the original plan rather than reveal what has changed. Under pressure, the workaround is predictable: capable people compensate, local decisions proliferate, and the formal record stops describing how the operation actually runs.

Define the consequences of each outcome and protect teams that bring disconfirming evidence early. The design should state what remains unknown. Acknowledged uncertainty can be bounded; disguised uncertainty is inherited by the operation.

Examples of resources released or reallocated, scope changed, experiments commissioned, and conditions attached to continuation. Where consequence is high, independent checking or cross-functional challenge reduces the risk that advocacy becomes its own proof.

Choose the path that preserves the most value and optionality, not the path that best protects the original narrative. Escalation is appropriate when authority is missing, not as a routine substitute for clear delegation.

A tabletop review of keep all five outcomes credible should use the observed break as its scenario: Organizations call every gate successful only when it passes, teaching teams to defend the original plan rather than reveal what has changed. Participants then apply the proposed response. They define the consequences of each outcome and protect teams that bring disconfirming evidence early. The review identifies where authority, information, time, or recovery would run out. It is useful only if it changes the boundary or confirms who will act when the condition appears.

The minimum record for keep all five outcomes credible is anchored in actual proof. It includes examples of resources released or reallocated, scope changed, experiments commissioned, and conditions attached to continuation. The record also carries the choice supported by that proof: Choose the path that preserves the most value and optionality, not the path that best protects the original narrative. Keeping evidence and decision together lets a later owner distinguish a deliberate residual risk from a premise that nobody examined. It also gives the team a clear reason to reopen the choice when the underlying facts change.

Across the work governed by portfolio boards, sponsors, PMO leaders, and investment owners, occurrences of this pattern should be compared rather than treated as unrelated project stories. The positive standard is: Advance is one valid outcome among several; redirection, narrowing, pausing, and stopping are evidence of governance working. If the same failure recurs, leaders can change delegation, capacity, intake, policy, or the intervention itself. That portfolio response is usually more valuable than asking each team to absorb another local workaround.

Two questions test keep all five outcomes credible. What evidence would show that the condition is no longer adequate? Would that evidence arrive early enough to preserve a feasible choice? The intended choice is clear: Choose the path that preserves the most value and optionality, not the path that best protects the original

narrative. A weak answer to either question calls for a smaller commitment, an earlier signal, or a safer fallback. A strong answer earns movement without claiming that uncertainty has disappeared.

Separate evidence owners from advocates

The sponsor may advocate for the investment, but critical evidence needs an owner accountable for its integrity. The practical failure occurs between the formal approval and the ordinary pressure of running the work.

Forecasts, readiness claims, and benefit assumptions are prepared only by the team whose funding depends on approval. The happy path can survive this weakness. Ordinary variation exposes it: a disputed input, absent owner, competing commitment, or case that falls outside the assumed rule.

Assign independent or cross-functional challenge for high-consequence premises and disclose assumptions rather than burying them in appendices. The smallest credible control is the one that changes behavior when the evidence is weak. Anything else is commentary.

Traceable sources, reference-class comparisons where available, sensitivity ranges, and documented dissent or uncertainty. Qualitative observation belongs beside quantitative measures when workflow behavior, adoption, or judgment explains why the number moved.

Commission a test or reduce the commitment when the forum cannot distinguish advocacy from evidence. A stop or rollback is evidence that the control works when it protects greater value elsewhere in the portfolio.

A tabletop review of separate evidence owners from advocates should use the observed break as its scenario: Forecasts, readiness claims, and benefit assumptions are prepared only by the team whose funding depends on approval. Participants then apply the proposed response. They assign independent or cross-functional challenge for high-consequence premises and disclose assumptions rather than burying them in appendices. The review identifies where authority, information, time, or recovery would run out. It is useful only if it changes the boundary or confirms who will act when the condition appears.

The minimum record for separate evidence owners from advocates is anchored in actual proof. It includes traceable sources, reference-class comparisons where available, sensitivity ranges, and documented dissent or uncertainty. The record also carries the choice supported by that proof: Commission a test or reduce the commitment when the forum cannot distinguish advocacy from evidence. Keeping evidence and decision together lets a later owner distinguish a deliberate residual risk from a premise that nobody examined. It also gives the team a clear reason to reopen the choice when the underlying facts change.

Across the work governed by portfolio boards, sponsors, PMO leaders, and investment owners, occurrences of this pattern should be compared rather than treated as unrelated project stories. The positive standard is: The sponsor may advocate for the investment, but critical evidence needs an owner accountable for its integrity. If the same failure recurs, leaders can change delegation, capacity, intake, policy, or the intervention itself. That portfolio response is usually more valuable than asking each team to absorb another local workaround.

Two questions test separate evidence owners from advocates. What evidence would show that the condition is no longer adequate? Would that evidence arrive early enough to preserve a feasible choice? The intended choice is clear: Commission a test or reduce the commitment when the forum cannot distinguish advocacy from evidence. A weak answer to either question calls for a smaller commitment, an earlier signal, or a safer fallback. A strong answer earns movement without claiming that uncertainty has disappeared.

Make conditions expire

A gate decision is based on evidence that can age; approval should identify the assumptions that would reopen it. The visible artifact is often complete before the operating condition behind it is credible.

A project treats a six-month-old approval as permission after costs, leadership, dependencies, or market conditions have materially changed. The organization may call this execution risk, yet the failure began when authority and evidence were separated at the decision point.

Attach triggers and a latest-valid date to important conditions, with a named owner responsible for bringing the work back. A useful test asks what ordinary operating pressure would break the premise and how that break would be detected before consequence spreads.

Trigger thresholds, dependency watch items, benefit assumptions, and evidence that the next review occurs before options disappear. Evidence should arrive before the latest useful decision date; perfect analysis delivered after options expire has little governance value.

Reopen, narrow, or pause when the basis for approval no longer describes the work. The operating owner should confirm that the decision can be executed with the people, systems, and time actually available.

A tabletop review of make conditions expire should use the observed break as its scenario: A project treats a six-month-old approval as permission after costs, leadership, dependencies, or market conditions have materially changed. Participants then apply the proposed response. They attach triggers and a latest-valid date to important conditions, with a named owner responsible for bringing the work back. The review identifies where authority, information, time, or recovery would run out. It is useful only if it changes the boundary or confirms who will act when the condition appears.

The minimum record for make conditions expire is anchored in actual proof. It includes trigger thresholds, dependency watch items, benefit assumptions, and evidence that the next review occurs before options disappear. The record also carries the choice supported by that proof: Reopen, narrow, or pause when the basis for approval no longer describes the work. Keeping evidence and decision together lets a later owner distinguish a deliberate residual risk from a premise that nobody examined. It also gives the team a clear reason to reopen the choice when the underlying facts change.

Across the work governed by portfolio boards, sponsors, PMO leaders, and investment owners, occurrences of this pattern should be compared rather than treated as unrelated project stories. The positive standard is: A gate decision is based on evidence that can age; approval should identify the assumptions that would reopen it. If the same failure recurs, leaders can change delegation, capacity, intake, policy, or the intervention itself. That portfolio response is usually more valuable than asking each team to absorb another local workaround.

Two questions test make conditions expire. What evidence would show that the condition is no longer adequate? Would that evidence arrive early enough to preserve a feasible choice? The intended choice is clear: Reopen, narrow, or pause when the basis for approval no longer describes the work. A weak answer to either question calls for a smaller commitment, an earlier signal, or a safer fallback. A strong answer earns movement without claiming that uncertainty has disappeared.

How the elements interact

Delegate real decision authority and Ask for evidence proportionate to commitment protect different parts of the same commitment. The gate forum needs explicit authority to continue, redirect, narrow, pause, or stop the next commitment. Evidence should answer the risk created by the next tranche, not satisfy the same checklist at every stage. This is where a management ritual either becomes a real control or reveals itself as administrative theater. If the operation encounters this failure—participants advise while the actual decision sits elsewhere, so the project interprets discussion as approval and moves on.—the response for ask for evidence proportionate to commitment cannot compensate for the missing delegate real decision authority obligation. Reviewers should reconcile the evidence for both elements before they accept the next action: advance reversible exploration with modest evidence; demand stronger proof before irreversible spend, dependency, or enterprise reliance.

Ask for evidence proportionate to commitment and Keep all five outcomes credible protect different parts of the same commitment. Evidence should answer the risk created by the next tranche, not satisfy the same checklist at every stage. Advance is one valid outcome among several; redirection, narrowing, pausing, and stopping are evidence of governance working. The issue is not a shortage of activity; it is a shortage of inspectable commitments. If the operation encounters this failure—teams produce large packs of low-value material while the decisive uncertainty—customer adoption, integration feasibility, capacity, or operating ownership—remains unresolved.—the response for keep all five outcomes credible cannot compensate for the missing ask for evidence proportionate to commitment obligation. Reviewers should reconcile the evidence for both elements before they accept the next action: choose the path that preserves the most value and optionality, not the path that best protects the original narrative.

Keep all five outcomes credible and Separate evidence owners from advocates protect different parts of the same commitment. Advance is one valid outcome among several; redirection, narrowing, pausing, and stopping are evidence of governance working. The sponsor may advocate for the investment, but critical evidence needs an owner accountable for its integrity. A strong operating model makes the hidden negotiation explicit while options are still available. If the operation encounters this failure—organizations call every gate successful only when it passes, teaching teams to defend the original plan rather than reveal what has changed.—the response for separate evidence owners from advocates cannot compensate for the missing keep all five outcomes credible obligation. Reviewers should reconcile the evidence for both elements before they accept the next action: commission a test or reduce the commitment when the forum cannot distinguish advocacy from evidence.

Separate evidence owners from advocates and Make conditions expire protect different parts of the same commitment. The sponsor may advocate for the investment, but critical evidence needs an owner accountable for its integrity. A gate decision is based on evidence that can age; approval should identify the assumptions that would reopen it. The work becomes governable when ownership, evidence, boundaries, and response are connected to a real choice. If the operation encounters this failure—forecasts, readiness claims, and benefit assumptions are prepared only by the team whose funding depends on approval.—the response for make conditions expire cannot compensate for the missing separate evidence owners from advocates obligation. Reviewers should reconcile the evidence for both elements before they accept the next action: reopen, narrow, or pause when the basis for approval no longer describes the work.

Make conditions expire and Delegate real decision authority protect different parts of the same commitment. A gate decision is based on evidence that can age; approval should identify the assumptions that would reopen it. The gate forum needs explicit authority to continue, redirect, narrow, pause, or stop the next commitment. The cost appears downstream, but the missing decision sits upstream. If the operation encounters this failure—a project

treats a six-month-old approval as permission after costs, leadership, dependencies, or market conditions have materially changed.—the response for delegate real decision authority cannot compensate for the missing make conditions expire obligation. Reviewers should reconcile the evidence for both elements before they accept the next action: if the forum cannot change money, capacity, scope, or reliance, redesign the forum rather than adding ceremony.

Put the model into the management cadence

A framework becomes useful when it changes the normal cadence of work. It should shape intake, funding, design, delivery, and operating review without creating a separate governance industry around itself. The following sequence is deliberately small:

EXHIBIT 2

Start with one decision, then calibrate from evidence

- 1 **Move 1**
Rewrite every gate agenda around the decision required, not the presentation sequence.
- 2 **Move 2**
Limit the core pack to the evidence that could change that decision.
- 3 **Move 3**
Record outcome, conditions, dissent, owner, and expiry in the same meeting.
- 4 **Move 4**
Audit whether capacity and funding actually moved after redirect, pause, or stop decisions.

- Rewrite every gate agenda around the decision required, not the presentation sequence.
- Limit the core pack to the evidence that could change that decision.
- Record outcome, conditions, dissent, owner, and expiry in the same meeting.
- Audit whether capacity and funding actually moved after redirect, pause, or stop decisions.

The first cycle should be treated as calibration. Reviewers should note which questions changed a decision, which evidence was unavailable, where authority was unclear, and which controls cost out of proportion to the exposure. That record improves a gate that earns the next commitment without pretending the first version will fit every workflow or investment.

Ownership should remain close to the consequence. The PMO or AI-governance function can maintain the method, surface cross-portfolio patterns, and challenge weak evidence. It should not absorb the business owner's accountability for the result or the executive's accountability for the commitment.

Measures that reveal whether the control works

- Commitments narrowed, redirected, paused, or stopped before avoidable exposure grew.
- Exceptions or assumptions discovered after the latest useful decision date.
- Scarce specialist and executive attention consumed by recurring unresolved conditions.
- Decisions reopened because the original evidence, boundary, or rationale was unclear.
- Operating outcomes and recovery performance after reliance expanded.
- Time from a material signal to a consequential decision.

These measures are diagnostic, not a universal scorecard. Their purpose is to identify where the operating model fails repeatedly and whether governance is preserving options earlier. A lower count can indicate improvement, but it can also indicate weak detection. Leaders should read the measures with case evidence and frontline observation.

The strongest counterargument

Strong gates can still become slow or political. The answer is not maximum documentation; it is a small number of consequential reviews, clear delegation, and evidence scaled to reversibility and exposure. Routine low-risk decisions should stay with the team.

A reasonable critic could also argue that experienced leaders already do this through judgment. Many do. The weakness is not the absence of judgment; it is the absence of a durable operating record when people, pressure, and conditions change. A gate that earns the next commitment makes the basis for judgment visible enough to challenge, execute, and revisit.

The shift

A stage gate governs only when it connects evidence to a consequential choice. If every project advances regardless of what the review finds, the gate is a status meeting wearing the language of control. The practical shift is from reporting activity to governing the conditions under which the organization relies on that activity.

That discipline does not make uncertainty disappear. It gives leaders a better place to stand: a named decision, evidence proportionate to consequence, an explicit boundary, a responsible owner, and a response when reality departs from the assumption. Those elements are enough to turn hidden operating risk into a choice the portfolio can make deliberately.

Notes and sources

[1] Stage-Gate International, "Gate Governance." <https://www.stage-gate.com/solutions/services/gate-governance/>

[2] Bent Flyvbjerg, "Quality Control and Due Diligence in Project Management," 2013. <https://arxiv.org/abs/1302.2544>

About the author

Marco Policani is an enterprise portfolio, PMO, and AI operating-governance leader. He builds portfolio governance systems where AI carries the analytical load and named humans own every decision — an approach documented across case studies, walkthroughs, and working governance guides on his portfolio site.

Portfolio & governance library: policani.net/governance · Full portfolio: policani.net · LinkedIn: linkedin.com/in/marcpolicani

This white paper may be shared freely with attribution. © 2026 Marco Policani.