

Stage Gates Work Only When They Can Kill Projects

How to turn project reviews into real choices about money and people



Marco Policani

Enterprise Portfolio · PMO · AI Operating Governance

policani.net · linkedin.com/in/marcpolicani

July 2026

EXECUTIVE SUMMARY

A stage gate earns its place only when it can continue, redirect, narrow, pause, or stop the next commitment of money and people to a project.

A stage gate that cannot stop anything is worse than no review: each pass makes a weak project look scrutinized while early doubts harden into sunk cost. Five properties turn the gate into a real decision about the next commitment: a forum with authority to act, evidence sized to that commitment, every outcome from continue to stop available in practice, key evidence checked by someone other than the project's advocates, and approval conditions that expire.

- 01** Judge your gates by what they changed in the last year: money, people, scope, or what other programs rely on. If nothing material changed, the gates are ceremony.
- 02** Ask for the evidence the next commitment actually depends on, not a standing checklist, and read a very long gate pack as a sign that one question is still open.
- 03** Give every approval the conditions that would reopen it and a date after which it lapses, so an old approval does not become permanent permission.

The meeting that can approve but not refuse

Ask a portfolio board when its last review changed a project's money, people, or scope. If no one can name a material case, the board may be tracking delivery rather than governing the investment.

A stage gate is meant to review evidence before a project moves from one stage to the next. In practice, many boards treat it as a broad choice to approve or reject the project as a whole.

The stage-gate method is well established in product development. Its own guidance describes gates as decision points, not progress updates. [1] Flyvbjerg's work on project due diligence also explains why leaders need an outside view. A forecast should be tested against what happened in similar work, not judged only by the team's plan. [2]

The distinction matters because many reviews can ask questions but cannot change the work. Leaders may request another slide or note a concern, yet the project keeps its funding, staff, and scope. Each pass then gives weak work the look of scrutiny without changing its path.

Large projects give leaders little room for that error. A McKinsey-Oxford study found that large IT projects delivered 56 percent less value than forecast on average. Review gates were outside the study's design, so it cannot assign that gap to weak governance. The result still shows what leaders can lose when early doubt hardens into a large commitment. [4]

Why momentum wins

Projects gather support before the evidence is complete: a sponsor makes a public case, a team works hard on the plan, and contracts are signed. Other programs begin to depend on the result, so stopping feels less like a choice and more like an admission.

EXHIBIT 1

Five properties separate a gate from a status meeting

- 1 **Put authority in the room**
The gate forum needs explicit authority to continue, redirect, narrow, pause, or stop the next commitment.
- 2 **Size evidence to the choice**
Evidence should answer the risk created by the next commitment, not satisfy the same checklist at every stage.
- 3 **Keep all five outcomes usable**
Continue is one valid outcome; redirect, narrow, pause, and stop must also work in practice.
- 4 **Separate advocacy from key evidence**
The sponsor may advocate for the investment, but critical evidence needs an owner accountable for its integrity.
- 5 **Make approval expire**
A gate decision is based on evidence that can age; approval should identify the assumptions that would reopen it.

Research on escalation of commitment helps explain this pattern. A meta-analysis covering 35 years of research found that personal responsibility and negative framing were linked with greater escalation. It also found a positive average sunk-cost effect, though the size varied widely across studies. [7]

A review must make change easier while the cost is still low. If approval is the only workable outcome, past effort will shape the next choice even when new evidence points elsewhere. This is a design problem, not a claim that sponsors cannot be trusted.

Investment theory adds a second reason to act early: an irreversible choice gives up the option to wait for more information, and that lost option is part of the cost. [8] A short test is cheap to change, while a large contract, a fixed launch date, or a dependency used by many teams is not.

Here is the shift: a gate should not approve the whole project. It should decide only the next commitment while the organization can still change course. That commitment may include money, scarce staff, a new dependency, or work that other programs will rely on.

Five design properties make that shift real.

1. Put real authority in the room

Once a gate decides only the next commitment, the review board needs written authority to continue, redirect, narrow, pause, or stop it. The board must be able to move more than words on a page.

The charter should answer four questions:

- What can this forum decide?
- Which choices still belong to a sponsor, resource owner, or executive committee?
- What attendance is required for a valid decision?
- What happens when two groups disagree?

The last question matters because security, architecture, finance, and executive sponsors may each control part of the choice. Without a clear path for conflict, a project can keep asking until someone says yes.

Test the authority with a modest case by having the forum set a real condition, such as a smaller scope or delayed hiring. If the project treats that condition as advice, the gate has no authority yet, so fix the charter before sending a larger decision through it.

2. Size the evidence to the next commitment

Authority alone does not say what evidence is enough. Start with the choice, not the template: state what the organization is about to commit, then ask which facts must be true for that choice to make sense.

A six-week test may need only a clear question, a safe limit, and a way to learn. A large spend needs stronger proof, as does a choice that other programs will rely on or that will be hard to undo.

Use a one-page evidence map at the front of the review pack. For each key fact, show:

- the claim the decision depends on;
- the test that was run;
- what the test found;
- the person who owns that evidence; and
- what remains unknown.

Everything else can sit behind that page as reference, which keeps the meeting on the few facts that could change the choice. A thick pack may be thorough, but it can also hide the one open question that matters.

The outside view belongs here: if five similar projects cost more or took longer than planned, the new estimate should explain why this case will differ. [2] Use that history to test the new estimate without treating a team's past results as punishment. Hope cannot be the only forecast.

3. Keep all five outcomes usable

Evidence can change a commitment only when the review board has more than one usable outcome. Continue is valid, along with redirect, narrow, pause, and stop. Teams will believe these choices are real only after they see them used.

Define the effect of each choice before the meeting. Ask where people go after a stop, how knowledge stays safe during a pause, and which promises change when scope narrows. If those answers do not exist, every choice except continue becomes a new fight.

Cooper and Edgett describe portfolio management as a live decision process. Projects may be sped up, stopped, or given lower priority, and resources may move among active work. [3] That movement is what makes a gate useful to the whole portfolio.

Redirect and narrow often do more work than a full stop. Evidence may show that one customer group has a need while another does not, that a platform needs two features rather than six, or that a rollout works at half the planned pace. A forum should be able to keep the part that still earns support and release the rest.

Protect the team that brings bad news because a group that finds a failed assumption before the next spend has saved money and time. If leaders treat that result as poor delivery, the next team will hide the signal, so credit the evidence and move people with care.

4. Separate advocacy from key evidence

Several outcomes are useless if the board cannot trust the evidence behind them. A sponsor should make the case for the project, but the same people should not own every forecast and readiness claim.

No bad faith is required for a conflict to matter because people see evidence through the goals they are trying to reach. The answer is a targeted independent check, not a second team that repeats all the work.

Finance can rebuild the benefit math, operations can test whether it has room to absorb the change, and architecture can review the integration claim. A customer group can check whether the stated need is real. Use these checks on facts that could change the decision.

Bain studied decision practices across more than 1,000 companies and reported a strong link between decision effectiveness and financial performance. The study was correlational and did not test one review method. [5] Its finding still supports a basic point: who decides and how the decision works deserve the same care as the project plan.

Record dissent in the decision with a short note that shows what a reviewer questioned, how the forum answered, and what would reopen the issue. Agreement has more meaning when the record also has room for doubt.

5. Make approval expire

Even a sound approval ages. A gate decision is based on evidence at one point in time, while costs, suppliers, sponsors, and demand can all change. Yet an old approval often stays in force as if the project had not changed.

EXHIBIT 2

The decision each property protects

Put authority in the room

If the forum cannot change money, capacity, scope, or reliance, redesign the forum rather than adding ceremony.

Size evidence to the choice

Advance reversible tests with modest evidence; demand stronger proof before irreversible spend, dependency, or enterprise reliance.

Keep all five outcomes usable

Choose the path that preserves the most value and future choice, not the path that best protects the original narrative.

Separate advocacy from key evidence

Commission a test or reduce the commitment when the forum cannot distinguish advocacy from evidence.

Make approval expire

Reopen, narrow, or pause when the basis for approval no longer describes the work.

Give each key condition a trigger and an end date. The trigger might be a cost increase, a supplier delay, a drop in use, or a missed safety result; the end date says when the work must return even if no trigger fires.

Keep the list to the three or four facts that carried the approval, and name one owner for each. That person brings the work back when the trigger is crossed.

For example, a release may depend on a supplier finishing an interface by June. The gate can allow the next step while setting a six-week delay as a trigger. If the supplier slips eight weeks, the team does not bury the change in a status report. The choice returns while unspent money and open options remain.

Expiry gives the other four properties a life between meetings. Authority, evidence, real outcomes, and independent checks should not vanish as soon as the board leaves the room.

Use gates to manage the portfolio

Every live review uses scarce leadership attention, so a board that reviews all work in the same way will soon skim the packs. That habit rewards polish over evidence.

Route low-risk, reversible work through a lighter path where the team files its evidence and reviewers can object without a meeting. Save live time for choices that are large, hard to reverse, or disputed.

Group similar choices when comparison helps. A board that reviews five requests for the same scarce engineers can ask which work deserves them most, a question that is hard to see when each project arrives alone.

A McKinsey study found that companies which moved an average of 56 percent of capital across business units over 15 years earned 30 percent higher total shareholder returns than companies that moved the least. The study examined capital shifts across business units, not project stops, and it cannot assign the return difference to one cause. Its finding makes the ability to move capital relevant to strategy. [6]

A gate has not finished when the minutes are approved; confirm where released money and people went, and how soon they became useful elsewhere.

What to do next

Do not redesign every review at once; choose one mid-sized decision with real stakes and room to recover.

Before the meeting:

- Write the next commitment in one sentence.
- List all five outcomes and what each would do to money, people, scope, and dependencies.
- Replace the standing checklist with a one-page evidence map.
- Confirm who can make the choice and who resolves a conflict.

During the meeting, record the decision, conditions, dissent, owner, trigger, and expiry date. After the meeting, check that money and people moved as directed.

Run the same test for two or three reviews, then ask what changed because this gate met. Look for a smaller scope, a changed release, a pause, a stop, or staff moved to stronger work.

If nothing material changed, fix the authority before adding more templates.

The strongest objections

A gate with real power can become slow or political. Use fewer live reviews, clear limits, and evidence matched to the size of the choice. Routine delivery updates should stay with the team.

Politics is a real risk because a board that can stop work can be used in a budget fight. Publish its authority, show the evidence, record dissent, and compare the case with similar work. Let an affected sponsor challenge the evidence or reopen the choice when a named condition changes, without bypassing the board. These steps make hidden tests harder to use.

Some leaders also fear that strong gates will make teams avoid bold ideas, but hidden standards create more fear than clear ones. A small, reversible test should face a light burden, while a large, fixed

commitment should face a high one. Teams can take smart risks when they know what proof the next step requires.

From reviewing projects to deciding commitments

A gate that cannot change a commitment was never governing it. Decide the next use of money, people, and trust instead of approving the project as a whole.

The aim is to change weak work sooner, while capacity still has somewhere better to go, rather than to stop more work. Projects that continue also gain something: an approval that came from a forum with the power to refuse it.

A gate earns its place when evidence can change the next commitment while the organization still has options.

Notes and sources

- [1] Stage-Gate International, "Gate Governance." <https://www.stage-gate.com/solutions/services/gate-governance/>
- [2] Bent Flyvbjerg, "Quality Control and Due Diligence in Project Management," 2013. <https://arxiv.org/abs/1302.2544>
- [3] Robert G. Cooper and Scott J. Edgett, "Portfolio Management: Fundamental for New Product Success," Stage-Gate International. <https://www.stage-gate.com/blog/portfolio-management-fundamental-for-new-product-success/>
- [4] Michael Bloch, Sven Blumberg, and Jurgen Laartz, "Delivering large-scale IT projects on time, on budget, and on value," McKinsey & Company, 2012. <https://www.mckinsey.com/capabilities/tech-and-ai/our-insights/delivering-large-scale-it-projects-on-time-on-budget-and-on-value>
- [5] Marcia Blenko, Michael Mankins, and Paul Rogers, "The Five Steps to Better Decisions," Bain & Company, August 6, 2013. <https://www.bain.com/insights/the-five-steps-to-better-decisions/>
- [6] Stephen Hall, Dan Lovullo, and Reinier Musters, "How to Put Your Money Where Your Strategy Is," McKinsey & Company, March 1, 2012. <https://www.mckinsey.com/capabilities/strategy-and-corporate-finance/our-insights/how-to-put-your-money-where-your-strategy-is>
- [7] Dustin J. Sleesman, Donald E. Conlon, Gerry McNamara, and Jonathan E. Miles, "Cleaning Up the Big Muddy: A Meta-Analytic Review of the Determinants of Escalation of Commitment," Academy of Management Journal 55(3), 2012, 541-562. [http://www.iot.ntnu.no/innovation/norsi-pims-courses/huber/Sleesman,%20Conlon%20&%20McNamara%20\(2012\).pdf](http://www.iot.ntnu.no/innovation/norsi-pims-courses/huber/Sleesman,%20Conlon%20&%20McNamara%20(2012).pdf)
- [8] Robert S. Pindyck, "Irreversibility, Uncertainty, and Investment," Journal of Economic Literature 29(3), 1991, 1110-1148; NBER Working Paper 3307, 1990. <https://www.nber.org/papers/w3307>

About the author

Marco Policani is an enterprise portfolio, PMO, and AI operating-governance leader. He builds portfolio governance systems where AI carries the analytical load and named humans own every decision — an approach documented across case studies, walkthroughs, and working governance guides on his portfolio site.

Portfolio & governance library: policani.net/governance · Full portfolio: policani.net · LinkedIn: linkedin.com/in/marcpolicani

This white paper may be shared freely with attribution. © 2026 Marco Policani.