
WHITE PAPER · PORTFOLIO & DELIVERY GOVERNANCE

Stage Gates Should Decide the Next Commitment, Not the Whole Project

Giving project reviews real power over money and people

Marco Policani

Enterprise Portfolio · PMO · AI Operating Governance

policani.net · linkedin.com/in/marcpolicani

July 2026

EXECUTIVE SUMMARY

A stage gate has real power only when it decides the next commitment of money, people, or contracts instead of approving the whole project, while the board can still change that commitment.

A review that approves the whole project can say yes once and then only watch. Contracts are signed, people are hired, and other programs plan around the project before the evidence arrives. This paper moves the decision to the next commitment of money, people, or contracts, sized so that it can still be changed. The board decides that step with written authority to act, evidence that fits its size, key facts that the sponsor does not own, and an approval that runs out.

- 01** Judge a gate by whether it could change what a project receives, not by how many projects it stops. If it could not, fix the charter first.
- 02** List the steps a project will take and mark each as cheap or costly to reverse. Approve the cheap work up to the next costly step, and review that step on its own.
- 03** Give every approval an end date, the three or four facts it rests on, and a trigger that sends the decision back to the board.

A review that can only say yes

A stage gate is a review where leaders decide whether a project moves to its next stage. Ask a portfolio board when its last review changed how much money, how many people, or how much scope a project received. If no one can name a case, the review may be tracking progress rather than governing the investment.

A board with well-chosen projects may rarely stop anything, so the test is not a count of stops. The test is whether the review could have changed what the project was given, and whether it ever did. It is a test of one board, not a claim about how common the problem is.

Robert Cooper and Scott Edgett, who developed the Stage-Gate method, describe the problem from the other side. They observe that projects often reach advanced development before anyone examines them closely, and that once work has begun

there is very little chance it will be killed.[1] That is a practitioner's observation, not a measured rate. It still names the problem this paper addresses.

A review fails in one of two ways. It can lack the authority to act: leaders ask for another slide or note a concern, and the project keeps its funding, staff, and scope. It can also decide the wrong thing, voting yes or no on the whole project, once, so that an early yes leaves the board little to decide later.

Cooper and Edgett's remedy is tougher go or kill decisions at the gates. Their method already commits money one stage at a time, which they call incremental commitment, and its gates decide the next stage along with its resources.[1] The open question is where one stage ends and the next begins. A stage drawn around a phase of work, such as design or build, can end after the point where reversing a step has become expensive. This paper argues that a gate should sit before the first step that is hard to reverse and decide that step on its own. That step is the next commitment: a step that is hard to reverse because it spends a sum the board would not want to lose, pulls scarce people from other work, signs a contract, or lets other programs plan around the result. Decided that way, every review is a fresh choice, and the board has something left to decide.

Why a yes to the whole project is hard to take back

A project can gather support before the evidence is in. A sponsor makes the public case, the team builds a plan, and contracts are signed. Other programs schedule work that depends on the project. After that, stopping means reversing all of it, which feels more like admitting a mistake than making a choice.

Economists describe a cost that a budget does not show. Pindyck's analysis of investment under uncertainty says that an irreversible choice gives up the chance to wait for more information, and that the lost chance is part of the cost.[2] His work is theory, not a study of project boards. It applies in a plain way here. A six-week test is cheap to change when new evidence arrives. A multi-year contract, a fixed launch date, or a platform that many teams build on is not.

People add a second pressure. A meta-analysis covering 35 years of research on escalation of commitment, the tendency to keep backing a failing course, found that personal responsibility was linked with greater escalation. It also found that the average effect of money already spent was positive, though its size varied widely across studies.[3] These are mostly studies of individual decisions, not boards. They help explain why sponsors and committees find it hard to reverse an earlier yes. They do not show that it always happens.

Both pressures are likely strongest after a yes, which is this paper's inference rather than a finding. By then contracts are signed, and the people who approved the project are the ones asked to reverse it. Authority alone does not fix a gate that meets at that point, because the board holds power over a decision already made. The fix starts with what the gate decides.

Decide the next commitment

Almost every project takes a series of steps. Some are cheap to reverse, such as a short test with a small team and no contract. The rest are commitments in the sense used here. A step is costly to reverse when undoing it would lose more than the board is willing to lose on one test, or breaks a plan that other teams depend on. The gate should sit before the first commitment, approve the cheap work that comes before it, and decide that commitment on its own, with the project's history in view.

EXHIBIT 1

How to find the next commitment

- ① **List the commitments**
Write down every step the project will take, in order, from tests to contracts and launch dates.
- ② **Mark the costly ones**
A short test is cheap to reverse. A multi-year contract or a shared platform is not.
- ③ **Approve up to the next one**
The board decides the cheap work before the next costly step, and nothing beyond it.
- ④ **Review the costly one alone**
Put stronger evidence in front of it, and ask whether its terms can be made smaller.

To find the next commitment, list the steps the project will take, in order, and mark each as cheap or costly to reverse. A multi-year contract, a fixed launch date, or a shared platform is costly. The gate approves the cheap work up to the next costly step. It decides that step on its own, with stronger evidence in front of it.

Here is one illustration of a rule that applies to any project. An insurer plans to rebuild its claims intake process. The plan calls for eighteen months, thirty engineers, and a contract with a software vendor. All figures are illustrative. In a whole-project review, the board approves the plan in month one. The vendor contract is signed in month

two. In month seven, after the team has built for all three customer groups, real use shows that one group does not need the new process. By then the contract, the hiring, and two other teams' schedules all assume the full plan. The board's real choice is between finishing and writing off the work, and finishing is the likely result.

In a next-commitment review, the first gate approves a ten-week trial with six people and no contract, run with each of the three customer groups. The trial shows in week ten that one group does not need the new process. The second gate decides the vendor contract, narrowed to the two groups the trial supported, with a clause that lets the insurer exit after the first release. Both versions learn the same thing about that group. In the second, the board learns it before the contract exists. The illustration assumes that a small trial can reveal this kind of finding. Where it cannot, the gate needs the strongest evidence it can get before the costly step.

Some commitments cannot be split. A vendor may quote a price only for a multi-year term. When that happens, treat that decision as the real gate. Ask for the strongest evidence there, and ask whether the terms can be made smaller with a shorter first term, phased pricing, or an exit clause.

Staging is not free. Each extra review takes leaders' time, a pause for a decision can delay the work, and a vendor may charge more for a shorter term. Waiting has a price as well as a value. Add a gate only where reversing later costs clearly more than the delay now. Route low-risk, reversible work through a lighter path, where the team files its evidence and reviewers can object without a meeting. Save live time for choices that are large, hard to reverse, or disputed.

What the board needs to decide a commitment

A board that decides only the next commitment still needs three things to make that decision real. Without them, the review is back to a formality.

The board can act, and each action has a defined effect

The board needs written authority to continue, redirect, narrow, pause, or stop the next commitment. Its charter should answer four questions:

- What can this board decide?
- Which choices still belong to a sponsor, a resource owner, or an executive committee?
- Who must attend for a decision to count?
- Who settles a disagreement between groups?

The last question matters because security, architecture, finance, and executive sponsors may each control part of the choice. Without an answer, a project can keep asking until someone says yes. To test the authority, have the board set a real condition on a modest decision, such as a smaller scope or a delayed hire. If the project treats that condition as advice, fix the charter before sending a larger decision through it.

Define the effect of each outcome before the meeting. Ask where people go after a stop, how the work is kept safe during a pause, and which promises change when scope narrows. If those answers do not exist, every choice except continue becomes a new fight.

Cooper and Edgett describe a portfolio as a list of active projects that is constantly updated and revised. Projects may be sped up, killed, or given lower priority, and resources are allocated and re-allocated among them.[4] Redirect and narrow can keep value that a full stop would lose. Evidence may show that one customer group has a need and another does not, or that a platform needs two features rather than six. A board that can keep the part that still earns support and release the rest has more to offer than one that can only say yes or no.

The evidence fits the size of the commitment

Authority is wasted if the board cannot tell whether the next step is justified. Start with the choice, not the template. State what the organization is about to commit, then ask which facts must be true for that choice to make sense. A six-week test needs a clear question, a safe limit, and a way to learn. A large spend needs stronger proof, as does a choice that other programs will depend on or that will be hard to undo.

Put a one-page evidence map at the front of the review pack. The project team writes it for this decision, and the next gate starts a new one. For each key fact, it shows the claim the decision depends on, the test that was run, what the test found, the person who owns that evidence, and what remains unknown. Everything else can sit behind that page as reference. A thick pack may be thorough, but it can also hide the one open question that matters.

For the key estimates, Flyvbjerg proposes comparing a forecast with the actual results of similar projects, because front-end estimates often diverge from what happens.[5] If five similar projects cost more or took longer than planned, the new estimate should explain why this case will differ.

Someone other than the sponsor owns the key evidence

Evidence the board cannot trust is no use either. A sponsor should make the case for the project, but the same people should not own every forecast and readiness claim. No bad faith is required. The link between personal responsibility and escalation in the research above suggests that a sponsor who is accountable for a project may be slower to act on evidence against it.

The answer is a targeted check, not a second team that repeats all the work. Finance can rebuild the benefit math. Operations can test whether it has room to absorb the change. Architecture can review the integration claim, and a customer group can confirm that the stated need is real. Use these checks only on facts that could change the decision, and record dissent in the decision itself with a note on what a reviewer questioned and what would reopen the issue. Who should decide after a project misses the result it agreed to reach is the subject of a companion paper, *Why the Kill Gate Doesn't Fire*.

Let the approval run out

A decision on the next commitment also has a time limit. Even a sound approval ages. It rests on evidence from one moment, while costs, suppliers, sponsors, and demand change. Yet an old approval often stays in force as if the project had not changed.

EXHIBIT 2

What each control protects

Authority and outcomes

The board can change what the next step receives, and every outcome has a defined effect.

Evidence that fits

Proof grows with the size of the next commitment and the difficulty of reversing it.

Key facts the sponsor does not own

Someone else checks the facts that could change the decision.

An approval that runs out

An end date and triggers bring the decision back while options remain.

Give each approval an end date and the three or four facts it rests on. Name one owner for each fact and a trigger that sends the decision back to the board. A trigger

might be a cost increase, a supplier delay, a change in leadership, or a drop in demand. The end date says when the work must return even if no trigger fires.

For example, a release may depend on a supplier finishing an interface by June. The board can approve the next step while setting a six-week delay as the trigger. If the supplier slips eight weeks, the team does not bury the change in a status report. The choice returns to the board while unspent money and open options remain.

The end date is what makes a commitment a step rather than a blank check. The board has approved this step, for this period, on these assumptions. When the date or a trigger arrives, the next review is already on the calendar. This control covers the assumptions behind an approval. A missed result target is a different trigger, and who should decide then is the question the companion paper answers.

Where this goes wrong

A board with real power can become slow or political. Keep the number of live reviews small, set clear limits on what each can decide, and match the evidence to the size of the choice. Routine delivery updates stay with the team.

Politics is a real risk, because a board that can stop work can be used in a budget fight. Publish its authority, show the evidence, record dissent, and compare the case with similar work. Let an affected sponsor challenge the evidence, or ask the board to reopen a decision when a named assumption changes. The challenge goes through the board and never around it.

Some leaders fear that strong gates will make teams avoid bold ideas. A small, reversible test should face a light burden, and a large, fixed commitment a high one. Teams can plan around a standard they can see.

What to do next

Do not redesign every review at once. Choose one mid-sized decision with real stakes and room to recover.

Before the meeting:

- Write the next commitment in one sentence: what money, which people, which contract or dependency, and until when.
- List the five outcomes and what each would do to money, people, scope, and dependencies.
- Replace the standing checklist with a one-page evidence map.

- Confirm who can make the choice and who settles a conflict.

During the meeting, record the decision, the assumptions behind it with their triggers and end date, any dissent, and the owner. After the meeting, check that money and people moved as directed, and how soon the freed people and money went to work on higher-value projects.

After two or three reviews, compare what the board approved at each review with what was requested. Look for a smaller scope, a delayed hire, a changed release date, a pause, a stop, or staff moved to stronger work. Do not score the board by its number of stops. A board that stops work to prove it can is as wrong as one that never does.

If nothing ever changes and the board could not have changed it, fix the authority and the size of the commitment before adding more templates. If the board could have changed something and the projects were sound, the review is working.

From approving projects to deciding commitments

A gate that decides the whole project approves it once and then watches. A gate that decides the next commitment gets to decide again, while money and people can still go elsewhere.

The aim is not more stops. It is a yes that means something because it could have been a no, a narrower scope, or a later date.

Notes and sources

[1] Robert G. Cooper and Scott J. Edgett, Stage-Gate and the Critical Success Factors for New Product Development, BPTrends. <https://bptrends.info/wp-content/publicationfiles/07-06-ART-Stage-GateForProductDev-Cooper-Edgett1.pdf>

[2] Robert S. Pindyck, "Irreversibility, Uncertainty, and Investment," Journal of Economic Literature 29(3), 1991, 1110–1148; NBER Working Paper 3307, 1990. <https://www.nber.org/papers/w3307>

[3] Dustin J. Sleesman, Donald E. Conlon, Gerry McNamara, and Jonathan E. Miles, "Cleaning Up the Big Muddy: A Meta-Analytic Review of the Determinants of Escalation of Commitment," Academy of Management Journal 55(3), 2012, 541–562. [http://www.iot.ntnu.no/innovation/norsi-pims-courses/huber/Sleesman,%20Conlon%20&%20McNamara%20\(2012\).pdf](http://www.iot.ntnu.no/innovation/norsi-pims-courses/huber/Sleesman,%20Conlon%20&%20McNamara%20(2012).pdf)

[4] Robert G. Cooper and Scott J. Edgett, "Portfolio Management: Fundamental for New Product Success," Stage-Gate International. <https://www.stage-gate.com/blog/portfolio-management-fundamental-for-new-product-success/>

[5] Bent Flyvbjerg, "Quality Control and Due Diligence in Project Management," 2013. <https://arxiv.org/abs/1302.2544>

About the author

Marco Policani is an enterprise portfolio, PMO, and AI operating-governance leader. He builds portfolio governance systems where AI carries the analytical load and named humans own every decision — an approach documented across case studies, walkthroughs, and working governance guides on his portfolio site.

Portfolio & governance library: policani.net/governance · Full portfolio: policani.net · LinkedIn: linkedin.com/in/marcpolicani

This white paper may be shared freely with attribution. © 2026 Marco Policani.