


Standardize the Component, Not the Project

How an EPMO helps teams share a useful check without copying a whole project



Marco Policani

Enterprise Portfolio · PMO · AI Operating Governance

policani.net · linkedin.com/in/marcpolicani

August 2026

EXECUTIVE SUMMARY

Share a piece of project practice only when evidence shows it travels, and give it an owner who can change or retire it.

A shared folder preserves files but does not tell the next team which practice to use. GAO found that NASA's lessons system gave no assurance that later missions applied what teams learned. A project-to-standard review asks which small part can travel, what may vary, and who keeps it current.

- 01** Standardize a small, reusable piece, such as an intake question, decision record, readiness check, or handoff rule, rather than the whole project.
- 02** Promote a piece only after another project has applied it in live work, someone owns it, its allowed variations are written down, workarounds are tracked, and there is a path to change or retire it.
- 03** When projects vary a standard, publish the valid variations and treat repeated departures as a signal to review the shared version.
- 04** The EPMO brings the right owners together, records the promotion decision, and watches patterns across projects. Subject owners approve the shared rule; project leads judge local fit.

CONTENTS

A project folder is evidence, not a standard
 The reusable unit is smaller than the project
 Promotion must be earned
 The project-to-standard review
 Circulation is not evidence of reuse
 A bounded operating illustration

Variation is design data
 A shared asset needs a lifecycle
 Keep the EPMO in its lane
 What this evidence cannot establish
 The standard is the managed component
 Notes and sources

A project folder is evidence, not a standard

In one release portfolio I supported, teams entered finance testing with different scopes, calendars, and test needs. Each handoff raised the same question: what had to be ready before testing could begin? Copying one project's plan would carry details that did not fit the next release. Leaving every team to decide alone

kept that question open. Which entry checks could travel without turning every release into the same project?

That distinction is easy to lose. A project succeeds, its documents are placed in a shared folder, and the organization calls the folder a standard. Another project copies the files, changes half of them, and creates a second version. Soon the repository contains several artifacts with similar names, uncertain owners, conflicting rules, and no clear way to know which one should be trusted. Activity has been preserved. Learning has not.

Rowe and Sikes's PMI conference paper shows what a folder misses. Teams should capture lessons as they work, study them, and turn them into action. Leaders should track agreed actions until they are done.[1] A folder can store the lesson. It cannot make a later team use it.

NASA had a lessons system and told managers to use it. Yet GAO found no assurance that lessons shaped later missions. It also found that teams did not routinely find, collect, or share lessons. GAO called for stronger management support, incentives, coordination, and easier ways to use what teams learned.[2]

The EPMO's opportunity is therefore smaller and more useful than standardizing projects. It is to notice which parts of project work keep recurring, decide whether those parts have earned reuse, and govern them as operating assets. The unit of learning is not the whole project plan. It is the component that can travel without carrying the entire project's context with it.

The reusable unit is smaller than the project

Projects combine common work with conditions that may never repeat. Sponsors, customers, and contracts differ. So do technologies, geographies, regulatory duties, risk levels, and delivery histories. Copying the project treats all of that context as if it were part of the reusable method. It produces a large template that local teams must either obey or quietly dismantle.

A component is narrower. It may be an intake question that exposes whether a sponsor can name the outcome. It may be a decision record that preserves the owner, evidence, alternatives, and next review point. It may be a source standard, a readiness check, an exception route, a measure definition, or a handoff rule. Each has a specific job. Each can be tested apart from the project that first produced it.

The boundary matters because reuse needs a stable core and an explicit edge. The core states what must remain true when the component moves. The edge states where local judgment begins. A decision record might always require a named decider and the evidence considered, while allowing each portfolio to add regulatory, commercial, or technical fields. A readiness check might always establish entry conditions, while letting teams vary the tests that prove those conditions.

Kirchmer's APQC report supports this approach. It calls for a shared pattern that allows differences when the work calls for them. Leaders also need a way to keep the pattern current.[3] The core and its approved local forms belong to one system.

This changes the EPMO's first question. Instead of asking, 'Which project template should everyone use?' it asks, 'Which operating component has repeated often enough that teams should not have to rediscover it?' The first question begins with control. The second begins with evidence.

Promotion must be earned

A useful document should not become a standard just because a leader likes it or one team did well. Nor should a PMO promote it to fill a template gap. Those signals show interest, not proof that it will help another team. Leaders must decide whether people can rely on the same part of the work in a new setting.

EXHIBIT 1

Promotion turns local learning into a managed component

01

Iterate

Keep the candidate close to the work while its purpose, boundary, and failure modes emerge.

02

Mature

Test the core in another setting and record what stayed stable, varied, or failed.

03

Canonize

Publish the supported core, owner, scope, variants, version, and known limits.

04

Lock

Route changes through an explicit acceptance, variant, rejection, or evidence decision.

05

Monitor

Watch reuse, exceptions, workarounds, obsolescence, and signals that justify retirement.

The decision can be managed through five states: iterate, mature, canonize, lock, and monitor. These are not industry maturity levels. They are Marco's operating model for deciding when local learning becomes a shared asset.

Iterate: keep the component close to the work

A new component begins as a local answer. The team can change it quickly because its purpose, boundary, and failure modes are still being discovered. The evidence at this stage is practical: Did it remove a repeated ambiguity? Did people understand it? Did it expose a decision sooner? Did it create extra work somewhere else? A useful first result is not enough to promote it.

Mature: test the core in another setting

A component matures when another project can use its core without copying the original project's assumptions. The second use is especially valuable because it separates the essential fields from local furniture. Differences are recorded, not treated as defects. The review asks whether the change is a legitimate variant, a sign that the core is wrong, or evidence that the candidate is too narrow to share.

Canonize: publish the supported core and its boundary

Canonization is the explicit promotion decision. The component receives a name, purpose, accountable owner, current version, and boundary of use. Approved variants and known limits sit beside the supported core. A canonized intake question does not become a complete intake process. A shared decision record does not replace the judgment that fills it.

Lock: place change under governance

Locking places the shared core under visible change control. Proposed changes follow a defined path. The owner can accept a change into the core, approve it as a variant, reject it with a reason, or return it for more evidence. Version history protects users from finding out too late that the same asset meant different things in different projects.

Monitor: watch use, exceptions, and obsolescence

A component remains a standard only while it remains useful. Monitoring looks for repeated exceptions, workarounds, unclear handoffs, abandoned fields, stale sources, and changes in the setting that produced it. The outcome is not always another version. It may be a narrower scope, a split into two components, a return to local use, or retirement.

The promotion sequence protects both sides of the argument. It prevents a project office from imposing an untested answer, and it prevents proven learning from remaining trapped inside the team that discovered it. Most candidates should spend time in iterate or mature. Some should never move further. A small number should earn enterprise reliance.

The project-to-standard review

The review needs to lead to a choice, not a new ceremony. It can fit into an existing portfolio or quality meeting. The proposed owner brings a short record of what teams tried. The group chooses one path: keep it local, ask for more proof, share it with clear limits, or retire it.

Seven questions keep that decision honest.

- Repeated component. What exact part of the work has appeared in more than one project? Name the component, not the program that produced it.
- Evidence of reuse. Where has another team used the core, and what became easier, clearer, safer, or more consistent? Record failures and abandoned uses as well as successful ones.
- Owner. Who can explain the asset, decide a change, communicate a new version, and retire it? A repository administrator is not automatically the operating owner.
- Approved variants. Which differences are legitimate because context changes, and which would break the component's purpose or evidence standard?
- Version and change path. How will users know the current version, submit a change, see the disposition, and understand whether they must migrate?
- Exception signal. Which repeated workaround, override, failure, or local edit should force the owner to review the core or its boundary?
- Retirement trigger. What would show that the component has become obsolete, redundant, too costly to maintain, or unsafe to rely on?

The questions are deliberately harder than 'Was the artifact uploaded?' They force the candidate to show a pattern, a use boundary, and a management path. They also make absence visible. A component with no owner is a document. A component with no evidence of reuse is a proposal. A component with no exception signal is likely to drift. A component with no retirement trigger is a permanent obligation created by accident.

The review record should be short enough to read at the table. It should show why the candidate was promoted, who owns it, which version is current, and what will reopen the decision. The approved scope and variants sit beside that decision. Longer narratives remain supporting evidence.

Circulation is not evidence of reuse

A common promotion error is to count distribution. A file was downloaded, linked from an intranet, attached to several project plans, or mentioned in training. Those events show exposure. They do not show that the component shaped the work. A team can attach a standard decision record and still

make the decision in an unrecorded side meeting. It can copy a readiness checklist after the release is already committed. It can cite an intake form while bypassing the questions that matter.

Evidence of reuse should connect the component to an observable part of the operating process. Did another team use the intake questions before accepting demand? Did the decision record reach the accountable owner while alternatives were still open? Did the source standard change what was accepted as evidence? Did the handoff rule identify the receiver and entry condition before work crossed the boundary? These observations do not have to be large metrics. They have to show use at the point where the component was designed to matter.

The evidence should also preserve friction. If a team used the core but needed a local field, that is variant evidence. If the team abandoned the component because it arrived after the decision, that is timing evidence. If users completed it but could not explain its purpose, that is a design failure. Promotion based only on successful adoption stories will canonize components whose hidden cost is pushed onto local teams.

A useful review separates exposure from operating evidence. Encountering the component establishes adoption; applying it at the intended point establishes use. The stronger evidence shows a change in a named ambiguity, handoff, or control condition. Continued maintenance and access then show whether reuse can last. The review needs enough of this chain to avoid confusing page views with operating value.

This keeps the EPMO away from a false numerical threshold. Ten downloads are not weaker than one hundred if ten projects used the asset in live decisions. Two observed uses may be informative when the settings are materially different. Five nearly identical uses inside one team may teach very little about transfer. The review weighs relevance, variation, and consequence. It does not promote on volume alone.

A bounded operating illustration

The finance-critical readiness work mentioned at the opening shows what this looks like without turning one engagement into a universal case. The recurring problem was late entry into validation. Teams arrived with different release scopes, system changes, finance calendars, and test needs. A single end-to-end project template would have buried those differences or forced the teams to rewrite most of it.

The reusable component was the entry architecture. It asked whether the workback schedule was approved and the validation criteria were explicit. It also checked that test cases existed before formal testing. Engineering and finance had to resolve the questions that would otherwise surface in executive review. Those conditions could travel. The detailed test cases, dates, owners, and evidence remained local.

Reuse changed the component. Early versions exposed where the language was too broad and where a finance calendar created a legitimate variant. A shared knowledge base made the current material easier to find. Operating value appeared when teams used the readiness conditions before the handoff.

A review of this example could not claim that one checklist made releases faster. It could show the entry problem and which checks teams used. It would name the local differences that mattered, the person who owned the readiness rules, and what would force another review. A policy change or repeated late handoff would matter more than file views.

The example also shows who should own the rule. A central PMO can keep the review record and spot patterns across teams. It cannot set a finance test rule on its own. A finance owner must judge what the shared rule says, tell teams when it changes, and retire it when it no longer helps.

Variation is design data

Local variation is often treated as noncompliance. That is a poor default. A project may vary because its sponsor has different authority, its data crosses a jurisdiction, its customer cannot accept the standard handoff, its technology has another failure mode, or its decision is harder to reverse. Those differences can be legitimate. They can also reveal that the shared core was defined too broadly.

PMI-sponsored researchers studied how six large organizations governed new ideas. They found a real tension between firm rules and room to adapt. Their report calls for rules that fit the setting and that people can understand and use in the same way.[4] A standard can do both if it says what may change and who may approve the change.

An approved variant is therefore part of the asset, not a tolerated deviation hidden in a project plan. It names the condition that justifies the difference, the part of the core that still applies, the owner who approved it, and the evidence that would end or broaden the variant. This turns local practice into information that the enterprise can learn from.

Exceptions play a similar role. One exception may be noise. The same exception across several projects is a signal. It may show that the core ignores a common dependency, that a field asks for evidence no team can obtain, or that users have found a safer route. The EPMO should not count exceptions only as breaches. It should aggregate them as design feedback.

This is also why a waiver log cannot sit apart from version governance. If exceptions accumulate but never reach the component owner, the enterprise maintains the appearance of a standard while every project works around it. The published rule and the operating reality separate. Review of exception patterns closes that gap.

Give each departure a governed status

Local differences need more than a free-text explanation. The review should distinguish an approved variant from an exception and from evidence that the core must change. An approved variant repeats under a named condition and preserves the component's purpose. An exception permits a bounded departure for one case, with an owner and an expiry or return condition. A proposed core change shows that the shared rule no longer handles the work it was meant to govern.

Those statuses create different management actions. Variants are published so the next team can select them deliberately. Exceptions are watched for recurrence; a cluster may reveal a missing variant or a failing core. Proposed changes enter the version path and require evidence from the settings they would affect. Treating all three as local customization would hide the signal that makes the component improve.

A shared asset needs a lifecycle

Promotion creates an obligation. Once teams are told to rely on a component, someone must keep it current, explain its limits, and manage the consequences of change. The asset now has users, dependencies, and switching costs. Treating promotion as the finish line hides that burden.

MIT CISR describes how AI work can move from pilots to wider use. It warns that no proven playbook covers that shift.[5] Once many teams rely on a shared part, someone must keep it useful. That takes clear roles, systems, and care over time.

The SEI model also stresses repeatable results, measures, sound operations, and steady improvement. [6] It is a maturity model, not a test of this paper's five steps. It supports the need to manage shared work, but it cannot prove that my steps will work everywhere.

This duty grows when teams share an AI control. A source rule, human review step, test routine, log rule, or path for raising a concern may touch many workflows. NIST warns that AI results can worsen after launch. Logs can split across systems, and human review can strain as use grows.[7] Owners should watch the risks their shared control is meant to catch, not follow one calendar for every case.

The same lifecycle logic applies outside AI, even when the signals differ. A financial readiness asset may change when close policy changes. An intake component may need revision when decision rights move. A handoff rule may fail after a vendor, customer, or platform change. The owner must be able to see those shifts and act before a trusted component becomes a stale control.

Keep the EPMO in its lane

The EPMO should guide the choice to share a component. It should not write every rule or approve every local change. Experts still own the content. Project leads still judge fit in their work. Risk, legal, security, finance, engineering, and operations keep their own authority. The EPMO can see patterns across projects and bring those owners together.

EXHIBIT 2

The EPMO connects the system without owning every decision

Project teams

Surface repeated components, evidence of use, local variants, and workarounds.

Promotion review

Decide whether to mature, canonize, revise, keep local, or retire the candidate.

Component owner

Own the supported core, approved variants, version path, and retirement decision.

Adopting teams

Use the current component, select valid variants, and report exceptions or failure.

The role begins with scanning projects for recurring questions, artifacts, workarounds, and failures. A promotion review then brings together the people who own the work and those who would reuse it. The EPMO keeps the core, approved variants, versions, and decisions visible, and watches for exception patterns that a single project may not see.

The EPMO should also resist two forms of empire building. It should not convert every useful local practice into an enterprise asset. Maintenance is real work, and low-frequency reuse may not justify it. It should not use the review to pull execution decisions upward. The purpose is to reduce repeated invention while preserving the project's authority over legitimate context.

Cadence follows the flow of candidates and the consequence of reliance. A large transformation portfolio may need a monthly review. A smaller environment may use quarterly or event-triggered decisions. High-risk components may need scheduled checks, while a low-risk checklist may change only when exceptions accumulate. The sources in this paper do not establish one correct timetable.

What this evidence cannot establish

The sources support lessons work that leads to action, shared rules that allow valid local differences, and ongoing care of common tools. They also show that a large lessons folder did not ensure later use.

They do not prove a fixed return for this approach. Nor do they prove that two uses are enough or that an EP MO should own every rule.

EXHIBIT 3

The evidence supports a practice boundary, not a fixed return

- 
Lessons practice
 PMI supports capture, analysis, retrieval, management action, and tracked follow-through.

- 
Context boundary
 PMI and APQC support codified practice with context-appropriate flexibility and variants.

- 
Lifecycle evidence
 MIT CISR, SEI, and NIST support managed capabilities, workflows, and monitoring.

- 
Contrary evidence
 GAO shows that an agency-wide lessons repository did not assure later application.

- 
Not established
 No source proves a promotion threshold, fixed return, universal cadence, or EP MO ownership.

Those gaps matter because a neat framework can create false confidence. Reuse evidence can be biased toward successful projects. Teams may comply with a component without using it in the decision. Owners may protect an asset they created. Exception counts may rise because reporting improved, not because the standard worsened. The review should treat those as interpretation problems, not hide them behind a score.

The framework is a governance test that makes the claims around a candidate visible and gives leaders a small set of valid dispositions. Thin evidence should keep the component local or send it back to mature. Restraint protects the standard from promotion by enthusiasm.

The standard is the managed component

Projects should not all mature the same way. Their evidence changes, dependencies move, and local conditions deserve judgment. Forcing them into one template confuses consistency with control and turns variation into paperwork.

The better enterprise move is to standardize the component that has earned reuse. Give it a stable core, a clear boundary, approved variants, an owner, a version path, an exception signal, and a retirement trigger. Then keep watching whether it still helps the work it was meant to serve.

A folder preserves what a project produced. A promotion decision establishes what the organization is prepared to rely on. The EPMO creates value in the distance between those two things.

Notes and sources

[1] Sandra F. Rowe and Sharon Sikes, Lessons Learned: Taking It to the Next Level, paper presented at PMI Global Congress 2006—North America, Project Management Institute, 2006. <https://www.pmi.org/learning/library/lessons-learned-next-level-Communicating-7991>

[2] U.S. Government Accountability Office, NASA: Better Mechanisms Needed for Sharing Lessons Learned, GAO-02-195, January 30, 2002. <https://www.gao.gov/products/gao-02-195>

[3] Mathias Kirchmer, Appropriate Process Standardization, APQC, January 23, 2024. <https://www.apqc.org/resource-library/resource-listing/appropriate-process-standardization>

[4] Michael Knapp, Catherine P. Killen, Chris Stevens, and Shankar Sankaran, Governance of Innovation in Portfolios, Programs, and Projects, PMI Sponsored Research, January 31, 2019. <https://www.pmi.org/learning/library/governance-innovation-projects-programs-portfolios-11796>

[5] Stephanie L. Woerner, Ina M. Sebastian, Peter Weill, and Evgeny Káganer, Grow Enterprise AI Maturity for Bottom-Line Impact, MIT Center for Information Systems Research, August 21, 2025. https://cisr.mit.edu/publication/2025_0801_EnterpriseAIMaturityUpdate_WoernerSebastianWeillKagane
r

[6] Ipek Ozkaya, Anita Carleton, Sebastián Echeverría, Robert Edman, John Haller, Erin Harper, Michael D. Konrad, Carol J. Smith, and Shawn Wray, AI Adoption Maturity Model v1.0, Carnegie Mellon University Software Engineering Institute, June 30, 2026. <https://www.sei.cmu.edu/projects/ai-adoption-maturity-model-improving-how-organizations-adopt-ai-solutions/>

[7] National Institute of Standards and Technology, NIST AI 800-4: Challenges to the Monitoring of Deployed AI Systems, March 9, 2026, updated March 18, 2026. <https://www.nist.gov/news-events/news/2026/03/new-report-challenges-monitoring-deployed-ai-systems>

About the author

Marco Policani is an enterprise portfolio, PMO, and AI operating-governance leader. He builds portfolio governance systems where AI carries the analytical load and named humans own every decision — an approach documented across case studies, walkthroughs, and working governance guides on his portfolio site.

Portfolio & governance library: policani.net/governance · Full portfolio: policani.net · LinkedIn: linkedin.com/in/marcpolicani

This white paper may be shared freely with attribution. © 2026 Marco Policani.